z/OS
2.5

z/OS SMP/E Reference

IBM

Note

Before using this information and the product it supports, read the information in "Notices" on page 463.

Contents

Chapter 3. Defining control statements in SMPPARM members............................ 103

Chapter 4. SMP/E data sets and files.. 113

Figures

xiii

Tables

About this document

Use this publication when you need to:

- Allocate SMP/E data sets
- Call SMP/E exit routines or service routines
- Code application programs that use SMP/E application programming interfaces
- Code SMP/E modification control statements
- Code SMPPARM members
- Define which utility programs SMP/E can call
- Transform data elements so that they can be packaged inline

SMP/E publications

The IBM SMP/E for z/OS, V3R6 publications are available as PDF files in the z/OS Internet library (www.ibm.com/servers/resourcelink/svc00100.nsf/pages/zosInternetLibrary).

Table 1 on page xxi lists the IBM SMP/E for z/OS, V3R6 publications and briefly describes each one.

For information about z/OS publications and more information about the IBM SMP/E for z/OS, V3R6 books, see *z/OS Information Roadmap*.

Table 1. Publications for IBM SMP/E for z/OS, V3R6

Title	Description
z/OS SMP/E Messages, Codes, and Diagnosis, GA32-0883	Explains SMP/E messages and return codes and the actions to take for each; and how to handle suspected SMP/E problems.
z/OS SMP/E Commands, SA23-2275	Explains SMP/E commands and processing in detail.
z/OS SMP/E Reference, SA23-2276	Explains SMP/E modification control statements, data sets, exit routines, and programming interfaces in detail and provides additional SMP/E reference material.
z/OS SMP/E User's Guide, SA23-2277	Describes how to use SMP/E to install programs and service.

How to send your comments to IBM

We invite you to submit comments about the z/OS product documentation. Your valuable feedback helps to ensure accurate and high-quality information.

Important: If your comment regards a technical question or problem, see instead

Submit your feedback by using the appropriate method for your type of comment or question:

Feedback on z/OS function
> If your comment or question is about z/OS itself, submit a request through the IBM RFE Community (www.ibm.com/developerworks/rfe/).

Feedback on IBM® Documentation function
> If your comment or question is about the IBM Documentation functionality, for example search capabilities or how to arrange the browser view, send a detailed email to IBM Documentation Support at ibmdocs@us.ibm.com.

Feedback on the z/OS product documentation and content
> If your comment is about the information that is provided in the z/OS product documentation library, send a detailed email to mhvrcfs@us.ibm.com. We welcome any feedback that you have, including comments on the clarity, accuracy, or completeness of the information.

> To help us better process your submission, include the following information:

- Your name, company/university/institution name, and email address
- The following deliverable title and order number: z/OS SMP/E Reference, SA23-2276-50
- The section title of the specific information to which your comment relates
- The text of your comment.

When you send comments to IBM, you grant IBM a nonexclusive authority to use or distribute the comments in any way appropriate without incurring any obligation to you.

IBM or any other organizations use the personal information that you supply to contact you only about the issues that you submit.

If you have a technical problem

If you have a technical problem or question, do not use the feedback methods that are provided for sending documentation comments. Instead, take one or more of the following actions:

- Go to the IBM Support Portal (support.ibm.com).
- Contact your IBM service representative.
- Call IBM technical support.

Summary of changes

This information includes terminology, maintenance, and editorial changes. Technical changes or additions to the text and illustrations for the current edition are indicated by a vertical line to the left of the change.

Note: IBM z/OS policy for the integration of service information into the z/OS product documentation library is documented on the z/OS Internet Library under IBM z/OS Product Documentation Update Policy (www-01.ibm.com/servers/resourcelink/svc00100.nsf/pages/ibm-zos-doc-update-policy?OpenDocument).

Summary of changes for z/OS SMP/E Reference Version 3 Release 7 in z/OS Version 2 Release 5 (V2R5)

New

The following content is new.

May 2023
A new statement was added to the "Attributes" description in the SMPWRK6 topic. For more information, refer to "SMPWRK6" on page 135.

March 2023
A topic was added to the GIMZIP chapter on preparing to sign GIMZIP packages. For more information, refer to "Preparing to sign GIMZIP packages" on page 426.

February 2023
For APAR IO28360, the GIMZIP service routine is updated to support digitally signing GIMZIP packages, and GIMUNZIP is updated to verify the digital signatures for signed GIMZIP packages. Refer to "GIMZIP packaging service routine" on page 419 and "GIMUNZIP file extraction service routine" on page 392 for more information.

Changed

The following content is changed.

- None.

Deleted

The following content is deleted.

- None.

Summary of changes for z/OS SMP/E Reference Version 3 Release 7 in z/OS Version 2 Release 4 (V2R4)

New

The following content is new.

March 2021 refresh

- New information added to the CALLLIBS subentry for LMOD entries, refer to "LMOD entry (distribution and target zone)" on page 205.

- For APAR IO28169, the new "contiguous" attribute is added to the list of attributes of the <FILEATTR> tag for MVS data sets, refer to "File attribute file" on page 435 for more information.
- For APAR IO26315, the new "unloadformat" attribute is added to the list of attributes of the <FILEATTR> tag for VSAM data sets, refer to "File attribute file" on page 435 for more information.

Prior to August 2020 refresh

- SMP/E V3R7 overview in *z/OS SMP/E User's Guide*

Changed

The following content is changed.

August 2020 refresh

The ++JAR and ++JARUPD MCS syntax diagrams were corrected, refer to "Adding or replacing a JAR element" on page 39 and "++JARUPD MCS" on page 45.

Prior to August 2020 refresh

- The SUP operand description for ++VER MCS was updated; for more information, refer to "++VER MCS" on page 92.

Deleted

The following content is deleted.

- None.

Changes made in SMP/E Version 3 Release 6

This document contains information that was previously presented in *SMP/E Reference*, SA23-2276-02, which supports z/OS Version 2 Release 2.

For SA23-2276-30

APAR IO26243, which allows a three character PTF value in the SMP/E level indicator.

For APAR IO25506, the samples for ORDER entry LIST outputs were updated; see "ORDER entry (global zone)" on page 250.

For APAR IO25475, new link edit attributes were added to ++MOD MCS; see "++MOD MCS" on page 60, "MOD entry (distribution and target zone)" on page 226 and "LMOD entry (distribution and target zone)" on page 205.

For SA23-2276-01

GIMZIP now includes the new `empty` attribute; see "GIMZIP packaging service routine" on page 419.

Updates were made to GIMUNZIP, including the addition of a new tag and new attributes; see "GIMUNZIP file extraction service routine" on page 392.

HTTPS is now provided as an additional method for downloading packages; see "GIMGTPKG service routine" on page 389.

For SA23-2276-16

Minor updates were made.

For SA23-2276-15

New information:

- "Determining the required size of SMPWKDIR" on page 401 was added.

Changed information:

- "++MOD MCS" on page 60 was updated with information about RMODE(31) and RMOD(31).
- "LMOD entry (distribution and target zone)" on page 205 was updated with information about RMODE(31) for the LKED ATTRIBUTES subentry.
- "MOD entry (distribution and target zone)" on page 226 was updated with information about RMODE(31) for the LKED ATTRIBUTES subentry.
- "UTILITY entry (global zone)" on page 289 was updated.
- Various updates were made to include the use of the HOLDDATA operand

Changes made in SMP/E Version 3 Release 5

For SA23-2276-14

Changed information:

- The SYSOUT subentry under the DDDEF entry has been enhanced. For details, see "DDDEF entry (distribution, target, and global zone)" on page 163.
- Information has been added to Chapter 4, "SMP/E data sets and files," on page 113.
- Information has been added to "GLOBALZONE entry (global zone)" on page 184.

For SA23-2276-13

Changed information:

- In Chapter 2, "SMP/E modification control statements," on page 5, the following sections have been updated for corrections:
  – "Hierarchical file system element MCS" on page 21
  – "++HOLD MCS" on page 30

For SA23-2276-12

New information:

- These sections were updated to include the new FIXCAT or HOLDFIXCAT subentry:
  – "++HOLD MCS" on page 30 and "++RELEASE MCS" on page 78
  – "HOLDDATA entry (global zone)" on page 197 and Table 18 on page 311
  – "OPTIONS entry (global zone)" on page 241 and Table 23 on page 315
  – Table 28 on page 320 and "SYSMOD entry (global zone)" on page 277
- A description has been added for SMPHRPT; see "SMPHRPT" on page 120
- The parameter COMPAT was added to "EXEC statement" on page 383.

Changed information:

- The UTILITY INPUT subentry of "LMOD entry (distribution and target zone)" on page 205 was updated to state the proper UCLIN syntax when the file is located in the UNIX file system.
- The description for the SOURCEID operand was updated in the following places to describe the long SOURCEID support:
  – "++ASSIGN MCS" on page 7
  – Table 28 on page 320, Table 29 on page 323, and the SOURCEID entry on page "SOURCEID SYSMOD entrysubentriesSOURCEID SOURCEIDSYSMOD subentrytarget zone SOURCEIDSYSMOD subentrydistribution zone" on page 270and "SOURCEID SYSMOD entrysubentriesSOURCEIDSOURCEIDSYSMOD subentryglobal zone" on page 281

- The SMPOUT DD and the PAGELEN entry ("OPTIONS entry (global zone)" on page 241 and Table 23 on page 315) were updated to include SMPHRPT as one of the applicable data sets.

Chapter 1. Syntax notation and rules

This chapter explains the syntax notation and rules for SMP/E modification control statements (MCSs) and OPCODE members used by SMP/E. It describes:

- How to read the notation used to show how control statements should be coded
- The rules to follow when coding control statements

How to read the syntax diagrams

Throughout this publication, the structure defined in this section is used in describing syntax:

- Read the syntax diagrams from left to right, from top to bottom, following the path of the line.

  The ►►── symbol indicates the beginning of a statement.

  The ──► symbol indicates that the statement syntax is continued on the next line.

  The ►── symbol indicates that a statement is continued from the preceding line.

  The ──►◄ symbol indicates the end of a statement.

- Required items appear on the horizontal line (main path).

  ►►─ STATEMENT ──── required_item ─►◄

- Optional items appear below the main path.

  ►►─ STATEMENT ──────────────────►◄
          └── optional_item ──┘

- If you can choose from two or more items, they appear in a vertical stack.

  If you **must** choose one of the items, one item of the stack appears on the main path.

  ►►─ STATEMENT ──┬── required_choice1 ──┬─►◄
           └── required_choice2 ──┘

  If choosing one of the items is optional, the entire stack appears below the main path.

  ►►─ STATEMENT ──────────────────────►◄
          ├── optional_choice1 ──┤
          └── optional_choice2 ──┘

  If one of the optional items is the default, it appears above the main path and the remaining choices will be shown as follows:

          ┌── default_choice1 ──┐
  ►►─ STATEMENT ──┼──────────────────┼─►◄
          ├── optional_choice2 ──┤
          └── optional_choice3 ──┘

- Keywords appear in uppercase (for example, PARM1). **They must be spelled exactly as shown.**
- Variables appear in lowercase italics (for example, *parmx*). They represent user-supplied names or values.

  ►►─ STATEMENT ──── *variable* ─►◄

- An arrow returning to the left above the main line indicates an item that can be repeated.

Syntax notation and rules

A repeat arrow above a stack indicates that you can make more than one choice from the stacked items, or repeat a single choice.

– A repeat arrow above a stack of keywords means that you can enter one or more of the keywords. However, **each keyword can be entered only once**.

– A repeat arrow above a variable means that you can enter one or more values for the variable. However, **each value can be entered only once**.

- If punctuation marks, parentheses, arithmetic operators, or other such symbols are shown, you must enter them as part of the syntax.

- Sometimes a single substitution represents a set of several parameters. For example, in the following diagram, the callout `Parameter Block` can be replaced by any of the interpretations of the subdiagram that is labeled `Parameter Block:`.

Parameter Block

Syntax rules for MCS and SMPPARM members

Follow these rules when you code SMP/E modification control statements (MCS) and SMPPARM member control statements:

- SMP/E input is case-sensitive. Use uppercase letters to enter all SMP/E keywords. Enter operands in the same case as the intended operand values. Enter the text within a comment in any case you prefer.

- Start each statement on a new logical 80-byte record.

  For MCSs, do the following:

  – Code the "++" for the MCS in columns 1 and 2.

  – Code the MCS name on the same line as the "++".

  For OPCODE member control statements, do the following:

  – Code the KEY=*xxx* operand first.

  – Do **not** continue OPCODE member control statements on a subsequent record.

  Note: Except for these restrictions, the SMP/E MCSs and OPCODE member control statements can begin and end anywhere up to and including column 72.

- You can code optional information in any order, except where noted in the syntax and operand descriptions.

- Separate operands and their values with a blank or comma.

  Note: Although the syntax diagrams show only commas when indicating the allowable separator characters for repeating values, one or more blank characters may be used instead to separate repeating values.

- You can continue a statement on more than one line. SMP/E assumes a statement is continued if it did not find a period (.) before column 73.

  Note:

1. OPCODE members are an exception—they cannot span multiple records.

2. If an operand's value must span multiple lines and that value is delimited by quotation marks, the value should extend up to and including column 72 and restart on column 1 of the next line. Put a quotation mark before the value and another after the value, but do **not** add extra quotation marks where the value spans lines. Blanks within the quoted value are considered to be part of the value, including any blanks at the beginning of a continuation line.

- Start comments with "/*" and end them with "*/". The first "*/" encountered after the initial "/*" will end the comment. A comment can appear anywhere within or after a statement, but should not start before a statement, nor begin in column 1. (When "/*" starts in column 1, it indicates the end of an input data set.) A comment after the ending period **must** start on the same line as the period. You cannot specify any additional operands or comments after that final comment. For example, you can code a comment like this:

```
SET      BDY(MVSTST1)         /* Comment after period
                                 continued on subsequent
                                 records is okay.       */.
```

However, you should **not** code a comment like this:

```
SET      BDY(MVSTST1)   .     /* Comment after period okay */
                              /* but this comment will give a
                                 syntax error */
```

This causes a syntax error at the start of the second comment after the period.

- Comments can be in single-byte characters (such as English alphanumeric characters) or in double-byte characters (such as Kanji).
- End each statement with a period.
- SMP/E completes processing for one statement before it starts processing the next one.
- SMP/E ignores columns 73 through 80. If data, such as a period, is specified beyond column 72, SMP/E ignores it and indicates an error in the statement after the one containing that data.

Syntax rules for XML statements

XML statements may be coded in the CLIENT, SERVER, SYSIN, file attribute, and package attribute files for use with the GIMZIP and GIMUNZIP service routines.

The following syntax rules apply to XML statements:

- SMP/E ignores columns 73 through 80.
- All tags have a starting and ending delimiter specified as *<keyword>* and *</keyword>*, respectively.
- Any tag that does not contain another tag (that is, nested tags) may have an ending delimiter of either **</keyword>** or just **/>**.

- Comments must must begin with "<!--" and end with "-->". All data between the "<!--" and the "-->" is ignored. Comments may not be placed inside a tag.
- Any text not contained within comment delimiters is syntax checked.
- Tags are case sensitive; attribute values may be mixed case.

Syntax notation and rules

- A tag is not required to start on a new line.
- XML tag names and attribute values may not contain the XML markup characters, '<', '>', and '&'.

Chapter 2. SMP/E modification control statements

Each SYSMOD processed by SMP/E is composed of two distinct types of data: instructions to SMP/E identifying the elements in the SYSMOD and how to install them, and the actual element replacements or updates. The instructions to SMP/E consist of a series of control statements, called modification control statements (or MCSs). This chapter describes the various MCSs that are processed by SMP/E.

Building SYSMODs (packaging)

Building SYSMODs ("packaging") includes combining the appropriate MCS statements with software elements to create one or more SYSMODs. Depending on the type of SYSMODs you are building and how you plan to distribute them, packaging can also involve putting the SYSMODs in the proper format on the distribution medium.

Although this book describes the syntax of SMP/E MCS statements, it does not contain all the information you need to use these statements for packaging SYSMODs.

- To package function SYSMODs and the associated service (PTF SYSMODs and APAR SYSMODs), you must use this book along with Standard Packaging Rules for z/OS-based Products (publibz.boulder.ibm.com/epubs/pdf/gimpkg80.pdf), which contains the rules, restrictions, and recommendations for packaging SYSMODs.
- To package USERMOD SYSMODs, use this book along with *z/OS SMP/E User's Guide* , which steps you through building a USERMOD and provides USERMOD examples that you might find helpful.

Note:

1. Each section describing an individual MCS has examples of SYSMODs containing that MCS. In the examples, the MCS being described is underlined. This is done only to make that MCS stand out; it does not imply that any special processing must be done to enter that data.

2. The examples of MCSs do not show the use of all the operands for each MCS. When you want to know how to use a particular operand for a specific MCS, first check the section describing that MCS. If the operand is not shown in an example there, check the index entry for the desired **operand** to see which other MCSs also contain that operand. Then check the examples under those MCSs. Examples of the use of an operand for one MCS can often illustrate its use for another MCS.

++APAR MCS

The ++APAR MCS identifies a service SYSMOD. This type of SYSMOD is a temporary corrective fix to the elements of target system and distribution libraries. All other MCSs for this SYSMOD follow this header MCS. For more information about packaging an APAR fix, see Standard Packaging Rules for z/OS-based Products (publibz.boulder.ibm.com/epubs/pdf/gimpkg80.pdf).

Syntax

++APAR MCS

```
►►── ++APAR( sysmod_id ) ──────────────────────────────────►
                    └─ DESCRIPTION( description ) ─┘

►──┬──────────────────────────────────┬──┬─ REWORK( level ) ─┬─ • ─►◄
   └─ FILES( number ) ─┬────────────────┘  └────────────────┘
                       └─ RFDSNPFX( relfile_prefix ) ─┘
```

Operands

DESCRIPTION

specifies a descriptive name to be associated with this SYSMOD.

- DESCRIPTION can also be specified DESC.

- The DESCRIPTION value can be in single-byte characters (such as English alphanumeric) or double-byte characters (such as Kanji).
- The DESCRIPTION value can contain up to 64 bytes of data, including blanks. Extra blanks, as well as leading and trailing blanks, are deleted. (For double-byte data, the 64-byte maximum includes all shift-in and shift-out characters, as well as the double-byte characters. Shift-in and shift-out pairs with only blanks between are compressed).
- The DESCRIPTION value can span multiple 80-byte records. Data must continue up to and including column 72 and begin in column 1 of the next line. All data past column 72 is ignored. For single-byte data, the break does not translate to a blank unless a blank is explicitly coded in column 72 of the first line or in column 1 of the second line. For double-byte data, each line must contain both the shift-out and shift-in characters.
- If DESCRIPTION is specified, it must contain at least one non-blank character.
- If parentheses are included in the text, they must be in matched pairs.

FILES
specifies the number of relative files belonging to this APAR fix. It can be a decimal number from 1 to 9999. For information about packaging SYSMODs in relative files, see Standard Packaging Rules for z/OS-based Products (publibz.boulder.ibm.com/epubs/pdf/gimpkg80.pdf).

Note:

1. Although SMP/E allows you to package APAR fixes in relative files, they are not generally packaged in this format.

2. If a packager uses a high-level qualifier on RELFILE data sets, the RFDSNPFX operand on the header MCS (not the RFPREFIX operand on the RECEIVE command) **must** be used to identify that high-level qualifier.

REWORK
specifies the level of this SYSMOD, which was reworked for minor changes. Up to eight numeric characters can be specified.

For SYSMODs supplied by IBM, the REWORK level is *yyyyddd,* where *yyyy* is the year the SYSMOD was reworked and *ddd* is the Julian date.

REWORK allows an updated SYSMOD to be automatically received again, as long as it is more recent than the version that has already been received. This takes the place of rejecting the SYSMOD and receiving it again.

Note: If a SYSMOD appears more than once in the SMPPTFIN data set, the first occurrence may be received. However, none of the subsequent versions of the SYSMOD are received, even if their rework level is higher than the one for the first version of the SYSMOD. (Message GIM40001E is issued for each of the subsequent versions of the SYSMOD.)

RFDSNPFX
identifies to SMP/E the prefix used in the relative file data set names for this SYSMOD. SMP/E uses this prefix when allocating data set names for the SYSMOD's relative files during RECEIVE processing.

- This operand can be specified only if the FILES operand is also specified.
- The RFDSNPFX value specified on the MCS statement must match the actual prefix used in the data set names for the associated relative files.

  For example, if the names of the relative files created for a SYSMOD start with "IBM", as in IBM.*sysmod_id*.F1, the header MCS statement for the SYSMOD must specify RFDSNPFX(IBM) so SMP/E knows which prefix to use when allocating the data set names for the SYSMOD's relative files during RECEIVE processing.

- Following standard data set naming conventions, the prefix can be from 1 to 8 alphanumeric or national ($, #, @) characters or a dash (–).

  To enable full Security Server protection for tape data sets and to keep the tape header within the 17-character limit (including periods), you should limit the prefix to 1 to 3 characters. If the name

exceeds the 17-character limit, only the rightmost 17 characters are written to the tape header label.

sysmod_id
> specifies a unique 7-character system modification identifier for the APAR fix. See "Naming conventions for SYSMODs" on page 458 for more information.

Examples

Here is an example of a SYSMOD containing a ++APAR statement for a temporary fix to module IFBMOD01. As the example shows, this fix is needed to answer APAR AZ12345 on an MVS™ system. The module must be at the service level provided by PTF UZ00004 for function FXY1040.

```
++APAR(AZ12345)              /* APAR type fix           */.
++VER(Z038) FMID(FXY1040)    /* for MVS product FXY1040 */
           PRE(UZ00004)      /* at this service level.  */.
++ZAP(IFBMOD01)              /* Change to one module    */
           DISTLIB(AOSFB)    /* in this DLIB.           */.
...
... IMASPZAP control statements
...
```

++ASSIGN MCS

The ++ASSIGN MCS assigns a source identifier (source ID) to one or more specified SYSMODs, as long as those SYSMODs are in the SMPPTS data set by the end of RECEIVE processing.

Syntax

++ASSIGN MCS

Operands

SOURCEID
> is a 1- to 64-character string identifying the source of the SYSMODs being processed. SMP/E associates the SOURCEID value with the SYSMODs named on the ++ASSIGN MCS. The SOURCEID value can consist of any nonblank character (X'41' through X'FE') except single quotation mark ('), asterisk (*), percent (%), comma (,), left parenthesis (() and right parenthesis ()).

TO
> specifies the SYSMODs with which the source ID is to be associated.

Usage notes

- The source ID specified on the ++ASSIGN statement is added to any source ID that was assigned to a specified SYSMOD by the RECEIVE command. It is also added to any source IDs currently associated with a specified SYSMOD that has already been received.
- ++ASSIGN statements are processed only when the SMPPTFIN data set is processed.
  - If the whole SMPPTFIN data set is processed, all ++ASSIGN statements are processed.
  - If only selected SYSMODs are processed, the ++ASSIGN statements for those SYSMODs are processed.

  ++ASSIGN statements are not processed when only the SMPHOLD data set is being processed.

- The SOURCEID and TO values must be validly specified and cannot be blank or null. For more information about source ID naming conventions, see "Naming conventions for source IDs" on page 457.
- The source ID is not assigned to any SYSMODs that are not in the global zone.

- The same SYSMOD cannot appear more than once on a single ++ASSIGN MCS.
- If the same SYSMOD appears on more than one ++ASSIGN MCS, all the source IDs are associated with the SYSMOD.
- The ++ASSIGN MCS is used in the SMPPTFIN data set and can be placed between, before, or after SYSMODs, ++FEATURE MCS, or ++PRODUCT MCS.

  It must be followed by one of the following: a ++APAR, ++FEATURE, ++FUNCTION, ++PRODUCT, ++PTF, or ++USERMOD MCS; another ++ASSIGN MCS; or an end-of-file. If one of these does not follow, SMP/E does not receive the SYSMOD being processed, and it skips the ++ASSIGN MCS.

Examples

Here are some examples of ++ASSIGN statements for SYSMODs from several preventive service levels that have been merged into the same ESO tape. A ++ASSIGN MCS has been placed between the groups of SYSMODs to identify their source:

```
++ASSIGN SOURCEID(PUT0701)   /* service level 0701.     */
TO(UZ12345,UZ12346).
++PTF(UZ12345)               /* PTF UZ12345             */.
++VER(Z038) FMID(HXP1100)    /* for MVS function HXP1100.*/.
++MOD(A) DISTLIB(DN554)      /* Update module A.        */.
A
++PTF(UZ12346)               /* PTF UZ12346             */.
++VER(Z038) FMID(HXP1200)    /* for MVS function HXP1200.*/.
++MOD(C) DISTLIB(DN554)      /* Update module C.        */.
C
        .
        .
        .
```

```
++ASSIGN SOURCEID(PUT0702)   /* service level 0702.     */
TO(UZ12347,UZ12348).
++PTF(UZ12347)               /* PTF UZ12347             */.
++VER(Z038) FMID(HXP1100)    /* for MVS function HXP1100.*/.
++IF FMID(HXP1200) THEN REQ(UZ12348).
++MOD(D) DISTLIB(DN554)      /* Update module D.        */.
D
++PTF(UZ12348)               /* PTF UZ12348             */.
++VER(Z038) FMID(HXP1200)    /* for MVS function HXP1200.*/.
++MOD(A) DISTLIB(AOS12)      /* Update module A.        */.
A
++MOD(B) DISTLIB(AOS12)      /* Update module B.        */.
B
        .
        .
        .
```

Data element MCS

Data element MCSs describe elements that are not macros, modules, or source. Data elements have the following characteristics:

- A data element must be a member of a PDS or PDSE (DSORG=PO) or reside in a sequential data set (DSORG=PS).
- The record format (RECFM) must be F, FA, FM, FS, FB, FBA, FBM, FBS, V, VA, VM, VB, VBA, VBM, VS, or VBS.
- The LRECL for a data element must be from 1 to 32,654.
- The records can be numbered or unnumbered.
- A VSAM data set can be a data element if it is in REPRO format. However, after the data is installed by SMP/E, you must run an AMS REPRO job to create the original form of the VSAM data.

Any type of data element may be installed in any distribution or target library that meets these requirements. For example, CLIST data elements can be installed into variable block target libraries.

When copying a data element during APPLY, ACCEPT, or RESTORE processing, it may be necessary for SMP/E to perform the copy itself, rather than invoking the copy utility. SMP/E performs the copy when:

- The target library or distribution library is a sequential data set.
- The data element must be reformatted to be compatible with the target or distribution library. (For more information about the reformatting of data elements, see the section on reformatting data elements in the APPLY command chapter in *z/OS SMP/E Commands*)

There are MCSs to replace data elements, just as there are MCSs to replace other types of elements. (There are no MCSs to update data elements.) Table 2 on page 9 shows the MCSs that can be used for data elements.

Table 2. MCS statements for data elements. If an element is provided in only one language, the *x*'s can be left off the MCS. If an element is provided in more than one language, replace the *x*'s with the appropriate value from Table 3 on page 11.

MCS	Description
++BOOK*xxx*	Online book member
++BSIND*xxx*	Index for an online publications library (bookshelf)
++CGM*xxx*	Graphics source for an online book
++CLIST	CLIST
++DATA	Data not covered by other types
++DATA1−++DATA5	IBM generic data types 1−5 These are for IBM use only, to define elements that are not covered by any existing data types.
++DATA6*xxx*	IBM generic data type 6 This is for IBM use only to define an element not covered by any existing data types.
++EXEC	EXEC
++FONT*xxx*	Printer Font Object Contents Architecture (FOCA) font
++GDF*xxx*	GDF graphics panel
++HELP*xxx*	Help information (for example, a member in SYS1.HELP or a dialog help panel)
++IMG*xxx*	Graphics image for an online book
++MSG*xxx*	Message member (such as for a dialog or for a message data set)
++PARM	PARMLIB member
++PNL*xxx*	Panel for a dialog
++PROBJ*xxx*	Printer object element
++PROC	Procedure in PROCLIB
++PRODXML	Product XML document

Table 2. MCS statements for data elements. If an element is provided in only one language, the*x*'s can be left off the MCS. If an element is provided in more than one language, replace the *x*'s with the appropriate value from Table 3 on page 11. *(continued)*

MCS	Description
++PRSRC*xxx*	Printer source element
++PSEG*xxx*	Graphics page segment for an online book
++PUBLB*xxx*	Online publications library (bookshelf)
++SAMP*xxx*	Sample data, program, or JCL in a data set for sample code
++SKL*xxx*	File skeleton for a dialog
++TBL*xxx*	Table for a dialog
++TEXT*xxx*	Text
++USER1−++USER5	User-defined data types 1–5 These are for user-defined elements that are not covered by any existing data types.
++UTIN*xxx*	General utility input
++UTOUT*xxx*	General utility output

Supporting several languages

Some types of elements, such as panels, messages, or text, may have to be translated into several languages. In these cases, the corresponding MCSs contain *xxx* to indicate which language is supported by a given element. Figure 1 on page 10 shows an example where product XX1 must provide both English and French support for a message module, a panel, a panel message, and a sample element.

```
++FUNCTION(FXX1101).              ++FUNCTION(FXX1102).

++VER(Z038) FMID(EXX1100).        ++VER(Z038) FMID(EXX1100).

++MOD(ZZZMOD0E)…      message     ++MOD(ZZZMOD0F)…
   DISTLIB(AZZZMOD1).    modules     DISTLIB(AZZZMOD1).

++PNLENU(ZZZPNL01)…   panels      ++PNLFRA(ZZZPNL01)…
   DISTLIB(AZZZPNLE)                DISTLIB(AZZZPNLF)
   SYSLIB(SZZZPNLE).                SYSLIB(SZZZPNLF).

++MSGENU(ZZZMSG01)…   dialog      ++MSGFRA(ZZZMSG01)…
   DISTLIB(AZZZMSGE)    messages     DISTLIB(AZZZMSGF)
   SYSLIB(SZZZMSGE).                SYSLIB(SZZZMSGF).

++SAMPENU(ZZZSMP01)…  samples     ++SAMPFRA(ZZZSMP01)…
   DISTLIB(AZZZSAME)                DISTLIB(AZZZSAMF)
   SYSLIB(SZZZSAME).                SYSLIB(SZZZSAMF).
```

Figure 1. Example of using data element MCSs

Note:

1. The message modules can be in the same distribution library, because the element names are different.

2. For the panels, dialog messages, and samples, there is a different element *type* for each language version of an element. Therefore, the element *name* can remain the same for all the languages in

which the element is supported. However, elements with the same name must be installed in different libraries. (SMP/E does not check whether different types of data elements have the same name. Likewise, SMP/E does not prevent elements with the same name from being installed in the same libraries.)

Table 3 on page 11 shows the *xxx* values that can be used when the MCS indicates the language being supported.

Table 3. National language identifiers used for language-unique elements			
Value	**Language**	**Value**	**Language**
ARA	Arabic	HEB	Hebrew
CHS	Simplified Chinese	ISL	Icelandic
CHT	Traditional Chinese	ITA	Italian (Italy)
DAN	Danish	ITS	Italian (Switzerland)
DES	German (Switzerland)	JPN	Japanese
DEU	German (Germany)	KOR	Korean
ELL	Greek	NLB	Dutch (Belgium)
ENG	English (United Kingdom)	NLD	Dutch (Netherlands)
ENP	Uppercase English	NOR	Norwegian
ENU	English (United States)	PTB	Portuguese (Brazil)
ESP	Spanish	PTG	Portuguese (Portugal)
FIN	Finnish	RMS	Rhaeto-Romanic
FRA	French (France)	RUS	Russian
FRB	French (Belgium)	SVE	Swedish
FRC	French (Canada)	THA	Thai
FRS	French (Switzerland)	TRK	Turkish

Syntax

The syntax to be used depends on the processing to be done for the element:

- Adding or replacing the element
- Deleting the element

Adding or replacing a data element
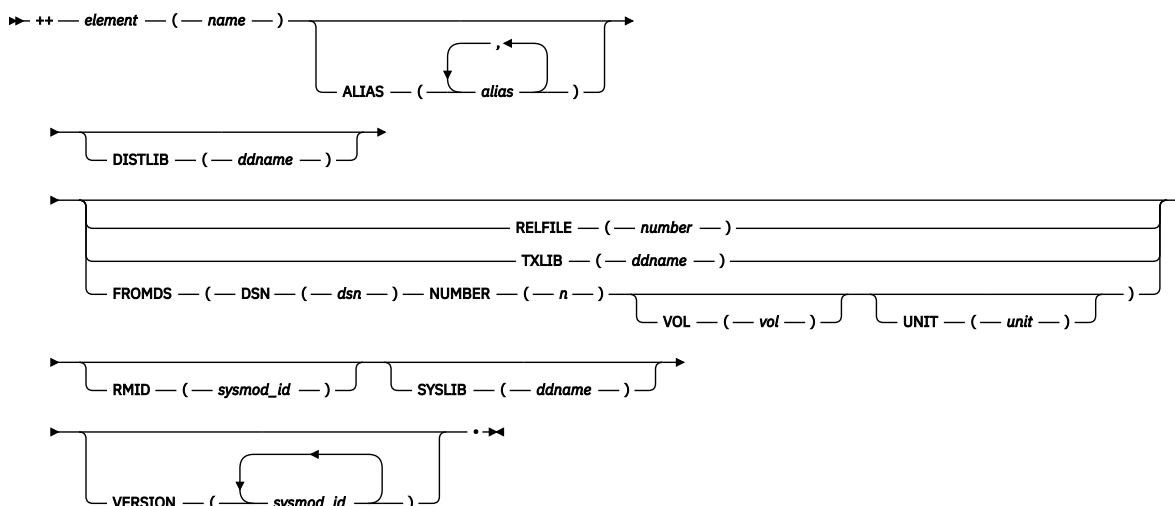
Data element MCS

Deleting a data element

Data element MCS

Operands

ALIAS
specifies the alias names for the data element in both the target and distribution libraries.

You can use ALIAS when data elements of the same type must be defined in the same zone and must have the same name for programming access. In this case, you can specify the common name on ALIAS and a unique name as the data element name.

DELETE
indicates that the element and all of its alias names are to be removed from the target libraries, the distribution libraries, and the SMP/E data sets.

Note:

1. DELETE is mutually exclusive with all other operands except DISTLIB and VERSION.

2. If the element statement is in a base function, you may want to use the DELETE operand on the ++VER MCS to delete the previous release, rather than on the element statement to delete a specific element.

3. Specification of the DELETE operand results in all alias names of the data element being deleted along with the data element identified.

DISTLIB
specifies the ddname of the distribution library for the data element.

Note:

1. DISTLIB must be specified if the element has not been previously recorded on the target zone or distribution zone.

2. If a data element entry already exists in the target zone or distribution zone and the value currently in that entry does not match that specified in the DISTLIB operand, the SYSMOD is not applied or accepted.

element
specifies the type of element. Table 2 on page 9 shows the MCSs used for the various element types.

FROMDS
identifies the partitioned data set that contains this element.

Note: The FROMDS operand and its DSN, NUMBER, VOL, and UNIT suboperands are included in the MCS generated by the BUILDMCS command. IBM does not intend the FROMDS operand to be used in manually coded MCS.

DSN
specifies the dsname of the FROMDS data set. The specified data set name must conform to standard data set naming conventions and cannot contain parentheses. The maximum length of the entire name is 44 characters (including the periods).

NUMBER
specifies a number that SMP/E is to use when assigning a name to the SMPTLIB data set associated with this FROMDS data set. (This is similar to the way the relative file number is used in RELFILE processing.)

VOL
specifies, for an uncataloged data set, the volume serial number of the volume containing the FROMDS data set. If specified, this volume identifier must be from 1 to 6 alphanumeric characters.

VOL may be omitted for a cataloged data set.

UNIT
specifies, for an uncataloged data set, the UNIT type containing the FROMDS data set. If specified, the UNIT value must be from 1 to 8 characters and must conform to standard UNIT naming conventions. SMP/E accepts any nonblank characters specified between the open and close parentheses, up to a maximum length of 8.

UNIT may be omitted for a cataloged data set.

Note: FROMDS is mutually exclusive with DELETE, RELFILE, and TXLIB.

name
specifies the name of the element to be replaced. The name can contain any alphanumeric characters and $, #, @, or hex C0.

RELFILE
specifies which relative file associated with the SYSMOD contains this element. This operand is required if you provide the element in RELFILE format, rather than inline or in a TXLIB data set.

Note: RELFILE is mutually exclusive with DELETE, FROMDS, and TXLIB.

RMID
specifies the last PTF that **replaced** this data element. This operand can be used only in a service-updated function, and the specified PTF must be integrated into the function.

Note: RMID is mutually exclusive with DELETE.

SYSLIB
specifies the ddname of the target library for the specified element.

Note: SYSLIB is mutually exclusive with DELETE.

TXLIB
specifies the ddname of the partitioned data set containing the element. This operand is required if the element is provided in a TXLIB data set rather than inline or in RELFILE format.

Note:

1. SMPTLIB cannot be used as a value on the TXLIB operand.
2. TXLIB is mutually exclusive with DELETE, FROMDS, and RELFILE.

VERSION
> specifies one or more function SYSMODs that currently contain the element. The function containing the data element statement takes over ownership of the element from the specified functions.
>
> When VERSION is specified on an element statement, it overrides any VERSION operand values specified on the ++VER MCS.

Usage notes

- If the element is packaged inline, it must immediately follow the data element MCS and must not contain any records that start with the characters ++. Neither FROMDS, nor RELFILE, nor TXLIB can be specified on the data element MCS.
- To be packaged inline, a data element must contain fixed-block-80 records. If the original format of the element is not fixed-block-80 records, you can use GIMDTS to transform the element into the required format before packaging it. Later, when SMP/E installs the element, it is changed back to its original format. For more information about using GIMDTS, see "GIMDTS: Data transformation service routine" on page 387.
- If the data element is packaged in a TXLIB data set, the ddname specified on the TXLIB operand is required during APPLY and ACCEPT processing.
- For information about packaging SYSMODs in RELFILE, TXLIB, or inline format, see Standard Packaging Rules for z/OS-based Products (publibz.boulder.ibm.com/epubs/pdf/gimpkg80.pdf).

Examples

The following examples are provided to help you use the data element MCS:

Example 1: Packaging a CLIST in a function

Here is an example of a SYSMOD containing a ++CLIST statement to install your CLIST and have SMP/E track any changes to it. It can be packaged as a function, as shown in the following example:

```
++FUNCTION(MYCLST1)        /* Function.            */.
++VER(Z038)                /* For MVS systems.     */.
++CLIST(CLIST1)            /* Install this CLIST   */
        TXLIB(NEWSMP)      /* from this TXLIB      */
        DISTLIB(AMACLIB)   /* into this DLIB and   */
        SYSLIB(ISPCLIB)    /* this target library. */.
```

Example 2: Packaging a renamed CLIST

Suppose that, for some reason, you need to rename CLIST1, which was introduced in "Example 1: Packaging a CLIST in a function" on page 14. The new name is to be CLISTX. You do not need to change anything else about the CLIST. Here is an example of a SYSMOD containing the data element statements needed to package this renamed CLIST:

```
++USERMOD(USR0001)         /* User modification    */.
++VER(Z038) FMID(MYCLST1)  /* to user application. */.
++CLIST(CLIST1)            /* Delete the original  */
        DELETE             /* CLIST.               */.
++CLIST(CLISTX)            /* Add the renamed CLIST */
        TXLIB(NEWSMP)      /* from this TXLIB      */
        DISTLIB(AMACLIB)   /* into this DLIB and   */
        SYSLIB(ISPCLIB)    /* this target library. */.
```

You must ensure that the renamed CLIST is in the TXLIB library (NEWSMP) used to provide SMP/E with the element.

++DELETE MCS

During APPLY processing, the ++DELETE MCS deletes a load module and any known alias or symbolic link names from a target library and the associated target zone. It can also delete alias or symbolic link names without deleting the load module itself.

During ACCEPT processing, if inline JCLIN was processed, the distribution zone's entries are updated to reflect the ++DELETE changes. The target libraries and target zone entries are not updated.

Note: This change is not reversible. You cannot restore the SYSMOD containing the ++DELETE MCS; therefore, you cannot recreate the deleted load module.

Syntax

++DELETE MCS

```
►►— ++DELETE — ( — name — ) →

        ►— SYSLIB — ( — ALL — ) ─────────────────────────────── • →◄
        │                        ┌────── , ◄──────┐              │
        │              └─ ALIAS — ( ─┴─ alias ─┴─ ) ─┘           │
        │              ┌──────── , ◄────────┐                    │
        └───────── SYSLIB — ( ─┴─ ddname ─┴─ ) ─────────────────┘
```

Operands

ALIAS
specifies that only the indicated alias or symbolic link names are to be deleted, but not the load module.

- An alias name can have up to 1023 characters.

  Although a load module residing in a PDSE can have an alias name greater than 8 characters, the ++DELETE statement cannot be used to delete such an alias value. To delete such an alias value without deleting the load module, you need to resupply JCLIN to define the load module without providing an ALIAS statement for the alias value to be deleted. Make sure to also include a ++MOD statement for a module in the load module to force SMP/E to relink the load module.

- Single apostrophes must be used as delimiters for an alias name in these cases:

  – The alias name contains a character that is not uppercase alphabetic, numeric, national ($, #, or @), slash (/), plus (+), hyphen, period, or ampersand (&).

  – The alias name is continued on another line in the control statement.

  Apostrophes used as delimiters do not count as part of the 1023-character limit.

- If an apostrophe is part of the alias name and is not a delimiter, it needs to be doubled. These two apostrophes count as 2 characters in the 1023-character limit.

- Characters specified must be in the range X'41' through X'FE'. Unlike the binder, SMP/E does not support shift-in (X'0F') and shift-out (X'0E') characters.

- When processing the ++DELETE statement, SMP/E uses the alias name as is, and does not enforce any rules the binder might be using as a result of the CASE execution parameter.

Note:

1. **Do not** specify this operand if you also want to delete the load module. ++DELETE without the ALIAS operand automatically deletes any alias or symbolic link names associated with the load module.

2. When ALIAS is specified, SMP/E checks the ALIAS control statements in the LMOD entry to verify that the specified name is actually an alias or symbolic link of the load module. For copied load

modules, instead of looking for ALIAS control statements, SMP/E checks the corresponding MOD entry's TALIAS subentries.

3. If a valid ALIAS value is specified, SMP/E deletes the alias from all known target libraries, including any associated side deck library, no matter what was coded for SYSLIB. SMP/E overrides the SYSLIB value with SYSLIB(ALL).

4. When you specify ALIAS to delete an alias for a load module, you must reflect this change using JCLIN. To do this, include a ++JCLIN statement with JCLIN data containing a link-edit step for the load module, with the alias deleted from the list of aliases on the link-edit ALIAS statement. This causes SMP/E to replace the alias list in the LMOD entry.

name
specifies the name of the load module to be deleted.

SYSLIB
specifies the ddname of the target library where the load module resides.

- If ALL is specified, the load module is deleted from all target libraries defined in the target zone.
- If a single ddname is specified, the load module is deleted only from the indicated target library.

  If the load module or alias is to be deleted from two target libraries, a second ddname can be specified. However, to make sure the module is deleted from both libraries, you should use ALL instead.

- For load modules with a SYSLIB allocation (or load modules having a CALLLIBS subentry), specifying ALL or deleting the load module from all the target libraries in which it resides, deletes the base version of the load module from the SMPLTS library.

- For load modules with a side deck library, specifying ALL or deleting the load module from all the target libraries in which it resides, deletes the associated definition side deck from the side deck library.

Note: If ALIAS is specified, SMP/E deletes the alias from all known target libraries, including any associated side deck library, no matter what was coded for SYSLIB. SMP/E overrides the SYSLIB value with SYSLIB(ALL).

Usage notes

- SYSLIB must always be specified to identify the affected target libraries.
- ALIAS should be specified **only** if alias names, and not load modules, are to be deleted.
- ++DELETE statements must follow any ++VER and ++IF statements and must precede any ++JCLIN or element MCSs.
- Regardless of the order in which ++MOVE, ++RENAME, and ++DELETE statements are coded in a SYSMOD, they are always processed in this order for APPLY and ACCEPT:

  1. ++MOVE
  2. ++RENAME
  3. ++DELETE

  Afterwards, the ++JCLIN statements are processed, and then the element statements are processed.

  Note: You cannot restore the SYSMOD containing the ++DELETE MCS.

Examples

The following examples are provided to help you use the ++DELETE MCS:

Example 1: Deleting a single load module

Here is an example of a SYSMOD containing a ++DELETE statement that deletes load module LMODA from SYS1.LINKLIB:

```
++PTF(UR01234)             /* Identify the PTF number. */.
++VER(Z038) FMID(HXY1300)  /* For MVS function HXY1300.*/.
++IF (ESY1300) THEN        /* If ESY1300 is installed  */
  REQ(UR12399)             /* UR12399 is required.     */.
++HOLD(UR01234)            /* Hold UR01234.            */
        FMID(HXY1300)      /* For MVS function HXY1300.*/
        SYSTEM             /* System hold              */
        REASON(DELETE)     /* because of ++DELETE.     */
        COMMENT(THIS DELETION OF LMODA FROM LINKLIB
              IS IRREVERSIBLE).
++DELETE (LMODA)           /* Delete load module LMODA */
  SYSLIB(LINKLIB)          /* from LINKLIB.            */.
++JCLIN                    /* JCLIN follows.           */.
          .
          .
          .
++MOD(MODAA) DISTLIB(AOSXX) /* Element MCS statements.  */.
```

Example 2: Deleting an alias from a load module

Assume that IBM has shipped you a PTF that deletes an alias for load module LMODA. Here is an example of a SYSMOD containing a ++DELETE statement that deletes an alias from LMODA:

```
++PTF(UR01235)             /* Identify the PTF number. */.
++VER(Z038) FMID(HXY1300)  /* For MVS function HXY1300.*/.
++HOLD(UR01235)            /* Hold UR01235.            */
        FMID(HXY1300)      /* For MVS function HXY1300.*/
        SYSTEM             /* System hold              */
        REASON(DELETE)     /* because of ++DELETE.     */
        COMMENT(THE DELETION OF THE ALIAS FOR LMODA
              IS IRREVERSIBLE).
++DELETE (LMODA)           /* Identify LMOD LMODA.     */
  SYSLIB(ALL)              /* Process all SYSLIBs.     */
  ALIAS(OTHNAME)           /* Identify alias.          */.
++JCLIN                    /* JCLIN to delete alias.   */.
          .
          .
          .
```

Example 3: Deleting an alias from a load module in a UNIX file system

Assume that IBM has shipped you a PTF that deletes an alias for load module BPXLMODB, which resides in a UNIX file system. Here is an example of a SYSMOD containing a ++DELETE statement that deletes an alias from BPXLMODB:

```
++PTF(UZ00440)             /* Identify the PTF number. */.
++VER(Z038) FMID(HOP1101)  /* For MVS function HOP1101.*/.
++HOLD(UZ00440)            /* Hold UZ00440.            */
        FMID(HOP1101)      /* For MVS function HOP1101.*/
        SYSTEM             /* System hold              */
        REASON(DELETE)     /* because of ++DELETE.     */
        COMMENT(THE DELETION OF THE ALIAS FROM
              BPXLMODB IS IRREVERSIBLE).
++DELETE (BPXLMODB)        /* Identify LMOD BPXLMODB.  */
  SYSLIB(ALL)              /* Process all SYSLIBs.     */
  ALIAS('../nicename')     /* Identify alias.          */.
++JCLIN                    /* JCLIN follows.           */.
          .
          .
          .
++MOD(BPXMODAA) DISTLIB(AOSXX) /* Element MCS statements.  */.
```

++FEATURE MCS

The ++FEATURE MCS is used to describe a set of function SYSMODs that are collectively referred to as a FEATURE. It introduces descriptive information about a new or replacement set of functions into the global zone. A ++FEATURE MCS may be associated with an orderable software feature.

Syntax

++FEATURE MCS

➤➤── ++FEATURE ── (── *name* ──) ── DESCRIPTION ── (── *description* ──) ──➤

```
         ┌───────────────────────┐      PRODUCT ── ( ── prodid ── , ── vv.rr.mm ── ) ──➤
         │          ┌──── , ◄───┐ │
         │          │           │ │
         └─ FMID ── ( ─┴─ fmid ─┴─ ) ─┘
```

```
  ┌──────────────────────────────┐    · ──➤◄
  └─ REWORK ── ( ── level ── ) ──┘
```

Operands

name
> is a 1- to 8-character feature name. It can contain uppercase alphabetic, numeric, or national characters ($, #, @).

DESCRIPTION
> describes the feature that represents this collection of function SYSMODs.

> - DESCRIPTION can also be specified as DESC.
> - The DESCRIPTION value can be in single-byte characters (such as English alphanumeric) or double-byte characters (such as Kanji).
> - The DESCRIPTION value can contain up to 64 bytes of data, including blanks. (For double-byte data, the 64-byte maximum includes all shift-in and shift-out characters, as well as the double-byte characters.) Extra blanks as well as leading and trailing blanks are deleted.
> - The DESCRIPTION value can span multiple 80-byte records. Data must continue up to and including column 72 and begin in column 1 of the next line. All data past column 72 is ignored. The break does not translate to a blank unless a blank is explicitly coded in column 72 of the first line or in column 1 of the second line.
> - The DESCRIPTION value cannot be only blanks.
> - If parentheses are included in the text, they must be in matched pairs.

FMID
> specifies the list of all function SYSMODs that make up this feature. Each FMID is 7 characters long and must be a valid SYSMOD ID. That is, it must contain uppercase alphabetic, numeric, or national characters ($, @, #). If multiple FMIDs are specified, they must be separated by commas.

PRODUCT
> identifies the *prodid* and *vv.rr.mm* of the product with which this feature is associated.

REWORK
> is the level of this feature, which was reworked for minor changes. Up to eight numeric characters can be specified.

> For IBM features, the REWORK level is *yyyyddd*, which is the year followed by the Julian date (for example, 2008110).

> REWORK allows an updated feature to be automatically received again, as long as it is more recent than the version that has already been received. This takes the place of rejecting the feature and receiving it again.

> **Note:** If a ++FEATURE statement appears more than once in the SMPPTFIN data set, the first occurrence may be received. However, none of the subsequent versions of the ++FEATURE statement are received, even if their rework level is higher than the one for the first version of the ++FEATURE statement.

Usage notes

- The ++FEATURE statements are processed only when the SMPPTFIN data set is processed.

- If the whole SMPPTFIN data set is processed, all ++FEATURE statements are processed.
- If only selected SYSMODs are processed, SMP/E processes the ++FEATURE statements that name at least one of the selected SYSMODs in their FMID operand list.

++FEATURE statements are not processed when only the SMPHOLD data set is being processed.

- The *name*, DESCRIPTION, and PRODUCT values are required and cannot be blank or null.
- The ++FEATURE MCS is used in the SMPPTFIN data set and can be placed between, before, or after SYSMODs, ++FEATURE MCS, or ++PRODUCT MCS.

If the function SYSMODs identified in the FMID operand of the ++FEATURE statement have not been previously received, the ++FEATURE statement must follow these function SYSMODs in the SMPPTFIN data set.

The ++FEATURE statement must be followed by one of the following: a ++APAR, ++ASSIGN MCS, ++FUNCTION, ++PRODUCT, ++PTF, ++USERMOD, another ++FEATURE MCS, or an end-of-file. If one of these does not follow, SMP/E skips the ++FEATURE MCS.

Example

Here is an example of a ++FEATURE MCS for the OS/390® product.

```
++FEATURE( OS325BAS )          /* Feature definition */
  DESCRIPTION( OS/390 Base Feature ) /* Description */
  FMID(HBB6605,HMP1B00,JBB66C5,…) /* FMID List    */
  PRODUCT (5647-A01,2.5.0)    /* Owning Product    */ .
```

++FUNCTION MCS

The ++FUNCTION MCS identifies a SYSMOD as a base function or dependent function. This type of SYSMOD introduces a new or replacement function into target system and distribution libraries. All other MCSs follow this header MCS statement. For more information about packaging a function, see Standard Packaging Rules for z/OS-based Products (publibz.boulder.ibm.com/epubs/pdf/gimpkg80.pdf).

Syntax

++FUNCTION MCS

▶▶── ++FUNCTION ──(── *sysmod_id* ──) ──────────────────────────────────▶
                                    └─ DESCRIPTION ──(── *description* ──) ─┘

▶──▶
   └─ FESN ──(── *fe_service_number* ──) ─┘

▶──▶
   └─ FILES ──(── *number* ──) ──┬──────────────────────────────────┬─┘
                                   └─ RFDSNPFX ──(── *relfile_prefix* ──) ─┘

▶──────────────────────────────── •─▶◀
   └─ REWORK ──(── *level* ──) ─┘

Operands

DESCRIPTION

specifies a descriptive name to be associated with this SYSMOD. For more information, see Standard Packaging Rules for z/OS-based Products (publibz.boulder.ibm.com/epubs/pdf/gimpkg80.pdf).

- DESCRIPTION can also be specified DESC.
- The DESCRIPTION value can be in single-byte characters (such as English alphanumeric) or double-byte characters (such as Kanji).

- The DESCRIPTION value can contain up to 64 bytes of data, including blanks. (For double-byte data, the 64-byte maximum includes all shift-in and shift-out characters, as well as the double-byte characters.) Extra blanks, as well as leading and trailing blanks are deleted.
- The DESCRIPTION value can span multiple 80-byte records. Data must continue up to and including column 72 and begin in column 1 of the next line. All data past column 72 is ignored. The break does not translate to a blank unless a blank is explicitly coded in column 72 of the first line or in column 1 of the second line.
- If DESCRIPTION is specified, it must contain at least one non-blank character.
- If parentheses are included in the text, they must be in matched pairs.

FESN

is a 7-character field engineering (FE) service number.

FILES

specifies the number of relative files belonging to this function. It can be a decimal number from 1 to 9999, and is used only if the function is packaged in relative files, rather than inline or in indirect libraries. For information about packaging SYSMODs in relative files, see Standard Packaging Rules for z/OS-based Products (publibz.boulder.ibm.com/epubs/pdf/gimpkg80.pdf).

Note:

1. Functions are generally packaged in relative files to improve SMP/E's performance when applying and accepting the SYSMOD.

2. If a packager uses a high-level qualifier on RELFILE data sets, the RFDSNPFX operand on the header MCS (not the RFPREFIX operand on the RECEIVE command) **must** be used to identify that high-level qualifier.

REWORK

is the level of this SYSMOD, which was reworked for minor changes. Up to eight numeric characters can be specified.

REWORK is generally used only for SYSMODs supplied by IBM that have been reworked for minor changes, such as for a service update or to use a ++MOVE, ++RENAME, or ++DELETE MCS. For these SYSMODs, the REWORK level is *yyyyddd*, which is the year followed by the Julian date (for example, 2008110).

REWORK allows an updated SYSMOD to be automatically received again, as long as it is more recent than the version that has already been received. This takes the place of rejecting the SYSMOD and receiving it again.

Note: If a SYSMOD appears more than once in the SMPPTFIN data set, the first occurrence may be received. However, none of the subsequent versions of the SYSMOD are received, even if their rework level is higher than the one for the first version of the SYSMOD. (Message GIM40001E is issued for each of the subsequent versions of the SYSMOD.)

RFDSNPFX

identifies to SMP/E the prefix used in the relative file data set names for this SYSMOD. SMP/E uses this prefix when allocating data set names for the SYSMOD's relative files during RECEIVE processing.

- This operand can be specified only if the FILES operand is also specified.
- The RFDSNPFX value specified on the MCS statement must match the actual prefix used in the data set names for the associated relative files.

  For example, if the names of the relative files created for a SYSMOD start with "IBM", as in IBM.*sysmod_id*.F1, the header MCS statement for the SYSMOD must specify RFDSNPFX(IBM) so SMP/E knows which prefix to use when allocating the data set names for the SYSMOD's relative files during RECEIVE processing.

- Following standard data set naming conventions, the prefix can be from 1 to 8 alphanumeric or national ($, #, @) characters or a dash (–).

  To enable full Security Server protection for tape data sets and to keep the tape header within the 17-character limit (including periods), you should limit the prefix to 1 to 3 characters. If the name

exceeds the 17-character limit, only the rightmost 17 characters are written to the tape header label.

sysmod_id
is a unique 7-character name, or SYSMOD ID, for the function. This ID is also called a function modification identifier (FMID). For more information, see "Naming conventions for SYSMODs" on page 458.

Usage notes

A function cannot contain statements that update elements (++MACUPD, ++SRCUPD, and ++ZAP).

Examples

Here is an example of a function SYSMOD to be created with a SYSMOD ID of JXY1040 that is dependent on function JXY1000. The elements and JCL input data are members of three unloaded partitioned data sets on a tape created using the relative file method:

```
++FUNCTION(JXY1040)       /* New function SYSMOD    */
          FILES(3)        /* in RELFILE format.     */
          RFDSNPFX(IBM)   /* RELFILE prefix for IBM.  */.
++VER(Z038) FMID(JXY1000) /* For MVS function JXY1000.*/.
++JCLIN     RELFILE(1)    /* JCLIN in RELFILE 1.      */.
++MOD(IFBMOD01)           /* This module            */
          DISTLIB(AOSFB)  /* for this DLIB          */
          RELFILE(2)      /* is in RELFILE 2.       */.
++MOD(IFBMOD02)           /* This module            */
          DISTLIB(AOSFB)  /* for this DLIB          */
          RELFILE(2)      /* is in RELFILE 2.       */.
++MAC(IFBMAC01)           /* This macro             */
          DISTLIB(AOSMAC) /* for this DLIB          */
          RELFILE(3)      /* is in RELFILE 3.       */.
```

Hierarchical file system element MCS

The hierarchical file system element MCSs describe elements located in a UNIX file system. Hierarchical file system elements can have any of the following characteristics:

- The record format (RECFM) must be F, FA, FM, FB, FBA, FBM, V, VA, VM, VB, VBA, or VBM.
- Elements with variable-length records cannot contain spanned records.
- The maximum LRECL is 32,654.
- The records can be numbered or unnumbered.

There are MCSs to replace hierarchical file system elements, just as there are MCSs to replace other types of elements. (There are no MCSs to update hierarchical file system elements.) Table 4 on page 21 shows the MCSs that can be used for hierarchical file system elements. Table 5 on page 22 shows the *xxx* values that can be used when the MCS indicates the language being supported.

Table 4. MCS statements for hierarchical file system elements. If an element is provided in only one language, the *xxx* can be left off the MCS. If an element is provided in more than one language, replace the *xxx* with the appropriate value from Table 5 on page 22.

MCS	Description
++AIX1 through ++AIX5	Elements to be used by an AIX® client.
++CLIENT1 through ++CLIENT5	Elements to be used by any client (intended for clients not described by other elements types.
++HFS*xxx*	Generic hierarchical file system element (data not covered by other types)

Table 4. MCS statements for hierarchical file system elements. If an element is provided in only one language, the *xxx* can be left off the MCS. If an element is provided in more than one language, replace the *xxx* with the appropriate value from . *(continued)*

MCS	Description
++OS21 through ++OS25	Elements to be used by an OS/2 client.
++SHELLSCR	UNIX shell script elements.
++UNIX1 through ++UNIX5	Elements to be used by a UNIX client.
++WIN1 through ++WIN5	Elements to be used by a Windows client.

Table 5. National language identifiers used for language-unique elements

Value	Language	Value	Language
ARA	Arabic	HEB	Hebrew
CHS	Simplified Chinese	ISL	Icelandic
CHT	Traditional Chinese	ITA	Italian (Italy)
DAN	Danish	ITS	Italian (Switzerland)
DES	German (Switzerland)	JPN	Japanese
DEU	German (Germany)	KOR	Korean
ELL	Greek	NLB	Dutch (Belgium)
ENG	English (United Kingdom)	NLD	Dutch (Netherlands)
ENP	Uppercase English	NOR	Norwegian
ENU	English (United States)	PTB	Portuguese (Brazil)
ESP	Spanish	PTG	Portuguese (Portugal)
FIN	Finnish	RMS	Rhaeto-Romanic
FRA	French (France)	RUS	Russian
FRB	French (Belgium)	SVE	Swedish
FRC	French (Canada)	THA	Thai
FRS	French (Switzerland)	TRK	Turkish

Syntax

The syntax to be used depends on the processing to be performed for the element:

- Adding or replacing the element
- Deleting the element

Adding or replacing a hierarchical file system element

Hierarchical file system element MCS

```
►►— ++ — element — ( — name — ) ─────────────────────────────────────►
                              ├─ BINARY ─┤  └─ DISTLIB — ( — ddname — ) ─┘
                              └─ TEXT ───┘

►─────────────────────────────────────────────────────────────────────►
    │                                ,◄                │
    └─ LINK — ( ─┬─ linkname ─┬─ )    └─ PARM — ( — character_string — ) ─┘
                 └────────────┘

►─────────────────────────────────────────────────────────────────────►
    │                   ┌─ RELFILE — ( — number — ) ────────────────┐
    │                   ├─ TXLIB — ( — ddname — ) ──────────────────┤
    └─ FROMDS — ( — DSN — ( — dsn — ) — NUMBER — ( — n — )
                        └─ VOL — ( — vol — ) ─┘ └─ UNIT — ( — unit — ) ─┘ )

►─────────────────────────────────────────────────────────────────────►
    └─ RMID — ( — sysmod_id — ) ─┘

►─────────────────────────────────────────────────────────────────────►
    └─ SHSCRIPT — ( — scriptname ────────────────── )
                      └─ PRE ─┘ └─ POST ─┘

►─────────────────────────────────────────────────────────────────────►
    │                    ,◄                       ,◄
    └─ SYMLINK — ( ─┬─ symlinkname ─┬─ ) — SYMPATH — ( ─┬─ sympathname ─┬─ )
                    └───────────────┘                   └───────────────┘

►─────────────────────────────────────────────────────────────────────►
    └─ SYSLIB — ( — ddname — ) ─┘
                                │                ,◄
                                └─ VERSION — ( ─┬─ sysmod_id ─┬─ ) ─┘
                                                └─────────────┘

►— · ►◄
```

Deleting a hierarchical file system element

Hierarchical file system element MCS

```
►►— ++ — element — ( — name — ) — DELETE ────────────────────────────►
                                         └─ DISTLIB — ( — ddname — ) ─┘

►─────────────────────────────────── · ►◄
    │                 ,◄
    └─ VERSION — ( ─┬─ sysmod_id ─┬─ )
                    └─────────────┘
```

Operands

BINARY
indicates that the hierarchical file system copy utility should install the element into a UNIX file system in *binary* mode. This means that the element is installed in its entirety as a data stream, with no breaks for logical records.

Note:

1. BINARY is mutually exclusive with TEXT.
2. When BINARY is specified on the element MCS, SMP/E sets the BINARY mode indicator in the hierarchical file system element entry. When TEXT is specified on the element MCS, SMP/E sets the TEXT mode indicator in the hierarchical file system element entry.

   If neither BINARY nor TEXT is specified on the element MCS, SMP/E uses the mode indicator in the hierarchical file system element entry to tell the HFS copy utility how to install the element.

   If neither BINARY nor TEXT is specified on the element MCS and there is no mode indicator in the hierarchical file system element entry, the HFS copy utility determines how to install the element.

3. SMP/E recommends the appropriate value BINARY or TEXT be specified to ensure that the HFS copy utility uses the correct mode. If no value is specified, the HFS copy utility chooses either binary or text mode based on the RECFM of the element to be copied, and it might choose incorrectly.

DELETE
specifies that the hierarchical file system element and all of its link names and symbolic link names are to be removed from the "target library" (UNIX file system) and the distribution library.

Note:

1. DELETE is mutually exclusive with all other operands except DISTLIB and VERSION.

2. If the element statement is in a base function, you may want to use the DELETE operand on the ++VER MCS to delete the previous release, rather than on the element statement to delete a specific element.

3. Specification of the DELETE operand results in all link names and symbolic link names of the element being deleted along with the element identified.

DISTLIB
specifies the ddname of the distribution library for the specified hierarchical file system element. During ACCEPT processing, SMP/E installs the hierarchical file system element into the distribution library as a member. (The distribution library must be a PDS or PDSE; it cannot be part of a UNIX file system.)

Note:

1. DISTLIB must be specified when the hierarchical file system element is first installed.

2. If an element entry already exists in the target zone or distribution zone and the value currently in that entry does not match that specified in the DISTLIB operand, the SYSMOD is not applied or accepted.

element
specifies the type of element. shows the MCSs used for the various hierarchical file system element types.

FROMDS
identifies the partitioned data set that contains this element.

Note: The FROMDS operand and its DSN, NUMBER, VOL, and UNIT suboperands are included in the MCS generated by the BUILDMCS command. IBM does not intend the FROMDS operand to be used in manually coded MCS.

DSN
specifies the dsname of the FROMDS data set. The specified data set name must conform to standard data set naming conventions and cannot contain parentheses. The maximum length of the entire name is 44 characters (including the periods).

NUMBER
specifies a number that SMP/E is to use when assigning a name to the SMPTLIB data set associated with this FROMDS data set. (This is similar to the way the relative file number is used in RELFILE processing.)

VOL
specifies, for an uncataloged data set, the volume serial number of the volume containing the FROMDS data set. If specified, this volume identifier must be from 1 to 6 alphanumeric characters.

VOL may be omitted for a cataloged data set.

UNIT
specifies, for an uncataloged data set, the UNIT type containing the FROMDS data set. If specified, the UNIT value must be from 1 to 8 characters and must conform to standard UNIT naming conventions. SMP/E accepts any nonblank characters specified between the open and close parentheses, up to a maximum length of 8.

UNIT may be omitted for a cataloged data set.

Note: FROMDS is mutually exclusive with DELETE, RELFILE, and TXLIB.

LINK
specifies the alternative names by which this hierarchical file system element can be known in a UNIX file system. The full name is produced by concatenating the specified linkname with the UNIX file system directory identified by the SYSLIB subentry. Each linkname is passed to the HFS copy utility as an execution parameter.

Note:

1. The linkname can be from 1 to 1023 characters.

2. A linkname can be enclosed in single apostrophes ('). A linkname **must** be enclosed in single apostrophes if any of the following is true:

   • The linkname contains lowercase alphabetic characters.

   • The linkname contains a character that is not uppercase alphabetic, numeric, national ($, #, @), slash (/), plus (+), hyphen, period, or ampersand (&).

   • The linkname spans more than one line in the control statement.

   The single apostrophes used to enclose a linkname (the delimiters) do not count as part of the 1023-character limit.

3. Any apostrophes specified as part of a linkname (not the delimiters) must be doubled.

   Double apostrophes count as two characters in the 1023-character limit.

4. The linkname can include characters X'40' through X'FE'.

5. LINK values are saved and passed to the HFS copy utility as follows:

   • If LINK is specified on the element MCS, any values previously saved in the element entry are overlaid.

   • If LINK is not specified on the element MCS and saved values exist in the hierarchical file system element entry, the saved values are passed to the HFS copy utility as execution parameters.

   If LINK is not specified on the element MCS and there are no saved values in the hierarchical file system element entry, no linknames are passed to the HFS copy utility.

name
specifies the name of the hierarchical file system element member. The name can contain any uppercase alphabetic, numeric, or national ($, #, @) character and can be 1 to 8 characters long.

PARM
specifies a character string that is to be passed to the hierarchical file system copy utility as an execution-time parameter. (The values that can be specified on the PARM operand, such as PATHMODE, are those accepted by the BPXCOPY utility. See *z/OS UNIX System Services Command Reference* for a description of BPXCOPY and the values it accepts.) The maximum length of this character string is 300 bytes of nonblank data. If any blanks are specified in the PARM value, they are deleted by SMP/E during processing and do not count toward the 300-byte maximum.

Note:

1. PARM is an optional operand.

2. The character string can be entered free-form, without regard to blanks (which are compressed out of the string), and can span multiple 80-byte records.

3. If parentheses are specified in the PARM value, there must always be a pair (left and right); otherwise, the results are unpredictable.

4. PARM values are saved and passed to the HFS copy utility as follows:

   • If PARM is specified on the element MCS, any values previously saved in the hierarchical file system element entry are overlaid.

- If PARM is not specified on the element MCS and saved values exist in the hierarchical file system element entry, the saved values are passed to the HFS copy utility as execution parameters.

  If PARM is not specified on the element MCS and there are no saved values in the hierarchical file system element entry, no parameters are saved from the element MCS or passed from the hierarchical file system element entry to the HFS copy utility.

5. If the UTILITY entry for the HFS copy utility specifies a PARM value, those parameters are passed to the utility in addition to any parameters saved in the hierarchical file system element entry.

RELFILE
> identifies which relative file associated with the SYSMOD contains this element. This operand is required if you provide the element in RELFILE format, rather than inline or in a TXLIB data set.

> **Note:**

> 1. The RELFILE value must be a decimal number from 1 to 9999.
> 2. RELFILE is mutually exclusive with FROMDS and TXLIB.

RMID
> specifies the last SYSMOD that **replaced** this element. This operand can be used only in a service-updated function, and the specified PTF must be integrated into the function.

SHSCRIPT
> specifies a UNIX shell script, *scriptname*, to be invoked when the element is installed in (or deleted from) a directory of a UNIX file system. *scriptname* can contain any uppercase alphabetic, numeric, or national ($, #, @) character and can be 1 to 8 characters long.

> A shell script is commonly used to complete the installation of an element. For example, if the hierarchical file system element is a TAR or PAX file, you can provide a shell script that performs the necessary steps to restore the file.

> *scriptname* must be the first value to follow the SHSCRIPT operand.

> You define the shell script to SMP/E through a ++SHELLSCR statement, which can be within the same SYSMOD as the hierarchical file system element, or within a SYSMOD that was processed previously. When the element is itself a SHELLSCR type, the SHSCRIPT operand must match the name of the element.

> You cannot define more than one shell script for an element.

> You can follow *scriptname* with either of two optional values, PRE and POST, to specify the point in SMP/E processing when the shell script is to be invoked. The following examples show how you can use the PRE and POST values:

> - To run the shell script before the element is copied to a UNIX file system directory, specify:

> ```
> SHSCRIPT(scriptname,PRE)
> ```

> - To run the shell script after the element is copied to a UNIX file system directory, specify:

> ```
> SHSCRIPT(scriptname,POST)
> ```

> or specify no value after *scriptname* to use POST by default.

> - To run the shell script both before and after the element is copied to a UNIX file system directory, specify:

> ```
> SHSCRIPT(scriptname,PRE,POST)
> ```

> If you do not specify a PRE or POST value for the shell script, SMP/E invokes the shell script after the element is installed in the directory.

> When the element is a shell script, PRE is not valid. Also, you cannot specify the SHSCRIPT operand with the DELETE operand.

SYMLINK
specifies a list of one or more symbolic links, which are file names that can be used as alternate names for referring to this element in a UNIX file system. Each linkname listed here is associated with a pathname listed in the SYMPATH operand. For more information about how the linknames and pathnames are associated, see the description of the SYMPATH operand and "Example 3: Packaging a SYSMOD with a symbolic link" on page 30.

The SYMLINK value specified should be a relative path value (that is, it does not start with a slash ["/"]). When the symbolic link is created, it is created relative to the pathname of the element's SYSLIB ddname.

SYMLINK must be specified if the SYMPATH operand is specified, otherwise it must be omitted.

A symbolic linkname can be from one to 1023 characters. Any characters in the range X'40' through X'FE' may be specified.

The value may be enclosed in single apostrophes. It **must** be enclosed in single apostrophes if:

- it is continued to the next line in the MCS, or
- it contains a character that is not uppercase alphabetic, numeric, national ($, #, @), slash (/), plus (+), hyphen, period, or ampersand (&).

If an apostrophe is a part of the symbolic linkname and is not a delimiter, then it must be doubled. These two apostrophes count as two characters against the 1023 character limit for a symbolic linkname. The single apostrophes used to enclose a symbolic linkname do not count against the 1023 character limit.

SYMPATH
specifies a list of one or more pathnames that are associated with symbolic links identified by the SYMLINK operand. The first pathname in the SYMPATH operand is associated with the first symbolic link in the SYMLINK operand, the second pathname with the second symbolic link, and so on. If there are more symbolic links listed than there are pathnames, then the last listed pathname is used for the remaining symbolic links. If more pathnames are specified than symbolic linknames, then the excess pathnames (at the end of the list) are ignored.

The SYMPATH value specified should be a relative path value. When the symbolic link is accessed, the system assumes the destination of that link (the SYMPATH value) is relative to that symbolic link (the SYMLINK value). For more information about how the pathnames and linknames are associated, see "Example 3: Packaging a SYSMOD with a symbolic link" on page 30.

SYMPATH must be specified if the SYMLINK operand is specified, otherwise it must be omitted.

A symbolic pathname can be one to 1023 characters. Any characters in the range X'40' through X'FE' may be specified.

The value may be enclosed in single apostrophes. It **must** be enclosed in single apostrophes, if:

- it is continued to the next line in the MCS, or
- it contains a character that is not uppercase alphabetic, numeric, national ($, #, @), slash (/), plus (+), hyphen, period, or ampersand (&).

If an apostrophe is a part of the symbolic pathname and is not a delimiter, then it must be doubled. These two apostrophes count as two characters against the 1023 character limit for a symbolic pathname. The single apostrophes used to enclose a symbolic linkname do not count against the 1023 character limit.

SYSLIB
specifies the ddname of the "target library" within the UNIX file system for the element.

During APPLY processing, the HFS copy utility installs the hierarchical file system element into a UNIX file system. During RESTORE processing, the HFS copy utility copies the hierarchical file system element from the distribution library member into a UNIX file system.

Note: SYSLIB must be specified when the hierarchical file system element is first installed.

TEXT

indicates that the hierarchical file system copy utility should install the element into a UNIX file system in *text* mode. This means that the element is installed with breakpoints for logical records.

Note:

1. TEXT is mutually exclusive with BINARY.

2. When TEXT is specified on the element MCS, SMP/E sets the TEXT mode indicator in the hierarchical file system element entry. When BINARY is specified on the element MCS, SMP/E sets the BINARY mode indicator in the hierarchical file system element entry.

   If neither BINARY nor TEXT is specified on the element MCS, SMP/E uses the mode indicator in the hierarchical file system element entry to tell the HFS copy utility how to install the element.

   If neither BINARY nor TEXT is specified on the element MCS and there is no mode indicator in the hierarchical file system element entry, the HFS copy utility determines how to install the element.

3. SMP/E recommends the appropriate value BINARY or TEXT be specified to ensure that the HFS copy utility uses the correct mode. If no value is specified, the HFS copy utility chooses either binary or text mode based on the RECFM of the element to be copied, and it might choose incorrectly.

TXLIB

is the ddname of the partitioned data set containing the hierarchical file system element. This operand is required if the hierarchical file system element is provided in a data set that the users have access to, rather than inline or in RELFILE format.

Note:

1. SMPTLIB cannot be used as a value on the TXLIB operand.

2. TXLIB is mutually exclusive with FROMDS and RELFILE.

VERSION

specifies one or more function SYSMODs that currently contain the element. The function containing the element MCS takes over ownership of the element from the specified functions.

When VERSION is specified on an element statement, it overrides any VERSION operand values that might be specified on the ++VER MCS.

Usage notes

- If the hierarchical file system element is packaged inline, it must immediately follow the hierarchical file system element MCS and must not contain any records starting with ++. Neither FROMDS, nor RELFILE, nor TXLIB can be specified on the hierarchical file system element MCS.

- To be packaged inline, a hierarchical file system element must contain fixed-block–80 records. If the original format of the element is not fixed-block–80 records, you can use GIMDTS to transform the element into the required format before packaging it. Later, when SMP/E installs the element, it is changed back to its original format. For more information about using GIMDTS, see "GIMDTS: Data transformation service routine" on page 387.

- If the hierarchical file system element is packaged in a TXLIB data set, the ddname specified in the TXLIB operand is required during APPLY and ACCEPT processing.

- For information about elements packaged in RELFILE format, see Standard Packaging Rules for z/OS-based Products (publibz.boulder.ibm.com/epubs/pdf/gimpkg80.pdf).

- A ++HFS MCS can be used to supply a pre-built program object to be placed into a UNIX file system. A user can do an OGET for an existing program object in the file system to cause the program object to be placed into an MVS data set as fixed length records. This data can then be packaged as a ++HFS element with the BINARY operand. The HFS copy utility can then copy the element into the file system as a binary entity, which can then be executed.

Examples

The following examples are provided to help you use the hierarchical file system element MCS:

-
-
-

Example 1: Packaging a hierarchical file system element in a function

Here is an example of a SYSMOD containing a hierarchical file system element statement to install your element and have SMP/E track any changes to it. The element:

- Uses the Japanese language
- Is to be installed in BINARY mode
- Can also be known by its linkname "USERHFS"
- Is to be processed by shell script "UNTAR" after the element is copied to a UNIX file system directory.

The element can be packaged in a function, as shown in the following example:

```
++FUNCTION(MYHFSEL)            /* Function.           */.
++VER(Z038)                    /* For MVS SRELs.      */.
++HFSJPN(HFSELEM1)             /* Install this element */
        TXLIB(NEWHFS)    /* from this TXLIB        */
        DISTLIB(ABPXLIB)  /* into this DLIB and     */
        SYSLIB(BPXLIB1)   /* this target library.   */
        BINARY            /* Use BINARY mode.       */
        LINK('../USERHFS') /* Define linkname.       */
        SHSCRIPT(UNTAR,POST) /* Untar the file with the UNTAR exec. */.
```

To install the function, you need to specify DDDEF entries or DD statements for the TXLIB, DISTLIB, and SYSLIB data sets. Remember that the SYSLIB data set is actually a pathname in a UNIX file system. For an example of defining a pathname in a DDDEF entry, see .

Example 2: Packaging a renamed hierarchical file system Element

Suppose that, for some reason, you need to rename HFSELEM1, which was introduced in . The new name is to be HFSELEMX. You do not need to change anything else about the hierarchical file system element. Here is an example of a SYSMOD containing the hierarchical file system element statements needed to package this renamed hierarchical file system element:

```
++USERMOD(USR0001)             /* User modification      */.
++VER(Z038) FMID(MYHFSEL)      /* to user application.   */.
++HFSJPN(HFSELEM1)             /* Delete the original    */
        DELETE            /* HFS element.           */.
++HFSJPN(HFSELEMX)             /* Add the renamed element*/
        TXLIB(NEWHFS)    /* from this TXLIB        */
        DISTLIB(ABPXLIB)  /* into this DLIB and     */
        SYSLIB(BPXLIB1)   /* this target library.   */
        TEXT              /* Use TEXT mode.         */
        LINK('../USERHFS') /* Define linkname.       */.
```

To install the USERMOD, you need to specify DDDEF entries or DD statements for the TXLIB, DISTLIB, and SYSLIB data sets. Remember that the SYSLIB data set is actually a pathname in a UNIX file system. For an example of defining a pathname in a DDDEF entry, see .

You must also ensure that the renamed hierarchical file system element is in the TXLIB library used to provide SMP/E with the element.

Example 3: Packaging a SYSMOD with a symbolic link

Here is an example of a SYSMOD containing a hierarchical file system element statement to install your element with a symbolic link. The element:

- Is to be installed in TEXT mode
- Can be known by its linkname "/usr/lpp/gskssl/include/gskssl.h"
- Can also be known by its symbolic link "/usr/include/gskssl.h".

The element can be packaged in a function, as shown in the following example:

```
++FUNCTION(SAMPLE)          /* Function.          */.
++VER(Z038)                 /* For MVS SRELs.     */.
++HFS(GSKAH010)             /* Install this element */
        RELFILE(2)          /* from a RELFILE        */
        DISTLIB(AGSKHFS)    /* into this DLIB and    */
        SYSLIB(SGSKHFS)     /* this target library.  */
        LINK('../include/gskssl.h')
                            /* Define linkname.      */
        SYMLINK('../../../../usr/include/gskssl.h')
                            /* Define symbolic link  */
        SYMPATH('../../usr/lpp/gskssl/include/gskssl.h')
                            /* Define path name.     */
        PARM(PATHMODE(0,6,4,4))
                            /* HFS copy utility parameters */
        TEXT                /* Use TEXT mode.        */.
```

To install the function, you need to specify DDDEF entries or DD statements for the DISTLIB and SYSLIB data sets. Remember that the SYSLIB data set is actually a pathname in a UNIX file system. In this example, the DDDEF for SYSLIB must be defined as follows:

```
UCLIN.
ADD DDDEF(SGSKHFS) PATH('/service/usr/lpp/gskssl/IBM/').
ENDUCLIN.
```

When SMP/E installs file GSKAH010 into the directory specified on the given DDDEF entry, the LINK and SYMLINK values will be resolved relative to the DDDEF directory. The file will then be known by the following absolute path names in a UNIX file system:

Table 6. Absolute path names

Source	Resolved absolute pathname
SMP/E element name	/service/usr/lpp/gskssl/IBM/GSKAH010
LINK name	/service/usr/lpp/gskssl/include/gskssl.h
SYMLINK name	/service/usr/include/gskssl.h

The contents of a symbolic link (SYMLINK) is its associated SYMPATH. In this case, the contents of the symbolic link is "../../usr/lpp/gskssl/include/gskssl.h". When a symbolic link is resolved, its content is resolved relative to the symbolic link itself. That is, the SYMPATH value is relative to the SYMLINK value. Therefore, in this case, the SYMPATH value resolves to the LINK name, "/service/usr/lpp/gskssl/include/gskssl.h".

++HOLD MCS

The ++HOLD MCS identifies a SYSMOD to be placed into exception SYSMOD status (signifying that special SMP/E processing is required before it can be applied or accepted). ++HOLD statements can occur within a SYSMOD (internal HOLDDATA), or they can be read directly from the SMPHOLD file during RECEIVE processing (external HOLDDATA). For additional information about processing the ++HOLD statements, see *z/OS SMP/E Commands*.

Syntax

++HOLD MCS

```
►►── ++HOLD( sysmod_id ) ── FMID( fmid ) ──┬── REASON ── ( reason_id ) ──────────┬──►
                                           └── REASON ── ( SYSTEM Reason IDs ) ──┘
```

```
►──┬──────────────────────── ERROR ─────────────────────────────┬──►
   │                     ┌──────── , ◄────────┐                  │
   ├── FIXCAT ── CATEGORY( ──┴── category ──┴── ) ──┬──────────────────────┬──┤
   │                                           └── RESOLVER( sysmod_id ) ──┘  │
   ├──────────────────────── SYSTEM ────────────────────────────┤
   └──────────────────────── USER ──────────────────────────────┘
```

```
►──┬──────────────────────────┬──┬────────────────┬──►
   └── CLASS( ── class ── ) ──┘  └── DATE( yyddd ) ──┘
```

```
►──┬──────────────────────────────────────────┬── . ──►◄
   ├── COMMENT( text ) ───────────────────────┤
   └── COMMENT( Enhanced HOLDDATA text ) ──────┘
```

SYSTEM reason IDs used by IBM

```
►►──┬── ACTION ──┬──►◄
    ├── AO ──────┤
    ├── DB2BIND ─┤
    ├── DDDEF ───┤
    ├── DELETE ──┤
    ├── DEP ─────┤
    ├── DOC ─────┤
    ├── DOWNLD ──┤
    ├── EC ──────┤
    ├── ENH ─────┤
    ├── EXIT ────┤
    ├── EXRF ────┤
    ├── FULLGEN ─┤
    ├── IOGEN ───┤
    ├── IPL ─────┤
    ├── MSGSKEL ─┤
    ├── MULTSYS ─┤
    └── RESTART ─┘
```
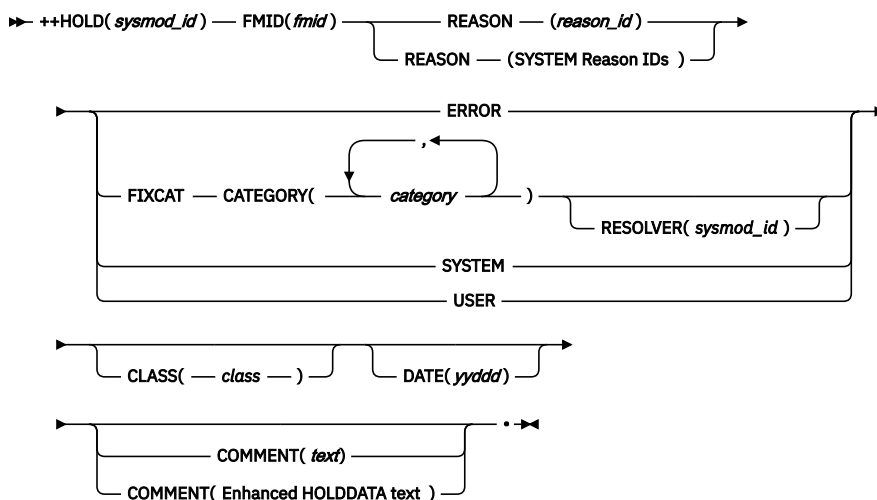
Enhanced HOLDDATA text used by IBM

```
►►── SMRTDATA ── ( ── CHGDT ── ( ── yymmdd ── ) ──┬─────────────────────────┬──►
                                                 └── FIX ── ( ── sysmod_id ── ) ──┘
```

```
►──┬───────────────────────────────────┬── ) ──►◄
   └── SYMP ── ( ── symptoms ── ) ──────┘
```

Operands

CATEGORY

a 1- to 64-character string that specifies a Fix Category. There can be one or more Fix Category values. A Fix Category value associates the reason ID APAR to a particular category of fixes. A Fix Category might be used to identify a group of APAR fixes required to support a particular hardware device, or to provide a particular software capability, similar to how a Preventive Service Planning Bucket (PSP-Bucket) identifies a group of APARs. Examples of Fix Categories used by IBM are

```
IBM.Device.2094.z/OS
```

and

```
IBM.HealthChecker.
```

The Fix Category values are used during RECEIVE command processing to create source IDs for the SYSMODs that will resolve the specified reason ID APAR. During APPLY and ACCEPT command processing the Fix Category values are used to determine if the HOLDDATA is applicable to the current command by comparing the values in the HOLDDATA to those in the active Fix Category interest list. The active Fix Category interest list is specified on the APPLY and ACCEPT command, or within the active OPTIONS entry.

A category value might contain any nonblank character in the range X'41' - X'FE' except single quotation mark ('), asterisk (*), percent (%), comma (,), left parenthesis ((), and right parenthesis ()).

CLASS
a 1- to 7-character string indicating an alternative reason to release an exception SYSMOD for processing. A class name is specified along with a reason ID to identify a condition when the reason ID need not be resolved. The same class name can be specified on any number of ++HOLD statements in any number of SYSMODs.

These are the specific values currently used by IBM:

Class
    Explanation

ERREL
The SYSMOD is held for an error reason ID but should be installed anyway. IBM has determined that the problem the SYSMOD resolves is significantly more critical than the error reflected by the holding APAR.

HIPER
The SYSMOD is held with a hold class of HIPER (High Impact)

PE
The SYSMOD is held with a hold class of "PTF in Error".

SECINT
The reason ID SYSMOD identifies a fix for a security or integrity error. HOLDDATA for security or integrity fixes is available through the z Systems Security Portal. Information on registration and accessing the z Systems Security Portal is available at Enterprise security (www.ibm.com/systems/z/solutions/enterprise-security.html). If you are already registered you can link directly to the IBM Resource Link® Security Alerts.

UCLREL
UCLIN needed for the SYSMOD has been handled by IBM and no longer requires your attention.

YR2000
Identifies PTFs that provide Year 2000 function, or fix a Year 2000-related problem.

For additional information, see .

COMMENT(text)
free-form text to be used to describe the problem identified by the REASON operand. The comments supplied are associated only with the reason ID supplied. The comments can be in single-byte characters (such as English alphanumeric characters) or double-byte characters (such as Kanji).

COMMENT(Enhanced HOLDDATA)
contains IBM–supplied Enhanced HOLDDATA, as follows:

SMRTDATA
indicates that the COMMENT operand contains Enhanced HOLDDATA.

    CHGDT
    specifies the date that the ++HOLD statement was last updated with Enhanced HOLDDATA. The date is specified as *yymmdd,* where *yy* is the last two digits of the year, *mm* is the month and *dd* is the date of the month.

FIX

identifies the SYSMOD that resolves the held SYSMOD.

SYMP

contains a description of the hold. The description might include symbols representing symptoms of the problem, a description of the problem in text form, or a combination of both. The symbols that represent symptoms identify how the problem affects the held SYSMOD.

Here is a list of symbols that might appear and their meanings. For a complete list of the symbols, see Enhanced HOLDDATA for z/OS (service.software.ibm.com/holdata/390holddata.html).

DAL

Data loss

DST2007

Daylight Saving Time support

EURO99

Euro currency symbol support

FUL

Major function loss

IPL

System outage (Requires IPL)

PRF

Performance

PRV

Pervasive

SYSPLXDS

SYSPLEX data sharing

XSYS

Cross-system migration, compatibility, or toleration

YR2000

Year 2000

B*n.n*, T*n.n*

The Common Vulnerability Scoring System (CVSS) Base and Temporal scores for the identified security or integrity fix. Further information about CVSS and a guide to scoring can be found on the FIRST (Forum of Incident Response and Security Teams (www.first.org) web site. Additional information is available at IBM System z Security Portal FAQ (www.ibm.com/common/ssi/cgi-bin/ssialias?htmlfid=ZSQ03054USEN).

Only the first 57 characters of the SYMP field will be displayed in the Exception SYSMOD Report generated by the REPORT ERRSYSMODS command. Parentheses can be used within the field, but they must be matched pairs.

DATE

specifies the date that the ++HOLD statement was generated. The date is specified as *yyddd*, where *yy* is the last two digits of the year and *ddd* is the Julian date.

ERROR, FIXCAT, SYSTEM, or USER

specifies the hold category into which the SYSMOD is to be put. At least one of the values must be specified.

ERROR

An APAR reported an error in the SYSMOD. The SYSMOD should not be applied or accepted until the APAR is resolved. A PTF held for this reason is also called a *program error PTF*, or PE-PTF. SMP/E automatically resolves the APAR and allows the SYSMOD to be applied or accepted when a SYSMOD that either matches or supersedes the APAR is also applied or accepted. Error holds can be read only from the SMPHOLD data set.

Note: ERROR can also be specified as ERR.

FIXCAT

An APAR provides a fix for the held SYSMOD and the fix is associated with one or more Fix Categories. It is optional whether the APAR will affect processing for the held SYSMOD, based on the APAR's Fix Categories and the Fix Categories of interest specified by the user. If any one or more Fix Categories for the APAR match any of those of interest to the user, then the held SYSMOD will not be applied or accepted until the APAR is resolved. The APAR is resolved when a SYSMOD that matches the APAR name, or a SYSMOD that supersedes the APAR, is applied or accepted. FIXCAT holds can be read only from the SMPHOLD data set.

See the FIXCAT operand for APPLY and ACCEPT command processing for details of specifying the Fix Categories of interest.

SYSTEM

Special action outside normal SMP/E processing is required for the SYSMOD. Examples are SYSMODs requiring a SYSGEN after they are installed, or SYSMODs requiring the installation of an associated engineering change (EC) level. System holds can appear in the SYSMOD itself or in the SMPHOLD data set.

Note: SYSTEM can also be specified as SYS.

USER

The SYSMOD requires special processing because of a decision you have made. User holds can be read only from the SMPHOLD data set.

FMID

specifies the FMID to which the held SYSMOD is applicable. For external HOLDDATA (a ++HOLD statement not within a SYSMOD), this information allows SMP/E to receive only those statements associated with FMIDs defined in the user's global zone. This operand is required.

REASON

a 1- to 7-character string used to help users identify the reason why the SYSMOD is being put into exception SYSMOD status. The reason IDs that can be specified depend on the type of hold.

Note: A SYSMOD can contain only one ++HOLD MCS for each required reason ID.

- An *error reason ID* is the APAR number that caused the SYSMOD to be placed in exception status.

- A *system reason ID* identifies some special processing the SYSMOD requires. Although SMP/E accepts any 1- to 7-character alphanumeric string as a system reason ID, these are the values currently used by IBM on ++HOLD statements:

ID

> **Explanation**

ACTION

The SYSMOD needs special handling before or during APPLY processing, ACCEPT processing, or both.

AO

The SYSMOD may require action to change automated operations procedures and associated data sets and user exits in products or in customer applications. The PTF cover letter describes any changes (such as to operator message text, operator command syntax, or expected actions for operator messages and commands) that can affect automation routines.

DB2BIND

A DB2® application REBIND is required for the SYSMOD to become effective.

DDDEF

Data set changes or additions as required.

DELETE

The SYSMOD contains a ++DELETE MCS, which deletes a load module from the system.

DEP

The SYSMOD has a software dependency.

DOC

The SYSMOD has a documentation change that should be read before the SYSMOD is installed.

DOWNLD
Code that is shipped with maintenance that needs to be downloaded.

DYNACT
The changes supplied by the SYSMOD may be activated dynamically without requiring an IPL. The HOLD statement describes the instructions required for dynamic activation. If those instructions are not followed, then an IPL is required for the SYSMOD to take effect.

EC
The SYSMOD needs a related engineering change.

ENH
The SYSMOD contains an enhancement, new option or function. The HOLD statement provides information to the user regarding the implementation and use of the enhancement.

EXIT
The SYSMOD contains changes that may affect a user exit. For example, the interface for an exit may be changed, an exit may need to be reassembled, or a sample exit may be changed.

EXRF
The SYSMOD must be installed in both the active and the alternative Extended Recovery Facility (XRF) systems at the same time to maintain system compatibility. (If you are not running XRF, you should bypass this reason ID.)

FULLGEN
The SYSMOD needs a complete system or subsystem generation to take effect.

IOGEN
The SYSMOD needs a system or subsystem I/O generation to take effect.

IPL
The SYSMOD requires an IPL to become effective. For example, the SYSMOD may contain changes to LPA or NUCLEUS, the changes may require a CLPA, or a failure to perform an IPL might lead to catastrophic results, such as could be caused by activation of a partial fix.

Note: If you plan to perform an IPL with CLPA after the SYSMOD has been applied, then no further investigation of the HOLD is required; simply bypass the IPL reason ID. However, if you are not planning to perform an IPL with CLPA, then the details of the HOLD statement must be investigated to determine what kind of actions are required to activate the SYSMOD.

MSGSKEL
This SYSMOD contains message changes that must be compiled for translated versions of the message changes to become operational on extended TSO consoles.

If you want to use translated versions of the messages, you must run the message compiler once for the library containing the English message outlines, and once for each additional language you want to be available on your system. For details, see *z/OS MVS Planning: Operations*.

If you want to use **only** the English version of the messages, you do not need to run the message compiler. You should bypass this reason ID.

MULTSYS
Identifies fixes that need to be applied to multiple systems, in one of three cases: preconditioning, coexistence, or exploitation.

RESTART
To become effective, the SYSMOD requires a special subsystem restart operation. The HOLD statement contains information regarding the required restart actions.

- A *user reason ID* is defined by the user.
- A *Fix Category reason ID* is the SYSMOD ID for the APAR that caused the SYSMOD to be placed into exception status. The APAR reason ID is associated to one or more categories of fixes.

For additional information, see "Naming conventions for HOLD reason IDs and HOLD classes" on page 455.

RESOLVER
Identifies the SYSMOD that resolves the held SYSMOD. More specifically, the resolving SYSMOD supersedes the reason ID APAR that caused the SYSMOD to be held.

sysmod_id
If the ++HOLD is being processed from the SMPHOLD data set, *sysmod_id* specifies the name of the SYSMOD to be placed into exception SYSMOD status.

If the ++HOLD is internal to a SYSMOD, then *sysmod_id* specifies either the SYSMOD that contains the ++HOLD or the originating SYSMOD that contained the ++HOLD and is now superseded by the current SYSMOD.

This operand is required.

Usage notes

- If the text within the COMMENT operand contains parentheses, they must be in matched pairs. For example:

```
COMMENT ( ADD 'PARM(COND(5))' ON THE EXEC )
```

This works, but the following statements result in an error:

```
COMMENT ( ADD OPENING PARENTHESES '(' )
```

or

```
COMMENT ( ADD CLOSING PARENTHESES ')' )
```

- The following restrictions apply only to ++HOLD MCSs packaged within a SYSMOD:

  - The only HOLD type allowed is SYSTEM.

  - ++HOLD statements must be placed after all ++VER and ++IF statements and before the first ++JCLIN or element statement.

  - The SYSMOD ID on the ++HOLD MCS must match either the SYSMOD containing the ++HOLD MCS, or the originating SYSMOD that contained the ++HOLD MCS and is now superseded by the current SYSMOD.

Examples

The following examples are provided to help you use the ++HOLD MCS:

Example 1: Noting a special documentation change

Here is an example of a SYSMOD containing a ++HOLD statement. PTF UZ12345, which is applicable to function FXY1040, introduces a documentation change that the system operator should know about.

```
++PTF(UZ12345)            /* PTF modification.      */.
++VER(Z038) FMID(FXY1040)  /* For MVS product FXY1040. */
                          /* No prerequisites.      */.++HOLD      (UZ12345)        /* Hold this
PTF        */
        FMID(FXY1040)    /* for this function      */
        SYSTEM           /* for system processing and*/
        REASON(DOC)      /* for doc change.        */
        COMMENT(this PTF changes operator console
                message xxx1233. A reply of U is
                required to this message ).
++MOD(IFBMOD01)          /* Change this module    */
        DISTLIB(AOS12)   /* in this DLIB.         */.
```

Example 2: Marking a PTF that is in error

When IBM discovers errors in PTFs it has already shipped, it provides information about these program error PTFs (PE-PTFs) in RETAIN. To make sure these PTFs are held, you can use the RETAIN information to build ++HOLD statements for the PTFs.

Note: In this example, because only ++HOLD statements are being shown, none of them are underlined.

```
++HOLD     (UZ12345)        /* Hold this PTF         */
           FMID(FXY1040)    /* for this function     */
           ERROR            /* for APAR fix.         */
           REASON(AZ00001)  /* APAR is AZ00001.      */
           COMMENT(this APAR causes loop)  /*      */.
++HOLD     (UZ12345)        /* Hold this PTF         */
           FMID(FXY1040)    /* for this function     */
           ERROR            /* for APAR fix.         */
           REASON(AZ00002)  /* APAR is AZ00002.      */
           COMMENT(this APAR causes 0C4)   /*      */.
++HOLD     (UZ12346)        /* Hold this PTF         */
           FMID(FXY1040)    /* for this function     */
           ERROR            /* for APAR fix.         */
           REASON(AZ00003)  /* APAR is AZ00003.      */
           COMMENT(incorrect output on console)  /*      */.
++HOLD     (UZ12347)        /* Hold this PTF         */
           FMID(FXY1040)    /* for this function     */
           ERROR            /* for APAR fix.         */
           REASON(AZ00004)  /* APAR is AZ00004.      */
                            /* Integrity - no comment.  */.
```

Example 3: Specifying a hold class

Here is an example of a SYSMOD containing a ++HOLD statement for a PTF requiring a change that applies only if a Year 2000-related fix is desired.

```
++PTF(UZ12345)                 /* PTF modification.      */.
++VER(Z038) FMID(FXY1040)   /* For MVS product FXY1040. */
           .                /* No prerequisites.      */
++HOLD     (UZ12345)        /* Hold this PTF         */
           FMID(FXY1040)    /* for this function     */
           ERROR            /* for APAR fix.         */
           REASON(AQ19558)  /* APAR is AQ19558.      */
           COMMENT(SMRTDATA(FIX(UQ21725)
                   SYMP(YR2000) CHGDT(981019)))
           CLASS(YR2000)    /* This change is required  */
                            /* only for Yr 2000 support.*/.
++MOD(IFBMOD01)             /* Change this module    */
           DISTLIB(AOS12)   /* in this DLIB.         */.
```

Example 4: Identifying APARs and their associated fix categories

Here are some examples of using the ++HOLD FIXCAT MCS to identify APARs and their associated fix categories.

```
++HOLD(HBB7730)                    /* Held FMID      */
  FMID(HBB7730)
  FIXCAT  /* Associates an APAR to a fix category */
  REASON(AA15968)                  /* The APAR       */
  CATEGORY(IBM.Device.2094)    /* A fix category  */
  RESOLVER(UA27113)                /* The fixing PTF  */.
```

++IF MCS

The ++IF MCS specifies requisites that must be installed with a SYSMOD if a certain function is installed. If the specified function is installed in the zone that SMP/E is trying to install this SYSMOD into, the requisite must also be installed in that zone; otherwise SMP/E does not install the SYSMOD containing the ++IF MCS. If the specified function is not yet installed in the zone where SMP/E is trying to install this SYSMOD, SMP/E saves the information from the ++IF MCS in case the specified function is installed later.

++IF statements are interpreted, reformatted, and placed in the target zone data set during APPLY processing and in the distribution zone data set during ACCEPT processing.

++IF statements can describe dependencies between two products, even if those products have different SRELs and are defined in different zones. The information from these statements is used by the REPORT command to identify cross-zone requisites—that is, conditional requisites needed in one zone because of ++IF statements in a SYSMOD installed in another zone.

A ++IF statement is associated with the ++VER MCS preceding it in the SYSMOD. Multiple ++IF statements can be specified following each ++VER MCS.

Syntax

++IF MCS

```
►►─ ++IF ── FMID ── ( ── sysmod_id ── ) ─┬────────┬─ REQ ── ( ─┬── sysmod_id ──┬─ ) ─►
                                          └─ THEN ─┘            └◄──── , ───────┘

    ►── · ─►◄
```

Operands

FMID
specifies the function that is to be checked to determine whether it is installed in one of the following:

- The target libraries (for APPLY processing)
- The distribution libraries (for ACCEPT processing)

This operand is required. It is not satisfied by superseded FMIDs.

REQ
specifies the SYSMODs that are needed if the function SYSMOD specified on the FMID operand of the ++IF MCS is installed. This operand is required.

THEN
specifies that the requisite for the ++IF MCS follows.

Usage notes

- The operands must be specified in the order shown in the syntax.
- The function specified for FMID must be different from the SYSMOD ID on the header statement.

Examples

PTF UZ00004 contains service for elements belonging to function FXY1040. If function FXY1050 is installed in the same zone, PTF UZ00005 must also be installed in that zone in order for PTF UZ00004 to be installed successfully. If function FXY1050 is not presently installed, PTF UZ00005 is not required; however, SMP/E saves the ++IF MCS in case function FXY1050 is installed in the future. Then, when FXY1050 is installed, PTF UZ00005 is considered an unsatisfied conditional requisite that must also be installed.

Here is an example of a SYSMOD containing a ++IF statement to define UZ00005 as a conditional requisite:

```
++PTF(UZ00004)              /* PTF SYSMOD.          */.
++VER(Z038) FMID(FXY1040)   /* For MVS product FXY1040. */.
++IF        FMID(FXY1050)   /* If FXY1050 is on, or     */
            THEN            /* when it's put on later,  */
            REQ(UZ00005)    /* UZ00005 is required.     */.
++MOD(IFBMOD01)             /* Replace this module      */
            DISTLIB(AOSFB)  /* in this DLIB.            */.
++MACUPD(IFBMAC01)          /* Update this macro        */
            DISTLIB(IFBMACS) /* in this DLIB.           */.
```

++JAR MCS

The ++JAR MCS describes a Java™ ARchive (JAR) file. The Java *jar* command is used to construct and update such files in a UNIX file system. JAR elements can have any of the following characteristics:

- The record format (RECFM) must be F, FA, FM, FB, FBA, FBM, V, VA, VM, VB, VBA, or VBM.
- Elements with variable-length records cannot contain spanned records.
- The maximum LRECL is 32,654.
- The records can be numbered or unnumbered.

There are MCS statements to replace JAR elements, just as there are MCS statements to replace other types of elements. The ++JARUPD MCS can be used to update JAR elements.

Syntax

The syntax to be used depends on the processing to be performed for the element:

- Adding or replacing the element
- Deleting the element

Adding or replacing a JAR element

++JAR MCS
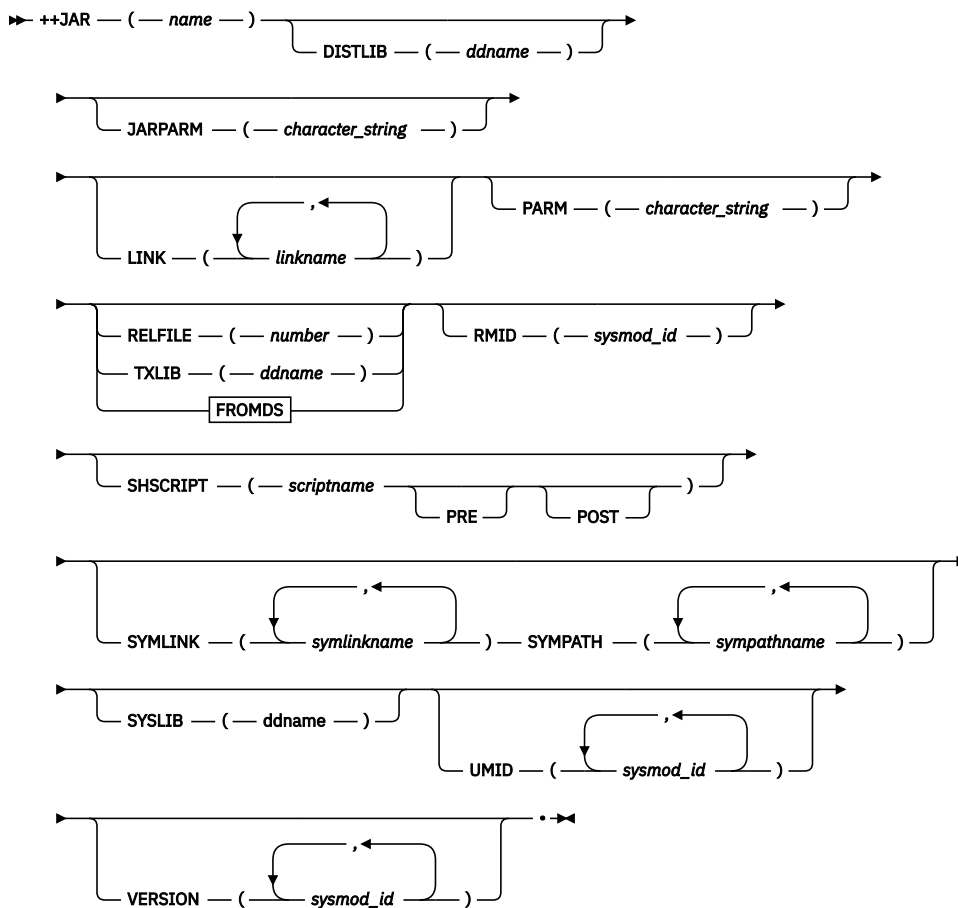
FROMDS

Deleting a JAR element

++JAR MCS

Operands

DELETE

specifies that the JAR element and all of its link names and symbolic link names are to be removed from the "target library" (UNIX file system) and the distribution library.

Note:

1. DELETE is mutually exclusive with all other operands except DISTLIB and VERSION.
2. If the element statement is in a base function, you may want to use the DELETE operand on the ++VER MCS to delete the previous release, rather than on the element statement to delete a specific element.
3. Specification of the DELETE operand results in all link names and symbolic link names of the element being deleted along with the element identified.

DISTLIB

specifies the ddname of the distribution library for the specified JAR element. During ACCEPT processing, SMP/E installs the JAR element into the distribution library as a member. (The distribution library must be a PDS or PDSE; it cannot be part of a UNIX file system.)

Note:

1. DISTLIB must be specified when the JAR element is first installed.
2. If an element entry already exists in the target zone or distribution zone and the value currently in that entry does not match that specified in the DISTLIB operand, the SYSMOD is not applied or accepted.

FROMDS

identifies the partitioned data set that contains this element.

Note: The FROMDS operand and its DSN, NUMBER, VOL, and UNIT suboperands are included in the MCS generated by the BUILDMCS command. IBM does not intend the FROMDS operand to be used in manually coded MCS.

DSN

specifies the dsname of the FROMDS data set. The specified data set name must conform to standard data set naming conventions and cannot contain parentheses. The maximum length of the entire name is 44 characters (including the periods).

NUMBER

specifies a number that SMP/E is to use when assigning a name to the SMPTLIB data set associated with this FROMDS data set. (This is similar to the way the relative file number is used in RELFILE processing.)

VOL

specifies, for an uncataloged data set, the volume serial number of the volume containing the FROMDS data set. If specified, this volume identifier must be from 1 to 6 alphanumeric characters.

VOL may be omitted for a cataloged data set.

UNIT
> specifies, for an uncataloged data set, the UNIT type containing the FROMDS data set. If specified, the UNIT value must be from 1 to 8 characters and must conform to standard UNIT naming conventions. SMP/E accepts any nonblank characters specified between the open and close parentheses, up to a maximum length of 8.
>
> UNIT may be omitted for a cataloged data set.
>
> **Note:** FROMDS is mutually exclusive with DELETE, RELFILE, and TXLIB.

JARPARM
> specifies a character string that is to be passed to the jar command as a command option when updating the JAR file. The maximum length of this character string is 300 bytes of data.
>
> Only those jar command options that require only the option indicator to be specified are supported, such as 0 and M. Options requiring additional input over and above the option indicator are not supported. Examples of unsupported options are the m and -C options.
>
> **Note:**
>
> 1. The character string can include any non-blank characters (characters x'41' through x'FE').
> 2. The JARPARM value is saved and passed to the jar command as follows:
>   - If JARPARM is specified on the element MCS, any value previously saved in the JAR element entry is overlaid and the new value is passed to the jar command.
>   - If JARPARM is not specified on the element MCS and a saved value exists in the JAR element entry, then the saved value is passed to the jar command.
>   - If JARPARM is not specified on the element MCS and there is no saved value in the JAR element entry, then no value is saved in the element entry and no value is passed to the jar command.

LINK
> specifies the alternative names by which this JAR element can be known in a UNIX file system. The full name is produced by concatenating the specified linkname with a UNIX file system directory identified by the SYSLIB subentry. Each linkname is passed to the HFS copy utility as an execution parameter.
>
> **Note:**
>
> 1. The linkname can be from 1 to 1023 characters.
> 2. A linkname can be enclosed in single apostrophes ('). A linkname **must** be enclosed in single apostrophes if any of the following is true:
>   - The linkname contains lowercase alphabetic characters.
>   - The linkname contains a character that is not uppercase alphabetic, numeric, national ($, #, @), slash (/), plus (+), hyphen, period, or ampersand (&).
>   - The linkname spans more than one line in the control statement.
>
>   The single apostrophes used to enclose a linkname (the delimiters) do not count as part of the 1023-character limit.
> 3. Any apostrophes specified as part of a linkname (not the delimiters) must be doubled.
>
>   Double apostrophes count as two characters in the 1023-character limit.
> 4. The linkname can include characters X'40' through X'FE'.
> 5. LINK values are saved and passed to the HFS copy utility as follows:
>   - If LINK is specified on the element MCS, any values previously saved in the element entry are overlaid.
>   - If LINK is not specified on the element MCS and saved values exist in the JAR element entry, the saved values are passed to the HFS copy utility as execution parameters.
>
>   If LINK is not specified on the element MCS and there are no saved values in the JAR element entry, no linknames are passed to the HFS copy utility.

name
specifies the name of the JAR element member. The name can contain any uppercase alphabetic, numeric, or national ($, #, @) character and can be 1 to 8 characters long.

PARM
specifies a character string that is to be passed to the hierarchical file system copy utility as an execution-time parameter. (The values that can be specified on the PARM operand, such as PATHMODE, are those accepted by the BPXCOPY utility. See *z/OS UNIX System Services Command Reference* for a description of BPXCOPY and the values it accepts.) The maximum length of this character string is 300 bytes of nonblank data. If any blanks are specified in the PARM value, they are deleted by SMP/E during processing and do not count toward the 300-byte maximum.

Note:

1. PARM is an optional operand.

2. The character string can be entered free-form, without regard to blanks (which are compressed out of the string), and can span multiple 80-byte records.

3. If parentheses are specified in the PARM value, there must always be a pair (left and right); otherwise, the results are unpredictable.

4. PARM values are saved and passed to the HFS copy utility as follows:

   - If PARM is specified on the element MCS, any values previously saved in the JAR element entry are overlaid.

   - If PARM is not specified on the element MCS and saved values exist in the JAR element entry, the saved values are passed to the HFS copy utility as execution parameters.

     If PARM is not specified on the element MCS and there are no saved values in the JAR element entry, no parameters are saved from the element MCS or passed from the JAR element entry to the HFS copy utility.

5. If the UTILITY entry for the HFS copy utility specifies a PARM value, those parameters are passed to the utility in addition to any parameters saved in the JAR element entry.

RELFILE
identifies which relative file associated with the SYSMOD contains this element. This operand is required if you provide the element in RELFILE format, rather than inline or in a TXLIB data set.

Note:

1. The RELFILE value must be a decimal number from 1 to 9999.

2. RELFILE is mutually exclusive with FROMDS and TXLIB.

RMID
specifies the last SYSMOD that **replaced** this element. This operand can be used only in a service-updated function, and the specified PTF must be integrated into the function.

SHSCRIPT
specifies a UNIX shell script, *scriptname,* to be invoked when the element is installed in (or deleted from) a directory of a UNIX file system. *scriptname* can contain any uppercase alphabetic, numeric, or national ($, #, @) character and can be 1 to 8 characters long.

A shell script is commonly used to complete the installation of an element. For example, if the JAR element requires an external link, you can provide a shell script that performs the necessary steps to create the external link for the JAR element.

scriptname must be the first value to follow the SHSCRIPT operand.

You define the shell script to SMP/E through a ++SHELLSCR statement, which can be within the same SYSMOD as the JAR element, or within a SYSMOD that was processed previously.

You cannot define more than one shell script for an element.

You can follow *scriptname* with either of two optional values, PRE and POST, to specify the point in SMP/E processing when the shell script is to be invoked. The following examples show how you can use the PRE and POST values:

- To run the shell script before the element is copied to a UNIX file system directory, specify:

```
SHSCRIPT(scriptname,PRE)
```

- To run the shell script after the element is copied to a UNIX file system directory, specify:

```
SHSCRIPT(scriptname,POST)
```

or specify no value after *scriptname* to use POST by default.

- To run the shell script both before and after the element is copied to a UNIX file system directory, specify:

```
SHSCRIPT(scriptname,PRE,POST)
```

If you do not specify a PRE or POST value for the shell script, SMP/E invokes the shell script after the element is installed in the directory.

You cannot specify the SHSCRIPT operand with the DELETE operand.

SYMLINK
specifies a list of one or more symbolic links, which are file names that can be used as alternate names for referring to this element in a UNIX file system. Each linkname listed here is associated with a pathname listed in the SYMPATH operand. For more information about how the linknames and pathnames are associated, see the description of the SYMPATH operand and "Example 3: Packaging a SYSMOD with a symbolic link" on page 30.

The SYMLINK value specified should be a relative path value (that is, it does not start with a slash ["/"]). When the symbolic link is created, it is created relative to the pathname of the element's SYSLIB ddname.

SYMLINK must be specified if the SYMPATH operand is specified, otherwise it must be omitted.

A symbolic linkname can be from one to 1023 characters. Any characters in the range X'40' through X'FE' may be specified.

The value may be enclosed in single apostrophes. It **must** be enclosed in single apostrophes if:

- it is continued to the next line in the MCS, or
- it contains a character that is not uppercase alphabetic, numeric, national ($, #, @), slash (/), plus (+), hyphen, period, or ampersand (&).

If an apostrophe is a part of the symbolic linkname and is not a delimiter, then it must be doubled. These two apostrophes count as two characters against the 1023 character limit for a symbolic linkname. The single apostrophes used to enclose a symbolic linkname do not count against the 1023 character limit.

SYMPATH
specifies a list of one or more pathnames that are associated with symbolic links identified by the SYMLINK operand. The first pathname in the SYMPATH operand is associated with the first symbolic link in the SYMLINK operand, the second pathname with the second symbolic link, and so on. If there are more symbolic links listed than there are pathnames, then the last listed pathname is used for the remaining symbolic links. If more pathnames are specified than symbolic linknames, then the excess pathnames (at the end of the list) are ignored.

The SYMPATH value specified should be a relative path value. When the symbolic link is accessed, the system assumes the destination of that link (the SYMPATH value) is relative to that symbolic link (the SYMLINK value). For more information about how the pathnames and linknames are associated, see "Example 3: Packaging a SYSMOD with a symbolic link" on page 30.

SYMPATH must be specified if the SYMLINK operand is specified, otherwise it must be omitted.

A symbolic pathname can be one to 1023 characters. Any characters in the range X'40' through X'FE' may be specified.

The value may be enclosed in single apostrophes. It **must** be enclosed in single apostrophes, if:

- it is continued to the next line in the MCS, or
- it contains a character that is not uppercase alphabetic, numeric, national ($, #, @), slash (/), plus (+), hyphen, period, or ampersand (&).

If an apostrophe is a part of the symbolic pathname and is not a delimiter, then it must be doubled. These two apostrophes count as two characters against the 1023 character limit for a symbolic pathname. The single apostrophes used to enclose a symbolic linkname do not count against the 1023 character limit.

SYSLIB
specifies the ddname of the "target library" within the UNIX file system for the element.

During APPLY processing, the HFS copy utility installs the JAR element into a UNIX file system. During RESTORE processing, the HFS copy utility copies the JAR element from the distribution library member into a UNIX file system.

Note: SYSLIB must be specified when the JAR element is first installed.

TXLIB
is the ddname of the partitioned data set containing the JAR element. This operand is required if the JAR element is provided in a data set that the users have access to, rather than inline or in RELFILE format.

Note:

1. SMPTLIB cannot be used as a value on the TXLIB operand.
2. TXLIB is mutually exclusive with FROMDS and RELFILE.

UMID
specifies the SYSMODs that have updated this JAR file since it was last replaced. This operand can be used only in a service-updated function, and the specified SYSMODs must be integrated into the function.

VERSION
specifies one or more function SYSMODs that currently contain the element. The function containing the element MCS takes over ownership of the element from the specified functions.

When VERSION is specified on an element statement, it overrides any VERSION operand values that might be specified on the ++VER MCS.

Usage notes

- If the JAR element is packaged inline, it must immediately follow the ++JAR MCS and must not contain any records starting with ++. Neither FROMDS, nor RELFILE, nor TXLIB can be specified on the ++JAR MCS when the element is packaged inline.
- To be packaged inline, a JAR file element must contain fixed-block-80 records. To achieve this, the original element can be extracted from a UNIX file system into a data set (using the TSO command OGET for example), and then GIMDTS can be used to transform the element into fixed-block-80 records.

Examples

Here is an example of the ++JAR MCS.

```
++JAR(ABCTTT) DISTLIB(AABCBIN) SYSLIB(SABCBIN) RELFILE(2)
  PARM(PATHMODE(0,6,4,4))
  LINK('../TicTacToe.jar')
  SYMLINK('../../../../../usr/lib/TicTacToe.jar')
  SYMPATH('../../usr/lpp/abc/bin/TicTacToe.jar').
```

++JARUPD MCS

The ++JARUPD MCS describes an update to a Java ARchive (JAR) file in a UNIX file system. The Java *jar* command is used to construct and update such files in a UNIX file system. A JARUPD element is itself a Java Archive file, but it contains only the component files to be added to or replaced in an existing JAR file, as opposed to a complete replacement for an existing JAR file.

Note:

1. To update JAR files, SMP/E must use the update (u) option of the *jar* command, which is provided in version 1.2 of the Java Development Kit (JDK). Therefore, in order to perform APPLY or ACCEPT command operations for a SYSMOD that contains JARUPD elements, SMP/E requires the IBM Developer Kit for OS/390, Java 2 Technology Edition, or its successor. Java 2 Technology Edition is a no-charge product available for users of OS/390 Version 2 Release 8 and above, or z/OS Version 1 Release 1 and above.

   Java must be available in the execution environment for SMP/E. The SMPJHOME DD statement or DDDEF entry should be used to specify the directory where Java resides. For example, if Java 1.4, which has been installed in the /usr/lpp/java/J1.4 directory, should be used, then the following DD statement should be specified:

   ```
   //SMPJHOME DD PATH='/usr/lpp/java/J1.4/',PATHDISP=KEEP
   ```

   For additional information about how to set the SMPJHOME value, see Chapter 4, "SMP/E data sets and files," on page 113.

2. There is no way to delete component files from a JAR file. ++JARUPD can be used only to add or replace component files in a JAR file. If component files must be deleted, then the entire JAR file must be replaced with another JAR file from which those component files have been removed.

Syntax

++JARUPD MCS
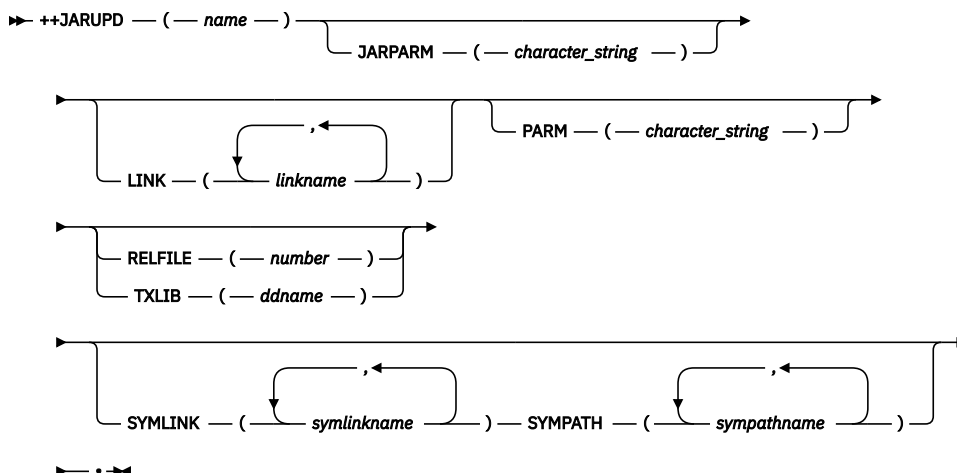
Operands

JARPARM

specifies a character string that is to be passed to the jar command as a command option when updating the JAR file. The maximum length of this character string is 300 bytes of data.

Only those jar command options that require only the option indicator to be specified are supported, such as 0 and M. Options requiring additional input over and above the option indicator are not supported. Examples of unsupported options are the m and -C options.

Note:

1. The character string can include any non-blank characters (characters x'41' through x'FE').
2. The JARPARM value is saved and passed to the jar command as follows:
   - If JARPARM is specified on the ++JARUPD, any value previously saved in the JAR element entry is overlaid and the new value is passed to the jar command.
   - If JARPARM is not specified on the ++JARUPD and a saved value exists in the JAR element entry, then the saved value is passed to the jar command.
   - If JARPARM is not specified on the ++JARUPD and there is no saved value in the JAR element entry, then no value is saved in the element entry and no value is passed to the jar command.

LINK
specifies the alternative names by which this JAR element can be known in a UNIX file system. The full name is produced by concatenating the specified linkname with a UNIX file system directory identified by the SYSLIB subentry. Each linkname is passed to the HFS copy utility as an execution parameter.

Note:

1. The linkname can be from 1 to 1023 characters.
2. A linkname can be enclosed in single apostrophes ('). A linkname must be enclosed in single apostrophes if any of the following is true:
   - The linkname contains lowercase alphabetic characters.
   - The linkname contains a character that is not uppercase alphabetic, numeric, national ($, #, @), slash (/), plus (+), hyphen, period, or ampersand (&).
   - The linkname spans more than one line in the control statement.

   The single apostrophes used to enclose a linkname (the delimiters) do not count as part of the 1023-character limit.
3. Any apostrophes specified as part of a linkname (not the delimiters) must be doubled.

   Double apostrophes count as two characters in the 1023-character limit.
4. The linkname can include characters X'40' through X'FE'.
5. LINK values are saved and passed to the HFS copy utility as follows:
   - If LINK is specified on the ++JARUPD, any values previously saved in the element entry are overlaid.
   - If LINK is not specified on the ++JARUPD and saved values exist in the JAR element entry, the saved values are passed to the HFS copy utility as execution parameters.

   If LINK is not specified on the ++JARUPD and there are no saved values in the JAR element entry, no linknames are passed to the HFS copy utility.

name
specifies the name of the JAR element member. The name can contain any uppercase alphabetic, numeric, or national ($, #, @) character and can be 1 to 8 characters long.

PARM
specifies a character string that is to be passed to the hierarchical file system copy utility as an execution-time parameter. (The values that can be specified on the PARM operand, such as PATHMODE, are those accepted by the BPXCOPY utility. See *z/OS UNIX System Services Command Reference* for a description of BPXCOPY and the values it accepts.) The maximum length of this character string is 300 bytes of nonblank data. If any blanks are specified in the PARM value, they are deleted by SMP/E during processing and do not count toward the 300-byte maximum.

Note:

1. PARM is an optional operand.
2. The character string can be entered free-form, without regard to blanks (which are compressed out of the string), and can span multiple 80-byte records.

3. If parentheses are specified in the PARM value, there must always be a pair (left and right); otherwise, the results are unpredictable.

4. PARM values are saved and passed to the HFS copy utility as follows:

- If PARM is specified on the ++JARUPD, any values previously saved in the JAR element entry are overlaid.

- If PARM is not specified on the ++JARUPD and saved values exist in the JAR element entry, the saved values are passed to the HFS copy utility as execution parameters.

  If PARM is not specified on the ++JARUPD and there are no saved values in the JAR element entry, no parameters are saved from the ++JARUPD or passed from the JAR element entry to the HFS copy utility.

5. If the UTILITY entry for the HFS copy utility specifies a PARM value, those parameters are passed to the utility in addition to any parameters saved in the JAR element entry.

RELFILE

identifies which relative file associated with the SYSMOD contains this element. This operand is required if you provide the element in RELFILE format, rather than inline or in a TXLIB data set.

Note:

1. The RELFILE value must be a decimal number from 1 to 9999.

2. RELFILE is mutually exclusive with TXLIB.

SYMLINK

specifies a list of one or more symbolic links, which are file names that can be used as alternate names for referring to this element in a UNIX file system. Each linkname listed here is associated with a pathname listed in the SYMPATH operand. For more information about how the linknames and pathnames are associated, see the description of the SYMPATH operand and "Example 3: Packaging a SYSMOD with a symbolic link" on page 30.

The SYMLINK value specified should be a relative path value (that is, it does not start with a slash ["/"]). When the symbolic link is created, it is created relative to the pathname of the element's SYSLIB ddname.

SYMLINK must be specified if the SYMPATH operand is specified, otherwise it must be omitted.

A symbolic linkname can be from one to 1023 characters. Any characters in the range X'40' through X'FE' may be specified.

The value may be enclosed in single apostrophes. It **must** be enclosed in single apostrophes if:

- it is continued to the next line in the MCS, or

- it contains a character that is not uppercase alphabetic, numeric, national ($, #, @), slash (/), plus (+), hyphen, period, or ampersand (&).

If an apostrophe is a part of the symbolic linkname and is not a delimiter, then it must be doubled. These two apostrophes count as two characters against the 1023 character limit for a symbolic linkname. The single apostrophes used to enclose a symbolic linkname do not count against the 1023 character limit.

SYMPATH

specifies a list of one or more pathnames that are associated with symbolic links identified by the SYMLINK operand. The first pathname in the SYMPATH operand is associated with the first symbolic link in the SYMLINK operand, the second pathname with the second symbolic link, and so on. If there are more symbolic links listed than there are pathnames, then the last listed pathname is used for the remaining symbolic links. If more pathnames are specified than symbolic linknames, then the excess pathnames (at the end of the list) are ignored.

The SYMPATH value specified should be a relative path value. When the symbolic link is accessed, the system assumes the destination of that link (the SYMPATH value) is relative to that symbolic link (the SYMLINK value). For more information about how the pathnames and linknames are associated, see "Example 3: Packaging a SYSMOD with a symbolic link" on page 30.

SYMPATH must be specified if the SYMLINK operand is specified, otherwise it must be omitted.

A symbolic pathname can be one to 1023 characters. Any characters in the range X'40' through X'FE' may be specified.

The value may be enclosed in single apostrophes. It **must** be enclosed in single apostrophes, if:

- it is continued to the next line in the MCS, or
- it contains a character that is not uppercase alphabetic, numeric, national ($, #, @), slash (/), plus (+), hyphen, period, or ampersand (&).

If an apostrophe is a part of the symbolic pathname and is not a delimiter, then it must be doubled. These two apostrophes count as two characters against the 1023 character limit for a symbolic pathname. The single apostrophes used to enclose a symbolic linkname do not count against the 1023 character limit.

TXLIB

is the ddname of the partitioned data set containing the JAR element. This operand is required if the JAR element is provided in a data set that the users have access to, rather than inline or in RELFILE format.

Note:

1. SMPTLIB cannot be used as a value on the TXLIB operand.
2. TXLIB is mutually exclusive with RELFILE.

Usage notes

- If the JAR update file is packaged inline, it must immediately follow the ++JARUPD MCS and must not contain any records starting with ++. Neither RELFILE nor TXLIB can be specified on the MCS when the element is packaged inline.
- To be packaged inline, a JAR file element must contain fixed-block-80 records. To do this, the original element can be extracted from a UNIX file system into a data set (using, for example, the TSO command OGET), and then GIMDTS can be used to transform the element into fixed-block-80 records.

Examples

Here is an example of the ++JARUPD MCS.

```
++JARUPD(ABCTTT) JARPARM(0M)
  PARM(PATHMODE(0,6,4,4))
  LINK('../TicTacToe.jar')
  SYMLINK('../../../../../usr/lib/TicTacToe.jar')
  SYMPATH('../../usr/lpp/abc/bin/TicTacToe.jar').
```

++JCLIN MCS

The ++JCLIN MCS identifies the start of the job control language (JCL) input that must be processed as part of installing the SYSMOD. The JCLIN input is used to identify the structure of the target system libraries to SMP/E so that the modules, macros, and source code being replaced can be installed correctly. Processing of ++JCLIN input during APPLY or ACCEPT is similar to the processing done for the JCLIN command. For an additional description of JCLIN processing, see *z/OS SMP/E Commands*.

Note: The ++JCLIN MCS does **not** cause SMP/E to update the target or distribution libraries; only the entries in the target and distribution zones are updated. These libraries are updated when SMP/E processes the elements in the SYSMOD. The element statements in the SYSMOD determine which elements should be installed and cause SMP/E to invoke the utilities that install the elements.

Syntax

++JCLIN MCS
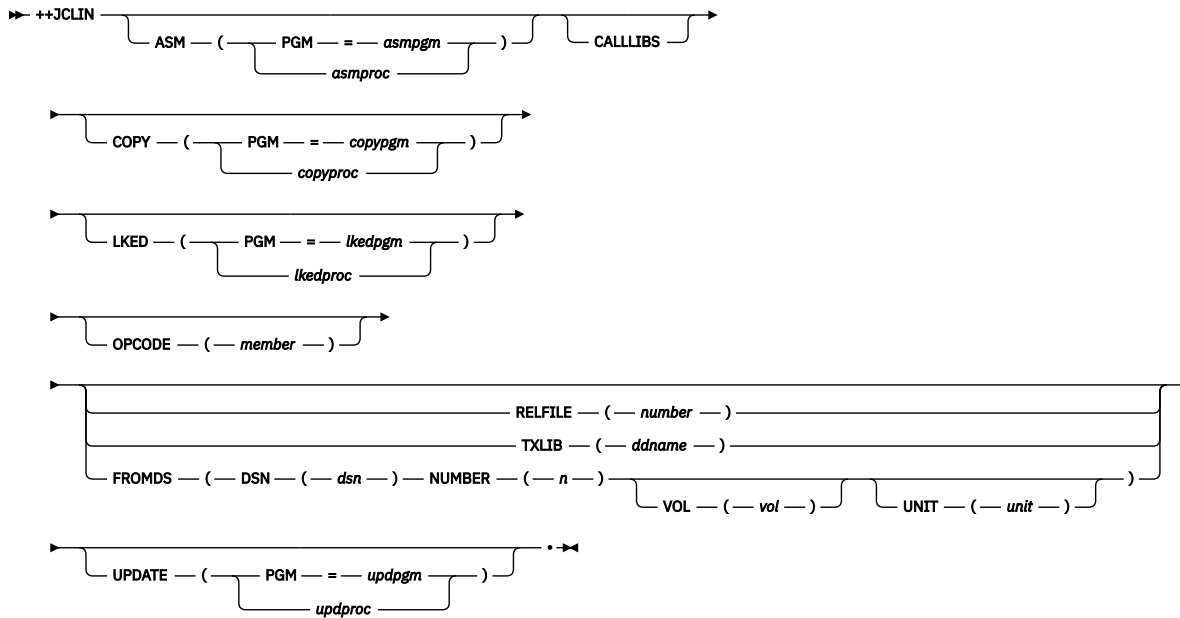
Operands

ASM
specifies the name of the assembler program or procedure that is used in the JCLIN data. This operand must be specified if the name is different from those recognized by SMP/E, which are the program names ASMBLR, ASMA90, IEUASM, IEV90, and IFOX00, and procedure name ASMS.

CALLLIBS
specifies that SMP/E is to process SYSLIB DD statements in JCLIN link-edit steps. SYSLIB DD statements are processed only if the CALLLIBS operand is specified on the JCLIN command or ++JCLIN MCS, or if //*CALLLIBS=YES is encountered after a job card preceding a link-edit step. If the CALLLIBS operand or the CALLLIBS comment is not specified, SMP/E ignores any SYSLIB DD statements it encounters.

COPY
specifies the name of the copy program or procedure that is used in the JCLIN data. This operand must be specified if the name is different from that recognized by SMP/E, which is the program name IEBCOPY.

FROMDS
identifies the partitioned data set that contains this element.

Note: The FROMDS operand and its DSN, NUMBER, VOL, and UNIT suboperands are included in the MCS generated by the BUILDMCS command. IBM does not intend the FROMDS operand to be used in manually coded MCS.

DSN
specifies the dsname of the FROMDS data set. The specified data set name must conform to standard data set naming conventions and cannot contain parentheses. The maximum length of the entire name is 44 characters (including the periods).

NUMBER
specifies a number that SMP/E is to use when assigning a name to the SMPTLIB data set associated with this FROMDS data set. (This is similar to the way the relative file number is used in RELFILE processing.)

VOL

specifies, for an uncataloged data set, the volume serial number of the volume containing the FROMDS data set. If specified, this volume identifier must be from 1 to 6 alphanumeric characters.

VOL may be omitted for a cataloged data set.

UNIT

specifies, for an uncataloged data set, the UNIT type containing the FROMDS data set. If specified, the UNIT value must be from 1 to 8 characters and must conform to standard UNIT naming conventions. SMP/E accepts any nonblank characters specified between the open and close parentheses, up to a maximum length of 8.

UNIT may be omitted for a cataloged data set.

Note: FROMDS is mutually exclusive with DELETE, RELFILE, and TXLIB.

LKED

specifies the name of the link-edit program or procedure that is used in the JCLIN data. This operand must be specified if the name is different from those recognized by SMP/E, which are the program names HEWL, HEWLH096, HEWLKED, IEWBLINK, IEWL, and LINKEDIT, and procedure name LINKS.

OPCODE

identifies an OPCODE member containing control statements identifying character strings to be treated as OPCODEs or macros.

PGM

specifies a program name (instead of a procedure name) for a utility.

RELFILE

specifies the relative position, within the files associated with this SYSMOD, of the file containing the JCLIN data as one of its members.

The member in the RELFILE data set containing the JCLIN input data must match the SYSMOD's SYSMOD ID; for example, the JCLIN for ++PTF(UZ11111) is supplied in member UZ11111 of the unloaded PDS identified by the RELFILE operand.

Note: RELFILE is mutually exclusive with FROMDS and TXLIB.

TXLIB

specifies the ddname of a partitioned data set containing the JCLIN data as one of its members.

The member of the TXLIB data set containing the JCLIN input data must match the SYSMOD's SYSMOD ID; for example, the JCLIN for ++PTF(UZ11111) is supplied in member UZ11111 of the data set identified by the TXLIB operand.

Note:

1. SMPTLIB cannot be used as a value on the TXLIB operand.
2. TXLIB is mutually exclusive with FROMDS and RELFILE.

UPDATE

specifies the name of the update program or procedure that is used in the JCLIN data. This operand must be specified if the name is different from that recognized by SMP/E, which is the program name IEBUPDTE.

Usage notes

- If the JCLIN data is packaged inline, it must immediately follow the ++JCLIN MCS and must not contain any records that start with ++. Neither FROMDS, nor RELFILE, nor TXLIB can be specified on the ++JCLIN MCS.

- If the JCLIN data is packaged in a TXLIB data set, the ddname specified on the TXLIB operand is required during APPLY and ACCEPT processing.

- For information about packaging SYSMODs in RELFILE format, see Standard Packaging Rules for z/OS-based Products (publibz.boulder.ibm.com/epubs/pdf/gimpkg80.pdf).

Examples

For these examples, assume function JXY1040 contains some inline JCLIN, one module (IFBMOD01), and one macro (IFBMAC01). The same format was used to package the JCLIN and the elements.

Note: Even though the examples are for a function, they are equally valid for PTFs, APARs, and USERMODs.

The following examples are provided to help you use the ++JCLIN MCS:

Example 1: ++JCLIN data packaged inline

Here is an example of a SYSMOD containing a ++JCLIN statement for packaging JCLIN inline:

```
++FUNCTION(JXY1040)         /* Function SYSMOD.       */.
++VER(Z038)                 /* For MVS, no requisites.  */.
++JCLIN                     /* JCLIN is required.       */.
//JOB     JOB 'accounting info',MSGLEVEL=(1,1)
//STEP1    EXEC PGM=IEWL
//SYSLMOD  DD DSN=SYS1.LINKLIB,DISP=SHR
//AOS12    DD DSN=SYS1.AOS12,DISP=SHR
//SYSLIN   DD *
 INCLUDE AOS12(IFBMOD01)
 ENTRY   IFBMOD01
 ALIAS   IFBMOD01A
 NAME    IFBMOD01(R)
/*
++MOD(IFBMOD01)            /* This module              */
        DISTLIB(AOS12)  /* for this DLIB.           */.
...
... object deck for IFBMOD01
...
++MAC(IFBMAC01)            /* This macro               */
        DISTLIB(AMACLIB) /* for this DLIB.           */.
...
... macro replacement for IFBMAC01
...
```

No special DD statements, other than those for the target and DLIBs, are required to process this SYSMOD.

During APPLY, SMP/E processes the ++JCLIN data, determines that module IFBMOD01 gets linked into load module IFBMOD01 (which has an alias), stores information in the target zone, and then invokes the proper utilities to install the module and macro.

Example 2: ++JCLIN data packaged in a RELFILE

Here is an example of a SYSMOD containing a ++JCLIN statement for packaging JCLIN in relative files. No special ++JCLIN operands are required.

```
++FUNCTION(JXY1040)         /* Function SYSMOD         */
        FILES(3)          /* in RELFILE format.      */.
++VER(Z038)                 /* For MVS, no requisites.  */.
++JCLIN                     /* JCLIN is required.       */
        RELFILE(1)        /* Data is in first RELFILE.*/.
++MOD(IFBMOD01)             /* This module             */
        DISTLIB(AOS12)  /* for this DLIB            */
        RELFILE(2)        /* is in second RELFILE.    */.
++MAC(IFBMAC01)             /* This macro              */
        DISTLIB(AMACLIB)/* for this DLIB            */
        RELFILE(3)        /* is in third RELFILE.     */.
```

The SMPTLIB DD statement is required for the RECEIVE command to load the RELFILE data sets from tape onto DASD, and at APPLY and ACCEPT, in order to access the data in the unloaded partitioned data sets on DASD.

Note: Different RELFILE data sets are required for the module and macro, because they are in different formats (modules must be in load module format, RECFM=U, while macros are RECFM=FB). Although a separate RELFILE data set was used to contain the ++JCLIN data, that data could have been put into the data set with the macros, because both have the same attributes (LRECL=80, RECFM=FB).

Example 3: ++JCLIN data packaged in a TXLIB with a user utility program name

Here is an example of a SYSMOD containing a ++JCLIN statement for packaging JCLIN in text libraries. The JCLIN contains assembly steps using a special link-edit utility, ALTIEWL.

```
++FUNCTION(JXY1040)        /* Function SYSMOD.       */.
++VER(Z038)                /* For MVS, no requisites. */.
++JCLIN                    /* JCLIN is required.      */
        TXLIB(LIB1040)  /* Data is in text library. */
        LKED(PGM=ALTIEWL)  /* Use other link-edit.  */.
++MOD(IFBMOD01)            /* This module             */
        DISTLIB(AOS12)  /* for this DLIB            */
        TXLIB(LIB1040)  /* is in text library.      */.
++MAC(IFBMAC01)            /* This macro              */
        DISTLIB(AMACLIB)/* for this DLIB            */
        TXLIB(LIB1040)  /* is in text library.      */.
```

A DD statement for LIB0104 is required during APPLY and ACCEPT in order for SMP/E to call the utilities to get the replacements installed into the operating system libraries. An example of the DD statement is:

```
//LIB1040  DD DSN=…
```

Note: In this case, the replacements for all three elements were put in one data set. That data set must contain three members: member JXY1040 must be the JCLIN data, member IFBMOD01 must contain the object deck for module IFBMOD01, and member IFBMAC01 must contain the replacement for macro IFBMAC01. Because all have the same formats (that is, RECFM=FB LRECL=80) they can be packaged in one data set.

Member JXY1040 looks as follows:

```
//JOB      JOB 'accounting info',MSGLEVEL=(1,1)
//STEP1    EXEC PGM=ALTIEWL
//SYSLMOD  DD DSN=SYS1.LINKLIB,DISP=SHR
//AOS12    DD DSN=SYS1.AOS12,DISP=SHR
//SYSLIN   DD *
 INCLUDE AOS12(IFBMOD01)
 ENTRY   IFBMOD01
 ALIAS   IFBMOD01A
 NAME    IFBMOD01(R)
 /*
```

Note the use of ALTIEWL on the EXEC statement.

++MAC MCS

The ++MAC MCS describes a single macro replacement. It must immediately precede the macro definition statements when they are within the SYSMOD.

Syntax

The syntax to be used depends on the processing to be done for the element:

- Adding or replacing the element
- Deleting the element

Adding or replacing a macro

++MAC MCS

```
►►── ++MAC ──(── name ──)──────────────────────────────────────►
                        └─ ASSEM ──(─┬─ module ─┬─)─┘
                                     └────,◄────┘

    ┌──────────────────────────────┐
────┤                              ├──────────────────────────►
    └┬─ DISTMOD ──(── ddname ──)─┬─┘
     └─ DISTSRC ──(── ddname ──)─┘

────┬─ DISTLIB ──(── ddname ──)─┬──────────────────────────────►
    └──────────────────────────────┘
                                 └─ MALIAS ──(─┬─ alias ─┬─)─┘
                                               └───,◄────┘

────┬──────────────────────────┬──────────────────────────────►
    └─ PREFIX ──(─┬─ prefix ─┬─)─┘
                  └───,◄─────┘

────┬─────────────────────────────────────────────────────────┬────►
    ├─ RELFILE ──(── number ──)─────────────────────────────────┤
    ├─ TXLIB ──(── ddname ──)───────────────────────────────────┤
    ├─ SSI ──(── code ──)───────────────────────────────────────┤
    └─ FROMDS ──(── DSN ──(── dsn ──)── NUMBER ──(── n ──)───────┘
                   └─ VOL ──(── vol ──)─┘ └─ UNIT ──(── unit ──)─┘ )

────┬─ RMID ──(── sysmod_id ──)─┬──┬─ SYSLIB ──(── ddname ──)─┬──►
    └──────────────────────────┘  └───────────────────────────┘

────┬──────────────────────────────┬───────────────────────────►
    └─ UMID ──(─┬─ sysmod_id ─┬─)────┘
                └────,◄───────┘

────┬──────────────────────────────┬── .──►◄
    └─ VERSION ──(─┬─ sysmod_id ─┬─)─┘
                   └────,◄───────┘
```

Deleting a macro

++MAC MCS

```
►►── ++MAC ──(── name ──)── DELETE ────────────────────────────►
                                    └─ DISTLIB ──(── ddname ──)─┘

────┬──────────────────────────────┬── .──►◄
    └─ VERSION ──(─┬─ sysmod_id ─┬─)─┘
                   └────,◄───────┘
```

Operands

ASSEM
> specifies the names of modules to be assembled in addition to those modules named as GENASM subentries in the MAC entry. The source for the assemblies is the ASSEM entry, SRC entry, or DISTSRC member whose name matches the specified ASSEM value. SMP/E looks for a match—first with an ASSEM entry, then with a SRC entry, and finally with an entry in the DISTSRC data set—and uses the first match it finds. The source must either be known to SMP/E at the time the ASSEM operand is encountered on the ++MAC statement, or be included in the same SYSMOD.

> **Note:**

> 1. APPLY and ACCEPT processing place the specified names into the SYSMOD entry created on the target zone and distribution zone.

2. If the object deck for the element specified on the ASSEM operand is also provided by the SYSMOD, the assembly may not occur. (See *z/OS SMP/E Commands* for more information.)

DELETE

specifies that the macro and alias names are to be removed from the target libraries, distribution libraries, and SMP/E data sets.

Note:

1. DELETE is mutually exclusive with all other operands except DISTLIB and VERSION.

2. If the element statement is in a base function, you may want to use the DELETE operand on the ++VER MCS to delete the previous release, rather than on the element statement to delete a specific element.

DISTLIB

specifies the ddname of the distribution library for the specified macro.

Note:

1. DISTLIB must be specified if the macro has not been previously recorded on the target zone or distribution zone. If a MAC entry already exists in the target zone or distribution zone and the value currently in that entry does not match that specified in the DISTLIB operand, the SYSMOD is not applied or accepted, unless that SYSMOD also used the ++MOVE MCS to change the DISTLIB to that new value.

2. You cannot use SYSPUNCH as the DISTLIB. It is used by SMP/E and other products to process assembled modules.

DISTMOD

specifies the ddname of the link-edit distribution library for those modules specified in the ASSEM operand. During ACCEPT processing, the object code from the assembler is link-edited to the library specified.

DISTSRC

specifies the ddname of the library containing the additional assembly or source to be assembled. The additional assembly or source must be specified in the ASSEM operand.

FROMDS

identifies the partitioned data set that contains this element.

Note: The FROMDS operand and its DSN, NUMBER, VOL, and UNIT suboperands are included in the MCS generated by the BUILDMCS command. IBM does not intend the FROMDS operand to be used in manually coded MCS.

DSN

specifies the dsname of the FROMDS data set. The specified data set name must conform to standard data set naming conventions and cannot contain parentheses. The maximum length of the entire name is 44 characters (including the periods).

NUMBER

specifies a number that SMP/E is to use when assigning a name to the SMPTLIB data set associated with this FROMDS data set. (This is similar to the way the relative file number is used in RELFILE processing.)

VOL

specifies, for an uncataloged data set, the volume serial number of the volume containing the FROMDS data set. If specified, this volume identifier must be from 1 to 6 alphanumeric characters.

VOL may be omitted for a cataloged data set.

UNIT

specifies, for an uncataloged data set, the UNIT type containing the FROMDS data set. If specified, the UNIT value must be from 1 to 8 characters and must conform to standard UNIT naming conventions. SMP/E accepts any nonblank characters specified between the open and close parentheses, up to a maximum length of 8.

UNIT may be omitted for a cataloged data set.

Note: FROMDS is mutually exclusive with DELETE, RELFILE, SSI, and TXLIB.

MALIAS
specifies the alias names for the macro in both the target system and the distribution libraries.

Note: MALIAS must be specified on the ++MAC MCS even if ALIAS was specified on the COPY SELECT statement. MALIAS is required for the RECEIVE command to properly handle aliases in RELFILEs. During RECEIVE processing, SMP/E copies RELFILEs to SMPTLIB data sets. All element members and their associated aliases are copied. Because JCLIN is not processed during RECEIVE, the aliases must be identified on the MCS to get RECEIVE to copy the aliases.

You can use MALIAS when two or more macros that must be defined in the same zone must have the same name for programming access. For example, you can use MALIAS if several products have a help macro whose name must match the name of the command processing module it describes. You can specify HELP on MALIAS and a unique element name as the macro name.

name
specifies the name of the macro member in the distribution library and, optionally, in the target system library. The name can contain any alphanumeric characters and $, #, @, or hex C0.

PREFIX
specifies the first characters (prefix) of the names of modules to be assembled in addition to those modules named as GENASM subentries in the target zone MAC entry. The prefix values must contain no more than 7 characters.

The full module names are determined by comparing the prefix with the target zone or distribution zone MOD entry names.

The source for the assembly is the ASSEM entry, SRC entry, or DISTSRC member whose name matches a MOD entry name beginning with one of the specified prefixes. SMP/E looks for a match— first with an ASSEM entry, then with a SRC entry, finally with an entry in the DISTSRC data set—and uses the first match it finds. The source must either be known to SMP/E at the time the PREFIX operand is encountered on the ++MAC statement, or be included in the same SYSMOD.

Note: If the object deck for an element selected by the PREFIX operand is also provided by the SYSMOD, the assembly may not occur. (See *z/OS SMP/E Commands* for more information.)

RELFILE
identifies which relative file associated with the SYSMOD contains this element.This operand is required if you provide the element in RELFILE format, rather than inline or in a TXLIB data set.

Note: RELFILE is mutually exclusive with FROMDS and TXLIB.

RMID
specifies the last SYSMOD that **replaced** this macro. This operand may be used only in a service-updated function, and the specified PTF must be integrated into the function.

SSI
specifies eight hexadecimal digits of system status information. This information is placed in the directory of the target system library or SMPMTS or SMPSTS during APPLY processing, and in the distribution library during ACCEPT processing, as four packed hexadecimal bytes of user data. See the IEBUPDTE program description in *z/OS DFSMSdfp Utilities*.

Note: This operand is ignored if text is located in a library, as is the case when either the FROMDS, RELFILE, or TXLIB operand is specified.

SYSLIB
specifies the ddname of the target library, if the macro exists in one. APPLY and RESTORE processing update this library.

Note: If a MAC entry already exists in the target zone or distribution zone and the value currently in that entry does not match that specified in the SYSLIB operand, SMP/E ignores the SYSLIB value in the SYSMOD being installed, unless that SYSMOD also used the ++MOVE MCS to change the SYSLIB to that new value.

TXLIB

is the ddname of the partitioned data set containing the macro.This operand is required if the macro is provided in a data set that the users have access to, rather than inline or in RELFILE format.

Note:

1. SMPTLIB cannot be used as a value on the TXLIB operand.

2. TXLIB is mutually exclusive with FROMDS and RELFILE.

UMID

specifies the SYSMODs that have **updated** this macro since it was last replaced. This operand can be used only in a service-updated function, and the specified PTFs must be integrated into the function.

VERSION

specifies one or more function SYSMODs that currently contain the element. The function containing the ++MAC MCS will take over ownership of the element from the specified functions.

When VERSION is specified on an element statement, it overrides any VERSION operand values that might be specified on the ++VER MCS.

Usage notes

- If the macro is packaged inline, it must immediately follow the ++MAC MCS and must not contain any records that start with ++. Neither FROMDS, nor RELFILE, nor TXLIB can be specified on the ++MAC MCS.

- If the macro is packaged in a TXLIB data set, the ddname specified in the TXLIB operand is required during APPLY and ACCEPT processing.

- For information about elements packaged in RELFILE format, see Standard Packaging Rules for z/OS-based Products (publibz.boulder.ibm.com/epubs/pdf/gimpkg80.pdf).

- If the ++MAC MCS is for an inner macro (a macro referred to by another macro instruction that resides in the macro library), the modules that require reassembly must be specified in the ASSEM operand list.

- If the macro resides in a target library (rather than the SMPMTS), the target library should be included in the SYSLIB DD concatenation for assemblies during APPLY processing.For additional information about SYSLIB requirements, see *z/OS SMP/E User's Guide* and *z/OS SMP/E Commands*.

Examples

The following examples are provided to help you use the ++MAC MCS.

Example 1: Replacing a macro through a USERMOD

Assume you want to replace a macro, named YOURMAC, that is part of a product you own, named YOURPROD. You have updated a copy of the macro in a text library, TSO.NEWSMP.MACLIB. Here is an example of a SYSMOD containing a ++MAC statement to build a USERMOD that installs the changes:

```
++USERMOD(USR0001)          /* User modification      */.
++VER(Z038) FMID(HUSR003)   /* for your product       */.
++MAC(YOURMAC)              /* to replace this macro. */
          TXLIB(NEWSMP)     /* Macro is in this TXLIB. */.
```

In this example, you have just applied another piece of service to the YOURPROD macro, but YOURPROD is still the owner of the macro. If you attempt to install some vendor-supplied service (that is, a PTF) to that macro, SMP/E issues an error message indicating that your user modification will be regressed, and will not install that service until the BYPASS(ID) operand is used.

Another method of installing the new macro is for you to assume ownership for the macro by using the VERSION operand. Assume you already have a user function, JXY1040, installed, and you want to transfer ownership of the SMP/E macro to your function. The following SYSMOD contains a ++MAC statement to do that:

```
++USERMOD(USR0001)        /* User modification     */.
++VER(Z038) FMID(JXY1040) /* for user application  */.
++MAC(GIMOPCDE)           /* to replace this macro. */
         VERSION() /* Version SMP/E.        */
         TXLIB(NEWSMP)  /* Macro is in this TXLIB. */.
```

If after the installation of this SYSMOD any subsequent IBM service modifies this macro, the replacement or update from the IBM service is not selected. It is your responsibility to provide continued modifications for the macro. Thus, this method of updating an element should be used carefully.

In both examples, because the new macro exists in a TXLIB, the following DD statement is required during APPLY and ACCEPT:

```
//NEWSMP    DD DSN=TSO.NEWSMP.MACLIB,DISP=SHR
```

Example 2: Deleting a macro

Assume you have installed one of your application programs as function HUSR001. The function contains macro USRMAC01, which is no longer required. Here is an example of a SYSMOD containing a ++MAC statement to delete the macro from the target and distribution libraries:

```
++USERMOD(USR0001)        /* User modification     */.
++VER(Z038) FMID(HUSR001) /* for user application. */.
++MAC(USRMAC01)           /* Delete this macro.    */
         DELETE           /*                       */.
```

Example 3: Adding a new macro

Assume you have installed one of your application programs as function HUSR001. A new macro USRMAC02 is required, and it has these requirements:

- The macro is to be put into SYS1.USRMACS in the target zone.

- The macro is to be put into SYS1.AUSRMACS in the distribution zone.

- The macro is to be installed with an alias of TERMINAL.

- After being installed, modules USRASM01, USRASM02, USRSRC01, and USRSRC02 are to be assembled. USRASM01 and USRASM02 already exist as assembler entries in the CSI, and USRSRC01 and USRSRC02 exist as source entries in the CSI.

Here is an example of a SYSMOD containing a ++MAC statement defining all this information:

```
++USERMOD(USR0002)        /* User modification     */.
++VER(Z038) FMID(HUSR001) /* for user application. */.
++MAC(USRMAC02)           /* Add this macro        */
         DISTLIB(AUSRMACS) /* to this DLIB          */
         SYSLIB(USRMACS) /* and to this tgt lib    */
         MALIAS(TERMINAL)/* with this alias.       */
         ASSEM(USRASM01, /* Assemble these modules. */
             USRASM02, /*                        */
             USRSRC01, /*                        */
             USRSRC02) /*                        */
                       /* Inline text follows.   */.
...
... macro USRMAC02 goes here
...
```

Example 4: Packaging a renamed macro

Suppose that, for some reason, you need to rename macro USRMAC02, which was introduced in . The new name is to be USRMACXX. You do not need to change anything else about the macro. Here is an example of a SYSMOD containing the ++MAC statements needed to package this renamed macro:

```
++USERMOD(USR0003)        /* User modification     */.
++VER(Z038) FMID(HUSR001) /* to user application.  */
         PRE(USR0002) /* Base on previous USERMOD.*/.
```

```
++MAC(USRMAC02)            /* Delete the original    */
        DELETE             /* macro.                 */.
++MAC(USRMACXX)            /* Add the renamed macro  */
        DISTLIB(AUSRMACS) /* to this DLIB            */
        SYSLIB(USRMACS)  /* and to this tgt lib     */
        MALIAS(TERMINAL)/* with this alias.          */
        ASSEM(USRASM01, /* Assemble these modules.   */
              USRASM02, /*                           */
              USRSRC01, /*                           */
              USRSRC02) /*                           */
                        /* Inline text follows.    */.
...
... macro USRMACXX goes here
...
```

++MACUPD MCS

The ++MACUPD MCS describes a single macro update within a PTF, an APAR fix, or a USERMOD. It must immediately precede the macro update statements within the SYSMOD.

Syntax

++MACUPD MCS

Operands

ASSEM
specifies the names of modules to be assembled in addition to those named as GENASM subentries in the MAC entry. The source for the assemblies is the ASSEM entry, SRC entry, or DISTSRC member whose name matches the specified ASSEM value. SMP/E looks for a match—first with an ASSEM entry, then with a SRC entry, and finally with an entry in the DISTSRC data set—and uses the first match it finds. The source must either be known to SMP/E at the time the ASSEM operand is encountered on the ++MACUPD statement, or be included in the same SYSMOD.

Note:

1. APPLY and ACCEPT processing place the specified names into the SYSMOD entry created on the target zone and distribution zone.

2. If the object deck for the element specified on the ASSEM operand is also provided by the SYSMOD, the assembly may not occur. For more information, see the section on assemblies in the APPLY command chapter in *z/OS SMP/E Commands*.

DISTLIB
specifies the ddname of the distribution library for the specified macro.

Note:

1. DISTLIB must be specified if the macro has not been previously recorded on the target zone or distribution zone. If a MAC entry already exists in the target zone or distribution zone, and the value currently in that entry does not match that specified in the DISTLIB operand, the SYSMOD is not applied or accepted.
2. You cannot use SYSPUNCH as the DISTLIB. It is used by SMP/E and other products to process assembled modules.

DISTMOD
specifies the ddname of the link-edit distribution library for the modules specified in the ASSEM operand. During ACCEPT processing, the object code from the assembler is link-edited to the library specified.

DISTSRC
specifies the ddname of the library containing the additional assembly or source to be assembled. The additional assembly or source must be specified in the ASSEM operand.

MALIAS
specifies the alias names for the macro in both the target system and distribution libraries.

You can use MALIAS when two or more macros that must be defined in the same zone must have the same name for programming access. For example, you can use MALIAS if several products have a help macro whose name must match the name of the command processing module it describes. You can specify HELP on MALIAS and a unique element name as the macro name.

name
specifies the name of the macro member in the distribution library and, optionally, in the target system library. The name can contain any alphanumeric characters and $, #, @, or hex C0.

PREFIX
specifies the first characters (prefix) of the names of modules to be assembled in addition to those modules named as GENASM subentries in the target zone MAC entry. The prefix values must be 7 characters or less.

The full module names are determined by comparing the prefix with the target zone or distribution zone MOD entry names.

The source for the assembly is the ASSEM entry, SRC entry, or DISTSRC member whose name matches a MOD entry name beginning with one of the specified prefixes. SMP/E looks for a match— first with an ASSEM entry, then with a SRC entry, and finally with an entry in the DISTSRC data set— and uses the first match it finds. The source must be either known to SMP/E at the time the PREFIX operand is encountered on the ++MACUPD statement, or be included in the same SYSMOD.

Note: If the object deck for the element specified on the PREFIX operand is also provided by the SYSMOD, the assembly may not occur. For more information, see the section on assemblies in the APPLY command chapter in *z/OS SMP/E Commands*.

SYSLIB
specifies the ddname of the target library, if the macro exists in one. APPLY and RESTORE processing update this library.

Usage notes

- If a SYSMOD containing a ++MACUPD statement attempts to change the ownership (FMID) of the element (with the VERSION operand), the SYSMOD cannot be installed.
- The changes for the macro must immediately follow the ++MACUPD MCS and must not contain any records that start with ++.
- If the macro resides in a target library (rather than the SMPMTS), that target library should be included in the SYSLIB DD concatenation for assemblies during APPLY processing. For additional information about SYSLIB requirements, see *z/OS SMP/E User's Guide* and *z/OS SMP/E Commands*.
- The only IEBUPDTE control statements allowed in a SYSMOD are ./ CHANGE and ./ ENDUP.

- The only IEBUPDTE CHANGE operand SMP/E checks is NAME, which must specify the same element as the ++MACUPD MCS. Other CHANGE operands may produce undesired results and are used at your own risk. For example, if you code UPDATE=INPLACE, SMP/E may update the distribution library. Once the distribution libraries are changed, there is no way to remove the updates.
- SMP/E does not support a continuation of the ./ CHANGE statement.
- SMP/E generates any ./ ALIAS statements needed and places them in the IEBUPDTE input data following the last text statement. The ./ ALIAS control statements are generated only for macro updates.
- When processing multiple updates to the same lines in a given macro, SMP/E uses the ./ CHANGE statement from the last update to the lines.
- If an APAR fix or USERMOD updates a macro that causes an assembly, SMP/E sets the ASSEMBLY indicator in the MOD entry for the assembled module. This can cause a problem when additional service that does not know about the macro change is installed at the same time as the APAR or USERMOD— for example, if you are installing your own USERMOD and IBM-supplied PTFs with the same APPLY command. In such cases, because the ASSEMBLY indicator is set, the module is reassembled, but does not contain the macro changes from the APAR or USERMOD. To prevent these assemblies, you can reset the ASSEMBLY indicator using UCLIN after installing the APAR or USERMOD.

Examples

Assume you want to update macro IFBMAC02, which resides in distribution library IFBMACS. Because of this change, module IFBSRC01 must be reassembled. Module IFBSRC01 exists as a source in distribution library SYS1.IFBSRC, and as an object module in distribution library SYS1.AOS23. The macro and the modules are part of JXY1040, a user-written function. Here is an example of a SYSMOD containing a ++MACUPD statement to make the necessary changes:

```
++PTF(USR0001)              /* Preventive service       */.
++VER(Z038) FMID(JXY1040)   /* for user product.        */.
++MACUPD(IFBMAC02)          /* Update this macro        */
          DISTLIB(IFBMACS)/* in this DLIB.              */
          ASSEM(IFBSRC01) /* Assemble this source.      */
          DISTSRC(IFBSRC) /* Source is here.            */
          DISTMOD(AOS23)  /* Assembled SRC goes here. */.
./ CHANGE name=IFBMAC02
... IEBUPDTE control cards and data
...
```

In this example, DD statements are required when the SYSMOD is applied to define the target libraries for the macro and the load module to be updated as a result of the assembly. For example, if the modules in SYS1.AOS23 (the assembled module's distribution library) were copied to SYS1.LINKLIB and the source in SYS1.IFBSRC (the source element's distribution library) were copied to SYS1.CHGLIB, the following DD statements are needed:

```
//LINKLIB  DD  DSN=SYS1.LINKLIB…
//CHGLIB   DD  DSN=SYS1.CHGLIB…
```

In this example, the following DD statements are needed when the SYSMOD is being accepted to define the distribution libraries:

```
//IFBMACS  DD  DSN=SYS1.IFBMACS…  (macro DLIB)
//IFBSRC   DD  DSN=SYS1.IFBSRC…   (source DLIB for assembly)
//AOS23    DD  DSN=SYS1.AOS23…    (DLIB for module assembled)
```

++MOD MCS

The ++MOD MCS describes a single module replacement. It must immediately precede the module definition statements when they are within the SYSMOD. You should use the ++MOD MCS when you want to provide the object form of a module. If you want to provide the source form and have it assembled when the SYSMOD is installed, use the ++SRC MCS instead.

Syntax

The syntax to be used depends on the processing to be done for the element:

- Adding or replacing the element
- Deleting the element

Adding or replacing a module

++MOD MCS

```
►►─ ++MOD ─ ( ─ name ─ ) ─────────────────────────────────────►

                    ┌──────── , ◄────────┐
              └─ CSECT ─ ( ─┴─ csect ─┴─ ) ─┘

   ┌─ DALIAS ─ ( ─ alias ─ ) ─┐   ┌─ DISTLIB ─ ( ─ ddname ─ ) ─┐
   │                          │
   │          ┌───── , ◄─────┐│
   └─ TALIAS ─ ( ─┴─ alias ─┴─ ) ─┘

        ┌──────── , ◄────────┐
   └─ LEPARM ─ ( ─┴─ leparm ─┴─ ) ─┘

                    ┌─ LKLIB ─ ( ─ ddname ─ ) ─┐
                    ├─ RELFILE ─ ( ─ number ─ ) ─┤
                    ├─ TXLIB ─ ( ─ ddname ─ ) ─┤
   └─ FROMDS ─ ( ─ DSN ─ ( ─ dsn ─ ) ─ NUMBER ─ ( ─ n ─ )
        └─ VOL ─ ( ─ vol ─ ) ─┘  └─ UNIT ─ ( ─ unit ─ ) ─┘ )

        ┌──── , ◄────┐
   └─ LMOD ─ ( ─┴─ lmod ─┴─ ) ─┘   └─ RMID ─ ( ─ sysmod_id ─ ) ─┘

        ┌──────── , ◄────────┐
   └─ UMID ─ ( ─┴─ sysmod_id ─┴─ ) ─┘

        ┌──────── , ◄────────┐
   └─ VERSION ─ ( ─┴─ sysmod_id ─┴─ ) ─┘ ─ . ─►◄
```

Deleting a module

++MOD MCS

```
►►─ ++MOD ─ ( ─ name ─ ) ─ DELETE ─────────────────────────────►

                    ┌──── , ◄────┐
              └─ CSECT ─ ( ─┴─ csect ─┴─ ) ─┘

   ┌─ DISTLIB ─ ( ─ ddname ─ ) ─┐
                    ┌──────── , ◄────────┐
              └─ VERSION ─ ( ─┴─ sysmod_id ─┴─ ) ─┘

►─ . ─►◄
```
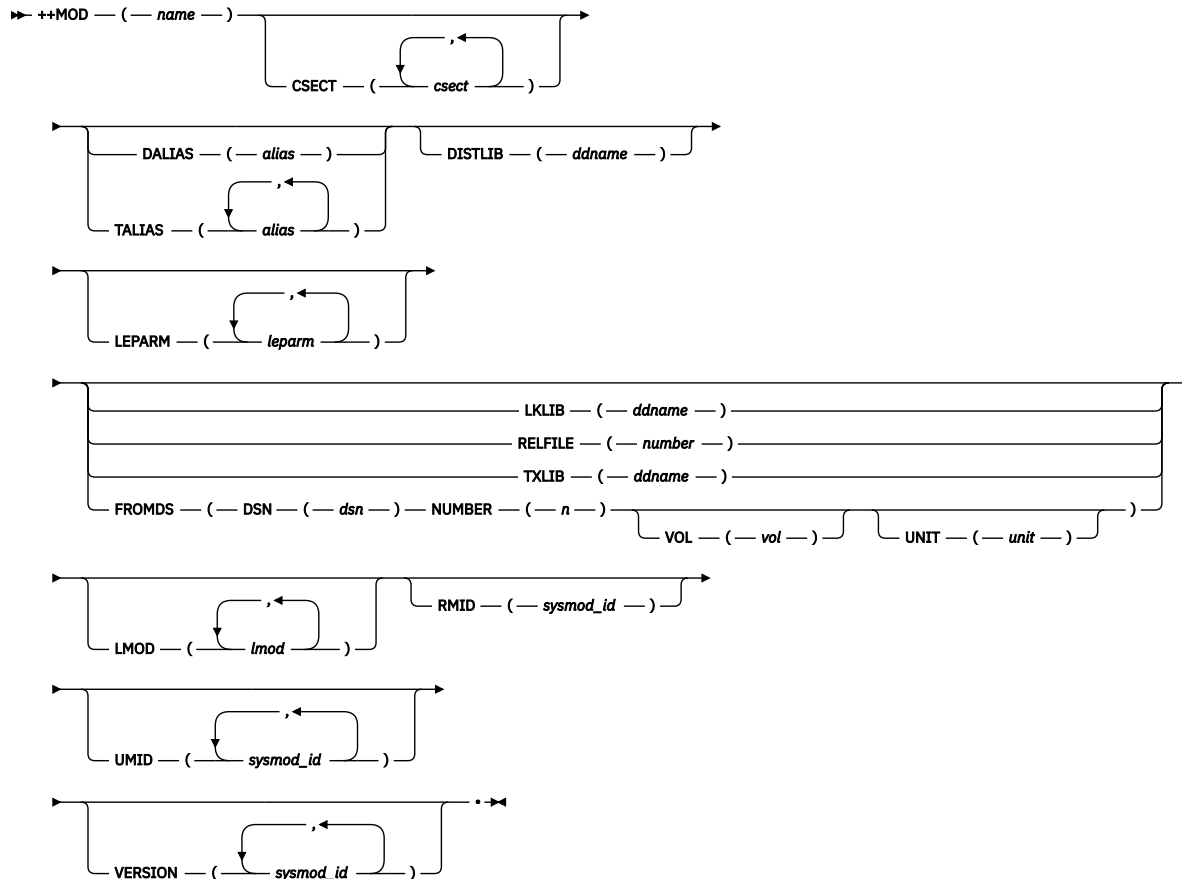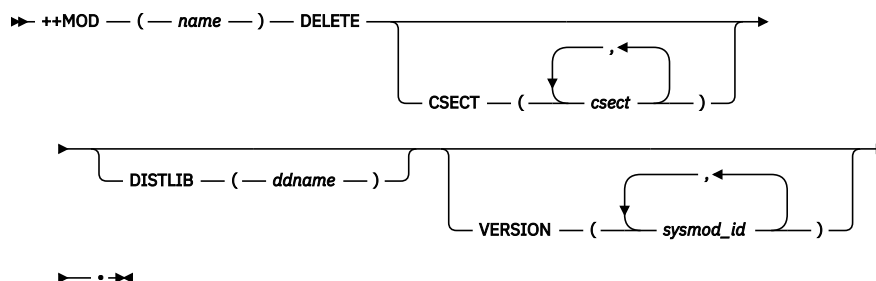
Operands

CSECT
> lists all the CSECTs contained in the module. This operand is required if the module contains more than one CSECT or if the CSECT name is different from the module name on the ++MOD MCS.

- If no CSECT name is specified, SMP/E assumes that the module contains only one CSECT, whose name matches the module name on the ++MOD MCS.
- If CSECT is specified, it must include all the CSECTs contained in the module, even if one of them has the same name as the module.

Note:

1. A CSECT name can contain from 1 to 8 characters. The name can contain any characters except the following:

   - Comma ,
   - Left parenthesis (
   - Right parenthesis)
   - Blank

2. Comments are not allowed within a CSECT name. For example, the following is not allowed:

   ```
   CSECT ( /* this is a csect name */ CSECT01)
   ```

   The comment is interpreted as part of the CSECT name, instead of a comment.

3. Even if CSECT is not specified on the ++MOD MCS used to create a MOD entry, CSECT information is saved if CSECT is specified on subsequent ++MOD statements that update the MOD entry.

DALIAS
is the alias name of a module that has an alias in the distribution library, but not in the target library. This might be used if the module is included under its alias name during system generation.

Note: DALIAS is mutually exclusive with TALIAS.

DELETE
indicates that the module and alias names are to be removed from the target libraries, distribution libraries, and SMP/E data sets.

In order to DELETE a module, the FMID on the SYSMOD must match the FMID of the module. If the module has no FMID associated with it, then the SYSMOD will delete the module. If the FMID of the module does not match the FMID on the SYSMOD, then the SYSMOD will not delete the module, because the specified FMID does not own the module.

Note: DELETE is mutually exclusive with all other operands except CSECT, DISTLIB, and VERSION.

DISTLIB
specifies the ddname of the distribution library for the specified module.

Note:

1. This operand must be specified if the module has not been previously recorded on the target zone or distribution zone. If a MOD entry already exists in the target zone or distribution zone and the value currently in that entry does not match that specified in the DISTLIB operand, the SYSMOD is not applied or accepted, unless that SYSMOD also used the ++MOVE MCS to change the DISTLIB to that new value.

2. You cannot use SYSPUNCH as the DISTLIB. It is used by SMP/E and other products to process assembled modules.

FROMDS
identifies the partitioned data set that contains this element.

Note: The FROMDS operand and its DSN, NUMBER, VOL, and UNIT suboperands are included in the MCS generated by the BUILDMCS command. IBM does not intend the FROMDS operand to be used in manually coded MCS.

DSN
specifies the dsname of the FROMDS data set. The specified data set name must conform to standard data set naming conventions and cannot contain parentheses. The maximum length of the entire name is 44 characters (including the periods).

NUMBER
specifies a number that SMP/E is to use when assigning a name to the SMPTLIB data set associated with this FROMDS data set. (This is similar to the way the relative file number is used in RELFILE processing.)

VOL
specifies, for an uncataloged data set, the volume serial number of the volume containing the FROMDS data set. If specified, this volume identifier must be from 1 to 6 alphanumeric characters.

VOL may be omitted for a cataloged data set.

UNIT
specifies, for an uncataloged data set, the UNIT type containing the FROMDS data set. If specified, the UNIT value must be from 1 to 8 characters and must conform to standard UNIT naming conventions. IBM SMP/E for z/OS, V3R6 accepts any nonblank characters specified between the open and close parentheses, up to a maximum length of 8.

UNIT may be omitted for a cataloged data set.

Note: FROMDS is mutually exclusive with DELETE, LKLIB, RELFILE, and TXLIB.

LEPARM
specifies link-edit utility attributes for the module. Any of the following values can be specified:
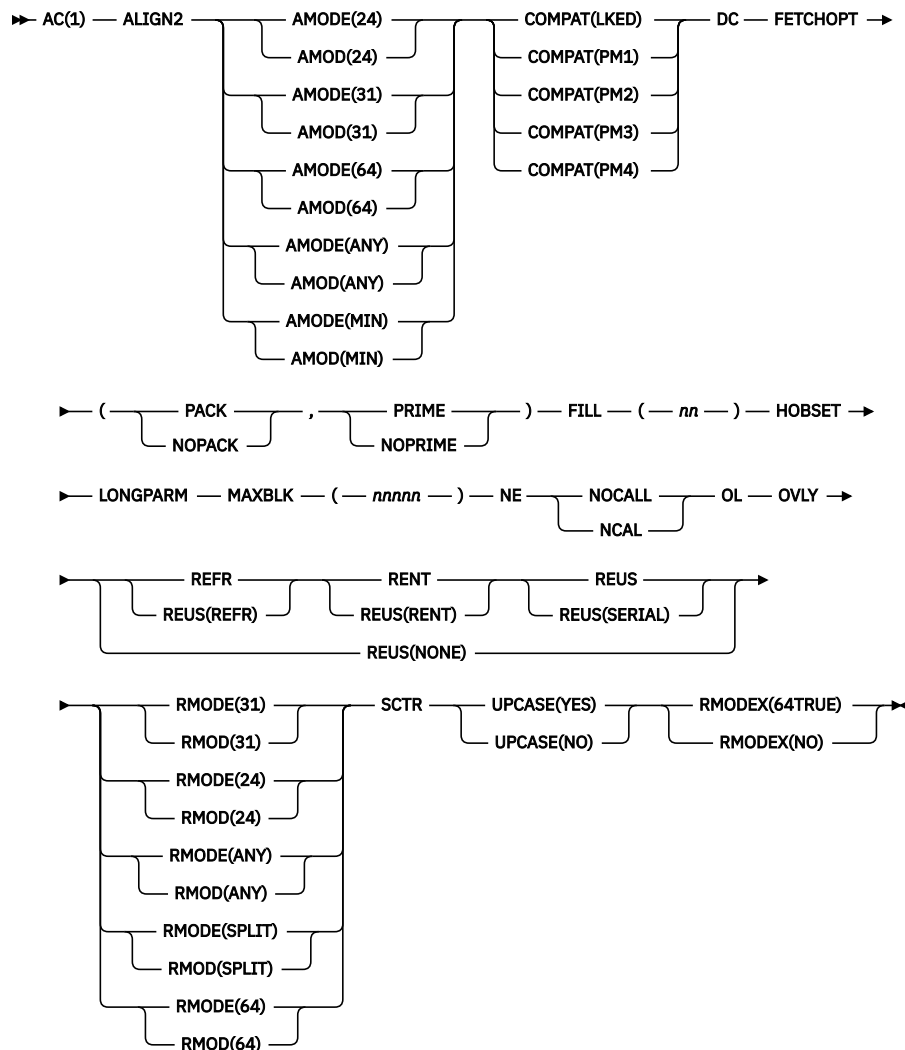
Note:

1. The LEPARM values from the ++MOD MCS are associated with a load module entry only if the module was copied, not link-edited, into the target libraries. (The COPY indicator is set in the load module entry.) If the load module was link-edited, JCLIN must be used to change its link-edit utility attributes.

2. During APPLY processing of a ++MOD MCS with LEPARMs, the LEPARM options are saved not in the MOD entry created, but in the LMOD entry. During ACCEPT processing, the MOD entry is created with the LEPARMs present. The target zone MOD entry can contain the LEPARM options through either UCLIN or the copying of the distribution zone to a target zone.

3. All LEPARM attributes may also be specified in the format '*attribute=value*'. For example, FILL(*nn*) may also be specified as FILL=*nn*.

4. The previously listed link-edit attributes are the only attributes that can be specified on the LEPARM operand. If any other attributes are specified, a syntax error will result during RECEIVE processing.

5. The LEPARM values of DCBS, LET, LIST, XCAL, and XREF are recognized by SMP/E, but are not saved. Specifying them on the ++MOD MCS does not cause them to be passed to the link-edit utility.

6. RMODE(31) is a synonym for RMODE(ANY).

For more information about how the LEPARM operand is processed, see the chapters on ACCEPT and APPLY processing in *z/OS SMP/E Commands*. These attributes are described in full in "MOD entry (distribution and target zone)" on page 226.

LKLIB
is the ddname of the partitioned data set containing the link-edited format of the object module. This operand is required if the module is provided in a data set, rather than inline or on a tape.

Note: LKLIB is mutually exclusive with FROMDS, RELFILE, and TXLIB.

LMOD
lists existing load modules that are to contain the module. If any of the names specified are not already LMOD subentries in the target zone MOD entry, they are added during APPLY processing.

Note:

1. LMOD can be used only to add a module to an existing load module.

2. LMOD cannot be used to create a new load module. Nor can it be used if any link-edit control statements must be added or changed to add the module to an existing load module. However, you can use JCLIN data to create a new load module and to add or change link-edit control statements.

3. If an LMOD entry does not exist for one of the load modules specified, sufficient information is not available to create one. Thus, when the MOD is to be link-edited during APPLY processing, an error message is issued, and no link-edit is performed for that load module.

name
specifies the name of the module in the distribution library and, optionally, in the target library. The name can contain any alphanumeric characters and $, #, @, or hex C0.

RELFILE
identifies which relative file associated with the SYSMOD contains this module.This operand is required if you provide the element in RELFILE format, rather than inline or in a LKLIB or TXLIB data set.

Note:

1. RELFILE is mutually exclusive with FROMDS, LKLIB, and TXLIB.

2. If an object module is provided in RELFILE format, it must be in link-edited format.

RMID
specifies the last SYSMOD that **replaced** this module. This operand can be used only in a service-updated function, and the specified PTF must be integrated into the function.

TALIAS

specifies one or more alias names for the module. The aliases exist in the distribution library and the target library.

You can use TALIAS for a module that was copied from a distribution library into a target library (defined by JCLIN data as a copied module), but not for one that is link-edited (defined by JCLIN data as a link-edited module).

Note: TALIAS must be specified on the ++MOD MCS even if ALIAS was specified on the COPY SELECT statement. TALIAS is required for the RECEIVE command to properly handle aliases in RELFILEs. During RECEIVE processing, SMP/E copies RELFILEs to SMPTLIB data sets. All element members and their associated aliases are copied. Because JCLIN is not processed during RECEIVE, the aliases must be identified on the MCS to get RECEIVE to copy the aliases.

Likewise, to specify an alias for a copied load module, you must use the TALIAS operand on the ++MOD statement for that load module.(To specify an alias for a link-edited load module, do not use TALIAS. You must identify that alias using an ALIAS link-edit control statement in the JCLIN that defined the load module. For more information, see *z/OS SMP/E Commands*.)

Note: TALIAS is mutually exclusive with DALIAS.

TXLIB

is the ddname of the partitioned data set containing an object module that has not been link-edited. This operand is required if the module is provided in a TXLIB data set rather than inline, in a FROMDS data set, in a LKLIB data set, or in RELFILE format.

Note:

1. SMPTLIB cannot be used as a value on the TXLIB operand.

2. TXLIB is mutually exclusive with FROMDS, LKLIB, and RELFILE.

UMID

specifies the SYSMODs that have **updated** this module since it was last replaced. This operand can be used only in a service-updated function, and the specified PTFs must be integrated into the function.

VERSION

specifies one or more function SYSMODs that currently contain the element. The function containing the ++MOD MCS takes over ownership of the element from the specified functions.

When VERSION is specified on an element statement, it overrides any VERSION operand values specified on the ++VER MCS.

Usage notes

- If the module is packaged inline, it must immediately follow the ++MOD MCS and must not contain any records that start with ++. Neither FROMDS, LKLIB, RELFILE, nor TXLIB may be specified.

- If the module is packaged in a TXLIB data set, the ddname specified on the TXLIB operand is required during APPLY and ACCEPT processing.

- If the module is in an LKLIB data set, the ddname specified in the LKLIB operand is required during APPLY and ACCEPT processing. Module replacements in an LKLIB data set must be in load module format (that is, processed by the link-edit utility).

- For information about packaging SYSMODs in RELFILE, TXLIB, or inline format, see Standard Packaging Rules for z/OS-based Products (publibz.boulder.ibm.com/epubs/pdf/gimpkg80.pdf).

- There are several ways to associate a module with a load module:

  - The DISTLIB library can be totally copied into the target library.

  - JCLIN can identify the module as part of one or more load modules.

  - The LMOD operand on the ++MOD MCS can indicate the associated load module.

  If SMP/E cannot identify the load module associated with a given module, it does not update the target libraries during APPLY processing. Instead, it issues warning message GIM43401W.

Examples

The following examples are provided to help you use the ++MOD MCS:

Example 1: Adding a new module to an existing load module

Module IFBMOD01 is a new module that is to be placed in the distribution library SYS1.AOSFB and is to be link-edited with the existing load module IEEFRQ in the target system library SYS1.LINKLIB. Module IFBMOD01 contains two CSECTs, IFBCST01 and IFBCST02. Here is an example of a SYSMOD containing a ++MOD statement to do this:

```
++USERMOD(USR0001)         /* User modification      */.
++VER(Z038) FMID(JXY1040)  /* to user application.   */.
++MOD(IFBMOD01)            /* Add this module        */
          DISTLIB(AOSFB)  /* to this DLIB at ACCEPT, */
          LMOD(IEEFRQ)    /* to this LMOD at apply.  */
          CSECT(IFBCST01  /* Module has two CSECTs.  */
                IFBCST02) /*                         */.
...
... object deck for IFBMOD01
...
```

The following DD statement is needed at APPLY time to define the operating system load module library:

```
//LINKLIB  DD  DSN=SYS1.LINKLIB,DISP=OLD
```

When the SYSMOD is accepted, the following DD statement is needed to define the distribution library for this module:

```
//AOSFB  DD  DSN=SYS1.AOSFB,DISP=OLD
```

Example 2: Specifying link-edit utility attributes with LEPARM

For this example, assume you have installed a product, FXY1040, packaged in RELFILE format. The package contained a module, IFBMOD01, that was identified by inline JCLIN as being installed as follows:

1. It was copied from the distribution library, AOS12, to the target system library, LPALIB.

2. It was linked with several other modules to form a load module, IFBLMDXX, in LINKLIB.

3. It was linked, by itself, to form a load module, named IFBLMDX1, in LINKLIB.

The JCLIN in the initial function was as follows:

```
//JOB      JOB 'accounting info',MSGLEVEL=(1,1)
//COPY1    EXEC PGM=IEBCOPY
//AOS12    DD DSN=SYS1.AOS12,DISP=SHR
//LPALIB   DD DSN=SYS1.LPALIB,DISP=SHR
//SYSIN    DD *
 COPY INDD=AOS12,OUTDD=LPALIB
 SELECT M=(IFBMOD01)
/*
//LINK1    EXEC PGM=IEWL,PARM='REUS'
//AOS12    DD DSN=SYS1.AOS12,DISP=SHR
//SYSLMOD  DD DSN=SYS1.LINKLIB,DISP=SHR
//SYSLIN   DD *
  INCLUDE AOS12(IFBMOD01)
  INCLUDE AOS12(IFBMOD0A,IFBMOD0B,IFBMOD0C)
  ENTRY    IFBMOD01
  NAME IFBLMDXX(R)
/*
//LINK2    EXEC PGM=IEWL
//AOS12    DD DSN=SYS1.AOS12,DISP=SHR
//SYSLMOD  DD DSN=SYS1.LINKLIB,DISP=SHR
//SYSLIN   DD *
  INCLUDE AOS12(IFBMOD01)
  NAME IFBLMDX1(R)
/*
```

The target zone now contains the following entries:

1. MOD entry for IFBMOD01, having LMOD subentries of IFBMOD01 (from the copy step), IFBLMDXX (from the first link step), and IFBLMDX1 (from the second link step).

2. LMOD entry for IFBMOD01, (created from the SELECT statement of the copy step), indicating that the load module was copied during installation. The LMOD entry does not have any link-edit utility parameters yet, because SMP/E copied IFBMOD01 from the RELFILE data sets, and, thus, had no need to obtain the link parameters.

3. LMOD entry for IFBLMDXX (created from the first link step), with link-edit attributes of REUS.

4. LMOD entry for IFBLMDX1 (created from the second link step), with link-edit attributes of STD (because no special parameters were specified).

Now assume a PTF to replace module IFBMOD01 is required. Module IFBMOD01 has link-edit utility attributes of REUS and RENT. Here is an example of a SYSMOD containing a ++MOD statement to do this:

```
++PTF(UZ12345)            /* PTF                    */.
++VER(Z038) FMID(FXY1040) /* for this function.     */.
++MOD(IFBMOD01)           /* Replace IFBMOD01       */
        DISTLIB(AOS12)    /* in this DLIB.          */
        LEPARM(RENT       /* Reentrant              */
               REUS)      /* and reusable.          */.
...
... object deck for IFBMOD01
...
```

When the PTF is applied, SMP/E processes the LEPARM as follows:

1. Because LMOD IFBMOD01 was copied from the DLIB module, it is updated using the LEPARM values from the ++MOD MCS.

2. LMOD IFBLMDXX link-edit attributes remain as they are. This is because the link-edit utility attributes of each module within the load module have no bearing on the link-edit utility attributes of the load module; and, in this case, they are not the same. One of the other modules in IFBLMDXX must have a more restrictive set of link-edit attributes, thus forcing the load module to have that restrictive set of attributes.

3. LMOD IFBLMDX1 link-edit attributes remain as they are, for the same reason that load module IFBLMDXX attributes did not change. This is true even though load module IFBLMDX1 is composed of only the one DLIB module. As long as the load module was identified by a link-edit step, SMP/E assumes that the load module may contain multiple DLIB modules.

   Note: The only way to change the link-edit attributes of a load module that was link-edited during initial installation is to provide JCLIN input to identify the new link-edit attributes.

Example 3: Packaging a renamed module

Suppose that, for some reason, you need to rename module IFBMOD01, which was introduced in "Example 1: Adding a new module to an existing load module" on page 66. The new name is to be IFBMODXX. You do not need to change anything else about the module. Here is an example of a SYSMOD containing the ++MOD statements needed to package this renamed module:

```
++USERMOD(USR0002)        /* User modification      */.
++VER(Z038) FMID(JXY1040) /* to user application.   */
         PRE(USR0001)     /* Base on previous USERMOD.*/.
++MOD(IFBMOD01)           /* Delete the original    */
        DELETE            /* module.                */.
++MOD(IFBMODXX)           /* Add the renamed module */
        DISTLIB(AOSFB)    /* to this DLIB at ACCEPT, */
        LMOD(IEEFRQ)      /* to this LMOD at apply.  */
        CSECT(IFBCST01    /* Module has two CSECTs.  */
              IFBCST02) /* */.
...
... object deck for IFBMODXX
...
```

Note: When packaging a renamed module, you must ensure that this module is included in the appropriate load modules, either through JCLIN or the LMOD operand on the ++MOD statement, as in this example.

The following DD statement is needed at APPLY time to define the operating system load module library:

```
//LINKLIB  DD  DSN=SYS1.LINKLIB,DISP=OLD
```

When the SYSMOD is accepted, the following DD statement is needed to define the distribution library for this module:

```
//AOSFB  DD  DSN=SYS1.AOSFB,DISP=OLD
```

Example 4: Deleting a module

If you need to delete a module, here is an example of a SYSMOD containing the ++MOD statement:

```
++USERMOD(USR0004)         /* User modification      */.
++VER(Z038)  FMID(JXY1040) /* to user application.   */.
++MOD(IFBMOD0Z)            /* Delete the module      */
     DELETE                /* from DLIB AOSFB - entry */
     DISTLIB(AOSFB)        /* for DLIB must be AOSFB. */.
```

Note: When packaging a SYSMOD to delete a module, no object deck is required. The previous example is a complete SYSMOD, and does not require any other SMP/E elements in the SYSMOD.

For more information about the processing done when a module is deleted, see the section on deleting elements in *z/OS SMP/E Commands*.

++MOVE MCS

The ++MOVE MCS moves a macro, a module, a source, or a load module (and any known aliases) from one library to another. The associated target or distribution zone is automatically updated to show that the entry has been deleted from the old library and added to the new one.

Note: You cannot use ++MOVE to move a load module from one path to another path in a UNIX file system if symbolic links have been defined for that load module.

Syntax

The syntax to be used depends on the type of library involved:

- Moving to a different distribution library (DISTLIB)
- Moving to a different target library (SYSLIB)

Moving to another DISTLIB

++MOVE MCS

▶▶— ++MOVE —(— *name* —)— DISTLIB —(— *ddname* —)— TODISTLIB —(— *ddname* —)—▶

```
    ┌─ MAC ─┐                          .─◄┤
    ├─ MOD ─┤
    └─ SRC ─┘  └─ FMID —(─┬─ sysmod_id ─┬─)─┘
                          └────◄────────┘
```

Moving to another SYSLIB

++MOVE MCS

▶▶— ++MOVE —(— *name* —)— SYSLIB —(— *ddname* —)— TOSYSLIB —(— *ddname* —)—▶

```
    ┌─ MAC ─┐                          .─◄┤
    ├─ SRC ─┤
    │        └─ FMID —(─┬─ sysmod_id ─┬─)─┘
    └──────── LMOD ───────────┘
```

Operands

DISTLIB
 specifies the ddname of the distribution library in which the member resides.

FMID
 specifies the FMID that owns the element. This is used when the current owner of the element is different from the FMID specified on the ++VER MCS. Up to 10 SYSMOD IDs can be specified.

 Note: FMID is mutually exclusive with LMOD.

MAC, MOD, SRC, or LMOD
 specifies the type of member to be moved.

 - If DISTLIB and TODISTLIB are specified, MAC, MOD, or SRC is valid.
 - If SYSLIB and TOSYSLIB are specified, MAC, SRC, or LMOD is valid.

 Note: LMOD is mutually exclusive with FMID.

name
 specifies the name of the element or load module to be moved.

SYSLIB
 specifies the ddname of the target library in which the member resides.

TODISTLIB
 specifies the ddname of the distribution library to which the member is to be moved. This must be specified if DISTLIB is specified.

TOSYSLIB
 specifies the ddname of the target library to which the member is to be moved. This must be specified if SYSLIB is specified.

Usage notes

- The member and library operands are required. You must specify the member name, the set of libraries affected by the move, and the member type.
- ++MOVE statements must follow any ++VER and ++IF statements and must precede any ++JCLIN or element MCSs.
- Regardless of the order in which ++MOVE, ++RENAME, and ++DELETE statements are coded in a SYSMOD, they are processed in the following order:
  - APPLY and ACCEPT
    1. ++MOVE
    2. ++RENAME
    3. ++DELETE
  - RESTORE
    1. ++RENAME
    2. ++MOVE

  Afterwards, ++JCLIN statements are processed, and then element statements.
- You must use the FMID operand to identify all possible owners of an element to be moved.

  Using the FMID operand does not imply that the owner of the element is being changed. You can change ownership only by specifying the VERSION operand on the ++VER MCS or on the element statement. Even if you are changing the owner of the element, you must specify the FMID operand, because the element is moved before its ownership is changed. For more information about the VERSION operand, see "++VER MCS" on page 92.

- A ++MOVE MCS can move a member to a given library from one library at a time. If a member exists in more than one library, you must use additional ++MOVE or ++DELETE statements to process the additional copies. For more information about ++DELETE statements, see "++DELETE MCS" on page 15.
- If a ++MOVE MCS changes the DLIB for an element that is also being replaced, the element MCS must specify the new DLIB.

Examples

Assume IBM ships you a PTF that moves module MODAA from its current distribution library to a new library, and also moves load module LMODA from its current target library to a new one. Here is an example of a SYSMOD containing ++MOVE statements that makes these changes:

```
++PTF(UR01234)            /* Identify the PTF number  */.
++VER(Z038) FMID(HXY1300)   /* for MVS function HXY1300.*/.
++IF (ESY1300) THEN        /* If ESY1300 is installed  */
  REQ(UR12399)             /* UR12399 is required.     */
++MOVE (MODAA)             /* Move module MODAA        */
  DISTLIB(AOS11)           /* from DLIB AOS11          */
  TODISTLIB(AOSXX) MOD     /* to DLIB AOSXX.           */.
++MOVE (LMODA)             /* Move load module LMODA   */
  SYSLIB(LINKLIB)          /* from LINKLIB             */
  TOSYSLIB(LPALIB) LMOD    /* to LPALIB.               */.
              .
              .
              .
++MOD(MODAA) DISTLIB(AOSXX) /* Element MCS statements.  */.
```

++NULL MCS

The ++NULL MCS is valid only in the SMPHOLD data set. It provides no SMP/E function other than allowing all service tapes to be built with the same format, even though a function may have no exception SYSMOD data for a particular month.

Syntax

++NULL MCS

►►— ++NULL — • —►◄

++PRODUCT MCS

The ++PRODUCT MCS is used to describe information about a product. It introduces descriptive information about a new or replacement product into the global zone.

Syntax

++PRODUCT MCS

►►— ++PRODUCT — (— *prodid* — , — *vv.rr.mm* —) — DESCRIPTION — (— *description* —) —►
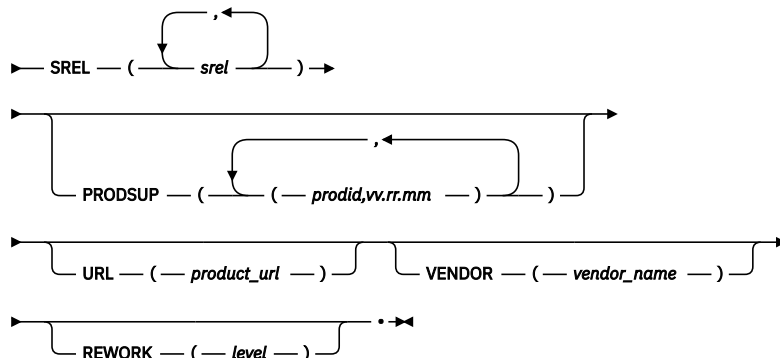
Operands

The PRODUCT is identified by the combination of the *prodid* and the *vv.rr.mm* values.

prodid
> is a 1- to 8-character product identifier. It can contain uppercase alphabetic, numeric, and national ($, #, @) characters. It may also contain one or more dashes (-). For IBM products, the *prodid* is assumed to be the IBM program product number (5647-A01, for example).

vv.rr.mm
> specifies the version, release, and modification level of this PRODUCT. It is 5 to 8 characters long in the form *vv.rr.mm*. The version (*vv*), release (*rr*), and modification (*mm*) values must be one or two numeric characters separated by a period ('.'). SMP/E will insert leading zeros to each section of the *vv.rr.mm* value that is one character long. That is, 2.5.0 will become 02.05.00 when stored.

DESCRIPTION
> a text description of the product.
>
> - DESCRIPTION can also be specified as DESC.
> - The DESCRIPTION value can be in single-byte characters (such as English alphanumeric) or double-byte characters (such as Kanji).
> - The DESCRIPTION value can contain up to 64 bytes of data, including blanks. (For double-byte data, the 64-byte maximum includes all shift-in and shift-out characters, as well as the double-byte characters.) Extra blanks as well as leading and trailing blanks are deleted.
> - The DESCRIPTION value can span multiple 80-byte records. Data must continue up to and including column 72 and begin in column 1 of the next line. All data past column 72 is ignored. The break does not translate to a blank unless a blank is explicitly coded in column 72 of the first line or in column 1 of the second line.
> - The DESCRIPTION value cannot be only blanks.
> - If parentheses are included in the text, they must be in matched pairs.

SREL
> specifies the system or subsystem releases on which the PRODUCT can be installed. Each SREL value must be four alphanumeric characters, usually one alphabetic character followed by three numeric characters. These are the systems and subsystems defined by IBM, with their SRELs:

> **System**
> > **SREL**
>
> **DB2**
> > P115
>
> **CICS®**
> > C150
>
> **IMS**
> > P115
>
> **MVS**
> > Z038
>
> **NCP**
> > P004

> The list of SRELs is used during RECEIVE processing to determine whether a PRODUCT should be received.

PRODSUP
> indicates which PRODUCTs are superseded (replaced) by this PRODUCT. It is a list of *prodid,vv.rr.mm* values for the PRODUCTs being superseded. The combination of *prodid* and *vv.rr.mm* determines the uniqueness of an entry in the PRODSUP operand list.

> The PRODSUP operand must **not** specify a *prodid,vv.rr.mm* combination matching the *prodid,vv.rr.mm* combination of the ++PRODUCT statement on which it is specified.

URL

specifies a uniform resource locator (URL) that can be used to obtain additional information about this product.

- The URL can be in single-byte characters (such as English alphanumeric) or double-byte characters (such as Kanji).

- The URL can contain up to 256 bytes of data, excluding blanks. (For double-byte data, the 256-byte maximum includes all shift-in and shift-out characters, as well as the double-byte characters.) All blanks including leading and trailing blanks are deleted.

- The URL can span multiple 80-byte records. Data must continue up to and including column 72 and begin in column 1 of the next line. All data past column 72 is ignored, including blanks. The break does not translate to a blank.

- The URL cannot be only blanks.

- If parentheses are included in the text, they must be in matched pairs.

VENDOR

specifies the name of the vendor supplying the product.

- The VENDOR can be in single-byte characters (such as English alphanumeric) or double-byte characters (such as Kanji).

- The VENDOR can contain up to 64 bytes of data, including blanks. (For double-byte data, the 64-byte maximum includes all shift-in and shift-out characters, as well as the double-byte characters.) Extra blanks as well as leading and trailing blanks are deleted.

- The VENDOR can span multiple 80-byte records. Data must continue up to and including column 72 and begin in column 1 of the next line. All data past column 72 is ignored. The break does not translate to a blank unless a blank is explicitly coded in column 72 of the first line or in column 1 of the second line.

- The VENDOR cannot be only blanks.

- If parentheses are included in the text, they must be in matched pairs.

REWORK

is the level of this ++PRODUCT MCS, which was reworked for minor changes. Up to eight numeric characters can be specified.

REWORK is generally used only for ++PRODUCT statements supplied by IBM that have been reworked for minor changes. For these ++PRODUCT statements, the REWORK level is *yyyyddd*, which is the year followed by the Julian date (for example, 2008110).

REWORK allows an updated ++PRODUCT MCS to be automatically received again, as long as it is more recent than the version that has already been received. This takes the place of rejecting the ++PRODUCT MCS and receiving it again.

Note: If a ++PRODUCT MCS appears more than once in the SMPPTFIN data set, the first occurrence may be received. However, none of the subsequent versions of the ++PRODUCT MCS are received, even if their rework level is higher than the one for the first version of the ++PRODUCT MCS.

Usage notes

- The ++PRODUCT statements are processed whenever the SMPPTFIN data set is processed. This is true whether only selected SYSMODs are being processed or the entire SMPPTFIN data set is being processed.

    ++PRODUCT statements are not processed when only the SMPHOLD data set is being processed.

- The *prodid*, *vv.rr.mm*, DESCRIPTION, and SREL values are required and cannot be blank or null.

- The ++PRODUCT MCS is used in the SMPPTFIN data set and can be placed between, before, or after SYSMODs, ++FEATURE MCS, or ++PRODUCT MCS. It must be followed by one of the following: a ++APAR, ++ASSIGN MCS, ++FEATURE MCS, ++FUNCTION, ++PTF, ++USERMOD, another ++PRODUCT

MCS, or an end-of-file. If one of these does not follow, SMP/E does not receive the SYSMOD being processed and it skips the ++PRODUCT MCS.

Example

Here is an example of a ++PRODUCT MCS for a product called OS/390.

```
++PRODUCT(5647-A01,2.5.0)        /* Product definition */
  DESCRIPTION(OS/390)               /* Description  */
  URL(http://www.S390.ibm.com/os390/) /* URL        */
  SREL(Z038)                        /* SREL value   */
  PRODSUP((5645-001,01.04.00))      /* Product sups */
  VENDOR(IBM)                       /* Product Vendor */.
```

++PROGRAM MCS

The ++PROGRAM MCS describes a program element (a pre-built load module or a program object). It must immediately precede the load module or program object when they are within the SYSMOD. Use the ++PROGRAM when you want to ship executables as program parts. If you want to provide the object form of the module, use the ++MOD MCS instead. JCLIN is not used to define the program element.

Syntax

The syntax to be used depends on the processing to be performed for the element:

- Adding or replacing the element
- Deleting the element

Adding or replacing a program element

Program element MCS

Deleting a program element

Program element MCS

Operands

ALIAS
specifies an alternate name for the program object or load module.

DELETE
specifies that the program element and all aliases are to be removed from the target library and the distribution library.

Note:

1. DELETE is mutually exclusive with all other operands except DISTLIB and VERSION.
2. If the element statement is in a base function, you may want to use the DELETE operand on the ++VER MCS to delete the previous release, rather than on the element statement to delete a specific element.
3. Specification of the DELETE operand results in all aliases of the program element being deleted along with the program element identified.

DISTLIB
specifies the ddname of the distribution library for the specified program element. During ACCEPT processing, SMP/E installs the program element into the distribution library as a member. (The distribution library must be a PDS or PDSE; it cannot be part of a UNIX file system.)

Note:

1. **DISTLIB** must be specified when the program element is first installed.
2. If a program element entry already exists in the target zone or distribution zone and the value currently in that entry does not match that specified in the DISTLIB operand, the SYSMOD is not applied or accepted.

FROMDS
identifies the partitioned data set that contains this element.

Note: The FROMDS operand and its DSN, NUMBER, VOL, and UNIT suboperands are included in the MCS generated by the BUILDMCS command. IBM does not intend the FROMDS operand to be used in manually coded MCS.

DSN
specifies the dsname of the FROMDS data set. The specified data set name must conform to standard data set naming conventions and cannot contain parentheses. The maximum length of the entire name is 44 characters (including the periods).

NUMBER
specifies a number that SMP/E is to use when assigning a name to the SMPTLIB data set associated with this FROMDS data set. (This is similar to the way the relative file number is used in RELFILE processing.)

VOL
specifies, for an uncataloged data set, the volume serial number of the volume containing the FROMDS data set. If specified, this volume identifier must be from 1 to 6 alphanumeric characters.

VOL may be omitted for a cataloged data set.

UNIT
specifies, for an uncataloged data set, the UNIT type containing the FROMDS data set. If specified, the UNIT value must be from 1 to 8 characters and must conform to standard UNIT naming conventions. SMP/E accepts any nonblank characters specified between the open and close parentheses, up to a maximum length of 8.

UNIT may be omitted for a cataloged data set.

Note: FROMDS is mutually exclusive with DELETE, RELFILE, and LKLIB.

LKLIB

is the ddname of the partitioned data set containing the program element. This operand is required if the program element is provided in a data set to which the users have access, rather than inline or in RELFILE format.

Note:

1. SMPTLIB cannot be used as a value on the LKLIB operand.
2. LKLIB is mutually exclusive with FROMDS and RELFILE.

name

specifies the name of the program element member. The name can contain any uppercase alphabetic, numeric, or national ($, #, @) character and can be 1 to 8 characters long.

RELFILE

identifies which relative file associated with the SYSMOD contains this element. This operand is required if you provide the element in RELFILE format, rather than inline or in a LKLIB data set.

Note:

1. The RELFILE value must be a decimal number from 1 to 9999.
2. RELFILE is mutually exclusive with FROMDS and LKLIB.

RMID

specifies the last SYSMOD that **replaced** this program element. This operand can be used only in a service-updated function, and the specified PTF must be integrated into the function.

SYSLIB

specifies the ddname of the target library for the program element. (The target library must be a PDS or PDSE; it cannot be part of a UNIX file system.)

During APPLY processing, the copy utility installs the program element into the target library. During RESTORE processing, the copy utility copies the program element from the distribution library member into the target library.

Note: SYSLIB must be specified when the program element is first installed.

VERSION

specifies one or more function SYSMODs that currently contain the element. The function containing the element MCS takes over ownership of the element from the specified functions.

When **VERSION** is specified on an element statement, it overrides any VERSION operand values that might be specified on the ++VER MCS.

Usage notes

- The target and distribution libraries for a program element must be a PDS (for pre-built load module) or a PDSE (for a program object).
- If the program element is packaged inline, it must immediately follow the ++PROGRAM MCS and must not contain any records starting with ++. Neither **FROMDS**, nor **RELFILE**, nor **LKLIB** can be specified on the ++PROGRAM MCS.
- To be packaged inline, a program element must be unloaded along with its aliases into a sequential data set and then transformed into the required fixed-block-80 record format before it is packaged (see "Examples" on page 76 for a discussion of doing this). Later, when SMP/E installs the element, it is changed back to its original format. For more information about using GIMDTS, see "GIMDTS: Data transformation service routine" on page 387.
- If the program element is packaged in a LKLIB data set, the ddname specified in the LKLIB operand is required during APPLY and ACCEPT processing.
- For information about packaging program objects or pre-built load modules, see Standard Packaging Rules for z/OS-based Products (publibz.boulder.ibm.com/epubs/pdf/gimpkg80.pdf).

Examples

Suppose you have a partitioned data set whose members contain a pre-built load module (LMODABC) and two aliases (ALIAS1 and ALIAS2). Further, suppose that you want to package LMODABC and its aliases inline in a PTF. The steps to package and install pre-built load module LMODABC in a PTF are as follows:

Step 1: Unload the partitioned data set member LMODABC and its aliases to create a sequential, VS format, data set.

```
//JOBx     JOB …
//STEP1    EXEC PGM=IEBCOPY,REGION=512K
//SYSPRINT DD  SYSOUT=*
//SYSUT3   DD  UNIT=SYSDA,SPACE=(TRK,(5,1))
//SYSUT4   DD  UNIT=SYSDA,SPACE=(TRK,(5,1))
//INPUT    DD  DSN=userid.TEST.LOAD,DISP=SHR
//OUTPUT   DD  DSN=userid.UNLOAD.DATA,DISP=(NEW,CATLG),
//             SPACE=(CYL,(20,10),RLSE),UNIT=SYSDA
//SYSIN    DD  *
  COPY    OUTDD=OUTPUT,INDD=INPUT
  SELECT  MEMBER=LMODABC
  SELECT  MEMBER=ALIAS1
  SELECT  MEMBER=ALIAS2
/*
```

Step 2: Transform the unloaded data set into fixed-block-80 type records by using the SMP/E provided service routine GIMDTS.

```
//JOBx     JOB …
//TFORM    EXEC PGM=GIMDTS
//* -----  GIMDTS IS AN SMP/E LOAD MODULE.
//SYSPRINT  DD SYSOUT=*
//*
//* ----- INPUT TO BE TRANSFORMED - RECFM=VS
//*
//SYSUT1    DD DSN=userid.UNLOAD.DATA,DISP=SHR
//*
//* ----- OUTPUT - RECFM=FB
//*
//SYSUT2    DD DSN=userid.FB80.DATA,DISP=OLD
```

Step 3: Place the fixed block 80 records inline in the PTF following the ++PROGRAM MCS that identifies the element and its aliases to SMP/E.

```
++PTF(PTF0001).
++VER(Z038) FMID(DC00001).
++PROGRAM(LMODABC) ALIAS(ALIAS1,ALIAS2) DISTLIB(TGTLIB)
    SYSLIB(VEND001).
  ⋮
  Place the data transformed by GIMDTS here.
  ⋮
```

Step 4: APPLY the PTF to install LMODABC and its aliases into the appropriate target library.

```
SET BDY(TGT1).
APPLY PTF(PTF0001).
```

During the APPLY step, SMP/E retransforms the inline data back into a variable spanned (VS) sequential data set. SMP/E then invokes the copy utility to copy the retransformed data into the target library. The program element and aliases are copied using a COPYMOD and SELECT statements.

After the APPLY, the target library contains LMODABC and its aliases.

++PTF MCS

The ++PTF MCS identifies a service SYSMOD. This type of SYSMOD can replace or update elements in target and distribution libraries, such as for a permanent correction, or it can add new elements. All other MCSs for this SYSMOD follow this header MCS. For more information about packaging a PTF, see Standard Packaging Rules for z/OS-based Products (publibz.boulder.ibm.com/epubs/pdf/gimpkg80.pdf).

Syntax

++PTF MCS

```
▶▶─ ++PTF ── ( ── sysmod_id ── ) ──────────────────────────────────────▶
                              └─ DESCRIPTION ── ( ── description ── ) ─┘

   ┌──────────────────────────────────────────────────┐
   └─ FILES ── ( ── number ── ) ──────────────────────┘
                              └─ RFDSNPFX ── ( ── relfile_prefix ── ) ─┘

   ┌────────────────────────────────┐ ·─▶◀
   └─ REWORK ── ( ── level ── ) ────┘
```

Operands

DESCRIPTION
specifies a descriptive name to be associated with this SYSMOD. For more information, see Standard Packaging Rules for z/OS-based Products (publibz.boulder.ibm.com/epubs/pdf/gimpkg80.pdf).

- DESCRIPTION can also be specified DESC.
- The DESCRIPTION value can be in single-byte characters (such as English alphanumeric) or double-byte characters (such as Kanji).
- The DESCRIPTION value can contain up to 64 bytes of data, including blanks. (For double-byte data, the 64-byte maximum includes all shift-in and shift-out characters, as well as the double-byte characters.) Extra blanks, as well as leading and trailing blanks are deleted.
- The DESCRIPTION value can span multiple 80-byte records. Data must continue up to and including column 72 and begin in column 1 of the next line. All data past column 72 is ignored. The break does not translate to a blank unless a blank is explicitly coded in column 72 of the first line or in column 1 of the second line.
- If DESCRIPTION is specified, it must contain at least one non-blank character.
- If parentheses are included in the text, they must be in matched pairs.

FILES
specifies the number of relative files belonging to this PTF. It can be a decimal number from 1 to 9999. For information about packaging SYSMODs in relative files, see Standard Packaging Rules for z/OS-based Products (publibz.boulder.ibm.com/epubs/pdf/gimpkg80.pdf).

Note:

1. Although SMP/E allows you to package PTFs in relative files, they are not generally packaged in this format.
2. If a packager uses a high-level qualifier on RELFILE data sets, the RFDSNPFX operand on the header MCS (not the RFPREFIX operand on the RECEIVE command) **must** be used to identify that high-level qualifier.

REWORK
specifies the level of this SYSMOD, which has been reworked for minor changes. Up to eight numeric characters can be specified.

For SYSMODs supplied by IBM, the REWORK level is *yyyyddd*, where *yyyy* is the year the SYSMOD was reworked and *ddd* is the Julian date.

REWORK allows an updated SYSMOD to be automatically received again, as long as it is more recent than the version that has already been received. This takes the place of rejecting the SYSMOD and receiving it again.

Note: If a SYSMOD appears more than once in the SMPPTFIN data set, the first occurrence may be received. However, none of the subsequent versions of the SYSMOD are received, even if their rework level is higher than the one for the first version of the SYSMOD. (Message GIM40001E is issued for each of the subsequent versions of the SYSMOD.)

RFDSNPFX
    identifies to SMP/E the prefix used in the relative file data set names for this SYSMOD. SMP/E uses this prefix when allocating data set names for the SYSMOD's relative files during RECEIVE processing.

- This operand can be specified only if the FILES operand is also specified.

- The RFDSNPFX value specified on the MCS statement must match the actual prefix used in the data set names for the associated relative files.

    For example, if the names of the relative files created for a SYSMOD start with "IBM", as in `IBM.`*sysmod_id*`.F1`, the header MCS statement for the SYSMOD must specify RFDSNPFX(IBM) so SMP/E knows which prefix to use when allocating the data set names for the SYSMOD's relative files during RECEIVE processing.

- Following standard data set naming conventions, the prefix can be from 1 to 8 alphanumeric or national ($, #, @) characters or a dash (–).

    To enable full Security Server protection for tape data sets and to keep the tape header within the 17-character limit (including periods), you should limit the prefix to 1 to 3 characters. If the name exceeds the 17-character limit, only the rightmost 17 characters are written to the tape header label.

sysmod_id
    specifies a unique 7-character system modification identifier for the PTF. For more information, see "Naming conventions for SYSMODs" on page 458.

Usage notes

If you want to update IBM-supplied code, you should use the ++USERMOD MCS rather than the ++PTF MCS. For more information, see "++USERMOD MCS" on page 90.

Examples

A PTF is required that replaces module IFBMOD01 for function FXY1040. The prerequisite service SYSMOD for the module is PTF UZ00004. The APAR incident fixed by this PTF is AZ12345. Here is an example of a SYSMOD containing a ++PTF statement to accomplish this:

```
++PTF(UZ00006)              /* Identify the PTF number. */.
++VER(Z038) FMID(FXY1040)   /* for MVS function FXY1040 */
          PRE(UZ00004)      /* with a prerequisite.     */
          SUP(AZ12345)      /* Fixes one APAR.          */.
++MOD(IFBMOD01)             /* This module             */
          DISTLIB(AOSFB)    /* from this DLIB.          */.
...
... object deck for IFBMOD01
...
```

++RELEASE MCS

The ++RELEASE MCS removes a previously held SYSMOD from *exception SYSMOD* status. ++RELEASE statements are processed by the RECEIVE command.

Syntax

++RELEASE MCS

SYSTEM reason IDs used by IBM

```
    ►►─── ACTION ───►◄
        ──── AO ────
        ─── DB2BIND ───
        ──── DDDEF ────
        ──── DELETE ────
        ───── DEP ─────
        ───── DOC ─────
        ──── DOWNLD ───
        ───── EC ─────
        ───── ENH ─────
        ───── EXIT ────
        ───── EXRF ────
        ──── FULLGEN ───
        ──── IOGEN ────
        ───── IPL ─────
        ─── MSGSKEL ───
        ─── MULTSYS ───
        ─── RESTART ───
```

Operands

ERROR, FIXCAT, SYSTEM, or USER

specifies the hold category from which the SYSMOD is to be removed. You must specify one of these categories.

ERROR

An APAR reported an error in the SYSMOD. The SYSMOD should not be applied or accepted until the APAR is resolved. A PTF held for this reason is also called a *program error PTF*, or PE-PTF. SMP/E automatically resolves the APAR and allows the SYSMOD to be applied or accepted when a SYSMOD that either matches or supersedes the APAR is also applied or accepted. Error holds can be read only from the SMPHOLD data set.

Note: ERROR can also be specified as ERR.

FIXCAT

An APAR provides a fix for the held SYSMOD and the fix is associated with one or more Fix Categories. It is optional whether the APAR will affect processing for the held SYSMOD, based on the APAR's Fix Categories and the Fix Categories of interest specified by the user. If any one or more Fix Categories for the APAR match any of those of interest to the user, then the held SYSMOD will not be applied or accepted until the APAR is resolved. The APAR is resolved when a SYSMOD that matches the APAR name, or a SYSMOD that supersedes the APAR, is applied or accepted. FIXCAT holds can be read only from the SMPHOLD data set.

See the FIXCAT operand for APPLY and ACCEPT command processing for details of specifying the Fix Categories of interest.

SYSTEM

Special action outside normal SMP/E processing is required for the SYSMOD. Examples are SYSMODs requiring a SYSGEN after they are installed, or SYSMODs requiring the installation of an associated engineering change (EC) level. System holds can appear in the SYSMOD itself or in the SMPHOLD data set.

Note: SYSTEM can also be specified as SYS.

USER

The SYSMOD requires special processing because of a decision you have made. User holds can be read only from the SMPHOLD data set.

DATE

specifies the date that the ++HOLD MCS was generated.

FMID
specifies the FMID to which the held SYSMOD is applicable. For external HOLDDATA (a ++HOLD statement not within a SYSMOD), this information allows SMP/E to receive only those statements associated with FMIDs defined in the user's global zone. This operand is required.

REASON
identifies the HOLD reason ID that is to be removed from the SYSMOD. This field is required.

- An *error reason ID* is the number of the APAR that caused the SYSMOD to be placed in exception status.
- A *fix category reason ID* is the SYSMOD ID for the APAR that caused the SYSMOD to be placed into exception status.
- A *system reason ID* is a 1- to 7-character string used to identify some special processing the SYSMOD requires. These are the specific values currently used by IBM:

  ID
  Explanation

  ACTION
  The SYSMOD needs special handling before or during APPLY processing, ACCEPT processing, or both.

  AO
  The SYSMOD may require action to change automated operations procedures and associated data sets and user exits in products or in customer applications. The PTF cover letter describes any changes (such as to operator message text, operator command syntax, or expected actions for operator messages and commands) that can affect automation routines.

  DB2BIND
  A DB2 application REBIND is required for the SYSMOD to become effective.

  DDDEF
  Data set changes or additions as required.

  DELETE
  The SYSMOD contains a ++DELETE MCS, which deletes a load module from the system.

  DEP
  The SYSMOD has a software dependency.

  DOC
  The SYSMOD has a documentation change that should be read before the SYSMOD is installed.

  DOWNLD
  Code that is shipped with maintenance that needs to be downloaded.

  DYNACT
  The changes supplied by the SYSMOD may be activated dynamically without requiring an IPL. The HOLD statement describes the instructions required for dynamic activation. If those instructions are not followed, then an IPL is required for the SYSMOD to take effect.

  EC
  The SYSMOD needs a related engineering change.

  ENH
  The SYSMOD contains an enhancement, new option or function. The HOLD statement provides information to the user regarding the implementation and use of the enhancement.

  EXIT
  The SYSMOD contains changes that may affect a user exit. For example, the interface for an exit may be changed, an exit may need to be reassembled, or a sample exit may be changed.

  EXRF
  The SYSMOD must be installed in both the active and the alternative Extended Recovery Facility (XRF) systems at the same time to maintain system compatibility. (If you are not running XRF, you should bypass this reason ID.)

FULLGEN

The SYSMOD needs a complete system or subsystem generation to take effect.

IOGEN

The SYSMOD needs a system or subsystem I/O generation to take effect.

IPL

The SYSMOD requires an IPL to become effective. For example, the SYSMOD may contain changes to LPA or NUCLEUS, the changes may require a CLPA, or a failure to perform an IPL might lead to catastrophic results, such as could be caused by activation of a partial fix.

Note: If you plan to perform an IPL with CLPA after the SYSMOD has been applied, then no further investigation of the HOLD is required; simply bypass the IPL reason ID. However, if you are not planning to perform an IPL with CLPA, then the details of the HOLD statement must be investigated to determine what kind of actions are required to activate the SYSMOD.

MSGSKEL

This SYSMOD contains message changes that must be compiled for translated versions of the message changes to become operational on extended TSO consoles.

If you want to use translated versions of the messages, you must run the message compiler once for the library containing the English message outlines, and once for each additional language you want to be available on your system. For details, see *z/OS MVS Planning: Operations*.

If you want to use **only** the English version of the messages, you do not need to run the message compiler. You should bypass this reason ID.

MULTSYS

Identifies fixes that need to be applied to multiple systems, in one of three cases: preconditioning, coexistence, or exploitation.

RESTART

To become effective, the SYSMOD requires a special subsystem restart operation. The HOLD statement contains information regarding the required restart actions.

- A *user reason ID* is defined by the user.

For additional information, see "Naming conventions for HOLD reason IDs and HOLD classes" on page 455.

sysmod_id

specifies that SMP/E is to remove the identified SYSMOD from *exception SYSMOD* status. This operand is required.

Usage notes

- ++RELEASE statements are not allowed within a SYSMOD. They are allowed only in SMPHOLD.

- ++RELEASE statements **unconditionally** remove a SYSMOD from exception status and should, therefore, be used with caution. To install a SYSMOD that is currently in exception status, you should probably not create and process a ++RELEASE statement, but rather use the appropriate BYPASS operand of the APPLY or ACCEPT command.

- ++RELEASE statements do not affect ++HOLD statements within a SYSMOD (internal HOLDDATA). However, SMP/E can ignore this internal HOLDDATA during APPLY or ACCEPT processing if `BYPASS(HOLDSYS)` or `BYPASS(HOLDUSER)` is specified.

Examples

The following examples are provided to help you use the ++RELEASE MCS:

Example 1: Removing a SYSMOD from HOLDUSER status

Here is an example of a ++HOLD statement that holds the PTF until after some event (such as new hardware) occurs:

```
++HOLD    (UZ12345)         /* Put this PTF            */
          FMID(FXY1040)     /* for this function       */
          USER              /* into hold user status   */
          REASON(CPU0A)     /* for CPU 0A update.      */
          COMMENT(I DO NOT WANT THIS TO GO ON
                  UNTIL AFTER THE CPU CHANGE)  /*       */.
```

When the CPU change is made, the following sample ++RELEASE statement allows the PTF to be installed:

```
++RELEASE (UZ12345)         /* Remove this PTF         */
          FMID(FXY1040)     /* for this function       */
          USER              /* from hold user status   */
          REASON(CPU0A)     /* for CPU 0A update.      */.
```

That PTF is now eligible for normal installation.

Example 2: Incorrect use of ++RELEASE

Assume the following ++HOLD MCS was processed as part of the normal preventive service installation:

```
++HOLD    (UZ12345)         /* Put this PTF            */
          FMID(FXY1040)     /* for this function       */
          ERROR             /* into hold error status  */
          REASON(AZ12345)   /* for the APAR.           */
          COMMENT(WHEN RUNNING PRODUCT XYZ and OC4
                  ABEND MAY OCCUR)  /*                  */.
```

You are running one system with product XYZ installed and one without product XYZ. The PTF provides a fix for another problem you are encountering on the system without product XYZ; so you want to install this PTF on that system. Here is an example of a ++RELEASE statement that lets you apply PTF UZ12345 without having to use the BYPASS operand:

```
++RELEASE (UZ12345)         /* Remove this PTF         */
          FMID(FXY1040)     /* for this function       */
          ERROR             /* from hold error status  */
          REASON(AZ12345)   /* for the APAR.           */.
```

The risk with this method of processing is that SMP/E no longer has any record of PTF UZ12345 being in exception status. Therefore, the next time any modifications are installed on the system with product XYZ installed, the PTF is installed, introducing a potential 0C4 problem into that system.

The correct way to install the PTF on the system without product XYZ is to use the following command:

```
SET       BDY(MVSTST1)        /* Process MVSTST1 tgt zone.*/.
APPLY     S(UZ12345)          /* Process the PTF.       */
          BYPASS(HOLDERR(AZ12345)) /* Bypass known error.*/.
```

Now the PTF is installed on one system, but SMP/E still remembers that it is in hold error status and does not allow it to be installed on any other system.

Note: Any other applicable operand (FORFMID, SOURCEID, and so on) can be used in place of the SELECT or S operand.

Example 3: A ++RELEASE statement with a FIXCAT HOLD

Here is an example of the use of the FIXCAT hold category on a ++RELEASE statement:

```
++RELEASE (HBB7720)         /* Remove this PTF                */
          FMID(HBB7720)     /* for this function             */
          REASON(AK18603)   /* for this APAR                 */
```

```
      FIXCAT            /* from fix category hold status  */
      DATE(06231)       /* on this date                   */.
```

++RENAME MCS

During APPLY processing, the ++RENAME MCS renames a load module in all the target libraries and the LMOD entry in the target zone. It also updates the LMOD subentry of the MOD entry for each module in the load module. For load modules having a CALLLIBS subentry (or having a SYSLIB allocation), SMP/E also renames the base version of the load module in the SMPLTS library. For load modules having a side deck library subentry, SMP/E also renames the definition side deck in the side deck library.

During ACCEPT processing, if inline JCLIN was processed, the distribution zone entries are updated to reflect the ++RENAME changes. The target libraries and target zone entries are not updated.

Note:

1. You cannot use ++RENAME to rename a load module for which symbolic links have been defined.

2. You cannot use ++RENAME to rename macros, modules, source, or other elements. SMP/E does, however, support these ways to rename an element:

   - **For elements other than modules,** you must use an element MCS statement to delete the original element, and another element MCS statement to reintroduce it using the new name. You can do both of these steps in the same SYSMOD. For examples, see the descriptions of the various element MCS statements in this chapter.

   - **For modules,** you must use a ++MOD statement to delete the original module, and another ++MOD statement to reintroduce it using the new name. In addition, you must ensure that the renamed module is included in the appropriate load modules, either through JCLIN or the LMOD operand on the ++MOD statement. You can do all of these steps in the same SYSMOD. For an example, see "Example 3: Packaging a renamed module" on page 67 in this chapter.

Syntax

++RENAME MCS

▶▶── ++RENAME ── (── *old_name* ──) ── TONAME ── (── *new_name* ──) ── • ──▶◀

Operands

old_name
    specifies the name of the load module to be renamed.

TONAME
    specifies the new name for the load module.

Usage notes

- There are no optional operands. You must specify the load module name and its new name.

- ++RENAME statements must follow any ++VER and ++IF statements, and must precede any ++JCLIN or element MCSs.

- Regardless of the order in which ++MOVE, ++RENAME, and ++DELETE statements are coded in a SYSMOD, they are processed in the following order:

  – APPLY and ACCEPT

     1. ++MOVE
     2. ++RENAME
     3. ++DELETE

  – RESTORE

     1. ++RENAME

    2. ++MOVE

++JCLIN statements are processed next, followed by element statements.

Examples

Here is an example of a SYSMOD containing a ++RENAME statement renaming load module LMODA:

```
++PTF(UR01234)              /* Identify the PTF number. */.
++VER(Z038) FMID(HXY1300)   /* For MVS function HXY1300.*/.
++IF (ESY1300) THEN         /* If ESY1300 is installed  */
  REQ(UR12399)              /* UR12399 is required.     */.
++RENAME (LMODA)            /* Rename load module LMODA */
  TONAME(LMODBB)            /* to LMODBB.               */.
++JCLIN                     /* JCLIN follows.           */.
        .
        .
        .
++MOD(MODAA) DISTLIB(AOSXX) /* Element MCS statements.  */.
```

++SRC MCS

The ++SRC MCS describes a single source replacement. It must immediately precede the source definition statements when they are within the SYSMOD.

You should use the ++SRC MCS when you want to provide the source form of a module and have it get assembled when the SYSMOD is installed. If you want to provide the object form of the module, use the ++MOD MCS instead.

Syntax

The syntax to be used depends on the processing to be done for the element:

- Adding or replacing the element
- Deleting the element

Adding or replacing source

++SRC MCS
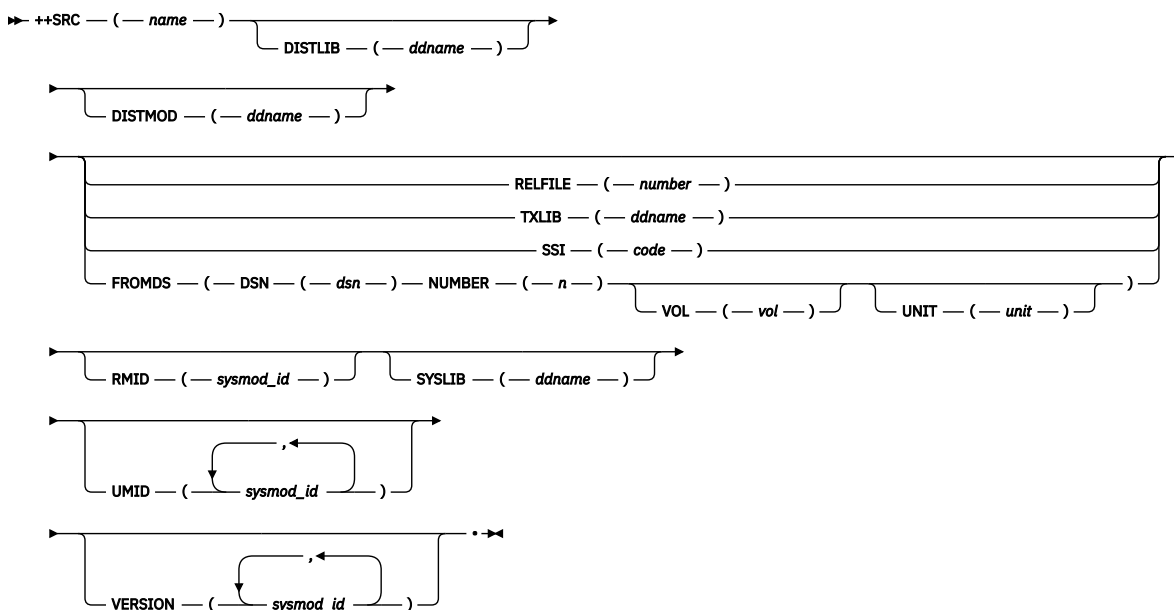
Deleting source

++SRC MCS

Operands

DELETE
indicates that the source is to be removed from the target libraries, distribution libraries, and SMP/E data sets.

Note: DELETE is mutually exclusive with all other operands except DISTLIB and VERSION.

DISTLIB
specifies the ddname of the distribution library for the specified source.

Note: DISTLIB must be specified if the source has not been previously recorded on the target zone or distribution zone. If a SRC entry already exists in the target zone or distribution zone, and the value currently in that entry does not match that specified in the DISTLIB operand, the SYSMOD is not applied or accepted, unless that SYSMOD also used the ++MOVE MCS to change the DISTLIB to that new value.

DISTMOD
specifies the ddname of the link-edit distribution library for the assembled source code. During ACCEPT processing, the object code from the assembler is link-edited to the library specified.

FROMDS
identifies the partitioned data set that contains this element.

Note: The FROMDS operand and its DSN, NUMBER, VOL, and UNIT suboperands are included in the MCS generated by the BUILDMCS command. IBM does not intend the FROMDS operand to be used in manually coded MCS.

DSN
specifies the dsname of the FROMDS data set. The specified data set name must conform to standard data set naming conventions and cannot contain parentheses. The maximum length of the entire name is 44 characters (including the periods).

NUMBER
specifies a number that SMP/E is to use when assigning a name to the SMPTLIB data set associated with this FROMDS data set. (This is similar to the way the relative file number is used in RELFILE processing.)

VOL
specifies, for an uncataloged data set, the volume serial number of the volume containing the FROMDS data set. If specified, this volume identifier must be from 1 to 6 alphanumeric characters.

VOL may be omitted for a cataloged data set.

UNIT
specifies, for an uncataloged data set, the UNIT type containing the FROMDS data set. If specified, the UNIT value must be from 1 to 8 characters and must conform to standard UNIT naming conventions. IBM SMP/E for z/OS, V3R6 accepts any nonblank characters specified between the open and close parentheses, up to a maximum length of 8.

UNIT may be omitted for a cataloged data set.

Note: FROMDS is mutually exclusive with DELETE, RELFILE, SSI, and TXLIB.

name
specifies the name of the source in the distribution library and, optionally, in the target library. The name can contain any alphanumeric characters and $, #, @, or hex C0.

RELFILE
identifies which relative file associated with the SYSMOD contains this element. This operand is required if you provide the element in RELFILE format, rather than inline or in a TXLIB data set.

Note:

1. SMPTLIB cannot be used as a value on the TXLIB operand.
2. RELFILE is mutually exclusive with FROMDS and TXLIB.

RMID
specifies the last SYSMOD that **replaced** this source. This operand can be used only in a service-updated function, and the specified PTF must be integrated into the function.

SSI
specifies eight hexadecimal digits of system status information. This information is placed in the directory of the target system library or SMPMTS or SMPSTS during APPLY processing, and in the distribution library during ACCEPT processing, as four packed hexadecimal bytes of user data. See the IEBUPDTE program description in *z/OS DFSMSdfp Utilities*.

Note: This operand is ignored if text is located in a library, as is the case when either the FROMDS, RELFILE, or TXLIB operand is specified.

SYSLIB
specifies the ddname of the target library, if the source module exists in one. APPLY and RESTORE processing update this library.

Note: If a SRC entry already exists in the target zone or distribution zone and the value currently in that entry does not match that specified in the SYSLIB operand, SMP/E ignores the SYSLIB value in the SYSMOD being installed, unless that SYSMOD also used the ++MOVE MCS to change the SYSLIB to that new value.

TXLIB
is the ddname of the partitioned data set containing the source. This operand is required if the module is provided in a data set the users have access to, rather than inline or in RELFILE format.

Note:

1. SMPTLIB cannot be used as a value on the TXLIB operand.
2. TXLIB is mutually exclusive with FROMDS and RELFILE.

UMID
specifies the UMIDs of the source. This operand can be used only in function SYSMODs and specifies the PTF service level of the source (the set of SYSMODs that have updated this source since it was last replaced).

VERSION
specifies one or more function SYSMODs that currently contain the element. The function containing the ++SRC MCS takes over ownership of the element from the specified functions.

When VERSION is specified on an element statement, it overrides any VERSION operand values that might be specified on the ++VER MCS.

Usage notes

- If the source is packaged inline, it must immediately follow the ++SRC MCS and must not contain any records starting with ++. Neither FROMDS, nor RELFILE, nor TXLIB can be specified.
- If the source is in a TXLIB data set, the ddname specified on the TXLIB operand is required during APPLY and ACCEPT processing.
- For information about packaging SYSMODs in RELFILE, TXLIB, or inline format, see Standard Packaging Rules for z/OS-based Products (publibz.boulder.ibm.com/epubs/pdf/gimpkg80.pdf).

Examples

The following examples are provided to help you use the ++SRC MCS:

Example 1: Adding a new source to the system

A replacement for the source IFBSRC01 is in a partitioned data set referred to by the ddname REPLACE. The distribution library for the source is SYS1.IFBSRC; SYS1.AOS23 is the distribution library for the module, IFBSRC01, resulting from the assembly of the source, IFBSRC01.

Here is an example of a SYSMOD containing a ++SRC statement that causes SMP/E to install the source code, assemble it, and save the resulting module in the DLIBs:

```
++USERMOD(USR0001)          /* User modification       */.
++VER(Z038) FMID(FXY1040)   /* for user function in MVS.*/.
++SRC(IFBSRC01)             /* Replace source          */
          DISTLIB(IFBSRC) /* in this DLIB.           */
          DISTMOD(AOS23)  /* MOD goes in this DLIB.  */
          TXLIB(REPLACE)  /* Replacement SRC is here. */.
```

The following DD statements are required when the SYSMOD is applied:

```
//REPLACE  DD DSN=…
//SMPSTS   DD DSN=SYS1.SMPSTS,DISP=OLD
```

plus whatever other DD statements are required based on which load modules the assembled source is to be linked to. These load modules should be identified by JCLIN, either as a separate step or within the SYSMOD itself. Assuming the load module was composed of only this one module, the following ++JCLIN MCS can be added to the SYSMOD after the ++VER statement.

```
++JCLIN.                    /* JCLIN to get SRC linked. */
//JOB1     JOB 'accounting info',MSGLEVEL=(1,1)
//STEP1    EXEC PGM=IEBCOPY
//AOS23    DD DSN=SYS1.AOS23,DISP=SHR
//LPALIB   DD DSN=SYS1.LPALIB,DISP=SHR
//SYSIN    DD *
 COPY INDD=AOS23,OUTDD=LPALIB
 SELECT M=(IFBSRC01)
/*
```

In this case, you also need the following DD statement when the SYSMOD is applied:

```
//LPALIB   DD DSN=SYS1.LPALIB,DISP=OLD
```

The following DD statements are required when the SYSMOD is accepted:

```
//REPLACE  DD DSN=…
//IFBSRC   DD DSN=SYS1.IFBSRC,DISP=OLD
//AOS23    DD DSN=SYS1.ASO23,DISP=OLD
```

Example 2: Packaging renamed source

Suppose that, for some reason, you need to rename source IFBSRC01, which was introduced in . The new name is to be USRSRCXX. You do not need to change anything else about the source. Because the original source was linked by JCLIN, you must provide updated JCLIN to link the renamed version of the source. Here is an example of a SYSMOD containing the ++SRC statements and JCLIN needed to package this renamed source:

```
++USERMOD(USR0002)          /* User modification       */.
++VER(Z038) FMID(FXY1040)   /* for user function in MVS.*/
          PRE(USR0001) /* Base on previous USERMOD.*/.
++SRC(IFBSRC01)             /* Delete the original     */
          DELETE          /* source.                 */.
++SRC(USRSRCXX)             /* Add renamed source      */
          DISTLIB(IFBSRC) /* to this DLIB.           */
          DISTMOD(AOS23)  /* MOD goes in this DLIB.  */
          TXLIB(REPLACE)  /* Replacement SRC is here. */.
…
```

```
++JCLIN.                      /* JCLIN to get SRC linked. */
//JOB1     JOB 'accounting info',MSGLEVEL=(1,1)
//STEP1    EXEC PGM=IEBCOPY
//AOS23    DD DSN=SYS1.AOS23,DISP=SHR
//LPALIB   DD DSN=SYS1.LPALIB,DISP=SHR
//SYSIN    DD *
 COPY INDD=AOS23,OUTDD=LPALIB

SELECT M=(USRSRCXX)
 /*
```

You must ensure that the renamed source is in the TXLIB library (REPLACE) used to provide SMP/E with the element.

The following DD statements are required when the SYSMOD is applied:

```
//REPLACE  DD DSN=…
//SMPSTS   DD DSN=SYS1.SMPSTS,DISP=OLD
//LPALIB   DD DSN=SYS1.LPALIB,DISP=OLD
```

The following DD statements are required when the SYSMOD is accepted:

```
//REPLACE  DD DSN=…
//IFBSRC   DD DSN=SYS1.IFBSRC,DISP=OLD
//AOS23    DD DSN=SYS1.AOS23,DISP=OLD
```

++SRCUPD MCS

The ++SRCUPD MCS describes a single set of source update statements within a PTF, an APAR fix, or a USERMOD. It must immediately precede the source update statements within the SYSMOD.

Syntax

++SRCUPD MCS

Operands

name
    specifies the name of the source in the distribution library and, optionally, in the target library. The name can contain any alphanumeric characters and $, #, @, or hex C0.

DISTLIB
    specifies the ddname of the distribution library for the specified source.

    Note: DISTLIB must be specified if the source has not been previously recorded on the target zone or distribution zone. If a SRC entry already exists in the target zone or distribution zone, and the value currently in that entry does not match that specified in the DISTLIB operand, the SYSMOD is not applied or accepted.

DISTMOD
    specifies the ddname of the link-edit distribution library for the assembled source code. During ACCEPT processing, the object code from the assembler is link-edited to the library specified.

SYSLIB
    specifies the ddname of the target library, if the source exists in one. APPLY and RESTORE processing update this library.

Usage notes

- If a SYSMOD containing a ++SRCUPD statement attempts to change the ownership (FMID) of the element (with the VERSION operand), the SYSMOD cannot be installed.
- The changes for the source must immediately follow the ++SRCUPD MCS and must not contain any records starting with ++.
- The only IEBUPDTE control statements allowed in a SYSMOD are ./ CHANGE and ./ ENDUP.
- The only IEBUPDTE CHANGE operand that SMP/E checks is NAME, which must specify the same element as the ++SRCUPD MCS. Other CHANGE operands may produce undesired results and are used at your own risk. For example, if you code UPDATE=INPLACE, SMP/E may update the distribution library. Once the distribution libraries are changed, there is no way to remove the updates.
- When processing multiple updates to the same lines in a given source module, SMP/E uses the ./ CHANGE statement from the last update to the lines.
- SMP/E does not support a continuation of the ./ CHANGE statement.

Examples

The following examples are provided to help you use the ++SRCUPD MCS:

Example 1: Updating an existing source

Here is an example of a SYSMOD containing a ++SRCUPD statement to make a modification to an IBM source module, in this case module JESMOD01, which is part of product EJS1102:

```
++USERMOD(MY00001)        /* User modification     */.
++VER(Z038) FMID(EJS1102) /* for MVS JES.          */
            PRE(UZ12345)  /* Current service level. */.
++SRCUPD(JESMOD01)        /* Update this JES MOD.  */
                          /* DISTLIB already known. */.
./ CHANGE NAME=JESMOD01
        LA   R1,MYPARM                    00001000
        BALR MYPGM                        00001100
MYPGM   EQU  *                            00500000
        …                                 00500100
        …                                 00500200
        B    R14                          00500300
```

Example 2: Making subsequent source updates

Assume that, after installing the modification given in the first example, you need to make another modification. Assume the following lines had to be changed:

```
        …                                 00500150
        BALR   MYPGM2                     00500160
MYPGM2  EQU    *                          00600000
        …                                 00600100
        …                                 00600200
        B      R14                        00600300
```

You have two choices as to the method of building the second modification:

1. The first method is to build the second SYSMOD to contain only the newly changed lines of code, and then specify the PRE operand in both the current IBM service level and the previous user modification.
2. The second method is to build the second SYSMOD to contain all the user modifications to this module; then specify the PRE operand in the current IBM service level and supersede the previous user modification.

Here is an example of a SYSMOD for the first method:

```
++USERMOD(MY00002)        /* User modification     */.
++VER(Z038) FMID(EJS1102) /* for MVS JES.          */
            PRE(UZ12345   /* Current service level */
                MY00001)  /* plus previous USERMOD. */.
++SRCUPD(JESMOD01)        /* Update this JES MOD.  */
```

```
                                  /* DISTLIB already known.   */.
./ CHANGE NAME=JESMOD01
       …                                          00500150
       BALR  MYPGM2                               00500160
MYPGM2 EQU   *                                    00600000
       …                                          00600100
       …                                          00600200
       B     R14                                  00600300
```

Here is an example of a SYSMOD for the second method:

```
++USERMOD(MY00002)        /* User modification        */.
++VER(Z038) FMID(EJS1102) /* for MVS JES.             */
           PRE(UZ12345)   /* Current service level.   */
           SUP(MY00001)   /* Supersede previous usermod.*/.
++SRCUPD(JESMOD01)        /* Update this JES MOD.     */
                          /* DISTLIB already known.   */.
./ CHANGE NAME=JESMOD01
       LA    R1,MYPARM                            00001000
       BALR MYPGM                                 00001100
MYPGM  EQU   *                                    00500000
       …                                          00500100
       BALR  MYPGM2                               00500160
       …                                          00500200
       B    R14                                   00500300
       …                                          00500150
MYPGM2 EQU   *                                    00600000
       …                                          00600100
       …                                          00600200
       B     R14                                  00600300
```

The advantages to the second method are that all modifications to one source are contained in a single SYSMOD, and if that SYSMOD has to be restored or re-applied, processing is much more efficient.

Note: The SMP/E ++SRCUPD MCS is the same for both methods; only the ++VER MCS and the actual update control cards change.

++USERMOD MCS

The ++USERMOD MCS identifies a user modification. This type of SYSMOD can be used to add user-defined functions or to replace or update elements for IBM-supplied code in the target or distribution libraries. All other MCSs for this SYSMOD follow this header MCS. For more information about packaging a USERMOD, see Standard Packaging Rules for z/OS-based Products (publibz.boulder.ibm.com/epubs/pdf/gimpkg80.pdf).

Syntax

++USERMOD MCS

```
►►─ ++USERMOD ── ( ── sysmod_id ── ) ──────────────────────────►
                         └─ DESCRIPTION ── ( ── description ── ) ─┘

    ┌─────────────────────────────────────────────────┐
────┴─ FILES ── ( ── number ── ) ────────────────────────┴──────►
                    └─ RFDSNPFX ── ( ── relfile_prefix ── ) ─┘

────┬──────────────────────────────── • ──►◄
    └─ REWORK ── ( ── level ── ) ─┘
```

Operands

DESCRIPTION

specifies a descriptive name to be associated with this SYSMOD. For more information, see Standard Packaging Rules for z/OS-based Products (publibz.boulder.ibm.com/epubs/pdf/gimpkg80.pdf).

- DESCRIPTION value can also be specified DESC.
- The DESCRIPTION value can be in single-byte characters (such as English alphanumeric) or double-byte characters (such as Kanji).

- The DESCRIPTION value can contain up to 64 bytes of data, including blanks. (For double-byte data, the 64-byte maximum includes all shift-in and shift-out characters, as well as the double-byte characters.) Extra blanks, as well as leading and trailing blanks are deleted.
- The DESCRIPTION value can span multiple 80-byte records. Data must continue up to and including column 72 and begin in column 1 of the next line. All data past column 72 is ignored. The break does not translate to a blank unless a blank is explicitly coded in column 72 of the first line or in column 1 of the second line.
- If DESCRIPTION is specified, it must contain at least one non-blank character.
- If parentheses are included in the text, they must be in matched pairs.

FILES

specifies the number of relative files belonging to this USERMOD. It can be a decimal number from 1 to 9999. For information about packaging SYSMODs in relative files, see Standard Packaging Rules for z/OS-based Products (publibz.boulder.ibm.com/epubs/pdf/gimpkg80.pdf).

Note:

1. Although SMP/E allows you to package USERMODs in relative files, they are not generally packaged in this format.

2. If a packager uses a high-level qualifier on RELFILE data sets, the RFDSNPFX operand on the header MCS (not the RFPREFIX operand on the RECEIVE command) **must** be used to identify that high-level qualifier.

REWORK

specifies the level of this SYSMOD, which was reworked for minor changes. Up to eight numeric characters can be specified.

For SYSMODs supplied by IBM, the REWORK level is *yyyyddd*, where *yyyy* is the year the SYSMOD was reworked and *ddd* is the Julian date.

REWORK allows an updated SYSMOD to be automatically received again, as long as it is more recent than the version that has already been received. This takes the place of rejecting the SYSMOD and receiving it again.

Note: If a SYSMOD appears more than once in the SMPPTFIN data set, the first occurrence may be received. However, none of the subsequent versions of the SYSMOD are received, even if their rework level is higher than the one for the first version of the SYSMOD. (Message GIM40001E is issued for each of the subsequent versions of the SYSMOD.)

RFDSNPFX

identifies to SMP/E the prefix used in the relative file data set names for this SYSMOD. SMP/E uses this prefix when allocating data set names for the SYSMOD's relative files during RECEIVE processing.

- This operand can be specified only if the FILES operand is also specified.
- The RFDSNPFX value specified on the MCS statement must match the actual prefix used in the data set names for the associated relative files.

  For example, if the names of the relative files created for a SYSMOD start with "IBM", as in IBM.*sysmod_id*.F1, the header MCS statement for the SYSMOD must specify RFDSNPFX(IBM) so SMP/E knows which prefix to use when allocating the data set names for the SYSMOD's relative files during RECEIVE processing.

- Following standard data set naming conventions, the prefix can be from 1 to 8 alphanumeric or national ($, #, @) characters or a dash (–).

  To enable full Security Server protection for tape data sets and to keep the tape header within the 17-character limit (including periods), you should limit the prefix to 1 to 3 characters. If the name exceeds the 17-character limit, only the rightmost 17 characters are written to the tape header label.

sysmod_id

specifies a unique 7-character system modification identifier for the USERMOD. For more information, see "Naming conventions for SYSMODs" on page 458.

Usage notes

If you have updated an element, SYSMODs other than functions cannot replace it unless you explicitly allow them to by bypassing MODID checking. However, if you install a function, it may overlay your user modifications. SMP/E issues a warning message when it detects this condition.

Examples

A source (IFBSRC02) owned by function SYSMOD FXY1040 is to be modified. Your modification requires a service level provided UZ00007; you are only updating, not replacing, the source. You have chosen a SYSMOD ID of MY00001. Here is an example of a SYSMOD containing a ++SRCUPD statement that accomplishes this:

```
++USERMOD(MY00001)          /* USERMOD number         */.
++VER(Z038) FMID(FXY1040)   /* for function FXY1040   */
          PRE(UZ00007)      /* at this service level. */.
++SRCUPD(IFBSRC02)          /* Update this source     */
          DISTLIB(IFBSRC) /* in this DLIB.          */.
./ CHANGE NAME=IFBSRC02
...
... update control cards
...
```

For additional examples of USERMODs, see "++SRCUPD MCS" on page 88 and *z/OS SMP/E User's Guide*.

++VER MCS

The ++VER MCS describes the environment required for receiving and installing a SYSMOD. A SYSMOD must contain a separate ++VER MCS for each environment to which it applies. At least one ++VER MCS must be present in a SYSMOD, and a maximum of 255 ++VER statements are allowed for each SYSMOD.

Syntax

++VER MCS
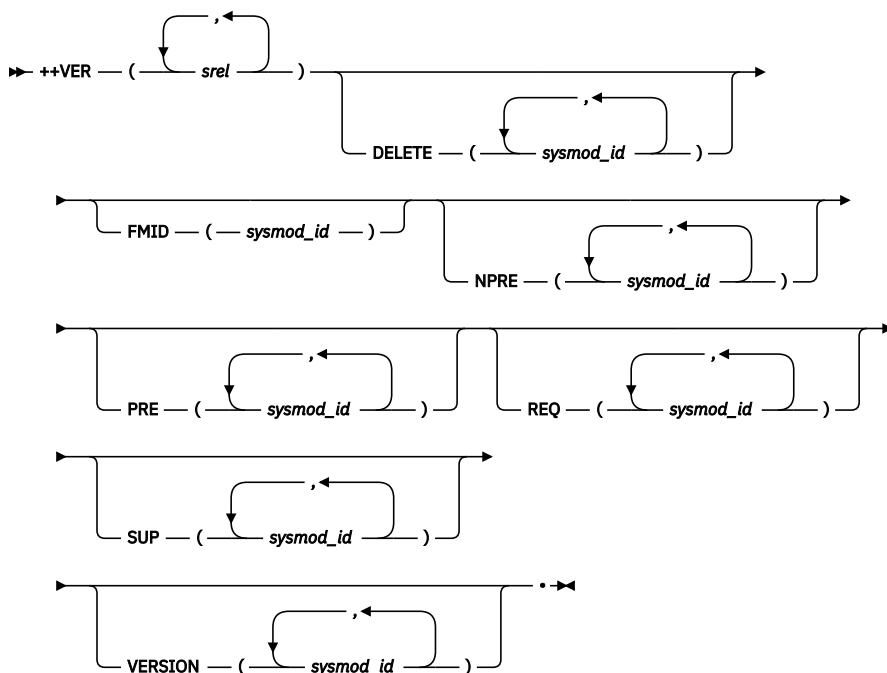
Operands

DELETE
indicates which function SYSMODs should be deleted when this function is installed. These functions are permanently deleted and cannot be restored.

DELETE can be specified only in function SYSMODs.

The same SYSMOD can be specified on both DELETE and SUP. This cleans up entries for the deleted function, and, at the same time, allows SYSMODs that name the deleted function as a requisite to still be installed.

SYSMODs specified in the DELETE operand do not have to be specified in VERSION operands of ++VER, ++MAC, ++SRC, or ++MOD statements.

For additional information about the effect of DELETE during APPLY and ACCEPT processing, see *z/OS SMP/E Commands*.

FMID
FMID identifies the function to which a SYSMOD applies. FMID must be specified for all SYSMODs except base functions.

The following considerations relate to the FMID operand:

- Unlike prerequisites specified by the PRE operand, the functional prerequisite specified by the FMID operand is satisfied only by the specified SYSMOD. It is not satisfied by another SYSMOD that supersedes that function.

- When specified on the ++VER MCS for a function, FMID defines the function as a dependent function. In this case, FMID indicates that the elements supplied by the dependent function SYSMOD are functionally higher than the specified base function.

  A function cannot be both a base function and a dependent function. Therefore, if a base function contains more than one ++VER MCS, none of them can specify the FMID operand. Likewise, if a dependent function contains more than one ++VER MCS, all of them must specify the FMID operand.

- When specified on the ++VER MCS for a non-function SYSMOD, FMID indicates the functional level of all elements in the SYSMOD.

- SMP/E RECEIVE processing does not receive a dependent function unless the FMID of the base is already present in the global zone or BYPASS(FMID) is specified. For more information about packaging of dependent SYSMODs, refer to RECEIVE processing in *z/OS SMP/E Commands*.

NPRE
indicates which function SYSMODs cannot exist in the same zone as this function. These are negative prerequisite SYSMODs. The current SYSMOD cannot be applied or accepted if any of the listed SYSMODs are already present.

This operand has no effect on RECEIVE eligibility.

NPRE can only be specified within a function SYSMOD.

PRE
indicates which SYSMODs are prerequisites for this SYSMOD. A prerequisite SYSMOD must either be already installed, or must be installed concurrently with this SYSMOD. For additional information about how prerequisites are resolved during APPLY and ACCEPT processing, see *z/OS SMP/E Commands*.

If a SYSMOD replaces an existing element, the PRE operand must specify the previous SYSMOD that replaced the element (RMID) and all the SYSMODs that have updated the element (UMIDs) since it was last replaced.

If a SYSMOD updates an existing element, the PRE operand must specify the previous SYSMOD that replaced the element. It should also specify the last SYSMOD that updated the element since then.

This operand has no effect on RECEIVE eligibility.

REQ

indicates which SYSMODs are requisites for this SYSMOD. The specified SYSMOD must either be already installed, or must be installed concurrently with this SYSMOD. If the specified SYSMOD also specifies this SYSMOD as a requisite, these two SYSMODs are corequisites, and neither can be installed independently; they must be installed within the same APPLY and ACCEPT command. For additional information about how requisites are resolved during APPLY and ACCEPT processing, see *z/OS SMP/E Commands*.

This operand has no effect on RECEIVE eligibility.

srel

specifies the system or subsystem release on which the SYSMOD can be installed. The SREL must contain four alphanumeric characters, usually one alphabetic character followed by three numeric characters. These are the systems and subsystems defined by IBM, with their SRELs:

System
   SREL

DB2
   P115

CICS
   C150

IMS
   P115

MVS
   Z038

NCP
   P004

The SREL is used during RECEIVE processing to determine whether a SYSMOD should be received. For more information about how the SREL operand is processed, see *z/OS SMP/E Commands*.

SUP

indicates which SYSMODs are superseded (contained in and replaced) by this SYSMOD. For example, it may specify one or more APARs fixed in the element modifications supplied with this SYSMOD.

For functions, the same SYSMOD can be specified on both DELETE and SUP. This cleans up entries for the deleted function, and, at the same time, allows SYSMODs that name the deleted function as a requisite to still be installed.

For PTFs, APARs, and USERMODs to properly supersede another SYSMOD, the following items from the superseded SYSMOD must be contained in either the superseding SYSMOD or a SYSMOD from the requisite set for the superseding SYSMOD:

- all the elements, unless the element is being deleted by the superseding SYSMOD.
- all the ++VER PRE, REQ, SUP, and VERSION operand values, unless the superseding SYSMOD is correcting any of these values from the superseded SYSMOD.
- all ++HOLD statements, unless the superseding SYSMOD is correcting any of these from the superseded SYSMOD.

VERSION

indicates functions whose elements should be considered functionally lower than the elements contained in this SYSMOD. It specifies one or more function SYSMODs that currently contain the element. The function containing the ++VER MCS takes over ownership of all the elements from the specified functions.

When VERSION is specified on an element statement, it overrides any VERSION operand values specified on the ++VER MCS.

Note: A SYSMOD containing an element update (++MACUPD, ++SRCUPD, or ++ZAP) cannot change the ownership of the element. The ownership can be changed (with the ++VER VERSION operand) only if the SYSMOD provides a replacement for the element.

Usage notes

- A SYSMOD cannot contain multiple ++VER statements that have identical SREL and FMID values, because SMP/E would not be able to determine which ++VER MCS to use in doing the remaining applicability checking during APPLY and ACCEPT processing.
- You cannot specify the same SYSMOD more than once on a single ++VER operand. Likewise, you generally cannot specify the same SYSMOD on more than one operand. However, you can specify the same SYSMOD on VERSION and another operand (except FMID). You can also specify the same SYSMOD on the DELETE and SUP operands.
- Corequisite SYSMODs (which are related through the REQ operand) that are applicable to the same FMID **cannot** have elements in common. Because the REQ operand implies no service hierarchy, SMP/E cannot determine which SYSMOD has the highest service level of the common elements. When the relationship between the SYSMODs containing the common elements is defined through the REQ operand, SMP/E issues an error or warning message.
  - If the requisites have an element in common and each contains a replacement for the element, ID check processing fails and neither of the requisites is installed, unless BYPASS(ID) is specified.
  - If the requisites have an element in common and each contains an update for the element, ID check processing issues a warning message for the requisites. Generally, the requisites are both installed. However, SMP/E does not allow multiple ZAPs for the same module to be processed by the same APPLY command. If this is the case, neither of the requisites is installed; they must be processed by separate APPLY commands.

Examples

The following examples are provided to help you use the ++VER MCS:

Example 1: Defining base and dependent functions

Assume you want to package one of your user applications as a function so you can use SMP/E in installing and maintaining it. You also have an optional enhancement to that product you want to package. The base function has two modules, USRMOD01 and USRMOD02, which are in link-edit format and reside in the library pointed to by the USRLIBXX DD statement. The optional enhancement changes USRMOD02 and depends on the base application being present. Here are examples of SYSMODs for these functions:

```
++FUNCTION(EUSR001)            /* Base application       */.
++VER(Z038)                    /* for MVS version.       */.
++MOD(USRMOD01)                /* Has this MOD           */
        DISTLIB(AOS12)  /* in this DLIB.          */
        LKLIB(USRLIBXX) /* Replacement is here    */.
++MOD(USRMOD02)                /* Has this MOD           */
        DISTLIB(AOS12)  /* in this DLIB.          */
        LKLIB(USRLIBXX) /* Replacement is here.   */.
```

```
++FUNCTION(FUSR011)            /* Dependent function     */.
++VER(Z038)                    /* for MVS version.       */
        FMID(EUSR001)   /* Base application must be
                               present.            */.
++MOD(USRMOD02)                /* Has this MOD           */
        DISTLIB(AOS12)  /* in this DLIB.          */
        LKLIB(USRLIBXX) /* Replacement is here.   */.
++MOD(USRMOD03)                /* Has this MOD           */
        DISTLIB(AOS12)  /* in this DLIB.          */
        LKLIB(USRLIBXX) /* Replacement is here.   */.
```

The dependent function, FUSR011, specifies the base function on its FMID operand. This means the base function must be present in order for the dependent function to be installed. The FMID operand also indicates that if the two functions have any elements in common, the version in the dependent product is the one that is to be installed. The dependent product has now assumed ownership of those elements.

Example 2: Defining intersecting dependent functions

Assume you want to add another dependent function for the base function defined in the preceding example. It works whether or not the other dependent function is installed. Here is an example of a SYSMOD containing a ++VER statement to define the relationship between the dependent functions:

```
++FUNCTION(FUSR012)        /* Dependent function    */.
++VER(Z038)                /* for MVS version.      */
        FMID(EUSR001)   /* Base application must be
                           present.            */
        VERSION(FUSR011)  /* If present is at higher
                             functional level.    */.
++MOD(USRMOD02)            /* Has this MOD          */
        DISTLIB(AOS12)  /* in this DLIB.         */
        LKLIB(USRLIBXX) /* Replacement is here.  */.
++MOD(USRMOD04)            /* Has this MOD          */
        DISTLIB(AOS12)  /* in this DLIB.         */
        LKLIB(USRLIBXX) /* Replacement is here.  */.
```

After this function is installed, module USRMOD02 is the version from function FUSR012, no matter what the former functional level was.

Example 3: Deleting a previous level of a function

Assume you want to provide a new level of the base function defined in the first example. It includes both of the dependent functions for the original base function. Here is an example of a SYSMOD containing a ++VER statement deleting the previous level of the function:

```
++FUNCTION(EUSR002)        /* New base application  */.
++VER(Z038)                /* for MVS version.      */
        DELETE(EUSR001) /* Delete prev level.    */.
++MOD(USRMOD01)            /* Has this MOD          */
        DISTLIB(AOS12)  /* in this DLIB.         */
        LKLIB(USRLIBXX) /* Replacement is here.  */.
++MOD(USRMOD02)            /* Has this MOD          */
        DISTLIB(AOS12)  /* in this DLIB.         */
        LKLIB(USRLIBXX) /* Replacement is here.  */.
++MOD(USRMOD03)            /* Has this MOD          */
        DISTLIB(AOS12)  /* in this DLIB.         */
        LKLIB(USRLIBXX) /* Replacement is here.  */.
++MOD(USRMOD04)            /* Has this MOD          */
        DISTLIB(AOS12)  /* in this DLIB.         */
        LKLIB(USRLIBXX) /* Replacement is here.  */.
```

After the new function is installed, all references to SYSMODs EUSR001, FUSR011, and FUSR012 are deleted from the target zone and distribution zone (except the SYSMOD entry for EUSR001, which indicates it was deleted by EUSR002).

This SYSMOD does not require the previous level of the base function to be installed. The DELETE operand just says that if that previous function was installed, SMP/E should delete it before installing the new level. For more information about delete processing for the APPLY and ACCEPT commands, see *z/OS SMP/E Commands*.

Example 4: Deleting a function without replacing it (dummy delete)

Assume you no longer need a particular function, and you want to delete it from your system. First, you must make sure that no other functions depend on the function you want to delete. Once you have done this, you need to define a dummy function SYSMOD that deletes the function you want to delete. You then receive, apply, and accept the dummy function, and run UCLIN to delete the SYSMOD entries for the deleted function and for the dummy function.

For example, assume you are ready to delete function MYFUNC1 using dummy function DELFUNC. MYFUNC1 is applicable to SREL Z038 and is installed in target zone TGT1 and distribution zone DLIB1. Here is an example of the dummy function:

```
++FUNCTION(DELFUNC)      /* Any valid unique SYSMOD ID.   */.
++VER(Z038)              /* For SREL Z038 (MVS products). */
     DELETE(MYFUNC1)     /* Deletes MYFUNC1.              */.
```

These are the commands you use to receive and install the dummy function, and to delete the SYSMOD entries for the dummy function and the deleted function:

```
SET     BDY(GLOBAL)     /* Set to global zone.      */.
RECEIVE S(DELFUNC)      /* Receive the function.    */.
SET     BDY(TGT1)       /* Set to applicable target. */.
APPLY   S(DELFUNC)      /* Apply to delete old      */
                        /* function.                */.
SET     BDY(DLIB1)      /* Set to applicable DLIB.  */.
ACCEPT  S(DELFUNC)      /* Accept to delete old     */
                        /* function.                */.
SET     BDY(TGT1)       /* Set to applicable target. */.
UCLIN.
DEL     SYSMOD(DELFUNC) /* Delete SYSMOD entries for */.
DEL     SYSMOD(MYFUNC1) /* dummy and old function.  */.
ENDUCL.
SET     BDY(DLIB1)      /* Set to applicable DLIB.  */.
UCLIN.
DEL     SYSMOD(DELFUNC) /* Delete SYSMOD entries for */.
DEL     SYSMOD(MYFUNC1) /* dummy and old function.  */.
ENDUCL                  /*                          */.
```

When you accept the dummy function, SMP/E automatically deletes the associated SYSMOD entry from the global zone and the MCS entry from the SMPPTS.

To complete the cleanup, you should also use the REJECT command to delete any SYSMODs and HOLDDATA applicable to the dummy function and the old function. In addition, you should delete the FMIDs from the GLOBALZONE entry to prevent SMP/E from receiving any SYSMODs or HOLDDATA applicable to either of those functions. Here are examples of the commands you can use to do this.

```
SET     BDY(GLOBAL)      /* Set to global zone.      */.
REJECT HOLDDATA NOFMID   /* Reject SYSMODs, HOLDDATA  */
       PRODUCT           /* PRODUCT information       */
       DELETEFMID        /* for the deleted functions.*/
       (DELFUNC MYFUNC1) /* Delete the FMIDs from the */
                         /* GLOBALZONE entry.         */.
```

Note: If you delete a function that used totally copied libraries, there will be DLIB entries left in the zone after the deletion, which may cause problems if this function is later reinstalled. You should run UCLIN to delete the leftover DLIB entries to ensure that SMP/E will create new DLIB entries that point to the proper libraries when a new copy of the function is installed.

++ZAP MCS

The ++ZAP MCS describes a module update within a PTF, APAR fix, or USERMOD. It must precede the IMASPZAP statements within the SYSMOD.

Syntax

++ZAP

Operands

DALIAS
> is the alias name of a module that has an alias in the distribution library, but not in the target library. This might be used if the module is included under its alias name during system generation.
>
> **Note:** DALIAS is mutually exclusive with TALIAS.

DISTLIB
> specifies the ddname of the distribution library for the specified module.
>
> **Note:** This operand must be specified if the module has not been previously recorded on the target zone or distribution zone. If a MOD entry already exists in the target zone or distribution zone and the value currently in that entry does not match that specified in the DISTLIB operand, the SYSMOD is not applied or accepted.

name
> specifies the name of the module member in the distribution library and, optionally, in the target system library. The name can contain any alphanumeric characters and $, #, @, or hex C0.
>
> **Note:** If the module to be updated has been assembled, specify the name of the assembled module, not the CSECT name.

TALIAS
> identifies all the alias names of a module that has aliases in both the target and distribution libraries.
>
> You can use TALIAS for a module that was copied from a distribution library into a target library (defined by JCLIN data as a copied module), but not for one that is link-edited (defined by JCLIN data as a link-edited module). TALIAS must be specified on the ++MOD MCS even if ALIAS was specified on the COPY SELECT statement.
>
> **Note:** TALIAS is mutually exclusive with DALIAS.

Usage notes

- If a SYSMOD containing a ++ZAP statement attempts to change the ownership (FMID) of the element (with the VERSION operand), the SYSMOD cannot be installed.
- The changes for the module must immediately follow the ++ZAP MCS and must not contain any records that start with ++.
- The only IMASPZAP control statements allowed in a SYSMOD are:

Control statements

ABSDUMP	DUMP	NAME	VER
ABSDUMPT	DUMPT	REP	VERIFY
BASE	IDRDATA	SETSSI	*(comment)

- An EXPAND control statement in link-edit utility format can be placed within IMASPZAP input to allow lengthening of control sections. The EXPAND statement must follow the NAME statement. For the syntax and description of the EXPAND statement, see *z/OS MVS Program Management: User's Guide and Reference*. All control statements must begin in or after column 2.
- Expand type IMASPZAP processing cannot be performed against a noneditable (NE) module.
- Any SETSSI statements placed in the input stream for expand-type IMASPZAP processing must be in a form acceptable to both IMASPZAP and the link-edit utility; that is, they must begin in column 2 or after. The SSI statements must follow the EXPAND statements.
- The name specified on the ++ZAP MCS must be the same as the name of the distribution library module. The CSECT value on the IMASPZAP NAME statement must be the same as the load module's CSECT name. That CSECT name is usually the same as the distribution library name, but it can be different. For example, if the module to be updated has been assembled, the ++ZAP statement should specify the name of the assembled module, not the CSECT name.

The `LIST LMOD` statement produces a target zone listing of link-edit utility control statements that might have changed the CSECT name of the member. A link-edit map may be helpful in other cases where the names differ.

- The IMASPZAP NAME statement can optionally be coded as follows:

```
NAME   csect-name

or

NAME   lmod-name   csect-name

or

NAME   lmod-name   csect-name   class-name
```

The coding of one operand assumes that operand to be a CSECT name for the module referred to in the ++ZAP statement. In this case, all load modules containing the module named in the ++ZAP statement are processed by IMASPZAP.

Two operands can be specified, in which case the second operand is assumed to be a CSECT name. The first operand is assumed to be a valid load module containing the module named in the ++ZAP statement. In this case, only the indicated load module is processed by IMASPZAP.

Three operands can also be specified, in which case the first two operands are as described in the previous paragraphs. The third operand indicates, for program objects only, the class of text that you want to modify. The default is B_text. If you want to omit the CSECT name and supply a class-name, code a single asterisk for the CSECT name, followed by the class-name.

- When using IMASPZAP on an assembled module, be careful: The modification identifier is updated, but not the modification of any associated macros.

It is not recommended that you use IMASPZAP to modify assembled modules. An assembled module modified by IMASPZAP does not cause updating of the distribution library during accept processing; therefore, a subsequently generated system does not contain the IMASPZAP modification.

A more satisfactory method of updating assembled modules is to update the macros that generate them.

- SMP/E processing does not save a backup copy of the nucleus during apply processing when the nucleus is modified by a SYSMOD containing a non-expand-type IMASPZAP modification.

- Only one ZAP can be applied to a module by a single APPLY command. If you need to install several ZAPs for a given module, each one must be packaged separately and installed by a separate APPLY command.

- If you need to install a ++ZAP change on a live system library and that library uses library lookaside (LLA), virtual lookaside facility (VLF), or both, you must do some setup work before installing the change. Otherwise, the change might not take effect, even after installation. See *z/OS MVS Initialization and Tuning Guide* for guidance on:

  - Removing libraries from LLA and VLF control
  - Refreshing LLA and VLF

  Note: Installing a change on a live system is **not** recommended.

Examples

The examples in this section are based on the load module structure shown in .

```
 Load Module Name     Load Module Name
   Module Name          Module Name
     CSECT Name           CSECT Name

 LMODA                LMODB
   MOD1                 MOD1
     CSECT1               CSECT1
     CSECT2               CSECT2
     CSECT3               CSECT3
   MOD2                 MOD2
     MOD2                 MOD2
   MOD3
     MOD3
   MOD4
     MOD4
```

Figure 2. Load module structure for ++ZAP examples

The following examples show how ++ZAP statements can be used to update modules, load modules, and CSECTs within modules:

Example 1: Changing all load modules that contain the same module

Assume you want to change CSECT2 in module MOD1, which is in both LMODA and LMODB. Here is an example of a SYSMOD that accomplishes this by specifying the CSECT name on the NAME statement without including any load module names:

```
++USERMOD(MYMOD01).
++VER(Z038) FMID(FXY1000).
++ZAP(MOD1).
  NAME CSECT2
  VER 000D FF4160
  REP 000D FE4160
```

Example 2: Changing the only load module that contains a given module

Assume you want to change CSECT MOD3 in module MOD3, which is only in LMODA. Here is an example of a SYSMOD that accomplishes this by specifying the CSECT name on the NAME statement without including any load module names:

```
++USERMOD(MYMOD02).
++VER(Z038) FMID(FXY1000).
++ZAP(MOD3).
  NAME MOD3
  VER 000A 00
  REP 000A FF
```

Example 3: Changing one of several load modules that contain a given module

Assume you want to change CSECT2 in module MOD1, which is in both LMODA and LMODB. You want to update only the version in LMODB. Here is an example of a SYSMOD that accomplishes this by specifying both the load module name and the CSECT name on the NAME statement.

```
++USERMOD(MYMOD03).
++VER(Z038) FMID(FXY1000).
++ZAP(MOD1).
  NAME LMODB CSECT2
  VER 0000 00
  REP 0000 FF
```

Example 4: Expanding a module

Assume you want to update CSECT3 with an EXPAND request. CSECT3 is in module MOD1, which is in both LMODA and LMODB. Here is an example of a SYSMOD that accomplishes this by specifying the CSECT name on the NAME statement and on an EXPAND statement.

```
++USERMOD(MYMOD04).
++VER(Z038) FMID(FXY1000).
++ZAP(MOD1).
  NAME CSECT3
  VER 000D FF
  REP 000D FE
  EXPAND CSECT3(4)
```

Chapter 3. Defining control statements in SMPPARM members

The SMPPARM data set can contain members that allow you to customize SMP/E as follows:

Member name
    Use

GIMDDALC
    Defines data sets to be dynamically allocated

GIMEXITS
    Defines exit routines

GIMOPCDE
    Defines macro and assembler operation codes

GIMDDALC control statements

SMPPARM member GIMDDALC is used to specify data sets that are to be dynamically allocated by SMP/E. The following types of data sets may be specified in GIMDDALC:

- Data sets to be allocated to a SYSOUT class (or to the terminal for foreground execution)
- Data sets to be allocated as temporary data sets
- SMPTLIB data sets

During any attempt to allocate a data set, SMP/E first looks for a DD statement specified in the job. If no DD statement was specified, SMP/E then looks for a DDDEF entry in the zone. If no DDDEF entry was found, SMP/E then checks to see if there is an SMPPARM data set. If so, and there is a GIMDDALC member within it, SMP/E looks for a GIMDDALC control statement for the corresponding ddname.

Syntax

SYSOUT data sets

SMPWRKn and SYSUTn data sets

SMPTLIB data sets

Operands

DD
identifies the ddname of the data set to be allocated.

- *ddname* can be from 1 to 8 alphanumeric (A through Z, and 0 through 9), or national characters (@,#, or $), and must start with an alphabetic or national character.
- The DD operand must be the first operand specified on a GIMDDALC control statement.

BLOCK(*size*), CYLINDERS, or TRACKS
specifies the space units to be used in allocating the data set: blocks, cylinders, or tracks. *size* is the size of each block to be allocated, and can be from 1 to 5 decimal digits. To specify the number of these units, use the SPACE operand.

BLOCK can be specified as BLK, CYLINDERS can be specified as CYL, and TRACKS can be specified as TRK.

DATACLAS
specifies the name of a data class to be used for allocating a data set managed by SMS.

The data class can be from 1 to 8 alphanumeric characters (A through Z and 0 through 9) or national characters (@, #, $).

DIR
specifies the number of directory blocks to allocate.

The number can contain from 1 to 4 decimal digits.

DSNTYPE
specifies the type of partitioned data set to be created.

LIBRARY
indicates a PDSE is to be created.

PDS
indicates a PDS is to be created.

MGMTCLAS
specifies the name of a management class to be used for allocating a data set managed by SMS.

The management class can be from 1 to 8 alphanumeric characters (A through Z and 0 through 9) or national characters (@, #, $).

SPACE
specifies the primary and secondary space allocation for a data set. To specify the units for the space allocation, use one of the BLOCK, CYLINDERS, or TRACKS operands.

- The *primary* and *secondary* values can contain from 1 to 4 decimal digits, and the two values must be separated by a comma or a blank.
- The units of *primary* and *secondary* space allocation for SMPTLIB data sets is always tracks.

STORCLAS
specifies the name of a storage class to be used for allocating a data set managed by SMS.

The storage class can be from 1 to 8 alphanumeric characters (A through Z and 0 through 9) or national characters (@, #, $).

SYSOUT

specifies the output class for SYSOUT data sets.

- The *class* value must be 1 alphabetic or numeric character (A through Z and 0 through 9).
- An asterisk (*) indicates the class from the OUTPUT DD statement is to be used. If no OUTPUT DD statement is found, the system default message class is used.
- TERM indicates if SMP/E is executing in the foreground, the data set is to be allocated to the terminal.
- SYSOUT is mutually exclusive with all other operands except DD.
- You can not specify SYSOUT for SMPTLIB data sets.

UNIT

specifies the UNIT type the data set is to reside on.

UNIT can be any nonblank 1 to 8 character string.

VOLUME

specifies the volume serial number of the volume the data set is to reside on.

The volume serial number can be from 1 to 6 alphanumeric characters (A through Z and 0 through 9).

Syntax notes

1. GIMDDALC control statements must each start on a new line.
2. GIMDDALC control statements may be continued on more than one line. SMP/E assumes a statement is continued if it did not find a period (.) before column 73.
3. SMP/E ignores columns 73 through 80. If data, such as a period, is specified beyond column 72, SMP/E will ignore it and indicate an error in the control statement following the one containing the data.
4. GIMDDALC control statements may contain comments. Comments start with /* and end with */. The first */ encountered after the initial /* will end the comment. A comment can appear anywhere within a statement. A comment after the ending period must start on the same line as the period, and you cannot start any additional comments after that final comment. For example, you can code a comment like this:

```
   DD(SMPTLIB) SPACE(5,5) DIR(20).  /* Comment after period continued
                                on subsequent lines is OK.    */
```

However, you cannot code a comment like this:

```
   DD(SMPTLIB) SPACE(5,5) DIR(20).  /* Comment after period OK,      */
                                /* but starting another comment
                                   causes a syntax error.       */
```

This causes a syntax error at the start of the second comment after the period.

5. Only one GIMDDALC control statement is allowed for a single data set.
6. At least one operand (other than DD) must be specified on each GIMDDALC control statement.
7. For SMPTLIB, only the SPACE, TRACKS, and DIR operands are allowed.

Sample GIMDDALC member

The following is a sample GIMDDALC member. The member is named GIMDDALC and is shipped with SMP/E in the SAMPLIB target library.

```
   /*
   **************************************************************************
   *                                                                       *
   *  Licensed Materials - Property of IBM                                 *
   *  5650-ZOS 5655-G44                                                    *
   *  (C) Copyright IBM Corp. 1999, 2013                                   *
   *                                                                       *
   **************************************************************************
```

GIMDDALC statements

```
*                                                                   *
*   GIMDDALC -- Sample GIMDDALC member for the SMPPARM data set     *
*                                                                   *
*********************************************************************

  This sample member is a model to assist the SMP/E user in constructing
  a GIMDDALC member specifically for their use.  A user should copy this
  sample into the GIMDDALC member of a data set.  That data set is then
  defined to SMP/E using the SMPPARM ddname either using a DD statement
  in the SMP/E job, or using a DDDEF entry.

  The GIMDDALC member contains information used to dynamically allocate
  three kinds of data sets:

    1. Data sets to be allocated to a SYSOUT class (or to the terminal
       for foreground execution),
    2. Data sets to be allocated as temporary data sets, and
    3. SMPTLIB data sets.

  During any attempt to allocate a ddname, SMP/E will first look for a
  DD statement specified in the job.  If no DD statement was specified,
  SMP/E will look for a DDDEF entry in the zone.  If no DDDEF entry was
  found, SMP/E will then go to member GIMDDALC in the SMPPARM data set.
  If there is an SMPPARM data set, and there is a GIMDDALC member within
  it, SMP/E will look for an appropriate control statement.

  NOTE: The allocation values in the control statements of this sample
        member are not intended to be recommendations for the actual
        values you should use and will not be appropriate for all users
        and all environments.

  See the "SMP/E Reference" for details on the syntax of the control
  statements and on the use of member GIMDDALC.
```

```
*********************************************************************
*                                                                   *
* CHANGE_ACTIVITY                                                    *
*                                                                   *
* ----------------------------------------------------------------- *
* FLAG REASON  RELEASE   DATE   ORG  DESCRIPTION                     *
* ---- ------- ------- -------- --- ------------------------------- *
* $L0= ACE0029  29.00  03/12/99 KJQ: ALTERNATE CUSTOMIZATION         *
*                                                                   *
*********************************************************************
*    CHANGE FLAG KEY: M@PNC                                          *
*    WHERE:                                                          *
*                                                                   *
* M= MULTIPLE    P= DCRS = D,E,F,G    N= 0-9,A-Z    C= A = ADD       *
*    FOR DELETE     RMPS = H,I,J,K    (USE '0' ONLY   C = CHANGE     *
*    FLAGS ONLY     LINE = L,M,N,O     WHEN ELEMENT   P = COPY       *
*                   PTMS = P,Q,R,S     IS BEING       M = MOVE       *
*                   APARS = 0-9        CREATED)       D = DELETE     *
*********************************************************************

    DEFINE ALLOCATIONS FOR SYSOUT DATA SETS
                                              */
DD(SMPOUT)   SYSOUT(2,TERM).
DD(SMPLIST)  SYSOUT(2).
DD(SMPRPT)   SYSOUT(2).
DD(SMPSNAP)  SYSOUT(2).
DD(SYSPRINT) SYSOUT(*).
DD(LNKPRINT) SYSOUT(*) /* LNKPRINT IS DEFINED IN THE LKED UTILITY ENTRY
                          AND WILL BE USED FOR LINK-EDIT OUTPUT      */.
DD(CPYPRINT) SYSOUT(*) /* CPYPRINT IS DEFINED IN THE COPY UTILITY ENTRY
                          AND WILL BE USED FOR IEBCOPY OUTPUT        */.
DD(SYSUDUMP) SYSOUT(2).
DD(SMPPUNCH) SYSOUT(B).
DD(SMPDEBUG) SYSOUT(2). /*

    DEFINE ALLOCATIONS FOR TEMPORARY DATA SETS
                                              */
DD(SMPWRK1)  BLOCK(3120) SPACE(364,380) DIR(111) UNIT(SYSALLDA).
DD(SMPWRK2)  BLOCK(3120) SPACE(364,380) DIR(111) UNIT(SYSALLDA).
DD(SMPWRK3)  BLOCK(3120) SPACE(364,380) DIR(111) UNIT(SYSALLDA).
DD(SMPWRK4)  BLOCK(3120) SPACE(364,380) DIR(111) UNIT(SYSALLDA).
DD(SMPWRK6)  BLOCK(3120) SPACE(364,380) DIR(111) UNIT(SYSALLDA).
DD(SYSUT1)   BLOCK(3120) SPACE(380,760)          UNIT(SYSALLDA).
DD(SYSUT2)   BLOCK(3120) SPACE(380,760)          UNIT(SYSALLDA).
DD(SYSUT3)   BLOCK(3120) SPACE(380,760)          UNIT(SYSALLDA).
DD(SYSUT4)   TRACKS      SPACE(1,1)              UNIT(SYSALLDA).
DD(SYSPUNCH) TRACKS      SPACE(25,10)   DIR(10)  UNIT(SYSALLDA).
```

```
DD(SMPTLOAD) TRACKS      SPACE(50,20)  DIR(16)  UNIT(SYSALLDA). /*

   DEFINE ALLOCATIONS FOR SMPTLIB DATA SETS.  THIS SAMPLE SMPTLIB
   ALLOCATION CORRESPONDS TO SMP/E'S DEFAULT SMPTLIB ALLOCATION.
                                                              */
DD(SMPTLIB)  TRACKS SPACE(0,0) DIR(0).
```

GIMEXITS control statements

SMP/E provides support for two exit points in SMP/E command processing. You can write exit routine programs that receive control at those exit points to:

- process statements from SMPPTFIN during the RECEIVE command, and
- control retry processing when data sets run out of space during ACCEPT, APPLY, GZONEMERGE, LINK LMODS, LINK MODULE, RECEIVE, or RESTORE processing.

The GIMEXITS member of the SMPPARM data set contains control statements used to identify the exit routine programs to be given control at those SMP/E exit points. It is the presence of the control statement that indicates to SMP/E that an exit routine program is to be called at the appropriate exit point during SMP/E processing. The statements not only identify the exit routine program, but also allow you to identify the data set that contains the exit routine program. This eliminates the need to copy the exit routines into the new MIGLIB or LINKLIB data sets; the exit routines can reside in almost any authorized data set you desire.

Syntax

GIMEXITS control statement

```
►►─ EXIT ─ ( ─┬─ RECEIVE ─┬─ ) ─ MODNAME ─ ( ─ name ─ ) ─►
              └─ RETRY ────┘

   ►─┬─────────────────────────────┬─ . ─►◄
     └─ DATASET ─ ( ─ dataset ─ ) ─┘
```

Operands

EXIT

identifies the SMP/E exit point that will pass control to the exit routine defined on the MODNAME operand of this control statement. Valid values are:

RECEIVE

indicates the RECEIVE command exit point will pass control to the exit routine.

RETRY

indicates the RETRY exit point will pass control to the exit routine.

The EXIT operand must be the first operand specified on a GIMEXITS control statement.

MODNAME

specifies the name of the exit routine to receive control during SMP/E execution at the identified exit point. The exit routine must be a load module residing in an authorized library. If the DATASET operand is not specified, SMP/E uses the system search order for programs to locate the exit routine. See the topic "The Search Order the System uses for Programs" in *z/OS MVS Initialization and Tuning Guide* for more information.

The name can be from 1 to 8 alphanumeric characters (A through Z and 0 through 9) or national characters (@, #, $). MODNAME is a required operand.

DATASET

specifies the data set name of a load library in which the exit routine resides. The data set must be cataloged and it must be an authorized library. If the DATASET operand is not specified, SMP/E uses the system search order for programs to locate the exit routine. See the topic "The Search Order the System uses for Programs" in *z/OS MVS Initialization and Tuning Guide* for more information.

The data set name must conform to standard naming conventions for data sets. Each part of the name must contain from 1 to 8 alphanumeric characters (A through Z and 0 through 9) or dash (-), separated from the other parts by a period (.), with no intervening blanks. The first character of each part must not be numeric or dash, and the maximum length of the entire name is 44 characters (including the periods).

Syntax notes

1. GIMEXITS control statements must each start on a new line.
2. GIMEXITS control statements may be continued on more than one line. SMP/E assumes a statement is continued if it did not find a period (.) before column 73.
3. SMP/E ignores columns 73 through 80. If data, such as a period, is specified beyond column 72, SMP/E will ignore it and indicate an error in the control statement following the one containing the data.
4. GIMEXITS control statements may contain comments. Comments start with /* and end with */. The first */ encountered after the initial /* will end the comment. A comment can appear anywhere within a statement. A comment after the ending period must start on the same line as the period, and you cannot start any additional comments after that final comment. For example, you can code a comment like this:

   ```
   EXIT(RECEIVE) MODNAME(MYRECEX).  /* Comment after period continued
                                       on subsequent lines is OK.    */
   ```

   However, you cannot code a comment like this:

   ```
   EXIT(RECEIVE) MODNAME(MYRECEX).  /* Comment after period OK,      */
                                    /* but starting another comment
                                       causes a syntax error.        */
   ```

   This causes a syntax error at the start of the second comment after the period.
5. Only one GIMEXITS control statement is allowed for a single exit point.

Sample GIMEXITS member

The following is a sample GIMEXITS member. The member is named named GIMEXITS and is shipped with SMP/E in the SAMPLIB target library.

```
/*
*************************************************************************
*                                                                       *
*  Licensed Materials - Property of IBM                                 *
*  5650-ZOS 5655-G44                                                    *
*  (C) Copyright IBM Corp. 1999, 2013                                   *
*                                                                       *
*************************************************************************
*                                                                       *
*  GIMEXITS -- Sample GIMEXITS member for the SMPPARM data set          *
*                                                                       *
*************************************************************************

  This sample member is a model to assist the SMP/E user in constructing
  a GIMEXITS member specifically for their use.  A user should copy this
  sample into the GIMEXITS member of a data set.  That data set is then
  defined to SMP/E using the SMPPARM ddname either using a DD statement
  in the SMP/E job, or using a DDDEF entry.

  The GIMEXITS member contains information used to identify the exit
  routines which are to get control during SMP/E processing at specific
  exit points:

  RECEIVE  The RECEIVE exit point allows you to scan statements in the
           SMPPTFIN data set during RECEIVE command execution.

  RETRY    The RETRY exit point enables you to control SMP/E's RETRY
           processing when a data set runs out of space during APPLY,
           ACCEPT, RESTORE, and LINK command execution.

  See the "SMP/E Reference for details on the syntax of the control
  statements and on the use of member GIMEXITS.

*************************************************************************
*                                                                       *
* CHANGE_ACTIVITY                                                        *
*                                                                       *
* --------------------------------------------------------------------- *
* FLAG REASON  RELEASE   DATE   ORG  DESCRIPTION                         *
* ---- ------- ------- -------- --- ----------------------------------   *
* $L0= ACE2900  29.00  03/12/99 KJQ: ALTERNATE CUSTOMIZATION             *
*                                                                       *
*************************************************************************
*    CHANGE FLAG KEY: M@PNC                                              *
*    WHERE:                                                              *
*                                                                       *
* M= MULTIPLE    P= DCRS = D,E,F,G    N= 0-9,A-Z    C= A = ADD           *
*    FOR DELETE     RMPS = H,I,J,K    (USE '0' ONLY    C = CHANGE        *
*    FLAGS ONLY     LINE = L,M,N,O     WHEN ELEMENT    P = COPY          *
*                   PTMS = P,Q,R,S    IS BEING        M = MOVE           *
*                   APARS = 0-9       CREATED)        D = DELETE         *
*************************************************************************
  */
EXIT(RECEIVE) MODNAME(MYRECEX) DATASET(SMPE.EXITS.LOAD). /*

          Define exit routine MYRECEX in data set SMPE.EXITS.LOAD
          to get control during SMP/E RECEIVE command processing.
                                                                    */
EXIT(RETRY)   MODNAME(MYRTYEX) DATASET(SMPE.EXITS.LOAD). /*

          Define exit routine MYRTYEX in data set SMPE.EXITS.LOAD
          to get control during SMP/E RETRY processing.
                                                                    */
```

Figure 3. Sample GIMEXITS member provided in SAMPLIB

OPCODE control statements

During JCLIN processing, SMP/E scans the SMPJCLIN input, which consists of various job steps calling system utilities. To determine the structure of the target system, SMP/E specifically looks at copy steps, link-edit steps, and assembly steps.

OPCODE statements

As SMP/E scans inline assembly steps, it looks at each assembler instruction to determine whether the instruction is a macro invocation or an OPCODE. If SMP/E determines that this is a macro invocation, it builds a MAC entry in the target zone and defines the connection between the macro and the assembly in which the macro was found. As a result, when that macro is later changed by the installation of a SYSMOD, SMP/E can cause all the assemblies that used that macro to be redone.

SMP/E uses GIMOPCDE members to determine whether an assembler instruction is a macro invocation or an OPCODE.

SMP/E provides you with a default set of OPCODE definitions that identifies all the known standard assembler OPCODEs. You can define your own members to identify additional OPCODE values or macro names by using the sample GIMOPCDE member SMP/E supplied for you. If you define your own member, you must allocate the SMPPARM data set. You can specify the name of one of these user-defined GIMOPCDE members on the JCLIN command or on the ++JCLIN statement to have SMP/E pick up this additional information. The data from this other GIMOPCDE member is merged with the data from the default set of OPCODE definitions. If duplicate data is specified, the user-specified member has priority.

SMP/E uses the following method to determine whether an assembler instruction is an OPCODE or a macro:

- SMP/E looks for a user-allocated SMPPARM data set.
- If SMPPARM is **not** found, SMP/E uses the default set of OPCODE definitions.

  If SMPPARM is found and there is a user-defined GIMOPCDE member specified on the JCLIN statement or in the ++JCLIN MCS, then SMP/E searches for the specified member in SMPPARM. If it finds that member, it will look first in that member for a definition of the character string.

- If the user-defined GIMOPCDE member specified is not found, SMP/E searches SMPPARM for the GIMOPCDE member. If it finds the GIMOPCDE member, it will look **only** in that member for a definition of the character string. (The GIMOPCDE member, if it exists in SMPPARM, will completely override SMP/E's default set of OPCODE definitions.)
- If the GIMOPCDE member is not found, SMP/E uses the default set of OPCODE definitions.
- If the character string is not defined in either the SMPPARM data set or the default set, SMP/E considers it to be a macro.

For additional information about how SMP/E processes assembler input as JCLIN, see .

Syntax

GIMOPCDE member control statements

```
►► KEY ─ = ─ name ─┬─ TYPE ─ = ─ OPCODE ─┬─ . ►◄
                   └─ TYPE ─ = ─ MACRO ──┘
```

Operands

KEY
identifies a character string for SMP/E to check for when scanning assembler entries during JCLIN processing, either during the JCLIN command or when applying a SYSMOD with inline JCLIN.

name can be any alphanumeric character string from one to eight characters in length.

TYPE
specifies how SMP/E treats that character string when it is encountered. The following values may be specified:

OPCODE
specifies that SMP/E treats the specified character string as a valid assembler OPCODE.

Note: OPCODE can also be specified as OP.

MACRO
indicates that SMP/E treats the specified character string as the name of a macro.

Note: MACRO can also be specified as MAC.

Note:

1. The TYPE operand can be specified only once per control statement.
2. If the TYPE operand is not specified, the default is TYPE=OPCODE.

Usage notes

- Either a blank or a comma can be used to separate the operands.
- Comments are permitted, subject to the next restriction.
- An OPCODE control statement **cannot** span multiple lines. The complete control statement, including comments, must be contained on a single line.
- If, during JCLIN processing, SMP/E encounters a string specified in one of the OPCODE control statements, and the OPCODE control statement specifies that the string is to be treated as a macro, SMP/E builds a MAC entry in the target zone and connects that MAC entry to the ASSEM entry built for the current assembly. For a further description of JCLIN processing as related to the GIMOPCDE members, see *z/OS SMP/E Commands*.

Examples

The following examples are provided to help you use OPCODE statements.

Example 1: Defining a new OPCODE for special assemblers

Some licensed programs use a version of the assembler supplied with their product rather than the standard assembler. Often, this assembler recognizes special OPCODEs. When SMP/E is processing JCLIN containing assembler steps that use these OPCODEs, you do not want SMP/E creating target zone macro entries for them. To avoid this situation, create a special GIMOPCDE member containing appropriate control statements to define those OPCODEs to SMP/E.

The following is an example of an GIMOPCDE member containing three control statements to define the three-character strings, *OP1*, *OPCD1*, and *OPCODE01* as special assembler OPCODEs:

```
KEY=OP1         TYPE=OPCODE.
KEY=OPCD1       TYPE=OP.
KEY=OPCODE01    TYPE=OPCODE.
```

Example 2: Overriding an SMP/E-defined OPCODE

You may provide your own GIMOPCDE member to override SMP/E's default set of OPCODE definitions by using the sample GIMOPCDE member provided for you. This sample member contains a control statement identifying *PUNCH* as a character string that should be treated as an assembler OPCODE. You may have constructed a *PUNCH* macro and now want SMP/E to recognize *PUNCH* as a macro name. This can be done by adding the following control statement to an GIMOPCDE member and then specifying that member on the JCLIN command or ++JCLIN statement:

```
KEY=PUNCH       TYPE=MACRO  /* override type to macro  */.
```

Chapter 4. SMP/E data sets and files

This chapter describes the data sets that are used to process SMP/E commands. These data sets are listed by the ddnames for the data sets required by each SMP/E command. According to the type of data set, you can define it with a DDDEF entry, a DD statement, or with GIMDDALC control statements in SMPPARM member GIMDDALC. For more information about DDDEF entries, see "DDDEF entry (distribution, target, and global zone)" on page 163. For more information about SMPPARM member GIMDDALC, see Chapter 3, "Defining control statements in SMPPARM members," on page 103. For more information about dynamic allocation, see *z/OS SMP/E User's Guide*.

CLIENT

ddname
> Specified by the user.

Use
> The CLIENT data set contains information about the local z/OS system such as how to navigate a local FTP firewall and HTTP proxy server, in addition to the location of Java application classes and Java debug options. This information is used by the RECEIVE FROMNETWORK and RECEIVE ORDER commands, as well as the GIMGTPKG service routine. The contents of the CLIENT data set are described in the RECEIVE chapter in *z/OS SMP/E Commands*.

Attributes
> Sequential or member of a partitioned data set; LRECL=80, RECFM=F or FB.

Device
> Direct access.

Note: This data set may reside in a UNIX file system. Specify PATHOPTS(ORDONLY) and either FILEDATA=TEXT or FILEDATA=BINARY on the DD statement if it is to reside in a UNIX file system. For simplicity, FILEDATA=TEXT is preferred, and ensure the newline character (X'15') marks the end of each line. If you use FILEDATA=BINARY, you must ensure that each line is padded with blanks to the 80-byte record length because there is no end of record marker in binary data.

Distribution library (DLIB)

ddname
> The ddname for a distribution library should match the low-level qualifier of the data set name. For example, the ddname for SYS1.AMACLIB should be AMACLIB.

Use
> Distribution libraries (DLIBs) contain updated versions of macros, source, and modules that were shipped by IBM and stored during ACCEPT processing. They are used during system generation to build the target libraries, so you should keep them at a tested functional and service level. They are also used during RESTORE processing to replace elements in the target libraries. You must provide a DDDEF entry or DD statement for each distribution library that is being processed.

Attributes
> Partitioned.

Device
> Direct access only.

INFILE

ddname
> The ddname for an INFILE data set cannot be the same as the ddname for any other data sets used by SMP/E. Other than this, there are no restrictions.

Use
The INFILE data set is a sequential data set containing a zone that is to be loaded by the ZONEIMPORT command. The ddname of this sequential data set is specified on the INFILE operand of the ZONEIMPORT command. The zone in the INFILE data set can be used in two ways:

- To recreate a zone that was destroyed
- To recreate a zone on another CSI data set

You must use the ZONEEXPORT command to create the INFILE data set. (This same data set is called the OUTFILE data set when it is created by the ZONEEXPORT command.)

Attributes
Sequential; no DCB parameters are required.

If the INFILE data set is on tape, there may be more than one volume for an exported data set. Remember to specify all the volume serial numbers on the INFILE DD statement.

Device
Tape or direct access.

Note: BLKSIZE must not exceed 32760.

Link library (LKLIB)

ddname
The ddname for a link library must match the LKLIB value on the element MCS. For example, the ddname for the link library on statement ++MOD(MODA) LKLIB(LIBX) must be LIBX.

Use
Link libraries contain replacements for object modules in link-edited format. They are used when the modules are provided in partitioned data sets rather than inline or in relative files.

Attributes
Partitioned.

Device
Direct access only.

ORDERSERVER

ddname
Specified by the user.

Use
The ORDERSERVER data set contains information about the IBM Automated Delivery Request server used by the RECEIVE ORDER command. The contents of the ORDERSERVER are described in the RECEIVE chapter in *z/OS SMP/E Commands*.

Attributes
Sequential or member of a partitioned data set; LRECL=80, RECFM=F or FB.

Device
Direct access.

Note: This data set may reside in a UNIX file system. Specify PATHOPTS(ORDONLY) and either FILEDATA=TEXT or FILEDATA=BINARY on the DD statement if it is to reside in a UNIX file system. For simplicity, FILEDATA=TEXT is preferred, and ensure the newline character (X'15') marks the end of each line. If you use FILEDATA=BINARY, you must ensure that each line is padded with blanks to the 80-byte record length because there is no end of record marker in binary data.

OUTFILE

ddname
> The ddname for an OUTFILE data set cannot be the same as the ddname for any other data sets used by SMP/E. Other than this, there are no restrictions.

Use
> The OUTFILE data set is a sequential data set containing a zone copied by the ZONEEXPORT command. The ddname of this sequential data set is specified on the OUTFILE operand. The zone in the OUTFILE data set can be used in two ways:
>
> - As a backup copy, to recreate a zone that was destroyed
> - As a transportable copy, to recreate a zone on another CSI data set
>
> You must use the ZONEIMPORT command to process the OUTFILE data set. (This same data set is called the INFILE data set when it is processed by the ZONEIMPORT command.)

Attributes
> Sequential; no DCB parameters are required.

Device
> Tape or direct access.

Note: BLKSIZE must not exceed 32760.

SERVER

ddname
> Specified by the user.

Use
> The SERVER data set contains information about a TCP/IP connected host running an FTP or HTTP(S) server.

Attributes
> Sequential or member of a partitioned data set; LRECL=80, RECFM=F or FB.

Device
> Direct access.

Note: This data set may reside in a UNIX file system. Specify PATHOPTS(ORDONLY) and either FILEDATA=TEXT or FILEDATA=BINARY on the DD statement if it is to reside in a UNIX file system. For simplicity, FILEDATA=TEXT is preferred, and ensure the newline character (X'15') marks the end of each line. If you use FILEDATA=BINARY, you must ensure that each line is padded with blanks to the 80-byte record length because there is no end of record marker in binary data.

SMPCLNT

ddname
> SMPCLNT

Use
> The SMPCLNT data set contains information about the local z/OS system such as how to navigate a local FTP firewall, the location of Java application classes and Java debug options. This information is used by the GIMGTPKG service routine, and has the same contents as the CLIENT data set. The contents of the CLIENT data set are described in the RECEIVE chapter in *z/OS SMP/E Commands*.

Attributes
> Sequential or member of a partitioned data set; LRECL=80, RECFM=F or FB.

Device
> Direct access.

Note: This data set may reside in a UNIX file system. Specify PATHOPTS(ORDONLY) and either FILEDATA=TEXT or FILEDATA=BINARY on the DD statement if it is to reside in a UNIX file system.

For simplicity, FILEDATA=TEXT is preferred, and ensure the newline character (X'15') marks the end of each line. If you use FILEDATA=BINARY, you must ensure that each line is padded with blanks to the 80-byte record length because there is no end of record marker in binary data.

SMPCNTL

ddname
SMPCNTL.

Use
The SMPCNTL data set contains the SMP/E commands to be processed.

Attributes
Sequential; LRECL=80, BLKSIZE=multiple of 80, RECFM=FB.

Device
Card, tape, direct access, or terminal.

Note:

1. BLKSIZE must not exceed 32760.

2. This data set may reside in a UNIX file system. Specify PATHOPTS(ORDONLY) and either FILEDATA=TEXT or FILEDATA=BINARY on the DD statement if it is to reside in a UNIX file system. For simplicity, FILEDATA=TEXT is preferred, and ensure the newline character (X'15') marks the end of each line. If you use FILEDATA=BINARY, you must ensure that each line is padded with blanks to the 80-byte record length because there is no end of record marker in binary data.

SMPCPATH

ddname
SMPCPATH

Use
The SMPCPATH (SMP/E classpath) is a directory in the UNIX file system that contains SMP/E Java application classes. The classes are used by RECEIVE ORDER processing to communicate with the IBM Automated Delivery Request server, and by the RECEIVE command and GIMGTPKG, GIMZIP, and GIMUNZIP service routines to calculate SHA-1 hash values when ICSF is not available. The default location for SMPCPATH is /usr/lpp/smp/classes/.

Attributes
Existing directory in the UNIX file system.

Device
Direct access only.

Note:

1. SMPCPATH can be specified using a DD statement or a DDDEF entry. If a value is not specified, the default path of /usr/lpp/smp/classes/ is used.

2. For additional information about how to specify the location of the Java runtime, see "Options that affect Java," in the "Preparing to use Internet service retrieval" chapter in *z/OS SMP/E User's Guide*.

SMPCSI

ddname
SMPCSI.

Use
The data set specified by the ddname SMPCSI is the CSI containing the global zone. (This data set is also known as the master CSI.) The CSI is the database used by SMP/E to record status and other information about the various target and distribution libraries being supported.

Attributes
VSAM; RECORDSIZE(24 143), KEYS(24 0).

Device

Direct access only.

Note:

1. The low-level qualifier of the data set name must be CSI.

2. If you have used IBM SMP/E for z/OS, V3R6 to update a CSI data set, you might not be able to process that data set with previous releases of SMP/E. For more information, see the migration section in *z/OS SMP/E User's Guide*.

3. When running on systems with the required level of Data Facility Product (DFP), SMP/E automatically takes advantage of the local shared resource (LSR) feature of VSAM. This reduces the number of times SMP/E must access data when it is reading CSI data sets. As a result, SMP/E performance is improved for commands such as APPLY, APPLY CHECK, ACCEPT, ACCEPT CHECK, and especially LIST.

4. CSI data sets should usually be allocated dynamically. However, you may want to use the batch local shared resources (BLSR) subsystem with expanded storage hiperspaces (instead of SMP/E's implementation of LSR) to improve SMP/E performance during APPLY and ACCEPT processing for a large number of changes.

   To do this, you use JCL statements instead of dynamic allocation to define the CSI data sets containing the zones required for processing. For each CSI data set, you need to provide two DD statements:

   • The first DD statement is used by SMP/E to enqueue on the data set (to protect it from simultaneous updates) and to trigger batch LSR. This DD statement uses the DDNAME parameter to point to the second DD statement.

   • The second DD statement is used by batch LSR to open the CSI data set.

   Both DD statements must specify the same DSN value. The following example shows how this can be done:

   ```
   //SMPCSI    DD DSN=dataset1.CSI,DISP=SHR,
   //          SUBSYS=(BLSR,'DDNAME=MYSMPCSI',
   //          'HBUFND=value','HBUFNI=value')
   //MYSMPCSI  DD DSN=dataset1.CSI,DISP=SHR
   //*
   //tgtzone   DD DSN=dataset2.CSI,DISP=SHR,
   //          SUBSYS=(BLSR,'DDNAME=MYTGT1',
   //          'HBUFND=value','HBUFNI=value')
   //MYTGT1    DD DSN=dataset2.CSI,DISP=SHR
   ```

   In this example, note that:

   • If the target zone and global zone exist in the same CSI, *dataset1* and *dataset2* refer to the same data set.

   • HBUFND and HBUFNI are used instead of BUFND and BUFNI to indicate the use of hiperspace in expanded storage (if available).

   • *tgtzone* is the name of the target zone specified on the SET BDY command to be processed.

   • The DSN parameter must be specified on all DD statements to ensure internal enqueue protection.

5. For more information about using the EXEC statement to specify the CSI data set containing the global zone, see *z/OS SMP/E User's Guide*.

6. For information about the target and distribution zones in a CSI data set, see "Zone statement" on page 139.

SMPDATA1

ddname

SMPDATA1.

Use

The SMPDATA1 data set is used to store library change file records. The library change file records are written by SMP/E during APPLY or RESTORE processing when the CHANGEFILE subentry of the

OPTIONS entry is set to YES. For more information about library change file records, see Chapter 8, "Library change file records," on page 359.

Attributes

Sequential; BLKSIZE=1069–32670, RECFM=VB, DISP=MOD.

- To have the system determine the optimal block size for the data set, specify BLKSIZE=0.
- If you specify a non-zero BLKSIZE that is less than 1069, SMP/E uses a default of 8800.
- DISP=MOD must be specified to maintain a cumulative history of SMP/E 'delta' processing.
- If a DISP other than MOD is specified by the user, SMP/E uses it.
- SMP/E uses LRECL=1065 in allocating the SMPDATA1 data set.

Device

Direct access only.

Note:

1. The SMPDATA1 data set can only be defined to SMP/E with a DD statement or a DDDEF.
2. The SMPDATA1 DD statement or DDDEF must be defined to each target zone to enable 'delta' processing to occur for that target zone.
3. Do not concatenate SMPDATA1 data sets.
4. Do not allocate the SMPDATA1 data set as a path in a UNIX file system.

The SMPDATA1 DDDEF can also be defined with the SMP/E Administration dialogs.

SMPDATA2

ddname

SMPDATA2.

Use

The SMPDATA2 data set is used to store the library change file records when the SMPDATA1 data set becomes full. This type of processing is called *spill processing*. The library change records are written by SMP/E during APPLY or RESTORE processing. For more information about library change file records, see Chapter 8, "Library change file records," on page 359.

Attributes

Sequential; BLKSIZE=1069–32760, RECFM=VB, DISP=MOD.

- To have the system determine the optimal block size for the data set, specify BLKSIZE=0.
- If you specify a non-zero BLKSIZE that is less than 1069, SMP/E uses a default of 8800.
- DISP=MOD must be specified.
- If a DISP other than MOD is specified, SMP/E will use it.
- SMP/E will use LRECL=1065 in allocating the SMPDATA2 data set.

Device

Direct access only.

Note:

1. The SMPDATA2 data set can only be defined to SMP/E with a DD statement or a DDDEF.
2. The SMPDATA2 DD statement or DDDEF must be defined to each target zone to enable spill processing to occur for that target zone.
3. Do not concatenate SMPDATA2 data sets.
4. Do not allocate the SMPDATA2 data set as a path in a UNIX file system.
5. The size of the SMPDATA2 data set will vary depending on how the user intends to manage it.
6. As is the case with SMPLOG data set, the user is responsible for managing the contents and space.

The SMPDATA2 DDDEF can also be defined with the SMP/E Administration dialogs.

SMPDEBUG

ddname
SMPDEBUG.

Use
The SMPDEBUG data set contains a dump that was requested by the DEBUG command. Depending on the operands specified, it may contain (1) a dump of SMP/E control blocks and storage areas associated with the specified dump points, or (2) a dump of the VSAM RPL control block and additional RPL information for the specified SMP/E function.

Attributes
Sequential; LRECL=121, BLKSIZE=multiple of 121, RECFM=FBA, DISP=MOD.

Device
SYSOUT, printer, direct access, tape, or terminal.

Note:

1. BLKSIZE must not exceed 32760.

2. This data set may reside in a UNIX file system. Specify FILEDATA=TEXT and PATHOPTS(OWRONLY,OAPPEND) on the DD statement for this data set, if it is to reside in a UNIX file system.

SMPDIR

ddname
SMPDIR.

Use
The SMPDIR directory identifies the name of a path to be allocated in a UNIX file system. The name is not a complete pathname; it is the directory under which the files produced by GIMZIP are stored.

Attributes
Existing directory within a UNIX file system.

Device
Direct access.

Note:

1. The SMPDIR directory must be defined with a DD statement.

2. The SMPDIR directory must be defined as a directory in a UNIX file system.

3. The size of the SMPDIR directory depends on the size of the package files produced by GIMZIP.

SMPDUMMY

ddname
SMPDUMMY

Use
SMPDUMMY is intended for use as the definition side deck library for a load module when the IMPORT statements associated with a DLL are not wanted. It allows the SYSDEFSD DD statement to be allocated as a DUMMY data set.

When any of the following SYSDEFSD DD statements is encountered in the link edit JCLIN input stream:

```
//SYSDEFSD DD DSN=SMPDUMMY,DISP=xxx
     (where SMPDUMMY may be the lowest-level qualifier of
      a multi-qualifier data set name)
          -or-
//SYSDEFSD DD DSN=NULLFILE
     (where NULLFILE must be a single-word parameter)
```

```
            -or-
  //SYSDEFSD DD DUMMY
```

the SIDE DECK LIBRARY subentry for the load module will be set to SMPDUMMY. When needed for processing, SMP/E will dynamically allocate SMPDUMMY as a DUMMY data set. Any existing SMPDUMMY DDDEF entry or GIMDDALC statement will be ignored. If SMPDUMMY was previously allocated outside of SMP/E, SMP/E will free the SMPDUMMY DD and reallocate it as a DUMMY data set.

Attributes
DUMMY.

Device
None.

Note: The SMPDUMMY ddname is intended only for association with the SIDE DECK LIBRARY subentry for load modules. SMP/E does not prohibit its specification for other library ddnames with the various input methods (MCS, UCLIN, JCLIN, and so on). However, SMPDUMMY will always be allocated by SMP/E as "DD DUMMY". Therefore, the results are unpredictable, in that they depend on each utility's ability to handle a DUMMY data set.

SMPHOLD

ddname
SMPHOLD.

This may refer to an actual data set, or it may refer to a file on a tape (such as file 4 on an ESO tape).

Use
SMPHOLD contains ++HOLD and ++RELEASE statements to be processed by the RECEIVE command.

Attributes
Sequential; LRECL=80, BLKSIZE=multiple of 80, RECFM=FB.

Device
Direct access or tape.

Note:

1. BLKSIZE must not exceed 32760.

2. This data set may reside in a UNIX file system. Specify PATHOPTS(ORDONLY) and either FILEDATA=TEXT or FILEDATA=BINARY on the DD statement if it is to reside in a UNIX file system. For simplicity, FILEDATA=TEXT is preferred, and ensure the newline character (X'15') marks the end of each line. If you use FILEDATA=BINARY, you must ensure that each line is padded with blanks to the 80-byte record length because there is no end of record marker in binary data.

SMPHRPT

ddname
SMPHRPT

Use
The SMPHRPT data set is an alternate report data set. If SMPHRPT is allocated, the HOLD reports that are generated during RECEIVE, APPLY, ACCEPT and REPORT SYSMODS processing are directed to this data set while other reports are directed to SMPRPT.

Attributes
Sequential; LRECL=121, BLKSIZE=multiple of 121, RECFM=FBA, DISP=MOD.

Device
SYSOUT, printer, direct access, tape, or terminal.

Note:

1. BLKSIZE cannot exceed 32760.

2. If SMPHRPT is not allocated, all report output goes to the SMPRPT data set. If SMPRPT is not allocated, all report output goes to the SMPOUT data set.

3. If SMPHRPT is allocated to a data set, the disposition must be MOD, because SMP/E opens and closes the SMPHRPT DD statement at each SET command. If the disposition is SHR or OLD, SMPHRPT contains only the reports from the last set of commands that is processed before the end of SMP/E processing.

4. The SMPHRPT ddname can be allocated to a file in the UNIX file system. Specify FILEDATA=TEXT and PATHOPTS(OWRONLY, OAPPEND) on the DD statement if the SMPHRPT ddname is allocated to a file in the UNIX file system. If OAPPEND is not specified, SMPHRPT contains only the reports from the last set of commands that is processed before the end of SMP/E processing.

SMPJCLIN

ddname
   SMPJCLIN.

Use
   The SMPJCLIN data set contains a job stream of assembly, link-edit, and copy job steps. This data is typically the stage 1 output from the most recent full or partial system generation, but it may be other data in a similar format, such as output from the SMP/E GENERATE command. This job stream is used as input to the JCLIN command to update or create entries in a target zone.

Attributes
   Sequential; LRECL=80, BLKSIZE=multiple of 80, RECFM=FB.

Device
   Card, tape, direct access, or terminal.

Note:

1. BLKSIZE must not exceed 32760.

2. This data set may reside in a UNIX file system. Specify PATHOPTS(ORDONLY) and either FILEDATA=TEXT or FILEDATA=BINARY on the DD statement if it is to reside in a UNIX file system. For simplicity, FILEDATA=TEXT is preferred, and ensure the newline character (X'15') marks the end of each line. If you use FILEDATA=BINARY, you must ensure that each line is padded with blanks to the 80-byte record length because there is no end of record marker in binary data.

SMPJHOME

ddname
   SMPJHOME

Use
   SMPJHOME (SMP/E JAVAHOME) is a directory in the UNIX file system that contains the Java runtime. SMP/E uses Java classes during RECEIVE ORDER processing to communicate with the IBM Automated Delivery Request server, and by the RECEIVE command and GIMGTPKG, GIMZIP, and GIMUNZIP service routines to calculate SHA-1 hash values when ICSF is not available.

   SMP/E requires the Java runtime directory when either of the following conditions applies.

   • When SMP/E is installing ++JARUPD elements, or

   • When a UNIX shell script is invoked during the installation of a file system element, and the shell script issues a Java command.

   No default location is assumed by SMP/E for SMPJHOME.

Attributes
   Existing directory in the UNIX file system.

Device
   Direct access only.

Note:

1. SMPJHOME can be specified using a DD statement or a DDDEF entry.

2. For additional information about how to specify the location of the SMP/E application classes, see "Options that affect Java," in the "Preparing to use internet service retrieval" chapter in *z/OS SMP/E User's Guide*.

SMPLIST

ddname
SMPLIST.

Use
The SMPLIST data set contains the output of all LIST commands.

Attributes
Sequential; LRECL=121, BLKSIZE=multiple of 121, RECFM=FBA, DISP=MOD.

Device
SYSOUT, printer, direct access, tape, or terminal.

Note:

1. BLKSIZE must not exceed 32760.

2. If SMPLIST is not defined, all LIST output goes to the SMPOUT data set.

3. If SMPLIST is allocated to a data set, the disposition must be MOD, because SMP/E opens and closes the SMPLIST DD statement at each SET command. If the disposition is SHR or OLD, SMPLIST contains only the output from the last set of commands processed before the end of SMP/E processing.

4. This data set may reside in a UNIX file system. Specify FILEDATA=TEXT and PATHOPTS(OWRONLY,OAPPEND) on the DD statement for this data set, if it is to reside in a UNIX file system. If OAPPEND is not specified, SMPLIST contains only the output from the last set of commands processed before the end of SMP/E processing.

SMPLOG

ddname
SMPLOG.

Use
The SMPLOG data set (LOG) contains time-stamped records of SMP/E processing. The records in this data set can be written automatically by SMP/E or added by the user through the LOG command. The data set also contains messages issued by SMP/E, as well as detailed information about data set allocation.

Attributes
Sequential; BLKSIZE=514-32000, RECFM=VB, DISP=MOD.

- The BLKSIZE value determines how many records are written to the LOG at one time. As a block is filled, it is written to SMPLOG.

  If BLKSIZE is less than 514, SMP/E uses the default of 3200.

- DISP=MOD must be specified to maintain a cumulative history of SMP/E processing.

- SMP/E uses LRECL=510 in allocating the SMPLOG data set.

Device
Tape or direct access.

Note:

1. **Important:** If the data set allocated to the SMPLOG DD statement has the wrong data set attributes, SMP/E will open the data set for output processing with acceptable attributes. This action will overlay the contents of the data set.

2. Each zone should have its own SMPLOG data set.

3. A secondary log data set (SMPLOGA) should be defined to hold log data when the SMPLOG data set is full. Otherwise, the extra log data is written to the SMPOUT data set, with the date and time stamp encrypted.

4. SMPLOG should be updated only by the LOG command or by processing for other SMP/E commands.

5. Because some messages are longer than in earlier releases of SMP/E, you may need to increase the size of any data sets used for SMP/E messages (such as SMPOUT). How much more space you need depends on the current size of these data sets and which messages are issued. To start, you may want to allocate new data sets twice the size of the old ones.

6. Make sure that DDDEF SMPCNTL is not NULL.

SMPLOGA

ddname
SMPLOGA.

Use
SMPLOGA is a backup LOG data set. If SMPLOGA is defined, it is used automatically when the SMPLOG data set is full. The LOG data set contains time-stamped records of SMP/E processing. The records in this data set can be written automatically by SMP/E or added by the user through the LOG command. The data set also contains messages issued by SMP/E, as well as detailed information about data set allocation.

Attributes
Sequential; BLKSIZE=514-32000, RECFM=VB, DISP=MOD.

- The BLKSIZE value determines how many records are written to the LOG at one time. As a block is filled, it is written to SMPLOG.

  If BLKSIZE is less than 514, SMP/E uses the BLKSIZE in effect for the SMPLOG data set.

  The BLKSIZE for the SMPLOGA data set must match the BLKSIZE for the related SMPLOG data set.

- DISP=MOD must be specified to maintain a cumulative history of SMP/E processing.

- SMP/E uses LRECL=510 in allocating the SMPLOGA data set.

Device
Tape or direct access.

Note:

1. **Important:** If the data set allocated to the SMPLOGA DD statement has the wrong data set attributes, SMP/E will open the data set for output processing with acceptable attributes. This action will overlay the contents of the data set.

2. Each zone should have its own SMPLOGA data set.

3. If the SMPLOGA data set is full, the extra log data is written to the SMPOUT data set, with the date and time stamp encrypted.

4. SMPLOGA should be updated only by the LOG command or by processing for other SMP/E commands.

5. Because some messages are longer than in earlier releases of SMP/E, you may need to increase the size of any data sets used for SMP/E messages (such as SMPOUT). How much more space you need depends on the current size of these data sets and which messages are issued. To start, you may want to allocate new data sets twice the size of the old ones.

SMPLTS

ddname
SMPLTS.

Use
The SMPLTS data set is used to maintain the base versions of load modules and program objects for which SYSLIB allocations have been specified. A base version of each load module or program object

includes only the modules that were explicitly defined for it by SMP/E MCS statements and JCLIN. Other modules that might be implicitly included in the execution-time load modules and program objects when they are bound are not stored in the base version copy in the SMPLTS.

SMP/E saves the base version of a load module or program object in SMPLTS **only** when the load module or program object is defined to SMP/E with CALLLIBS and the JCLIN specifies SYSLIB allocations (which are used to create the CALLLIBS subentry lists in the corresponding LMOD entries) and:

- the load module or program object contains XZMOD subentries (indicating cross-zone modules), or
- the target zone was created before SMP/E V3R2 and the UPGRADE command has not yet been used to upgrade it.

SMP/E uses the load modules and program objects in the SMPLTS as input when binding the execution-time load modules and program objects into their specified target libraries.

In some cases, SMP/E may also need to create a temporary member in the SMPLTS to resolve certain warning conditions identified by the binder. This temporary member (if created) is deleted from the SMPLTS after a successful link-edit into the target library.

Attributes
Partitioned (DSNTYPE=PDS or DSNTYPE=LIBRARY); RECFM=U, DISP=OLD, BLKSIZE greater than or equal to 6144.

If you allocate the SMPLTS as a PDS, IBM recommends a block size of 32760 to minimize the use of DASD space for this data set.

Device
Direct access only.

Note:

1. Each target zone must have its own SMPLTS data set that is not shared by any other target zone.
2. The SMPLTS data set is required for APPLY, LINK LMODS, LINK MODULE, and RESTORE processing.
3. The BLKSIZE of the SMPLTS data set needs to be 1024 or greater to support load modules that specify the SCTR attribute. A large BLKSIZE is recommended to provide maximum space efficiency.
4. The SMPLTS data set is eligible for RETRY processing, in the same manner as other target libraries, after an *x*37 ABEND. The size of an SMPLTS data set varies depending on the number of load modules that specify a CALLLIBS subentry list in the target zone associated with the SMPLTS.
5. If the SMPLTS data set is to contain program objects, it needs to be allocated as a PDSE.

SMPMTS

ddname
SMPMTS.

Use
The SMPMTS data set (MTS) is a target library for macros that exist only in a distribution library. This data set allows the current version of these macros to be used for assemblies during APPLY processing. For more information about MTS entries, see .

Attributes
Partitioned; LRECL=80, BLKSIZE=multiple of 80, RECFM=FB, DISP=OLD.

Device
Direct access only.

Note:

1. Each target zone must have its own SMPMTS data set that is not shared by any other target zone. This SMPMTS data set may be used with the related distribution zone.

2. For APPLY processing, the SMPMTS data set must be allocated with enough space to hold all the system generation macros, as well as any other macros that do not reside in a target library. This is because functions are now packaged with a complete set of system generation macros.

3. The SMPMTS data set must be in the SYSLIB concatenation for APPLY and RESTORE processing. It can be in the SYSLIB concatenation for ACCEPT processing. For information about SYSLIB concatenation requirements, see *z/OS SMP/E User's Guide*.

SMPNTS

ddname
SMPNTS.

Use

The SMPNTS (SMP/E Network Temporary Store) is a directory of UNIX file system files that are used for temporary storage of network transported packages that were received during SMP/E RECEIVE processing.

Attributes
Existing directory within a UNIX file system. A UNIX file system directory name is concatenated with the appropriate subdirectories and file names to create complete pathnames.

- The directory name can be from 1 to 255 characters.

- The directory name must begin and end with a slash (/).

- In addition to the required delimiters (/), a directory name must also be enclosed in single apostrophes (') if any of the following is true:

  – The directory name contains lowercase alphabetic characters

  – The directory name contains a character that is not uppercase alphabetic, numeric, or national ($, #, or @), slash (/), plus (+), hyphen, period, or ampersand (&).

- The apostrophes must be outside the required delimiters, as in '/directory name/', not /'directory name'/.

- The single apostrophes used to enclose the directory name (the delimiters) do not count as part of the 255-character limit.

- Any apostrophes specified as part of the directory (not the delimiters) must be doubled.

  Double apostrophes count as two characters in the 255-character limit.

- The directory name can include characters through X'40' and X'FE'.

- Do not use symbolic substitution.

Device
Direct access only.

Note:

1. The SMPNTS can be defined to SMP/E only with a DD statement or a DDDEF.

2. Do not allocate the SMPNTS as anything other than a directory in a UNIX file system.

3. The size of the SMPNTS directory depends on the size of the packages received from the network and stored there.

SMPOBJ

ddname
SMPOBJ.

Use
The SMPOBJ data set is used primarily for source-maintained products. It contains preassembled modules that can be used instead of reassembling those modules. These modules must be in load module format—that is, in the same format as modules residing in the distribution library.

Attributes
Partitioned; RECFM=U, DISP=SHR.

Device
Direct access only.

SMPOUT

ddname
SMPOUT.

Use
The SMPOUT data set contains messages issued during SMP/E processing, as well as dumps of the VSAM RPL, if any dumps were taken. It might also contain LIST output and reports if the SMPHRPT, SMPLIST, and SMPRPT data sets are not defined.

Attributes
Sequential; LRECL=121, BLKSIZE=multiple of 121, RECFM=FBA, DISP=MOD.

Device
SYSOUT, printer, direct access, tape, or terminal.

Note:

1. BLKSIZE must not exceed 32760.

2. If SMPOUT is allocated to a data set, the disposition must be MOD, because SMP/E opens and closes the SMPOUT DD statement at each SET command. If the disposition is SHR or OLD, SMPOUT contains only the output from the last set of commands processed before the end of SMP/E processing.

3. If SMPOUT is being allocated for the GIMGTPKG, GIMUNZIP, GIMXSID, or GIMZIP service routine, and you want the resulting SMP/E messages to be formatted to an 80–character length for easier browsing on a terminal, then specify LRECL=81 and BLKSIZE as a multiple of 81.

4. This data set may reside in a UNIX file system. Specify FILEDATA=TEXT and PATHOPTS(OWRONLY,OAPPEND) on the DD statement for this data set, if it is to reside in a UNIX file system. If OAPPEND is not specified, SMPOUT contains only the output from the last set of commands processed before the end of SMP/E processing.

SMPPARM

ddname
SMPPARM.

Use
The SMPPARM data set contains members that allow you to customize SMP/E as follows:

Member name
    Use

GIMDDALC
Defines data sets to be dynamically allocated

GIMEXITS
Defines exit routines

GIMOPCDE
Defines macro and assembler operation codes

Note: The File Allocation Reports for the APPLY and ACCEPT commands include status information for the SMPPARM data set, even if an SMPPARM data set is not allocated. SMP/E does this as a reminder to users who may have intended to supply one or more SMPPARM members. If you do not intend to supply any SMPPARM members, you may ignore this status information for the SMPPARM data set.

- For the GIMDDALC member:

  During SET processing, SMP/E determines whether SMPPARM member GIMDDALC was provided. If so, it processes the GIMDDALC control statements contained within the member. For more

information about GIMDDALC members, see Chapter 3, "Defining control statements in SMPPARM members," on page 103.

- For the GIMEXITS member:

During RECEIVE, APPLY, ACCEPT, RESTORE and LINK processing, SMP/E will determine if SMPPARM member GIMEXITS has been provided. If so, it will process the EXITS control statements contained within the member. For more information about GIMEXITS members, see Chapter 3, "Defining control statements in SMPPARM members," on page 103.

- For the GIMOPCDE member:

During JCLIN processing, SMP/E determines whether assembler instructions are macro invocations or OPCODEs. SMP/E contains a default set of OPCODE definitions, which identify standard assembler OPCODEs. If you do not want to use this default set, you can define your own by using the sample GIMOPCDE member supplied to you. For more information about GIMOPCDE members, see Chapter 3, "Defining control statements in SMPPARM members," on page 103.

Attributes
Partitioned; LRECL=80, BLKSIZE=multiple of 80, RECFM=FB.

Device
Direct access only.

SMPPTFIN

ddname
SMPPTFIN.

This ddname can refer to an actual data set or to a file on a tape (such as file 1 on an ESO tape).

Use
The SMPPTFIN data set contains SYSMODs and ++ASSIGN statements to be processed by the RECEIVE command.

Attributes
Sequential; LRECL=80, BLKSIZE=multiple of 80, RECFM=FB.

Device
Card, tape, direct access, or terminal.

Note:

1. BLKSIZE must not exceed 32760.

2. If the SMPPTFIN data set is inline, you must make sure that the combination of characters used for the delimiter does not occur in the input itself.

   For example, if DD * is specified for SYSIN and the SMPPTFIN data set contains the characters $$, then $$ cannot be used as the default delimiter for the input or for the delimiter specified on the DLM parameter. (The output of SMP/E service routine GIMDTS contains the characters $$.)

3. If you want to receive from multiple product tapes, you cannot concatenate them on a single SMPPTFIN DD statement. Instead, you must process each tape in a separate step, using separate SMPPTFIN DD statements.

   Refer to the documentation supplied with the tape for exact information about how to code the SMPPTFIN DD statement. For example, with a product tape, the program directory contains this information.

4. This data set may reside in a UNIX file system. Specify PATHOPTS(ORDONLY) and either FILEDATA=TEXT or FILEDATA=BINARY on the DD statement if it is to reside in a UNIX file system. For simplicity, FILEDATA=TEXT is preferred, and ensure the newline character (X'15') marks the end of each line. If you use FILEDATA=BINARY, you must ensure that each line is padded with blanks to the 80-byte record length because there is no end of record marker in binary data.

SMPPTS

ddname
> SMPPTS.

Use
> The SMPPTS data set (PTS) is used as a repository for SYSMODs. It contains one member for each SYSMOD that was received. Each member is called an MCS entry and is an exact copy of the SYSMOD as it was received from the SMPPTFIN data set. The name of an MCS entry matches the SYSMOD ID of the SYSMOD it contains. For more information, see "MCS entry (SMPPTS)" on page 225.

> **Note:** SYSMODs in the SMPPTS data set may be stored by SMP/E in a compact format. Specifically, inline element data within SYSMODs may be compacted during RECEIVE or GZONEMERGE command processing. This compact format helps to reduce the space requirements of the SMPPTS data set. The compacted data is automatically expanded when needed during APPLY and ACCEPT command processing. See the description of the COMPACT subentry in "OPTIONS entry (global zone)" on page 241 for information about how to specify whether SYSMODs should be compacted.

> To view the original uncompacted form of any SYSMOD in the SMPPTS data set, you can use the SMP/E Query Dialogs to display the MCS entry in the Global zone for a SYSMOD. You can also use the GIMCPTS service routine to expand one or more compacted SYSMODs (see "GIMCPTS: SYSMOD compaction service routine" on page 385 for more information about GIMCPTS).

Attributes
> Partitioned; LRECL=80, BLKSIZE=multiple of 80, RECFM=FB, DISP=OLD.

> It is recommended that SMPPTS be allocated as a PDSE (DSNTYPE=LIBRARY).

> If you plan to run concurrent jobs, specify DISP=SHR instead of DISP=OLD.

> Unless SMPPTS spill data sets are defined, the SMPPTS must be large enough to contain all the SYSMODs that are to be received from the SMPPTFIN data set.

> Do not concatenate SMPPTS data sets. If you need multiple data sets for the SMPPTS, use SMPPTS spill data sets.

> You **must** specify a valid data set name for the SMPPTS. NULLFILE and DD DUMMY are invalid for the SMPPTS.

Device
> Direct access only.

SMPPTS spill data set

ddname
> SMPPTS1 through SMPPTS99.

Use
> SMPPTS spill data sets can be used to store SYSMODs when the SMPPTS data set becomes full. This type of processing is called *spill processing*. SMPPTS spill data sets are used by SMP/E in the same way as the primary SMPPTS data set is used. For more information, see "SMPPTS" on page 128.

Attributes
> Partitioned; LRECL=80, BLKSIZE=multiple of 80, RECFM=FB, DISP=OLD.

> It is recommended that SMPPTS spill data sets be allocated as a PDSE (DSNTYPE=LIBRARY).

> If you plan to run concurrent jobs, specify DISP=SHR instead of DISP=OLD.

> Do not concatenate SMPPTS spill data sets.

> The first SMPPTS spill data set must be specified with a ddname of SMPPTS1, the second SMPPTS2, and so on, up to a maximum of SMPPTS99. Do not skip any ddnames in this sequence; if a spill data set is omitted, SMP/E ignores any data sets that may follow the omitted data set. (For example, if you specify only SMPPTS1 and SMPPTS3, then SMP/E uses only SMPPTS1 and ignores SMPPTS3.)

You **must** specify a valid data set name for an SMPPTS data set. NULLFILE and DD DUMMY are invalid for SMPPTS spill data sets.

Device

Direct access only.

SMPPUNCH

ddname

SMPPUNCH.

Use

The SMPPUNCH data set contains output from various SMP/E commands. This output generally consists of commands or control statements.

- For BUILDMCS, it contains the complete superseding function SYSMODs.
- For GENERATE, it contains a job stream for building target libraries.
- For REPORT CROSSZONE, it contains commands for installing cross-zone requisites.
- For REPORT ERRSYSMODS, it contains commands for installing SYSMODs that resolve the error hold reason IDs for exception SYSMODs.
- For REPORT SOURCEID, it contains commands for listing SYSMODs associated with the source IDs that were found in the specified zones.
- For REPORT SYSMODS, it contains commands for installing SYSMODs from the input zone that are applicable to the comparison zone.
- For UNLOAD, it contains UCLIN statements for recreating the entries that were unloaded.

Attributes

Sequential; LRECL=80, BLKSIZE=multiple of 80, RECFM=FB, DISP=MOD.

Device

Card, tape, or direct access.

Note:

1. BLKSIZE must not exceed 32760.

2. If SMPPUNCH is allocated to a data set, the disposition must be MOD, because SMP/E opens and closes the SMPPUNCH DD statement at each SET command. If the disposition is SHR or OLD, SMPPUNCH contains only the output from the last set of commands processed before the end of SMP/E processing.

3. For the UNLOAD command, SMPPUNCH should be allocated to a direct access data set or to tape, because the volume of output is large. You may also want a large BLKSIZE; the larger the BLKSIZE, the fewer the times SMP/E must do I/O.

4. This data set may reside in a UNIX file system. Specify FILEDATA=TEXT and PATHOPTS(OWRONLY,OAPPEND) on the DD statement for this data set, if it is to reside in a UNIX file system. If OAPPEND is not specified, SMPPUNCH contains only the output from the last set of commands processed before the end of SMP/E processing.

SMPRPT

ddname

SMPRPT.

Use

The SMPRPT data set contains the reports produced during SMP/E processing.

Attributes

Sequential; LRECL=121, BLKSIZE=multiple of 121, RECFM=FBA, DISP=MOD.

Device

SYSOUT, printer, direct access, tape, or terminal.

Note:

1. BLKSIZE must not exceed 32760.

2. If SMPRPT is not defined, all report output goes to the SMPOUT data set.

3. If SMPRPT is allocated to a data set, the disposition must be MOD, because SMP/E opens and closes the SMPRPT DD statement at each SET command. If the disposition is SHR or OLD, SMPRPT contains only the reports from the last set of commands processed before the end of SMP/E processing.

4. This data set may reside in a UNIX file system. Specify FILEDATA=TEXT and PATHOPTS(OWRONLY,OAPPEND) on the DD statement for this data set, if it is to reside in a UNIX file system. If OAPPEND is not specified, SMPRPT contains only the reports from the last set of commands processed before the end of SMP/E processing.

SMPSCDS

ddname
SMPSCDS.

Use
The SMPSCDS data set (SCDS) contains backup copies of target zone entries that are created during APPLY processing. These backup copies are made before the entries are (1) changed by inline JCLIN, a ++MOVE MCS, or a ++RENAME MCS, or (2) deleted by an element MCS with the DELETE operand. The backup copies are used during RESTORE processing to return the entries to the way they were before APPLY processing. For more information about BACKUP entries, see "BACKUP entries (SMPSCDS)" on page 156.

Attributes
Partitioned; LRECL=80, BLKSIZE=multiple of 80, RECFM=FB, DISP=OLD.

Device
Direct access only.

Note: Each target zone must have its own SMPSCDS data set; that data set must be unique to that target zone. This SMPSCDS data set must also be used with the related distribution zone.

SMPSNAP

ddname
SMPSNAP.

Use
The SMPSNAP data set is used for snap dump output. When a severe error occurs, such as an abend or severe VSAM return code, SMP/E requests a snap dump of its storage before doing any error recovery. In addition, the DEBUG command may request a snap dump of SMP/E storage when specified messages are issued, or it may request a snap dump of control blocks and storage areas associated with a specified dump point.

Attributes
Sequential.

Device
SYSOUT, printer, direct access, tape, or terminal.

Note: BLKSIZE must not exceed 32760.

SMPSRVR

ddname
SMPSRVR

Use
The SMPSRVR data set contains information about a TCP/IP connected host running an FTP or HTTP(S) server. It is used by the GIMGTPKG service routine.

Attributes
Sequential or member of a partitioned data set; LRECL=80, RECFM=F or FB.

Device
Direct access.

Note: This data set may reside in a UNIX file system. Specify PATHOPTS(ORDONLY) and either FILEDATA=TEXT or FILEDATA=BINARY on the DD statement if it is to reside in a UNIX file system. For simplicity, FILEDATA=TEXT is preferred, and ensure the newline character (X'15') marks the end of each line. If you use FILEDATA=BINARY, you must ensure that each line is padded with blanks to the 80-byte record length because there is no end of record marker in binary data.

SMPSTS

ddname
SMPSTS.

Use
The SMPSTS data set (STS) is a target library for source that exists only in a distribution library. This data set allows the current version of these modules to be used for assemblies during APPLY processing.

Attributes
Partitioned; LRECL=80, BLKSIZE=multiple of 80, RECFM=FB, DISP=OLD.

Device
Direct access only.

Note: Each target zone must have its own SMPSTS data set, which may not be shared by any other target zone. This SMPSTS data set can also be used with the related distribution zone.

SMPTLIB

ddname
SMPTLIB.

Use
SMPTLIB data sets (TLIBs) are used as temporary storage for relative files that are loaded from SMPPTFIN during RECEIVE processing. They are deleted when the associated SYSMOD is deleted by REJECT, RESTORE, or ACCEPT processing.

You can have SMP/E dynamically allocate the TLIB data sets, or you can allocate them yourself before RECEIVE processing. For information about how SMP/E allocates the TLIBs, see the RECEIVE Command chapter in *z/OS SMP/E Commands*. (Regardless of how the SMPTLIB data sets are allocated, they do not appear in the File Allocation report.)

Note:

1. No DD statement or DDDEF entry is required for the SMPTLIB data sets if they are preallocated and cataloged.

2. If you allocate and catalog the SMPTLIB data sets yourself, make sure they are allocated on the volume specified in the catalog.

3. If you allocate the SMPTLIB data sets yourself but do not catalog them, make sure they are allocated on the volume specified in the SMPTLIB DD statement or DDDEF entry being used for the command you are processing.

4. If you need to specify a unit that is not SYSALLDA, and the unit is not set by use of a STORCLAS or an ACS filter routine, then you must use a DDDEF entry instead of a DD statement to allocate SMPTLIB data sets and specify the unit value in the DDDEF entry.

5. If you are using SMS to manage your data sets, you can set up the unit, volume, and space allocation through a STORCLAS or an ACS filter routine, instead of specifying them on a DD statement or DDDEF entry.

6. If you are using SMS to manage your data sets, do not specify dummy volumes on a DD statement or in a DDDEF entry for the SMPTLIB allocation. Dummy volumes used as indicators to ACS routines for SMS class selection can cause operator mount messages to be issued for the non-existent dummy volumes. This is because SMP/E first attempts to allocate the SMPTLIB data sets as if they already exist on the volume, which can cause operator mount messages to be issued for the dummy volume.

7. SMPTLIB data sets should not be allocated as PDSEs, because IEBCOPY does not support copying an unloaded PDS load library from tape to a PDSE load library on DASD.

8. DDDEFs are not required in target or distribution zones for SMPTLIB data sets, unless you have uncataloged the SMPTLIB data sets from the global zone (this is not recommended). In this case, only the VOLUME and UNIT subentries are accepted as input on UCLIN statements for target or distribution zone SMPTLIB DDDEFs.

Attributes
Partitioned.

Here are two examples of SMPTLIB DD statements:

```
/SMPTLIB DD UNIT=SYSALLDA,VOL=SER=SMPVOL
```

```
//SMPTLIB DD UNIT=SYSALLDA,
//         VOL=SER=(SMPVL1,SMPVL2)
```

If you use DDDEFs to have SMP/E dynamically allocate the SMPTLIB data sets, you cannot specify the initial or final disposition. SMP/E determines the disposition based on the command it is processing.

Device
Direct access only.

SMPTLOAD

ddname
SMPTLOAD.

Use
SMP/E may need to use an SMPTLOAD data set when installing program elements that were packaged inline. SMPTLOAD is required when applying or accepting a program element and either:

- the destination data set (target or distribution library) is a PDS and the unloaded inline data represents a PDSE that contains program elements
- the destination data set is a PDSE and the unloaded inline data set has RECFM=U and is not a PDSE.

SMP/E dynamically allocates a new SMPTLOAD data set whenever one is needed.

Attributes
Partitioned (DSNTYPE=PDS or DSNTYPE=LIBRARY), DISP=(NEW,DELETE).

Device
Direct access.

Note:

1. If the unloaded inline data represents a PDSE, SMP/E allocates SMPTLOAD with DSNTYPE=LIBRARY. Otherwise, SMP/E allocates SMPTLOAD with DSNTYPE=PDS.

2. All other allocation information for the SMPTLOAD is obtained from the SMPTLOAD DDDEF or, if a DDDEF does not exist, from GIMDDALC control statements found in SMPPARM member GIMDDALC. If a DD statement is supplied for SMPTLOAD, it is freed and the allocation for SMPTLOAD is done using the DDDEF or the GIMDDALC control statement found in SMPPARM member GIMDDALC.

SMP/E uses the following information, if it is supplied by the user, to dynamically allocate an SMPTLOAD data set:

- Primary and secondary space amounts

- Allocation units for primary and secondary space
- Number of directory blocks
- A unit designation
- A volume designation
- A STORCLAS designation
- A MGMTCLAS designation
- A DATACLAS designation

If SMP/E cannot get the required allocation information from a DDDEF or a GIMDDALC control statement in the GIMDDALC member of SMPPARM, then the allocation fails.

3. If there is an SMPTLOAD DDDEF, SMP/E tries to allocate the SMPTLOAD data set, even if the DDDEF does not supply any of the previously listed information. This is because SMP/E always tries to allocate the SMPTLOAD data set new with either a DSNTYPE of PDS or LIBRARY, which may be sufficient if an ACS routine exists to supply other needed allocation values.

4. SMPTLOAD does not appear in the File Allocation Report.

5. SMP/E will not use an SMPTLOAD data set specified with JCL. If an SMPTLOAD DD statement is specified, SMP/E dynamically frees the SMPTLOAD ddname prior to allocating a new SMPTLOAD data set.

SMPWKDIR

ddname
SMPWKDIR.

Use
The SMPWKDIR directory identifies the name of a directory in a UNIX file system. SMPWKDIR is used to store temporary files created during SMP/E processing. If SMPWKDIR is not specified, SMP/E will use one of the following directories for temporary files, depending on the command or service:

- APPLY and ACCEPT commands: the system /tmp directory
- RECEIVE command: the system /tmp directory and the package directory of the SMPNTS directory specified on a DD statement or DDDEF entry.
- GIMZIP and GIMUNZIP service routines: the package directory specified on the SMPDIR DD statement.

Attributes
Existing directory in a UNIX file system.

Device
Direct access only.

Note: The SMPWKDIR can be defined for the RECEIVE, APPLY, and ACCEPT commands using either a DD statement or a DDDEF entry. It can be defined for the GIMZIP and GIMUNZIP service routines using a DD statement only.

SMPWRK1

ddname
SMPWRK1.

Use
The SMPWRK1 data set is used as temporary storage for macro updates and replacements that will be processed by the IEBUPDTE and IEBCOPY programs. During APPLY and ACCEPT processing, SMP/E places the input in this data set before calling the utility.

Attributes
Partitioned; LRECL=80, BLKSIZE=multiple of 80, RECFM=FB, DISP=(NEW,DELETE).

Specifying DISP=(NEW,DELETE) minimizes space loss problems. SMPWRK1 is generally needed only for the duration of the SMP/E job step. To keep the data set longer than that, you must use a different DISP value and compress the data set to reclaim space.

The SMS-related options DATACLAS, MGMTCLAS, STORCLAS, and DSNTYPE provide additional data set support. For more information, refer to *z/OS DFSMS Using Data Sets*.

Device
Direct access only.

Note:

1. BLKSIZE must not exceed 32760.
2. If the BLKSIZE value is omitted or inapplicable, SMP/E uses a default of 0 which allows the system to determine the most appropriate block size.

SMPWRK2

ddname
SMPWRK2.

Use
The SMPWRK2 data set is used as temporary storage for source updates and source replacements that will be processed by the IEBUPDTE and IEBCOPY programs. During APPLY and ACCEPT processing, SMP/E places the input in this data set before calling the utility.

Attributes
Partitioned; LRECL=80, BLKSIZE=multiple of 80, RECFM=FB, DISP=(NEW,DELETE).

Specifying DISP=(NEW,DELETE) minimizes space loss problems. SMPWRK2 is generally needed only for the duration of the SMP/E job step. To keep the data set longer than that, you must use a different DISP value and compress the data set to reclaim space.

The SMS-related options DATACLAS, MGMTCLAS, STORCLAS, and DSNTYPE provide additional data set support. For more information, refer to *z/OS DFSMS Using Data Sets*.

Device
Direct access only.

Note:

1. BLKSIZE must not exceed 32760.
2. If the BLKSIZE value is omitted or inapplicable, SMP/E uses a default of 0 which allows the system to determine the most appropriate block size.

SMPWRK3

ddname
SMPWRK3.

Use
The SMPWRK3 data set is used as temporary storage for object modules supplied by a SYSMOD, object modules created by assemblies, and IMASPZAP control cards following ++ZAP statements.

Attributes
Partitioned; LRECL=80, BLKSIZE=multiple of 80, RECFM=FB, DISP=(NEW,DELETE).

Specifying DISP=(NEW,DELETE) minimizes space loss problems. SMPWRK3 is generally needed only for the duration of the SMP/E job step. To keep the data set longer than that, you must use a different DISP value and compress the data set to reclaim space.

If you want to save assembled object modules until SMP/E has successfully applied or accepted the SYSMODs that caused the assemblies, specify DISP=(NEW,KEEP). This allows SMP/E to reuse the assembled object modules if the APPLY or ACCEPT command fails. For more information about reusing assemblies, see *z/OS SMP/E Commands*.

Note: If SMPWRK3 is a permanent data set, make sure to specify OLD as the initial disposition. Do not specify SHR. An initial disposition of SHR may cause an abend.

The SMS-related options DATACLAS, MGMTCLAS, STORCLAS, and DSNTYPE provide additional data set support. For more information, refer to *z/OS DFSMS Using Data Sets*.

Device
Direct access only

Note:

1. BLKSIZE must not exceed 32760.
2. If the BLKSIZE value is omitted or inapplicable, SMP/E uses a default of 0 which allows the system to determine the most appropriate block size.

SMPWRK4

ddname
SMPWRK4.

Use
The SMPWRK4 data set is used as temporary storage for IMASPZAP and link-edit input containing EXPAND control statements.

Attributes
Partitioned; LRECL=80, BLKSIZE=multiple of 80, RECFM=FB, DISP=(NEW,DELETE).

Specifying DISP=(NEW,DELETE) minimizes space loss problems. SMPWRK4 is generally needed only for the duration of the SMP/E job step. To keep the data set longer than that, you must use a different DISP value and compress the data set to reclaim space.

The SMS-related options DATACLAS, MGMTCLAS, STORCLAS, and DSNTYPE provide additional data set support. For more information, refer to *z/OS DFSMS Using Data Sets*.

Device
Direct access only.

Note:

1. BLKSIZE must not exceed 32760.
2. If the BLKSIZE value is omitted or inapplicable, SMP/E uses a default of 0 which allows the system to determine the most appropriate block size.

SMPWRK6

ddname
SMPWRK6.

Use
The SMPWRK6 data set is used during APPLY and ACCEPT processing as temporary storage for inline replacements for data elements. It is also used as temporary storage for inline updates for JAR elements. SMP/E places the input in this data set so it can be directly accessed and installed by the copy utility or SMP/E.

Attributes
Partitioned; LRECL=80, BLKSIZE=multiple of 80, RECFM=FB, DISP=(NEW,DELETE).

It is recommended that SMPWRK6 be allocated as a PDSE (DSNTYPE=LIBRARY).

Specifying DISP=(NEW,DELETE) minimizes space loss problems. SMPWRK6 is generally needed only for the duration of the SMP/E job step. To keep the data set longer than that, you must use a different DISP value and compress the data set to reclaim space.

The SMS-related options DATACLAS, MGMTCLAS, STORCLAS, and DSNTYPE provide additional data set support. For more information, refer to *z/OS DFSMS Using Data Sets*.

Device
    Direct access only.

Note:

1. BLKSIZE must not exceed 32760.
2. If the BLKSIZE value is omitted or inapplicable, SMP/E uses a default of 0 which allows the system to determine the most appropriate block size.

SMPnnnnn

ddname
    SMP*nnnnn*, where *nnnnn* is a number from 00000 through 99999.

Use
    SMP/E allocates SMP*nnnnn* data sets for its own internal processing. To avoid processing errors, do not assign such a ddname to any data set.

SYSIN

ddname
    SYSIN.

Use
    Contains package control statements used as input to the GIMZIP and GIMUNZIP service routines.

Attributes
    Sequential or a member of a partitioned data set; LRECL=80

Device
    Direct access or SYSIN file.

Note:

1. SYSIN can be defined to GIMZIP and GIMUNZIP only with a DD statement.
2. The size of the input data set will vary depending on how you intend to manage it.
3. You are responsible for managing the contents and space of the input data set.
4. This data set may reside in a UNIX file system. Specify PATHOPTS(ORDONLY) and either FILEDATA=TEXT or FILEDATA=BINARY on the DD statement if it is to reside in a UNIX file system. For simplicity, FILEDATA=TEXT is preferred, and ensure the newline character (X'15') marks the end of each line. If you use FILEDATA=BINARY, you must ensure that each line is padded with blanks to the 80-byte record length because there is no end of record marker in binary data.

SYSLIB

ddname
    SYSLIB.

Use
    The SYSLIB data set is a concatenation of macro libraries that are to be used by the assembler utility.

    For APPLY and RESTORE processing, the data sets should be concatenated in this order:

    1. SMPMTS
    2. MACLIB
    3. MODGEN
    4. Target system macro libraries (such as libraries specified for SYSLIB on the ++MAC statement)
    5. Distribution macro libraries (such as libraries specified for DISTLIB on the ++MAC statement)

    For more information about the proper SYSLIB concatenation for ACCEPT processing, see *z/OS SMP/E User's Guide*.

Attributes

Partitioned; LRECL=80, BLKSIZE=multiple of 80, RECFM=FB.

Device

Direct access only.

SYSPRINT

ddname

SYSPRINT.

You can specify different SYSPRINT data sets for each of the utilities that SMP/E calls. This can be done with the PRINT value in the appropriate UTILITY entry. For more information, see "UTILITY entry (global zone)" on page 289.

Use

The SYSPRINT data set contains output from the utilities called by SMP/E.

Attributes

Sequential; DISP=MOD.

Do not specify LRECL, BLKSIZE, or RECFM unless they are compatible with the attributes used by the utilities called.

If SYSPRINT is allocated to a data set, the disposition must be MOD, because SMP/E opens and closes the SYSPRINT DD statement at each SET command. If the disposition is SHR or OLD, SYSPRINT contains only the output from the last set of commands processed before the end of SMP/E processing.

Device

SYSOUT, printer, direct access, or tape.

SYSOUT or a tape is recommended, because SYSPRINT might be opened with different DCB attributes by the various utilities and service aids called by SMP/E.

Note:

1. BLKSIZE must not exceed 32760.

2. How you specify the SYSPRINT data set can affect whether a listing of the utility output is produced. For example, no listing is produced if the PRINT value in the UTILITY entry specifies either of the following:

   - A DDDEF for a DUMMY data set
   - A DDDEF for a data set that is sent to a SYSOUT class that suppresses output

SYSPUNCH

ddname

SYSPUNCH.

Use

The SYSPUNCH data set is a temporary data set containing object modules assembled by running the job stream produced by system generation or the GENERATE command. These modules are not installed in the distribution libraries at ACCEPT time.

Attributes

Partitioned; LRECL=80, BLKSIZE=multiple of 80, RECFM=FB.

Do not specify a DISP value. Instead, let SMP/E use its defaults:

- DISP=(NEW,PASS) for the first job generated
- DISP=(OLD,PASS) for any subsequent jobs

Device

Direct access only.

SYSUT1, SYSUT2, and SYSUT3

ddname
SYSUT1, SYSUT2, and SYSUT3.

Use
These are scratch data sets for SMP/E and the utilities it calls. They can be used instead of the following data sets when certain utilities are called:

- SYSIN

  - For invocations of the copy utility
  - For some invocations of update utility
  - For some invocations of the *x*37 RETRY COMPRESS utility

- SYSLIN for invocations of the link-edit utility
- SYSUT2 for invocations of the assembler utility

SYSUT1 and SYSUT2 are also used by the GIMDTS service routine:

- SYSUT1 points to the data set containing the input in its original format.
- SYSUT2 points to the data set containing the transformed output.

Attributes
Sequential; DISP=(NEW,DELETE)

Device
Direct access only.

Note:

1. BLKSIZE must not exceed 32760.
2. Do not specify BLKSIZE=0. SMP/E does not support system-determined block size (SDB) for this data set.
3. If the BLKSIZE value is omitted or invalid, SMP/E uses a default of 3200.
4. When processing partitioned data sets, GIMUNZIP dynamically allocates SYSUT1 as a large format sequential data set (DSNTYPE=LARGE).

SYSUT4

ddname
SYSUT4.

Use
This data set can be used instead of the SYSIN data set when certain utilities are called:

- For invocations of the *x*37 RETRY COMPRESS utility
- For some invocations of the assembler utility

It can also be used to contain the sequential unload form of RELFILE data sets for RECEIVE FROMNETWORK and RECEIVE FROMNTS command processing.

Attributes
Sequential; DISP=(NEW,DELETE)

Device
Direct access only.

Note:

1. BLKSIZE must not exceed 32760.
2. Do not specify BLKSIZE=0. SMP/E does not support system-determined block size (SDB) for this data set.

3. If the BLKSIZE value is omitted or invalid, SMP/E uses a default of 3200.

Target library

ddname
The ddname for a target library should match the low-level qualifier of the data set name. For example, the ddname for SYS1.LINKLIB should be LINKLIB.

Use
Target libraries contain updated versions of macros, source modules, and load modules that were stored during APPLY or RESTORE processing. They are the libraries used for your running system. You must provide a DDDEF entry or DD statement for each target library that is being processed.

Attributes
Partitioned.

Device
Direct access only.

Text library (TXLIB)

ddname
The ddname for a text library must match the TXLIB value on the ++JCLIN or element MCS. For example, the ddname for the text library on statement ++MOD(MODA) TXLIB(LIBX) must be LIBX.

Use
Text libraries contain JCLIN input or replacements for macros, source, or object modules that have not been link-edited. They are used when the JCLIN or elements are provided in partitioned data sets rather than inline or in relative files.

Attributes
Partitioned.

Device
Direct access only.

Zone statement

ddname
The ddname must match the name of the target or distribution zone.

Use
The *zone* DD statement is used by SMP/E to access an individual target or distribution zone in a CSI data set. For example, to have SMP/E access target zone MVSTGT1, you can provide a DD statement like this one:

```
//MVSTGT1  DD DSN=SMPE.SMPCSI.CSI,DISP=SHR
```

If you do not provide a DD statement for a target or distribution zone, SMP/E allocates the zone dynamically using the ZONEINDEX information in the GLOBALZONE entry. Also note that, while DD statements may be used to override the ZONEINDEX information, they are not a substitute for a zoneindex. A zoneindex is always required for a zone.

Attributes
VSAM; RECORDSIZE(24 143), KEYS(24 0).

Device
Direct access only.

Note:

1. The low-level qualifier of the data set name must be CSI.

2. If you have used IBM SMP/E for z/OS, V3R6 to update a zone, you might not be able to process that zone with previous releases of SMP/E. For more information, see the migration section in *z/OS SMP/E User's Guide*.

3. When running on systems with the required level of DFP, SMP/E automatically takes advantage of the local shared resource (LSR) feature of VSAM. This reduces the number of times SMP/E must access data when it is reading CSI data sets. As a result, SMP/E performance is improved for commands such as APPLY, APPLY CHECK, ACCEPT, ACCEPT CHECK, and especially LIST.

4. CSI data sets should usually be allocated dynamically. However, you may want to use the batch local shared resources (BLSR) subsystem with expanded storage hiperspaces (instead of SMP/E's implementation of LSR) to improve SMP/E performance during APPLY and ACCEPT processing for a large number of changes. For more information about BLSR, see the notes section in "SMPCSI" on page 116.

5. For information about the CSI data set containing the global zone, see "SMPCSI" on page 116.

Chapter 5. SMP/E data set entries

This chapter describes the entries in the various data sets SMP/E uses. It discusses the following:

- How these data sets are organized
- How the entries in these data sets are organized
- How to create, update, and obtain information about these entries

Note: The LIST, UCLIN, and UNLOAD commands used to process these entries are described in *z/OS SMP/E Commands*.

How the data sets are organized

SMP/E uses the following data sets as a database for its processing: the CSI, PTS, SCDS, MTS, and STS. It is important to understand how SMP/E uses each of these data sets and how they are related.

The following is a description of each data set:

- The **CSI** data set is a VSAM data set that serves as the primary data set for SMP/E. SMP/E divides the CSI into multiple partitions using the VSAM key structure. Each partition is referred to as a *zone*.

  There are three types of zones:

  Global zone
  A single global zone is used to record information about SYSMODs that have been received. The global zone also contains information enabling SMP/E to access the other two types of zones, and information allowing you to tailor parts of SMP/E processing.

  Target zones
  One or more target zones are used to record information about the status and structure of the operating system (or target) libraries.

  Distribution zones
  One or more distribution zones are used to record information about the status and structure of the distribution libraries.

  In the CSI data set is pictured as one data set. In fact, you can group all your zones within one VSAM data set (up to 32766 zones per data set), or divide them up into multiple VSAM data sets, even one zone per data set. The choice is yours; it is based on such factors as:

  – Which packs the associated target and distribution libraries are on. It is advisable to keep the associated zone on the same pack as the libraries controlled from the zone, so when the pack is dumped, the data and description are kept synchronized.

  – The organization of your shop. Totally separate service organizations for different products may require many separate data sets.

  illustrates the same zone relationship as , with the CSI spread across multiple VSAM data sets.

- The **PTS** is used strictly as a storage data set for SYSMODs. A SYSMOD is read from the SMPPTFIN data set and stored directly on the PTS without any modifications or SMP/E information. It is, therefore, related to the global zone in that both data sets contain information about the received SYSMODs. You can, therefore, look at the global zone and the PTS as a pair of data sets that must be processed (such as scratched, saved, or modified) concurrently.

- The **SCDS** is used by SMP/E to store backup copies of modified target zone entries during apply processing of a SYSMOD with inline JCLIN. Thus, the SCDS is directly related to a specific target zone, and each target zone must have its own SCDS.

- The **MTS** is a library in which SMP/E stores updated copies of macros during apply when no other target macro library is identified. The MTS is, therefore, related to a specific target zone, and each target zone must have its own MTS data set.
- The **STS** is a library in which SMP/E stores updated copies of source during APPLY when no other target source library is identified. The STS is, therefore, related to a specific target zone, and each target zone must have its own STS data set.

The relationships among SMP/E data sets and zones are shown in <u>Figure 4 on page 142</u> for a single CSI data set, and in <u>Figure 5 on page 144</u> for multiple CSI data sets.


```
//SMPCSI    DD DSN=SMPE.SMPCSI.CSI,DISP=SHR
//TGT1      DD DSN=SMPE.SMPCSI.CSI,DISP=SHR
//TGT2      DD DSN=SMPE.SMPCSI.CSI,DISP=SHR
//DLB1      DD DSN=SMPE.SMPCSI.CSI,DISP=SHR
//DLB2      DD DSN=SMPE.SMPCSI.CSI,DISP=SHR
```

Figure 4. Single-CSI structure

How data set entries are organized

Within the global zone, target zone, and distribution zone, SMP/E keeps many types of entries. These can be divided into two categories:

- Those that are used to control SMP/E processing. These consist of:
  - GLOBALZONE definition entry
  - TARGETZONE definition entry
  - DLIBZONE definition entry
  - OPTIONS entries
  - UTILITY entries
  - DDDEF entries
  - FMIDSET entries
  - ZONESET entries

```
//SMPCSI    DD DSN=SMPE.SMPCSI.CSI,DISP=SHR
//TGT1      DD DSN=TARGET.SMPCSI.CSI,DISP=SHR
//TGT2      DD DSN=TARGET.SMPCSI.CSI,DISP=SHR
//DLB1      DD DSN=DLIB.SMPCSI.CSI,DISP=SHR
//DLB2      DD DSN=DLIB.SMPCSI.CSI,DISP=SHR
```

SMPE.SMPCSI.CSI
- GLOBAL zone ⟷ PTS

TARGET.SMPCSI.CSI
- TGT1 zone ⟷ SCDS   MTS   STS
- TGT2 zone ⟷ SCDS   MTS   STS
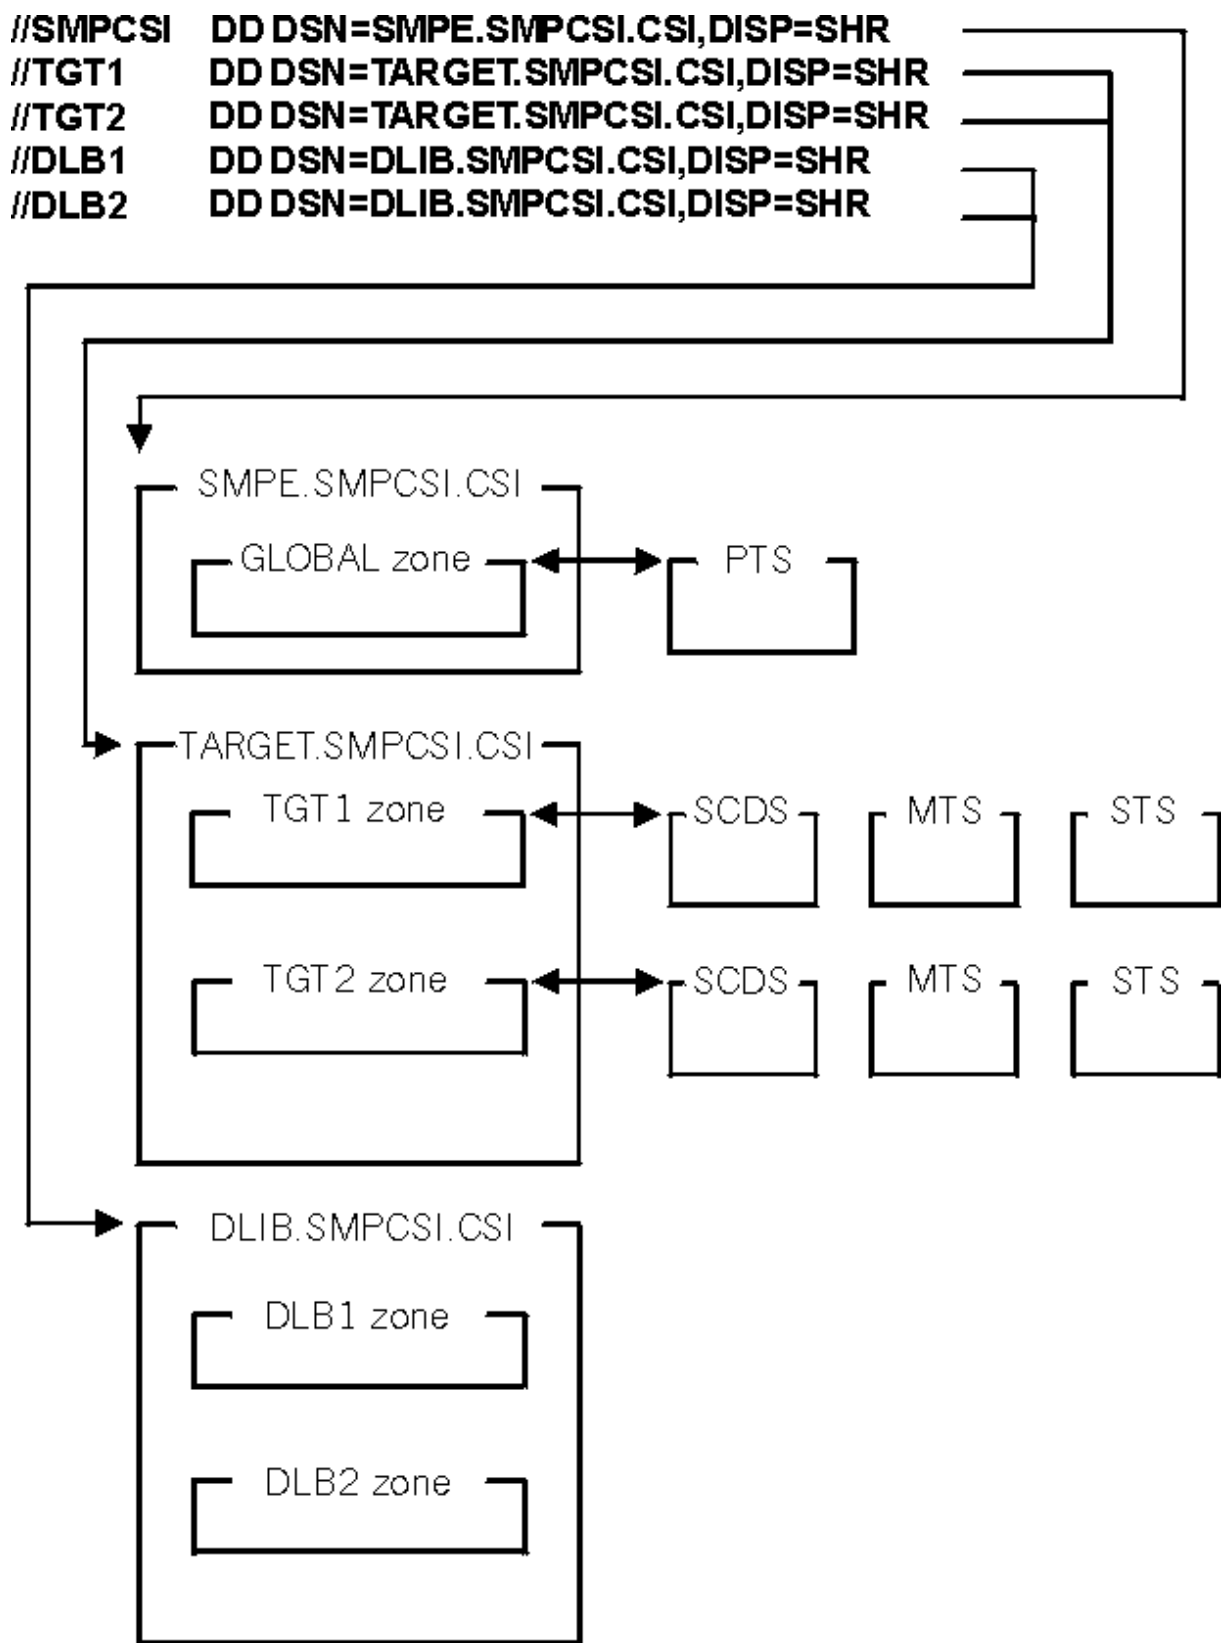
DLIB.SMPCSI.CSI
- DLB1 zone
- DLB2 zone

Figure 5. Multiple-CSI structure

- Those that are used to define the status and structure of the target libraries and distribution libraries. These consist of:

- ASSEM entries
- Data element entries
- DLIB entries
- hierarchical file system element entries
- LMOD entries
- MAC entries
- MOD entries
- SRC entries
- SYSMOD entries

   Note: The SYSMOD entries also contain information referred to as HOLDDATA.

The following sections describe the entry relationships for the entries that control processing, and entries that define status and structure.

Entries that control processing

The starting point for all SMP/E processing is the SMPCSI DD statement. This DD statement directs SMP/E to the CSI data set containing the **GLOBALZONE entry**. The next step after obtaining the GLOBALZONE entry depends on which zone you direct SMP/E to process. (For more information about identifying which zone SMP/E is to process, see the SET command chapter in *z/OS SMP/E Commands*.

Processing the global zone

Once the GLOBALZONE entry has been obtained and you direct SMP/E to process the global zone (for instance to receive SYSMODs), the OPTIONS subentry in the GLOBALZONE entry directs SMP/E to the correct **OPTIONS entry** to use. The OPTIONS entry contains information about the SMP/E processing options to be used. During processing, SMP/E may have to invoke one of the system utilities. Another entry, the **UTILITY entry**, is used to define to SMP/E information about the utility program to invoke, the parameters to pass to it, and the return code to expect from it. SMP/E finds the UTILITY entry through the OPTIONS entry that names the UTILITY entry. Thus, for SMP/E to use the correct utility program, you must define **both** the UTILITY entry that describes the utility program and the OPTIONS entry that names that UTILITY entry.

The other processing entries in the global zone are the DDDEF, FMIDSET, and ZONESET entries.

- **DDDEF entries** provide information used by SMP/E to dynamically allocate data sets. These entries are not connected to any other processing entry. They are accessed individually, by name, as SMP/E needs information to dynamically allocate a specific data set. They are not shown in Figure 6 on page 146.
- **FMIDSET entries** define sets of FMIDs. Various other commands, such as APPLY and ACCEPT, can then process SYSMODs by FMIDSET.
- **ZONESET entries** define sets of zones. The REPORT command can then check for SYSMODs in the zones defined in the ZONESET.

Each of these entries (GLOBALZONE, OPTIONS, UTILITY, DDDEF, FMIDSET, and ZONESET) exists in the global zone.

Figure 6 on page 146 illustrates how the various global zone processing entries are related to one another. Subsequent sections of this chapter provide a more detailed explanation of each entry and the data it contains.

Figure 6. Global zone: relationships between entries that control processing

Processing a target zone or a distribution zone

A target zone and a distribution zone are processed in very similar ways. The only difference is the name of the controlling entry. Therefore, they are covered together, and the differences are pointed out as appropriate.

When you direct SMP/E to process a target zone (for instance to apply a PTF) or a distribution zone (for instance to accept that PTF), SMP/E accesses the **ZONEINDEX subentries** in the GLOBALZONE entry. Those subentries list the target zones and distribution zones that you have defined, including the zone type, target or DLIB, and the name of the CSI data set on which they reside.

After SMP/E has determined that the zone specified is valid, it uses the CSI data set name specified to dynamically allocate a DD statement to access the required data set (optionally, SMP/E looks for a DD statement equal to the zone name). SMP/E can now access that DD statement to obtain further processing information for the zone.

The first entry accessed is the **TARGETZONE entry** (for target zones) or the **DLIBZONE entry** (for distribution zones). Each of these, in turn, directs SMP/E to the correct OPTIONS entry to use, which in turn directs SMP/E to the correct UTILITY entries. Target and distribution zones also contain DDDEF entries.

- The **OPTIONS and UTILITY entries** serve the same purpose for the target zone and distribution zone as for the global zone. However, the OPTIONS and UTILITY entries used to process a target zone or distribution zone are defined not in the target or distribution zone, but in the global zone that points to the target or distribution zone.
- TARGETZONE and DLIBZONE entries contain a **RELATED subentry**, which identifies a related zone. For a target zone, the RELATED subentry identifies the distribution zone from which this target zone was

built. For a distribution zone, the RELATED subentry identifies the target zone that was built from these distribution libraries.

- **DDDEF entries** in the target zone and distribution zone are not connected to any other entries. They provide information used by SMP/E to dynamically allocate data sets.

Figure 7 on page 148 shows how the various target zone, distribution zone, and global zone processing entries are related. Subsequent sections of this chapter provide a more detailed explanation of each entry and the data it contains. Because neither the FMIDSET entries nor the DDDEF entries are connected to any other distribution zone or target zone entry, they are not shown in this figure.

Entries that define status and structure

Once the processing control information for the specified zone has been determined, SMP/E uses the status and structure entries within the various zones to determine what should be done.

Processing the global zone

Processing the global zone is fairly simple, in that only one structure and status entry exists, SYSMOD. **SYSMOD entries** are used to determine whether a SYSMOD has been received already. Various indicators within the SYSMOD entry also indicate whether the SYSMOD has been applied or accepted. The SYSMOD entries, although not directly connected to the PTS **MCS entries**, are implicitly connected, in that SMP/E assumes that there is a one-to-one relationship between global zone SYSMOD entries and PTS MCS entries.

Processing the target zone

Processing the target zone is much more complex than processing the global zone, because the operations performed are much more complex and require more entry types to direct SMP/E. The primary purpose of the target zone entries is to enable SMP/E to apply new function and service to the target system libraries; thus, the APPLY command is used to describe the relationship between the various target zone entries.

The starting point for applying a SYSMOD is actually the **global zone SYSMOD entries**. From the global zone SYSMOD entry, SMP/E obtains information about the relationship between this SYSMOD and other SYSMODs. SMP/E then uses other global zone SYSMOD entries and the **target zone SYSMOD entries** to resolve these relationships.

Once the eligible SYSMODs have been determined, the actual MCS statements from the PTS are accessed and lead SMP/E to the other target zone entries:

- ++MAC and ++MACUPD statements lead to target zone MAC entries.
- ++SRC and ++SRCUPD statements lead to target zone SRC entries.
- ++MOD and ++ZAP statements lead to target zone MOD and LMOD entries.
- Data element statements lead to target zone data element entries.
- ++*hfs_element* statements lead to target zone hierarchical file system element entries.
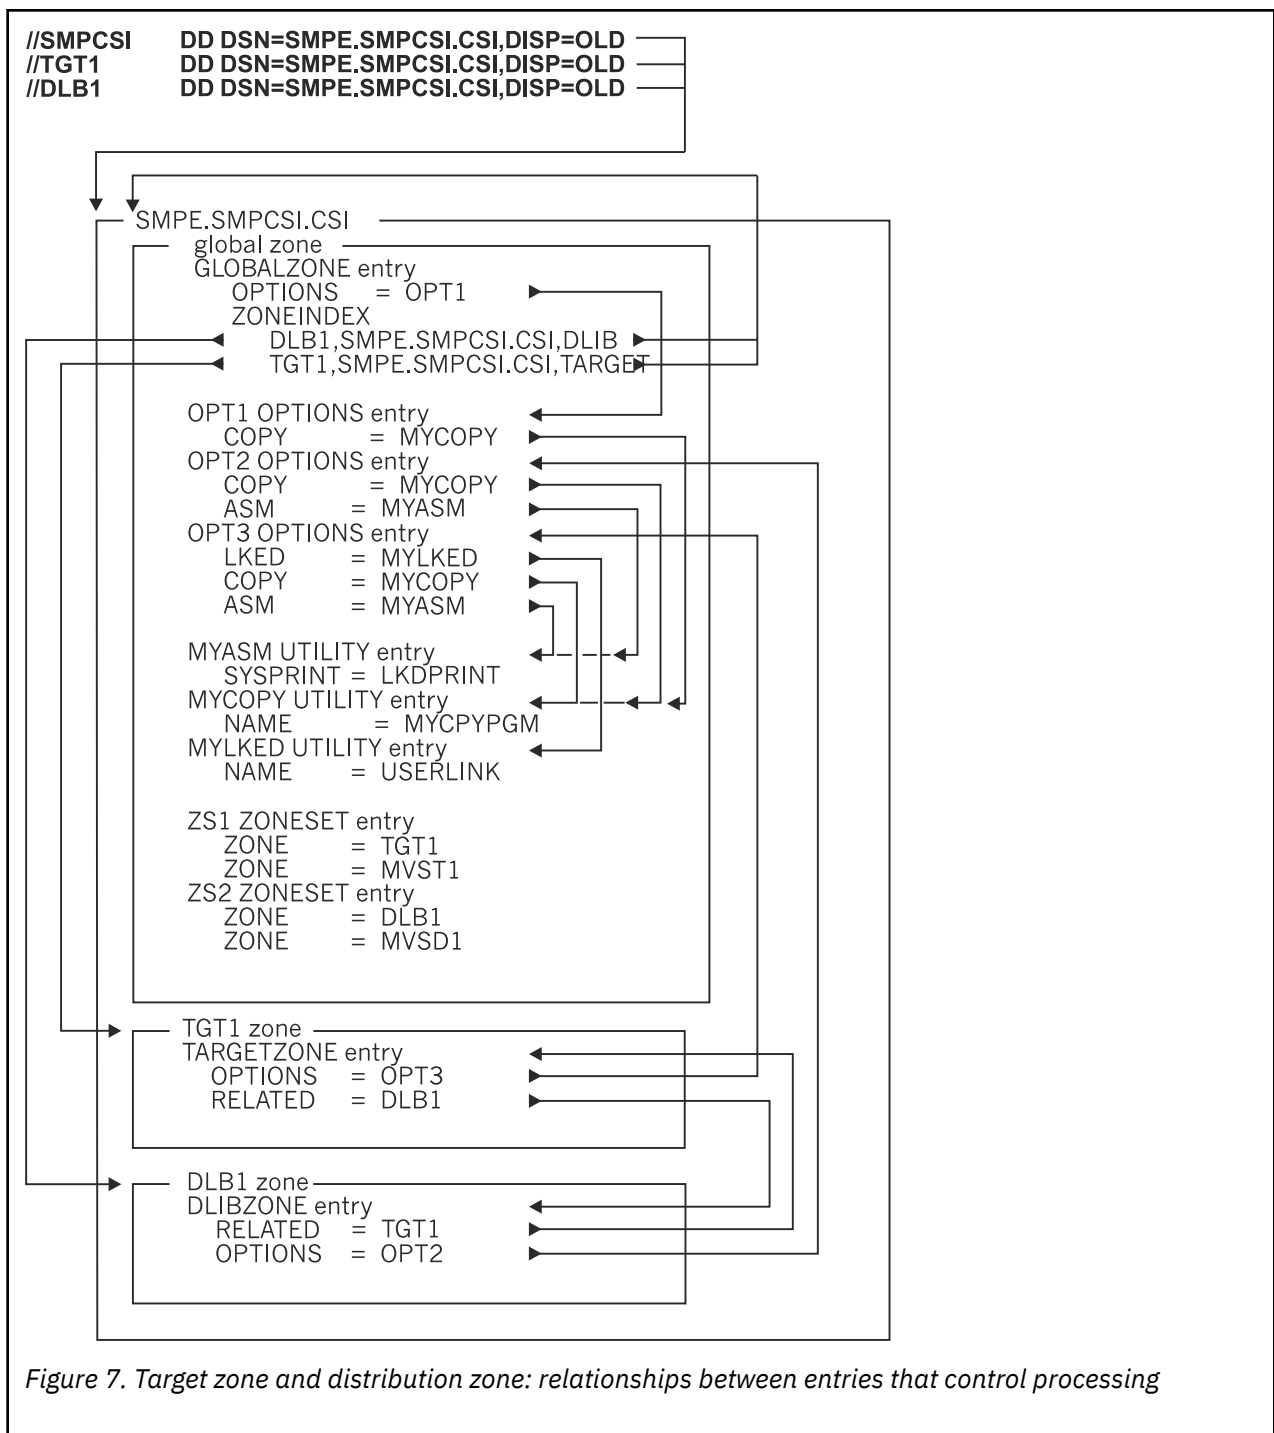
```
//SMPCSI     DD DSN=SMPE.SMPCSI.CSI,DISP=OLD
//TGT1       DD DSN=SMPE.SMPCSI.CSI,DISP=OLD
//DLB1       DD DSN=SMPE.SMPCSI.CSI,DISP=OLD

         SMPE.SMPCSI.CSI
            global zone
              GLOBALZONE entry
                  OPTIONS    = OPT1
                  ZONEINDEX
                      DLB1,SMPE.SMPCSI.CSI,DLIB
                      TGT1,SMPE.SMPCSI.CSI,TARGET

              OPT1 OPTIONS entry
                  COPY       = MYCOPY
              OPT2 OPTIONS entry
                  COPY       = MYCOPY
                  ASM        = MYASM
              OPT3 OPTIONS entry
                  LKED       = MYLKED
                  COPY       = MYCOPY
                  ASM        = MYASM

              MYASM UTILITY entry
                  SYSPRINT = LKDPRINT
              MYCOPY UTILITY entry
                  NAME       = MYCPYPGM
              MYLKED UTILITY entry
                  NAME       = USERLINK

              ZS1 ZONESET entry
                  ZONE       = TGT1
                  ZONE       = MVST1
              ZS2 ZONESET entry
                  ZONE       = DLB1
                  ZONE       = MVSD1


            TGT1 zone
              TARGETZONE entry
                  OPTIONS    = OPT3
                  RELATED    = DLB1


            DLB1 zone
              DLIBZONE entry
                  RELATED    = TGT1
                  OPTIONS    = OPT2
```

Figure 7. Target zone and distribution zone: relationships between entries that control processing

If the SYSMOD replaces or modifies a macro (++MAC or ++MACUPD in SYSMOD), SMP/E uses the target zone **MAC entry** to determine the functional and service level of the macro and to determine whether any assemblies must be redone as a result of the macro update. This is done by checking the GENASM subentries in the target zone MAC entry. The GENASM subentries contain the names of either ASSEM entries or SRC entries in the target zone. In either case, SMP/E accesses the appropriate entry and uses the information stored there to perform the required assemblies.

If the SYSMOD replaces or modifies source (++SRC or ++SRCUPD in SYSMOD), SMP/E uses the target zone **SRC entry** to determine the functional and service level of the source and the library containing the source code, and then performs the required assembly.

After an assembly has been done (either using data from the ASSEM entry or SRC entry), SMP/E knows that the resulting object deck must be link-edited into the target libraries somewhere. Although there is

no direct connection (that is, no subentry value present) between the ASSEM or SRC entries and a **MOD entry**, there is an implicit relationship. SMP/E assumes that for each ASSEM and SRC entry, there exists a MOD entry with the same name. Thus, after performing the assembly, SMP/E accesses the corresponding MOD entry to determine where to install the object deck. This leads us to the same point as if a MOD were supplied in the SYSMOD.

If the SYSMOD replaces or modifies a module (++MOD or ++ZAP in SYSMOD), or if an assembly was done for a SRC or ASSEM entry, SMP/E uses the target zone **MOD entry** to determine the functional and service level of the module and the load modules into which the module should be linked. Information about load modules is kept in the LMOD subentries in the MOD entry.

The LMOD subentries within the MOD entry lead SMP/E to the target zone **LMOD entries**. These represent load modules that exist in the target libraries. LMOD entries contain all the information necessary either to relink the load module or to superzap it.

If the SYSMOD replaces a data element (a data element MCS is in the SYSMOD), SMP/E uses the target zone **data element entry** to determine the functional and service level of the data element. It then gets the data element installed in the appropriate target library.

If the SYSMOD replaces a hierarchical file system element (a ++*hfs_element* MCS is in the SYSMOD), SMP/E uses the target zone **hierarchical file system element entry** to determine the functional and service level of the hierarchical file system element. It then gets the hierarchical file system element installed in the appropriate target library (which is actually in a UNIX file system). If a shell script has been defined for the element, SMP/E passes control to the shell script to allow it to perform any necessary post-installation processing.

After all the updating, assembling, linking, and so on, are done, you arrive back almost at the starting point, the SYSMOD entry. The difference is that this is the target zone SYSMOD entry rather than the global zone SYSMOD entry. The target zone SYSMOD entry contains all the information about what has been done as a result of installing this SYSMOD.

That leaves us with one target zone entry not discussed, the **DLIB entry**. The DLIB entries are not connected to any other entry. They are used by SMP/E to keep information about libraries that are totally copied during product installation, and are used during APPLY processing to create the appropriate element and LMOD entries for elements coming from the copied DLIB.

Figure 8 on page 150 shows how the various SMP/E status and structure entries are related. Subsequent sections of this chapter provide a more detailed explanation of each entry and the data it contains.
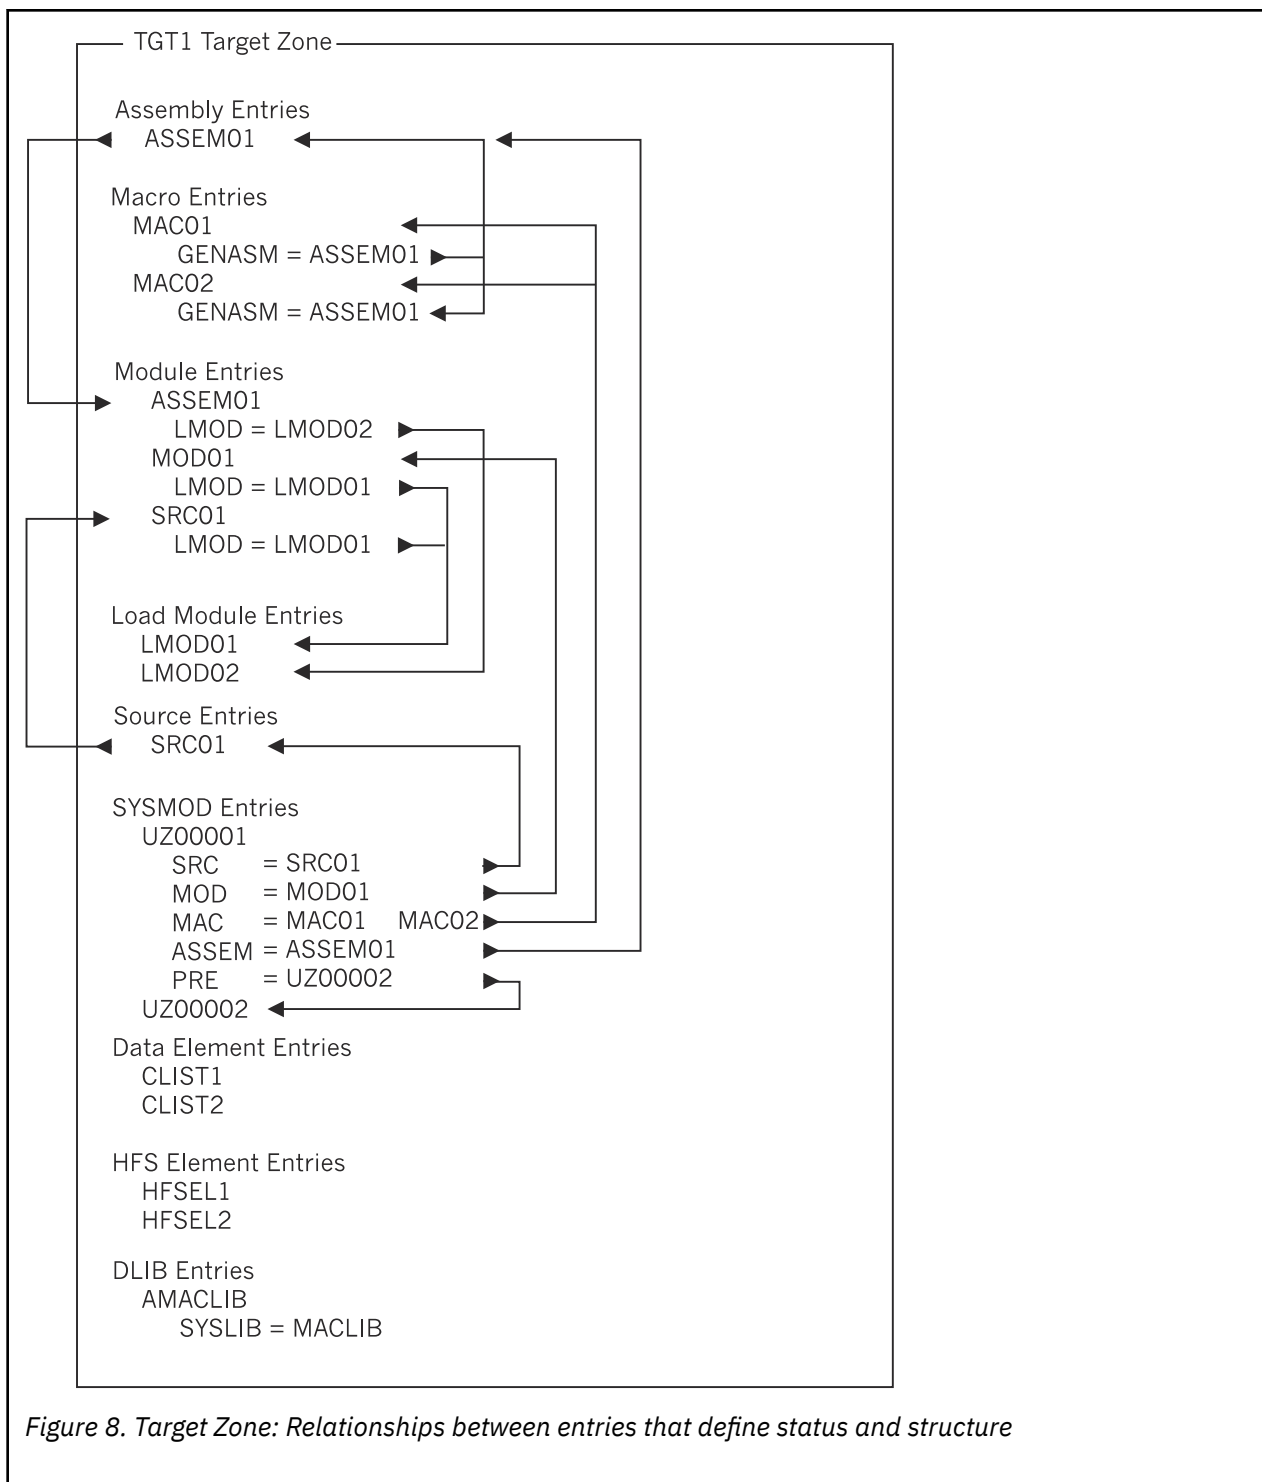
```
┌─ TGT1 Target Zone ────────────────────────────────────────┐
│                                                            │
│     Assembly Entries                                       │
│        ASSEM01                                             │
│                                                            │
│     Macro Entries                                          │
│        MAC01                                               │
│           GENASM = ASSEM01                                 │
│        MAC02                                               │
│           GENASM = ASSEM01                                 │
│                                                            │
│     Module Entries                                         │
│        ASSEM01                                             │
│           LMOD = LMOD02                                    │
│        MOD01                                               │
│           LMOD = LMOD01                                    │
│        SRC01                                               │
│           LMOD = LMOD01                                    │
│                                                            │
│     Load Module Entries                                    │
│        LMOD01                                              │
│        LMOD02                                              │
│     Source Entries                                         │
│        SRC01                                               │
│                                                            │
│     SYSMOD Entries                                         │
│        UZ00001                                             │
│           SRC    = SRC01                                   │
│           MOD    = MOD01                                   │
│           MAC    = MAC01     MAC02                         │
│           ASSEM  = ASSEM01                                 │
│           PRE    = UZ00002                                 │
│        UZ00002                                             │
│     Data Element Entries                                   │
│        CLIST1                                              │
│        CLIST2                                              │
│                                                            │
│     HFS Element Entries                                    │
│        HFSEL1                                              │
│        HFSEL2                                              │
│                                                            │
│     DLIB Entries                                           │
│        AMACLIB                                             │
│           SYSLIB = MACLIB                                  │
│                                                            │
└────────────────────────────────────────────────────────────┘
```

Figure 8. Target Zone: Relationships between entries that define status and structure

Note: Lines to the left of the figure are implicit connections based on the assumption that entries with equal names will be found.

Processing the distribution zone

Processing the distribution zone is similar to processing the target zone, but more complex than processing the global zone. The primary purpose of the distribution zone entries is to enable SMP/E to accept new function and service to the distribution libraries; it is also used to remove a SYSMOD from the system. The ACCEPT command will be used to describe the relationship between the various distribution zone entries.

The starting point for accepting a SYSMOD is actually the **global zone SYSMOD entry**. From that entry, SMP/E obtains information about the relationship between this SYSMOD and other SYSMODs. SMP/E then uses other global zone SYSMOD entries and the **distribution zone SYSMOD entries** to resolve these relationships. In addition, SMP/E checks the target zone SYSMOD entries to ensure that the selected SYSMODs have been applied.

After the eligible SYSMODs have been determined, the various MCS statements lead SMP/E to the other distribution zone entries:

- ++MAC and ++MACUPD statements lead to distribution zone MAC entries.
- ++SRC and ++SRCUPD statements lead to distribution zone SRC entries.
- ++MOD and ++ZAP statements lead to distribution zone MOD entries.
- Data element statements lead to distribution zone data element entries.
- ++*hfs_element* statements lead to distribution zone hierarchical file system element entries.

If the SYSMOD replaces or modifies a macro (++MAC or ++MACUPD in SYSMOD), SMP/E uses the distribution zone **MAC entry** to determine the functional and service level of the macro and to determine whether any assemblies must be redone as a result of the macro update. This is done by checking the GENASM subentries in the target zone MAC entry. The names in the list can be names of either target zone ASSEM or SRC entries. In either case, SMP/E accesses the appropriate entry and uses the information stored there to perform the required assemblies.

If the SYSMOD replaces or modifies a source (++SRC or ++SRCUPD in SYSMOD), SMP/E uses the distribution zone **SRC entry** to determine the functional and service level of the source and the library containing the source code, and then performs the required assembly.

After an assembly has been done (with data from either the ASSEM entry or the SRC entry), SMP/E knows that the resulting object deck may be link-edited into the distribution libraries somewhere. Although there is no direct connection (that is, no subentry value present) between the ASSEM or SRC entries and a **MOD entry**, there is an implicit relationship. SMP/E assumes that for each ASSEM and SRC entry there exists a MOD entry with the same name. Thus, after performing the assembly, SMP/E accesses the corresponding MOD entry to determine where to install the object deck. This leads us to the same point as if a MOD were supplied in the SYSMOD.

If the SYSMOD replaces or modifies a module (++MOD or ++ZAP in SYSMOD), or if an assembly was done for a SRC or ASSEM entry, SMP/E uses the distribution zone **MOD entry** to determine the functional and service level of the module, the DLIB into which it must be linked, the load modules into which the module should be linked, and the link-edit attributes that should be used. Information about load modules is kept in the LMOD subentries in the MOD entry.

The LMOD subentries within the MOD entry lead SMP/E to the distribution zone **LMOD entries**. These represent load modules that exist in the target libraries. LMOD entries contain all the information necessary, either to relink the load module or to update it using IMASPZAP.

If the SYSMOD replaces a data element (a data element MCS is in the SYSMOD), SMP/E uses the distribution zone **data element entry** to determine the functional and service level of the data element. It then gets the data element installed in the appropriate distribution library.

If the SYSMOD replaces a hierarchical file system element (a ++*hfs_element* MCS is in the SYSMOD), SMP/E uses the distribution zone **hierarchical file system element entry** to determine the functional and service level of the hierarchical file system element. It then gets the hierarchical file system element installed in the appropriate distribution library.

After all the updating, assembling, linking, and so on, are done, you arrive back almost at the starting point, the SYSMOD entry. The difference is that this is the distribution zone SYSMOD entry, which contains all the information about what was done as a result of installing this SYSMOD.

shows how the various SMP/E status and structure entries are related. Subsequent sections of this chapter provide a more detailed explanation of each entry and the data it contains.

```
                    ┌── DLIB1 Distribution Zone ──────────────────────────┐
                    │                                                      │
                    │    Assembly Entries                                  │
              ┌────◄│       ASSEM01  ◄──────────────┐  ◄────────────┐      │
              │     │                               │               │      │
              │     │    Macro Entries              │               │      │
              │     │       MAC01                   │               │      │
              │     │          GENASM = ASSEM01 ►───┘               │      │
              │     │       MAC02               ◄─────────────────┐ │      │
              │     │          GENASM = ASSEM01 ◄──────┐          │ │      │
              │     │                                  │          │ │      │
              │     │    Module Entries                │          │ │      │
              └───►│       ASSEM01                     │          │ │      │
                    │          LMOD = LMOD02 ►─────┐   │          │ │      │
                    │       MOD01              ◄───┼───┤          │ │      │
                    │          LMOD = LMOD01 ►─────┤   │          │ │      │
                 ┌─►│       SRC01                  │   │          │ │      │
                 │  │          LMOD = LMOD01 ►─────┤   │          │ │      │
                 │  │                              │   │          │ │      │
                 │  │    Load Module Entries       │   │          │ │      │
                 │  │       LMOD01           ◄─────┘   │          │ │      │
                 │  │       LMOD02           ◄─────────┘          │ │      │
                 │  │    Source Entries                           │ │      │
                 └─◄│       SRC01    ◄───────────────┐            │ │      │
                    │                                │            │ │      │
                    │    SYSMOD Entries              │            │ │      │
                    │       UZ00001                  │            │ │      │
                    │          SRC   = SRC01    ►────┘            │ │      │
                    │          MOD   = MOD01    ►─────────────────┘ │      │
                    │          MAC   = MAC01   MAC02 ►──────────────┘      │
                    │          ASSEM = ASSEM01  ►─────────────────────────┘
                    │          PRE   = UZ00002  ►────┐
                    │       UZ00002             ◄────┘
                    │    Data Element Entries
                    │       CLIST1
                    │       CLIST2
                    │
                    │    HFS Element Entries
                    │       HFSEL1
                    │       HFSEL2
                    │
                    │    DLIB Entries
                    │       AMACLIB
                    │          SYSLIB = MACLIB
                    └──────────────────────────────────────────────────────
```
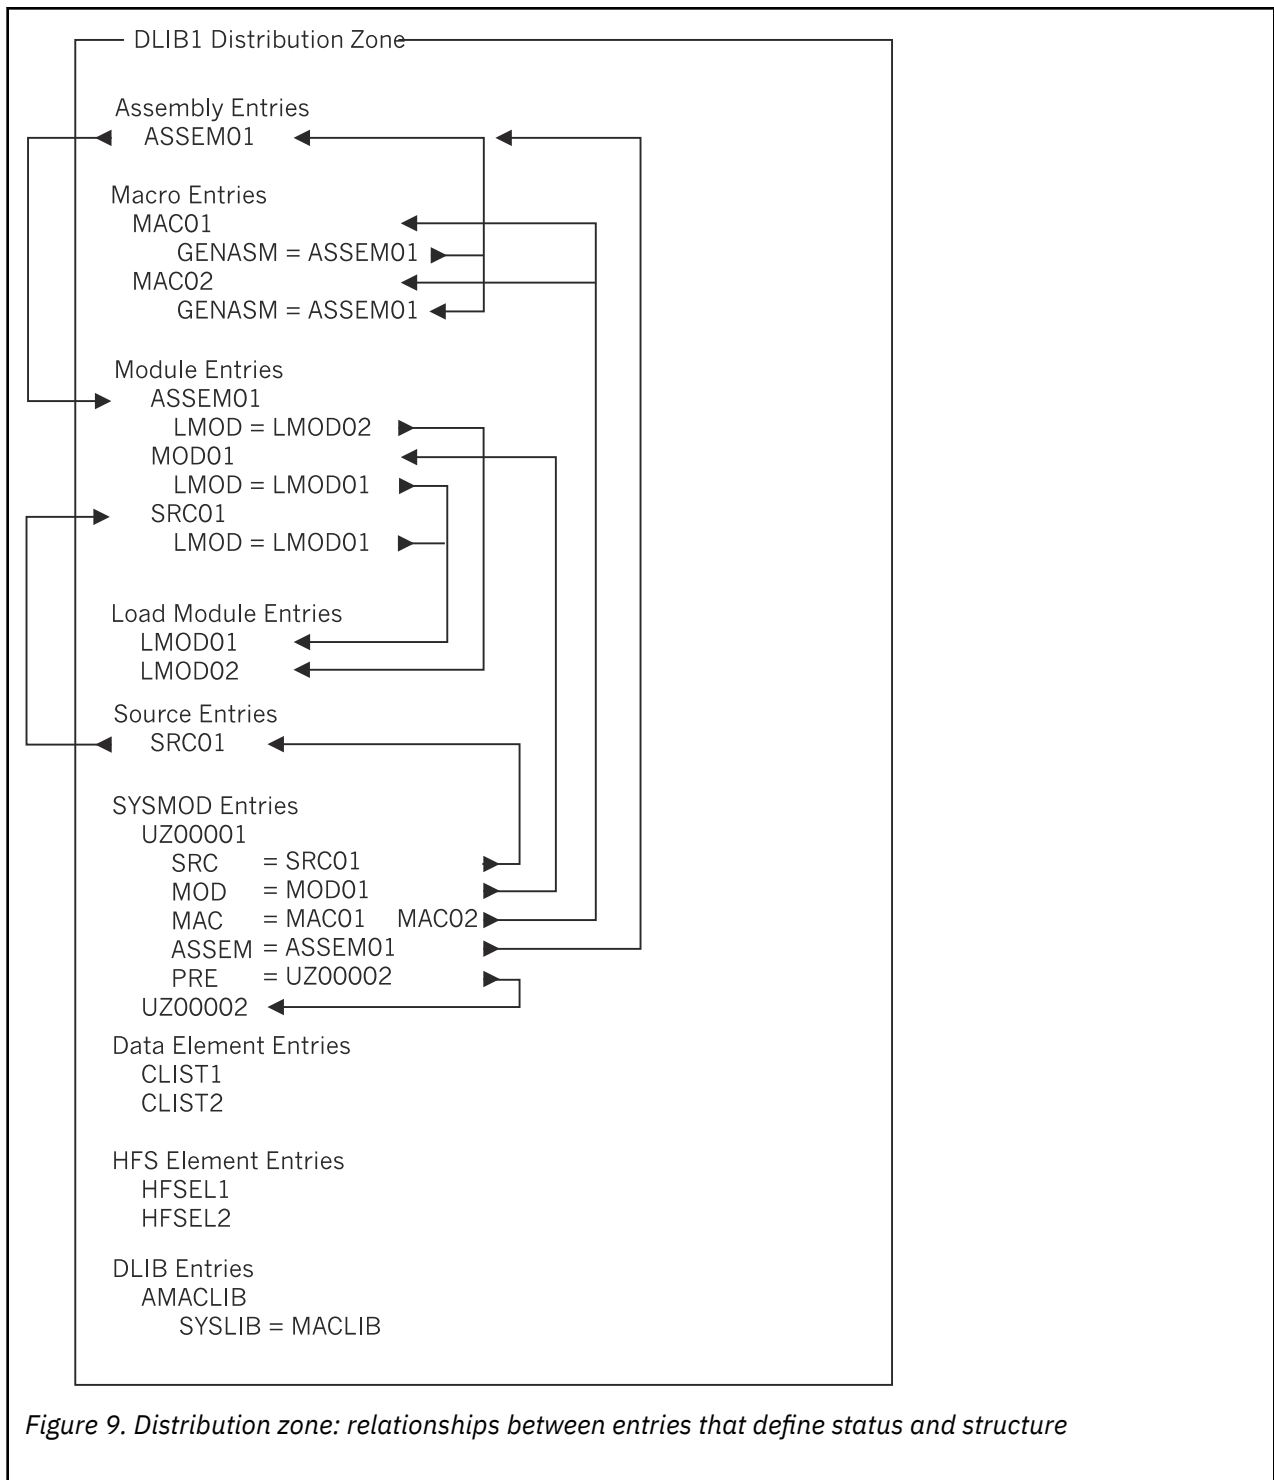
Figure 9. Distribution zone: relationships between entries that define status and structure

Note: Lines to the left of the figure are implicit connections based on the assumption that entries with equal names will be found.

The remaining sections of this chapter deal with each of the specific entries in the various SMP/E data sets.

ASSEM entry (distribution and target zone)

The ASSEM entry contains assembler statements that can be assembled to create an object module. It is created during JCLIN processing when SMP/E encounters an assembler step with inline assembler

input. When the module is reassembled using the statements in the ASSEM entry, SMP/E copies those statements into the SMPWRK2 data set, and then assembles the module.

If a macro is invoked in the assembly, the ASSEM entry is pointed to by the GENASM subentry in the MAC entry created for that macro. As a result, when that macro is updated, SMP/E can reassemble the affected module using the statements in the ASSEM entry. For additional information, see the "Processing" section in the JCLIN command chapter in *z/OS SMP/E Commands*.

Subentries

These are the subentries for the ASSEM entry as they appear in the LIST output:

name
> is the name of the ASSEM entry. It corresponds to the name of the module that can be reassembled by use of that ASSEM entry.
>
> The name can contain from 1 to 8 alphanumeric characters.

ASSEMBLER INPUT
> is the actual assembler statements that were saved for this module during JCLIN processing. These statements are passed to the assembler whenever this module must be reassembled.
>
> The UCL (Update Control Language) operands are ++ASMIN and ++ENDASMIN.
>
> • An ASSEM entry must contain at least the ++ASMIN and ++ENDASMIN statements, plus the associated assembler statements.
> • The ++ASMIN and ++ENDASMIN statements must start in column 1.
> • No other operands can start on the same line as the ++ASMIN statement.
> • If you specify the ++ASMIN statement, you must also specify the ++ENDASMIN statement.

LASTUPD
> identifies the cause of the last change to this ASSEM entry.
>
> The UCL operand is LASTUPD(*value*). This subentry can contain one of the following values:
>
> **JCLIN**
> > indicates that the change was made during JCLIN command processing.
>
> **UCLIN**
> > indicates that the change was made as a result of UCLIN processing.
>
> *sysmod_id*
> > indicates that the change was made during the installation of the indicated SYSMOD.
> >
> > The SYSMOD ID must contain 7 alphanumeric characters.

LASTUPD TYPE
> indicates how the entry was last changed.
>
> The UCL operand is LASTUPDTYPE(*value*). This subentry can contain one of the following values:
>
> **ADD**
> > The entry was added.
>
> **UPD**
> > The entry was updated.

LIST Examples

To list all the ASSEM entries in a particular zone, you can use the following commands:

```
SET      BDY(TGT1)            /* Set to requested zone.   */.
LIST     ASSEM                /* List all ASSEM entries.  */.
```

To list specific ASSEM entries, you can use these commands:

```
     SET     BDY(TGT1)         /* Set to requested zone.   */.
     LIST    ASSEM(ASSEM01     /* List only these two      */
                  ASSEM02)     /* entries.                 */.
```

The format of the LIST output for each ASSEM entry is the same for both of these commands. The only difference is the number of ASSEM entries listed. Figure 10 on page 154 is an example of LIST output for ASSEM entries.

```
PAGE nnnn  - NOW SET TO zzzzzz ZONE nnnnnnn  DATE mm/dd/yy TIME hh:mm:ss SMP/E 36.nn SMPLIST OUTPUT

TGT1           ASSEMBLER ENTRIES


  NAME

ASSEM01   LASTUPD         = JXY1102 TYPE=ADD
          ASSEMBLER INPUT = ...
                            ... assembler statements
                            ....

ASSEM02   LASTUPD         = JXY1121 TYPE=UPD
          ASSEMBLER INPUT = ...
                            ... assembler statements
                            ...
```

Figure 10. ASSEM entry: sample LIST output

You can use the LIST command to find the MAC entries for macros that are called by the assembler statements in this ASSEM entry. To include the names of these MAC entries in the LIST output, you can use the XREF operand, as shown in these commands:

```
     SET     BDY(TGT1)         /* Set to requested zone.   */.
     LIST    ASSEM             /* List all ASSEM entries   */
             XREF              /* and macros that use them. */.
```

Note:

1. You can use XREF in either mass mode or select mode.

2. SMP/E obtains the data included for the XREF operand by checking the GENASM subentries in all the MAC entries. Because this data is not contained in the ASSEM entry itself, you cannot use UCLIN to change it in the ASSEM entry.

Figure 11 on page 154 is an example of the LIST output produced when the XREF operand is used.

```
PAGE nnnn  - NOW SET TO zzzzzz ZONE nnnnnnn  DATE mm/dd/yy TIME hh:mm:ss SMP/E 36.nn SMPLIST OUTPUT

TGT1           ASSEMBLER ENTRIES


  NAME

ASSEM01   LASTUPD         = JXY1102 TYPE=ADD
          ASSEMBLER INPUT = ...
                            ... assembler statements
                            ...
          MACROS USED     = NAME      FMID
                            MAC01     JXY1102
                            MAC02     JXY1121

ASSEM02   LASTUPD         = JXY1121 TYPE=UPD
          ASSEMBLER INPUT = ...
                            ... assembler statements
                            ...
          MACROS USED     = NAME      FMID
                            MAC01     JXY1102
                            MAC03     JXY1121
```

Figure 11. ASSEM entry: sample LIST output when XREF is specified

UNLOAD Examples

To dump the ASSEM entries in UCL format, you can use the UNLOAD command. To unload all the ASSEM entries in a particular zone, you can use the following commands:

```
SET     BDY(TGT1)          /* Set to requested zone.    */.
UNLOAD  ASSEM              /* Unload all ASSEM entries.  */.
```

To unload specific ASSEM entries, you can use these commands:

```
SET     BDY(TGT1)          /* Set to requested zone.    */.
UNLOAD  ASSEM(ASSEM01      /* Unload only these two      */
             ASSEM02)      /* entries.                   */.
```

The format of the UNLOAD output for each ASSEM entry is the same for both of these commands. The only difference is the number of ASSEM entries listed. Figure 12 on page 155 is an example of UNLOAD output for ASSEM entries.

```
UCLIN .
REP       ASSEM          ( ASSEM01  )
          LASTUPD        ( JXY1102  )
          LASTUPDTYPE    ( ADD )
++ASMIN
...
... assembler statements
...
++ENDASMIN
                              .

REP       ASSEM          ( ASSEM02  )
          LASTUPD        ( JXY1121  )
          LASTUPDTYPE    ( UPD )
++ASMIN
...
... assembler statements
...
++ENDASMIN
                              .
ENDUCL.
```

Figure 12. ASSEM entry: sample UNLOAD output

UCLIN Examples

You can use the ADD, REP, and DEL UCL statements to change subentries in the ASSEM entry. When you use UCLIN to update an ASSEM entry, keep these points in mind:

• After the UCLIN changes are done, the ASSEM entry must contain at least ++ASMIN and ++ENDASMIN statements, plus the associated assembler input. Otherwise, there is not enough information in the entry to assemble the associated module.

• The input following the ++ASMIN statement replaces the existing assembler input in the ASSEM entry.

• If a DEL statement deletes all the existing subentries in the entry, SMP/E deletes the entire entry.

• When SMP/E processes a DEL statement, it does not compare any assembler input after the ++ASMIN statement with the input that is currently in the ASSEM entry. It just deletes the existing assembler input. This causes an error, because there is now insufficient data in the ASSEM entry.

The following examples are provided to help you use the ASSEM entry.

Example 1: Deleting an ASSEM entry

The main use of UCLIN for an ASSEM entry is to delete the entry. Here is an example:

```
SET     BDY(TGT1)          /* Set to target zone.    */.
UCLIN                      /*                        */.
```

```
DEL      ASSEM(ASSEM01)      /* Delete the entry.      */.
ENDUCL                       /*                        */.
```

Example 2: Adding a new ASSEM entry

To create an ASSEM entry, you should run the JCLIN command rather than use UCLIN. (For examples of how JCLIN creates ASSEM entries, see *z/OS SMP/E Commands*. However, you can also use UCLIN to create ASSEM entries. For example, you can use the following commands to create a new ASSEM entry and to update an **existing** MAC entry to show that the macro is used in the new assembly:

```
SET      BDY(TGT1)           /* Set to target zone.    */.
UCLIN                        /*                        */.
ADD      ASSEM(ASSEM99)      /* New ASSEM entry.       */
++ASMIN                      /* Assembler data.        */
ASSEM99  CSECT               /*                        */
         MAC99    0,0,1      /*                        */
         END ASSEM99         /*                        */
++ENDASMIN                   /* End of assembler data. */
                             /* End of adding ASSEM.   */.
ADD      MAC(MAC99)          /* Modify macro entry used */
         GENASM(ASSEM99)     /* to indicate use in new */
                             /* assembly entry.        */.
ENDUCL                       /*                        */.
```

BACKUP entries (SMPSCDS)

BACKUP entries are a collection of target zone entries that are copied to the SMPSCDS during APPLY processing before they are updated by inline JCLIN, a ++MOVE MCS, or a ++RENAME MCS, or before they are deleted by an element MCS with the DELETE operand. A BACKUP entry is also created for a MOD entry if a SYSMOD being installed provides a ++MOD statement that either changes the distribution library for the module or adds the module to an existing load module.

- As SMP/E processes the inline JCLIN for a given SYSMOD, it determines which entries will be affected by that JCLIN. Before making the changes, it saves a copy of each of those entries on the SMPSCDS.

- Likewise, as SMP/E processes the ++MOVE, ++RENAME, and element MCSs in a given SYSMOD, it determines which entries will be updated or deleted. Before updating or deleting the entries, it saves a copy of each of those entries on the SMPSCDS.

Besides saving copies of the affected entries, SMP/E also saves a SYSMOD entry on the SMPSCDS to indicate which entries were added by JCLIN, updated by JCLIN or an MCS statement in the SYSMOD, or deleted by an MCS statement in the SYSMOD. Each entry is associated with only one SYSMOD. The entries associated with a SYSMOD are called the BACKUP entries for that SYSMOD. BACKUP entries consist of:

- A SYSMOD entry indicating what entries were added, deleted, or updated
- ASSEM entries for updated target zone ASSEM entries
- JAR entries for deleted JAR entries
- LMOD entries for updated target zone LMOD entries
- MAC entries for updated or deleted target zone MAC entries
- MOD entries for updated or deleted target zone MOD entries
- SRC entries for updated or deleted target zone SRC entries
- Data element entries for deleted target zone data element entries
- Hierarchical file system entries for deleted target zone hierarchical file system entries
- DLIB entries for updated target zone DLIB entries

SMP/E provides access to the BACKUP entries as a group—for example, through the LIST command—but it does not provide access to the individual entries or subentries.

Subentries

The subentries in the BACKUP entries are the same as those in the various entry types that are copied. For more information, see the section in this chapter describing each entry.

LIST Examples

To list the BACKUP entries for all the SYSMODs in the SMPSCDS, you can use the following commands:

```
SET       BDY(TGT1)             /* Set to requested tgt zone. */.
LIST      BACKUP                /* List all BACKUP entries.   */.
```

To list the BACKUP entries for a specific SYSMOD, you can use these commands:

```
SET       BDY(TGT1)             /* Set to requested tgt zone. */.
LIST      BACKUP(UZ12345,       /* List only BACKUP entries   */
                 UZ12346)       /* for these two SYSMODs.     */.
```

The format of the LIST output for each group of BACKUP entries is the same for both of these commands. The only difference is the number of SYSMODs for which BACKUP entries are listed.

When you list the BACKUP entries for a SYSMOD, the first entry in the output is a summary entry for the SYSMOD, which indicates the date and time the SYSMOD was applied, as well as which entries were added or updated as a result of applying the SYSMOD. This is followed by a listing of all the existing target zone entries affected by this SYSMOD, before they were updated. Nothing is listed for added entries, because no entry existed before the SYSMOD was installed.

Figure 13 on page 158 is a partial example of LIST output for BACKUP entries. It shows the summary records, but not all the backup copies of the entries modified by the SYSMOD. This is because the format of those copies is the same as the format for the original target zone entries.

UCLIN Examples

You can use the DEL UCL statement to delete BACKUP entries from the SMPSCDS. This can be helpful if you plan to do an APPLY followed by ACCEPT when several target libraries have been created from the same distribution library.

When a SYSMOD is accepted into a distribution zone, the entries associated with the SYSMOD are automatically deleted from the SMPSCDS for the RELATED target zone. However, even if the SYSMOD was also applied to other target zones created from the same distribution zone, SMP/E does not clean up the SMPSCDS data sets for the other target zones.

To delete the entries from these data sets, you can accept the SYSMOD and name these other target zones as the RELATED zone. However, this updates the distribution library each time, which is time-consuming and can use up space in the distribution library data set.

Instead, you can use the DEL command to delete these entries without updating the distribution library. To determine which entries to specify, check the SMPLOG data set to see which ones SMP/E deleted during ACCEPT processing.

Note: You can also use the CLEANUP command to delete BACKUP entries without specifying them individually. For more information about the CLEANUP command, see *z/OS SMP/E Commands*.

```
PAGE nnnn  - NOW SET TO zzzzzz ZONE nnnnnnn  DATE mm/dd/yy TIME hh:mm:ss SMP/E 36.nn SMPLIST OUTPUT

SMPSCDS      BACKUP ENTRIES

  NAME

UZ12345   DATE/TIME APP   = 07.100 08:15:00
          ASSEM   (ADD)  = ASSEM01   ASSEM02   ASSEM03
          LMOD    (ADD)  = LMOD01    LMOD02
          MACRO   (ADD)  = MAC01
          MOD     (ADD)  = MOD01     MOD02     MOD03     MOD04
          SRC     (ADD)  = SRC01     SRC02     SRC03
          DLIB    (ADD)  = DLIB01
          ASSEM   (DEL)  = ASSEM04   ASSEM05
          LMOD    (DEL)  = LMOD03    LMOD04    LMOD05
          MACRO   (DEL)  = MAC05     MAC09
          MOD     (DEL)  = MOD10     MOD11     MOD12
          SRC     (DEL)  = SRC11
          ASSEM   (UPD)  = ASSEM91   ASSEM92   ASSEM93
          LMOD    (UPD)  = LMOD9A    LMOD9B
          MACRO   (UPD)  = MAC99
          MOD     (UPD)  = MOD99
          SRC     (UPD)  = SRC99
          DLIB    (UPD)  = DLIB99
...
... deleted entries for ASSEM04 ASSEM05
...                      LMOD03  LMOD05  LMOD05
...                      MAC05   MAC09
...                      MOD10   MOD11   MOD12
...                      SRC11
... updated entries for ASSEM91 ASSEM92 ASSEM93
...                      LMOD9A  LMOD9B
...                      MAC99
...                      MOD99
...                      SRC99
...                      DLIB99
... would follow here - format is as in sample
... for each target zone entry
...
```

Figure 13. BACKUP entries: sample LIST output

Example: Deleting BACKUP entries

Assume BACKUP entries exist on the SCDS for SYSMODs UZ12345 and UZ12346. Those SYSMODs have been accepted, but the BACKUP entries have not been deleted from the SCDS. This can happen when multiple systems are supported off one set of DLIBs, where you accept from only one of the target systems. Here is an example of UCL you could use to delete the BACKUP entries from the SCDS:

```
SET       BDY(TGT1)          /* Set to TGT1 zone.          */.
UCLIN                        /*                            */.
DEL       BACKUP(UZ12345)    /* Delete the BACKUP entries. */.
DEL       BACKUP(UZ12346)    /* Delete the BACKUP entries. */.
ENDUCL                       /*                            */.
```

Note: You must specify a separate UCL statement for each BACKUP entry to be deleted.

Data element entry (distribution and target zone)

The data element entry describes an element that is not a macro, module, or source (for example, a CLIST or a sample procedure). Data elements may exist in distribution or target libraries. A data element entry is created the first time you install a SYSMOD that contains an MCS for a data element that does not yet have an entry in the CSI data set.

SMP/E records the function and service level of the data element in the entry. Once a data element entry exists, it is updated as subsequent SYSMODs that affect the data element are installed.

Table 2 on page 9 shows the types of entries used for data elements. Some types of elements, such as panels, messages, or text, may have been translated into several languages. In these cases, the entry type contains *xxx*, which represents the language used for the element. (If an element was not translated, the entry type does not contain any *xxx* value.) Table 3 on page 11 shows the *xxx* values and the languages they represent.

Subentries

These are the subentries for the data element entry as they appear in the LIST output:

name
is the name of the data element represented by the entry. It can contain from 1 to 8 alphanumeric characters and $, #, @, or hex C0.

ALIAS
specifies a list of alias names for the element.

The UCL operand is ALIAS(*name…*).

Each alias name can contain from 1 to 8 alphanumeric characters.

DISTLIB
specifies the ddname of the distribution library for the data element.

The UCL operand is DISTLIB(*ddname*).

- The ddname can contain from 1 to 8 alphanumeric characters.
- The DISTLIB subentry is required. Without it, SMP/E cannot process any changes for the data element.

FMID
specifies the functional owner of this data element. The functional owner is the last function SYSMOD that replaced this element.

The UCL operand is FMID(*sysmod_id*).

The SYSMOD ID must contain 7 alphanumeric characters.

LASTUPD
identifies the cause of the last change to this data element entry.

The UCL operand is LASTUPD(*value*). This subentry can contain one of the following values:

UCLIN
indicates that the change was made as a result of UCLIN processing.

sysmod_id
indicates that the change was made during the installation of the indicated SYSMOD.

The SYSMOD ID must contain 7 alphanumeric characters.

LASTUPD TYPE
indicates how the entry was last changed.

The UCL operand is LASTUPDTYPE(*value*). This subentry may contain one of the following values:

ADD
The entry was added.

UPD
The entry was updated.

RMID
identifies the last SYSMOD that **replaced** this data element. Any subsequent SYSMOD that modifies this data element must have a defined relationship (such as PRE or SUP) with this SYSMOD.

The UCL operand is RMID(*sysmod_id*).

- The SYSMOD ID must contain 7 alphanumeric characters.

Data element entry (distribution and target zone)

- If RMID is not specified but FMID is, SMP/E sets the RMID value to the specified FMID.

SYSLIB

specifies the ddname of the target library for the data element.

The UCL operand is SYSLIB(*ddname*).

- You can specify only one SYSLIB value.
- The ddname can contain from 1 to 8 alphanumeric characters.

LIST Examples

To list all the data element entries of a given type (such as CLIST) in a particular zone, you can use the following commands:

```
SET       BDY(TGT1)            /* Set to requested zone.  */.
LIST      CLIST                /* List all CLIST entries. */.
```

To list specific CLIST entries, you can use these commands:

```
SET       BDY(TGT1)            /* Set to requested zone.  */.
LIST      CLIST(CLIST1         /* List only these two     */
             CLIST2)           /* entries.                */.
```

The format of the LIST output for each data element entry is the same for both of these commands. The only difference is the number of data element entries listed. Figure 14 on page 160 is an example of LIST output for data element entries.

```
PAGE nnnn   - NOW SET TO zzzzzz ZONE nnnnnnn  DATE mm/dd/yy TIME hh:mm:ss SMP/E 36.nn SMPLIST OUTPUT

TGT1           CLIST ENTRIES


  NAME

CLIST1    LASTUPD        = MYCLST1 TYPE=ADD
          LIBRARIES      = DISTLIB=AMACLIB    SYSLIB=ISPCLIB
          FMID           = MYCLST1
          RMID           = MYCLST1

CLIST2    LASTUPD        = MYCPTF1 TYPE=ADD
          LIBRARIES      = DISTLIB=AMACLIB    SYSLIB=ISPCLIB
          FMID           = MYCLST1
          RMID           = MYCPTF1
```

Figure 14. Data element entry: sample LIST output

By specifying the FORFMID operand, you can reduce the number of data element entries listed. When FORFMID is specified, SMP/E lists a data element entry only if its FMID matches one of the FMIDs specified on the FORFMID operand. For example, to list CLIST entries whose FMIDs are defined in FMIDSET TP or else are MYCLST1, you can use these commands:

```
SET       BDY(TGT1)            /* Set to target zone.     */.
LIST      CLIST                /* List all CLIST entries  */
          FORFMID(TP           /* for the TP FMIDSET       */
               MYCLST1)        /* and FMID MYCLST1.       */.
```

You can use the LIST command to find out the names of all SYSMODs that have modified data elements. To include the names of these SYSMODs in the LIST output, you can use the XREF operand, as shown in these commands:

```
SET       BDY(TGT1)            /* Set to requested zone.  */.
LIST      CLIST                /* List all CLIST entries  */
          XREF                 /* and related SYSMODs.    */.
```

Note:

1. You can use XREF in either mass mode or select mode.
2. SMP/E obtains the data included for the XREF operand by checking for entries for this data element in all the SYSMOD entries. Because this data is not contained in the data element entry itself, you cannot use UCLIN to change it in the data element entry.

Figure 15 on page 161 is an example of the LIST output produced when the XREF operand is used.

```
PAGE nnnn  - NOW SET TO zzzzzz ZONE nnnnnnn  DATE mm/dd/yy TIME hh:mm:ss SMP/E 36.nn SMPLIST OUTPUT

TGT1          CLIST ENTRIES


   NAME

CLIST1    LASTUPD        = MYCLST1 TYPE=ADD
          LIBRARIES      = DISTLIB=AMACLIB    SYSLIB=ISPCLIB
          FMID           = MYCLST1
          RMID           = MYCLST1
          SYSMOD HISTORY = SYSMOD   TYPE      DATE   MCS    --STATUS--
                           MYCLST1  FUNCTION 07.100  CLIST    APP    ACC

CLIST2    LASTUPD        = MYCPTF1 TYPE=ADD
          LIBRARIES      = DISTLIB=AMACLIB    SYSLIB=ISPCLIB
          FMID           = MYCLST1
          RMID           = MYCPTF1
          SYSMOD HISTORY = SYSMOD   TYPE       DATE   MCS    --STATUS--
                           MYCPTF1  PTF       07.150  CLIST    APP    ACC
```

Figure 15. Data element entry: sample LIST output when XREF is specified

UNLOAD Examples

To dump the data element entries in UCL format, you can use the UNLOAD command. To unload all the CLIST entries in a particular zone, you can use the following commands:

```
SET     BDY(TGT1)          /* Set to requested zone.   */.
UNLOAD  CLIST              /* Unload all CLIST entries. */.
```

To unload specific CLIST entries, you can use these commands:

```
SET     BDY(TGT1)          /* Set to requested zone.   */.
UNLOAD  CLIST(CLIST1       /* Unload only these two    */
            CLIST2)        /* entries.                 */.
```

The format of the UNLOAD output for each data element entry is the same for both of these commands. The only difference is the number of data element entries listed. Figure 16 on page 162 is an example of UNLOAD output for CLIST entries.

Data element entry (distribution and target zone)

```
 UCLIN .
 REP       CLIST           ( CLIST1   )
           LASTUPD         ( MYCLST1  )
           LASTUPDTYPE     ( ADD )
           DISTLIB         ( AMACLIB  )
           SYSLIB          ( ISPCLIB  )
           FMID            ( MYCLST1 )
           RMID            ( MYCLST1 )
                              .
 REP       CLIST           ( CLIST2   )
           LASTUPD         ( MYCPTF1  )
           LASTUPDTYPE     ( ADD )
           DISTLIB         ( AMACLIB  )
           SYSLIB          ( ISPCLIB  )
           FMID            ( MYCLST1 )
           RMID            ( MYCPTF1 )
                              .
 ENDUCL.
```

Figure 16. Data element entry: sample UNLOAD output

By specifying the FORFMID operand, you can reduce the number of data element entries unloaded. When FORFMID is specified, SMP/E unloads a data element entry only if its FMID matches one of the FMIDs specified on the FORFMID operand. For example, to unload CLIST entries whose FMIDs are either defined in FMIDSET TP or are MYCLST1, you can use these commands:

```
SET      BDY(TGT1)          /* Set to target zone.       */.
UNLOAD   CLIST              /* Unload all CLIST entries   */
         FORFMID(TP         /* for the TP FMIDSET         */
                MYCLST1)    /* and FMID MYCLST1.          */.
```

UCLIN Examples

You can use the ADD, REP, and DEL UCL statements to change subentries in the data element entry. After the UCLIN changes are made, the data element entry must contain at least the following subentries:

• DISTLIB

• FMID

• RMID

Otherwise, there is not enough information in the entry to process the data element. If any of these subentries are missing, SMP/E does not make the requested UCL updates to the entry, and the entry remains as it was before the UCL command.

Example: Adding a new data element entry

Assume you have installed CLIST3 outside of SMP/E, but now you want to start using SMP/E to track changes to that CLIST. Here is an example of the UCL statements you would use to define entries for that CLIST in the appropriate target and distribution zones:

```
SET      BDY(TGT1)          /* Set to target zone.      */.
UCLIN                       /*                          */.
ADD      CLIST(CLIST3)      /* Define new CLIST entry.  */
         DISTLIB(AMACLIB)   /* Define DLIB.             */
         SYSLIB(ISPCLIB)    /* System library.          */
         FMID(MYCLST1)      /* Functional owner.        */.
ENDUCL                      /*                          */.
SET      BDY(DLB1)          /* Now do same to DLIB.     */.
UCLIN                       /*                          */.
ADD      CLIST(CLIST3)      /* Define new CLIST entry.  */
         DISTLIB(AMACLIB)   /* Define DLIB.             */
                            /* no SYSLIB info in DLIB.  */
         FMID(MYCLST1)      /* Functional owner.        */.
ENDUCL                      /*                          */.
```

DDDEF entry (distribution, target, and global zone)

The DDDEF entry contains the information SMP/E needs to dynamically allocate a specific data set. With DDDEF entries, you do not have to provide a DD statement for every data set SMP/E may need to process a particular command. When SMP/E determines that it needs a specific data set, it looks for a DD statement that it can use to allocate that data set. If there is no DD statement, SMP/E checks whether the current zone contains a DDDEF entry for that data set. If so, it uses the information in the DDDEF entry to dynamically allocate the data set.

For more information about dynamically allocating data sets, see *z/OS SMP/E User's Guide*.

Note:

1. In a job with multiple SET commands, if you use DDDEF entries that specify SYSOUT for SMP/E output (such as SMPOUT or SMPRPT), SMP/E produces multiple SYSOUT data sets. This can cause undesirable results; for example, the output could appear to be out of sequence from one SET command to the next. Therefore, when you run such a job, you may prefer to use DD statements instead of DDDEF entries for SMP/E output data sets.

2. SMP/E does not strictly enforce rules for which subentries you can specify in DDDEF entries for specific data sets. To prevent possible allocation errors, refer to *z/OS MVS JCL User's Guide*.

Subentries

These are the subentries for the DDDEF entry as they appear in the LIST output: For more information concerning DDDEF entry syntax, see the UCLIN command chapter in *z/OS SMP/E Commands*.

name
> is the ddname of the data set to be allocated.
>
> - The name can contain from 1 to 8 alphabetic (A through Z), national (@, #, or $), or numeric (0 through 9) characters. The first character must be alphabetic or national.
> - Other than checking for duplicate DDDEF names in a given zone, SMP/E does not check whether the specified name is associated with another data set. For example, SMP/E does not check whether a DDDEF name is the same as a zone name. You must do this checking yourself to avoid undesired results.

BLK(*size*), CYL, or TRK
> specifies the space units to be used in allocating the data set: blocks, cylinders, or tracks.
>
> The UCL operand is BLOCK(*size*), CYLINDERS, or TRACKS.
>
> - *size* is the size, in decimal, of each block to be allocated. To specify the number of blocks to be allocated, use SPACE.
> - You can specify either the long form or the short form of these operands:
>   - BLOCK or BLK
>   - CYLINDERS or CYL
>   - TRACKS or TRK
> - These operands are mutually exclusive with each other and with OLD, MOD, SHR, CONCAT, PATH, and SYSOUT.

CATALOG, DELETE, or KEEP
> specifies the final disposition (DISP) of the data set.
>
> The UCL operand is CATALOG, DELETE, or KEEP.
>
> - These operands are mutually exclusive with each other and with CONCAT and PATH.
> - You cannot specify a final disposition for SMPTLIB data sets. SMP/E automatically specifies a final disposition based on the command being processed. For more information, see "SMPTLIB" on page 131.

CONCAT
identifies one or more DDDEF entries, **existing in the same zone**, that should be concatenated during SMP/E processing.

The UCL operand is CONCAT(*name...*).

- The DDDEF names can contain from 1 to 8 alphabetic (A through Z), national (@, #, or $), or numeric (0 through 9) characters. The first character must be alphabetic or national.
- There must be a DDDEF entry for each specified name. SMP/E does not check other sources (such as DD statements or GIMDDALC control statements in SMPPARM member GIMDDALC) to get the information needed to allocate the data sets.
- SMP/E allows you to specify up to 123 names. However, the actual number of partitioned data sets that can be concatenated depends on the operating system you are running under. To determine the maximum number of data sets you can concatenate, see *z/OS DFSMS Using Data Sets*.
- CONCAT is required for concatenated data sets.
- CONCAT cannot specify any DDDEF entries that contain the CONCAT operand. (You cannot nest concatenated DDDEF entries.)
- CONCAT is mutually exclusive with all other operands.

  Note: When a DDDEF that specifies a disposition of OLD is allocated as part of a CONCAT, a disposition of SHR is used for the allocation.

DATACLAS
specifies the name of a data class to be used for allocating a new data set managed by SMS.

The UCL operand is DATACLAS(*name*).

- The data class name is defined by the storage administrator at your installation. This value can contain from 1 to 8 alphanumeric characters (A through Z and 0 through 9) or national characters (@, #, $) and must start with either an alphabetic or national character.
- You can specify this operand only if SMS manages your storage.
- Specify this operand only when you are allocating a new data set. SMS ignores this parameter if it is specified for an existing data set.
- This operand is mutually exclusive with PATH and CONCAT.

DATASET
is the name of the data set to be allocated.

The UCL operand is DATASET(*dsname*).

- You can specify either DATASET or DA.
- The data set name must conform to standard naming conventions for data sets. Each part of the name must contain from 1 to 8 characters, separated from the other parts by a period (.), with no intervening blanks. The maximum length of the entire name is 44 characters (including the periods).
- The data set name itself cannot contain parentheses.
- To define a dummy data set, omit the DATASET operand and specify an initial DISP of NEW.
- DATASET is required if OLD, SHR, or MOD is specified.
- DATASET is mutually exclusive with CONCAT, PATH, and SYSOUT.
- You cannot use DATASET to specify the data set name of an SMPTLIB data set. Instead, you can use the DSPREFIX operand in either the SMPTLIB DDDEF entry or in the OPTIONS entry used to process those data sets.
- SMP/E does not check whether the specified data set name is unique within the zone. For example, SMP/E does not check whether the data set name was also defined in another DDDEF entry in the same zone, or whether the data set name defines the CSI data set containing the zone. You must do this checking yourself to avoid undesired results.

DIR
specifies the number of directory blocks to allocate.

The UCL operand is `DIR(`*nnnn*`)`.

- The number specified can contain from 1 to 4 decimal digits.
- DIR is mutually exclusive with OLD, MOD, SHR, CONCAT, and PATH.

DSNTYPE
specifies the type of partitioned data set to be created.

The UCL operand is `DSNTYPE(LIBRARY)` or `DSNTYPE(PDS)`.

LIBRARY
specifies that a PDSE (which must be an SMS-managed data set) is to be created.

PDS
specifies that a PDS is to be created.

You can also specify DSNTYPE in either a data class or a member of SYS1.PARMLIB.

Note:

1. This operand is mutually exclusive with PATH and CONCAT.
2. When SMP/E RECEIVE processing allocates a new SMPTLIB data set, it uses the original DSNTYPE of the corresponding RELFILE data set. If SMP/E cannot determine the original DSNTYPE of the corresponding RELFILE data set, SMP/E uses the DSNTYPE value specified in the SMPTLIB DDDEF entry.

DSPREFIX
specifies the data set prefix to be used to construct the full data set name for SMPTLIB data sets. For more information about names for SMPTLIB data sets, see *z/OS SMP/E Commands*.

The UCL operand is `DSPREFIX(`*prefix*`)`.

- The prefix can contain from 1 to 26 alphanumeric characters.
- The prefix must follow standard conventions for naming data sets.
- Instead of specifying DSPREFIX in the SMPTLIB DDDEF entry, you can specify it in the OPTIONS entry that is in effect when you receive RELFILEs into the SMPTLIB data sets.

  If you do not specify a data set prefix in the SMPTLIB DDDEF entry or in the appropriate OPTIONS entry, no prefix is included when SMP/E assigns a name to the SMPTLIB data sets.

- DSPREFIX is mutually exclusive with CONCAT and PATH.

Note: This subentry exists only in the global zone.

MGMTCLAS
specifies the name of a management class to be used for allocating a new data set managed by SMS.

The UCL operand is `MGMTCLAS(`*name*`)`.

- The management class name is defined by the storage administrator at your installation. This value can contain from 1 to 8 alphanumeric characters (A through Z and 0 through 9) or national characters (@, #, $) and must start with either an alphabetic or national character.
- You can specify this operand only if SMS manages your storage.
- Specify this operand only when you are allocating a new data set. SMS ignores this parameter if it is specified for an existing data set.
- This operand is mutually exclusive with PATH and CONCAT.

MOD, NEW, OLD, or SHR
specifies the initial disposition (DISP) of the data set.

The UCL operand is `MOD`, `NEW`, `OLD`, or `SHR`.

- These operands are mutually exclusive with each other and with CONCAT and PATH.

  OLD, SHR, and MOD are also mutually exclusive with SYSOUT, PROTECT, BLOCK, CYLINDERS, TRACKS, DIR, and SPACE.

DDDEF entry (distribution, target, and global zone)

- You cannot specify an initial disposition for SMPTLIB data sets. SMP/E automatically specifies an initial disposition, which is based on the command being processed. For more information, see "SMPTLIB" on page 131.

PATH

identifies the name of the path to be allocated in a UNIX file system. The name is not a complete pathname; it is a directory. This value is concatenated with the appropriate element or load module name to create a complete pathname.

The UCL operand is PATH(*pathname*).

- PATH is mutually exclusive with all other DDDEF entry operands.
- The pathname can be from 1 to 255 characters.
- The pathname must begin and end with a slash (/).
- In addition to the required delimiters (/), a pathname must also be enclosed in single apostrophes (')
  if any of the following is true:
  - The pathname contains lowercase alphabetic characters.
  - The pathname contains a character that is not uppercase alphabetic, numeric, national ($, #, or @), slash (/), plus (+), hyphen, period, or ampersand (&).
  - The pathname spans more than one line in the UCL control statement.

  The apostrophes must be outside the required delimiters, as in *'/pathname/'*, not */'pathname'/*.

  The single apostrophes used to enclose a pathname (the delimiters) do not count as part of the 255-character limit.
- Any apostrophes specified as part of a pathname (not the delimiters) must be doubled.

  Double apostrophes count as two characters in the 255-character limit.
- The pathname can include characters X'40' through X'FE'.
- Because symbolic substitution is not detected by SMP/E, it is not supported by SMP/E.

PROTECT

specifies that the z/OS SecureWay Security Server PROTECT option should be used when a new data set is first allocated. If PROTECT is specified and z/OS SecureWay Security Server is installed, the data set allocated by SMP/E will be protected by z/OS SecureWay Security Server.

You can also use PROTECT to indicate that an existing data set is protected by z/OS SecureWay Security Server. In this case, because SMP/E does not allocate the data set, it does not check the PROTECT indicator. However, you can use it to keep a record of which data sets have been protected with z/OS SecureWay Security Server.

The UCL operand is PROTECT.

- PROTECT is mutually exclusive with CONCAT, PATH, OLD, MOD, and SHR.

SPACE

specifies the primary and secondary space allocation for new data sets.

The UCL operand is SPACE(*prime,second*).

- Each value must contain from 1 to 4 decimal digits, and the two values must be separated by a comma or a blank.
- SPACE is mutually exclusive with CONCAT, PATH, OLD, MOD, and SHR.

STORCLAS

specifies the name of a storage class used for allocating a new data set managed by SMS.

The UCL operand is STORCLAS(*name*).

- The storage class name is defined by the storage administrator at your installation. This value can contain from 1 to 8 alphanumeric characters (A through Z and 0 through 9) or national characters (@, #, $) and must start with either an alphabetic or national character.

- You can specify this operand only if SMS manages your storage.
- Specify this operand only when you are allocating a new data set. SMS ignores this parameter if it is specified for an existing data set.
- This operand is mutually exclusive with PATH and CONCAT.

SYSOUT

specifies the output class for SYSOUT data sets.

The UCL operand is SYSOUT(*value*).

- If you specify a class, it must be 1 alphabetic or numeric character (A through Z or 0 through 9).
- If you specify *, the output class depends on references to an OUTPUT JCL statement.
    - If there is an implicit or explicit reference to an OUTPUT JCL statement, the output is written to the same class as the CLASS parameter on the OUTPUT statement.
    - If there is no reference to an OUTPUT JCL statement, the output is written to the same class as the one specified as MSGCLASS on the JOB card.
- SYSOUT is mutually exclusive with CATALOG, DELETE, KEEP, DIR, DSNTYPE, UNIT, SPACE, VOLUME, BLOCK, CYLINDER, TRACK, CONCAT, PATH, and DATASET.
- You cannot specify SYSOUT for SMPTLIB data sets.

UNIT

specifies the UNIT type the data set resides on if it is not cataloged.

The UCL operand is UNIT(*type*).

- If the data set is not cataloged, you must specify UNIT (unless it is not cataloged because of SMS). If you specify UNIT for a cataloged data set, it overrides the value in the catalog.
- The UNIT value can contain from 1 to 8 characters and should conform to standard UNIT naming conventions.

    SMP/E accepts any nonblank characters specified between the open and close parentheses, up to a maximum length of 8.

- UNIT is mutually exclusive with CONCAT and PATH.

VOLUME

specifies the volume serial number of the volume that the data set resides on if not cataloged.

The UCL operand is VOLUME(*volid...*).

- If the data set is not cataloged, you must specify VOLUME (unless it is not cataloged because of SMS). If you specify VOLUME for a cataloged data set, it overrides the value in the catalog.
- The volume identifier can contain from 1 to 6 alphanumeric characters.
- For SMPTLIB data sets, you can specify up to five volume serial numbers. All the volumes must have the same UNIT type. For other data sets, you can specify only one volume serial number.
- VOLUME is mutually exclusive with CONCAT and PATH.

WAIT=YES or WAIT= NO

indicates whether SMP/E should wait for the data set to be allocated if the volume is not mounted or if the data set is already in use. Not waiting causes allocation to fail for the data set.

The UCL operand is WAITFORDSN.

- You can specify either WAITFORDSN or WAIT.
- WAIT is mutually exclusive with CONCAT and PATH.
- If you do not specify a value, the default is not to wait.
- WAIT is not related to the PROCESS parameter specified on the EXEC statement. PROCESS affects how long a job should wait for a data set before being run. For more information, see .

LIST Examples

To list all the DDDEF entries in a particular zone, you can use the following commands:

```
SET      BDY(TGT1)          /* Set to requested target. */.
LIST     DDDEF              /* List all DDDEF entries.  */.
```

To list specific DDDEF entries in a particular zone, you can use these commands:

```
SET      BDY(TGT1)          /* Set to requested target. */.
LIST     DDDEF(SMPMTS       /* List only these three    */
             MACLIB         /* entries.                 */
             SYSLIB)        /*                          */.
```

The format of the LIST output for each DDDEF entry is the same for both of these commands. The only difference is the number of DDDEF entries listed.

Figure 17 on page 168 and Figure 18 on page 169 are examples of LIST output for DDDEF entries.

```
PAGE nnnn  - NOW SET TO zzzzzz ZONE nnnnnnn  DATE mm/dd/yy TIME hh:mm:ss SMP/E 36.nn SMPLIST OUTPUT

TGT1            DDDEF ENTRIES


   NAME

AMACLIB   DATASET        = SYS1.AMACLIB
          VOLUME         = DLIB01
          UNIT           = 3380
          SHR

BPXLIB1   PATH           = '/path_name1/'

CMACLIB   DATASET        = SYS1.MACLIB
          SHR

MACLIB    DATASET        = SYS1.MACLIB
          OLD

SMPMTS    DATASET        = SYS1.SMPMTS
          OLD

SMPLOG    DATASET        = SYS1.SMPLOG
          MOD

SYSLIB    CONCAT         = SMPMTS    CMACLIB    AMACLIB

SMPOUT    SYSOUT         = A

SMPWRK1   UNIT           = SYSDA
          SPACE          = (25,25)
          DIR            = 25
          ALLOC          = TRK
          NEW
          DELETE

SMPWRK2   DATACLAS       = FB80CLAS
          MGMTCLAS       = SMPEMCLS
          STORCLAS       = SMPESCLS

SMPTLIB   VOLUME         = DLIB01    DLIB02    DLIB03
          UNIT           = 3380
```

Figure 17. DDDEF entry: sample LIST output for a target zone

```
PAGE nnnn  - NOW SET TO zzzzzz ZONE nnnnnnn  DATE mm/dd/yy TIME hh:mm:ss SMP/E 36.nn SMPLIST OUTPUT

GLOBAL        DDDEF ENTRIES


   NAME

SMPPTS    DATASET         = SYS1.SMPPTS
          OLD

SMPLOG    DATASET         = SYS1.GLOBAL.SMPLOG
          MOD

SMPOUT    SYSOUT          = A

SMPWRK1   UNIT            = SYSDA
          SPACE           = (25,25)
          DIR             = 25
          ALLOC           = TRK
          NEW
          DELETE

SMPWRK2   DATACLAS        = FB80CLAS
          MGMTCLAS        = SMPEMCLS
          STORCLAS        = SMPESCLS

SMPTLIB   VOLUME          = DLIB01    DLIB02    DLIB03
          UNIT            = 3380
          DSPREFIX        = C87MVSP
          PROTECT
```

Figure 18. DDDEF entry: sample LIST output for a global zone

UNLOAD Examples

To dump the DDDEF entries in UCL format, you can use the UNLOAD command. To unload all the DDDEF entries in a particular zone, you can use the following commands:

```
SET     BDY(TGT1)           /* Set to requested target.  */.
UNLOAD  DDDEF               /* Unload all DDDEF entries.  */.
```

To unload specific DDDEF entries in a particular zone, you can use these commands:

```
SET     BDY(TGT1)           /* Set to requested target.  */.
UNLOAD  DDDEF(SMPMTS         /* Unload only these three   */
            MACLIB          /* entries.                  */
            SYSLIB)         /*                           */.
```

The format of the UNLOAD output for each DDDEF entry is the same for both of these commands. The only difference is the number of DDDEF entries unloaded.

Note: You can use the UNLOAD command only for target and distribution zones, not for the global zone.

is an example of UNLOAD output for DDDEF entries.

```
UCLIN .
REP       DDDEF             ( AMACLIB  )
          DATASET           ( SYS1.AMACLIB                                      )
          VOLUME            ( DLIB01 )
          UNIT              ( 3380    )
          SHR
                            .
REP       DDDEF             ( BPXLIB1  )
          PATH              ( '/path_name1/' )
                            .
REP       DDDEF             ( CMACLIB  )
          DATASET           ( SYS1.MACLIB                                       )
          SHR
                            .
REP       DDDEF             ( MACLIB   )
          DATASET           ( SYS1.MACLIB                                       )
          OLD
                            .
REP       DDDEF             ( SMPDATA1 )
          DATASET           ( MVSTGT1.SMPDATA1  )                               )
          MOD
                            .
REP       DDDEF             ( SMPDATA2 )
          DATASET           ( MVSTGT1.SMPDATA2  )                               )
          MOD
                            .
REP       DDDEF             ( SMPMTS   )
          DATASET           ( SYS1.SMPMTS                                       )
          OLD
                            .
REP       DDDEF             ( SMPLOG   )
          DATASET           ( SYS1.SMPLOG                                       )
          MOD
                            .
REP       DDDEF             ( SYSLIB   )
          CONCAT            ( SMPMTS    CMACLIB   AMACLIB  )
                            .
                            .
ENDUCL.
```

Figure 19. DDDEF entry: sample UNLOAD output

UCLIN Examples

You can use the ADD, REP, and DEL UCL statements to change subentries in a DDDEF entry. When you use UCLIN to update a DDDEF entry, keep these points in mind:

- If a DEL statement deletes all the existing subentries in the entry, SMP/E deletes the entire entry.
- When SMP/E dynamically allocates a concatenation list, the order of the concatenation is the same as that specified in the DDDEF entry. Therefore, you cannot use ADD statements to update a DDDEF entry that already contains a concatenation list. SMP/E would not know the correct position for the new data.

For example, suppose you have a DDDEF entry for the SYSLIB data set that concatenates SMPMTS and MACLIB. You **cannot** use the following commands to add XYZMACS at the end of the list:

```
SET      BDY(DLIB1)          /* Set to DLIB1 zone.   */.
UCLIN                        /*                      */.
ADD      DDDEF(SYSLIB)       /*                      */
         CONCAT(XYZMACS)     /*                      */.
ENDUCL                       /*                      */.
```

Instead, you can use these commands to replace the entire concatenation list:

```
SET      BDY(DLIB1)          /* Set to DLIB1 zone.   */.
UCLIN                        /*                      */.
REP      DDDEF(SYSLIB)       /*                      */
         CONCAT(SMPMTS,      /* Replace entire list. */
                MACLIB,      /*                      */
                XYZMACS)     /*                      */.
ENDUCL                       /*                      */.
```

The following examples are provided to help you use the DDDEF entry.

Example 1: Defining global zone DDDEFs for cataloged data sets

For this example, assume you want to define a DDDEF entry for the SMPPTS and SMPLOG. The data set names are SYS1.SMPPTS and SYS1.GLOBAL.SMPLOG, and both data sets are cataloged on the system that you are running on. Here is an example of how to define the DDDEF entries:

```
SET      BDY(GLOBAL)      /* Set to global zone.   */.
UCLIN                     /*                       */.
ADD      DDDEF(SMPPTS)    /* DDDEF for PTS.        */
         DA(SYS1.SMPPTS)  /* Data set is cataloged. */
         OLD              /* DISP=OLD.             */
ADD      DDDEF(SMPLOG )   /* DDDEF for LOG.        */
         DA(SYS1.SMPLOG)  /* Data set is cataloged. */
         MOD              /* DISP=MOD.             */.
ENDUCL                    /*                       */.
```

Example 2: Defining DLIB zone DDDEFs for cataloged data sets

Assume you want SMP/E to dynamically allocate distribution libraries AOS12 and ASAMPLIB during ACCEPT. Assuming these libraries are cataloged on the system you are running SMP/E on, you can use the following commands to define the DDDEF entries:

```
SET      BDY(DLIB1)        /* Set to DLIB zone.    */.
UCLIN                      /*                      */.
ADD      DDDEF(AOS12)      /* DLIB ddname.         */
         DA(SYS1.AOS12)    /* Assume cataloged.    */
         OLD               /* OLD for update.      */.
ADD      DDDEF(ASAMPLIB)   /* DLIB ddname.         */
         DA(SYS1.ASAMPLIB) /* Assume cataloged.    */
         OLD               /* OLD for update.      */.
ENDUCL                     /*                      */.
```

Note: Both data sets have a disposition of OLD. This is because these data sets are used for output when SMP/E calls the system utilities, and you want only one job to update the data sets at a time.

Example 3: Defining target zone DDDEFs for cataloged data sets

Assume you want SMP/E to dynamically allocate data sets MODGEN, MACLIB, SMPMTS, and SMPLOG during APPLY. Assuming these data sets are cataloged on the system that you are running SMP/E on, you can use the following commands to define the DDDEF entries:

```
SET      BDY(TGT1)        /* Set to target zone.   */.
UCLIN                     /*                       */.
ADD      DDDEF(MODGEN )   /* For concatenation.    */
         DA(SYS1.MODGEN)  /* Data set is cataloged. */
         SHR              /* SHR for read.         */.
ADD      DDDEF(MACLIB )   /* In case updated.      */
         DA(SYS1.MACLIB)  /* Data set is cataloged. */
         OLD              /* OLD during update.    */.
ADD      DDDEF(SMPMTS )   /* For update.           */
         DA(SYS1.SMPMTS)  /* Data set is cataloged. */
         OLD              /* OLD for update.       */.
ADD      DDDEF(SMPLOG )   /* For update.           */
         DA(SYS1.SMPLOG)  /* Data set is cataloged. */
         MOD              /* MOD for log.          */.
ENDUCL                    /*                       */.
```

Note: MACLIB and SMPMTS have a disposition of OLD. This is because these data sets are used for output when SMP/E calls the system utilities, and you want only one job to update the data sets at a time.

Example 4: Defining a DDDEF for a noncataloged data set

During APPLY processing, SMP/E may need to refer to some distribution library data sets not cataloged on the system that SMP/E is running on. For example, suppose AMODGEN is a distribution macro library

not cataloged on the running system, and you want to use AMODGEN in the SYSLIB concatenation for assemblies. You can use the following commands to define the DDDEF entry:

```
SET     BDY(TGT1)           /* Set to target zone.      */.
UCLIN                       /*                          */.
ADD     DDDEF(AMODGEN )     /* For concatenation.       */
        DA(SYS1.AMODGEN)    /* Data set not cataloged.  */
        VOLUME(DLIB01)      /*                          */
        UNIT(3330-1)        /*                          */
        SHR                 /* SHR for read.            */.
ENDUCL                      /*                          */.
```

Note: This data set has a disposition of SHR. This is because it is not updated during APPLY, but rather is used just for input.

Example 5: Defining a concatenated DDDEF entry

After you have defined DDDEF entries for all the data sets to be included in the SYSLIB concatenation, you can define the DDDEF entry that concatenates the data sets. This entry must specify the concatenation ddname and the order of concatenation. The ddname is the name of the DDDEF entry, and the order of concatenation is the order in which the DDDEF names are specified on the CONCAT operand.

You can use the following commands to define the DDDEF entry:

```
SET     BDY(TGT1)           /* Set to target zone.       */.
UCLIN                       /*                           */.
ADD     DDDEF(SYSLIB)       /* ddname for concatenation. */
        CONCAT(             /* Concatenation order is    */
            SMPMTS          /* SMPMTS first,             */
            MACLIB          /* then MACLIB,              */
            MODGEN          /* then MODGEN,              */
            AMODGEN         /* then AMODGEN.             */
            )               /* End of list.              */.
ENDUCL                      /*                           */.
```

The SYSLIB concatenation is the same as if the following JCL were used:

```
//SYSLIB   DD DSN=SYS1.SMPMTS,DISP=OLD
//         DD DSN=SYS1.MACLIB,DISP=OLD
//         DD DSN=SYS1.MODGEN,DISP=OLD
//         DD DSN=SYS1.AMODGEN,DISP=SHR,
//            UNIT=3330-1,VOL=SER=DLIB01
```

Note: SYS1.MACLIB was allocated as DISP=OLD, even though it is not updated during APPLY processing. This is to limit access by any other job while SMP/E is running. If you wanted to use a disposition of SHR when MACLIB is part of the concatenation, you can use the following commands to define an additional DDDEF entry for MACLIB and change the SYSLIB DDDEF entry:

```
UCLIN                       /*                           */.
ADD     DDDEF(CMACLIB)      /* For read access.          */
        DA(SYS1.MACLIB)     /* Data set is cataloged.    */
        SHR                 /* Change DISP to SHR.       */.
ADD     DDDEF(SYSLIB)       /* ddname for concatenation. */
        CONCAT(             /* concatenation order is    */
            SMPMTS          /* SMPMTS first,             */
            CMACLIB         /* ***** NOTE CHANGE *****    */
            MODGEN          /* then MODGEN,              */
            AMODGEN         /* then AMODGEN.             */
            )               /* End of list.              */.
ENDUCL                      /*                           */.
```

This corresponds to the following JCL:

```
//SYSLIB   DD DSN=SYS1.SMPMTS,DISP=OLD
//         DD DSN=SYS1.CMACLIB,DISP=SHR
//         DD DSN=SYS1.MODGEN,DISP=OLD
//         DD DSN=SYS1.AMODGEN,DISP=SHR,
//            UNIT=3330-1,VOL=SER=DLIB01
```

SYS1.MACLIB is now allocated with DISP=SHR so other jobs can still access it.

Example 6: Defining DDDEFs for temporary data sets

Assume you want SMP/E to dynamically allocate the SMPOUT and SMPWRK1 data sets during ACCEPT. You can use the following commands to define the DDDEF entries:

```
SET      BDY(DLIB1)        /* Set to DLIB zone.         */.
UCLIN                      /*                           */.
ADD      DDDEF(SMPOUT)     /* SMPOUT ddname.            */
         SYSOUT(A)         /* SYSOUT class.             */.
ADD      DDDEF(SMPWRK1)    /* SMPWRK1 ddname.           */
         NEW               /* New data set.             */
         DELETE            /* Delete when finished.     */
         TRACKS            /* Allocate in tracks.       */
         SPACE(25,25)      /* prime and secondary space. */
         DIR(25)           /* 25 directory blocks.      */
         UNIT(SYSDA)       /* Allocate on SYSDA.        */.
ADD      DDDEF(SMPWRK2)    /* SMPWRK2 ddname.           */
         DATACLAS(FB80CLAS) /* Add DATACLAS name.       */
         MGMTCLAS(SMPEMCLS) /* Add MGMTCLAS name.       */
         STORCLAS(SMPESCLS) /* Add STORCLAS name.       */
ENDUCL                     /*                           */.
```

Example 7: Defining a global zone DDDEF for SMPTLIB data sets

Assume you want SMP/E to dynamically allocate the SMPTLIB data sets when it receives SYSMODs packaged in relative files. In addition, you want these data sets to be protected by z/OS SecureWay Security Server. You can use the following commands to define the DDDEF entry:

```
SET      BDY(GLOBAL)       /* Set to global zone.       */.
UCLIN                      /*                           */.
ADD      DDDEF(SMPTLIB)    /* SMPTLIB ddname.           */
         DATACLAS(FB80CLAS) /* Add DATACLAS name.       */
         DSPREFIX(C87MVSP) /* Data set prefix.          */
         MGMTCLAS(SMPEMCLS) /* Add MGMTCLAS name.       */
         STORCLAS(SMPESCLS) /* Add STORCLAS name.       */
         PROTECT           /* Request SecureWay Security */
                           /* Server protection.        */.
ENDUCL                     /*                           */.
```

Note: No initial or final disposition is specified. This is because when SMP/E dynamically allocates SMPTLIB data sets, it automatically specifies the initial and final disposition based on the command it is processing.

Example 8: Defining a DLIB zone DDDEF for SMPTLIB data sets

Assume you want SMP/E to dynamically allocate the SMPTLIB data sets when it accepts SYSMODs packaged in relative files. You can use the following commands to define the DDDEF entry:

```
SET      BDY(DLIB1)        /* Set to DLIB zone.         */.
UCLIN                      /*                           */.
ADD      DDDEF(SMPTLIB)    /* SMPTLIB ddname.           */
         VOLUME(DLIB01     /* Volumes used - DLIB01     */
                DLIB02     /* DLIB02                    */
                DLIB03)    /* DLIB03.                   */
         UNIT(3380)        /* Allocate on 3380.         */.
ENDUCL                     /*                           */.
```

Note: No initial or final disposition is specified. This is because when SMP/E dynamically allocates SMPTLIB data sets, it automatically specifies the initial and final disposition based on the command it is processing.

Example 9: Protecting data sets

Assume you have previously defined your SMPMTS and SMPSTS data sets outside of SMP/E and protected them with z/OS SecureWay Security Server. You now want to have these dynamically allocated, and you

also want to keep a record that they are protected by z/OS SecureWay Security Server. You can use the following commands to define the DDDEF entries:

```
SET     BDY(TGT1)            /* Set to target zone.                   */.
UCLIN                        /*                                       */.
ADD     DDDEF(SMPMTS)        /* SMPMTS ddname.                        */
        DA(SYS1.SMPMTS)      /* Data set is cataloged.                */
        OLD                  /* OLD for update.                       */
        PROTECT              /* Protected by SecureWay Security Server.*/.
ADD     DDDEF(SMPSTS)        /* SMPSTS ddname.                        */
        DA(SYS1.SMPSTS)      /* Data set is cataloged.                */
        OLD                  /* OLD for update.                       */
        PROTECT              /* Protected by SecureWay Security Server.*/.
ENDUCL                       /*                                       */.
```

Example 10: Defining pathnames in a UNIX file system

Assume a product you plan to add to your system contains load modules and elements that will reside in a UNIX file system. You want to have the pathnames dynamically allocated instead of having DD statements for them in a cataloged procedure. You can use the following commands to define the DDDEF entry for each pathname:

```
----+----1----+----2----   …   ----5----+----6----+----7--
SET     BDY(TGT1)            /* Set to target zone.     */.
UCLIN                        /*                         */.
ADD     DDDEF(BPXLIB1)       /* Specify DDDEF name.     */
        PATH('/path_name1/') /* Define pathname.        */
                             /*                         */.
ADD     DDDEF(BPXLIB2)       /* Specify DDDEF name.     */
        PATH('/path_name2/This/pathname/is/an/example/of/a/
very/long/pathname/It/shows/a/long/name/')
                             /* define pathname.        */
                             /*                         */.
ENDUCL                       /*                         */.
```

DLIB entry (distribution and target zone)

The DLIB entry describes a distribution library that was totally copied to a target library. It is created during JCLIN processing when SMP/E encounters a COPY or COPYMOD statement without any SELECT statements. The INDD value on the COPY or COPYMOD statement becomes the name of the DLIB entry, and the OUTDD value on the COPY or COPYMOD statement becomes the SYSLIB value in the DLIB entry. For additional information, see the "Processing" section of the JCLIN command chapter in *z/OS SMP/E Commands*.

SMP/E uses the DLIB entry to determine where elements and load modules should be applied, as follows:

- **Element entries other than MOD.** If there is no SYSLIB value either in the element entry or on the element MCS, SMP/E looks for a DLIB entry whose name matches the DISTLIB value for the element. SMP/E applies the element to the library indicated by the SYSLIB value in the DLIB entry, and adds that SYSLIB value to the element entry.

- **MOD entries.** For MOD entries, processing is slightly different, because MOD entries do not contain a SYSLIB subentry.

  - For a MOD entry that points to an LMOD entry, SMP/E checks whether the DISTLIB value in the MOD entry matches the name of a DLIB entry. If there is a match, and if the SYSLIB value in the DLIB entry is different from the ones in the LMOD entry, the SYSLIB value in the DLIB entry is used in addition to the LMOD's SYSLIB values.

  - For a MOD entry that does not point to an LMOD entry, SMP/E checks whether the DISTLIB value in the MOD entry matches the name of a DLIB entry. If there is a match, SMP/E assumes a load module with the same name as the MOD entry should exist in the library specified by the SYSLIB value in the DLIB entry.

    - If there is already an LMOD entry with the same name as the MOD entry, and that LMOD entry does not already contain two SYSLIB values, SMP/E adds the SYSLIB value from the DLIB entry to the LMOD entry.

- If there is no LMOD entry with the same name as the MOD entry, SMP/E creates one and adds the SYSLIB value from the DLIB entry to the LMOD entry.

Lastly, SMP/E updates the MOD entry by adding an LMOD subentry for the LMOD entry that was updated or created.

Subentries

These are the subentries for the DLIB entry as they appear in the LIST output:

name
is the name of the distribution library represented by the DLIB entry.

The name can contain from 1 to 8 alphanumeric characters.

LASTUPD
identifies the cause of the last change to this DLIB entry.

The UCL operand is LASTUPD(*value*). This subentry can contain one of the following values:

JCLIN
indicates that the change was made during JCLIN command processing.

UCLIN
indicates that the change was made as a result of UCLIN processing.

sysmod_id
indicates that the change was made during the installation of the indicated SYSMOD.

The SYSMOD ID must contain 7 alphanumeric characters.

LASTUPD TYPE
indicates how the entry was last changed.

The UCL operand is LASTUPDTYPE(*value*). This subentry can contain one of the following values:

ADD
The entry has been added.

UPD
The entry has been updated.

SYSTEM LIBRARIES
identifies the ddnames of the target libraries into which the distribution library should be copied.

The UCL operand is SYSLIB(*ddname...*).

- A DLIB entry used for modules can contain one or two SYSLIB values.
- A DLIB entry used for macros, source, or data elements should contain only one SYSLIB value.
- The ddnames can contain from 1 to 8 alphanumeric characters.

LIST Examples

To list all the DLIB entries in a particular zone, you can use the following commands:

```
SET     BDY(TGT1)           /* Set to requested zone.  */.
LIST    DLIB                /* List all DLIB entries.  */.
```

To list specific DLIB entries, you can use these commands:

```
SET     BDY(TGT1)           /* Set to requested zone.  */.
LIST    DLIB(AMACLIB        /* List only these two     */
             MACDLB01)      /* entries.                */.
```

The format of the LIST output for each DLIB entry is the same for both of these commands. The only difference is the number of DLIB entries listed. is an example of LIST output for DLIB entries.

DLIB entry (distribution and target zone)

```
PAGE nnnn  - NOW SET TO zzzzzz ZONE nnnnnnn  DATE mm/dd/yy TIME hh:mm:ss SMP/E 36.nn SMPLIST OUTPUT

TGT1         DLIB ENTRIES


  NAME

AMACLIB   SYSTEM LIBRARY  = MACLIB

MACDLB01  SYSTEM LIBRARY  = MACLIB     MACTGT01

MODDLB01  LASTUPD         = JXY1102 TYPE=ADD
          SYSTEM LIBRARY  = LPALIB     MODTGT01

MODDLB02  LASTUPD         = JXY1121 TYPE=UPD
          SYSTEM LIBRARY  = LPALIB     MODTGT02
```

Figure 20. DLIB entry: sample LIST output

UNLOAD Examples

To dump the DLIB entries in UCL format, you can use the UNLOAD command. To unload all the DLIB entries in a particular zone, you can use the following commands:

```
SET     BDY(TGT1)           /* Set to requested zone.   */.
UNLOAD  DLIB                /* Unload all DLIB entries. */.
```

To unload specific DLIB entries, you can use these commands:

```
SET     BDY(TGT1)           /* Set to requested zone.   */.
UNLOAD  DLIB(AMACLIB        /* Unload only these two    */
            MACDLB01)       /* entries.                 */.
```

The format of the UNLOAD output for each DLIB entry is the same for both of these commands. The only difference is the number of DLIB entries listed. is an example of UNLOAD output for DLIB entries.

```
UCLIN .
REP       DLIB             ( AMACLIB  )
          SYSLIB           ( MACLIB   )
                             .
REP       DLIB             ( MACDLB01 )
          SYSLIB           ( MACLIB     MACTGT01 )
                             .
REP       DLIB             ( MODDLB01 )
          LASTUPD          ( JXY1102 )
          LASTUPDTYPE      ( ADD )
          SYSLIB           ( LPALIB     MODTGT01 )
                             .
REP       DLIB             ( MODDLB02 )
          LASTUPD          ( JXY1121  )
          LASTUPDTYPE      ( UPD )
          SYSLIB           ( LPALIB     MODTGT02 )
                             .
ENDUCL .
```

Figure 21. DLIB entry: sample UNLOAD output

UCLIN Examples

You can use the ADD, REP, and DEL UCL statements to change subentries in the DLIB entry. When you use UCLIN to update a DLIB entry, keep these points in mind:

- After the UCLIN changes are made, the DLIB entry must contain at least a SYSLIB subentry. Otherwise, there is not enough information in the entry to indicate where elements should be installed.

- If a DEL statement deletes all the existing subentries in the entry, SMP/E deletes the entire entry.

Example: Changing the destination of a copied library

Assume you are doing a system generation and have completed the following steps:

1. Performed stage 1 system generation

2. Allocated new target zone

3. Copied distribution zone to new target zone

4. Ran SMP/E JCLIN command

You now want to change the target library into which a module distribution library has been totally copied. Assume that the DLIB was MODDLB01, that it was previously copied to LINKLIB, and that it is now copied to LPALIB. You can do this with the following UCL statements:

```
SET      BDY(TGT1)        /* Set to target zone.      */.
UCLIN                     /*                          */.
REP      DLIB(MODDLB01)   /* Module DLIB.             */
         SYSLIB(LPALIB)   /* Now copied to LPALIB.    */
                          /* No longer in LINKLIB.    */.
ENDUCL                    /*                          */.
```

If the modules should **also** be copied into LPALIB, the following UCL can be used:

```
SET      BDY(TGT1)        /* Set to target zone.      */.
UCLIN                     /*                          */.
ADD      DLIB(MODDLB01)   /* Module DLIB.             */
         SYSLIB(LPALIB)   /* Now copied to LPALIB.    */
                          /* SYSLIB added, not replaced.*/.
ENDUCL                    /*                          */.
```

In both cases, you must take care of copying the modules from the distribution library. SMP/E does not process elements when it records UCL changes.

Now that the UCL changes are made, when SMP/E installs SYSMODs that update MODDLB01, SMP/E checks the DLIB entry for MODDLB01, determines that the modules are copied to LPALIB, and creates or modifies the LMOD entries (with the same name as the MOD entries) to indicate a copy to LPALIB.

Note: This technique works before you actually install any service. If you want to change (not add) another library in the DLIB copy list, you must find all the MOD, MAC, SRC, and LMOD entries with a SYSLIB value matching the old target library, and then change that value using UCLIN. If you do not do this, SMP/E continues to copy entries with the old SYSLIB value to the old library (and also, in some cases, to the new library).

DLIBZONE entry (distribution zone)

The DLIBZONE entry contains information SMP/E uses to process a specific distribution zone and the associated distribution libraries. It is created by UCLIN and must be defined before you can do any other processing for that distribution zone.

Subentries

These are the subentries for the DLIBZONE entry as they appear in the LIST output:

name
> is the name of the distribution zone. You assign the name when the zone is created.
>
> The name can contain from one to seven alphanumeric characters (A through Z, 0 through 9) or national characters ($, #, @). The first character must be alphabetic.

ACCJCLIN
> indicates that JCLIN is to be saved in the distribution zone whenever a SYSMOD containing inline JCLIN is accepted.

DLIBZONE entry (distribution zone)

The UCL operand is ACCJCLIN.

- ACCJCLIN should be specified only for zones in which all the products have been installed with SMP/E Release 3, or later.
- To save inline JCLIN for a product at ACCEPT time, you must first accept that product using SMP/E Release 3, or later, and have ACCJCLIN set in the DLIBZONE entry. From then on, you must keep ACCJCLIN set in the DLIBZONE entry. This ensures that any time you accept service for that product, its JCLIN is updated in the distribution zone. For more information, see the section on inline JCLIN in the ACCEPT command chapter in *z/OS SMP/E Commands*.

OPTIONS

is the name of the OPTIONS entry in the global zone that should be used when processing this distribution zone. For more information, see "OPTIONS entry (global zone)" on page 241.

The UCL operand is OPTIONS(*name*).

- The name can contain from one to eight alphanumeric characters.
- This name can be overridden by using the OPTIONS parameter on the SET command. For more information about the SET command, see *z/OS SMP/E Commands*.
- If no OPTIONS entry name is specified, SMP/E uses a set of default utility values when processing this distribution zone. For more information, see "OPTIONS entry (global zone)" on page 241.

RELATED

is the name of the target zone to which this distribution zone is related. A distribution zone is related to the target zone that was built from these distribution libraries, such as during system generation.

The UCL operand is RELATED(*zone*).

- The zone name can contain from one to seven alphanumeric characters.
- Although the entry can be defined without this subentry, you **must** define the subentry before you can install any SYSMODs into the distribution libraries.

SREL

lists the system releases to be supported within this distribution zone.

The UCL operand is SREL(*srel...*).

- The SREL must contain four alphanumeric characters, usually one alphabetic character followed by three numeric characters. These are the SRELs defined by IBM:

   System
       SREL

   DB2
       P115

   CICS
       C150

   IMS
       P115

   MVS
       Z038

   NCP
       P004

- Although the entry can be defined without this subentry, you **must** define the subentry before you can install any SYSMODs in the distribution libraries.

Note: Although you can support multiple products with different SREL values from one distribution zone, those products are still subject to all other restrictions related to combining products in one zone. The most common reason for not being able to combine products is common element names. For example, modules or macros with the same name are found in both products, but reside in different libraries.

UPGLEVEL

indicates the highest SMP/E release level that is allowed to make incompatible changes to the distribution zone. Before making an incompatible change to the distribution zone, SMP/E will check the UPGLEVEL value for that zone. If the release level of SMP/E is higher than the zone's UPGLEVEL value, SMP/E will not make the incompatible change.

The UPGLEVEL value is in the form *vr.pp*, where *vr* represents the version and release of SMP/E and *pp* represents the PTF level of SMP/E.

There is no UCL support for this subentry. When a zone is created by SMP/E using the UCLIN command or the Administration dialog, SMP/E sets the UPGLEVEL subentry value for that zone to the level of SMP/E used to create the zone. The UPGRADE command is used to change the UPGLEVEL subentry value for a zone.

ZDESC

is a user-written description for this zone.

The UCL operand is ZONEDESCRIPTION(*text*).

- The zone description can be in single-byte characters (such as English alphanumeric characters) or in double-byte characters (such as Kanji).
- The zone description can contain up to 500 bytes of data, including blanks. (For double-byte data, the 500-byte maximum includes all shift-in and shift-out characters, as well as the double-byte characters.) Extra blanks are deleted. All data beyond column 72 is ignored, including blanks.
- The zone description cannot be only blanks.
- If parentheses are included in the text, they must be in matched pairs.

LIST Examples

To list the DLIBZONE entry for a particular distribution zone, you can use the following commands:

```
SET      BDY(DLIB1)            /* Set to requested DLIB.   */.
LIST     DLIBZONE              /* List DLIBZONE entry.     */.
```

Figure 22 on page 179 is an example of LIST output for a DLIBZONE entry.

```
PAGE nnnn  - NOW SET TO zzzzzz ZONE nnnnnnn  DATE mm/dd/yy TIME hh:mm:ss SMP/E 36.nn SMPLIST OUTPUT

DLIB1         DZONE ENTRY


  NAME

DLIB1    DZONE          = DLIB1
         ZDESC          = ZONE DESCRIPTION FOR DLIB1 ZONE
         RELATED        = TGT2
         SREL           = Z038
         OPTIONS        = OPTDLIB2
         ACCJCLIN
```

Figure 22. DLIBZONE entry: sample LIST output

UCLIN Examples

You can use the ADD, REP, and DEL UCL statements to change subentries in the DLIBZONE entry. When you use UCLIN to update a DLIBZONE entry, remember that if a DEL statement deletes all the existing subentries in the entry, SMP/E deletes the entire entry.

The following examples are provided to help you use the DLIBZONE entry.

Example 1: Defining a DLIBZONE entry

Assume you are about to add a new function to your system and want to define a separate distribution zone for it named DLIB2. The OPTIONS entry you plan to use is OPTDLIB2, and the SREL for the SYSMOD is Z038. The distribution zone is in data set SMPE.SMPCSI.CSI. After installing the function in the distribution libraries, you plan to do a system generation to build a set of target libraries. The target zone for those libraries is TGT2. The following UCLIN can be used to define the new zone.

```
SET       BDY(GLOBAL)        /* Set to global zone.      */.
UCLIN                        /* UCLIN for GZONE entry    */.
ADD       GZONE              /* to set up                */
          ZONEINDEX(         /* index for new zone.      */
                  (DLIB2,SMPE.SMPCSI.CSI,DLIB) /*        */
                )            /*                          */
                             /*                          */.
ENDUCL                       /* End global zone update.  */.
```

```
SET       BDY(DLIB2)         /* Now define new zone.     */.
UCLIN                        /* UCLIN to define it.      */.
ADD       DLIBZONE(DLIB2)    /* Identify name.           */
          OPTIONS(OPTDLIB2)  /* OPTIONS entry to use.    */
          SREL(Z038)         /* SREL for MVS.            */
          RELATED(TGT2)      /* For this tgt library.    */
          ZDESC(
               THIS DISTRIBUTION ZONE IS FOR
                SREL Z038
               )             /* Zone description.        */
                             /*                          */.
ENDUCL                       /*                          */.
```

Note: Even if the OPTIONS entry or TARGETZONE entry has not yet been defined, you can still refer to it in the DLIBZONE entry. However, you must create the entry before you process this distribution zone. For examples of defining OPTIONS entries, see "OPTIONS entry (global zone)" on page 241. For examples of defining TARGETZONE entries, see "TARGETZONE entry (target zone)" on page 286.

Example 2: Formatting a zone description

Assume you enter the following zone description with the first line ending in column 72 and the second line starting in column 1:

```
----+----1----+----2----  …  ----5----+----6----+----7--
SET       BDY(DLIB1)         /* Set to DLIB zone.        */.
UCLIN                        /* UCLIN for DZONE entry    */.
ADD       DZONE(DLIB1)       /* to set up.               */
          ZDESC(             THIS IS THE DESCRIPTION FOR
THE DLIB1 ZONE)
                             /* End of zone description. */.
ENDUCL                       /* End DLIB zone update.    */.
```

Because there is no blank between the word ending in column 72 and the next word starting in column 1, SMP/E runs the two together.

The words in a zone description, even words that end in column 72, must be separated by blanks. To format the zone description in this example correctly, you can put a blank at the beginning of the second line:

```
----+----1----+----2----  …  ----5----+----6----+----7--
SET       BDY(DLIB1)         /* Set to DLIB zone.        */.
UCLIN                        /* UCLIN for DZONE entry    */.
ADD       DZONE(DLIB1)       /* to set up.               */
          ZDESC(             THIS IS THE DESCRIPTION FOR
 THE  DLIB1 ZONE)
                             /* End of zone description. */.
ENDUCL                       /* End DLIB zone update.    */.
```

Because there is a blank explicitly coded between the word ending in column 72 and the word starting in column 1, SMP/E does not run the words together.

FEATURE entry (global zone)

The FEATURE entry describes a collection of function SYSMODS called a feature. An SMP/E feature may be associated with an orderable software feature.

Subentries

name
is the name of this feature.

DESCRIPTION
describes the software feature.

The UCL operand is DESCRIPTION(_description_).

FMID
specifies the FMIDs that make up this software feature. There can be multiple FMID subentries associated with a single feature.

The UCL operand is FMID(_fmid,..._).

PRODUCT
identifies the product with which this feature is associated.

The UCL operand is PRODUCT(_prodid,vv.rr.mm_).

RECDATE
specifies the date on which the feature was received.

There is no UCL support for this subentry.

RECTIME
specifies the time at which the feature was received.

There is no UCL support for this subentry.

UCLDATE
specifies the date on which the feature was last processed using the UCLIN command.

There is no UCL support for this subentry.

UCLTIME
specifies the time at which the feature was last processed using the UCLIN command.

There is no UCL support for this subentry.

REWORK
identifies the level of this feature, which was received again for minor changes.

There is no UCL support for this subentry.

LIST Examples

To list all the FEATURE entries in a particular zone, you can use the following commands:

```
SET      BDY(GLOBAL)        /* Set to global zone.      */.
LIST     FEATURE            /* List all FEATURE entries.*/.
```

To list specific FEATURE entries, you can use these commands:

```
SET      BDY(GLOBAL)        /* Set to global zone.      */.
LIST     FEATURE(OS3250BA   /* List only these two      */
               OS3250LD)    /* entries.                 */.
```

The format of the LIST output for each FEATURE entry is the same for both of these commands. The only difference is the number of FEATURE entries listed. shows an example of LIST output for FEATURE entries. FEATURE entries are listed alphanumerically by name.

```
PAGE nnnn  - NOW SET TO zzzzzz ZONE nnnnnnn  DATE mm/dd/yy TIME hh:mm:ss SMP/E 36.nn SMPLIST OUTPUT

GLOBAL  FEATURE  ENTRIES


  NAME

OS3250BA  DESCRIPTION     = OS/390 Base
          DATE/TIME REC   = 07.267  12:55:39
                     UCL  = 07.031  02:25:45
          PRODUCT         = 5647-A01  02.05.00
          FMID            = HBB6605   HMP1B00

OS3250DD  DESCRIPTION     = OpenEdition DCE User Privacy DES
          DATE/TIME REC   = 07.267  12:55:41
                     UCL  = 07.031  02:25:45
          PRODUCT         = 5647-A01  02.05.00
          FMID            = JMB3125

OS3250LD  DESCRIPTION     = Language Environment Decryption
          DATE/TIME REC   = 07.267  12:55:42
                     UCL  = 07.031  02:25:45
          PRODUCT         = 5647-A01  02.05.00
          FMID            = JMWL755
```

Figure 23. FEATURE entry: sample LIST output

By specifying the FORFMID operand, you can reduce the number of FEATURE entries listed. When FORFMID is specified, SMP/E lists a FEATURE entry only if its FMID matches one of the FMIDs specified on the FORFMID operand. For example, to list FEATURE entries whose FMIDs either are defined in FMIDSET ZOSSET or are HMP1E00, you can use these commands:

```
SET     BDY(GLOBAL)        /* Set to global zone.     */.
LIST    FEATURE            /* List all FEATURE entries */
        FORFMID(ZOSSET     /* for the ZOSSET FMIDSET   */
               )  /* and FMID .        */.
```

UCLIN Examples

You can use the ADD, REP, and DEL UCL statements to change subentries in the FEATURE entry. When you use UCLIN to update a FEATURE entry, keep in mind that after the UCLIN changes are made, the FEATURE entry must contain at least the DESCRIPTION and PRODUCT subentries.

Example: Adding a FEATURE entry

Assume you have a feature, named Sample Feature, for which you want to create a description and put in an entry named SAMPFEAT. You could set up a FEATURE entry as follows:

```
SET BDY(GLOBAL)                      /* Set to global zone      */.
UCLIN                                /* Start UCLIN processing  */.
  ADD FEATURE(SAMPFEAT)              /* Identify the feature:   */
      DESCRIPTION(Sample Feature)    /* - Name                  */
      PRODUCT(1234-ABC,02.05.00)     /* - Owning product        */
      FMID(XYZ1234)                  /* - FMID                  */.
ENDUCL                               /* End UCLIN processing    */.
```

FMIDSET entry (global zone)

The FMIDSET entry defines a group of FMIDs for use in limiting the SYSMODs processed by an SMP/E command. For example, you can specify an FMIDSET on the FORFMID operand of the APPLY command to process only SYSMODs applicable to one of the FMIDs in the FMIDSET.

Subentries

These are the subentries for the FMIDSET entry as they appear in the LIST output:

name
>   is the name of the FMIDSET.
>
>   The name can contain from 1 to 8 alphanumeric characters.

FMID
>   lists the function SYSMODs (that is, FMIDs) that are to be part of this FMIDSET.
>
>   The UCL operand is FMID(*sysmod_id...*).
>
>   The SYSMOD ID must contain 7 alphanumeric characters.

LIST Examples

To list all the FMIDSET entries in a global zone, you can use the following commands:

```
SET      BDY(GLOBAL)          /* FMIDSET in global only.   */.
LIST     FMIDSET              /* list all FMIDSET entries.  */.
```

To list specific FMIDSET entries in a global zone, you can use these commands:

```
SET      BDY(GLOBAL)          /* FMIDSET in global only.   */.
LIST     FMIDSET(ALLSET       /* List only these two       */
                 XXSET)       /* entries.                  */.
```

The format of the LIST output for each FMIDSET entry is the same for both of these commands. The only difference is the number of FMIDSET entries listed.

Figure 24 on page 183 is an example of LIST output for FMIDSET entries.

```
PAGE nnnn  - NOW SET TO zzzzzz ZONE nnnnnnn  DATE mm/dd/yy TIME hh:mm:ss SMP/E 36.nn SMPLIST OUTPUT

GLOBAL        FMIDSET ENTRIES


  NAME

ALLSET     FMID          = FXX1102   JXX1121   JXX1122   JXX1123
                          FYY1102   JYY1121   JYY1122   JYY1123
                          FZZ1102   JZZ1121   JZZ1122   JZZ1123

XXSET      FMID          = FXX1102   JXX1121   JXX1122   JXX1123

YYSET      FMID          = FYY1102   JYY1121   JYY1122   JYY1123

ZZSET      FMID          = FZZ1102   JZZ1121   JZZ1122   JZZ1123
```

Figure 24. FMIDSET entry: sample LIST output

UCLIN Examples

You can use the ADD, REP, and DEL UCL statements to change subentries in an FMIDSET entry. When you use UCLIN to update an FMIDSET entry, keep these points in mind:

- After the UCLIN changes are made, the FMIDSET entry must contain at least an FMID subentry. Otherwise there is not enough information in the entry for SMP/E to use the entry.

- If a DEL statement deletes all the existing subentries in the entry, SMP/E deletes the entire entry.

The following examples are provided to help you use the FMIDSET entry.

Example 1: Defining an FMIDSET entry

Assume you have four function SYSMODs installed, JXX1102, JXX1121, JXX1122, and JXX1123, and that when installing service you would like to be able to install just the PTFs for these products. You could set up an FMIDSET entry as follows:

```
SET       BDY(GLOBAL)        /* FMIDSETs only in global. */.
UCLIN                        /*                          */.
ADD       FMIDSET(XXSET)     /* Define FMIDSET name.     */
          FMID(FXX1102       /* Define function SYSMODs  */
               JXX1121       /* to be part of set.       */
               JXX1122       /*                          */
               JXX1123)      /*                          */
                             /*                          */.
ENDUCL                       /*                          */.
```

You can use the FORFMID operand on the APPLY as follows:

```
SET       BDY(TGT1)          /* Set for APPLY to TGT1.   */.
APPLY     FORFMID(XXSET)     /* Apply XX PTFs.           */.
```

Example 2: Modifying an FMIDSET entry

Once an FMIDSET has been defined, it can be modified as new function SYSMODs are installed or as existing function SYSMODs are replaced. For example, starting with the FMIDSET from the previous example, assume you install a new feature, JXX1124, that deletes JXX1121, and that you want to modify the FMIDSET definition. You can do this as follows:

```
SET       BDY(GLOBAL)        /* FMIDSETs only in global. */.
UCLIN                        /*                          */.
REP       FMIDSET(XXSET)     /* Define FMIDSET name.     */
          FMID(FXX1102       /* Define function SYSMODs. */
                             /* JXX1121 left out.        */
               JXX1122       /*                          */
               JXX1123       /*                          */
               JXX1124)      /* JXX1124 added.           */
                             /*                          */.
ENDUCL                       /*                          */.
```

You can also do the following:

```
SET       BDY(GLOBAL)        /* FMIDSETs only in global. */.
UCLIN                        /*                          */.
DEL       FMIDSET(XXSET)     /* Define FMIDSET name.     */
          FMID(JXX1121)      /* Delete JXX1121 from set. */
                             /*                          */.
ADD       FMIDSET(XXSET)     /* Define FMIDSET name.     */
          FMID(JXX1124)      /* Add new FMID.            */
                             /*                          */.
ENDUCL                       /*                          */.
```

The result in both cases is the same. The choice of which method to use most likely depends on the number of FMIDs defined in the FMIDSET.

GLOBALZONE entry (global zone)

The GLOBALZONE entry contains processing-related information for SMP/E. It is also used by SMP/E as an index to target and distribution zones, either in the same CSI or a different CSI data set. The GLOBALZONE entry is created by UCLIN and must be defined before you can do any other processing for that global zone.

Subentries

These are the subentries for the GLOBALZONE entry as they appear in the LIST output:

FMID

    lists the function SYSMODs for which SMP/E is to receive service.

    The UCL operand is FMID(*sysmod_id...*).

    The SYSMOD ID must contain seven alphanumeric characters.

OPTIONS
> is the name of the OPTIONS entry in the global zone that should be used in processing this global zone. For more information, see "OPTIONS entry (global zone)" on page 241.
>
> The UCL operand is OPTIONS(*name*).
>
> - The name can contain from one to eight alphanumeric characters.
> - This name can be overridden by using the OPTIONS parameter on the SET command. For more information about the SET command, see *z/OS SMP/E Commands*.
> - If no OPTIONS entry name is specified, SMP/E uses a set of default utility values when processing the global zone. For more information, see "OPTIONS entry (global zone)" on page 241.

SREL
> lists the system releases to be supported in this global zone.
>
> The UCL operand is SREL(*srel...*).
>
> The SREL must contain four alphanumeric characters, usually one alphabetic character followed by three numerics. These are the SRELs defined by IBM:
>
> **System**
> > **SREL**
>
> **DB2**
> > P115
>
> **CICS**
> > C150
>
> **IMS**
> > P115
>
> **MVS**
> > Z038
>
> **NCP**
> > P004
>
> **Note:** Although you can support multiple products with different SREL values from one global zone, those products are still subject to all other restrictions related to combining products in one zone. The most common reason for not being able to combine products is common element names. For example, modules or macros with the same name are found in both products, but reside in different libraries.

UPGLEVEL
> indicates the highest SMP/E release level that is allowed to make incompatible changes to the global zone. Before making an incompatible change to the global zone, SMP/E will check the UPGLEVEL value for that zone. If the release level of SMP/E is higher than the zone's UPGLEVEL value, SMP/E will not make the incompatible change.
>
> The UPGLEVEL value is in the form *vr.pp*, where *vr* represents the version and release of SMP/E and *pp* represents the PTF level of SMP/E.
>
> There is no UCL support for this subentry. When a zone is created by SMP/E using the UCLIN command or the Administration dialog, SMP/E sets the UPGLEVEL subentry value for that zone to the level of SMP/E used to create the zone. The UPGRADE command is used to change the UPGLEVEL subentry value for a zone.

ZDESC
> is a user-written description for this zone.
>
> The UCL operand is ZONEDESCRIPTION(*text*).
>
> - The zone description can be in single-byte characters (such as English alphanumeric characters) or in double-byte characters (such as Kanji).

- The zone description can contain up to 500 bytes of data, including blanks. (For double-byte data, the 500-byte maximum includes all shift-in and shift-out characters, as well as the double-byte characters.) Extra blanks are deleted. All data beyond column 72 is ignored, including blanks.
- The zone description cannot be only blanks.
- If parentheses are included in the text, they must be in matched pairs.

ZONEINDEX

identifies all the target zones and distribution zones associated with this global zone. This list is used by SMP/E to determine the data set in which the zone resides and the type of zone.

For ADD and REP statements, the UCL operand is ZONEINDEX((*name,dsn,type*)...).

Note: You cannot use REP to replace a single zone in the ZONEINDEX list. REP will cause the entire list of zones to be replaced. Instead, you must first delete the zone from the ZONEINDEX list, and then add the zone to the list with the new SMPCSI data set name.

For DEL statements, the UCL operand is ZONEINDEX((*name*)...).

- ZONEINDEX can also be specified as ZINDEX.
- *name* is the name of the zone. The name can contain from one to seven alphanumeric characters (A through Z, 0 through 9) or national characters ($, #, @). The first character must be alphabetic.
- *dsn* is the fully qualified name of the data set in which the zone resides.
- *type* is the zone type, either TARGET or DLIB.

The ZONEINDEX is only a pointer to a zone. Changes you make to a ZONEINDEX do not affect the associated zone. Therefore, you cannot add, rename, or delete a zone by simply adding, replacing, or deleting its ZONEINDEX. However, these are some other commands you can use:

- **Adding a zone:** To add a zone, you can use the following commands, depending on the particular situation:
  - UCLIN to add a ZONEINDEX, then UCLIN to define the DZONE or TZONE entry
  - ZONECOPY
  - ZONERENAME
  - ZONEEXPORT, UCLIN for the ZONEINDEX, and ZONEIMPORT
- **Renaming a zone:** To rename a zone, you must use the following command:
  - ZONERENAME
- **Deleting a zone:** To delete a zone, you can use the following commands, depending on the particular situation:
  - ZONEDELETE
  - ZONEEXPORT

For more information about any of these commands, see the chapter on that command.

LIST Examples

To list all the GLOBALZONE entry, you can use the following commands:

```
SET     BDY(GLOBAL)        /* Set to global zone.    */.
LIST    GLOBALZONE         /* List GLOBALZONE entry.  */.
```

is an example of LIST output for a GLOBALZONE entry.

```
PAGE nnnn  - NOW SET TO zzzzzz ZONE nnnnnnn  DATE mm/dd/yy TIME hh:mm:ss SMP/E 36.nn SMPLIST OUTPUT

GLOBAL        GZONE ENTRY


  NAME

GLOBAL    ZONEINDEX      = TGT1     TARGET   SYS1.TGT.SMPCSI.CSI
                           TGT2     TARGET   SYS1.TGT.SMPCSI.CSI
                           DLIB1    DLIB     SYS1.DLIB.SMPCSI.CSI
                           DLIB2    DLIB     SYS1.DLIB.SMPCSI.CSI
          ZDESC          = ZONE DESCRIPTION FOR GLOBAL ZONE
          SREL           = Z038     I115       P115       R020
          FMID           = JXY1102  FXY1121    FXY1122    FXY1123
          OPTIONS        = OPTGBL
```

Figure 25. GLOBALZONE entry: sample LIST output

UCLIN Examples

You can use the ADD, REP, and DEL UCL statements to change subentries in the GLOBALZONE entry. When you use UCLIN to update a GLOBALZONE entry, keep these points in mind:

- After the UCLIN changes are made, the GLOBALZONE entry must contain at least one of these subentries: FMID, OPTIONS, SREL, or ZONEINDEX. Otherwise, there is not enough information in the entry for SMP/E to use the entry.

- If a DEL statement deletes all the existing subentries in the entry, SMP/E deletes the entire entry.

The following examples are provided to help you use the GLOBALZONE entry.

Example 1: Defining the GLOBALZONE entry

To perform any SMP/E processing, you must first define the GLOBALZONE entry. Given the following assumptions:

- You are supporting functions JXY1102, JXY1121, JXY1122, and JXY1123 from this global zone.

- You are supporting system releases Z038, P115, and R020, and I121.

- You are eventually going to define two target zones, TGT1 and TGT2, to exist on SYS1.TGT.SMPCSI.CSI, and two distribution zones, DLIB1 and DLIB2, to exist on SYS1.DLIB.SMPCSI.CSI.

- The OPTIONS entry used is OPTGBL.

The following UCLIN can be used to define the GLOBALZONE entry:

```
SET      BDY(GLOBAL)       /* Set to global zone.    */.
UCLIN                      /*                        */.
ADD      GLOBALZONE        /* Define global now.     */
         SREL(Z038,P115,   /* Identify SRELs.        */
             R020,I115)    /*                        */
         FMID(JXY1102,     /* Identify functions.    */
             JXY1121,      /*                        */
             JXY1122,      /*                        */
             JXY1123)      /*                        */
         OPTIONS(OPTGBL)   /* OPTIONS entry to be used.*/
         ZONEINDEX(        /*                        */
                 (TGT1,SYS1.TGT.SMPCSI.CSI,TARGET) /* */
                 (TGT2,SYS1.TGT.SMPCSI.CSI,TARGET) /* */
                 (DLIB1,SYS1.DLIB.SMPCSI.CSI,DLIB) /* */
                 (DLIB2,SYS1.DLIB.SMPCSI.CSI,DLIB) /* */
                 )         /*                        */
                           /*                        */.
ENDUCL                     /*                        */.
```

Note: Even though the OPTIONS entry has not been set up yet, you can still refer to it in the GLOBALZONE entry. The same is true for the target zone name. However, the OPTIONS entry must be created before the

global zone is processed. For examples of setting up the OPTIONS entries, see "OPTIONS entry (global zone)" on page 241.

Example 2: Deleting two ZONEINDEX entries

Assume you decide not to create the TGT1 target zone and the DLIB1 distribution zone defined in the ZONEINDEX of example 1. You should delete the entries for those two zones by using UCL as follows:

```
SET     BDY(GLOBAL)        /* Set to global zone.    */.
UCLIN                      /*                        */.
DEL     GLOBALZONE         /* Define global now.     */
        ZONEINDEX(         /*                        */
              (TGT1)       /* Delete TGT1 entry.     */
              (DLIB1)      /* Delete DLIB1 entry.    */
            )              /*                        */
                           /*                        */.
ENDUCL                     /*                        */.
```

Note: The CSI data set name and the zone type operands are not required for a delete, although the open and close parentheses around the zones are still required.

Example 3: Formatting a zone description

Assume you enter the following zone description with the first line ending in column 72 and the second line starting in column 1:

```
----+----1----+----2----  …  ----5----+----6----+----7--
SET     BDY(GLOBAL)        /* Set to global zone.    */.
UCLIN                      /* UCLIN for GZONE entry  */.
ADD     GZONE              /* to set up.             */
        ZDESC(             THIS IS THE DESCRIPTION FOR
THE GLOBAL ZONE)
                           /* End of zone description. */.
ENDUCL                     /* End global zone update.  */.
```

Because there is no blank between the word ending in column 72 and the next word starting in column 1, SMP/E runs the two together.

The words in a zone description, even words that end in column 72, must be separated by blanks. To format the zone description in this example correctly, you can put a blank at the beginning of the second line:

```
----+----1----+----2----  …  ----5----+----6----+----7--
SET     BDY(GLOBAL)        /* Set to global zone.    */.
UCLIN                      /* UCLIN for GZONE entry  */.
ADD     GZONE              /* to set up.             */
        ZDESC(             THIS IS THE DESCRIPTION FOR
 THE GLOBAL ZONE)
                           /* End of zone description. */.
ENDUCL                     /* End global zone update.  */.
```

Because there is a blank explicitly coded between the word ending in column 72 and the word starting in column 2, SMP/E does not run the words together.

Example 4: Renaming an SMPCSI data set and updating a ZONEINDEX

Assume you have a target zone named TGT1 that resides in a data set named SMPE.TEST.TGT1.CSI, and you decide to rename the data set to SYS1.PROD.TGT1.CSI. After renaming the data set, you must use UCLIN to update the ZONEINDEX for the target zone in the global zone, as follows:

```
SET     BDY(GLOBAL)        /* Set to global zone.    */.
UCLIN                      /* UCLIN for GZONE entry. */.
  DEL GZONE
  ZINDEX( (TGT1) )         /* Delete old TGT1 zone.  */.
  ADD GZONE
  ZINDEX( (TGT1, SYS1.PROD.TGT1.CSI, TARGET) )
                           /* Supply new TGT1 zone.  */.
ENDUCL                     /* End global zone update. */.
```

You cannot use REP to replace a single zone in the ZONEINDEX list. REP will cause the entire list of zones to be replaced. Instead, you must first delete the zone from the ZONEINDEX list, and then add the zone to the list with the new SMPCSI data set name.

Hierarchical file system element entry (distribution and target zone)

The hierarchical file system element entry describes an element that exists in a distribution library or a UNIX file system. A hierarchical file system entry is created the first time you install a SYSMOD containing an MCS for a hierarchical file system element that does not yet have an entry in the CSI data set.

SMP/E records the function and service level of the hierarchical file system element in the entry. Once a hierarchical file system element entry exists, it is updated as subsequent SYSMODs affecting the hierarchical file system element are installed.

Table 4 on page 21 shows the types of entries used for hierarchical file system elements. Some types of elements may be translated into several languages. In these cases, the entry type contains *xxx*, which represents the language used for the element. (If an element was not translated, the entry type does not contain any *xxx* value. Table 5 on page 22 shows the *xxx* values and the languages they represent.

Subentries

These are the subentries for the hierarchical file system element entries as they appear in the LIST output:

name
is the name of the hierarchical file system element represented by the entry.

The name can contain 1–8 uppercase alphabetic, numeric, or national ($, #, @) characters.

BINARY or TEXT
indicates the installation mode to be used when the HFS copy utility is invoked to install the element into a UNIX file system.

The UCL operand is BINARY or TEXT.

- *Binary* mode means that the element is installed in its entirety as a data stream, with no breaks for logical records.
- *Text* mode means that the element is installed with breakpoints for logical records.
- BINARY and TEXT are mutually exclusive.
- If there is no mode indicator in the hierarchical file system element entry, the HFS copy utility determines how to install the element.

DISTLIB
specifies the ddname of the distribution library for the hierarchical file system element.

The UCL operand is DISTLIB(*ddname*).

- The ddname can contain any uppercase alphabetic, numeric, or national ($, #, @) character, and can be 1–8 characters long.
- The DISTLIB subentry is required. Without it, SMP/E cannot process any changes for the hierarchical file system element.

FMID
identifies the functional owner of this hierarchical file system element. The functional owner is the last function SYSMOD that replaced this element.

The UCL operand is FMID(*sysmod_id*).

The SYSMOD ID must contain 7 uppercase alphabetic, numeric, or national ($, #, @) characters.

LASTUPD
identifies the cause of the last change to this hierarchical file system element entry.

The UCL operand is LASTUPD(*value*). This subentry can contain one of the following values:

UCLIN
> indicates that the change was made as a result of UCLIN processing.

sysmod-id
> indicates that the change was made during the installation of the indicated SYSMOD.
>
> The SYSMOD ID must contain 7 uppercase alphabetic, numeric, or national ($, #, @) characters.

LASTUPD TYPE
> indicates how the entry was last changed.
>
> The UCL operand is LASTUPDTYPE(*value*). This subentry can contain one of the following values:

ADD
> The entry was added.

UPD
> The entry was updated.

LINK
> specifies a list of alternative names by which this hierarchical file system element can be known in a UNIX file system. The full name is produced by concatenating the specified linkname with the UNIX file system directory identified by the SYSLIB subentry.
>
> In LIST output, linknames are always enclosed in single apostrophes. If an apostrophe is part of a linkname, it is always shown as two consecutive apostrophes in LIST output.
>
> The UCL operand is LINK(*linkname...*).
>
> - The linkname can be from 1 to 1023 characters.
> - A linkname can be enclosed in single apostrophes ('). A linkname **must** be enclosed in single apostrophes if any of the following is true:
>   - The linkname contains lowercase alphabetic characters.
>   - The linkname contains a character that is not uppercase alphabetic, numeric, national ($, #, or @), slash (/), plus (+), hyphen, period, or ampersand (&).
>   - The linkname spans more than one line in the control statement.
>
>   The single apostrophes used to enclose a linkname (the delimiters) do not count as part of the 1023-character limit.
> - Any apostrophes specified as part of a linkname (not the delimiters) must be doubled.
>
>   Double apostrophes count as two characters in the 1023-character limit.
> - The linkname can include characters X'40' through X'FE'.

PARM
> specifies a character string that is to be passed to the HFS copy utility as an execution-time parameter. The maximum length of this character string is 300 bytes of nonblank data. If any blanks are specified in the PARM value, they are deleted by SMP/E during processing and do not count toward the 300-byte maximum.
>
> The UCL operand is PARM(*character_string*).
>
> - PARM is an optional operand.
> - The character string can be entered free-form, without regard to blanks (which are compressed out of the string), and can span multiple 80-byte records.
> - If parentheses are specified in the PARM value, there must always be a pair (left and right); otherwise, the results are unpredictable.

RMID
> identifies the last SYSMOD that **replaced** this hierarchical file system element. Any subsequent SYSMOD that modifies this hierarchical file system element must have defined a relationship (such as PRE or SUP) with this SYSMOD.

The UCL operand is RMID(*sysmod_id*).

- The SYSMOD ID must contain 7 uppercase alphabetic, numeric, or national ($, #, @) characters.
- If RMID is not specified but FMID is, SMP/E sets the RMID value to the specified FMID.

SHSCRIPT
specifies a UNIX shell script to run when this element is copied to (or deleted from) a directory in a UNIX file system.

The UCL operand is SHSCRIPT(*scriptname*). This subentry optionally contains one or both of the following values, which specify the point in SMP/E processing when the shell script is run:

PRE
The shell script is run before the element is copied to the UNIX file system directory.

POST
The shell script is run after the element is copied to the UNIX file system directory.

POST is the default; the shell script is run after the element is copied to the directory.

SYMLINK
specifies a list of one or more symbolic links, which are file names that can be used as alternate names for referring to this element in a UNIX file system. Each symbolic link name listed here is associated with a path name listed in the SYMPATH entry. See the description of the SYMPATH entry for more information about how the symbolic link names and path names are associated.

The UCL operand is SYMLINK(*symlinkname...*).

- A hierarchical file system element entry that contains a SYMLINK entry must contain a matching SYMPATH entry. SMP/E will reject any UCLIN command that would violate this condition.
- A symbolic linkname can be from one to 1023 characters. Any characters in the range X'40' through X'FE' may be specified.
- A symbolic linkname can be enclosed in single apostrophes ('). A symbolic linkname **must** be enclosed in single apostrophes if any of the following is true:
  - The symbolic linkname contains lowercase alphabetic characters.
  - The symbolic linkname contains a character that is not uppercase alphabetic, numeric, national ($, #, or @), slash (/), plus (+), hyphen, period, or ampersand (&).
  - The symbolic linkname spans more than one line in the control statement.

  The single apostrophes used to enclose a symbolic linkname (the delimiters) do not count as part of the 1023-character limit.
- Any apostrophes specified as part of a symbolic linkname (not the delimiters) must be doubled. Double apostrophes count as two characters in the 1023-character limit.

SYMPATH
specifies a list of one or more pathnames that are associated with symbolic links identified by the SYMLINK operand. The first pathname in the SYMPATH operand is associated with the first symbolic link in the SYMLINK operand, the second pathname with the second symbolic link, and so on. If there are more symbolic links listed than there are pathnames, then the last listed pathname is used for the remaining symbolic links. If more pathnames are specified than symbolic linknames, then the excess pathnames (at the end of the list) are ignored. The UCL operand is SYMPATH(*sympathname...*).

- A hierarchical file system element entry that contains a SYMPATH entry must contain a matching SYMLINK entry. is specified, otherwise it must be omitted.
- A symbolic pathname can be from one to 1023 characters. Any characters in the range X'40' through X'FE' may be specified.
- A symbolic pathname can be enclosed in single apostrophes ('). A symbolic pathname **must** be enclosed in single apostrophes if any of the following is true:
  - The symbolic pathname contains lowercase alphabetic characters.

- The symbolic pathname contains a character that is not uppercase alphabetic, numeric, national ($, #, or @), slash (/), plus (+), hyphen, period, or ampersand (&).
- The symbolic pathname spans more than one line in the control statement.

The single apostrophes used to enclose a symbolic pathname (the delimiters) do not count as part of the 1023-character limit.

- Any apostrophes specified as part of a symbolic pathname (not the delimiters) must be doubled. Double apostrophes count as two characters in the 1023-character limit.

SYSLIB

specifies the ddname of the "target library" within a UNIX file system for the hierarchical file system element.

The UCL operand is SYSLIB(*ddname*).

- Only one SYSLIB value can be specified.
- The ddname can contain any uppercase alphabetic, numeric, or national ($, #, @) character, and can be 1 to 8 characters long.
- The SYSLIB subentry is required. Without it, SMP/E cannot process any changes for the hierarchical file system element.

LIST Examples

To list all the generic hierarchical file system element entries in a particular zone that support the Spanish language, you can use the following commands:

```
SET       BDY(TGT1)           /* Set to requested zone.   */.
LIST      HFSESP              /* List all generic HFS     */
                              /* element entries (Spanish)*/.
```

To list specific UNIX1 element entries, you can use these commands:

```
SET       BDY(TGT1)           /* Set to requested zone.   */.
LIST      UNIX1(UNXEL1        /* List only these          */
             UNXEL2           /* three                    */
             UNXEL3)          /* entries.                 */.
```

The format of the LIST output for each hierarchical file system element entry is the same for both of these commands. The only difference is the number of hierarchical file system element entries listed. shows an example of LIST output for UNIX1 element entries.

```
PAGE nnnn  - NOW SET TO zzzzzz ZONE nnnnnnn  DATE mm/dd/yy TIME hh:mm:ss SMP/E 36.nn SMPLIST OUTPUT

TGT1    UNIX1    ENTRIES


   NAME

UNXEL1   LASTUPD         = UNXFUNC TYPE=ADD
         LIBRARIES       = DISTLIB=APOSIXL1  SYSLIB=UNXTGTL1
         BINARY
         FMID            = UNXFUNC
         RMID            = UNXPTF1
         SHSCRIPT        = INSTAL01,PRE
         LINK            = 'linkname_1'
                         = 'linkname_2'

UNXEL2   LASTUPD         = UNXFUNC TYPE=ADD
         LIBRARIES       = DISTLIB=APOSIXL2  SYSLIB=UNXTGTL2
         TEXT
         FMID            = UNXFUNC
         RMID            = UNXFUNC
         LINK            = 'linkname_3'
                         = 'linkname_4'
         SHSCRIPT        = INSTAL02,POST
UNXEL3   LASTUPD         = UNXFUNC   TYPE=ADD
         LIBRARIES       = DISTLIB=APOSIXL2  SYSLIB=UNXTGTL2
         TEXT
         FMID            = UNXFUNC
         RMID            = UNXPTF2
         LINK            = 'linkname_5'
                           'linkname_6'
         SHSCRIPT        = INSTAL03,PRE,POST
         PARM            = This_is_another_sample_character_string_specified_as_the_value_of_PARM.__It_is_a_maximum____
                           length_character_string_for_this_subentry.__I.e.,_it_is_300_characters_long.__The_string____
                           has_no_blanks_in_it.__If_you_see_something_that_looks_like_a_blank,_it's_not.  The_previous_
                           2_characters_are_X'41'…
```

Figure 26. Hierarchical file system element entry: sample LIST output

By specifying the FORFMID operand, you can reduce the number of hierarchical file system element entries listed. When FORFMID is specified, SMP/E lists a hierarchical file system element entry only if its FMID matches one of the FMIDs specified on the FORFMID operand. For example, to list hierarchical file system element entries whose FMIDs either are defined in FMIDSET HFS or are HFSFUNC, you can use these commands:

```
SET     BDY(TGT1)          /* Set to target zone.     */.
LIST    HFS                /* List all hierarchical   */
                           /* file system element     */
                           /* entries                 */
        FORFMID(HFS        /* for the HFS FMIDSET      */
            HFSFUNC)       /* and FMID HFSFUNC.        */.
```

You can use the LIST command to find out the names of all SYSMODs that have modified a hierarchical file system element. To include the names of these SYSMODs in the LIST output you can use the XREF operand, as shown in these commands:

```
SET     BDY(TGT1)          /* Set to requested zone.  */.
LIST    HFS                /* List all hierarchical   */
                           /* file system element     */
                           /* entries                 */
        XREF               /* and related SYSMODs.    */.
```

Note:

1. XREF can be used either in mass mode or in select mode.

2. SMP/E obtains the data included for the XREF operand by checking all the SYSMOD entries for subentries for this hierarchical file system element. Because this data is not contained in the hierarchical file system element entry itself, you cannot use UCLIN to change it in the hierarchical file system element entry.

is an example of the LIST output produced when the XREF operand is used.

Hierarchical file system element entry (distribution and target zone)

```
PAGE nnnn  - NOW SET TO zzzzzz ZONE nnnnnnn  DATE mm/dd/yy TIME hh:mm:ss SMP/E 36.nn SMPLIST OUTPUT

TGT1   UNIX1    ENTRIES


   NAME

UNXEL1   LASTUPD        = UNXFUNC TYPE=ADD
         LIBRARIES      = DISTLIB=APOSIXL1  SYSLIB=UNXTGTL1
         BINARY
         FMID           = UNXFUNC
         RMID           = UNXPTF1
         LINK           = 'linkname_1'
                        = 'linkname_2'
         SHSCRIPT       = INSTAL01,PRE
         SYSMOD HISTORY = SYSMOD   TYPE     DATE    MCS     --STATUS--
                          UNXFUNC  FUNCTION 07.100  UNX     APP    ACC
                          UNXPTF1  PTF      07.120  UNX     APP    ACC

UNXEL2   LASTUPD        = UNXFUNC TYPE=ADD
         LIBRARIES      = DISTLIB=APOSIXL2  SYSLIB=UNXTGTL2
         TEXT
         FMID           = UNXFUNC
         RMID           = UNXFUNC
         LINK           = 'linkname_3'
                        = 'linkname_4'
         SHSCRIPT       = INSTAL02,POST
         SYSMOD HISTORY = SYSMOD   TYPE     DATE    MCS     --STATUS--
                          UNXFUNC  FUNCTION 07.100  UNX     APP    ACC
```

Figure 27. Hierarchical file system element entry: sample LIST output when XREF is specified

To list the UNIX shell script (SHELLSCR) element entries for a particular zone, you can use the following commands:

```
SET     BDY(TGT1)          /* Set to requested zone.  */.
LIST    SHELLSCR           /* List all UNIX shell     */
                           /* script element entries  */.
```

Figure 28 on page 194 shows an example of the LIST output for SHELLSCR element entries.

```
PAGE nnnn  - NOW SET TO TARGET ZONE zzzzzz  DATE mm/dd/yy TIME hh:mm:ss SMP/E 36.nn SMPLIST OUTPUT

TGT1   SHELLSCR   ENTRIES


   NAME

INSTAL01  LASTUPD        = FUNC001 TYPE=ADD
          LIBRARIES      = DISTLIB=HFSDLIB  SYSLIB=UNXTGTL1
          TEXT
          FMID           = FUNC001
          RMID           = FUNC001

INSTAL02  LASTUPD        = FUNC002 TYPE=UPD
          LIBRARIES      = DISTLIB=HFSDLIB  SYSLIB=UNXTGTL1
          TEXT
          FMID           = FUNC002
          RMID           = FUNC002


INSTAL03  LASTUPD        = FUNC001 TYPE=ADD
          LIBRARIES      = DISTLIB=HFSDLIB  SYSLIB=UNXTGTL1
          TEXT
          FMID           = FUNC001
          RMID           = FUNC001
```

Figure 28. Hierarchical file system element entry: sample LIST output for SHELLSCR entries

UNLOAD Examples

To dump the hierarchical file system element entries in UCL format, you can use the UNLOAD command. For example, to unload all the OS21 element entries in a particular zone, you can use the following commands:

```
SET       BDY(TGT1)           /* Set to requested zone.    */.
UNLOAD    OS21                /* Unload all OS21 element entries. */.
```

To unload specific OS21 element entries, you can use these commands:

```
SET       BDY(TGT1)           /* Set to requested zone.    */.
UNLOAD    OS21(OS2EL1         /* Unload only these         */
             OS2EL2           /* three                     */
             OS2EL3)          /* entries.                  */.
```

The format of the UNLOAD output for each OS21 element entry is the same for both of these commands. The only difference is the number of element entries listed. Figure 29 on page 195 is an example of UNLOAD output for OS21 element entries.

```
UCLIN .
REP         OS21              ( OS2EL1   )
            LASTUPD           ( OS2FUNC  )
            LASTUPDTYPE       ( ADD )
            DISTLIB           ( OS2DSTL1 )
            SYSLIB            ( OS2TGTL1 )
            FMID              ( OS2FUNC )
            RMID              ( OS2PTF1 )
            LINK              ( 'linkname_1'
                                'linkname_2' )
            SHSCRIPT          ( INST0S2,POST )
                              .
REP         OS21              ( OS2EL2   )
            LASTUPD           ( OS2FUNC  )
            LASTUPDTYPE       ( ADD )
            DISTLIB           ( OS2DSTL2 )
            SYSLIB            ( OS2TGTL2 )
            FMID              ( OS2FUNC )
            RMID              ( OS2FUNC )
            LINK              ( 'linkname_3'
                                'linkname_4' )
            SHSCRIPT          ( INST0S2,POST )
                              .
REP         OS21              ( OS2EL3 )
            LASTUPD           ( OS2FUNC )
            LASTUPDTYPE       ( ADD )
            DISTLIB           ( APOSIXL2 )
            SYSLIB            ( OS2TGTL2 )
            TEXT
            FMID              ( OS2FUNC )
            RMID              ( OS2PTF2 )
            LINK              ( 'linkname_5'
                                'linkname_6' )
            PARM              ( This_is_another_sample_character_string_spec
                                ified_as_the_value_of_PARM.__It_is_a_maximum
                                ____length_character_string_for_this_subentr
                                y.__I.e.,_it_is_300_characters_long.__The_st
                                ring____has_no_blanks_in_it.__If_you_see_som
                                ething_that_looks_like_a_blank,_it's_not.  T
                                he_previous_2_characters_are_X'41'…
                                )
                              .
ENDUCL.
```

Figure 29. OS21 element entry: sample UNLOAD output

By specifying the FORFMID operand, you can reduce the number of OS21 element entries unloaded. When FORFMID is specified, SMP/E unloads an OS21 element entry only if its FMID matches one of the FMIDs specified on the FORFMID operand. For example, to unload OS21 element entries whose FMIDs either are defined in FMIDSET OS2 or are OS2FUNC, you can use these commands:

```
SET     BDY(TGT1)           /* Set to target zone.        */.
UNLOAD  OS21                /* Unload all OS21 element entries */
        FORFMID(OS21        /* for the OS21 FMIDSET       */
              OS2FUNC)      /* and FMID OS2FUNC.          */.
```

UCLIN Examples

You can use the ADD, REP, and DEL UCL statements to change subentries in the hierarchical file system element entry. After the UCLIN changes are made, the hierarchical file system element entry must contain at least the following subentries:

• DISTLIB

• FMID

• RMID

• SYSLIB

Otherwise, there is not enough information in the entry to process the element. If any of these subentries are missing, SMP/E does not make the requested UCL updates to the entry, and the entry remains as it was before the UCL command.

The following examples are provided to help you use the hierarchical file system element entries.

Example 1: Adding a new hierarchical file system element entry

Defining a new hierarchical file system element entry with UCL is very seldom required; generally, hierarchical file system element entries are created from the information specified on the hierarchical file system element MCSs contained in the SYSMODs when the SYSMODs are installed. If, however, you want to use UCL in defining a new hierarchical file system element entry, the following is an example of the minimum information you should provide:

```
SET     BDY(TGT1)           /* Set to target zone.      */.
UCLIN                       /*                          */.
ADD     HFS(HFS01)          /* Define new entry.        */
        DISTLIB(AHFSLIB)    /* Define DLIB,             */
        SYSLIB(HFSLIB)      /* target library.          */
        FMID(ZUSR001)       /* Functional owner (in this
                               example a user function). */
        RMID(ZUSR001)       /* Same value as FMID.      */
                            /*                          */.
ENDUCL                      /*                          */.
SET     BDY(DLB1)           /* Now do same to DLIB.     */.
UCLIN                       /*                          */.
ADD     HFS(HFS01)          /* Define new HFS entry.    */
        DISTLIB(AHFSLIB)    /* Define DLIB,             */
        SYSLIB(HFSLIB)      /* target library.          */
        FMID(ZUSR001)       /* Functional owner (in this
                               example a user function). */
        RMID(ZUSR001)       /* Same value as FMID.      */
                            /*                          */.
ENDUCL                      /*                          */.
```

Note: If the RMID value had not been specified, it would have defaulted to the FMID value.

Example 2: Defining a linkname for an existing hierarchical file system element

The following shows how to inform SMP/E of new linkname for an existing hierarchical file system element:

```
SET     BDY(TGT1)           /* Set to target zone.    */.
UCLIN                       /*                        */.
ADD     HFS(HFS01)          /* Existing entry.        */
        LINK('../userlk')   /* New linkname.          */
                            /*                        */.
ADD     HFS(HFS02)          /* Existing entry.        */
        LINK('myname')      /* New linkname.          */
```

```
                              /*                    */.
ENDUCL                        /*                    */.
SET       BDY(DLB1)           /* Now do same thing to
                                 appropriate DLIB.  */.
UCLIN                         /*                    */.
ADD       HFS(HFS01)          /* Existing entry.    */
          LINK('../userlk')   /* New linkname.      */
                              /*                    */.
ADD       HFS(HFS02)          /* Existing entry.    */
          LINK('myname')      /* New linkname.      */
                              /*                    */.
ENDUCL                        /*                    */.
```

Note: UCLIN does not create a linkname in the UNIX file system; that must be done outside of SMP/E using standard utilities. The UCLIN changes ensure that, when the hierarchical file system element is subsequently modified, the hierarchical file system element is updated under both its primary name and its alternative name in the UNIX file system.

Example 3: Adding a hierarchical file system element entry with a PARM subentry

If you want to use UCL in adding a hierarchical file system element entry with a PARM subentry, the following example shows you how this can be done.

```
SET       BDY(TGT1).          /* Set to target zone.              */
UCLIN     .                   /*                                  */
ADD       HFS(HFSEL3) LASTUPD(HFSFUNC) LASTUPDTYPE(ADD)
          DISTLIB(APOSIXL2) SYSLIB(HFSTGTL2) FMID(HFSFUNC) TEXT
          RMID(HFSPTF2) LINK('linkname_5' 'linkname_6' )
          PARM(This_is_another_sample_character_string_specified_as_the_
value_of_PARM.__It_is_a_maximum____length_character_string_for_this_
subentry.__I.e.,_it_is_300_characters_long.__The_string____has_no_
blanks_in_it.__If_you_see_something_that_looks_like_a_blank,_it's_not.
  The_previous_2_characters_are_X'41'…) .
ENDUCL  .                     /*                                  */
```

HOLDDATA entry (global zone)

The HOLDDATA entry contains ++HOLD statements that either were received from SMPHOLD (external HOLDDATA) or were within a SYSMOD that was received (internal HOLDDATA). It is a record of the ++HOLD statements and is not used by SMP/E to control exception SYSMOD processing.

The HOLDDATA entry is separate from the SYSMOD entry. However, when SMP/E saves ++HOLD statements to create HOLDDATA entries, it also creates HOLDDATA subentries in the associated global zone SYSMOD entry. (These subentries appear in the LIST output as HOLDERROR, HOLDFIXCAT, HOLDSYSTEM, or HOLDUSER plus the hold reason IDs.) SMP/E uses these HOLDDATA subentries to control exception SYSMOD processing.

LIST Examples

To list all the HOLDDATA entries in a global zone, you can use the following commands:

```
SET       BDY(GLOBAL)        /* Set to requested zone.    */.
LIST      HOLDDATA           /* List all HOLDDATA entries. */.
```

To list HOLDDATA entries for a specific type of hold, code one or more of the following operands: HOLDERROR, HOLDSYSTEM, HOLDUSER. Here is an example:

```
SET       BDY(GLOBAL)        /* Set to requested zone.    */.
LIST      HOLDDATA           /* List HOLDDATA entries     */.
          HOLDSYSTEM         /* for system holds          */.
```

To list specific SYSMOD entries in the global zone along with the associated ++MCSs (HOLDDATA), you can use these commands:

HOLDDATA entry (global zone)

```
SET      BDY(GLOBAL)        /* Set to requested zone.   */.
LIST     SYSMOD(UZ12345     /* List only these two      */.
             UZ56789)       /* entries                  */.
         HOLDDATA           /* plus HOLDDATA for them.  */.
```

The output from the LIST requests differs in the amount of information presented. If you list just the HOLDDATA entries, only the HOLDDATA itself is presented, as in Figure 30 on page 198. If you request the HOLDDATA associated with SYSMOD entries, the HOLDDATA is included in the SYSMOD entry, as in Figure 31 on page 199. (For more information about listing HOLDDATA, see the description of the HOLDDATA operand of the LIST command in *z/OS SMP/E Commands*.)

Note:

1. The ++HOLD statements displayed in the LIST output are exactly as received by SMP/E. They are listed in alphanumeric order by reason ID.

2. If HOLDDATA for a particular SYSMOD has been received but the SYSMOD itself has not yet been received, only the hold information is shown when the SYSMOD entry is listed. For an example, see SYSMOD UZ56789 in Figure 31 on page 199.

3. HOLDDATA entries for system holds show either INT or EXT to indicate the source of the ++HOLD statement. INT means that the ++HOLD statement was contained in the held SYSMOD and can be resolved only with the BYPASS operand on the APPLY or ACCEPT command. EXT means that the ++HOLD statement was obtained from another source, such as SMPHOLD, and can be resolved by either the BYPASS operand or a ++RELEASE statement.

4. You cannot use UCLIN to add, update, or delete HOLDDATA entries or subentries. However, you can use the REJECT command to delete HOLDDATA entries. For more information about the REJECT command, see *z/OS SMP/E Commands*.

```
PAGE nnnn  - NOW SET TO zzzzzz ZONE nnnnnnn  DATE mm/dd/yy TIME hh:mm:ss SMP/E 36.nn SMPLIST OUTPUT

GLOBAL        HOLDDATA ENTRIES


  NAME

HBB7730   HOLDFIXCAT     = AA13644    ++HOLD(HBB7730) FMID(HBB7730) REASON(AA13644) FIXCAT
                                          RESOLVER(UA27033) CLASS(PSP) CATEGORY(IBM.Device.2094).
UZ12345   HOLDERROR      = AZ00001    ++HOLD(UZ12345) ERROR
                                          REASON(AZ00001) FMID(HBB2102)
                                          COMMENT(SMRTDATA(CHGDTE(071105) SYMP(PRV,IPL))).
          HOLDERROR      = AZ00002    ++HOLD(UZ12345) ERROR
                                          REASON(AZ00002) FMID(HBB2102)
                                          COMMENT(SMRTDATA(CHGDTE(071106) SYMP(FUL))).
          HOLDSYSTEM(INT) = DOC       ++HOLD(UZ12345) SYSTEM
                                          REASON(DOC) FMID(HBB2102)
                                          COMMENT(NEW MSG).
          HOLDSYSTEM(EXT) = UCLIN     ++HOLD(UZ12345) SYSTEM
                                          REASON(UCLIN) FMID(HBB2102)
                                          COMMENT(UCLIN REQUIRED).
          HOLDUSER       = INUSE      ++HOLD(UZ12345) USER
                                          REASON(INUSE) FMID(HBB2102)
                                          COMMENT(IM MODIFYING).

UZ56789   HOLDERROR      = AZ00023    ++HOLD(UZ56789) ERROR
                                          REASON(AZ00023) FMID(HBB2102)
                                          COMMENT(SMRTDATA(CHGDTE(071110) SYMP(ABENDS with E37))).
          HOLDERROR      = AZ00024    ++HOLD(UZ56789) ERROR
                                          REASON(AZ00024) FMID(HBB2102)
                                          COMMENT(SMRTDATA(CHGDTE(071113) SYMP(DAL) FIX(UW92458))).
```

Figure 30. HOLDDATA entry: sample LIST output when SYSMOD is not specified

```
PAGE nnnn  - NOW SET TO zzzzzz ZONE nnnnnnn  DATE mm/dd/yy TIME hh:mm:ss SMP/E 36.nn SMPLIST OUTPUT

GLOBAL        SYSMOD ENTRIES


  NAME

UZ12345   TYPE          = PTF
          STATUS        = REC
          DATE/TIME REC = 07.100  08:00:00
          SREL   VER(001) = Z038
          DELETE VER(001) = HBB2102
          SUPING VER(001) = AZ11111   AZ11112   AZ11113
          MAC           = MAC01     MAC02     MAC03
          MOD           = MOD01     MOD02     MOD03     MOD04
          SRC           = SRC01     SRC02
          HOLDERROR     = AZ00001   ++HOLD(UZ12345) ERROR
                                      REASON(AZ00001) FMID(HBB2102)
                                      COMMENT(SMRTDATA(CHGDTE(071105) SYMP(PRV,IPL))).
          HOLDERROR     = AZ00002   ++HOLD(UZ12345) ERROR
                                      REASON(AZ00002) FMID(HBB2102)
                                      COMMENT(SMRTDATA(CHGDTE(071106) SYMP(FUL))).
          HOLDSYSTEM(INT) = DOC     ++HOLD(UZ12345) SYSTEM
                                      REASON(DOC) FMID(HBB2102)
                                      COMMENT(NEW MSG).
          HOLDSYSTEM(EXT) = UCLIN   ++HOLD(UZ12345) SYSTEM
                                      REASON(UCLIN) FMID(HBB2102)
                                      COMMENT(UCLIN REQUIRED).
          HOLDUSER      = INUSE     ++HOLD(UZ12345) USER
                                      REASON(INUSE) FMID(HBB2102)
                                      COMMENT(IM MODIFYING).

UZ56789   HOLDERROR     = AZ00023   ++HOLD(UZ56789) ERROR
                                      REASON(AZ00023) FMID(HBB2102)
                                      COMMENT(SMRTDATA(CHGDTE(071110) SYMP(ABENDS with E37))).
          HOLDERROR     = AZ00024   ++HOLD(UZ56789) ERROR
                                      REASON(AZ00024) FMID(HBB2102)
                                      COMMENT(SMRTDATA(CHGDTE(071113) SYMP(DAL) FIX(UW92458))).
```

Figure 31. HOLDDATA listed for SYSMOD entry: sample LIST output when SYSMOD is specified

JAR entry (target and distribution zone)

The JAR entry describes a Java Archive (JAR) file that resides in a UNIX file system and a distribution library. A JAR entry is created the first time a SYSMOD is installed that contains a ++JAR MCS.

Subentries

These are the subentries for the JAR entry as they appear in the LIST output:

name
is the name of the Java ARchive element represented by the JAR entry.

The name can contain 1–8 uppercase alphabetic, numeric, or national ($, #, @) characters.

DISTLIB
specifies the ddname of the distribution library for the JAR element.

The UCL operand is DISTLIB(*ddname*).

- The ddname can contain any uppercase alphabetic, numeric, or national ($, #, @) character, and can be 1–8 characters long.
- The DISTLIB subentry is required. Without it, SMP/E cannot process any changes for the JAR element.

FMID
identifies the functional owner of this JAR element. The functional owner is the last function SYSMOD that replaced this JAR element.

The UCL operand is FMID(*sysmod_id*).

The SYSMOD ID must contain 7 uppercase alphabetic, numeric, or national ($, #, @) characters.

JARPARM
specifies a character string that is to be passed to the jar command as an option string when updating the JAR file. The maximum length of this character string is 300 bytes of data.

Any jar command options specified here will be passed to the jar command in addition to the options supplied by SMP/E, which are uvf (u indicates that the JAR file is to be updated, v produces verbose jar command output, and f indicates that the JAR file to be updated is specified on the command line rather than through *stdin*).

Only those jar command options that require only the option indicator to be specified are supported in the JARPARM entry, such as 0 and M. Options requiring additional input over and above the option indicator are not supported. Examples of unsupported options are the m and -C options. If such options are specified, the jar command will likely fail, since SMP/E does not prohibit their use.

LASTUPD
identifies the cause of the last change to this JAR entry.

The UCL operand is LASTUPD(*value*). This subentry can contain one of the following values:

UCLIN
indicates that the change was made as a result of UCLIN processing.

sysmod-id
indicates that the change was made during the installation of the indicated SYSMOD.

The SYSMOD ID must contain 7 uppercase alphabetic, numeric, or national ($, #, @) characters.

LASTUPDTYPE
indicates how the entry was last changed.

The UCL operand is LASTUPDTYPE(*value*). This subentry can contain one of the following values:

ADD
The entry was added.

UPD
The entry was updated.

LINK
specifies a list of alternative names by which this JAR element can be known in the UNIX file system.

In LIST output, linknames are always enclosed in single apostrophes. If an apostrophe is part of a linkname, it is always shown as two consecutive apostrophes in LIST output.

The UCL operand is LINK(*linkname...*).

- The linkname can be from 1 to 1023 characters.
- A linkname can be enclosed in single apostrophes ('). A linkname **must** be enclosed in single apostrophes if any of the following is true:
  - The linkname contains lowercase alphabetic characters.
  - The linkname contains a character that is not uppercase alphabetic, numeric, national ($, #, or @), slash (/), plus (+), hyphen, period, or ampersand (&).
  - The linkname spans more than one line in the control statement.

  The single apostrophes used to enclose a linkname (the delimiters) do not count as part of the 1023-character limit.
- Any apostrophes specified as part of a linkname (not the delimiters) must be doubled.

  Double apostrophes count as two characters in the 1023-character limit.
- The linkname can include characters X'40' through X'FE'.

PARM
specifies a character string that is to be passed to the copy utility as an execution-time parameter. The maximum length of this character string is 300 bytes of nonblank data. If any blanks are specified in the PARM value, they are deleted by SMP/E during processing and do not count toward the 300-byte maximum.

The UCL operand is PARM(*character_string*).

- PARM is an optional operand.
- The character string can be entered free-form, without regard to blanks (which are compressed out of the string), and can span multiple 80-byte records.
- If parentheses are specified in the PARM value, there must always be a pair (left and right); otherwise, the results are unpredictable.

RMID
identifies the last SYSMOD that **replaced** this JAR element. Any subsequent SYSMOD that modifies this JAR element must have defined a relationship (such as PRE or SUP) with this SYSMOD.

The UCL operand is RMID(*sysmod_id*).

- The SYSMOD ID must contain 7 uppercase alphabetic, numeric, or national ($, #, @) characters.
- If RMID is not specified but FMID is, SMP/E sets the RMID value to the specified FMID.

SHSCRIPT
specifies a UNIX shell script to run when this JAR element is copied to (or deleted from) a directory in the UNIX file system.

The UCL operand is SHSCRIPT(*scriptname*). This subentry optionally contains one or both of the following values, which specify the point in SMP/E processing when the shell script is run:

PRE
The shell script is run before the JAR element is copied to a UNIX file system directory.

POST
The shell script is run after the JAR element is copied to a UNIX file system directory.

POST is the default; the shell script is run after the JAR element is copied to the directory.

SYMLINK
specifies a list of one or more symbolic links, which are file names that can be used as alternate names for referring to this JAR element in the UNIX file system. Each symbolic link name listed here is associated with a path name listed in the SYMPATH entry. See the description of the SYMPATH entry for more information about how the symbolic link names and path names are associated.

The UCL operand is SYMLINK(*symlinkname...*).

- A JAR entry that contains a SYMLINK entry must contain a matching SYMPATH entry. SMP/E will reject any UCLIN command that would violate this condition.
- A symbolic linkname can be from one to 1023 characters. Any characters in the range X'40' through X'FE' may be specified.
- A symbolic linkname can be enclosed in single apostrophes ('). A symbolic linkname **must** be enclosed in single apostrophes if any of the following is true:
  - The symbolic linkname contains lowercase alphabetic characters.
  - The symbolic linkname contains a character that is not uppercase alphabetic, numeric, national ($, #, or @), slash (/), plus (+), hyphen, period, or ampersand (&).
  - The symbolic linkname spans more than one line in the control statement.

  The single apostrophes used to enclose a symbolic linkname (the delimiters) do not count as part of the 1023-character limit.
- Any apostrophes specified as part of a symbolic linkname (not the delimiters) must be doubled. Double apostrophes count as two characters in the 1023-character limit.

SYMPATH
specifies a list of one or more pathnames that are associated with symbolic links identified by the SYMLINK operand. The first pathname in the SYMPATH operand is associated with the first symbolic link in the SYMLINK operand, the second pathname with the second symbolic link, and so on. If there are more symbolic links listed than there are pathnames, then the last listed pathname is used for the remaining symbolic links. If more pathnames are specified than symbolic linknames, then the excess pathnames (at the end of the list) are ignored. The UCL operand is SYMPATH(*sympathname...*).

- A JAR entry that contains a SYMPATH entry must contain a matching SYMLINK entry. is specified, otherwise it must be omitted.
- A symbolic pathname can be from one to 1023 characters. Any characters in the range X'40' through X'FE' may be specified.
- A symbolic pathname can be enclosed in single apostrophes ('). A symbolic pathname **must** be enclosed in single apostrophes if any of the following is true:
  - The symbolic pathname contains lowercase alphabetic characters.
  - The symbolic pathname contains a character that is not uppercase alphabetic, numeric, national ($, #, or @), slash (/), plus (+), hyphen, period, or ampersand (&).
  - The symbolic pathname spans more than one line in the control statement.

  The single apostrophes used to enclose a symbolic pathname (the delimiters) do not count as part of the 1023-character limit.
- Any apostrophes specified as part of a symbolic pathname (not the delimiters) must be doubled. Double apostrophes count as two characters in the 1023-character limit.

SYSLIB
specifies the ddname of the "target library" within a UNIX file system for the JAR element.

The UCL operand is SYSLIB(*ddname*).

- Only one SYSLIB value can be specified.
- The ddname can contain any uppercase alphabetic, numeric, or national ($, #, @) character, and can be 1 to 8 characters long.
- The SYSLIB subentry is required. Without it, SMP/E cannot process any changes for the JAR element.

UMID
identifies all of the SYSMODs that have updated this JAR file since it was last replaced. Any subsequent SYSMOD that updates or replaces this JAR file must have a defined relationship (such as PRE or SUP) with all SYSMODs in the UMID subentry.

LIST Examples

To list all the JAR entries in a particular zone, you can use the following commands:

```
SET      BDY(TGT1)     /* Set to zone.         */.
LIST     JAR           /* List all JAR entries.  */.
```

To list specific JAR entries, you can use these commands:

```
SET      BDY(TGT1)     /* Set to zone.         */.
LIST     JAR(JAREL1,   /* List only these two  */
             JAREL2)   /* entries.             */.
```

The format of the LIST output for each JAR entry is the same for both of these commands. The only difference is the number of JAR entries listed.

is an example of LIST output for a JAR entry.

```
PAGE nnnn  - NOW SET TO zzzzzz ZONE nnnnnnn  DATE mm/dd/yy TIME hh:mm:ss SMP/E 36.nn SMPLIST OUTPUT

SMPPTS       MCS ENTRIES


   NAME

ABCTTT         LASTUPD        = HABC100  TYPE=ADD
               LIBRARIES      = DISTLIB=AABCBIN    SYSLIB=SABCBIN
               FMID           = HABC100
               RMID           = HABC100
               UMID           = OW12345   OW54321    OW34567
               SHSCRIPT       = ABCSCRPT,PRE,POST
               PARM           = PATHMODE(0,6,4,4)
               JARPARM        = 0M
               SYMLINK        = '../../../../../usr/lib/TicTacToe.jar'
               SYMPATH        = '../../usr/lpp/ttt/bin/TicTacToe.jar'
               LINK           = '../TicTacToe.jar'
```

Figure 32. JAR entry: sample LIST output

By specifying the FORFMID operand, you can reduce the number of JAR entries listed. When FORFMID is specified, SMP/E lists a JAR entry only if its FMID matches one of the FMIDs specified on the FORFMID operand. For example, to list JAR entries whose FMIDs either are defined in FMIDSET JAR or are JARFUNC, you can use these commands:

```
SET      BDY(TGT1)          /* Set to target zone.     */.
LIST     JAR                /* List all JAR entries    */
         FORFMID(JAR        /* for the JAR FMIDSET     */
              JARFUNC)      /* and FMID JARFUNC.       */.
```

You can use the LIST command to find out the names of all SYSMODs that have modified a JAR element. To include the names of these SYSMODs in the LIST output you can use the XREF operand, as shown in these commands:

```
SET      BDY(TGT1)          /* Set to requested zone.  */.
LIST     JAR                /* List all JAR entries    */
         XREF               /* and related SYSMODs.    */.
```

Note:

1. XREF can be used either in mass mode or in select mode.

2. SMP/E obtains the data included for the XREF operand by checking all the SYSMOD entries for subentries for this JAR element. Because this data is not contained in the JAR entry itself, you cannot use UCLIN to change it in the JAR entry.

is an example of the LIST output produced when the XREF operand is used.

```
PAGE nnnn  - NOW SET TO zzzzzz ZONE nnnnnnn  DATE mm/dd/yy TIME hh:mm:ss SMP/E 36.nn SMPLIST OUTPUT

TGT1    JAR      ENTRIES

   NAME

 ABCTTT      LASTUPD        = JAR0001  TYPE=ADD
             LIBRARIES      = DISTLIB=DLIB3      SYSLIB=SLSRBIN
             FMID           = JAR0001
             RMID           = JAR0001
             PARM           = PATHMODE(0,7,7,5)
             JARPARM        = 0M
             LINK           = '../TicTacToe.jar'
             UMID           = PT00002
             SYSMOD HISTORY = SYSMOD    TYPE        DATE    MCS        ---------- STAT
                             JAR0001  FUNCTION  07.100  JAR        APP
                             PT00002  PTF       07.100  JARUPD     APP
```

Figure 33. JAR entry: sample LIST output when XREF is specified

UCLIN Examples

You can use the ADD, REP, and DEL UCL statements to change subentries in the JAR entry. After the UCLIN changes are made, the JAR entry must contain at least the following subentries:

- DISTLIB
- FMID
- RMID
- SYSLIB

Otherwise, there is not enough information in the entry to process the element. If any of these subentries are missing, SMP/E does not make the requested UCL updates to the entry, and the entry remains as it was before the UCL command.

The following examples are provided to help you use the JAR entry:

```
Example 1

  SET BDY(TGT)                     /* Set to target zone       */.
  UCLIN                            /* Start UCLIN processing    */.
    ADD JAR(ABCTTT) JARPARM(0M)    /* Add JARPARMs              */.
  ENDUCL                           /* End UCLIN processing      */.

Example 2

  SET BDY(TGT)                     /* Set to target zone       */.
  UCLIN                            /* Start UCLIN processing    */.
    REP JAR(ABCTTT) JARPARM(M)     /* Replace with new JARPARMs */.
                   UMID(OW12345,OW54321) /* and new UMIDs       */.
  ENDUCL                           /* End UCLIN processing      */.

Example 3

  SET BDY(TGT)                     /* Set to target zone       */.
  UCLIN                            /* Start UCLIN processing    */.
    DEL JAR(ABCTTT) JARPARM()      /* Delete JARPARMs           */.
  ENDUCL                           /* End UCLIN processing      */.
```

In the first example, a string of options is added to the JARPARM subentry in the JAR entry ABCTTT. In the second example, the existing JARPARM subentry value for ABCTTT is replaced with a new value and the UMID subentry list is replaced with a new list of values. In the third example, the JARPARM subentry is deleted from the JAR entry for entry name ABCTTT.

UNLOAD Examples

To dump all the JAR entries in UCL format, you can use the UNLOAD command:

```
SET     BDY(TGT1)    /* Set to zone.            */.
UNLOAD  JAR          /* Unload all JAR entries.  */.
```

To dump specific JAR entries, you can use these commands:

```
SET     BDY(TGT1)    /* Set to zone.            */.
UNLOAD  JAR(JAREL1,  /* Unload only these two    */
           JAREL2)   /* entries.                 */.
```

The format of the UNLOAD output for each JAR entry is the same for both of these commands. The only difference is the number of JAR entries unloaded.

is an example of the output created by the UNLOAD command for a JAR entry:

```
    UCLIN .
    REP        JAR            ( ABCTTT   )
               LASTUPD        ( HABC100 )
               LASTUPDTYPE    ( ADD )
               DISTLIB        ( AABCBIN  )
               SYSLIB         ( SABCBIN  )
               FMID           ( HABC100 )
               RMID           ( HABC100 )
               UMID           ( OW12345   OW54321   OW34567 )
               SHSCRIPT       ( ABCSCRPT,PRE,POST )
               PARM           ( PATHMODE(0,6,4,4) )
               JARPARM        ( 0M )
               SYMLINK        ( '../../../../../usr/lib/TicTacToe.jar' )
               SYMPATH        ( '../../usr/lpp/ttt/bin/TicTacToe.jar' )
               LINK           ( '../TicTacToe.jar' )
                                .
    ENDUCL.
```

Figure 34. Example UNLOAD output for JAR entry

By specifying the FORFMID operand, you can reduce the number of JAR element entries unloaded. When FORFMID is specified, SMP/E unloads an JAR element entry only if its FMID matches one of the FMIDs specified on the FORFMID operand. For example, to unload JAR element entries whose FMIDs either are defined in FMIDSET JAR or are JARFUNC, you can use these commands:

```
SET      BDY(TGT1)           /* Set to target zone.           */.
UNLOAD   JAR                 /* Unload all JAR element entries */
         FORFMID(JAR         /* for the JAR FMIDSET           */
                 JARFUNC)    /* and FMID JARFUNC.             */.
```

LMOD entry (distribution and target zone)

The LMOD entry contains all the information needed to replace or update a given load module. This includes information such as whether the load module is link-edited or copied during the system generation process, any link-edit statements required to relink the load module, the link-edit attributes of the load module, and the libraries in which it resides. An LMOD entry is generally created by one of the following methods:

- **Installing a SYSMOD that adds the load module.** LMOD entries can be created when a SYSMOD is installed. SMP/E builds an LMOD entry if it encounters a ++MOD statement for a load module that does not yet have an LMOD entry and if the distribution library for the module was totally copied during system generation. For more information about copied load modules, see "DLIB entry (distribution and target zone)" on page 174 and *z/OS SMP/E Commands*.

- **Processing JCLIN.** LMOD entries can be created during JCLIN processing when SMP/E scans the copy and link-edit steps. At the same time, SMP/E builds MOD entries for modules that are linked or copied to the load module. For more information, see "MOD entry (distribution and target zone)" on page 226 and the "Processing" section of the JCLIN command chapter in *z/OS SMP/E Commands*.

Subentries

These are the subentries for the LMOD entry as they appear in the LIST output:

name
    is the name of the load module described by the LMOD entry.

    The name can contain from 1 to 8 alphanumeric characters.

CALLLIBS
    specifies one or more DDDEF entries, existing in the same zone, that compose the SYSLIB allocation to be used when the load module is link-edited.

    The UCL operand is CALLLIBS(*name...*).

- The DDDEF names can contain from 1 to 8 alphabetic (A through Z), national (@, #, or $), or numeric (0 through 9) characters. The first character must be alphabetic or national.
- There must be a DDDEF entry for each specified name. SMP/E does not check other sources (such as DD statements or GIMDDALC control statements in SMPPARM member GIMDDALC) to get the information needed to allocate the data sets.
- SMP/E does not enforce any limit on the number of names that can be specified in a CALLLIBS subentry list. However, DFSMS/MVS sets its own limits on how many data sets can be concatenated. The actual limit depends on the kind of data sets to be concatenated (partitioned data sets or PDSEs) and the total number of extents for partitioned data sets. Refer to *z/OS DFSMS Using Data Sets* for information about calculating this limitation.
- The order in which the libraries are specified is important because it indicates the order in which the SYSLIB concatenation is built.

COPY
is a special SMP/E indicator meaning that the load module was copied during system generation, and that there is a one-to-one correspondence between the distribution library module (the MOD entry) and the target system load module (the LMOD entry). This information is used during APPLY processing to determine whether the LEPARM values from a ++MOD statement are applicable to the load module. It is also used during delete and compress processing to determine whether a load module can be deleted before the new modules are installed.

The UCL operand is COPY.

LASTUPD
identifies the cause of the last change to this LMOD entry.

Note: If a given UCLIN command specifies only cross-zone subentries, this field is not changed.

The UCL operand is LASTUPD(*value*). This subentry can contain one of the following values:

JCLIN
indicates that the change was made during JCLIN command processing.

UCLIN
indicates that the change was made as a result of UCLIN processing.

sysmod_id
indicates that the change was made during the installation of the indicated SYSMOD. The SYSMOD did one of the following:

- Contained inline JCLIN that affected the LMOD entry.
- Renamed the load module using the ++RENAME statement.
- Moved the load module to a new target system library using the ++MOVE statement.

A SYSMOD that does not do any of the previously listed actions will not cause an update to the LASTUPD subentry. For example, a SYSMOD that supplies one or more modules (++MOD) that cause the load module to be link edited does not necessarily cause the LMOD entry's LASTUPD subentry to be updated. The LASTUPD subentry only indicates when the LMOD entry was updated, not when the module content of a load module is updated.

The SYSMOD ID must contain 7 alphanumeric characters.

LASTUPD TYPE
indicates how the entry was last changed.

Note: If a given UCLIN command specifies only cross-zone subentries, this field is not changed.

The UCL operand is LASTUPDTYPE(*value*). This subentry can contain one of the following values:

ADD
The entry was added.

DEL
A subentry in the entry was deleted.

MOV
> The load module was moved.

REN
> The load module was renamed.

UPD
> The entry was updated.

LKED ATTRIBUTES
> identifies the link-edit attributes that must be used when this load module is link-edited. SMP/E supports the following link-edit attributes. For more information, see *z/OS MVS Program Management: User's Guide and Reference*.

> **AC=1**
>> specifies that the AC=1 parameter (which is the authorization code) is to be passed to the link-edit utility when the load module is link-edited.
>>
>> The UCL operand is AC=1.

> **ALIASES**
>> specifies that the ALIASES parameter is to be passed to the link-edit utility when the load module is link-edited.
>>
>> The UCL operand is ALIASES(ALL).

> **ALIGN2**
>> specifies that the ALIGN2 parameter (alignment on a 2KB boundary) is to be passed to the link-edit utility when the load module is link-edited.
>>
>> The UCL operand is ALIGN2 or ALN2.

> **AMODE=24**
>> specifies that the AMODE=24 parameter is to be passed to the link-edit utility when the load module is link-edited.
>>
>> The UCL operand is AMODE=24 or AMOD=24.

> **AMODE=31**
>> specifies that the AMODE=31 parameter is to be passed to the link-edit utility when the load module is link-edited.
>>
>> The UCL operand is AMODE=31 or AMOD=31.

> **AMODE=64**
>> specifies that the AMODE=64 parameter is to be passed to the link-edit utility when the load module is link-edited.
>>
>> The UCL operand is AMODE=64 or AMOD=64.

> **AMODE=ANY**
>> specifies that the AMODE=ANY parameter is to be passed to the link-edit utility when the load module is link-edited.
>>
>> The UCL operand is AMODE=ANY or AMOD=ANY.

> **AMODE=MIN**
>> specifies that the AMODE=MIN parameter is to be passed to the link-edit utility when the load module is link-edited.
>>
>> The UCL operand is AMODE=MIN or AMOD=MIN.

> **CALL**
>> indicates that the CALL parameter has been specified for a load module in a JCLIN link-edit step.
>>
>> **Note:** If the load module being link-edited contains a CALLLIBS subentry list, SMP/E uses NCAL for the link to the SMPLTS library, and ignores the CALL subentry in the LMOD entry.

For more information about link-edit parameters, see the discussion of "Link-edit utility" under the PARM subentry of "UTILITY entry (global zone)" on page 289. Also see the APPLY command chapter in *z/OS SMP/E Commands*.

The UCL operand is CALL.

CASE
specifies that the CASE parameter, which controls case sensitivity and folding, is to be passed to the link-edit utility when the load module is link-edited.

The UCL operand is CASE(UPPER|MIXED).

- CASE(UPPER) and CASE(MIXED) are mutually exclusive.

COMPAT
specifies that the COMPAT parameter is to be passed to the link-edit utility when the load module is link-edited.

The UCL operand is COMPAT=LKED|PM1|PM2|PM3|PM4.

DC
specifies that the DC parameter, which is the downward compatible attribute, is to be passed to the link-edit utility when the load module is link-edited.

The UCL operand is DC.

DYNAM
specifies that the DYNAM parameter is to be passed to the link-edit utility when the load module is link-edited.

The UCL operand is DYNAM(DLL).

FETCHOPT
specifies that the FETCHOPT parameter is to be passed to the link-edit utility when the load module is link-edited.

The UCL operand is FETCHOPT(PACK|NOPACK, PRIME|NOPRIME).

FILL
specifies that the FILL parameter is to be passed to the link-edit utility when the load module is link-edited.

The UCL operand is FILL(nn), where nn is the 2-character representation of a single hex byte (00 - FF).

HOBSET
specifies that the HOBSET parameter is to be passed to the link-edit utility when the load module is link-edited.

The UCL operand is HOBSET.

LONGPARM
specifies that the LONGPARM parameter is to be passed to the link-edit utility when the load module is link-edited.

The UCL operand is LONGPARM.

MAXBLK
specifies that the MAXBLK parameter is to be passed to the link-edit utility when the load module is link-edited.

The UCL operand is MAXBLK($nnnnn$), where $nnnnn$ is a number between 256 and 32760.

NE
specifies that the NE parameter, which is the noneditable attribute, is to be passed to the link-edit utility when the load module is link-edited.

The UCL operand is NE.

NOCALL
> specifies that the NCAL parameter is to be passed to the link-edit utility when the load module is link-edited.

> **Note:** If the load module being link-edited contains a CALLLIBS subentry list, SMP/E uses CALL for the link to the actual target library, and ignores the NOCALL subentry in the LMOD entry.

> For more information about link-edit parameters, see the discussion of "Link-edit utility" under the PARM subentry of "UTILITY entry (global zone)" on page 289. Also see the APPLY command chapter in *z/OS SMP/E Commands*.

> The UCL operand is NOCALL or NCAL.

OL
> specifies that the OL parameter is to be passed to the link-edit utility when the load module is link-edited.

> The UCL operand is OL.

OVLY
> specifies that the OVLY parameter, which specifies that the load module is in overlay structure, is to be passed to the link-edit utility when the load module is link-edited.

> The UCL operand is OVLY.

REFR
> specifies that the REFR parameter, which is the refreshable attribute, is to be passed to the link-edit utility when the load module is link-edited.

> The UCL operand is REFR.

RENT
> specifies that the RENT parameter, which indicates that the load module is re-entrant, is to be passed to the link-edit utility when the load module is link-edited.

> The UCL operand is RENT.

REUS
> specifies that the REUS parameter, which indicates that the load module is reusable, is to be passed to the link-edit utility when the load module is link-edited.

> The UCL operand is REUS.

REUS(NONE)
> specifies that the REUS(NONE) parameter, which indicates that the load module cannot be reused, is to be passed to the link-edit utility when the load module is link-edited.

> The UCL operand is REUS(NONE).

RMODE=24
> specifies that the RMODE=24 parameter is to be passed to the link-edit utility when the load module is link-edited.

> The UCL operand is RMODE=24 or RMOD=24.

RMODE=31
> specifies that the RMODE=ANY parameter is to be passed to the link-edit utility when the load module is link-edited. (RMODE=31 is a synonym for RMODE=ANY.)

> The UCL operand is RMODE=31 or RMOD=31.

RMODE=64
> Specifies that the RMODE=64 parameter is to be passed to the link-edit utility when the load module is link-edited.

> The UCL operand is RMODE=64 or RMOD=64.

RMODE=ANY
specifies that the RMODE=ANY parameter is to be passed to the link-edit utility when the load module is link-edited.

The UCL operand is RMODE=ANY or RMOD=ANY.

RMODE=SPLIT
specifies that the RMODE=SPLIT parameter is to be passed to the link-edit utility when the load module is link-edited.

The UCL operand is RMODE=SPLIT or RMOD=SPLIT.

RMODEX=64TRUE
Specifies that the RMODEX=64TRUE parameter is to be passed to the link-edit utility when the load module is link-edited.

The UCL operand is RMODEX=64TRUE.

RMODEX=NO
Specifies that the RMODEX=NO parameter is to be passed to the link-edit utility when the load module is link-edited.

The UCL operand is RMODEX=NO.

SCTR
specifies that the SCTR parameter, which indicates that the load module can be scatter-loaded, is to be passed to the link-edit utility when the load module is link-edited.

The UCL operand is SCTR.

STD
is a special SMP/E indication that the load module should be link-edited with none of the previously listed attributes.

STD appears in LIST and UNLOAD output in these cases:

- STD was specified on a UCL statement for this entry.
- No link-edit attributes were specified for this entry. SMP/E uses the standard system link-edit parameters.
- The LMOD entry was created by a JCLIN copy step.

If this attribute is specified when a load module is to be link-edited and no UTILITY entry is being used, SMP/E passes the default link-edit parameters to the link-edit utility. For more information about link-edit parameters, see the APPLY command chapter in *z/OS SMP/E Commands*.

The UCL operand is STD.

UPCASE
specifies that the UPCASE parameter, which indicates how the Binder should process symbol names, is to be passed to the link-edit utility when the load module is link-edited.

The UCL operand is UPCASE(YES|NO).

- UPCASE(YES) and UPCASE(NO) are mutually exclusive.

LKED CONTROL
contains all the link-edit control cards necessary to relink-edit this load module.

Note: Relink-editing a load module is triggered by a replacement for one of the modules within that load module. SMP/E performs only CSECT replaces when relink-editing the load module; it does not completely rebuild the load module from the distribution libraries. Therefore, it is not necessary to save all the link-edit control statements in the LMOD entry; for example, the INCLUDE statements are not saved.

The UCL operands are ++LMODIN and ++ENDLMODIN.

- The ++LMODIN and ++ENDLMODIN statements must start in column 1.

- No other operands can start on the same line as the ++LMODIN statement.
- The ++ENDLMODIN statement must be specified if the ++LMODIN statement is specified.
- The link-edit control statements must follow SMP/E's coding conventions. For more information, see the "General JCLIN Coding Conventions" section of the JCLIN command chapter in *z/OS SMP/E Commands*.
- Comment statements are ignored and are not saved in the LMOD entry. (Any link-edit control statement with an asterisk in column 1 is considered a comment statement.)

MODDEL

specifies the modules that were once part of this load module but have been deleted. If the modules are reinstalled, they are linked back into this load module.

The UCL operand is MODDEL(*module...*).

The module names can contain from 1 to 8 alphanumeric characters.

RETURN CODE

specifies the maximum acceptable return code from the link edit utility for this specific load module. If the return code from the link edit utility is higher than this value, SMP/E normally assumes the link edit operation failed.

The UCL operand is RC(*rc*), where the *rc* value can be any decimal number from 0 to 16.

SIDE DECK LIBRARY

specifies the library to contain the definition side deck.

The definition side deck may be created as a member of a partitioned data set (PDS or PDSE), or as a file within a directory in a UNIX file system. The definition side deck is created by the link edit utility when the DYNAM(DLL) binder option is specified for DLL load modules that export symbols. If a load module does not export symbols, or if the DYNAM(DLL) option is not specified, the definition side deck is not created. The member or file name of the definition side deck is the same as the load module name. In addition, the definition side deck will have the same aliases as those specified for the load module.

The UCL operand is SIDEDECKLIB(*ddname*), where *ddname* is the ddname of the library where the load module's definition side deck will reside.

Note:

1. This subentry name appears in the LIST output as SIDE DECK LIB.
2. When this subentry value is SMPDUMMY, the SIDE DECK LIBRARY will be allocated as a DUMMY data set, and the definition side deck will not be created by the link edit utility.

SYSTEM LIBRARIES

specifies the target system libraries in which this load module resides.

The UCL operand is SYSLIB(*ddname...*).

- There can be up to two SYSLIB values in an LMOD entry.
- The ddnames can contain from 1 to 8 alphanumeric characters.

Note: This subentry name appears in the LIST output as SYSTEM LIBRARY.

UTILITY INPUT

specifies one or more files to be included when this load module is link-edited. The ddname of the library where the file resides is also indicated.

The UCL operand is UTIN((*name,ddname*)...), where *name* is a data set member name or relative UNIX filename, and *ddname* is the ddname of the library where the file resides.

- The *name* can be from 1 to 1023 characters in length.
- The *name* can include any nonblank characters X'41' through X'FE'.

- The *name* cannot begin or end with a slash ('/') because any UNIX filename specification must be relative to the directory associated with the *ddname*.
- A *name* can be enclosed in single quotation marks ('). A *name* value must be enclosed in single quotation marks if any of the following are true:
  - The *name* contains a character that is not uppercase alphabetic, numeric, national ($,#, or @), slash (/), plus (+), hyphen (-), period (.) or ampersand (&).
  - The *name* spans more than one line in the UCLIN statement.

  The single quotation marks used to enclose a *name* (the delimiters) do not count as part of the 1023-character limit.
- Any quotation marks specified as part of a *name* (not the delimiters) must be doubled. Double quotation marks count as two characters in the 1023 character limit.
- If the *name* value spans lines, it must extend through column 72 of the first line and begin in column 1 of the subsequent line.

The combination of *name* and *ddname* determines the uniqueness of a subentry in the UTILITY INPUT subentry list. Only one subentry value for a given *name* and *ddname* combination is saved in the UTILITY INPUT subentry list.

XZMOD
specifies one or more modules that were added to the load module by the LINK MODULE command. The name of the zone supplying each module is also indicated.

The UCL operand is XZMOD((*module,zone*)...).

- The zone name specified for an XZMOD subentry cannot match the name of the set-to zone.
- The XZMOD subentry is added to an LMOD entry automatically during LINK MODULE command processing. However, it is almost never removed automatically, except during JCLIN processing. For more information, see the "Cross-Zone Relationships" section of the JCLIN command chapter in *z/OS SMP/E Commands*.
- An LMOD entry can contain XZMOD and XZMODP subentries without any other subentries. Such an LMOD entry represents a *stub* load module, which is retained after a load module with cross-zone relationships is deleted. If the deleted load module is later reinstated, SMP/E uses the information in the stub load module to issue messages reminding the user to use the LINK MODULE command to restore those previous cross-zone relationships, if they are still appropriate.
- If UCLIN is used to update an existing LMOD entry and **only** cross-zone subentries (XZMOD and XZMODP) are changed, SMP/E does not update the LASTUPD and LASTUPDTYPE subentries.

XZMODP
indicates that the load module contains one or more modules from another zone and that XZMOD subentries exist in this LMOD entry.

The UCL operand is XZMODP.

- You never need to specify the XZMODP subentry on a UCL statement. SMP/E automatically determines the setting of XZMODP, according to whether the LMOD entry contains XZMOD subentries.
- You cannot add the XZMODP subentry to an LMOD entry that does not contain XZMOD subentries.
- You cannot delete the XZMODP subentry from an LMOD entry containing XZMOD subentries.
- An LMOD entry can contain XZMOD and XZMODP subentries without any other subentries.
- If UCLIN is used to update an existing LMOD entry and **only** cross-zone subentries are changed (XZMOD and XZMODP), SMP/E does not update the LASTUPD and LASTUPDTYPE subentries.

LIST Examples

To list all the LMOD entries in a particular zone, you can use the following commands:

```
SET       BDY(TGT1)            /* Set to requested zone.   */.
LIST      LMOD                 /* List all LMOD entries.   */.
```

To list specific LMOD entries, you can use these commands:

```
SET       BDY(TGT1)            /* Set to requested zone.   */.
LIST      LMOD(LMOD01          /* List only these two      */
              LMOD02)          /* entries.                 */.
```

The format of the LIST output for each LMOD entry is the same for both of these commands. The only difference is the number of LMOD entries listed. The code sample below is an example of LIST output for LMOD entries.

```
PAGE nnnn  - NOW SET TO zzzzzz ZONE nnnnnnn  DATE mm/dd/yy TIME hh:mm:ss SMP/E 36.nn SMPLIST OUTPUT

TGT1     LMOD      ENTRIES


  NAME

BPXLMOD1  LASTUPD        = HBPX001 TYPE=ADD
          SYSTEM LIBRARY  = BPXLIB1
          LKED ATTRIBUTES = RENT,CASE(MIXED)
          LKED CONTROL    =  ENTRY BPXMOD01
                             ALIAS ../FRIENDLY/NAME/THAT/DOES/NOT/NEED/QUOTES
                             ALIAS '../friendly/name/with a comma,/which/requires/quotes'
                             ALIAS '../friendly/name/with/a/''/which/requires/quotes'

LMOD01    LASTUPD        = JXY1102 TYPE=ADD
          SYSTEM LIBRARY  = LINKLIB
          RETURN CODE     = 4
          LKED ATTRIBUTES = STD
          MODDEL          = MOD05       MOD06
          LKED CONTROL    =  ENTRY MOD01
                             ALIAS LMOD01A1

LMOD02    LASTUPD        = JXY1102 TYPE=ADD
          SYSTEM LIBRARY  = LINKLIB    PPLIB01
          LKED ATTRIBUTES = RENT,REUS,AC=1

LMOD03    LASTUPD        = JXY1102 TYPE=ADD
          SYSTEM LIBRARY  = LINKLIB
          LKED ATTRIBUTES = RENT,REUS,AC=1
          LKED CONTROL    =  ALIAS MOD03
                             *** CHANGE/REPLACE STMTS FOR MOD03 FROM DLIB AOS12
                             CHANGE MOD03C1(NEW03C1)
                             CHANGE MOD03C2(NEW03C2)

LMOD04    LASTUPD        = LINK     TYPE=UPD
          SYSTEM LIBRARY  = LINKLIB
          RETURN CODE     = 4
          LKED ATTRIBUTES = STD
          XZMODP
          XZMOD           =  CICSMOD1 FROM ZONE CICS1
                          =  CICSMOD2 FROM ZONE CICS1
                          =  CICSMOD4 FROM ZONE CICS1
                          =  IMSMOD1  FROM ZONE IMS1
          LKED CONTROL    =  ENTRY MOD04
                             ALIAS LMOD04A1

LMOD05    XZMODP
          XZMOD           =  CICSMOD1 FROM ZONE CICS1
                          =  CICSMOD2 FROM ZONE CICS1
                          =  IMSMOD1  FROM ZONE IMS1

LMOD06    LASTUPD        = JCLIN    TYPE=ADD
          SYSTEM LIBRARY  = LINKLIB
          LKED ATTRIBUTES = SCTR,MAXBLK(6160),AC=1,AMOD=MIN


LMOD07    LASTUPD        = JCLIN    TYPE=ADD
          SYSTEM LIBRARY  = LINKLIB
          LKED ATTRIBUTES = FETCHOPT(PACK,PRIME),NOCALL
          MODDEL          = ISPLINK    MOD01

LMOD08    LASTUPD        = HXY0001 TYPE=ADD
          SYSTEM LIBRARY  = APPLLIB
          LKED ATTRIBUTES = RENT,REUS
```

LMOD entry (distribution and target zone)

```
           CALLLIBS        = PLIBASE    CSSLIB

LOTSOPRM  LASTUPD         = UCLIN    TYPE=UPD
          SYSTEM LIBRARY  = LINKLIB
          LKED ATTRIBUTES =
RENT,REUS,SCTR,OVLY,XCAL,REFR,DC,NE,AC=1,ALIGN2,AMODE=31,CASE(MIXED),OL,MAXBLK(32760),

FETCHOPT(NOPACK,NOPRIME),COMPAT=LKED,NCAL,FILL(00),HOBSET,RMODE=SPLIT,ALIASES(ALL),
                          DYNAM(DLL),UPCASE(YES)
          LKED CONTROL    =  ENTRY DL1MOD2

MODULE1   LASTUPD         = UCLIN    TYPE=ADD
          SYSTEM LIBRARY  = LINKLIB
          LKED ATTRIBUTES = STD
          MODDEL          = MOD1       MOD2

MODUL25   LASTUPD         = UCLIN    TYPE=UPD
          SYSTEM LIBRARY  = LINKLIB
          LKED ATTRIBUTES = STD
```

You can use the LIST command to find out the names of all the distribution library modules that are link-edited into this load module. To include the names of these modules in the LIST output you can use the XREF operand, as shown in these commands:

```
SET     BDY(TGT1)          /* Set to requested zone.   */.
LIST    LMOD               /* List all LMOD entries    */
        XREF               /* and MODs in them.        */.
```

Note:

1. XREF can be used either in mass mode or in select mode.

2. SMP/E obtains the data included for the XREF operand by checking the LMOD subentries in all the MOD entries. Because this data is not contained in the LMOD entry itself, you cannot use UCLIN to change it in the LMOD entry.

The code sample below is an example of the LIST output produced when the XREF operand is used.

```
PAGE nnnn  - NOW SET TO zzzzzz ZONE nnnnnnn  DATE mm/dd/yy TIME hh:mm:ss SMP/E 36.nn SMPLIST OUTPUT

 TGT1    LMOD      ENTRIES


    NAME

BPXLMOD1  LASTUPD         = HBPX001 TYPE=ADD
          SYSTEM LIBRARY  = BPXLIB1
          LKED ATTRIBUTES = RENT,CASE(MIXED)
          LKED CONTROL    =  ENTRY BPXMOD01
                             ALIAS ../FRIENDLY/NAME/THAT/DOES/NOT/NEED/QUOTES
                             ALIAS '../friendly/name/with a comma,/which/requires/quotes'
                             ALIAS '../friendly/name/with/a/''/which/requires/quotes'
          MODULES         = NAME       FMID
                            BPXMOD01   HBPX001

LMOD01    LASTUPD         = JCLIN     TYPE=ADD
          SYSTEM LIBRARY  = SLIB01
          LKED ATTRIBUTES = COPY,RENT,REUS,SCTR,OVLY,EFR,DC,NE,AC=1,ALIGN2,AMODE=24,RMODE=24
          MODULES         = NAME       FMID
                            MODULE2    HXZ1234

LMOD02    LASTUPD         = JCLIN     TYPE=ADD
          SYSTEM LIBRARY  = SLIB02
          LKED ATTRIBUTES = COPY,RENT,REUS,SCTR,OVLY,EFR,DC,NE,AC=1,ALIGN2,AMODE=24,RMODE=24
          MODULES         = NAME       FMID
                            MODULE2    HXZ1234

LMOD03    LASTUPD         = JXY1102   TYPE=ADD
          SYSTEM LIBRARY  = LINKLIB
          LKED ATTRIBUTES = RENT,REUS,AC=1
          LKED CONTROL    =  ALIAS MOD3
                             *** CHANGE/REPLACE STMTS FOR MOD03 FROM DLIB AOS12
                               CHANGE MOD03C1(NEW03C1)
                               CHANGE MOD03C2(NEW03C2)
          MODULES         = NAME       FMID
                            MOD03      JXY1121

LMOD04    LASTUPD         = LINK      TYPE=UPD
```

```
                SYSTEM LIBRARY  = LINKLIB
                RETURN CODE     = 4
                LKED ATTRIBUTES = STD
                XZMODP
                XZMOD           =  CICSMOD1 FROM ZONE CICS1
                                =  CICSMOD2 FROM ZONE CICS1
                                =  CICSMOD4 FROM ZONE CICS1
                                =  IMSMOD1  FROM ZONE IMS1
                LKED CONTROL    =  ENTRY MOD04
                                =  ALIAS LMOD04A1
                MODULES         = NAME      FMID
                                = MODULE4   HXZ1234

LMOD05    XZMODP
                XZMOD           =  CICSMOD1 FROM ZONE CICS1
                                =  CICSMOD2 FROM ZONE CICS1
                                =  IMSMOD1  FROM ZONE IMS1


LMOD06    LASTUPD         = JCLIN    TYPE=ADD
                SYSTEM LIBRARY  = LINKLIB
                LKED ATTRIBUTES = SCTR,MAXBLK(6160),AC=1,AMOD=MIN
                MODULES         = NAME      FMID
                                  MODULE5   HXZ1234

LMOD07    LASTUPD         = JCLIN    TYPE=ADD
                SYSTEM LIBRARY  = LINKLIB
                LKED ATTRIBUTES = FETCHOPT(PACK,PRIME),NOCALL
                MODDEL          = ISPLINK   MOD01
                MODULES         = NAME      FMID
                                  MODULE6   JXY0001

LMOD08    LASTUPD         = HXY0001 TYPE=ADD
                SYSTEM LIBRARY  = APPLLIB
                LKED ATTRIBUTES = RENT,REUS
                CALLLIBS        = PLIBASE   CSSLIB
                MODULES         = NAME      FMID
                                  MODULE7   HXY0001
                                  MODULE8   HXY0001

 MODULE1   LASTUPD         = UCLIN    TYPE=UPD
                SYSTEM LIBRARY  = LINKLIB
                LKED ATTRIBUTES = STD
                MODDEL          = MOD1      MOD2
                MODULES         = NAME      FMID
                                  MODULE1   ??????

 MODUL25   LASTUPD         = UCLIN    TYPE=UPD
                SYSTEM LIBRARY  = LINKLIB
                LKED ATTRIBUTES = STD
                MODULES         = NAME      FMID
                                  MODUL25   HXZ1234
```

UNLOAD Examples

To dump the LMOD entries in UCL format, you can use the UNLOAD command. To unload all the LMOD entries in a particular zone, you can use the following commands:

```
SET      BDY(TGT1)            /* Set to requested zone.   */.
UNLOAD   LMOD                 /* Unload all LMOD entries. */.
```

To unload specific LMOD entries, you can use these commands:

```
SET      BDY(TGT1)            /* Set to requested zone.   */.
UNLOAD   LMOD(LMOD01          /* Unload only these two    */
             LMOD02)          /* entries.                 */.
```

The format of the UNLOAD output for each LMOD entry is the same for both of these commands. The only difference is the number of LMOD entries listed. The following is an example of UNLOAD output for LMOD entries.

```
UCLIN .
REP       LMOD            ( DL1LMOD  )
          LASTUPD         ( DL1F001  )
          LASTUPDTYPE     ( ADD )
          SYSLIB          ( LINKLIB  )
```

```
              SIDEDECKLIB     ( SGOSSD   )
              /* LEPARM    */   RENT REUS RMOD=ANY DYNAM(DLL)
              CALLLIBS        ( SCEELOAD   SEDCBASE )
              /* UTILITY INPUT */
              UTIN            (
                              ( GOSLMOD1 SGOSSD  )
                              ( GOSLMOD2 SGOSSD  )
                                              )
++LMODIN
  ENTRY DL1MOD1
++ENDLMODIN
                            .

REP        LMOD            ( LMOD01   )
           LASTUPD         ( JXY1102 )
           LASTUPDTYPE     ( ADD )
           SYSLIB          ( LINKLIB  )
           RC              ( 4   )
           /* LEPARM    */   STD
           MODDEL          ( MOD05    MOD06  )
++LMODIN
 ENTRY MOD01
 ALIAS LMOD01A1
++ENDLMODIN
                            .
REP        LMOD            ( LMOD02   )
           LASTUPD         ( JXY1102 )
           LASTUPDTYPE     ( ADD )
           SYSLIB          ( LINKLIB   PPLIB01  )
           /* LEPARM    */   RENT REUS AC=1
                            .
REP        LMOD            ( LMOD03   )
           LASTUPD         ( JXY1102 )
           LASTUPDTYPE     ( ADD )
           SYSLIB          ( LINKLIB  )
           /* LEPARM    */   RENT REUS AC=1
++LMODIN
 ALIAS MOD03
 CHANGE MOD03C1(NEW03C1)
 CHANGE MOD03C2(NEW03C2)
 INCLUDE AOS12   (MOD03   )
++ENDLMODIN
                            .

REP        LMOD            ( LMOD04   )
           LASTUPD         ( LINK    )
           LASTUPDTYPE     ( UPD )
           SYSLIB          ( LINKLIB  )
           RC              ( 4   )
           /* LEPARM    */   STD
           /* CROSS-ZONE*/   XZMODP
           XZMOD           (
                           ( CICSMOD1 CICS1  )
                           ( CICSMOD2 CICS1  )
                           ( IMSMOD1  IMS1   )
                                           )
++LMODIN
 ENTRY MOD01
 ALIAS LMOD04A1
++ENDLMODIN
                            .
REP        LMOD            ( LMOD05   )
           /* CROSS-ZONE*/   XZMODP
           XZMOD           (
                           ( CICSMOD1 CICS1  )
                           ( CICSMOD2 CICS1  )
                           ( IMSMOD1  IMS1   )
                                           )
                            .
REP        LMOD            ( LMOD06   )
           LASTUPD         ( JCLIN    )
           LASTUPDTYPE     ( ADD )
           SYSLIB          ( LINKLIB  )
           /* LEPARM    */   SCTR MAXBLK(6160) AC=1 AMOD=MIN
                            .
REP        LMOD            ( LMOD07   )
           LASTUPD         ( JCLIN    )
           LASTUPDTYPE     ( ADD )
           SYSLIB          ( LINKLIB  )
           MODDEL          ( ISPLINK  MOD01  )
           /* LEPARM    */   FETCHOPT(PACK,PRIME) NOCALL
```

```
                           .
REP       LMOD           ( LMOD08   )
          LASTUPD        ( HXY0001 )
          LASTUPDTYPE    ( ADD )
          SYSLIB         ( APPLLIB  )
          /* LEPARM    */   RENT REUS
          CALLLIBS       ( PLIBASE  CSSLIB  )
                           .
```

```
REP       LMOD           ( LOTSOPRM )
          LASTUPD        ( UCLIN    )
          LASTUPDTYPE    ( UPD )
          SYSLIB         ( LINKLIB  )
          /* LEPARM    */   RENT REUS SCTR OVLY XCAL REFR DC NE AC=1
                            ALIGN2 AMODE=31 CASE(MIXED) OL MAXBLK(32760)
                            FETCHOPT(NOPACK,NOPRIME) COMPAT=LKED NCAL
                            FILL(00) HOBSET RMODE=SPLIT ALIASES(ALL)
                            DYNAM(DLL) UPCASE(YES)
++LMODIN
  ENTRY DL1MOD2
++ENDLMODIN
                           .
REP       LMOD           ( MOD04    )
          LASTUPD        ( JXY1102 )
          LASTUPDTYPE    ( ADD )
          SYSLIB         ( LPALIB   )
          /* LEPARM    */   COPY STD
                           .
ENDUCL.
```

UCLIN Examples

You can use the ADD, REP, and DEL UCL statements to change subentries in the LMOD entry. When you use UCLIN to update an LMOD entry, keep these points in mind:

- After the UCLIN changes are made, the LMOD entry must contain at least a SYSLIB subentry. Otherwise, there is not enough information in the entry to indicate where the load module should be installed.
- The input following the ++LMODIN statement replaces all existing link-edit control cards in the LMOD entry. This is different from JCLIN processing, where all control cards are replaced except CHANGE and REPLACE, which are merged with the existing CHANGE and REPLACE control cards.
- If a DEL statement deletes all the existing subentries in the entry, SMP/E deletes the entire entry.
- When SMP/E processes a DEL statement, it does not compare any control cards after the ++LMODIN statement with the control cards that are currently in the LMOD entry. It deletes all the existing control cards.

The following examples are provided to help you use the LMOD entry.

Example 1: Adding a new LMOD entry

The correct method of adding a new LMOD entry is through JCLIN, so that in addition to the LMOD entry being created, SMP/E also ensures that all related MOD entries are updated. If, however, you want to define a new load module entry using UCL, the following is an example of the minimum information you should provide.

```
SET       BDY(TGT1)          /* Set to target zone.    */.
UCLIN                        /*                        */.
ADD       LMOD(LMOD01)       /* Define new LMOD entry. */
          SYSLIB(LPALIB      /* System library 1.      */
                 PPLIB01)    /* System library 2.      */
          RENT REUS AC=1     /* All link attributes.   */
++LMODIN                     /* All link-edit stmts.   */
  INCLUDE AOS12(MOD01)
  INCLUDE AOS12(MOD02)
  ENTRY MOD01
  ALIAS LMOD01A1
  NAME LMOD01(R)
++ENDLMODIN
                             /*                        */.
REP       MOD(MOD01)         /* Now fix MOD entries.   */
          DISTLIB(AOS12)     /* MOD DLIB.              */
```

```
            FMID(FXY1102)      /* Functional owner.      */
            LMOD(LMOD01)       /* Connect to LMOD.       */
                               /*                        */.
REP         MOD(MOD02)         /* Now fix MOD entries.   */
            DISTLIB(AOS12)     /* MOD DLIB.              */
            FMID(FXY1102)      /* Functional owner.      */
            LMOD(LMOD01)       /* Connect to LMOD.       */
                               /*                        */.
ENDUCL                         /*                        */.
```

Note: In this example, the entire set of control cards needed to link the load module was specified. MOD entries for modules within the load module were also updated.

Example 2: Changing the link-edit attributes of an LMOD

As in the previous example, the correct way to change any part of a LMOD entry is through the use of JCLIN. However, at times, assuming a knowledge of SMP/E and UCL processing, you may decide to use UCL to do it. For this example, assume you have a load module, LMOD99, with link-edit attributes of RENT REUS, and that you have made changes so that it is no longer reentrant but is now authorized and will fill uninitialized areas. This change can be made as follows:

```
SET         BDY(TGT1)          /* Set to target zone.    */.
UCLIN                          /*                        */.
DEL         LMOD(LMOD99)       /* Existing LMOD entry.   */
            RENT.              /* Delete RENT attribute. */
ADD         LMOD(LMOD99)       /* Existing LMOD entry.   */
            AC=1 FILL(00)      /* Add authorization,     */
                               /*  and fill uninitialized */
                               /*  areas to hex byte 00.  */.
ENDUCL                         /*                        */.
```

Example 3: Deleting a MODDEL subentry

Suppose load module LMOD03 for product A included module MOD01 from product B, and both products are installed in zone ZOSZN. Product B deleted MOD01 without replacing it, and as a result, SMP/E created a MODDEL subentry for MOD01 in the LMOD entry for LMOD03. You have decided to reintroduce your own version of MOD01, but do not want it relinked into LMOD03. To prevent this relinking, delete the MODDEL subentry from the LMOD entry, as shown in the following example: :

```
SET         BDY(ZOSZN)          /* Set to zone.              */.
UCLIN                           /* Start UCLIN processing.   */.
DEL         LMOD(LMOD03)        /* Identify LMOD entry.      */
            MODDEL(MOD01)       /* Delete MODDEL subentry.   */.
ENDUCL                          /*                           */.
```

Example 4: Completing cross-zone updates

If SMP/E could not complete cross-zone updates for APPLY or RESTORE processing, you may need to use UCLIN to make the remaining changes. (SMP/E issues messages and reports to tell you what cross-zone work you may need to finish.) Here are some examples of using UCLIN to change XZMOD subentries in LMOD entries. Make sure you check the messages and reports to determine whether any additional changes are needed for cross-zone subentries in MOD entries or TARGETZONE entries.

- **Adding XZMOD subentries:** This example adds module IGCXXX from zone ZXXX and module IGCABC from zone ZABC to the IEANUC01 LMOD entry. If this LMOD entry previously had no XZMOD subentries, SMP/E automatically sets XZMODP to indicate that the LMOD entry now contains XZMOD subentries.

```
SET BDY(ZOSZN)                 /* Set to z/OS zone          */.
UCLIN                          /* start UCLIN processing    */.
ADD LMOD(IEANUC01)             /* identify LMOD entry       */
    XZMOD((IGCXXX,ZXXX),       /* add module from ZXXX zone */
          (IGCABC,ZABC))       /* add module from ZABC zone */.
ENDUCL                         /* end UCLIN processing      */.
```

- **Replacing XZMOD subentries:** This example replaces all the XZMOD subentries in LMOD entry CICSDIAG with a single subentry for module ISPLINK from zone ZOSZN. If this LMOD entry previously

had no XZMOD subentries, SMP/E automatically sets XZMODP to indicate that the LMOD entry now contains XZMOD subentries.

```
SET BDY(CICS1)              /* set to CICS zone            */.
UCLIN                       /* start UCLIN processing      */.
REP LMOD(CICSDIAG)          /* identify LMOD entry         */
    XZMOD((ISPLINK,ZOSZN))   /* replace list with one value  */.
ENDUCL                      /* end UCLIN processing        */.
```

- **Deleting XZMOD subentries:** This example shows how to delete a single XZMOD subentry or all XZMOD subentries.

  – For LMOD entry IMSDIAG1, one XZMOD subentry is deleted. If no XZMOD values are left, SMP/E turns the XZMODP indicator off to indicate that the LMOD entry no longer has XZMOD subentries.

  – For LMOD entry IMSDIAG2, all XZMOD subentries are deleted. SMP/E turns the XZMODP indicator off to indicate that the LMOD entry no longer contains XZMOD subentries.

```
SET BDY(IMS1)              /* set to IMS zone            */.
UCLIN                      /* start UCLIN processing     */.
DEL LMOD(IMSDIAG1)         /* identify LMOD entry        */
    XZMOD((ISPLINK,ZOSZN))   /* delete ISPLINK reference    */
DEL LMOD(IMSDIAG2)         /* identify LMOD entry        */
    XZMOD()               /* delete entire list         */.
ENDUCL                     /* end UCLIN processing       */.
```

Example 5: Adding a CALLLIBS subentry list to an LMOD entry

You can use either the ADD or the REP statement to add a CALLLIBS subentry list to an LMOD entry, depending on whether the CALLLIBS already exists or not. The order in which the libraries are specified is important because it indicates the order in which the SYSLIB concatenation is built. This is shown in the examples that follow.

- **Adding a CALLLIBS subentry:** This example adds a CALLLIBS subentry list containing PLIBASE and APPLIB to LMOD entry LMOD04.

```
SET BDY(ZOSZN)                /* set to z/OS zone              */.
UCLIN                         /* start UCLIN processing        */.
ADD LMOD(LMOD04)              /* identify LMOD entry           */
    CALLLIBS(PLIBASE,APPLIB) /* add CALLLIBS subentry         */.
ENDUCL                        /* end UCLIN processing          */.
```

- **Adding to an existing CALLLIBS subentry:** Suppose LMOD entry LMOD05 has a CALLLIBS subentry list containing PLIBASE and APPLIB, and CSSLIB is to be added to this list. The entire CALLLIBS subentry must be replaced as shown in the following example: :

```
SET BDY(ZOSZN)                /* set to z/OS zone              */.
UCLIN                         /* start UCLIN processing        */.
REP LMOD(LMOD05)              /* identify LMOD entry           */
    CALLLIBS(PLIBASE,APPLIB, /*                               */
            CSSLIB)          /*replace entire CALLLIBS        */.
ENDUCL.                       /* end UCLIN processing          */.
```

Note: If an ADD statement is used to try to update an existing CALLLIBS subentry list in an LMOD entry, SMP/E issues an error message.

Example 6: Deleting link-edit control statements

Suppose you have installed a SYSMOD that was packaged incorrectly and added unwanted link-edit control statements to the LMOD entry for an existing load module. As a result, you now need to delete those statements from the LMOD entry. Here is an example of the UCLIN statements you can use to make this change:

```
SET BDY(ZOSZN)                /* set to z/OS zone              */.
UCLIN                         /* start UCLIN processing        */.
DEL LMOD(LMODXZ)              /* identify LMOD entry           */
++LMODIN                      /* delete all link-edit          */
```

```
++ENDLMODIN                      /* control statements.       */.
ENDUCL                          /* end UCLIN processing        */.
```

Using the REP command instead of the DEL command gives you the same result. There is no need to specify the control statements to be deleted. In fact, you cannot selectively delete link-edit control statements from an LMOD entry—you can only delete them all.

Example 7: Adding a UTIN and SIDEDECKLIB subentry to an LMOD

In this example, a utility input subentry list, containing DSOMCMP1 from library SDSOMSD and DSOMCMP2 from library SDSOMSD, is added to the LMOD entry for load module LMOD1. In addition, a side deck library subentry of SGOSSD is added to the LMOD entry for load module LMOD5.

```
SET BDY(TGT)                    /* Set to target zone.         */.
UCLIN                           /* Start UCLIN processing.     */.
ADD LMOD(LMOD1)                 /* Identify LMOD entry.        */
  UTIN((DSOMCMP1,SDSOMSD)       /* Add utility input…          */
       (DSOMCMP2,SDSOMSD))      /* …subentry list.             */.
ADD LMOD(LMOD5)                 /* Identify LMOD entry.        */
  SIDEDECKLIB(SGOSSD)           /* Add side deck library.      */.
ENDUCL                          /* End UCLIN processing.       */.
```

Example 8: Deleting a UTIN subentry from an LMOD

In this example, DSOMCMP1 is deleted from the utility input subentry for LMOD entry LMOD2, and the entire utility input subentry list is removed from LMOD entry LMOD3.

```
SET BDY(TGT)                    /* Set to target zone.         */.
UCLIN                           /* Start UCLIN processing.     */.
DEL LMOD(LMOD2)                 /* Identify LMOD entry.        */
  UTIN((DSOMCMP1,SDSOMSD))      /* Remove DSOMCMP1 from        */
                                /*  utility input subentry.    */.
DEL LMOD(LMOD3)                 /* Identify LMOD entry.        */
  UTIN()                        /* Delete all utility input.   */.
ENDUCL                          /* End UCLIN processing.       */.
```

MAC entry (distribution and target zone)

The MAC entry describes a macro that exists in the distribution or target macro libraries. A MAC entry is generally created by one of the following methods:

- **Installing a SYSMOD that contains the macro.** MAC entries are created the first time you install a SYSMOD containing a ++MAC statement for a macro that does not yet have a MAC entry.

- **Processing JCLIN.** MAC entries can be built during JCLIN processing when SMP/E scans the assembler statements for inline assemblies to determine which macros are used in that assembly. The name of the assembly is kept in the MAC entry, so when the macro is updated, SMP/E can reassemble the required modules.

  MAC entries can also be built when SMP/E scans copy steps and finds a SELECT statement that specifies TYPE=MAC.

  For additional information, see the "Processing" section of the JCLIN command chapter in *z/OS SMP/E Commands*.

SMP/E records the function and service level of each macro in the MAC entry, as well as information about how that macro affects the structure of the distribution or target libraries and modules.

Once a MAC entry exists for a macro, it is updated as subsequent SYSMODs are installed.

Subentries

These are the subentries for the MAC entry as they appear in the LIST output:

name
    is the name of the macro represented by the MAC entry.

The name can contain from 1 to 8 alphanumeric characters and $, #, @, or hex C0.

DISTLIB

is the ddname of the macro distribution library.

The UCL operand is DISTLIB(*ddname*).

- The ddname can contain from 1 to 8 alphanumeric characters.
- The DISTLIB subentry is not required when the MAC entry is first defined. For example, during JCLIN processing SMP/E builds MAC entries without a DISTLIB value. However, SMP/E can add a DISTLIB value later when it processes the ++MAC statement. A DISTLIB value is needed at some time to process any changes for the macro.

FMID

identifies the functional owner of this macro. The functional owner is the last function SYSMOD that replaced this macro.

The UCL operand is FMID(*sysmod_id*).

The SYSMOD ID must contain 7 alphanumeric characters.

GENASM

identifies those assemblies that have to be done during APPLY each time this macro is modified. These assemblies must exist as either ASSEM or SRC entries in the target zone.

The UCL operand is GENASM(*name...*).

The names can contain from 1 to 8 alphanumeric characters.

LASTUPD

identifies the cause of the last change to this MAC entry.

The UCL operand is LASTUPD(*value*). This subentry can contain one of the following values:

JCLIN

indicates that the change was made during JCLIN command processing.

UCLIN

indicates that the change was made as a result of UCLIN processing.

sysmod-id

indicates that the change was made during the installation of the indicated SYSMOD.

The SYSMOD ID must contain 7 alphanumeric characters.

LASTUPD TYPE

indicates how the entry was last changed.

The UCL operand is LASTUPDTYPE(*value*). This subentry can contain one of the following values:

ADD

The entry was added.

MOV

The entry was moved.

UPD

The entry was updated.

MALIAS

identifies any alias names for this macro.

The UCL operand is MALIAS(*name...*).

The alias names can contain from 1 to 8 alphanumeric characters.

RMID

identifies the last SYSMOD that **replaced** this macro. Any subsequent SYSMOD that modifies this macro must have a defined relationship (such as PRE or SUP) with this SYSMOD.

The UCL operand is RMID(*sysmod_id*).

MAC entry (distribution and target zone)

- The SYSMOD ID must contain 7 alphanumeric characters.
- If RMID is not specified, but FMID is, SMP/E sets the RMID value to the specified FMID.

SYSLIB

is the ddname of the target system macro library.

The UCL operand is SYSLIB(*ddname*).

- Only one SYSLIB value may be specified.
- The ddname can contain from 1 to 8 alphanumeric characters.

UMID

identifies all those SYSMODs that have **updated** this macro since it was last replaced. Any subsequent SYSMOD that modifies this macro must have a defined relationship (such as PRE or SUP) with all these SYSMODs.

The UCL operand is UMID(*sysmod_id...*).

The SYSMOD ID must contain 7 alphanumeric characters.

LIST Examples

To list all the MAC entries in a particular zone, you can use the following commands:

```
SET      BDY(TGT1)            /* Set to requested zone.  */.
LIST     MAC                  /* List all MAC entries.   */.
```

To list specific MAC entries, you can use these commands:

```
SET      BDY(TGT1)            /* Set to requested zone.  */.
LIST     MAC(MAC01            /* List only these two     */
             MAC02)           /* entries.                */.
```

The format of the LIST output for each MAC entry is the same for both of these commands. The only difference is the number of MAC entries listed. shows an example of LIST output for MAC entries.

```
PAGE nnnn  - NOW SET TO zzzzzz ZONE nnnnnnn  DATE mm/dd/yy TIME hh:mm:ss SMP/E 36.nn SMPLIST OUTPUT

TGT1           MACRO ENTRIES


  NAME

MAC01      LASTUPD        = JXY1102 TYPE=ADD
           LIBRARIES      = DISTLIB=DLIBMAC1  SYSLIB=MACLIB01
           FMID           = JXY1102
           RMID           = JXY1102
           MALIAS         = TERMINAL  TERM     T
           GENASM         = ASSEM01   ASSEM02  SRC01     SRC02

MAC02      LASTUPD        = JXY1000 TYPE=UPD
           LIBRARIES      = DISTLIB=DLIBMAC1  SYSLIB=MACLIB01
           FMID           = JXY1121
           RMID           = UZ00010
           UMID           = UZ00014   UZ00015
           GENASM         = ASSEM01   SRC02
```

Figure 35. MAC entry: sample LIST output

By specifying the FORFMID operand, you can reduce the number of MAC entries listed. When FORFMID is specified, SMP/E lists a MAC entry only if its FMID matches one of the FMIDs specified on the FORFMID operand. For example, to list MAC entries whose FMIDs either are defined in FMIDSET TP or are JXY1102, you can use these commands:

```
SET      BDY(TGT1)            /* Set to target zone.     */.
LIST     MAC                  /* List all macro entries  */
```

```
              FORFMID(TP          /* for the TP FMIDSET      */
                     JXY1102)     /* and FMID JXY1102.       */.
```

You can use the LIST command to find out the names of all SYSMODs that have modified a macro. To include the names of these SYSMODs in the LIST output, you can use the XREF operand, as shown in these commands:

```
SET       BDY(TGT1)           /* Set to requested zone.  */.
LIST      MAC                 /* List all macro entries  */
          XREF                /* and related SYSMODs.    */.
```

Note:

1. XREF can be used either in mass mode or in select mode.

2. SMP/E obtains the data included for the XREF operand by checking for MAC and MACUPD entries for this macro in all the SYSMOD entries. Because this data is not contained in the MAC entry itself, you cannot use UCLIN to change it in the MAC entry.

Figure 36 on page 223 is an example of the LIST output produced when the XREF operand is used.

```
PAGE nnnn  - NOW SET TO zzzzzz ZONE nnnnnnn  DATE mm/dd/yy TIME hh:mm:ss SMP/E 36.nn SMPLIST OUTPUT

TGT1          MACRO ENTRIES


  NAME

MAC01     LASTUPD         = JXY1102 TYPE=ADD
          LIBRARIES       = DISTLIB=DLIBMAC1  SYSLIB=MACLIB01
          FMID            = JXY1102
          RMID            = JXY1102
          MALIAS          = TERMINAL  TERM      T
          GENASM          = ASSEM01   ASSEM02   SRC01    SRC02
          SYSMOD HISTORY  = SYSMOD   TYPE       DATE    MCS     --STATUS--
                            JXY1102  FUNCTION 07.100  MAC     APP    ACC

MAC02     LASTUPD         = JXY1000 TYPE=UPD
          LIBRARIES       = DISTLIB=DLIBMAC1  SYSLIB=MACLIB01
          FMID            = JXY1121
          RMID            = UZ00010
          UMID            = UZ00014   UZ00015
          GENASM          = ASSEM01   SRC02
          SYSMOD HISTORY  = SYSMOD   TYPE       DATE    MCS     --STATUS--
                            JXY1102  FUNCTION 07.100  MAC     APP    ACC
                            JXY1121  FUNCTION 07.150  MAC     APP    ACC
                            UZ00010  PTF      07.150  MAC     APP
                            UZ00014  PTF      07.160  MACUPD  APP
                            UZ00015  PTF      07.161  MACUPD  APP
```

Figure 36. MAC entry: sample LIST output when XREF is specified

UNLOAD Examples

To dump the MAC entries in UCL format, you can use the UNLOAD command. To unload all the MAC entries in a particular zone, you can use the following commands:

```
SET       BDY(TGT1)           /* Set to requested zone.  */.
UNLOAD    MAC                 /* Unload all MAC entries.  */.
```

To unload specific MAC entries, you can use these commands:

```
SET       BDY(TGT1)           /* Set to requested zone.  */.
UNLOAD    MAC(MAC01           /* Unload only these two   */
              MAC02).         /* entries.                */.
```

The format of the UNLOAD output for each MAC entry is the same for both of these commands. The only difference is the number of MAC entries listed. Figure 37 on page 224 is an example of UNLOAD output for MAC entries.

```
UCLIN .
REP       MAC              ( MAC01    )
          LASTUPD          ( JXY1102  )
          LASTUPDTYPE      ( ADD )
          DISTLIB          ( DLIBMAC1 )
          SYSLIB           ( MACLIB01 )
          FMID             ( JXY1102 )
          RMID             ( JXY1102 )
          MALIAS           ( TERMINAL  TERM      T       )
          GENASM           ( ASSEM01   ASSEM02   SRC01     SRC02    )
                              .
REP       MAC              ( MAC02    )
          LASTUPD          ( JXY1121  )
          LASTUPDTYPE      ( UPD )
          DISTLIB          ( DLIBMAC1 )
          SYSLIB           ( MACLIB01 )
          FMID             ( JXY1121 )
          RMID             ( UZ00010 )
          UMID             ( UZ00014   UZ00015  )
          GENASM           ( ASSEM01   SRC02    )
                              .
ENDUCL.
```

Figure 37. MAC entry: sample UNLOAD output

By specifying the FORFMID operand, you can reduce the number of MAC entries unloaded. When FORFMID is specified, SMP/E unloads a MAC entry only if its FMID matches one of the FMIDs specified on the FORFMID operand. For example, to unload MAC entries whose FMIDs either are defined in FMIDSET TP or are JXY1102, you can use these commands:

```
SET      BDY(TGT1)        /* Set to target zone.      */.
UNLOAD   MAC              /* Unload all macro entries */
         FORFMID(TP       /* for the TP FMIDSET       */
                 JXY1102) /* and FMID JXY1102.        */.
```

UCLIN Examples

You can use the ADD, REP, and DEL UCL statements to change subentries in the MAC entry. After the UCLIN changes are made, the MAC entry must contain at least the following subentries:

- FMID
- RMID

Otherwise, there is not enough information in the entry to process the macro. If any of these subentries are missing, SMP/E does not make the requested UCL updates to the entry, and the entry remains as it was before the UCL command.

The following examples are provided to help you use the MAC entry.

Example 1: Adding a new MAC entry

Defining a new macro entry with UCL is very seldom required; generally, MAC entries are created from the information specified on the ++MAC statements contained in the SYSMODs when SYSMODs are installed. If, however, you want to use UCL in defining a new macro entry, the following is an example of the minimum information you should provide:

```
SET      BDY(TGT1)        /* Set to target zone.        */.
UCLIN                     /*                            */.
ADD      MAC(MAC01)       /* Define new macro entry.    */
         DISTLIB(AMACLIB) /* Define DLIB,               */
         SYSLIB(MACLIB)   /* system library
                             (never the SMPMTS or MTS). */
         FMID(ZUSR001)    /* Functional owner (in this
                             example a user function).  */
                          /*                            */.
ENDUCL                    /*                            */.
SET      BDY(DLB1).       /* Now do same to DLIB.       */
```

```
UCLIN                        /*                       */.
ADD     MAC(MAC01)           /* Define new macro entry.  */
        DISTLIB(AMACLIB)     /* Define DLIB.          */
                             /* No SYSLIB info in DLIB.  */
        FMID(ZUSR001)        /* Functional owner (in this
                                example a user function). */.
                             /*                       */.
ENDUCL                       /*                       */.
```

Example 2: Defining an alias for an existing macro

The following defines the method of adding an alias to an existing macro:

```
SET     BDY(TGT1)            /* Set to target zone.    */.
UCLIN                        /*                       */.
ADD     MAC(MAC01)           /* Existing macro entry.  */
        MALIAS(MAC01AL1)     /* New alias name.       */
                             /* End of adding macro.   */.
ENDUCL                       /*                       */.
SET     BDY(DLB1)            /* Now do same thing to
                                appropriate DLIB.     */.
UCLIN                        /*                       */.
ADD     MAC(MAC01)           /* Existing macro entry.  */
        MALIAS(MAC01AL1)     /* New alias name.       */
                             /* End of adding macro.   */.
ENDUCL                       /*                       */.
```

Note: This UCL does not create an alias entry in the target or distribution libraries; that must be done outside of SMP/E using standard utilities. This ensures that, when the macro is subsequently modified, both the major entry and the alias entry are updated.

MCS entry (SMPPTS)

The MCS entry is a copy of a SYSMOD exactly as it was received from the SMPPTFIN data set. The MCS entry is in the SMPPTS data set, which is used as a warehouse for SYSMODs. When SMP/E receives a SYSMOD, it stores the SYSMOD as a separate member in the SMPPTS. The member name matches the SYSMOD ID, and each member is an MCS entry. SMP/E also creates a SYSMOD entry in the global zone to describe the SYSMOD that was received. Thus, the MCS entry and the global zone SYSMOD entry are closely related.

When SMP/E accepts or applies SYSMODs, it gets them from the MCS entries in the SMPPTS. An MCS entry is generally kept in the SMPPTS until the associated SYSMOD is accepted; then the entry is deleted.

- You may want SMP/E to save the MCS entries after ACCEPT processing, for example, if you plan to do a system generation. To do this, specify NOPURGE in the OPTIONS entry that is in effect during ACCEPT processing.
- Likewise, you may want to save the MCS entries after RESTORE processing. To do this, specify NOREJECT in the OPTIONS entry that is in effect during RESTORE processing.

Subentries

The MCS entry contains no SMP/E data and appears to the system as a member of a normal partitioned data set.

LIST Examples

To list all the MCS entries in the SMPPTS, you can use the following commands:

```
SET     BDY(GLOBAL)          /* Set to global.        */.
LIST    MCS                  /* List all MCS entries.  */.
```

To list specific MCS entries, you can use these commands:

```
SET      BDY(GLOBAL)       /* Set to global.        */.
LIST     MCS(UZ12345,       /* List only these two   */
         UZ12346)          /* entries.              */.
```

The format of the LIST output for each MCS entry is the same for both of these commands. The only difference is the number of MCS entries listed.

Figure 38 on page 226 is an example of LIST output for MCS entries.

```
PAGE nnnn  - NOW SET TO zzzzzz ZONE nnnnnnn  DATE mm/dd/yy TIME hh:mm:ss SMP/E 36.nn SMPLIST OUTPUT

SMPPTS        MCS ENTRIES


   NAME

UZ12345   MCS             = ++PTF(UZ12345).
                            ++VER(Z038) FMID(JXY1102).
                            ++MOD(XYMOD01) DISTLIB(AOS12).

UZ12346   MCS             = ++PTF(UZ12346).
                            ++VER(Z038) FMID(FXY1102).
                            ++JCLIN.
                            //JOB      JOB 'accounting info',MSGLEVEL=(1,1)
                            //STEP1    EXEC PGM=IEBCOPY
                            //MACLIB   DD DSN=SYS1.MACLIB,DISP=SHR
                            //AMACLIB  DD DSN=SYS1.AMACLIB,DISP=SHR
                            //SYSIN    DD *
                             COPY INDD=AMACLIB,OUTDD=MACLIB
                            /*
                            ++MAC(MAC01) DISTLIB(AMACLIB).
```

Figure 38. MCS entry: sample LIST output

As the example for UZ12346 shows, SMP/E includes inline JCLIN when listing the MCS entries.

You can use various SYSMOD-related LIST operands to limit which MCS entries are listed. For more information, see "SYSMOD entry (distribution and target zone)" on page 265 and "SYSMOD entry (global zone)" on page 277.

In addition to the LIST command, you can use standard system utility programs (such as IEBGENER, IEBPTPCH, IEHLIST, and so on) or products such as ISPF to display these entries or information about the data set.

UCLIN Examples

You cannot use UCLIN to add, update, or delete MCS entries. However, you can use the REJECT command to delete MCS entries. For more information about the REJECT command, see *z/OS SMP/E Commands*.

Do not use system utility programs to update MCS entries. The global zone SYSMOD entry is coordinated with the MCS entry. Any changes to the MCS entry made outside of SMP/E may get these entries out of synchronization and can result in unpredictable results when the associated SYSMOD is processed.

MOD entry (distribution and target zone)

The MOD entry describes a particular module, its function and service level, and how it relates to a load module in the target library. A MOD entry is generally created by one of the following methods:

- **Installing a SYSMOD that contains the module.** MOD entries are created the first time you install a SYSMOD that contains a ++MOD statement for a module that does not yet have a MOD entry. If the module comes from a copied distribution library, SMP/E also builds an LMOD entry with the same name. For additional information about copied libraries and the creation of LMOD entries, see "DLIB entry (distribution and target zone)" on page 174 and *z/OS SMP/E Commands*.

A MOD entry can also be created when you install a SYSMOD that causes the assembled source to be linked to the distribution library. This can happen when a SYSMOD contains both a ++MOD and a ++SRC statement for the same module, or when the DISTMOD operand is specified on the ++SRC statement.

- **Processing JCLIN.** MOD entries can be built during JCLIN processing when SMP/E scans the link-edit and copy steps. A MOD entry built during JCLIN processing is always associated with an LMOD entry. Thus, when the module is changed, SMP/E can determine which load modules are affected. For additional information, see the "Processing" section of the JCLIN command chapter in *z/OS SMP/E Commands*.

Subentries

These are the subentries for the MOD entry as they appear in the LIST output:

name
    is the name of the module represented by the MOD entry.

    The name can contain from 1 to 8 alphanumeric characters and $, #, @, or hex C0.

ASSEMBLE
    indicates that the source for this module must always be assembled, even if the object module is supplied in the SYSMOD.

    The UCL operand is ASSEMBLE.

CSECT
    specifies the CSECTs present in this module.

    The UCL operand is CSECT(*name*...).

- The CSECT subentry is not required. However, if CSECT is missing, SMP/E assumes that the module contains only one CSECT, whose name matches the module name.
- A CSECT name can contain from 1 to 8 characters.

    The name can contain any characters except the following:

    – Comma ,
    – Left parenthesis (
    – Right parenthesis)
    – Blank

- Comments are not allowed within a CSECT name. For example, the following is not allowed:

```
CSECT ( /* this is a csect name */ CSECT01)
```

    The comment is interpreted as part of the CSECT name, instead of a comment.

- The list of CSECT names must include all of the CSECTs in the module, even if one of the CSECTs matches the module name.
- Once a ++MOD with the CSECT operand is processed, SMP/E saves the CSECT information in the MOD entry. From then on, SMP/E uses that saved information if it is not supplied on subsequent SYSMODs.
- If a module is changed from multiple CSECTs to a single CSECT that matches the module name, the CSECT operand must be specified with one name in order to get SMP/E to store that information in the MOD entry.

DALIAS
    specifies the alias name for the module, where the alias exists only in the distribution library.

    The UCL operand is DALIAS(*name*).

    The DALIAS name can contain from 1 to 8 alphanumeric characters.

DISTLIB
    is the ddname of the module distribution library.

MOD entry (distribution and target zone)

The UCL operand is DISTLIB(*ddname*).

The ddname can contain from 1 to 8 alphanumeric characters.

FMID
identifies the functional owner of this module. The functional owner is the last function SYSMOD that replaced this module.

The UCL operand is FMID(*sysmod_id*).

- The SYSMOD ID must contain 7 alphanumeric characters.
- Some MOD entries, specifically those associated with a system generation assembly, may have no functional owner. The DISTLIB for these modules is SYSPUNCH. They may have either no FMID, or one of the following values:
  - If the module has no functional owner but was assembled during system generation, ∗SYSGEN appears as the FMID.
  - If the module has no functional owner and was not assembled during system generation, ??????? appears as the FMID.

LASTUPD
identifies the cause of the last change to this MOD entry.

Note: If a given UCLIN command specifies only cross-zone subentries, this field is not changed.

The UCL operand is LASTUPD(*value*). This subentry can contain one of the following values:

JCLIN
indicates that the change was made during JCLIN command processing.

UCLIN
indicates that the change was made as a result of UCLIN processing.

sysmod-id
indicates that the change was made during installation of the specified SYSMOD. The SYSMOD did one of the following:

- Contained inline JCLIN that affected the module
- Changed the distribution library for the module through the DISTLIB operand on the ++MOD statement
- Added the module to an existing load module through the LMOD operand on the ++MOD statement

The SYSMOD ID must contain 7 alphanumeric characters.

LASTUPD TYPE
indicates how the entry was last changed.

Note: If a given UCLIN command specifies only cross-zone subentries, this field is not changed.

The UCL operand is LASTUPDTYPE(*value*). This subentry can contain one of the following values:

ADD
The entry was added.

DEL
A subentry in the entry was deleted.

MOV
The module was moved.

UPD
The entry was updated.

LKED ATTRIBUTES
identifies the link-edit attributes that must be used when this module is link-edited. SMP/E supports the following link-edit attributes. For more information, see *z/OS MVS Program Management: User's Guide and Reference*.

AC=1

specifies that the AC=1 parameter, which is the authorization code, is to be passed to the link-edit utility when the module is link-edited.

The UCL operand is AC=1.

ALIGN2

specifies that the ALIGN2 parameter (alignment on a 2KB boundary) is to be passed to the link-edit utility when the module is link-edited.

The UCL operand is ALIGN2 or ALN2.

AMODE=24

specifies that the AMODE=24 parameter is to be passed to the link-edit utility when the module is link-edited.

The UCL operand is AMODE=24 or AMOD=24.

AMODE=31

specifies that the AMODE=31 parameter is to be passed to the link-edit utility when the module is link-edited.

The UCL operand is AMODE=31 or AMOD=31.

AMODE=64

specifies that the AMODE=64 parameter is to be passed to the link-edit utility when the module is link-edited.

The UCL operand is AMODE=64 or AMOD=64.

AMODE=ANY

specifies that the AMODE=ANY parameter is to be passed to the link-edit utility when the module is link-edited.

The UCL operand is AMODE=ANY or AMOD=ANY.

AMODE=MIN

specifies that the AMODE=MIN parameter is to be passed to the link-edit utility when the module is link-edited.

The UCL operand is AMODE=MIN or AMOD=MIN.

COMPAT

specifies that the COMPAT parameter is to be passed to the link-edit utility when the module is link-edited.

The UCL operand is COMPAT=LKED|PM1|PM2|PM3|PM4.

DC

specifies that the DC parameter, which is the downward compatible attribute, is to be passed to the link-edit utility when the module is link-edited.

The UCL operand is DC.

FETCHOPT

specifies that the FETCHOPT parameter is to be passed to the link-edit utility when the module is link-edited.

The UCL operand is FETCHOPT(PACK|NOPACK, PRIME|NOPRIME).

FILL

specifies that the FILL parameter is to be passed to the link-edit utility when the module is link-edited.

The UCL operand is FILL(*nn*), where *nn* is the 2-character representation of a single hex byte (00 - FF).

HOBSET
specifies that the HOBSET parameter is to be passed to the link-edit utility when the module is link-edited.

The UCL operand is HOBSET.

LONGPARM
specifies that the LONGPARM parameter is to be passed to the link-edit utility when the module is link-edited.

The UCL operand is LONGPARM.

MAXBLK
specifies that the MAXBLK parameter is to be passed to the link-edit utility when the module is link-edited.

The UCL operand is MAXBLK(*nnnnn*), where *nnnnn* is a number between 256 and 32760.

NE
specifies that the NE parameter, which is the noneditable attribute, is to be passed to the link-edit utility when the module is link-edited.

The UCL operand is NE.

NOCALL
specifies that the NCAL parameter is to be passed to the link-edit utility when the module is link-edited.

The UCL operand is NOCALL or NCAL.

OL
specifies that the OL parameter is to be passed to the link-edit utility when the module is link-edited.

The UCL operand is OL.

OVLY
specifies that the OVLY parameter, which specifies that the module is in overlay structure, is to be passed to the link-edit utility when the module is link-edited.

The UCL operand is OVLY.

REFR
specifies that the REFR parameter, which is the refreshable attribute, is to be passed to the link-edit utility when the module is link-edited.

The UCL operand is REFR.

RENT
specifies that the RENT parameter, which indicates that the module is reentrant, is to be passed to the link-edit utility when the module is link-edited.

The UCL operand is RENT.

REUS
specifies that the REUS parameter, which indicates that the module is reusable, is to be passed to the link-edit utility when the module is link-edited.

The UCL operand is REUS.

REUS(NONE)
specifies that the REUS(NONE) parameter, which indicates that the module cannot be reused, is to be passed to the link-edit utility when the module is link-edited.

The UCL operand is REUS(NONE).

RMODE=24
specifies that the RMODE=24 parameter is to be passed to the link-edit utility when the module is link-edited.

The UCL operand is RMODE=24 or RMOD=24.

RMODE=31

specifies that the RMODE=ANY parameter is to be passed to the link-edit utility when the module is link-edited. (RMODE=31 is a synonym for RMODE=ANY.)

The UCL operand is RMODE=31 or RMOD=31.

RMODE=64

specifies that the RMODE=64 parameter is to be passed to the link-edit utility when the module is link-edited.

The UCL operand is RMODE=64 or RMOD=64.

RMODE=ANY

specifies that the RMODE=ANY parameter is to be passed to the link-edit utility when the module is link-edited.

The UCL operand is RMODE=ANY or RMOD=ANY.

RMODE=SPLIT

specifies that the RMODE=SPLIT parameter is to be passed to the link-edit utility when the module is link-edited.

The UCL operand is RMODE=SPLIT or RMOD=SPLIT.

RMODEX=64TRUE

specifies that the RMODEX=64TRUE parameter is to be passed to the link-edit utility when the module is link-edited.

The UCL operand is RMODEX=64TRUE.

RMODEX=NO

specifies that the RMODEX=NO parameter is to be passed to the link-edit utility when the module is link-edited.

The UCL operand is RMODEX=NO.

SCTR

specifies that the SCTR parameter, which indicates that the module can be scatter-loaded, is to be passed to the link-edit utility when the module is link-edited.

The UCL operand is SCTR.

STD

is a special SMP/E indication that the module should be link-edited with none of the previously listed attributes.

When this indicator is present and a link-edit is to be done, SMP/E will pass the link-edit utility only those parameters specified in the appropriate link-edit UTILITY entry.

The UCL operand is STD.

UPCASE

specifies that the UPCASE parameter, which indicates how the Binder should process symbol names, is to be passed to the link-edit utility when the module is link-edited.

The UCL operand is UPCASE(YES|NO).

• UPCASE(YES) and UPCASE(NO) are mutually exclusive.

LMOD

specifies the names of the load modules into which this module was copied or link-edited on the target system.

The UCL operand is LMOD(*name...*).

• The load module names can contain from 1 to 8 alphanumeric characters.

MOD entry (distribution and target zone)

- If there are no LMOD names in a MOD entry, SMP/E assumes that the module was not selected when the SYSMOD containing the module was initially installed. Therefore, during APPLY processing, SMP/E will not link or copy the module into any target library.

RMID
identifies the last SYSMOD that **replaced** this module. Any subsequent SYSMOD that modifies this module must have a defined relationship (such as PRE or SUP) with this SYSMOD.

The UCL operand is RMID(*sysmod_id*).

- The SYSMOD ID must contain 7 alphanumeric characters.
- If RMID is not specified, but FMID is, SMP/E sets the RMID subentry to the FMID value.
- RMID is not required for a module assembled during system generation. The DISTLIB for these modules is SYSPUNCH.

RMIDASM
specifies that the last replacement (RMID) to the module was done by a SYSMOD that caused an assembly of the module as a result of a source or macro modification.

The UCL operand is RMIDASM.

TALIAS
specifies one or more alias names for the module, where the alias exists in the distribution library and, for copied modules, also in the target library.

The UCL operand is TALIAS(*name*...).

The alias names can contain from 1 to 8 alphanumeric characters.

UMID
identifies all those SYSMODs that have **updated** this module since it was last replaced. Any subsequent SYSMOD that modifies this module must have a defined relation (such as PRE or SUP) with all these SYSMODs.

The UCL operand is UMID(*sysmod_id*...).

The SYSMOD ID must contain 7 alphanumeric characters.

XZLMOD
specifies one or more load modules in other zones into which this module was added by the LINK MODULE command. The name of the zone supplying each module is also indicated.

The UCL operand is XZLMOD((*load module,zone*)...).

- The zone name specified for an XZLMOD subentry cannot match the name of the set-to zone.
- An entry can contain XZLMOD and XZLMODP subentries without any other subentries.
- If UCLIN is used to update an existing MOD entry and **only** cross-zone subentries are changed (XZLMOD and XZLMODP), SMP/E does not update the LASTUPD and LASTUPDTYPE subentries.
- The XZLMOD subentry is added to a MOD entry automatically during LINK MODULE command processing. However, it is **never** automatically removed.

XZLMODP
indicates that this module has been linked into one or more load modules controlled by a different target zone, and that XZLMOD subentries exist in this MOD entry.

The UCL operand is XZLMODP.

- It is never necessary to specify the XZLMODP subentry on a UCL statement. SMP/E automatically determines the setting of XZLMODP, according to whether the MOD entry contains XZLMOD subentries.
- You cannot add the XZLMODP subentry to a MOD entry that does not contain XZLMOD subentries.
- You cannot delete the XZLMODP subentry from a MOD entry containing XZLMOD subentries.
- An entry can contain XZLMOD and XZLMODP subentries without any other subentries.

- If UCLIN is used to update an existing MOD entry and **only** cross-zone subentries are changed (XZLMOD and XZLMODP), SMP/E does not update the LASTUPD and LASTUPDTYPE subentries.

LIST Examples

To list all the MOD entries in a particular zone, you can use the following commands:

```
SET       BDY(TGT1)            /* Set to requested zone.   */.
LIST      MOD                  /* List all MOD entries.    */.
```

To list specific MOD entries, you can use these commands:

```
SET       BDY(TGT1)            /* Set to requested zone.   */.
LIST      MOD(MOD01            /* List only these two      */
              MOD02)           /* entries.                 */.
```

The format of the LIST output for each MOD entry is the same for both of these commands. The only difference is the number of MOD entries listed. The following examples show LIST output for MOD entries. Figure 39 on page 234 does not have cross-zone subentries. Figure 40 on page 235 does.

MOD entry (distribution and target zone)

```
PAGE nnnn  - NOW SET TO zzzzzz ZONE nnnnnnn  DATE mm/dd/yy TIME hh:mm:ss SMP/E 36.nn SMPLIST OUTPUT

TGT1          MODULE ENTRIES

   NAME

ASSEM01   LASTUPD         = JCLIN   TYPE=ADD
          LIBRARIES       = DISTLIB=SYSPUNCH
          LMOD            = LMOD99A   LMOD99B

ASSEM02   LASTUPD         = JCLIN   TYPE=ADD
          LIBRARIES       = DISTLIB=SYSPUNCH
          LMOD            = ASSEM02

DL1MOD1   LASTUPD         = DL1F001 TYPE=ADD
          LIBRARIES       = DISTLIB=DLIB3
          FMID            = DL1F001
          RMID            = DL1F001
          LMOD            = DL1LMOD

MOD01     LASTUPD         = JXY1102 TYPE=ADD
          LIBRARIES       = DISTLIB=AOS12
          FMID            = JXY1102
          RMID            = JXY1102
          CSECT           = MOD01C1   MOD01C2   MOD01C3
          LMOD            = LMOD01

MOD02     LASTUPD         = JXY1102 TYPE=ADD
          LIBRARIES       = DISTLIB=AOS12
          FMID            = JXY1102
          RMID            = JXY1102
          CSECT           = MOD02C1
          LMOD            = LMOD01   LMOD02

MOD03     LASTUPD         = JXY1121 TYPE=UPD
          LIBRARIES       = DISTLIB=AOS12
          FMID            = JXY1121
          RMID            = JXY1121
          UMID            = UZ00010   UZ00014
          LMOD            = LMOD03

MOD04     LASTUPD         = JXY1121 TYPE=UPD
          LIBRARIES       = DISTLIB=AOS12
          ASSEMBLE
          FMID            = JXY1121
          RMID            = JXY1121
                            RMIDASM
          LMOD            = MOD04

MOD05     LASTUPD         = JXY0001 TYPE=ADD
          LIBRARIES       = DISTLIB=DLIB1
          FMID            = JXY0001
          RMID            = JXY0001

MOD06     LASTUPD         = JXY0001 TYPE=ADD
          LIBRARIES       = DISTLIB=DLIB1
          FMID            = JXY0001
          RMID            = JXY0001
```

Figure 39. MOD entry: sample LIST output (no cross-zone subentries)

```
PAGE nnnn  - NOW SET TO zzzzzz ZONE nnnnnnn  DATE mm/dd/yy TIME hh:mm:ss SMP/E 36.nn SMPLIST OUTPUT

NOW SET TO TARGET ZONE CICS1

CICS1         MODULE ENTRIES

  NAME

CICSMOD1  LASTUPD        = LINK    TYPE=UPD
          LIBRARIES      = DISTLIB=ARESLIB
          FMID           = HCI1703
          RMID           = HCI1703
          LMOD           = CICSMOD1
          XZLMODP
          XZLMOD         = LMOD01   IN ZONE ZOSZNA
                           LMOD02   IN ZONE ZOSZNA
                           LMOD01   IN ZONE ZOSZNB

CICSMOD2  LASTUPD        = LINK    TYPE=UPD
          LIBRARIES      = DISTLIB=ARESLIB
          FMID           = HCI1703
          RMID           = HCI1703
          LMOD           = CICSMOD2
          XZLMODP
          XZLMOD         = LMOD01   IN ZONE ZOSZNA
                           LMOD02   IN ZONE ZOSZNA
                           LMODAA01 IN ZONE ZOSZNB
                           LMOD01   IN ZONE ZOSZNB

CICSMOD3  LASTUPD        = HCI1703 TYPE=ADD
          LIBRARIES      = DISTLIB=ARESLIB
          FMID           = HCI1703
          RMID           = HCI1703
          LMOD           = CICSLMD2

CICSMOD4  XZLMODP
          XZLMOD         = LMOD01   IN ZONE ZOSZNA
                           LMOD01   IN ZONE ZOSZNB
```

Figure 40. MOD entry: sample LIST output (cross-zone entries)

By specifying the FORFMID operand, you can reduce the number of MOD entries listed. When FORFMID is specified, SMP/E lists a MOD entry only if its FMID matches one of the FMIDs specified on the FORFMID operand. For example, to list MOD entries whose FMIDs are either defined in FMIDSET TP or are JXY1102, you can use these commands:

```
SET     BDY(TGT1)          /* Set to target zone.   */.
LIST    MOD                /* List all MOD entries  */
        FORFMID(TP         /* for the TP FMIDSET    */
              JXY1102)     /* and FMID JXY1102.     */.
```

You can also use the LIST command to find out the names of all SYSMODs that have updated a module. To include the names of these SYSMODs in the LIST output, you can use the XREF operand, as shown in these commands:

```
SET     BDY(TGT1)          /* Set to requested zone.  */.
LIST    MOD                /* List all module entries */
        XREF               /* and related SYSMODs.    */.
```

Note:

1. XREF can be used either in mass mode or in select mode.
2. SMP/E obtains the data included for the XREF operand by checking all the SYSMOD entries to find which SYSMODs contained ++MOD or ++ZAP statements for this module. Because this data is not contained in the MOD entry itself, you cannot use UCLIN to change it in the MOD entry.

is an example of the LIST output produced when the XREF operand is used.

MOD entry (distribution and target zone)

```
PAGE nnnn  - NOW SET TO zzzzzz ZONE nnnnnnn  DATE mm/dd/yy TIME hh:mm:ss SMP/E 36.nn SMPLIST OUTPUT

TGT1          MODULE ENTRIES


   NAME

ASSEM01   LASTUPD         = JCLIN    TYPE=ADD
          LIBRARIES       = DISTLIB=SYSPUNCH
          LMOD            = LMOD99A   LMOD99B

ASSEM02   LASTUPD         = JCLIN    TYPE=ADD
          LIBRARIES       = DISTLIB=SYSPUNCH
          LMOD            = ASSEM02
MOD01     LASTUPD         = JXY1102 TYPE=ADD
          LIBRARIES       = DISTLIB=AOS12
          FMID            = JXY1102
          RMID            = JXY1102
          CSECT           = MOD01C1   MOD01C2   MOD01C3
          LMOD            = LMOD01
          SYSMOD HISTORY  = SYSMOD    TYPE      DATE    MCS     --STATUS--
                            JXY1102  FUNCTION 07.100  MOD     APP    ACC
MOD02     LASTUPD         = JXY1102 TYPE=ADD
          LIBRARIES       = DISTLIB=AOS12
          FMID            = JXY1102
          RMID            = JXY1102
          CSECT           = MOD02C1
          LMOD            = LMOD01    LMOD02
          SYSMOD HISTORY  = SYSMOD    TYPE      DATE    MCS     --STATUS--
                            JXY1102  FUNCTION 07.100  MOD     APP    ACC
MOD03     LASTUPD         = JXY1121 TYPE=UPD
          LIBRARIES       = DISTLIB=AOS12
          FMID            = JXY1121
          RMID            = JXY1121
          UMID            = UZ00010   UZ00014
          LMOD            = LMOD03
          SYSMOD HISTORY  = SYSMOD    TYPE      DATE    MCS     --STATUS--
                            JXY1102  FUNCTION 07.100  MOD     APP    ACC
                            JXY1121  FUNCTION 07.150  MOD     APP    ACC
                            UZ00010  PTF      07.150  ZAP     APP
                            UZ00014  PTF      07.160  ZAP     APP


MOD04     LASTUPD         = JXY1121 TYPE=UPD
          LIBRARIES       = DISTLIB=AOS12
          ASSEMBLE
          FMID            = JXY1121
          RMID            = JXY1121
                            RMIDASM
          LMOD            = MOD04
          SYSMOD HISTORY  = SYSMOD    TYPE      DATE    MCS     --STATUS--
                            JXY1102  FUNCTION 07.100  MOD     APP    ACC
                            JXY1121  FUNCTION 07.150  ASSEM   APP    ACC
```

Figure 41. MOD entry: sample LIST output when XREF is specified

UNLOAD Examples

To dump the MOD entries in UCL format, you can use the UNLOAD command. To unload all the MOD entries in a particular zone, you can use the following commands:

```
SET     BDY(TGT1)          /* Set to requested zone.  */.
UNLOAD  MOD                /* Unload all MOD entries. */.
```

To unload specific MOD entries, you can use these commands:

```
SET     BDY(TGT1)          /* Set to requested zone.  */.
UNLOAD  MOD(MOD01          /* Unload only these two   */
           MOD02)          /* entries.                */.
```

The format of the UNLOAD output for each MOD entry is the same for both of these commands. The only difference is the number of MOD entries unloaded. The following examples show UNLOAD output for MOD entries. Figure 42 on page 237 does not have cross-zone subentries. Figure 43 on page 238 does.

```
UCLIN .
REP        MOD                ( ASSEM01  )
           LASTUPD            ( JCLIN    )
           LASTUPDTYPE        ( ADD )
           DISTLIB            ( SYSPUNCH )
           LMOD               ( LMOD99A   LMOD99B  )
                              .
REP        MOD                ( ASSEM02  )
           LASTUPD            ( JCLIN    )
           LASTUPDTYPE        ( ADD )
           DISTLIB            ( SYSPUNCH )
           LMOD               ( ASSEM02  )
                              .
REP        MOD                ( DL1MOD1  )
           LASTUPD            ( DL1F001 )
           LASTUPDTYPE        ( ADD )
           DISTLIB            ( DLIB3    )
           FMID               ( DL1F001 )
           RMID               ( DL1F001 )
           LMOD               ( DL1LMOD  )
                              .
REP        MOD                ( MOD01    )
           LASTUPD            ( JXY1102  )
           LASTUPDTYPE        ( ADD )
           DISTLIB            ( AOS12    )
           FMID               ( JXY1102 )
           RMID               ( JXY1102 )
           CSECT              ( MOD01C1   MOD01C2    MOD01C3  )
           LMOD               ( LMOD01   )
                              .
REP        MOD                ( MOD02    )
           LASTUPD            ( JXY1102  )
           LASTUPDTYPE        ( ADD )
           DISTLIB            ( AOS12    )
           FMID               ( JXY1102 )
           RMID               ( JXY1102 )
           CSECT              ( MOD02C1 )
           LMOD               ( LMOD01    LMOD02   )
                              .
REP        MOD                ( MOD03    )
           LASTUPD            ( JXY1121  )
           LASTUPDTYPE        ( UPD )
           DISTLIB            ( AOS12    )
           FMID               ( JXY1121 )
           RMID               ( JXY1121 )
           UMID               ( UZ00010   UZ00014  )
           LMOD               ( LMOD03   )
                              .
REP        MOD                ( MOD04    )
           LASTUPD            ( JXY1121  )
           LASTUPDTYPE        ( UPD )
           DISTLIB            ( AOS12    )
           ASSEMBLE
           FMID               ( JXY1121 )
           RMID               ( JXY1121 )
                                RMIDASM
           LMOD               ( MOD04    )
                              .
REP        MOD                ( MOD05    )
           LASTUPD            ( JXY0001  )
           LASTUPDTYPE        ( ADD )
           DISTLIB            ( DLIB1    )
           FMID               ( JXY0001 )
           RMID               ( JXY0001 )
                              .
ENDUCL.
```

Figure 42. MOD entry: sample UNLOAD output (no cross-zone subentries)

MOD entry (distribution and target zone)

```
UCLIN .
REP      MOD              ( CICSMOD1 )
         LASTUPD          ( HCI1703  )
         LASTUPDTYPE      ( ADD )
         DISTLIB          ( ARESLIB  )
         FMID             ( HCI1703 )
         RMID             ( HCI1703 )
         LMOD             ( CICSMOD1 )
         /* CROSS-ZONE*/   XZLMODP
         XZLMOD           (
                           ( LMOD01   ZOSZNA   )
                           ( LMOD01   ZOSZNB   )
                                               )
                 .
REP      MOD              ( CICSMOD2 )
         LASTUPD          ( HCI1703  )
         LASTUPDTYPE      ( ADD )
         DISTLIB          ( ARESLIB  )
         FMID             ( HCI1703 )
         RMID             ( HCI1703 )
         LMOD             ( CICSMOD2 )
         /* CROSS-ZONE*/   XZLMODP
         XZLMOD           (
                           ( LMOD01   ZOSZNA   )
                           ( LMOD01   ZOSZNB   )
                           ( LMODAA01 ZOSZNB   )
                                               )
                 .
REP      MOD              ( CICSMOD3 )
         LASTUPD          ( HCI1703  )
         LASTUPDTYPE      ( ADD )
         DISTLIB          ( ARESLIB  )
         FMID             ( HCI1703 )
         RMID             ( HCI1703 )
         LMOD             ( CICSLMD2 )
                 .
REP      MOD              ( CICSMOD4 )
         /* CROSS-ZONE*/   XZLMODP
         XZLMOD           (
                           ( LMOD01   ZOSZNA   )
                           ( LMOD01   ZOSZNB   )
                                               )
                 .
ENDUCL.
```

Figure 43. MOD entry: sample UNLOAD output (cross-zone subentries)

By specifying the FORFMID operand, you can reduce the number of MOD entries unloaded. When FORFMID is specified, SMP/E unloads a MOD entry only if its FMID matches one of the FMIDs specified on the FORFMID operand. For example, to unload MOD entries whose FMIDs are either defined in FMIDSET TP or are JXY1102, you can use these commands:

```
SET      BDY(TGT1)           /* Set to target zone.      */.
UNLOAD   MOD                 /* Unload all MOD entries    */
         FORFMID(TP          /* for the TP FMIDSET        */
                 JXY1102)    /* and FMID JXY1102.         */.
```

UCLIN Examples

You can use the ADD, REP, and DEL UCL statements to change subentries in the MOD entry. After the UCLIN changes are done, the MOD entry must contain at least the following subentries:

• DISTLIB

• FMID

• RMID

Otherwise, there is not enough information in the entry to process the module. If any of the required subentries are missing, SMP/E does not make the requested UCL updates to the entry, and the entry remains as it was before the UCL command.

The following examples are provided to help you use the MOD entry.

Example 1: Adding a new MOD entry

To create a MOD entry, you should use one of the following methods rather than UCLIN:

- JCLIN, if a new module is being added to a new load module. For examples of how JCLIN creates MOD entries, see *z/OS SMP/E Commands*.
- The LMOD operand on the ++MOD statement, if a new module is being added to an existing load module and no changes are needed in the link-edit control statements other than the INCLUDE for the new module. For an example of how to do this, see "++MOD MCS" on page 60.

You could make the same changes with the UCLIN command. However, you should make sure the changes are correct and complete. For example, assume you want to define two new modules, MOD99A and MOD99B. These are linked together to form load module LMOD99AB, which has an entry point of MOD99A and exists in LINKLIB. The FMID of both new modules is ZUSR001, a user-created function. You could use these commands to create the MOD and LMOD entries:

```
SET       BDY(TGT1)         /* Set to target zone.     */.
UCLIN                       /*                         */.
ADD       MOD(MOD99A)       /* Define new module entry. */
          DISTLIB(USRDLIB1) /* Define DLIB.            */
          LMOD(LMOD99AB)    /* Load module.            */
          FMID(ZUSR001)     /* Functional owner.       */
                            /*                         */.
ADD       MOD(MOD99B)       /* Define new module entry. */
          DISTLIB(USRDLIB1) /* Define DLIB.            */
          LMOD(LMOD99AB)    /* Load module.            */
          FMID(ZUSR001)     /* Functional owner.       */
                            /*                         */.
ADD       LMOD(LMOD99AB)    /* Now define LMOD.        */
          SYSLIB(LINKLIB)   /*                         */
          RENT REUS         /* Attributes.             */
++LMODIN                    /* Link statements.        */
  INCLUDE USRDLIB1(MOD99A,MOD99B)
  ENTRY MOD99A
++ENDLMODIN                 /*                         */
                            /*                         */.
ENDUCL                      /*                         */.
```

Example 2: Forcing assembly of a module

Assume you have a macro that is used by a source, and you want to make sure the source is always assembled when the macro is changed. You can do this by setting the ASSEMBLE indicator as follows:

```
SET     BDY(DLIB1)        /* Set to DLIB zone.      */.
UCLIN                     /*                        */.
ADD     MOD(SRC01)        /* MOD entry for source.  */
        ASSEMBLE          /* Always assemble it.    */
                          /*                        */.
ENDUCL                    /*                        */.
```

Example 3: Completing cross-zone updates

If SMP/E could not complete cross-zone updates for APPLY or RESTORE processing, you may need to use UCLIN to make the remaining changes. (SMP/E issues messages and reports to tell you what cross-zone work you may need to finish.) Here are some examples of using UCLIN to change XZLMOD subentries in MOD entries. Make sure to check the messages and reports to determine whether any additional changes are needed for cross-zone subentries in LMOD entries or TARGETZONE entries.

- **Adding XZLMOD subentries:** This example adds LMOD IEANUC01 to the IGCXXX MOD entry in zone IMS1. If this MOD entry previously had no XZLMOD subentries, SMP/E automatically sets XZLMODP to indicate that the MOD entry now contains XZLMOD subentries.

```
SET BDY(IMS1)                 /* set to IMS zone         */.
UCLIN                         /* start UCLIN processing  */.
ADD MOD(IGCXXX)               /* identify MOD entry       */
```

```
    XZLMOD((IEANUC01,ZOSZNA))  /* add load module from zone ZOSZNA */.
ENDUCL                         /* end UCLIN processing          */.
```

- **Replacing XZLMOD subentries:** This example replaces all the XZLMOD subentries in MOD entry
  ISPLINK with two XZLMOD subentries. If this MOD entry previously had no XZLMOD subentries, SMP/E
  automatically sets XZLMODP to indicate that the MOD entry now contains XZLMOD subentries.

```
SET BDY(ZOSZNA)                  /* set to first zone.          */.
UCLIN                   /* start UCLIN processing        */.
REP MOD(ISPLINK)        /* identify MOD entry            */
    XZLMOD((CICSDIAG,CICS1)  /* replace list with two values  */
          (IMSDIAG1,IMS1))  /*                              */.
ENDUCL                  /* end UCLIN processing          */.
```

- **Deleting XZLMOD subentries:** This example deletes an XZLMOD subentry from the ISPLINK subentry.
  If this causes the XZLMOD subentry list to become empty, SMP/E automatically turns XZLMODP off to
  indicate that the MOD entry no longer contains XZLMOD subentries.

```
SET BDY(ZOSZNB)                  /* set to second zone.         */.
UCLIN                   /* start UCLIN processing        */.
DEL MOD(ISPLINK)        /* identify MOD entry            */
    XZLMOD((IMSDIAG2,IMS1))  /* delete IMSDIAG2 reference    */.
ENDUCL                  /* end UCLIN processing          */.
```

MTSMAC entry (SMPMTS)

The MTSMAC entry is a copy of a macro that resides only in a distribution library but is needed temporarily
during APPLY processing. The MTSMAC entry is in the SMPMTS data set, which serves as a target macro
library for such macros.

When SMP/E applies the SYSMODs that affect these macros, it calls utility programs to store the macros
on the SMPMTS. This way, the most current service level of each macro is available for use in assemblies.
After SMP/E has accepted all the SYSMODs that affect these macros, it deletes the associated MTSMAC
entries from the SMPMTS.

Note: If you specify SAVEMTS in the OPTIONS entry that is in effect during ACCEPT processing, SMP/E
will not delete MTSMAC entries from the SMPMTS after the SYSMODs that affect those macros have been
successfully accepted.

Subentries

The MTSMAC entry contains no SMP/E data and appears to the system as a member of a normal macro
library.

LIST Examples

You cannot use SMP/E to list the MTSMAC entries. However, you can use standard system utility programs
(such as IEBGENER, IEBPTPCH, and IEHLIST) or products such as ISPF to display these entries or
information about the data set.

UCLIN Examples

You can use the DEL UCL statement to delete an MTSMAC entry from the SMPMTS. This can be helpful
if you plan to do an APPLY followed by ACCEPT when several target libraries have been created from the
same distribution library.

When a SYSMOD is accepted into a distribution zone, the entries associated with it are automatically
deleted from the SMPMTS for the RELATED target zone. However, even if the SYSMOD was also applied to
other target zones created from the same distribution zone, SMP/E does not clean up the SMPMTS data
sets for the other target zones.

To delete the entries from these data sets, you can accept the SYSMOD and name these other target
zones as the RELATED zone. However, this would update the distribution library each time; this is time-
consuming and could use up space in the distribution library data set.

Instead, you can use the DEL command to delete these entries without updating the distribution library. To determine which entries to specify, check the SMPLOG data set to see which ones SMP/E deleted during ACCEPT processing.

Note: You can also use the CLEANUP command to delete MTSMAC entries without specifying them individually. For more information about the CLEANUP command, see *z/OS SMP/E Commands*.

Example: Deleting an MTSMAC entry

Assume that you have two target zones, TGT1 and TGT2, generated off the same distribution zone, DLB1. During ACCEPT processing of a SYSMOD, SMP/E has deleted MTSMAC MAC01 and MAC02 from the SMPMTS data set associated with target zone TGT2. After performing the ACCEPT, you want to delete the same macro from the SMPMTS associated with target zone TGT1. Assume either that you have a cataloged procedure for TGT1 with the correct SMPMTS specified, or that you have set up the correct DDDEF entries. The following UCLIN can be used to delete the MTSMAC entry:

```
SET       BDY(TGT1)         /* Set to TGT zone.      */.
UCLIN                       /*                       */.
DEL       MTSMAC(MAC01)     /* Delete the macro.     */.
DEL       MTSMAC(MAC02)     /* Delete the macro.     */.
ENDUCL                      /*                       */.
```

Note: One UCL statement is required for each MTSMAC entry to be deleted.

You can make the same changes by use of system utilities. However, the SMPLOG will not reflect the processing done.

OPTIONS entry (global zone)

The OPTIONS entry defines processing options that are to be used for an SMP/E command or set of commands. Although OPTIONS entries exist only in the global zone, they are also used to process commands for target and distribution zones. There are two ways you can specify the OPTIONS entry that should be in effect when you are processing a zone:

- In the GLOBALZONE, TARGETZONE, and DLIBZONE entries. The OPTIONS entry specified here is the default OPTIONS entry for that zone.
- On the SET command. The name specified on the SET command overrides the default OPTIONS name.

When SMP/E processes a command, it checks these sources to determine which OPTIONS entry should be in effect for that command. If SMP/E cannot find a reference to an OPTIONS entry, it will use defaults for some of the subentries. The following section describes these defaults, as well as the other values you can specify for each subentry.

Subentries

These are the subentries for the OPTIONS entry as they appear in the LIST output:

name
> is the name of the OPTIONS entry.
>
> The name can contain from 1 to 8 alphanumeric characters.

AMS
> is the name of the UTILITY entry that SMP/E is to use to obtain information when calling the access method services (AMS) utility.
>
> The UCL operand is AMS(*name*).
>
> - The name can contain from 1 to 8 alphanumeric characters.
> - If no entry name is specified, SMP/E uses a default UTILITY entry. For details, see Table 7 on page 290.

ASM
is the name of the UTILITY entry that SMP/E is to use to obtain information when calling the assembler utility.

The UCL operand is ASM(*name*).

- The name can contain from 1 to 8 alphanumeric characters.
- If no entry name is specified, SMP/E uses a default UTILITY entry. For details, see Table 7 on page 290.

CHANGEFILE
specifies whether library change file records should be generated and written to the SMPDATA1 and SMPDATA2 data sets during APPLY and RESTORE command processing.

The UCL operand is CHANGEFILE(YES|NO), where:

YES
indicates that library change file records should be generated during APPLY or RESTORE processing.

NO
indicates that library change file records should **not** be generated during APPLY or RESTORE processing. This is the default for APPLY and RESTORE processing.

COMP
is the name of the UTILITY entry that SMP/E is to use to obtain information when calling the utility program to compress data sets.

The UCL operand is COMP(*name*).

- The name can contain from 1 to 8 alphanumeric characters.
- If no entry name is specified, SMP/E uses a default UTILITY entry. For details, see Table 7 on page 290.

COMPACT
specifies whether inline element data within SYSMODs in the SMPPTS data set should be compacted. The element data is normally compacted during the GZONEMERGE and RECEIVE commands.

The UCL operand is COMPACT(YES|NO), where:

YES
indicates inline element data within SYSMODs in the SMPPTS should be compacted to reduce the space requirements of the SMPPTS during the GZONEMERGE and RECEIVE commands. The element data is expanded as needed during ACCEPT and APPLY command processing. YES is the default.

NO
indicates inline element data within SYSMODs in the SMPPTS should not be compacted during the GZONEMERGE and RECEIVE commands. The element data will reside in the SMPPTS data set in its original form.

COPY
is the name of the UTILITY entry that SMP/E is to use to obtain information when calling the copy utility.

The UCL operand is COPY(*name*).

- The name can contain from 1 to 8 alphanumeric characters.
- If no entry name is specified, SMP/E uses a default UTILITY entry. For details, see Table 7 on page 290.

DSPREFIX
specifies the data set prefix to be used to construct the full data set name when SMPTLIB data sets are being allocated for RELFILEs. For more information about names for SMPTLIB data sets, see *z/OS SMP/E Commands*.

The UCL operand is DSPREFIX(*prefix*).

- The prefix can contain from 1 to 26 alphanumeric characters.
- The prefix must follow standard naming conventions for data sets.
- If a DDDEF entry is in effect for the SMPTLIB data sets, the DSPREFIX value in that DDDEF entry overrides the DSPREFIX value specified in the OPTIONS entry in effect.
- If no DSPREFIX value is specified in either the OPTIONS entry or the DDDEF entry, no high-level qualifier is assigned to the SMPTLIB data sets.
- If the DSPREFIX value is the same as the RFDSNPFX value for the RELFILE data sets that are being processed, the SMPTLIB data sets cannot be allocated in either of these cases:
  - The RELFILE data sets are on DASD.
  - The RELFILE data sets are on tape and are cataloged.

DSSPACE
specifies the primary and secondary space allocation (in tracks) and the number of directory blocks to be allocated for each SMPTLIB data set. After the data set is loaded, unused space is freed. For more information about SMPTLIB data sets, see *z/OS SMP/E Commands*.

The UCL operand is DSSPACE(*prime*, *secondary*, *directory*).

- Each value can contain from 1 to 4 numeric characters.
- If a DDDEF entry is in effect for the SMPTLIB data sets, the SPACE and DIR values in that DDDEF entry override the DSSPACE values specified in the OPTIONS entry in effect.
- These values should be specified in the appropriate OPTIONS or DDDEF entries before you receive a relative file. Otherwise, SMP/E cannot allocate any new SMPTLIB data sets.

EXRTYDD
specifies the list of ddnames that are not eligible for retry processing after an *x*37 abend. EXRTYDD is used with RETRYDDN in order to exclude a subset of libraries from retry processing.

The UCL operand is EXRTYDD(*ddname*...).

- The ddnames can contain from 1 to 8 alphanumeric characters.
- If a ddname is specified in both the EXRTYDD list and the RETRYDDN list, the ddname is excluded from retry processing.
- If no ddnames are specified for RETRYDDN, no retry processing will be done for any libraries. The ddnames specified on EXRTYDD are ignored.
- If ALL is specified on EXRTYDD, it is treated as just another ddname; it does **not** exclude all ddnames from retry processing. To exclude all libraries from retry processing, do **one** of the following instead:
  - Specify RETRY(NO) on the SMP/E command being processed.
  - Do not specify a RETRYDDN list in the OPTIONS entry that is in effect for the SMP/E command being processed.
  - Do not have an OPTIONS entry in effect for the SMP/E command being processed.

FIXCAT
specifies the list of fix categories. During APPLY, ACCEPT, and REPORT MISSINGFIX processing, this subentry, or the FIXCAT operand on the command, identifies the fix categories of interest and is used to determine which FIXCAT HOLDDATA entries will affect command processing.

The UCL operand is FIXCAT(*category*...).

- A Fix Category value is case sensitive, can be 1- to 64-characters in length, can contain any nonblank character in the range X'41' - X'FE' except single quotation mark ('), comma (,), left parenthesis ((), and right parenthesis ()), and can be specified in two ways:
  - Explicitly, by fully specifying a particular fix category value. For example, IBM.Device.zIIP. In this case, all HOLDDATA associated with this Fix Category is applicable to command processing.

- Implicitly, by partially specifying a fix category value using any number of asterisks (*) as global characters and percent signs (%) as placeholders.
  - A single asterisk indicates that zero or more characters can occupy that position. Here are some examples:

    IBM.Device*
      In this example, all HOLDDATA associated with a fix category that begins with the character string IBM.Device is applicable.

    ***z/OS**
      In this example, all HOLDDATA associated with a fix category that ends with the character string z/OS is applicable.

    IBM*z/OS
      In this example, all HOLDDATA associated with a fix category that begins with the character string IBM and ends with the character string z/OS is applicable.

- A single percent sign indicates that any one single character can occupy that position. For example, IBM.Device.20%4 indicates that HOLDDATA associated with any of the following fix categories is applicable: IBM.Device.2084, IBM.Device.2094, and IBM.Device.20t4. HOLDDATA associated with a fix category of IBM.Device.20914, however, is not applicable.

The following examples are all acceptable fix categories:

```
IBM.Device.zIIP
*
IBM.Function*
IBM.Device.20%4.*
*.HealthChecker
```

HFSCOPY
is the name of the UTILITY entry that SMP/E is to use to obtain information when calling the HFS copy utility.

This is the utility used to install hierarchical file system elements and must meet the same program-to-program interface as BPXCOPY.

The UCL operand is HFSCOPY(*name*).

- The name can contain from 1 to 8 uppercase alphanumeric characters.
- If no HFSCOPY subentry is specified in the OPTIONS entry, SMP/E uses a default UTILITY entry. For details, see Table 7 on page 290.

IOSUP
is the name of the UTILITY entry that SMP/E is to use to obtain information when calling the IEHIOSUP utility program to process maintenance for an OS/VS1 system.

The UCL operand is IOSUP(*name*).

- The name can contain from 1 to 8 alphanumeric characters.

LKED
is the name of the UTILITY entry that SMP/E is to use to obtain information when calling the link-edit utility.

The UCL operand is LKED(*name*).

- The name can contain from 1 to 8 alphanumeric characters.
- If no entry name is specified, SMP/E uses a default UTILITY entry. For details, see Table 7 on page 290.

MSGWIDTH
specifies whether the message output (SMPOUT only) should be formatted in 80 or 120 character widths.

The UCL operand is MSGWIDTH(*value*). The subentry can contain one of the following values:

80
indicates that the message text should be formatted to an 80 character width.

120
indicates that the message text should be formatted to a 120 character width. This is the default.

MSGFILTER
specifies whether messages issued to SMPOUT during APPLY, ACCEPT, and RESTORE processing should be filtered to contain only the following:

- Messages with a severity of Warning or higher.
- Messages with a severity of Informational that add clarity or additional information to a previously issued higher severity message.
- Messages GIM20501I and GIM20502I.

The UCL operand is `MSGFILTER(YES|NO)`, where:

YES
indicates that the number of messages issued to SMPOUT should be in a filtered format for APPLY, ACCEPT, and RESTORE processing.

NO
indicates that the number of messages issued to SMPOUT should be in a unfiltered format for APPLY, ACCEPT, and RESTORE processing. This is the default.

NOPURGE
indicates that after SMP/E accepts SYSMODs, it should not delete the associated global zone SYSMOD entries, SMPPTS MCS entries, or SMPTLIB data sets.

The UCL operand is `NOPURGE`.

- The default is for `NOPURGE` not to be specified. In this case, SMP/E deletes the entries after the SYSMOD has been successfully accepted. (This is true only if the SYSMOD has been applied and `BYPASS(APPLYCHECK)` was not specified on the ACCEPT command.) The associated HOLDDATA, including internal SYSTEM HOLDDATA, is not deleted when the SYSMOD is deleted.
- Although this operand can be specified in any OPTIONS entry, it is effective only when the OPTIONS entry is used during ACCEPT processing.

NOREJECT
specifies that the global zone SYSMOD entry and the associated MCS entry should not be deleted after the SYSMOD is restored.

The UCL operand is `NOREJECT`.

- The default is for `NOREJECT` not to be specified. In this case, SMP/E deletes the entries after the SYSMOD has been successfully restored. The associated HOLDDATA, including internal SYSTEM HOLDDATA, is not deleted when the SYSMOD is deleted.
- Although this operand can be specified in any OPTIONS entry, it is effective only when the OPTIONS entry is used during RESTORE processing.

ORDERRET
Indicates the retention period, in days, that ORDER entries are kept in the global zone before being deleted. During RECEIVE ORDER processing, an ORDER entry will be deleted from the global zone if either of the following conditions occurs:

- The ORDER entry has a status of DOWNLOADED, and the difference between the current date and the ORDER entry's DOWNLDATE subentry value is greater than the OPTIONS entry ORDER RETENTION subentry value.
- The ORDER entry has a status of ERROR or PENDING and the difference between the current date and the ORDER entry's ORDERDATE subentry value is greater than the OPTIONS entry ORDER RETENTION subentry value.

When an ORDER entry is deleted from the global zone, SMP/E also deletes the order package stored in the SMPNTS.

The UCL operand is ORDERRET(*nnnn*).

- The value can contain from 1 to 4 numeric characters from 0 to 9999.
- If no value is specified, the default is 180 days.

PAGELEN
specifies the page length for the SMPLIST, SMPHRPT, SMPOUT, and SMPRPT data sets.

The UCL operand is PAGELEN(*nnnn*).

- The value can contain from 1 to 4 numeric characters.
- If no value is specified, the default is 60.

PEMAX
specifies the maximum number of subentries that can be present in any CSI entry. Most often, the largest entry is a SYSMOD entry.

The UCL operand is PEMAX(*nnnn*).

- The value can contain from 1 to 4 numeric characters, with a value range of 1 to 9999.
- If no value is specified, the default is 32767. Therefore, if you want a value higher than 9999, you must use the default.

RECEXZGRP
specifies a list of zones or zonesets to be excluded during APPLYCHECK and ACCEPTCHECK processing during the SYSMOD selection phase of RECEIVE processing.

The UCL operand is RECEXZGRP(*value*...).

- Each *value* is a zone or zoneset name. Each name can contain from 1 to 8 alphanumeric characters.
- The zone or zonesets can contain both target and distribution zones.

  Note: The value specified cannot be GLOBAL nor ALLZONES.
- If a zone is specified in both the RECEXZGRP list and the RECZGRP list, the zone is excluded from processing.
- If no zones are specified for RECZGRP, the zones specified on RECEXZGRP are ignored.

RECZGRP
specifies a list of zones and zonesets eligible for APPLYCHECK and ACCEPTCHECK processing during the SYSMOD selection phase of RECEIVE processing. Any SYSMOD that has been applied or accepted into any of the zones specified by this list will not be received by the RECEIVE command, unless the SYSMOD is specified on the SELECT operand of the RECEIVE command or the list is overridden by the BYPASS(ACCEPTCHECK) or BYPASS(APPLYCHECK) operands, or a new list is specified by the ZONEGROUP operand.

The UCL operand is RECZGRP(*value*...).

- Each *value* is a zone or zoneset name. Each name can contain from 1 to 8 alphanumeric characters.
- The zone or zonesets can contain both target and distribution zones.

  Note: The value specified cannot be GLOBAL.
- ALLZONES can be specified to indicate that all zones defined by a ZONEINDEX subentry in the GLOBALZONE entry are eligible. When ALLZONES is specified, all other values specified are ignored.
- If no zones or zonesets are specified, no APPLYCHECK or ACCEPTCHECK processing is performed unless the ZONEGROUP operand is specified on the RECEIVE command.

RETRY
is the name of the UTILITY entry that SMP/E is to use to obtain information when calling the utility program to compress a data set after an *x*37 abend.

The UCL operand is RETRY(*name*).

- The name can contain from 1 to 8 alphanumeric characters.

- If no entry name is specified, SMP/E uses a default UTILITY entry. For details, see Table 7 on page 290.

RETRYDDN
    specifies the list of ddnames eligible for retry processing after an *x*37 abend.

    The UCL operand is RETRYDDN(*ddname*...).

- The ddnames can contain from 1 to 8 alphanumeric characters.
- ALL can be specified to indicate that all ddnames are eligible.
- If no ddnames are specified, no retry processing will be done, even if RETRY was specified on the command being processed.
- Retry processing is attempted for partitioned data sets (PDSs) and partitioned data sets extended (PDSEs) that experience an *x*37 abend, indicating that they have run out of space.
- For pointers on how to set up the desired retry processing, see *z/OS SMP/E User's Guide*.

SAVEMTS
    indicates that MTSMAC entries should not be deleted from the SMPMTS after the SYSMODs that affect those macros have been successfully accepted.

    The UCL operand is SAVEMTS.

- If SAVEMTS is not specified, SMP/E deletes the entries after all the SYSMODs that affect the macros have been successfully accepted. (This is true only if the SYSMOD has been applied and BYPASS(APPLYCHECK) was not specified on the ACCEPT command.)
- Although this operand can be specified in any OPTIONS entry, it is effective only when the OPTIONS entry is used during ACCEPT processing.

SAVESTS
    indicates that STSSRC entries should not be deleted from the SMPSTS after the SYSMODs that affect the source have been successfully accepted.

    The UCL operand is SAVESTS.

- If SAVESTS is not specified, SMP/E deletes the entries after all the SYSMODs that affect the source have been successfully accepted. (This is true only if the SYSMOD has been applied and BYPASS(APPLYCHECK) was not specified on the ACCEPT command.)
- Although this operand can be specified in any OPTIONS entry, it is effective only when the OPTIONS entry is used during ACCEPT processing.

SUPPHOLD
    specifies a list of HOLD reason IDs for which the HOLDDATA card image should not be displayed on the following reports:

- Unresolved HOLD Reason Report
- Bypassed HOLD Reason Report
- SYSMOD Comparison HOLDDATA Report

    The UCL operand is SUPPHOLD(*value*...), where:

value
    is a HOLD reason ID, which is from 1 to 7 alphanumeric characters, #, @, and $. The first character cannot start with 0-9.

UPDATE
    is the name of the UTILITY entry that SMP/E is to use to obtain information when calling the update utility.

    The UCL operand is UPDATE(*name*).

- The name can contain from 1 to 8 alphanumeric characters.
- If no entry name is specified, SMP/E uses a default UTILITY entry. For details, see Table 7 on page 290.

ZAP

specifies the name of the UTILITY entry that SMP/E is to use to obtain information when calling the superzap utility.

The UCL operand is ZAP(*name*).

- The name can contain from 1 to 8 alphanumeric characters.
- If no entry name is specified, SMP/E uses a default UTILITY entry with NAME(IMASPZAP), RC(4), PRINT(SYSPRINT), and PARM().

LIST Examples

To list all the OPTIONS entries in a global zone, you can use the following commands:

```
SET      BDY(GLOBAL)        /* Set to requested zone.    */.
LIST     OPTIONS            /* List all OPTIONS entries.  */.
```

To list specific OPTIONS entries in a global zone, you can use these commands:

```
SET      BDY(GLOBAL)        /* Set to requested zone.    */.
LIST     OPTIONS(OPT1,      /* List only these three     */
             TGT1,          /* entries.                  */
              DLIB1)        /*                           */.
```

The format of the LIST output for each OPTIONS entry is the same for both of these commands. The only difference is the number of OPTIONS entries listed.

Figure 44 on page 248 is an example of LIST output for OPTIONS entries.

```
PAGE nnnn   - NOW SET TO GLOBAL ZONE    DATE mm/dd/yy TIME hh:mm:ss SMP/E 36.nn SMPLIST OUTPUT

GLOBAL          OPTIONS ENTRIES


ALLOPTS    AMS              = AMS
           ASM              = ASM
           COMP             = COMPRESS
           COPY             = COPY
           HFSCOPY          = BPXCOPY
           IOSUP            = IOSUP
           LKED             = LINK
           RETRY            = COMPRESS
           UPDATE           = UPDATE
           ZAP              = ZAP
           PAGELEN          = 60
           PEMAX            = 9999
           NOPURGE
           COMPACT SMPPTS   = YES
           NOREJECT
           DSSPACE          = PRI=0100   SEC=0050   DIR=0200
           MSGFILTER        = YES
           MSGWIDTH         = 80
           CHANGEFILE       = NO
           ORDER RETENTION  = 90
           DSPREFIX         = SMPE.SMPTLIB
           RETRYDDN         = ALL
           EXRTYDD          = LINKLIB    LPALIB
           RECZGRP          = ALLZONES
           RECZGRP          = DLIB       TGT
           SUPPHOLD         = DOC        IPL
           FIXCAT              = IBM.Device.zIPP
                            *.HealthChecker
```

Figure 44. OPTIONS entry: sample LIST output

UCLIN Examples

You can use the ADD, REP, and DEL UCL statements to change subentries in an OPTIONS entry. When you use UCLIN to update an OPTIONS entry, remember that if a DEL statement deletes all the existing subentries in the entry, SMP/E deletes the entire entry.

The following examples are provided to help you use the OPTIONS entry.

Example 1: Connecting an OPTIONS entry to UTILITY entries

Assume that you have set up three UTILITY entries: IEUASM, IEWL, and MYX37. For examples of setting up these entries, see "UTILITY entry (global zone)" on page 289. Now you want to define an OPTIONS entry, OPT1, that will be used during APPLY processing. The following UCL can be used to set up that entry:

```
SET       BDY(GLOBAL)        /* Set to global zone.      */.
UCLIN                        /*                          */.
ADD       OPTIONS(OPT1)      /* New OPTIONS entry.       */
          ASM(IEUASM)        /* Connect to assembler data. */
          LKED(IEWL)         /* Connect to link data.    */
          RETRY(MYX37)       /* Connect to retry.        */
                             /*                          */.
ENDUCL                       /*                          */.
```

To use the OPTIONS entry, you need either to define it as the default OPTIONS entry for the desired GLOBALZONE, DLIBZONE, or TARGETZONE entry, or to specify it on the OPTIONS operand of the SET command for the zone you want to process with these values.

Example 2: Changing the SMPOUT page length

Assume that you want to modify the OPTIONS entry created in "Example 1: Connecting an OPTIONS entry to UTILITY entries" on page 249 to indicate that your printer now prints 120 lines per page. The following UCL will do this:

```
SET       BDY(GLOBAL)        /* Set to global zone.     */.
UCLIN                        /*                         */.
ADD       OPTIONS(OPT1)      /* Add because page length is
                                not there yet.          */
          PAGELEN(120)       /* Change page length.     */
                             /*                         */.
ENDUCL                       /*                         */.
```

Example 3: Preparing to receive RELFILEs

To receive a new function packaged in RELFILE format, you have to define the amount of space to allocate the SMPTLIB data sets and, optionally, the prefix to be used in building the data set names. Assume that the program directory indicates that 300 tracks of primary space should be used, 50 tracks of secondary space, and 25 directory blocks, and that you want the data set names to start with SMP.RELFILE.M9801. The following UCL can be used to add this information to the OPTIONS entry created in "Example 1: Connecting an OPTIONS entry to UTILITY entries" on page 249:

```
SET       BDY(GLOBAL)                 /* Set to global zone. */.
UCLIN                                 /*                     */.
ADD       OPTIONS(OPT1)               /* Add because data is
                                         not there yet.      */
          DSSPACE(300,50,25)          /* Space allocation.   */
          DSPREFIX(SMP.RELFILE.M9801) /* Prefix.             */
                                      /*                     */.
ENDUCL                                /*                     */.
```

Example 4: Identifying libraries for retry processing

During processing of commands that update product libraries (such as ACCEPT, APPLY, RESTORE, and LINK), utilities called by SMP/E may issue *x*37 abends when the libraries they are updating run out of

space. SMP/E can attempt to recover from such out-of-space errors if RETRY(YES) was specified on the command being processed **and** if a RETRYDDN list is available in the OPTIONS entry that is in effect.

Assume that you want SMP/E to attempt this retry processing for all libraries **except** LINKLIB, MIGLIB, and NUCLEUS. The following UCL can be used to specify the desired ddnames in the OPTIONS entry created in "Example 1: Connecting an OPTIONS entry to UTILITY entries" on page 249:

```
SET     BDY(GLOBAL)                /* Set to global zone. */.
UCLIN                              /*                     */.
ADD     OPTIONS(OPT1)              /* Add because data is
                                      not there yet.     */
        RETRYDDN(ALL)              /* Retry all ddnames   */
        EXRTYDD(LINKLIB,MIGLIB,NUCLEUS) /* except these.   */
                                   /*                     */.
ENDUCL                            /*                     */.
```

ORDER entry (global zone)

The ORDER entry describes a HOLDDATA or PTF order initiated with the RECEIVE ORDER command. When SMP/E sends an order request to the IBM Automated Delivery Request server and the server accepts the order, SMP/E creates an ORDER entry in the global zone. The ORDER entry is used to record information about the order so that SMP/E can query the server for status of orders that have not been completed. Once the processing for an order has been completed by the server, SMP/E can download the HOLDDATA or PTF package for the order and store the package files in the SMPNTS directory.

Subentries

These are the subentries for the ORDER entry as they appear in the LIST output:

name
> is the name of the ORDER entry.
>
> The name can contain from 1 to 8 alphanumeric characters.

CONTENT
> the HOLDDATA or PTF content for this order. The value can be one of the following:
>
> **ALL**
> > indicates a request for all available PTFs.
>
> **APARS**
> > indicates a request for PTFs that resolve the APARS specified in the APARS subentry.
>
> **CRITICAL**
> > indicates a request for all available PTFs that resolve a critical problem. A critical problem is a high impact pervasive (HIPER) or a PTF in error (PE).
>
> **HOLDDATA**
> > indicates a request for HOLDDATA only. A HOLDDATA package contains the last 2-years worth of Enhanced HOLDDATA for the entire z/OS software platform. For more information, see Enhanced HOLDDATA for z/OS (service.software.ibm.com/holdata/390holddata.html).
>
> **PTFS**
> > indicates a request for the PTFs indicated in the PTFS subentry.
>
> *sourceid*
> > indicates a request for all the recommended PTFs identified with the indicated Recommended Service Update SOURCEID (RSU*yymm*), and PTFs that resolve a critical problem (HIPER or PE). Recommended service includes PTFs through the indicated RSU level, which is the most current RSU level at the time the order is created.

STATUS
> indicates the status of the HOLDDATA or PTF order. Status values can be any one of the following:

PENDING
the order request has been submitted to the server but the order's package has not yet been downloaded.

DOWNLOADED
the order package has been downloaded and stored in the SMPNTS directory.

ERROR
the server reported an uncorrectable error with the processing of this order.

APARS
identifies the APARS for which resolving PTFs were requested.

PTFS
identifies the specific PTFs that were requested.

ORDERDATE
indicates the date on which the order request was sent to the server.

ORDERTIME
indicates the time at which the order request was sent to the server.

DOWNLDATE
indicates the date on which the order's package was downloaded and stored in the SMPNTS directory.

DOWNLTIME
indicates the time at which the order's package was downloaded and stored in the SMPNTS directory.

ORDERID
the order identifier assigned by the IBM Automated Delivery Request server when the order request was sent. The value is used to correlate ORDER entries in the global zone with orders being processed by the server.

USERID
the userid that submitted the SMP/E job which created the ORDER entry.

PKGID
indicates the subdirectory name within the SMPNTS directory that contains the package associated with the order.

ZONES
indicates the target zones used to represent the software inventory associated with this order.

ORDERSERVER
contains the <ORDERSERVER> tags used to identify the IBM Automated Delivery Request server scheduled to fulfill the order request. This information is obtained from the ORDERSERVER data set specified in the RECEIVE ORDER command when the ORDER entry was created.

This subentry exists only when the order has a STATUS of PENDING. This subentry is deleted once the order package has been downloaded.

LIST Examples

To list all the ORDER entries in a global zone, you can use the following commands:

```
SET     BDY(GLOBAL)        /* Set to requested zone.    */.
LIST    ORDER              /* List all ORDER entries.  */.
```

To list specific ORDER entries in a global zone, you can use these commands:

```
SET     BDY(GLOBAL)        /* Set to requested zone.   */.
LIST    ORDER(ORD00034,    /* List only these three    */
            ORD00035,      /* entries.                 */
             ORD00036)     /*                          */.
```

The format of the LIST output for each ORDER entry is the same for both of these commands. The only difference is the number of ORDER entries listed.

Figure 45 on page 252 is an example of LIST output for ORDER entries. In this example, processing for ORDER entries ORD00034 and ORD00035 is completed and the packages have been downloaded. The order described by ORDER entry ORD00036 is in the PENDING state and the package has not yet been downloaded.

```
PAGE nnnn  - NOW SET TO GLOBAL ZONE   DATE mm/dd/yy TIME hh:mm:ss SMP/E 36.nn SMPLIST OUTPUT

GLOBAL        ORDER ENTRIES


  NAME

ORD00034    CONTENT          = RSU0608
          STATUS         = DOWNLOADED
          DATE/TIME ORDER = 06.258  09:11:17
          DATE/TIME DOWNL = 06.258  09:52:11
          USERID         = JOHNDOE
          ORDERID        = 0012341119
          PKGID          = ORD00034-15September2006-09.47.02.258
          ZONES          = ZOS14


ORD00035    CONTENT          = PTFS
          STATUS         = DOWNLOADED
          DATE/TIME ORDER = 06.272  15:02:47
          DATE/TIME DOWNL = 06.272  15:19:38
          USERID         = JOHNDOE
          ORDERID        = 0123456785
          PKGID          = ORD00035-29September2006-15.14.23.272
          PTFS           = UQ12345   UQ98765
          ZONES          = ZOS14


ORD00036    CONTENT          = ALL
          STATUS         = PENDING
          DATE/TIME ORDER = 06.281  10:01:53
          USERID         = JOHNDOE
          ORDERID        = 0234567893
          ZONES          = ZOS14
          ORDER SERVER   = <ORDERSERVER
                 url="https:eccgw01.boulder.ibm.com/services/projects/ecc/ws/"
                 keyring="MRWKYRNG"
                 certificate="SMPE Certificate">
               </ORDERSERVER>
```

Figure 45. ORDER entry: sample LIST output

Note: An alternative url for the IBM Automated Delivery Request server is `https://eccgw02.rochester.ibm.com/services/projects/ecc/ws`.

UCLIN Examples

You can use the DEL UCL statement to delete an ORDER entry as shown in the following example. UCLIN does not support either adding entries or updating subentries for an ORDER entry.

Example: Deleting an ORDER entry

The following UCL can be used to delete an ORDER entry:

```
SET     BDY(GLOBAL)        /* Set to global zone   */.
UCLIN                      /* Start UCLIN processing    */.
  DEL     ORDER(ORD00035)  /* Delete ORDER entry        */.
ENDUCL                     /* End UCLIN processing      */.
```

ORDER entries may reside only in the GLOBALZONE.

PRODUCT entry (global zone)

The PRODUCT entry describes a software product. A product is known to SMP/E by the combination of its *prodid* and the *vv.rr.mm* values. PRODUCT entries reside only in the global zone.

Subentries

prodid
>   specifies the product identifier for the product described in the PRODUCT entry. For IBM products, it is recommended that the *prodid* be the IBM program product number (for example, "5647-A01").

vv.rr.mm
>   specifies the version, release and modification level for the product described in the PRODUCT entry.

DESCRIPTION
>   describes the product.
>
>   The UCL operand is DESCRIPTION(*description*).

SREL
>   specifies the system or subsystem releases on which the PRODUCT can be installed. There can be multiple SREL subentries associated with a PRODUCT.
>
>   The UCL operand is SREL(*srel,...*).

PRODSUP
>   specifies the PRODUCTs that are superseded by this PRODUCT. Multiple PRODSUP subentries can be associated with a PRODUCT. The combination of *prodid* and *vv.rr.mm* determines the uniqueness of a subentry in the PRODSUP subentry list. Only one subentry value for a given *prodid* and *vv.rr.mm* combination is saved in the PRODSUP subentry list.
>
>   The UCL operand is PRODSUP((*prodid,vv.rr.mm*),...).

URL
>   specifies the URL that can be accessed to obtain additional information about the product.
>
>   The UCL operand is URL(*product_url*).

VENDOR
>   specifies the name of the vendor supplying the product.
>
>   The UCL operand is VENDOR(*vendor_name*).

RECDATE
>   specifies the date on which the PRODUCT was received.
>
>   There is no UCL support for this subentry.

RECTIME
>   specifies the time at which the PRODUCT was received.
>
>   There is no UCL support for this subentry.

UCLDATE
>   specifies the date on which the PRODUCT was last processed using the UCLIN command.
>
>   There is no UCL support for this subentry.

UCLTIME
>   specifies the time at which the PRODUCT was last processed using the UCLIN command.
>
>   There is no UCL support for this subentry.

REWORK
>   identifies the level of this PRODUCT, which was received again for minor changes.
>
>   There is no UCL support for this subentry.

LIST Examples

To list all the PRODUCT entries in a particular zone, you can use the following commands:

```
SET       BDY(GLOBAL)          /* Set to global zone.      */.
LIST      PRODUCT              /* List all PRODUCT entries.*/.
```

To list specific PRODUCT entries, you can use these commands:

```
SET       BDY(GLOBAL)          /* Set to global zone.      */.
LIST      PRODUCT(5647-A01)    /* List all PRODUCT entries */
                               /* with prodid of 5647-A01  */.
LIST      PRODUCT((5645-001,01.03.00))
                               /* List PRODUCT entry with  */
                               /* prodid of 5645-001       */
                               /* and VRM of 01.03.00      */.
LIST      PRODUCT(5668-949,(5645-001,1.2.0))
                               /* List all PRODUCT entries */
                               /* with prodid of 5668-949, */
                               /* plus PRODUCT entry with  */
                               /* prodid of 5645-001       */
                               /* and VRM of 01.02.00      */.
```

The format of the LIST output for each PRODUCT entry is the same for all of these commands. The only difference is the number of PRODUCT entries listed. shows an example of LIST output for PRODUCT entries. PRODUCT entries are listed alphanumerically by *prodid*. PRODUCT entries with the same *prodid* value are further sorted by *vv.rr.mm*.

```
PAGE nnnn  - NOW SET TO zzzzzz ZONE nnnnnnn  DATE mm/dd/yy TIME hh:mm:ss SMP/E 36.nn SMPLIST OUTPUT

GLOBAL  PRODUCT  ENTRIES


  NAME

5647-A01  VRM             = 02.09.00
          DESCRIPTION     = OS/390
          REWORK          = 1999215
          URL             = http://www.s390.ibm.com/os390/
          DATE/TIME REC   = 07.267  12:55:39
                    UCL   = 08.331  02:25:45
          VENDOR          = IBM
          PRODSUP         = 5645-001  02.07.00  5668-949  01.08.00
          SREL            = Z038
```
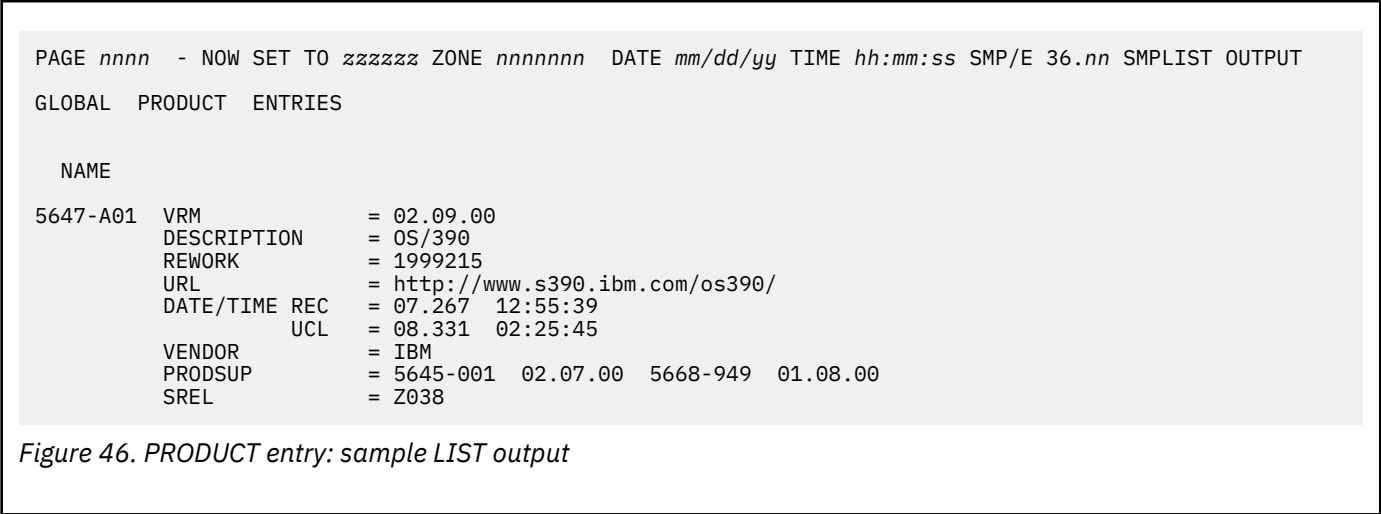
Figure 46. PRODUCT entry: sample LIST output

Note:

1. The entry name value for the PRODUCT entries formatted to the left of the subentry names is the prodid value.

2. If the URL value extends past column 120, it is continued on the next line.

3. The display for each PRODSUP subentry occupies two columns. The first column is the *prodid* and the second column is the *vv.rr.mm*. Four subentries can fit on a row.

UCLIN Examples

You can use the ADD, REP, and DEL UCL statements to change subentries in the PRODUCT entry. When you use UCLIN to update a PRODUCT entry, keep in mind that after the UCLIN changes are made, the PRODUCT entry must contain at least the DESCRIPTION and SREL subentries.

Example: Adding a PRODUCT entry

Assume you have a product, SAMPLE V1R2, for which you want to create a description and put in an entry named "1234-567,1.2.0". You could set up a PRODUCT entry as follows:

```
SET BDY(GLOBAL)                        /* Set to global zone       */.
UCLIN                                  /* Start UCLIN processing   */.
  ADD PRODUCT(1234-567,1.2.0)          /* Identify the product     */
      DESCRIPTION(SAMPLE)              /* - Name                   */
      VENDOR(IBM)                      /* - Vendor                 */
      PRODSUP((1234-456,01.04.00))     /* - Product sups           */
      SREL(Z038)                       /* - SREL                   */
      URL(http://www.ibm.com/)         /* URL                      */.
  ENDUCL                               /* End UCLIN processing     */.
```

PROGRAM entry (distribution and target zone)

The PROGRAM entry describes a program element (a pre-built load module or a program object). Program elements may exist in distribution or target libraries. A PROGRAM entry is created the first time you install a SYSMOD that contains a ++PROGRAM MCS for a program element that does not yet have a PROGRAM entry.

SMP/E records the function and service level of the program element in the entry. Once a PROGRAM entry exists, it is updated as subsequent SYSMODs that affect the program element are installed.

Subentries

These are the subentries for the PROGRAM entry as they appear in the LIST output:

name
is the name of the program element represented by the entry. It can contain from 1 to 8 alphanumeric characters and $, #, @, or hex C0.

ALIAS
specifies a list of alias names for the element.

The UCL operand is ALIAS(*name...*).

Each alias name can contain from 1 to 8 alphanumeric characters.

DISTLIB
specifies the ddname of the distribution library for the program element.

The UCL operand is DISTLIB(*ddname*).

- The ddname can contain from 1 to 8 alphanumeric characters.
- The DISTLIB subentry is required. Without it, SMP/E cannot process any changes for the program element.

FMID
specifies the functional owner of this program element. The functional owner is the last function SYSMOD that replaced this element.

The UCL operand is FMID(*sysmod_id*).

The SYSMOD ID must contain 7 alphanumeric characters.

LASTUPD
identifies the cause of the last change to this PROGRAM entry.

The UCL operand is LASTUPD(*value*). This subentry can contain one of the following values:

UCLIN
indicates that the change was made as a result of UCLIN processing.

sysmod_id
indicates that the change was made during the installation of the indicated SYSMOD.

The SYSMOD ID must contain 7 alphanumeric characters.

LASTUPD TYPE

indicates how the entry was last changed.

The UCL operand is LASTUPDTYPE(*value*). This subentry may contain one of the following values:

ADD

The entry was added.

UPD

The entry was updated.

RMID

identifies the last SYSMOD that **replaced** this program element. Any subsequent SYSMOD that modifies this program element must have a defined relationship (such as PRE or SUP) with this SYSMOD.

The UCL operand is RMID(*sysmod_id*).

- The SYSMOD ID must contain 7 alphanumeric characters.
- If RMID is not specified but FMID is, SMP/E sets the RMID value to the specified FMID.

SYSLIB

specifies the ddname of the target library for the program element.

The UCL operand is SYSLIB(*ddname*).

- You can specify only one SYSLIB value.
- The ddname can contain from 1 to 8 alphanumeric characters.

LIST Examples

To list all the PROGRAM entries in a particular zone, you can use the following commands:

```
SET      BDY(TGT1)             /* Set to requested zone.   */.
LIST     PROGRAM               /* List all PROGRAM entries.*/.
```

To list specific PROGRAM entries, you can use these commands:

```
SET      BDY(TGT1)             /* Set to requested zone.   */.
LIST     PROGRAM(PGM1          /* List only these two      */
             PGM2)             /* entries.                 */.
```

The format of the LIST output for each PROGRAM entry is the same for both of these commands. The only difference is the number of PROGRAM entries listed. is an example of LIST output for PROGRAM entries.

```
PAGE nnnn  - NOW SET TO zzzzzz ZONE nnnnnnn  DATE mm/dd/yy TIME hh:mm:ss SMP/E 36.nn SMPLIST OUTPUT

TGT1           PROGRAM ENTRIES


  NAME

PGM1       LASTUPD        = MYPROG1 TYPE=ADD
           LIBRARIES      = DISTLIB=ADSTLIB    SYSLIB=SPGMLIB
           FMID           = MYPROG1
           RMID           = MYPROG1

PGM2       LASTUPD        = PGMPTF1 TYPE=ADD
           LIBRARIES      = DISTLIB=ADSTLIB    SYSLIB=SPGMLIB
           FMID           = MYPROG1
           RMID           = PGMPTF1
```

Figure 47. PROGRAM entry: sample LIST output

By specifying the FORFMID operand, you can reduce the number of PROGRAM entries listed. When FORFMID is specified, SMP/E lists a PROGRAM entry only if its FMID matches one of the FMIDs specified on the FORFMID operand. For example, to list PROGRAM entries whose FMIDs are defined in FMIDSET ABC or else are MYPROG1, you can use these commands:

```
SET      BDY(TGT1)       /* Set to target zone.    */.
LIST     PROGRAM         /* List all PROGRAM entries */
         FORFMID(ABC     /* for the ABC FMIDSET      */
              MYPROG1)   /* and FMID MYPROG1.        */.
```

You can use the LIST command to find out the names of all SYSMODs that have modified program elements. To include the names of these SYSMODs in the LIST output, you can use the XREF operand, as shown in these commands:

```
SET      BDY(TGT1)       /* Set to requested zone.   */.
LIST     PROGRAM         /* List all PROGRAM entries */
         XREF            /* and related SYSMODs.     */.
```

Note:

1. You can use XREF in either mass mode or select mode.

2. SMP/E obtains the data included for the XREF operand by checking for entries for this program element in all the SYSMOD entries. Because this data is not contained in the PROGRAM entry itself, you cannot use UCLIN to change it in the PROGRAM entry.

Figure 48 on page 257 is an example of the LIST output produced when the XREF operand is used.

```
PAGE nnnn   - NOW SET TO zzzzzz ZONE nnnnnnn  DATE mm/dd/yy TIME hh:mm:ss SMP/E 36.nn SMPLIST OUTPUT

TGT1           PROGRAM ENTRIES


  NAME

PGM1      LASTUPD       = MYPROG1 TYPE=ADD
          LIBRARIES     = DISTLIB=ADSTLIB    SYSLIB=SPGMLIB
          FMID          = MYPROG1
          RMID          = MYPROG1
          SYSMOD HISTORY = SYSMOD   TYPE      DATE    MCS    --STATUS--
                          MYPROG1  FUNCTION 07.100  PROGRAM   APP    ACC

PGM2      LASTUPD       = PGMPTF1 TYPE=ADD
          LIBRARIES     = DISTLIB=ADSTLIB    SYSLIB=SPGMLIB
          FMID          = MYPROG1
          RMID          = PGMPTF1
          SYSMOD HISTORY = SYSMOD   TYPE      DATE    MCS    --STATUS--
                          PGMPTF1  PTF      07.150  PROGRAM   APP    ACC
```

Figure 48. PROGRAM entry: sample LIST output when XREF is specified

UNLOAD Examples

To dump the PROGRAM entries in UCL format, you can use the UNLOAD command. To unload all the PROGRAM entries in a particular zone, you can use the following commands:

```
SET      BDY(TGT1)       /* Set to requested zone.     */.
UNLOAD   PROGRAM         /* Unload all PROGRAM entries.*/.
```

To unload specific PROGRAM entries, you can use these commands:

```
SET      BDY(TGT1)       /* Set to requested zone.    */.
UNLOAD   PROGRAM(PGM1    /* Unload only these two     */
         PGM2)           /* entries.                  */.
```

The format of the UNLOAD output for each PROGRAM entry is the same for both of these commands. The only difference is the number of PROGRAM entries listed. Figure 49 on page 258 is an example of UNLOAD output for PROGRAM entries.

```
UCLIN .
REP       PROGRAM         ( PGM1     )
          LASTUPD         ( MYPROG1  )
          LASTUPDTYPE     ( ADD )
          DISTLIB         ( ADSTLIB  )
          SYSLIB          ( SPGMLIB  )
          FMID            ( MYPROG1 )
          RMID            ( MYPROG1 )
                          .
REP       PROGRAM         ( PGM2     )
          LASTUPD         ( PGMPTF1  )
          LASTUPDTYPE     ( ADD )
          DISTLIB         ( ADSTLIB  )
          SYSLIB          ( SPGMLIB  )
          FMID            ( MYPROG1 )
          RMID            ( PGMPTF1 )
                          .
ENDUCL.
```

Figure 49. PROGRAM entry: sample UNLOAD output

By specifying the FORFMID operand, you can reduce the number of PROGRAM entries unloaded. When FORFMID is specified, SMP/E unloads a PROGRAM entry only if its FMID matches one of the FMIDs specified on the FORFMID operand. For example, to unload PROGRAM entries whose FMIDs are either defined in FMIDSET ABC or are MYPROG1, you can use these commands:

```
SET      BDY(TGT1)        /* Set to target zone.        */.
UNLOAD   PROGRAM          /* Unload all PROGRAM entries */
         FORFMID(ABC      /* for the ABC FMIDSET        */
                MYPROG1)  /* and FMID MYPROG1.          */.
```

UCLIN Examples

You can use the ADD, REP, and DEL UCL statements to change subentries in the PROGRAM entry. After the UCLIN changes are made, the PROGRAM entry must contain at least the following subentries:

- DISTLIB

- FMID

- RMID

Otherwise, there is not enough information in the entry to process the program element. If any of these subentries are missing, SMP/E does not make the requested UCL updates to the entry, and the entry remains as it was before the UCL command.

Example: Adding a new PROGRAM entry

Assume you have installed PGM3 outside of SMP/E, but now you want to start using SMP/E to track changes to that program element. Here is an example of the UCL statements you would use to define entries for that program element in the appropriate target and distribution zones:

```
SET      BDY(TGT1)           /* Set to target zone.     */.
UCLIN                        /*                         */.
ADD      PROGRAM(PGM3)       /* Define new PROGRAM entry.*/
         DISTLIB(ADSTLIB)    /* Define DLIB.            */
         SYSLIB(SPGMLIB)     /* System library.         */
         FMID(MYPROG1)       /* Functional owner.       */.
ENDUCL                       /*                         */.
SET      BDY(DLB1)           /* Now do same to DLIB.    */.
UCLIN                        /*                         */.
ADD      PROGRAM(PGM3)       /* Define new PROGRAM entry.*/
         DISTLIB(ADSTLIB)    /* Define DLIB.            */
                             /* no SYSLIB info in DLIB. */
```

```
        FMID(MYPROG1)      /* Functional owner.      */.
ENDUCL                     /*                        */.
```

SRC entry (distribution and target zone)

The SRC entry describes source that exists in the distribution or target libraries. (SMP/E assumes that for each SRC entry in a particular zone there exists a MOD entry with the same name.) There are two ways a SRC entry can be created:

- **Installing a SYSMOD that contains the source.** SRC entries are created the first time you install a SYSMOD that contains a ++SRC statement for source that does not yet have a SRC entry.
- **Processing JCLIN.** SRC entries can be built during JCLIN processing when SMP/E scans the assembler step and determines that the assembler input is a member of a partitioned data set.

  SRC entries can also be built when SMP/E scans copy steps and finds a SELECT statement that specifies TYPE=SRC.

SMP/E records the function and service level of the source in the SRC entry, as well as information about how that source affects the structure of the distribution or target libraries and modules. Once a SRC entry exists for source, it is updated as subsequent SYSMODs that affect the source are installed.

Subentries

These are the subentries for the SRC entry as they appear in the LIST output:

name
> is the name of the source represented by the SRC entry.
>
> The name can contain from 1 to 8 alphanumeric characters and $, #, @, or hex C0.

DISTLIB
> specifies the ddname of the distribution library for the source.
>
> The UCL operand is DISTLIB(*ddname*).
>
> The ddname can contain from 1 to 8 alphanumeric characters.

FMID
> identifies the functional owner of this source. The functional owner is the last function SYSMOD that replaced this module.
>
> The UCL operand is FMID(*sysmod_id*).
>
> The SYSMOD ID must contain 7 alphanumeric characters.

LASTUPD
> identifies the cause of the last change to this SRC entry.
>
> The UCL operand is LASTUPD(*value*). This subentry can contain one of the following values:
>
> **JCLIN**
>> indicates that the change was made during JCLIN command processing.
>
> **UCLIN**
>> indicates that the change was made as a result of UCLIN processing.
>
> *sysmod_id*
>> indicates that the change was made during the installation of the indicated SYSMOD.
>>
>> The SYSMOD ID must contain 7 alphanumeric characters.

LASTUPD TYPE
> indicates how the entry was last changed.
>
> The UCL operand is LASTUPDTYPE(*value*). This subentry can contain one of the following values:
>
> **ADD**
>> The entry was added.

SRC entry (distribution and target zone)

MOV
>The entry was moved.

UPD
>The entry was updated.

RMID
>identifies the last SYSMOD that **replaced** this source. Any subsequent SYSMOD that modifies this module must have a defined relationship (such as PRE or SUP) with this SYSMOD.
>
>The UCL operand is RMID(*sysmod_id*).
>
>- The SYSMOD ID must contain 7 alphanumeric characters.
>- If RMID is not specified, but FMID is, SMP/E sets the RMID value to the specified FMID.

SYSLIB
>specifies the ddname of the target library for the source.
>
>The UCL operand is SYSLIB(*ddname*).
>
>- Only one SYSLIB value can be specified.
>- The ddname can contain from 1 to 8 alphanumeric characters.

UMID
>identifies all the SYSMODs that have **updated** this source since it was last replaced. Any subsequent SYSMOD that modifies this module must have a defined relationship (such as PRE or SUP) with all these SYSMODs.
>
>The UCL operand is UMID(*sysmod_id*...).
>
>The SYSMOD ID must contain 7 alphanumeric characters.

LIST Examples

To list all the SRC entries in a particular zone, you could use the following commands:

```
SET     BDY(TGT1)           /* Set to requested zone.  */.
LIST    SRC                 /* List all SRC entries.   */.
```

To list specific SRC entries, you could use these commands:

```
SET     BDY(TGT1)           /* Set to requested zone.  */.
LIST    SRC(SRC01           /* List only these two     */
            SRC02)          /* entries.                */.
```

The format of the LIST output for each SRC entry is the same for both of these commands. The only difference is the number of SRC entries listed. is an example of LIST output for SRC entries.

```
PAGE nnnn  - NOW SET TO zzzzzz ZONE nnnnnnn  DATE mm/dd/yy TIME hh:mm:ss SMP/E 36.nn SMPLIST OUTPUT

TGT1           SOURCE ENTRIES

   NAME

SRC01     LASTUPD        = JXY1102 TYPE=ADD
          LIBRARIES      = DISTLIB=ASRCLIB    SYSLIB=SRCLIB
          FMID           = JXY1102
          RMID           = JXY1102

SRC02     LASTUPD        = JXY1000 TYPE=UPD
          LIBRARIES      = DISTLIB=ASRCLIB    SYSLIB=SRCLIB
          FMID           = JXY1121
          RMID           = UZ00010
          UMID           = UZ00014    UZ00015
```

Figure 50. SRC entry: sample LIST output
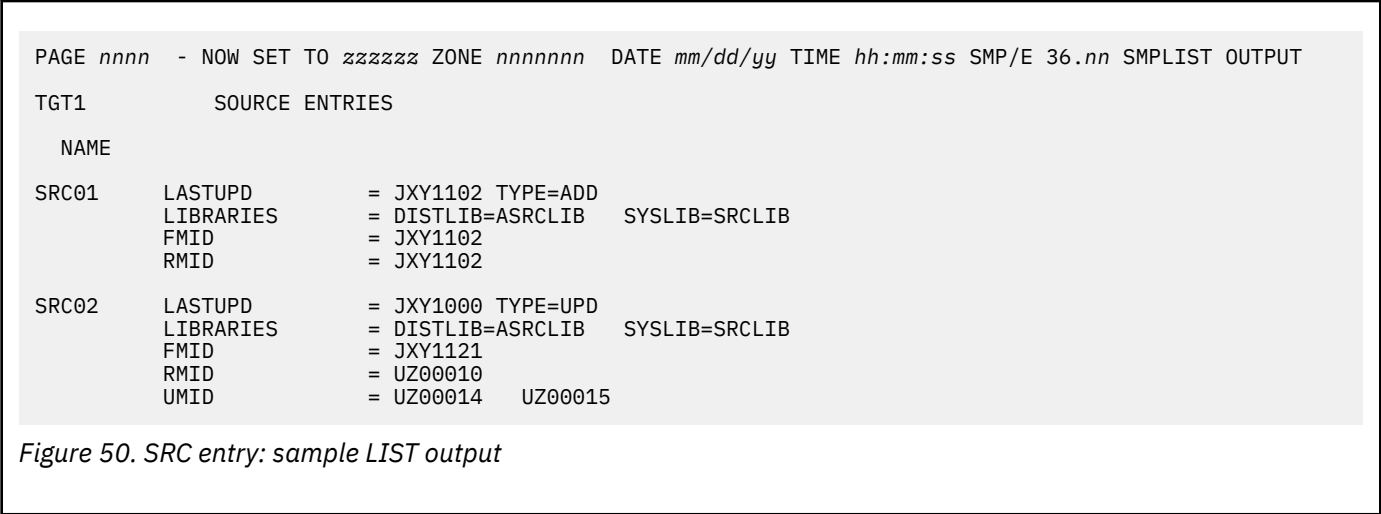
By specifying the FORFMID operand, you can reduce the number of SRC entries listed. When FORFMID is specified, SMP/E lists a SRC entry only if its FMID matches one of the FMIDs specified on the FORFMID operand. For example, to list SRC entries whose FMIDs either are defined in FMIDSET TP or are JXY1102, you could use these commands:

```
SET       BDY(TGT1)          /* Set to target zone.    */.
LIST      SRC                /* List all source entries */
          FORFMID(TP         /* for the TP FMIDSET       */
                 JXY1102)    /* and FMID JXY1102.       */.
```

You can also use the LIST command to find out the name of every SYSMOD that has modified source. To include the names of these SYSMODs in the LIST output, you can use the XREF operand, as shown in these commands:

```
SET       BDY(TGT1)          /* Set to requested zone.  */.
LIST      SRC                /* List all source entries */
          XREF               /* and related SYSMODs.    */.
```

Note:

1. XREF can be used either in mass mode or in select mode.

2. SMP/E obtains the data included for the XREF operand by checking for SRC and SRCUPD entries for this module in all the SYSMOD entries. Because this data is not contained in the SRC entry itself, you cannot use UCLIN to change it in the SRC entry.

Figure 51 on page 261 is an example of the LIST output produced when the XREF operand is used.

```
PAGE nnnn  - NOW SET TO zzzzzz ZONE nnnnnnn  DATE mm/dd/yy TIME hh:mm:ss SMP/E 36.nn SMPLIST OUTPUT

TGT1           SOURCE ENTRIES


  NAME

SRC01     LASTUPD        = JXY1102 TYPE=ADD
          LIBRARIES      = DISTLIB=ASRCLIB    SYSLIB=SRCLIB
          FMID           = JXY1102
          RMID           = JXY1102
          SYSMOD HISTORY = SYSMOD   TYPE      DATE   MCS     --STATUS--
                           JXY1102  FUNCTION 07.100  SRC      APP   ACC

SRC02     LASTUPD        = JXY1000 TYPE=UPD
          LIBRARIES      = DISTLIB=ASRCLIB    SYSLIB=SRCLIB
          FMID           = JXY1121
          RMID           = UZ00010
          UMID           = UZ00014   UZ00015
          SYSMOD HISTORY = SYSMOD   TYPE      DATE   MCS     --STATUS--
                           JXY1102  FUNCTION 07.100  SRC      APP   ACC
                           JXY1121  FUNCTION 07.150  SRC      APP   ACC
                           UZ00010  PTF      07.150  SRC      APP
                           UZ00014  PTF      07.160  SRCUPD   APP
                           UZ00015  PTF      07.161  SRCUPD   APP
```

Figure 51. SRC entry: sample LIST output when XREF is specified

UNLOAD Examples

To dump the SRC entries in UCL format, you can use the UNLOAD command. To unload all the SRC entries in a particular zone, you could use the following commands:

```
SET       BDY(TGT1)          /* Set to requested zone.  */.
UNLOAD    SRC                /* Unload all SRC entries. */.
```

To unload specific SRC entries, you could use these commands:

```
SET      BDY(TGT1)          /* Set to requested zone.   */.
UNLOAD   SRC(SRC01          /* Unload only these two    */
            SRC02)          /* entries.                 */.
```

The format of the UNLOAD output for each SRC entry is the same for both of these commands. The only difference is the number of SRC entries listed. Figure 52 on page 262 is an example of UNLOAD output for SRC entries.

```
UCLIN .
REP        SRC             ( SRC01     )
           LASTUPD         ( JXY1102   )
           LASTUPDTYPE     ( ADD )
           DISTLIB         ( ASRCLIB   )
           SYSLIB          ( SRCLIB    )
           FMID            ( JXY1102 )
           RMID            ( JXY1102 )
                            .
REP        SRC             ( SRC02     )
           LASTUPD         ( JXY1121   )
           LASTUPDTYPE     ( UPD )
           DISTLIB         ( ASRCLIB   )
           SYSLIB          ( SRCLIB    )
           FMID            ( JXY1121 )
           RMID            ( UZ00010 )
           UMID            ( UZ00014    UZ00015   )
                            .
ENDUCL.
```

Figure 52. SRC entry: sample UNLOAD output

By specifying the FORFMID operand, you can reduce the number of SRC entries unloaded. When FORFMID is specified, SMP/E unloads a SRC entry only if its FMID matches one of the FMIDs specified on the FORFMID operand. For example, to unload SRC entries whose FMIDs either are defined in FMIDSET TP or are JXY1102, you could use these commands:

```
SET      BDY(TGT1)          /* Set to target zone.      */.
UNLOAD   SRC                /* Unload all source entries*/
         FORFMID(TP         /* for the TP FMIDSET       */
                 JXY1102)   /* and FMID JXY1102.        */.
```

UCLIN Examples

You can use the ADD, REP, and DEL UCL statements to change subentries in the SRC entry. After the UCLIN changes are done, the SRC entry must contain at least the following subentries:

- DISTLIB
- FMID
- RMID

Otherwise, there is not enough information in the entry to process the macro. If any of these subentries are missing, SMP/E does not make the requested UCL updates to the entry, and the entry remains as it was before the UCL command.

The following examples are provided to help you use the SRC entry.

Example 1: Adding a new SRC entry

Assume that you have a user application installed and want to support it with SMP/E. To do this, you must record the modules in the target and distribution zones. You should identify the MOD and LMOD entries for your application by running the JCLIN function, using the appropriate link and copy steps as input. The source elements in your application will have to be entered through UCLIN. Assume that the FMID you want to assign the product is ZUSR001 and that no service level is to be recorded. The following UCL statements should be used for each such source:

```
SET       BDY(TGT1)          /* Set to target zone.        */.
UCLIN                        /*                            */.
ADD       SRC(SRC01)         /* Define new source entry.   */
          DISTLIB(ASRCLIB)   /* Define DLIB,               */
          SYSLIB(SRCLIB)     /* system library
                                (never the SMPSTS or STS). */
          FMID(ZUSR001)      /* Functional owner (in this
                                example a user function).  */
                             /*                            */.
ENDUCL                       /*                            */.
SET       BDY(DLB1)          /* Now do same to DLIB.       */.
UCLIN                        /*                            */.
ADD       SRC(SRC01)         /* Define new source entry.   */
          DISTLIB(ASRCLIB)   /* Define DLIB,               */
          SYSLIB(SRCLIB)     /* system library
                                (never the SMPSTS or STS). */
          FMID(ZUSR001)      /* Functional owner (in this
                                example a user function).  */.
```

Example 2: Recording the application of a corrective fix

Assume that you have installed a corrective fix, AZ12345, to source SRC01, outside of SMP/E, and that now you want to record that fix after it has been installed. (This is not necessary if the fix was initially installed with SMP/E.) Three updates must be made:

1. Record that the fix, AZ12345, is in the system.

2. Record that the fix is on the source.

3. Record that the source has been reassembled and installed.

The following UCL can be used to record these changes:

```
SET       BDY(TGT1)          /* Set to target zone.      */.
UCLIN                        /*                          */.
ADD       SYSMOD(AZ12345)    /* Add SYSMOD entry.        */
          APAR               /* Corrective fix.          */
          APPDATE(100)     /* Date applied.          */
          APPTIME(08:00:00)  /* Time applied.            */
          FMID(FXY1102)      /* Functional owner.        */
          SRCUPD(SRC01)      /* Updated SRC01.           */
          ASSEM(SRC01)       /* Assembled it too.        */
                             /*                          */.
REP       SRC(SRC01)         /* Update SRC01 to          */
          UMID(AZ12345)      /* add update ID of APAR.   */
                             /*                          */.
REP       MOD(SRC01)         /* Update MOD entry         */
          RMID(AZ12345)      /* with new replacement ID. */
          RMIDASM            /* Was assembled.           */
                             /*                          */.
ENDUCL                       /*                          */.
SET       BDY(DLIB1)         /* Set to DLIB zone.        */.
UCLIN                        /*                          */.
ADD       SYSMOD(AZ12345)    /* Add SYSMOD entry.        */
          APAR               /* Corrective fix.          */
          RECDATE(100)     /* Date received.            */
          RECTIME(08:00:00)  /* Time received.           */
          INSDATE(100)     /* Date accepted.            */
          INSTIME(09:00:00)  /* Time accepted.           */
          FMID(FXY1102)      /* Functional owner.        */
          SRCUPD(SRC01)      /* Updated SRC01.           */
          ASSEM(SRC01)       /* Assembled it too.        */
                             /*                          */.
REP       SRC(SRC01)         /* Update SRC01 to          */
          UMID(AZ12345)      /* add update ID of APAR.   */
                             /*                          */.
REP       MOD(SRC01)         /* Update MOD entry         */
          RMID(AZ12345)      /* with new replacement ID. */
          RMIDASM            /* Was assembled.           */
                             /*                          */.
ENDUCL                       /*                          */.
```

STSSRC entry (SMPSTS)

The STSSRC entry is a copy of source that resides only in a distribution library but is needed temporarily during APPLY processing. The STSSRC entry is in the SMPSTS data set, which serves as a target source library for such modules.

When SMP/E applies the SYSMODs that affect these source, it calls utility programs to store the modules on the SMPSTS. This way, the most current service level of each module is available for use in assemblies. After SMP/E has accepted all the SYSMODs that affect these source modules, it deletes the associated STSSRC entries from the SMPSTS.

Note: If you specify SAVESTS in the ORDER entry that is in effect during ACCEPT processing, SMP/E will not delete STSSRC entries from the SMPSTS after the SYSMODs that affect those source have been successfully accepted.

Subentries

The STSSRC entry contains no SMP/E data and appears to the system as a member of a normal source library.

LIST Examples

You cannot use SMP/E to list the STSSRC entries. However, you can use standard system utility programs (such as IEBGENER, IEBPTPCH, and IEHLIST) or products such as ISPF to display these entries or information about the data set.

UCLIN Examples

You can use the DEL UCL statement to delete an STSSRC entry from the SMPSTS. This can be helpful if you plan to do an APPLY followed by ACCEPT when several target libraries have been created from the same distribution library.

When a SYSMOD is accepted into a distribution zone, the entries associated with it are automatically deleted from the SMPSTS for the related target zone. However, even if the SYSMOD was also applied to other target zones created from the same distribution zone, SMP/E does not clean up the SMPSTS data sets for the other target zones.

To delete the entries from these data sets, you could accept the SYSMOD and name these other target zones as the related zone. However, this would update the distribution library each time; this is time-consuming and could use up space in the distribution library data set.

Instead, you can use the DEL command to delete these entries without updating the distribution library. To determine which entries to specify, check the SMPLOG data set to see which ones SMP/E deleted during ACCEPT processing.

Note: You can also use the CLEANUP command to delete STSSRC entries without specifying them individually. For more information, see the CLEANUP command in *z/OS SMP/E Commands*.

Example: Deleting an STSSRC entry

Assume that you have two target zones, TGT1 and TGT2, generated off the same distribution zone, DLB1. During ACCEPT processing of a SYSMOD, SMP/E has deleted STSSRC SRC01 and SRC02 from the SMPSTS data set associated with target zone TGT2. After performing the ACCEPT, you want to delete the same source from the SMPSTS associated with target zone TGT1. Assume either that you have a cataloged procedure for TGT1 with the correct SMPSTS specified, or that you have set up the correct DDDEF entries. You can use the following UCLIN to delete the STSSRC entry:

```
SET     BDY(TGT1)          /* Set to TGT1 zone.      */.
UCLIN                      /*                        */.
DEL     STSSRC(SRC01)      /* Delete the source.     */.
DEL     STSSRC(SRC02)      /* Delete the source.     */.
ENDUCL                     /*                        */.
```

Note: One UCL statement is required for each STSSRC entry to be deleted.

You can make the same changes by using system utilities; however, the SMPLOG will not reflect the processing done.

SYSMOD entry (distribution and target zone)

The SYSMOD entry in a distribution zone or a target zone describes a SYSMOD that has been installed in the corresponding distribution library or target library. This SYSMOD entry contains the same information as the global zone SYSMOD entry, except that it has information from only one ++VER statement, the one to install the SYSMOD.

When SMP/E installs a SYSMOD, it uses SYSMOD entries in the distribution or target zone to do the following:

- Determine the functional level of the system. SMP/E checks which function SYSMODs have been installed, and then uses that information to determine which service SYSMODs may be applicable.

- Determine the service level of the system. SMP/E checks which service SYSMODs have been installed.

- Make sure the requisites are satisfied for each SYSMOD to be installed. SMP/E checks whether a SYSMOD entry exists in the distribution or target zone for each requisite.

Subentries

These are the subentries for the SYSMOD entry as they appear in the LIST output:

sysmod_id
    is the SYSMOD identification.

    The SYSMOD ID must contain 7 alphanumeric characters.

ACCEPT
    indicates that the SYSMOD has been successfully accepted.

    The UCL operand is ACCEPT, ACPT, or ACC.

    Note: This subentry exists only in the distribution zone. It is required in distribution zone SYSMOD entries.

APAR
    indicates that this SYSMOD is an APAR, which provides a corrective fix to a problem.

    The UCL operand is APAR.

    APAR, FUNCTION, PTF, and USERMOD are mutually exclusive. If none of these operands is specified, PTF is the default.

APPLY
    indicates that the SYSMOD has been successfully applied.

    The UCL operand is APPLY, APPL, or APP.

    Note: This subentry exists only in the target zone. It is required in target zone SYSMOD entries.

ASSEM
    lists the assemblies done during the installation of this SYSMOD.

    The UCL operand is ASSEM(*name...*).

    The name can contain from 1 to 8 alphanumeric characters.

BYPASS
    indicates that the BYPASS operand was specified when this SYSMOD was installed.

    The UCL operand is BYPASS.

CIFREQ
    lists the conditional requisites that must be installed when this function SYSMOD is installed.

Note: The data specified is used by SMP/E only when present in a function SYSMOD.

The UCL operand is CIFREQ((*causer,req*)...).

- *causer* is the SYSMOD that specified this function SYSMOD on the FMID operand of an ++IF statement. *req* is the SYSMOD specified on the REQ operand as the conditional requisite associated with this function SYSMOD.
- The *causer* and *req* fields must contain 7 alphanumeric characters each.
- The CIFREQ operand is mutually exclusive with all other UCL operands. It will not cause LASTUPD and LASTUPDTYPE to be updated.

DELBY

specifies the SYSMOD that deleted this SYSMOD.

The UCL operand is DELBY(*sysmod_id*).

- The SYSMOD ID must contain 7 alphanumeric characters.
- This subentry is valid only for function SYSMODs.
- The DELBY operand is mutually exclusive with all other UCL operands.

DELETE

lists the SYSMODs deleted by this SYSMOD.

The UCL operand is DELETE(*sysmod_id*...).

The SYSMOD ID must contain 7 alphanumeric characters.

DELLMOD

indicates that the SYSMOD contained a ++DELETE statement.

The UCL operand is DELLMOD.

DESCRIPTION

specifies the descriptive name to be associated with this SYSMOD.

The UCL operand is DESCRIPTION(*description*).

DLMOD

lists the load modules deleted by ++DELETE statements contained in this SYSMOD.

The UCL operand is DLMOD(*name*...).

element

lists the data element replacements contained in the SYSMOD.

The UCL operand is *element*(*name*...).

- The name can contain from 1 to 8 alphanumeric characters.
- In place of *element*, specify one of the replacement values shown in Table 2 on page 9.
- Some types of elements, such as panels, messages, or text, may have been translated into several languages. In these cases, the *element* operand contains *xxx*, which represents the language used for the element. (If an element was not translated, the *element* operand does not contain any *xxx* value.) Table 3 on page 11 shows the *xxx* values and the languages they represent.

ELEMMOV

indicates that the SYSMOD contained a ++MOVE statement.

The UCL operand is ELEMMOV.

EMOVE

lists the elements and load modules that were moved by ++MOVE statements contained in this SYSMOD.

The UCL operand is EMOVE(*name*...).

ERROR

indicates that an error has occurred during the processing of this SYSMOD.

The UCL operand is ERROR.

- This operand can also be specified as ERR.
- If the RESTORE subentry is set, the RESDATE operand **must** be specified, and the RESTIME and ERROR operands **should** be specified (SMP/E will automatically set it otherwise).

FEATURE
lists the names of the software features that contain this SYSMOD. The identified software features correspond to FEATURE entries in the global zone.

The UCL operand is FEATURE(*name*,...).

FESN
identifies the field engineering (FE) service number.

The UCL operand is FESN(*string*).

The string must contain 7 alphanumeric characters.

FMID
identifies the function SYSMOD to which this SYSMOD is applicable.

The UCL operand is FMID(*sysmod_id*).

The SYSMOD ID must contain 7 alphanumeric characters.

FUNCTION
indicates that this SYSMOD is a function, which introduces a new function into the system.

The UCL operand is FUNCTION.

APAR, FUNCTION, PTF, and USERMOD are mutually exclusive. If none of these operands is specified, PTF is the default.

hfs_element
lists the hierarchical file system element replacements in the SYSMOD.

The UCL operand is *hfs-element*(*name*...).

The name can contain 1 to 8 uppercase alphabetic, numeric, or national ($, #, @) characters.

IFREQ
lists the conditional requisites that were installed with this SYSMOD.

The UCL operand is IFREQ(*sysmod_id*...).

INSTALLDATE
specifies the date on which this SYSMOD was installed.

The UCL operand is INSTALLDATE(*yyddd*).

- This operand can also be specified as INSDATE. For a distribution zone SYSMOD entry, it can be specified as ACCDATE. For a target zone SYSMOD entry, it can be specified as APPDATE.
- The *yyddd* must contain 5 numeric characters. SMP/E does not check whether the specified numbers are valid.
- For the distribution zone SYSMOD entry, INSTALLDATE is the date the SYSMOD was accepted. For the target zone SYSMOD entry, INSTALLDATE is the date the SYSMOD was applied.

INSTALLTIME
specifies the time at which this SYSMOD was installed.

The UCL operand is INSTALLTIME(*hh*:*mm*:*ss*).

- This operand can also be specified as INSTIME. For a distribution zone SYSMOD entry, it can be specified as ACCTIME. For a target zone SYSMOD entry, it can be specified as APPTIME.
- The *hh*, *mm*, and *ss* must contain 2 numeric characters each. The ":" must be coded as specified. SMP/E does not check whether the specified numbers are valid.

- For a distribution zone SYSMOD entry, INSTALLTIME is the time the SYSMOD was accepted. For a target zone SYSMOD entry, INSTALLTIME is the time the SYSMOD was applied.

JAR

lists the JAR file replacements (++JAR) supplied by the SYSMOD.

The UCL operand is JAR.

JARUPD

lists the JAR file updates (++JARUPD) supplied by the SYSMOD.

The UCL operand is JARUPD.

JCLIN

indicates that the SYSMOD contained inline JCLIN.

The UCL operand is JCLIN.

LASTSUP

specifies the most recent SYSMOD that superseded this SYSMOD. All previous superseding SYSMODs are saved in the SUPBY subentry list.

The UCL operand is LASTSUP(*sysmod_id*).

The SYSMOD ID must contain 7 alphanumeric characters.

LASTUPD

identifies the cause of the last change to this entry.

The UCL operand is LASTUPD(UCLIN), indicating that the change was made as a result of UCLIN processing.

LASTUPDTYPE

identifies the last type of update made to this entry.

The UCL operand is LASTUPDTYPE(*value*). This subentry can contain one of the following values:

ADD

The entry was added.

UPD

The entry was updated.

MAC

lists the macro replacements (++MAC statements) in the SYSMOD.

The UCL operand is MAC(*name...*).

The name can contain from 1 to 8 alphanumeric characters.

MACUPD

lists the macro updates (++MACUPD statements) in the SYSMOD.

The UCL operand is MACUPD(*name...*).

The name can contain from 1 to 8 alphanumeric characters.

MOD

lists the module replacements (++MOD statements) in the SYSMOD.

The UCL operand is MOD(*name...*).

The name can contain from 1 to 8 alphanumeric characters.

NPRE

lists negative prerequisite SYSMODs (that is, SYSMODs that must not be present in the system at the same time as this SYSMOD).

The UCL operand is NPRE(*sysmod_id...*).

- The SYSMOD ID must contain 7 alphanumeric characters.

- This subentry is valid only for function SYSMODs.

PRE
lists the prerequisite SYSMODs (that is, SYSMODs that must be present before this SYSMOD can be installed).

The UCL operand is PRE(*sysmod_id...*).

The SYSMOD ID must contain 7 alphanumeric characters.

PROGRAM
lists the program element replacements (++PROGRAM statements) in the SYSMOD.

The UCL operand is PROGRAM(*name...*).

The name can contain from 1 to 8 alphanumeric characters.

PTF
indicates that this SYSMOD is a PTF, which provides preventive service fixes.

The UCL operand is PTF.

APAR, FUNCTION, PTF, and USERMOD are mutually exclusive. If none of these operands is specified, PTF is the default.

RECDATE
specifies the date on which this SYSMOD was received.

The UCL operand is RECDATE(*yyddd*).

The *yyddd* must contain 5 numeric characters. SMP/E does not check whether the specified numbers are valid.

RECTIME
specifies the time at which this SYSMOD was received.

The UCL operand is RECTIME(*hh*:*mm*:*ss*).

The *hh*, *mm*, and *ss* must contain 2 numeric characters each. The ":" must be coded as specified. SMP/E does not check whether the specified numbers are valid.

REGEN
indicates how the SYSMOD was installed in the target libraries.

- In a DLIB zone SYSMOD entry, REGEN is not important. It is automatically set for all SYSMODs when they are accepted.

- In a target zone SYSMOD entry, if REGEN is set, it indicates that SYSGEN was used to install the SYSMOD in the target libraries. (The REGEN indicator was carried over when the distribution zone was copied into the target zone. This is generally done as part of SYSGEN.)

  If REGEN is not set, it indicates that the APPLY command was used to install the SYSMOD in the target libraries.

The UCL operand is REGEN or RGN.

RENLMOD
indicates that the SYSMOD contained a ++RENAME statement.

The UCL operand is RENLMOD.

REQ
lists requisite SYSMODs (that is, SYSMODs that must be installed concurrent with this SYSMOD).

The UCL operand is REQ(*sysmod_id...*).

The SYSMOD ID must contain 7 alphanumeric characters.

RESDATE
specifies the date on which this SYSMOD was restored.

The UCL operand is RESDATE(*yyddd*).

- This subentry exists only in the target zone.
- The *yyddd* must contain 5 numeric characters. SMP/E does not check whether the specified numbers are valid.
- If a SYSMOD is marked "RESTORE", the RESDATE operand **must** be specified, and the RESTIME and ERROR operands **should** be specified (SMP/E will automatically set it otherwise).

RESTIME
specifies the time that this SYSMOD was restored.

The UCL operand is RESTIME(*hh*:*mm*:*ss*).

- This subentry exists only in the target zone.
- The *hh*, *mm*, and *ss* must contain 2 numeric characters each. The ":" must be coded as specified. SMP/E does not check whether the specified numbers are valid.
- If a SYSMOD is marked "RESTORE", the RESDATE operand **must** be specified, and the RESTIME and ERROR operands **should** be specified (SMP/E will automatically set it otherwise).

RESTORE
indicates that a RESTORE attempt has been made for this SYSMOD. The RESTORE was not successful; otherwise, the SYSMOD entry would have been deleted.

The UCL operand is RESTORE.

- This subentry exists only in the target zone.
- This operand can also be specified as REST or RES.
- If a SYSMOD is marked "RESTORE", the RESDATE operand **must** be specified, and the RESTIME and ERROR operands **should** be specified (SMP/E will automatically set it otherwise).

REWORK
identifies the level of the SYSMOD, which was received again for minor changes.

The UCL operand is REWORK(*level*).

- Up to 8 numeric characters can be specified.
- For SYSMODs supplied by IBM, the REWORK level is *yyyyddd*, where *yyyy* is the year the SYSMOD was reworked and *ddd* is the Julian date.
- SMP/E does not check whether this data is valid.

RLMOD
indicates the load modules renamed by ++RENAME statements in this SYSMOD.

The UCL operand is RLMOD(*name...*).

SOURCEID
lists the character strings assigned to this SYSMOD during RECEIVE. These values might have been specified by the user on the RECEIVE command, included inline on the ++ASSIGN statement, or assigned by SMP/E when processing the contents of an order.

The UCL operand is SOURCEID(*source_id...*).

- The source ID can be from 1 to 64 characters in length and contain any nonblank character (X'41' through X'FE') except single quotation mark ('), asterisk (*), percent (%), comma (,), left parenthesis ((), and right parenthesis ()).
- A source ID value cannot span lines.
- A source ID value is case sensitive.

SRC
lists the source replacements (++SRC statements) in the SYSMOD.

The UCL operand is SRC(*name...*).

The name can contain from 1 to 8 alphanumeric characters.

SRCUPD
lists the source updates (++SRCUPD statements) in the SYSMOD.

The UCL operand is SRCUPD(*name...*).

The name can contain from 1 to 8 alphanumeric characters.

SUPBY
lists the SYSMODs that superseded this SYSMOD. For functions, this includes SYSMODs that both deleted and superseded this SYSMOD.

Note:

1. The most recent SYSMOD to supersede this SYSMOD is not included in the SUPBY list. It is saved in the LASTSUP field.

2. The SUPBY field may appear as "SUPBY(IN SYSMD)". For example, this is the case if the superseding SYSMODs were installed separately instead of on the same APPLY or ACCEPT command.

The UCL operand is SUPBY(*sysmod_id...*).

- This operand can also be specified as SUP.
- The SYSMOD ID must contain 7 alphanumeric characters.

SUPING
lists the SYSMODs superseded by this SYSMOD.

The UCL operand is SUPING(*sysmod_id...*).

The SYSMOD ID must contain 7 alphanumeric characters.

SZAP
lists the module superzaps (++ZAP statements) in the SYSMOD.

The UCL operand is SZAP(*name...*).

The name can contain from 1 to 8 alphanumeric characters.

UCLDATE
specifies the date on which this SYSMOD was last modified through UCLIN.

The UCL operand is UCLDATE(*yyddd*).

The *yyddd* must contain 5 numeric characters. SMP/E does not check whether the specified numbers are valid.

UCLTIME
specifies the time that this SYSMOD was last modified through UCLIN.

The UCL operand is UCLTIME(*hh*:*mm*:*ss*).

The *hh*, *mm*, and *ss* must contain 2 numeric characters each. The ":" must be coded as specified. SMP/E does not check whether the specified numbers are valid.

USERMOD
indicates that this SYSMOD is a USERMOD, which puts a user modification in the system.

The UCL operand is USERMOD.

APAR, FUNCTION, PTF, and USERMOD are mutually exclusive. If none of these operands is specified, PTF is the default.

VERNUM
specifies the relative number of the ++VER statement used when this SYSMOD was installed.

The UCL operand is VERNUM(*nnn*).

- *nnn* can contain from 1 to 3 numeric characters.

SYSMOD entry (distribution and target zone)

- When updating an existing entry, you should not specify VERNUM. This causes SMP/E to assume the same VERNUM value as in the current entry.
- If you do not specify VERNUM when adding a new entry, SMP/E assumes a VERNUM value of 1.
- The VERNUM values are kept in each entry built from information from the ++VER statements. (For example, subentries such as PRE and REQ have unique VERNUM values.) If all entries do not have the same VERNUM, an error will result.

VERSION

lists the function SYSMODs that are versioned by this SYSMOD. Versioning indicates that, if there are any elements in common between the SYSMODs listed and this SYSMOD, this SYSMOD's elements are at a higher functional level, and are thus the ones that should be installed.

The UCL operand is VERSION(*sysmod_id*...).

The SYSMOD ID must contain 7 alphanumeric characters.

XZAP

lists the module superzaps in the SYSMOD (++ZAP statements) that contain an EXPAND statement (indicating that the module should be expanded before it is updated).

The UCL operand is XZAP(*name*...).

The name can contain from 1 to 8 alphanumeric characters.

LIST Examples

To list all the SYSMOD entries in a particular zone, you could use the following commands:

```
SET     BDY(TGT1)           /* Set to requested zone.   */.
LIST    SYSMOD              /* List all SYSMOD entries. */.
```

To list specific SYSMOD entries, you could use these commands:

```
SET     BDY(TGT1)           /* Set to requested zone.   */.
LIST    SYSMOD(UZ00001      /* List only these two      */
               UZ00002)     /* entries.                 */.
```

The format of the LIST output for each SYSMOD entry is the same for both of these commands. The only difference is the number of SYSMOD entries listed. and are examples of LIST output for SYSMOD entries.

```
PAGE nnnn  - NOW SET TO zzzzzz ZONE nnnnnnn  DATE mm/dd/yy TIME hh:mm:ss SMP/E 36.nn SMPLIST OUTPUT

DLIB1        SYSMOD ENTRIES


   NAME

HZY1102  TYPE           = DELETED
         DELBY          = HZY2102

HZY2102  TYPE           = FUNCTION
         DESCRIPTION    = Sample function
         FEATURE        = SAMPRG1
         STATUS         = REC  ACC
         FMID           = HZY2102
         JCLIN          = YES
         FESN           = 1234567
         DATE/TIME REC  = 07.100  08:00:00
                   INS  = 07.102  08:08:00
         DELETE VER(001) = HZY1102   JZY1121   JZY1122
         NPRE   VER(001) = HZZ1102
         SUPING VER(001) = AZ11111   AZ11112   AZ11113
         VERS   VER(001) = HYY1102   JYY1121   JYY1122   JYY1123
         MAC            = MAC01     MAC02     MAC03
         MOD            = MOD01     MOD02     MOD03     MOD04
         SRC            = SRC01     SRC02

HZY2121  TYPE           = FUNCTION
         DESCRIPTION    = Sample function 2
         STATUS         = REC  ACC
         FMID           = HZY2121
         JCLIN          = YES
         FESN           = 1234567
         DATE/TIME REC  = 06.100  08:30:00
                   INS  = 06.108  08:38:00
         DELETE VER(001) = HZY1102   JZY1121   JZY1122
         FMID   VER(001) = HZY2102
         SUPING VER(001) = AZ11121   AZ11122   AZ11123
         MAC            = MAC01
         MOD            = MOD01     MOD02
         SRC            = SRC01     SRC03     SRC04

JZY1121  TYPE           = DELETED
         DELBY          = HZY2102

JZY1122  TYPE           = DELETED
         DELBY          = HZY2102
```

Figure 53. SYSMOD entry: sample LIST output for a distribution zone

SYSMOD entry (distribution and target zone)

```
PAGE nnnn  - NOW SET TO zzzzzz ZONE nnnnnnn  DATE mm/dd/yy TIME hh:mm:ss SMP/E 36.nn SMPLIST OUTPUT

TGT1         SYSMOD ENTRIES


   NAME

AZ99001   TYPE           = APAR
          DESCRIPTION    = Sample APAR
          STATUS         = REC  APP
          FMID           = HZY2102
          DATE/TIME REC  = 07.100  08:00:00
                    INS  = 07.100  08:08:00
          LASTSUP        = UZ00010
          ZAP            = MOD01

AZ99002   TYPE           = SUPERSEDED
          LASTSUP        = UZ00010

LUS0001   TYPE           = USERMOD
          DESCRIPTION    = Sample user modification
          STATUS         = REC  APP
          FMID           = HZY2102
          DATE/TIME REC  = 06.100  08:00:00
                    INS  = 06.100  08:09:00
          PRE   VER(001) = UZ00010
          MACUPD         = MAC02

UZ00008   TYPE           = PTF
          DESCRIPTION    = Sample PTF
          STATUS         = REC  APP
          FMID           = HZY2102
          DATE/TIME REC  = 06.100  08:00:00
                    INS  = 06.100  08:08:00
          LASTSUP        = UZ00010
          SUPBY(IN SYSMD) = UZ00009

UZ00009   TYPE           = SUPERSEDED
          LASTSUP        = UZ00010

UZ00010   TYPE           = PTF
          STATUS         = REC  APP
          FMID           = HZY2102
          DATE/TIME REC  = 06.100  08:00:00
                    INS  = 06.100  08:10:00
          PRE   VER(001) = UZ00008    UZ00007
          REQ   VER(001) = UZ00040
          SUPING VER(001) = AZ99001    AZ99002    UZ00009
          MAC            = MAC01
          MACUPD         = MAC02
          MOD            = MOD01
          SRCUPD         = SRC01

UZ00011   TYPE           = PTF
          STATUS         = REC  APP
          FMID           = HXY2102
          DATE/TIME REC  = 07.250  08:00:00
                    INS  = 07.250  08:10:00
          SOURCEID       = ABC0706    PUT0704    XAU3380
          PRE   VER(001) = UZ00010
          MOD            = MOD01
```

Figure 54. SYSMOD entry: sample LIST output for a target zone

By specifying various operands, you can reduce the number of SYSMOD entries listed. If you specify any of these operands on the LIST command, SMP/E automatically assumes that SYSMOD entries are to be processed, regardless of whether the SYSMOD operand was also specified. If you specify more than one of these operands on the same LIST command, only the SYSMODs that meet all the specified conditions are processed. For more information about these operands, see *z/OS SMP/E Commands*.

You can also use the LIST command to find any other SYSMODs that specify this SYSMOD in their DEL, PRE, REQ, NPRE, SUP, or VERSION lists. To include the names of these SYSMODs in the LIST output, you can use the XREF operand, as shown in these commands:

```
SET     BDY(DLIB1)          /* Set to requested zone.    */.
LIST    SYSMOD              /* List all SYSMOD entries   */
        XREF                /* and SYSMOD that hit them. */.
```

Note: XREF can be used either in mass mode or in select mode.

is an example of the LIST output produced when the XREF operand is used.

```
PAGE nnnn  - NOW SET TO zzzzzz ZONE nnnnnnn  DATE mm/dd/yy TIME hh:mm:ss SMP/E 36.nn SMPLIST OUTPUT

DLIB1        SYSMOD ENTRIES


  NAME

UZ12345   TYPE          = PTF
          STATUS        = REC
          FMID          = HZY2100
          DATE/TIME REC = 06.150  08:00:00
                   INS  = 06.160  08:08:00
          SUPING VER(001) = AZ11111   AZ11112   AZ11113
          MAC           = MAC01     MAC02     MAC03
          MOD           = MOD01     MOD02     MOD03     MOD04
          SRC           = SRC01     SRC02
          NPREBY    (XREF) = UZ00001   UZ00002
          PREBY     (XREF) = UZ00003   UZ00004
          REQBY     (XREF) = UZ00005   UZ00006
          VERSIONBY (XREF) = UZ00007   UZ00008
          DELBY     (XREF) = UZ00009   UZ00010
          IFREQBY   (XREF) = UZ00011   UZ00012
          SUPBY     (XREF) = UZ00013   UZ00014
```

Figure 55. SYSMOD entry: sample LIST output when XREF is specified

Note: Some of the *xxx*BY subentries listed in this example are not actually valid for the type of SYSMOD shown (that is, PTF). For example, a SYSMOD does not specify VERSION, NPRE, or DELETE for a PTF SYSMOD. Those subentries are included here only to show the format of the output.

UNLOAD Examples

To dump the SYSMOD entries in UCL format, you can use the UNLOAD command. To unload all the SYSMOD entries in a particular zone, you could use the following commands:

```
SET     BDY(TGT1)           /* Set to requested zone.      */.
UNLOAD  SYSMOD              /* Unload all SYSMOD entries. */.
```

To unload specific SYSMOD entries, you could use these commands:

```
SET     BDY(TGT1)           /* Set to requested zone.    */.
UNLOAD  SYSMOD(UZ00001      /* Unload only these two     */
               UZ00002)     /* entries.                  */.
```

The format of the UNLOAD output for each SYSMOD entry is the same for both of these commands. The only difference is the number of SYSMOD entries listed. is an example of UNLOAD output for one SYSMOD.

```
UCLIN .
REP       SYSMOD          ( HZY2102  )
          /* TYPE    */   FUNCTION
          DESCRIPTION     ( Sample Function )
          FEATURE         ( SAMPRG1 )
          /* STATUS   */   ACC
          RECDATE         ( 07100 )
          RECTIME         ( 08:08:00 )
          INSDATE         ( 07101 )
          INSTIME         ( 10:00:34 )
          VERNUM          ( 002 )
          PRE             ( UZ00001   UZ00002   UZ00003 )
          REQ             ( UZ00004   UZ00005   UZ00006 )
          SUPING          ( AZ00001   AZ00002   AZ00003 )
          MOD             ( MOD01     MOD02     )
          MAC             ( MAC01     )
          SRC             ( SRC01     SRC02     SRC03   )
                          .
REP       SYSMOD          ( HZY2102 )
          CIFREQ          (
                            ( UZ00087 UZ00088 )
                            ( UZ00090 UZ00090 )
                          )
                          .
ENDUCL.
```

Figure 56. SYSMOD entry: sample UNLOAD output

By specifying various operands, you can reduce the number of SYSMOD entries unloaded. If you specify any of these operands on the UNLOAD command, SMP/E automatically assumes that SYSMOD entries are to be processed. If you specify more than one of these operands on the same UNLOAD command, only those SYSMODs that meet all the specified conditions are processed. For more information about these operands, see *z/OS SMP/E Commands*.

UCLIN Examples

You can use the ADD, REP, and DEL UCL statements to change subentries in the SYSMOD entry. Generally, after UCLIN changes are done, the SYSMOD entry must contain at least these subentries, unless the entire entry has been deleted:

- ACCEPT or APPLY
- APAR, FUNCTION, PTF, or USERMOD
- INSTALLDATE
- RECDATE
- FMID

Otherwise, there is not enough information in the entry to process the SYSMOD. If any of the required subentries are missing, SMP/E does not make the requested UCL updates to the entry, and the entry remains as it was before the UCL command.

Note: The entry for a deleted SYSMOD can contain just the DELBY subentry. The entry for a superseded SYSMOD can contain just the SUPBY subentry. The entry for a SYSMOD that is both deleted and superseded can contain just the DELBY and SUPBY subentries.

The following examples are provided to help you use the SYSMOD entry.

Example 1: Creating a SYSMOD entry

Assume that you have installed an APAR fix outside of SMP/E and now want to get that fix recorded in the distribution zone. The APAR number was AZ12345, and was an update to a source SRC01. The following UCL should be used:

```
SET       BDY(DLIB1)         /* Set to DLIB zone.        */.
UCLIN                        /*                          */.
```

```
ADD     SYSMOD(AZ12345)    /* Specify SYSMOD.        */
        SRCUPD(SRC01)      /* Changed this source.   */
        RECDATE(100)     /* Received on this date.   */
        INSDATE(100)     /* Accepted same day.       */
        ASSEM(SRC01)       /* Assembled source.      */
        FMID(FZY2102)      /* Functional owner.      */
                           /*                      */.
ADD     SRC(SRC01)         /* Now update source entry */
        UMID(AZ12345)      /* with update ID.        */
                           /*                      */.
REP     MOD(SRC01)         /* Now update MOD         */
        RMID(AZ12345)      /* with new replacement ID */
        RMIDASM            /* from assembly.         */
                           /*                      */.
ENDUCL                     /*                      */.
```

Example 2: Removing the ERROR indicator

Assume that during an APPLY you encountered an error that caused a PTF to be marked as in error. After looking at the output you determine that the PTF actually installed correctly, and rather than reinstall the PTF you decide to make the appropriate changes to make the PTF look installed. The following UCL can be used:

```
SET     BDY(TGT1)          /* Set to target zone.    */.
UCLIN                      /*                      */.
DEL     SYSMOD(UZ12345)    /* Specify SYSMOD.        */
        RESTORE            /* Delete restore.        */
        RESDATE()          /* Delete restore date.   */
        RESTIME()          /* Delete restore time.   */
                           /*                      */.
ADD     SYSMOD(UZ12345)    /* Specify SYSMOD.        */
        APPLY              /* Add apply info.        */
        INSDATE(100)     /*                      */
        INSTIME(08:00:00)  /*                      */
                           /*                      */.
ENDUCL                     /*                      */.
```

Note: This method is very prone to errors. The preceding example gets the SYSMOD marked as having been applied; however, you have not made all the changes necessary to get the rest of the target zone entries coordinated. Those changes include updating the RMID and UMID fields of all the elements affected by the PTF, storing superseded SYSMOD entries, and updating the global zone SYSMOD entry.

If you are not extremely familiar with SMP/E internals and how to complete the update process, the recommended method is to reapply the PTF.

SYSMOD entry (global zone)

The global zone SYSMOD entry describes a SYSMOD that exists as an MCS entry in the SMPPTS. It contains information that SMP/E obtained when it received that SYSMOD. Because SMP/E processing is designed to keep the global zone SYSMOD entry and the MCS entry synchronized, you should use only SMP/E commands to update these entries. If you use system utilities to update the MCS entries, you will get unpredictable results when processing the corresponding SYSMOD. For more information about MCS entries, see "MCS entry (SMPPTS)" on page 225.

When SMP/E processes SYSMODs, it uses the global zone SYSMOD entry to determine which SYSMODs are applicable or whether the SYSMODs you selected are applicable. In addition, after a SYSMOD has been successfully applied or accepted, SMP/E records in the SYSMOD entry the names of the zones to which the SYSMOD was applied or accepted.

Note: Although this information is saved in the global zone SYSMOD entry, it is used only for reporting purposes. SMP/E does not use it during APPLY, ACCEPT, or RESTORE processing to determine the status of a SYSMOD. Instead, SMP/E uses SYSMOD entries in the target and distribution zones to determine whether a SYSMOD has been applied or accepted.

A SYSMOD entry is generally kept in the global zone until the associated SYSMOD is accepted; then the entry is deleted. You may want SMP/E to save the global zone SYSMOD entries after ACCEPT processing (for example, if you plan to do a system generation). To do this, specify NOPURGE in the ORDER entry that

is in effect during ACCEPT processing. Likewise, you may want to save the global zone SYSMOD entries after RESTORE processing. To do this, specify NOREJECT in the ORDER entry that is in effect during RESTORE processing.

Subentries

These are the subentries for the global zone SYSMOD entry as they appear in the LIST output. Only a few of these subentries can be changed with UCLIN. The description of each entry indicates whether UCLIN may be used and, if so, what the correct syntax is.

sysmod_id
> is the SYSMOD identification.
>
> The SYSMOD ID must contain 7 alphanumeric characters.

ACCEPT ZONE
> lists the distribution zones into which the SYSMOD has been successfully accepted.
>
> The UCL operand is ACCID(*zone...*).
>
> The name can contain from 1 to 7 alphanumeric characters.

APAR
> indicates that this SYSMOD is an APAR, which provides a corrective fix to a problem.
>
> There is no UCL support for this subentry in the global zone.

APPLY ZONE
> lists the target zones to which the SYSMOD has been successfully applied.
>
> The UCL operand is APPID(*zone...*).
>
> The name can contain from 1 to 7 alphanumeric characters.

DELLMOD
> indicates that the SYSMOD contained a ++DELETE statement.
>
> The UCL operand is DELLMOD.

DELETE
> lists the SYSMODs deleted by this SYSMOD.
>
> There is no UCL support for this subentry in the global zone.

DESCRIPTION
> specifies the descriptive name to be associated with this SYSMOD.
>
> The UCL operand is DESCRIPTION(*description*).

DLMOD
> lists the load modules deleted by ++DELETE statements contained in this SYSMOD.
>
> The UCL operand is DLMOD(*name...*).

element
> lists the data element replacements in the SYSMOD.
>
> There is no UCL support for this subentry in the global zone.
>
> In place of *element*, you will see one of the values shown in Table 2 on page 9.
>
> Some types of elements, such as panels, messages, or text, may have been translated into several languages. In these cases, the *element* operand contains *xxx*, which represents the language used for the element. (If an element was not translated, the *element* operand does not contain any *xxx* value.) Table 3 on page 11 shows the *xxx* values and the languages they represent.

ELEMMOV
> indicates that the SYSMOD contained a ++MOVE statement.
>
> The UCL operand is ELEMMOV.

EMOVE
lists the elements and load modules that were moved by ++MOVE statements contained in this SYSMOD.

ERROR
indicates that an error occurred when this SYSMOD was received.

There is no UCL support for this subentry in the global zone.

FEATURE
lists the names of the software features that contain this SYSMOD. The identified software features correspond to FEATURE entries in the global zone.

The UCL operand is FEATURE(*name,...*).

FESN
identifies the field engineering (FE) service number.

There is no UCL support for this subentry in the global zone.

FMID
identifies the function SYSMOD to which this SYSMOD is applicable.

There is no UCL support for this subentry in the global zone.

FUNCTION
indicates that this SYSMOD is a function, which introduces a new function into the system.

There is no UCL support for this subentry in the global zone.

hfs_element
lists the hierarchical file system element replacements (++hfs_element statements) in the SYSMOD.

There is no UCL support for this subentry in the global zone.

HOLDERROR
lists the error hold reason IDs in effect for this SYSMOD.

The actual ++HOLD statement associated with each reason ID (HOLDDATA) is not in the SYSMOD entry. However, if the LIST command specified HOLDDATA and SYSMOD, the listing of the SYSMOD entry includes the HOLDDATA. For more information, see information about listing HOLDDATA with the LIST command in *z/OS SMP/E Commands*.

There is no UCL support for this subentry in the global zone.

HOLDFIXCAT
lists the fix category hold reason IDs in effect for this SYSMOD.

There is no UCL support for this subentry in the global zone.

HOLDSYSTEM
lists the system hold reason IDs in effect for this SYSMOD.

The actual ++HOLD statement associated with each reason ID (HOLDDATA) is not in the SYSMOD entry. However, if the LIST command specified HOLDDATA and SYSMOD, the listing of the SYSMOD entry includes the HOLDDATA. For more information, see information about listing HOLDDATA with the LIST command in *z/OS SMP/E Commands*.

HOLDSYSTEM subentries show either INT or EXT to indicate the source of the ++HOLD statement.

* INT means that the ++HOLD statement was contained in the held SYSMOD. The held SYSMOD can be installed only if the BYPASS operand is specified on the APPLY or ACCEPT command.
* EXT means that the ++HOLD statement was obtained from another source, such as SMPHOLD. The held SYSMOD can be installed only if the BYPASS operand is specified on the APPLY or ACCEPT command or if the hold is removed by a ++RELEASE statement.

There is no UCL support for this subentry in the global zone.

HOLDUSER
lists the user hold reason IDs in effect for this SYSMOD.

The actual ++HOLD statement associated with each reason ID (HOLDDATA) is not in the SYSMOD entry. However, if the LIST command specified HOLDDATA and SYSMOD, the listing of the SYSMOD entry includes the HOLDDATA. For more information, see information about listing HOLDDATA with the LIST command in *z/OS SMP/E Commands*.

There is no UCL support for this subentry in the global zone.

JAR
lists the JAR file replacements (++JAR) supplied by the SYSMOD.

There is no UCL support for this subentry in the global zone.

JARUPD
lists the JAR file updates (++JARUPD) supplied by the SYSMOD.

There is no UCL support for this subentry in the global zone.

JCLIN
indicates that the SYSMOD contained inline JCLIN.

There is no UCL support for this subentry in the global zone.

MAC
lists the macro replacements (++MAC statements) in the SYSMOD.

There is no UCL support for this subentry in the global zone.

MACUPD
lists the macro updates (++MACUPD statements) in the SYSMOD.

There is no UCL support for this subentry in the global zone.

MOD
lists the module replacements (++MOD statements) in the SYSMOD.

There is no UCL support for this subentry in the global zone.

NPRE
lists negative prerequisite SYSMODs (that is, SYSMODs that must not be present in the system at the same time as this SYSMOD).

There is no UCL support for this subentry in the global zone.

PRE
lists prerequisite SYSMODs (that is, SYSMODs that must be present before this SYSMOD can be installed).

There is no UCL support for this subentry in the global zone.

PROGRAM
lists the program element replacements (++PROGRAM statements) in the SYSMOD.

There is no UCL support for this subentry in the global zone.

PTF
indicates that this SYSMOD is a PTF, which provides preventive service fixes.

There is no UCL support for this subentry in the global zone.

RECDATE
specifies the date on which this SYSMOD was received.

There is no UCL support for this subentry in the global zone.

RECTIME
specifies the time at which this SYSMOD was received.

There is no UCL support for this subentry in the global zone.

RENLMOD
    indicates that the SYSMOD contained a ++RENAME statement.

    The UCL operand is RENLMOD.

REQ
    lists requisite SYSMODs (that is, SYSMODs that must be installed concurrent with this SYSMOD).

    There is no UCL support for this subentry in the global zone.

REWORK
    identifies the level of the SYSMOD, which was received again for minor changes.

    For SYSMODs supplied by IBM, the REWORK level is *yyyyddd*, where *yyyy* is the year the SYSMOD was reworked and *ddd* is the Julian date.

    There is no UCL support for this subentry in the global zone.

RLMOD
    indicates the load modules renamed by ++RENAME statements in this SYSMOD.

    The UCL operand is RLMOD(*name...*).

SOURCEID
    lists the character strings assigned to this SYSMOD during RECEIVE. These values might have been specified by the user on the RECEIVE command, included inline on the ++ASSIGN statement, or assigned by SMP/E when processing the contents of an order.

    The UCL operand is SOURCEID(*source_id...*).

    - The source ID can be from 1 to 64 characters in length and contain any nonblank character (X'41' through X'FE') except single quotation mark ('), asterisk (*), percent (%), left parenthesis ((), and right parenthesis ()).
    - A source ID value cannot span lines.
    - A source ID value is case sensitive.

SRC
    lists the source replacements (++SRC statements) in the SYSMOD.

    There is no UCL support for this subentry in the global zone.

SRCUPD
    lists the source updates (++SRCUPD statements) in the SYSMOD.

    There is no UCL support for this subentry in the global zone.

SREL
    is the system or subsystem release specified on the indicated ++VER MCS in the SYSMOD; for example, SREL VER(001) = Z038. This SREL value matches an SREL value defined in the GLOBALZONE entry that was used when receiving the SYSMOD. These are the systems and subsystems defined by IBM, with their SRELs:

System
    SREL
DB2
    P115
CICS
    C150
IMS
    P115
MVS
    Z038
NCP
    P004

There is no UCL support for this subentry in the global zone.

SUPING
lists the SYSMODs superseded by this SYSMOD.

There is no UCL support for this subentry in the global zone.

SZAP
lists the module superzaps (++ZAP statements) in the SYSMOD.

There is no UCL support for this subentry in the global zone.

TLIBPREFIX
is the high-level data set name qualifier of the SMPTLIB data sets used to receive this SYSMOD, which was packaged in RELFILEs. This is the DSPREFIX value that was used during RECEIVE processing.

The UCL operand is TLIBPREFIX(*prefix*).

- The prefix can contain from 1 to 26 alphanumeric characters.
- The prefix must follow standard naming conventions for data sets.

Note: If the TLIBPREFIX subentry is deleted from the entry for a SYSMOD packaged in RELFILE format, the SYSMOD cannot be applied or accepted until the TLIBPREFIX subentry is re-created with the current prefix for the SMPTLIB data sets.

USERMOD
indicates that this SYSMOD is a USERMOD, which puts a user modification into the system.

There is no UCL support for this subentry in the global zone.

VERSION
lists the function SYSMODs that are versioned by this SYSMOD. Versioning indicates that, if there are any elements in common between the SYSMODs listed and this SYSMOD, this SYSMOD's elements are at a higher functional level, and are therefore the ones that should be installed.

There is no UCL support for this subentry in the global zone.

XZAP
lists the module superzaps (++ZAP statements) in the SYSMOD that contain an EXPAND statement (indicating that the module should be expanded before it is updated).

There is no UCL support for this subentry in the global zone.

LIST Examples

To list all the SYSMOD entries in a global zone, you can use the following commands:

```
SET      BDY(GLOBAL)          /* Set to requested zone.   */.
LIST     SYSMOD               /* List all SYSMOD entries. */.
```

To list specific SYSMOD entries in a global zone, you can use these commands:

```
SET      BDY(GLOBAL)          /* Set to requested zone.   */.
LIST     SYSMOD(UZ00001       /* List only these two      */
              UZ00002)        /* entries.                 */.
```

To list specific SYSMOD entries in the global zone along with the associated HOLDDATA, you can use these commands:

```
SET      BDY(GLOBAL)          /* Set to requested zone.    */.
LIST     SYSMOD(UZ12345       /* List only these two       */
              UZ56789)        /* entries                   */
         HOLDDATA             /* plus HOLDDATA for them.   */.
```

The format of the LIST output for each SYSMOD entry is the same for all of these commands. The only differences are the number of SYSMOD entries listed and the amount of information presented for each entry (HOLDDATA causes additional information to be listed).

Figure 57 on page 283, Figure 58 on page 284, and Figure 59 on page 285 are examples of LIST output for SYSMOD entries.

Note: If HOLDDATA for a particular SYSMOD has been received but the SYSMOD itself has not yet been received, only the hold information for the SYSMOD is listed. For an example, see SYSMOD UZ56789 in Figure 59 on page 285.

```
PAGE nnnn  - NOW SET TO zzzzzz ZONE nnnnnnn  DATE mm/dd/yy TIME hh:mm:ss SMP/E 36.nn SMPLIST OUTPUT

GLOBAL        SYSMOD ENTRIES


  NAME

JXY2102   TYPE           = FUNCTION
          STATUS         = REC
          JCLIN          = YES
          FESN           = 1234567
          TLIBPREFIX     = SMP.RELFILE
          DATE/TIME REC  = 06.100  08:00:00
          APPLY ZONE     = TGT1      TGT2
          ACCEPT ZONE    = DLIB1     DLIB2
          SREL   VER(001) = Z038
          DELETE VER(001) = JXY1102   JXY1121   JXY1122
          NPRE   VER(001) = HZZ1102
          SUPING VER(001) = AZ11111   AZ11112   AZ11113
          VERS   VER(001) = HYY1102   JYY1121   JYY1122   JYY1123
          MAC            = MAC01     MAC02     MAC03
          MOD            = MOD01     MOD02     MOD03     MOD04
          SRC            = SRC01     SRC02

JXY2121   TYPE           = FUNCTION
          STATUS         = REC
          JCLIN          = YES
          FESN           = 1234567
          TLIBPREFIX     = SMP.RELFILE
          DATE/TIME REC  = 06.100  08:30:00
          APPLY ZONE     = TGT1      TGT2
          ACCEPT ZONE    = DLIB1     DLIB2
          SREL   VER(001) = Z038
          DELETE VER(001) = JXY1102   JXY1121   JXY1122
          FMID   VER(001) = JXY2102
          SUPING VER(001) = AZ11121   AZ11122   AZ11123
          MAC            = MAC01
          MOD            = MOD01     MOD02
          SRC            = SRC01     SRC03     SRC04

AZ99801   TYPE           = APAR
          STATUS         = REC
          DATE/TIME REC  = 06.100  08:00:00
          APPLY ZONE     = TGT1      TGT2
          ACCEPT ZONE    = DLIB1     DLIB2
          SREL   VER(001) = Z038
          FMID   VER(001) = JXY2102
          ZAP            = MOD01
```

Figure 57. SYSMOD entry: sample LIST output for a global zone (Example 1)

SYSMOD entry (global zone)

```
PAGE nnnn  - NOW SET TO zzzzzz ZONE nnnnnnn  DATE mm/dd/yy TIME hh:mm:ss SMP/E 36.nn SMPLIST OUTPUT

GLOBAL        SYSMOD ENTRIES


  NAME

LUS0001   TYPE          = USERMOD
          STATUS        = REC
          DATE/TIME REC = 06.100  08:00:00
          APPLY ZONE    = TGT1      TGT2
          ACCEPT ZONE   = DLIB1     DLIB2
          SREL  VER(001) = Z038
          FMID  VER(001) = JXY2102
          PRE   VER(001) = UZ00010
          MACUPD        = MAC02

UZ00010   TYPE          = PTF
          STATUS        = REC
          DATE/TIME REC = 06.100  08:00:00
          APPLY ZONE    = TGT1      TGT2
          ACCEPT ZONE   = DLIB1     DLIB2
          SREL  VER(001) = Z038
          FMID  VER(001) = JXY2102
          PRE   VER(001) = UZ00008   UZ00007
          REQ   VER(001) = UZ00040
          SUPING VER(001) = AZ99801   AZ99802    UZ00009
          MAC           = MAC01
          MACUPD        = MAC02
          MOD           = MOD01
          SRC           = SRC01
          SRCUPD        = SRC01

UZ00011   TYPE          = PTF
          STATUS        = REC
          DATE/TIME REC = 06.150  08:00:00
          SOURCEID      = ABC0706   PUT0704    XAU3380
          APPLY ZONE    = TGT1      TGT2
          ACCEPT ZONE   = DLIB1     DLIB2
          SREL  VER(001) = Z038
          FMID  VER(001) = JXY2102
          PRE   VER(001) = UZ00010
          MOD           = MOD01
```

Figure 58. SYSMOD entry: sample LIST output for a global zone (Example 2)

```
PAGE nnnn  - NOW SET TO zzzzzz ZONE nnnnnnn  DATE mm/dd/yy TIME hh:mm:ss SMP/E 36.nn SMPLIST OUTPUT

GLOBAL        SYSMOD ENTRIES


   NAME

UZ12345   TYPE          = PTF
          STATUS        = REC
          DATE/TIME REC = 06.100  08:00:00
          SREL   VER(001) = Z038
          DELETE VER(001) = HBB2102
          SUPING VER(001) = AZ11111   AZ11112   AZ11113
          MAC           = MAC01     MAC02     MAC03
          MOD           = MOD01     MOD02     MOD03     MOD04
          SRC           = SRC01     SRC02
          HOLDERROR     = AZ00001   ++HOLD(UZ12345) ERROR
                                       REASON(AZ00001) FMID(HBB2102)
                                       COMMENT(SMRTDATA(CHGDTE(071105) SYMP(PRV,IPL))).
          HOLDERROR     = AZ00002   ++HOLD(UZ12345) ERROR
                                       REASON(AZ00002) FMID(HBB2102)
                                       COMMENT(SMRTDATA(CHGDTE(071106) SYMP(FUL))).
          HOLDSYSTEM(INT) = DOC     ++HOLD(UZ12345) SYSTEM
                                       REASON(DOC) FMID(HBB2102)
                                       COMMENT(NEW MSG).
          HOLDSYSTEM(EXT) = UCLIN   ++HOLD(UZ12345) SYSTEM
                                       REASON(UCLIN) FMID(HBB2102)
                                       COMMENT(UCLIN REQUIRED).
          HOLDUSER      = INUSE     ++HOLD(UZ12345) USER
                                       REASON(INUSE) FMID(HBB2102)
                                       COMMENT(IM MODIFYING).

UZ56789   HOLDERROR     = AZ00023   ++HOLD(UZ56789) ERROR
                                       REASON(AZ00023) FMID(HBB2102)
                                       COMMENT(SMRTDATA(CHGDTE(071110) SYMP(ABENDS with E37))).
          HOLDERROR     = AZ00024   ++HOLD(UZ56789) ERROR
                                       REASON(AZ00024) FMID(HBB2102)
                                       COMMENT(SMRTDATA(CHGDTE(071113) SYMP(DAL) FIX(UW92458))).
```

Figure 59. SYSMOD entry: sample LIST output when HOLDDATA is specified

UCLIN Examples

You can use the ADD, REP, and DEL UCL statements to change subentries in a SYSMOD entry. When you use UCLIN to update a SYSMOD entry, remember that if a DEL statement deletes all the existing subentries in the entry, SMP/E deletes the entire entry.

The following examples are provided to help you use the SYSMOD entry.

Example 1: Changing the SOURCEID of a SYSMOD

Assume that you received a SYSMOD and assigned it a SOURCEID value of DATALINK, and that you now want to change the SOURCEID value to PUT0701 so that the SYSMOD will be installed with that service level. The following UCL can be used to change the SOURCEID value:

```
SET      BDY(GLOBAL)        /* Set to global zone.     */.
UCLIN                       /*                         */.
REP      SYSMOD(UZ00001)    /* Specify SYSMOD.         */
         SOURCEID(PUT01)  /* Change SOURCEID value.   */
                            /*                         */.
ENDUCL                      /*                         */.
```

Example 2: Indicating that a SYSMOD was applied

Assume that you have received some service, accepted it with the BYPASS(APPLYCHECK) operand, and then performed a system generation. That service is now actually in the target libraries, and you would like that recorded in the global zone SYSMOD entry. Assume that the name of your target zone is TGT1. The following UCL can be used:

```
SET       BDY(GLOBAL)       /* Set to global zone.     */.
UCLIN                       /*                         */.
ADD       SYSMOD(UZ00001)   /* Specify SYSMOD.         */
          APPID(TGT1)       /* Applied to this zone.   */
                            /*                         */.
ENDUCL                      /*                         */.
```

Note: One UCL statement is required for each SYSMOD.

TARGETZONE entry (target zone)

The TARGETZONE entry contains information SMP/E uses to process a specific target zone and the associated target libraries. It is created by UCLIN and must be defined before you can do any other processing for that target zone.

Subentries

These are the subentries for the TARGETZONE entry as they appear in the LIST output:

name
> is the name of the target zone. You assign the name when the zone is created.
>
> The name can contain from one to seven alphanumeric characters (A–Z, 0–9) or national characters ($, #, @). The first character must be alphabetic.

OPTIONS
> is the name of the OPTIONS entry in the global zone that should be used in processing this target zone. For more information, see "OPTIONS entry (global zone)" on page 241.
>
> The UCL operand is OPTIONS(*name*).
>
> - The name can contain from one to eight alphanumeric characters.
> - This name can be overridden by using the OPTIONS parameter on the SET command. For more information, see *z/OS SMP/E Commands*.
> - If no OPTIONS entry name is specified, SMP/E uses a set of default utility values when processing this target zone. For more information, see "OPTIONS entry (global zone)" on page 241.

RELATED
> is the name of the distribution zone to which this target zone is related. A target zone is related to the distribution zone used to build the target libraries, such as during system generation.
>
> The UCL operand is RELATED(*zone*).
>
> - The zone name can contain from one to seven alphanumeric characters.
> - Although the entry can be defined without this subentry, you **must** define the subentry before you can install any SYSMODs in the target libraries.

SREL
> lists the system releases to be supported in this target zone.
>
> The UCL operand is SREL(*srel...*).
>
> - The SREL must contain four alphanumeric characters, usually one alphabetic character followed by three numeric characters. These are the SRELs defined by IBM:
>
>   **System**
>   > **SREL**
>
>   **DB2**
>   > P115
>
>   **CICS**
>   > C150
>
>   **IMS**
>   > P115

MVS
> Z038

NCP
> P004

- Although the entry can be defined without this subentry, you **must** define the subentry before you can install any SYSMODs in the target libraries.

Note: Although you can support multiple products with different SREL values from one target zone, those products are still subject to all other restrictions related to combining products in one zone. The most common reason for not being able to combine products is common element names. For example, modules or macros with the same name are found in both products, but reside in different libraries.

TIEDTO
> specifies other target zones that either:
>
> - Supplied modules for load modules controlled by this target zone
> - Control load modules that have been link-edited with modules supplied by this target zone
>
> The UCL operand is TIEDTO(*zone...*).
>
> **Note:** TIEDTO subentries are added automatically during LINK MODULE command processing. However, they are **never** automatically deleted.

UPGLEVEL
> indicates the highest SMP/E release level that is allowed to make incompatible changes to the target zone. Before making an incompatible change to the target zone, SMP/E will check the UPGLEVEL value for that zone. If the release level of SMP/E is higher than the zone's UPGLEVEL value, SMP/E will not make the incompatible change.
>
> The UPGLEVEL value is in the form *vr.pp*, where *vr* represents the version and release of SMP/E and *pp* represents the PTF level of SMP/E.
>
> There is no UCL support for this subentry. When a zone is created by SMP/E using the UCLIN command or the Administration dialog, SMP/E sets the UPGLEVEL subentry value for that zone to the level of SMP/E used to create the zone. The UPGRADE command is used to change the UPGLEVEL subentry value for a zone.

XZLINK
> specifies whether APPLY and RESTORE processing in another zone should automatically update load modules in this zone when cross-zone modules previously added to those load modules by the LINK MODULE command are changed.
>
> The UCL operand is XZLINK(*value*).
>
> - This subentry can contain one of the following values:
>
>   **DEFERRED**
>   > Cross-zone load modules controlled by this zone should **not** be automatically updated when modules previously included in them by the LINK MODULE command are updated or deleted.
>   >
>   > To make sure that the modules are synchronized with the cross-zone load modules and that the cross-zone information in SMP/E entries is correct, some combination of the following commands must be run later (depending on how the module changes affect the load modules):
>   >
>   > - LINK MODULE command, to include modules in the affected load modules
>   > - Link-edit (outside of SMP/E), to delete modules from the affected load modules
>   > - UCLIN command, to update cross-zone subentries as necessary
>
>   **AUTOMATIC**
>   > Cross-zone load modules controlled by this zone should be automatically updated when modules previously included in them by the LINK MODULE command are updated or deleted.
>
> - If XZLINK is not specified, SMP/E uses the default value of DEFERRED.

- XZLINK does not affect processing of the LINK command.
- The XZLINK(DEFERRED) value is listed only when the TARGETZONE entry contains TIEDTO records.

ZDESC
is a user-written description for this zone.

The UCL operand is ZONEDESCRIPTION(*text*).

- The zone description can be in single-byte characters (such as English alphanumeric characters) or in double-byte characters (such as Kanji).
- The zone description can contain up to 500 bytes of data, including blanks. (For double-byte data, the 500-byte maximum includes all shift-in and shift-out characters, as well as the double-byte characters.) Extra blanks are deleted. All data beyond column 72 is ignored, including blanks.
- The zone description cannot be only blanks.
- If parentheses are included in the text, they must be in matched pairs.

LIST Examples

To list the TARGETZONE entry for a particular target zone, you can use the following commands:

```
SET       BDY(TGT1)            /* Set to requested target. */.
LIST      TARGETZONE           /* List TARGETZONE entry.   */.
```

Figure 60 on page 288 is an example of LIST output for a TARGETZONE entry.

```
PAGE nnnn  - NOW SET TO zzzzzz ZONE nnnnnnn  DATE mm/dd/yy TIME hh:mm:ss SMP/E 36.nn SMPLIST OUTPUT

TGT1           TZONE ENTRY


  NAME

TGT1      TZONE            = TGT1
          ZDESC            = ZONE DESCRIPTION FOR TGT1 ZONE
          RELATED          = DLIB1
          SREL             = P115      R020      Z038
          OPTIONS          = OPTTGT1
          XZLINK           = AUTOMATIC
          TIEDTO           = CICS1     IMS1
```

Figure 60. TARGETZONE entry: sample LIST output

UCLIN Examples

You can use the ADD, REP, and DEL UCL statements to change subentries in the TARGETZONE entry. When you use UCLIN to update a TARGETZONE entry, remember that if a DEL statement deletes all the existing subentries in the entry, SMP/E deletes the entire entry.

The following examples are provided to help you use the TARGETZONE entry.

Example 1: Defining a TARGETZONE entry

Assume that you are about to build a new set of target libraries and thus want to define a new target zone named TGT2. The distribution zone that you have done the system generation from is DLIB1, the OPTIONS entry that will be set up will be called OPTTGT2, and the target zone will contain MVS (Z038) and two additional programs, each with its own SREL value (P010 and R020). The target zone will exist in data set SMPE.SMPCSI.CSI. Use the following UCLIN to define the new zone.

```
SET       BDY(GLOBAL)          /* Set to global zone.    */.
UCLIN                          /* UCLIN for GZONE entry  */.
ADD       GZONE                /* to set up              */
          ZONEINDEX(           /* index for new zone.    */
                  (TGT2,SMPE.SMPCSI.CSI,TARGET) /*    */
                  )            /*                        */
```

```
                          /*                      */.
ENDUCL                    /* End global zone update. */.
SET       BDY(TGT2)       /* Now define new zone.  */.
UCLIN                     /* UCLIN to define it.   */.
ADD       TARGETZONE(TGT2) /* Identify name.        */
          OPTIONS(OPTTGT2) /* OPTIONS entry to use. */
          SREL(Z038,       /* SRELs for MVS and     */
               P010,       /* two other             */
               R020)       /* programs.             */
          RELATED(DLIB1)   /* Generated from DLIB.  */
                          /*                      */.
ENDUCL                    /*                      */.
```

Note: Even though the OPTIONS entry has not been set up yet, you can still refer to it in the TARGETZONE entry. Prior to processing this target zone the OPTIONS entry must be created. For examples of setting up the OPTIONS entries, see "OPTIONS entry (global zone)" on page 241.

Example 2: Formatting a zone description

Assume that you enter the following zone description with the first line ending in column 72 and the second line starting in column 1:

```
----+----1----+----2----  …  ----5----+----6----+----7--
SET       BDY(TGT1)           /* Set to tgt zone.      */.
UCLIN                         /* UCLIN for TZONE entry */.
ADD       TZONE(TGT1)         /* to set up.            */
          ZDESC(                 THIS IS THE DESCRIPTION FOR
THE TGT1 ZONE)
                              /* End of zone description. */.
ENDUCL                        /* End tgt zone update.  */.
```

Because there is no blank between the word ending in column 72 and the next word starting in column 1, SMP/E will run the two together.

Words in a zone description, even words that end in column 72, must be separated by a blank. To format the zone description in this example correctly, you can put a blank at the beginning of the second line:

```
----+----1----+----2----  …  ----5----+----6----+----7--
SET       BDY(TGT1)           /* Set to tgt zone.      */.
UCLIN                         /* UCLIN for TZONE entry */.
ADD       TZONE(TGT1)         /* to set up.            */
          ZDESC(                 THIS IS THE DESCRIPTION FOR
 THE TGT1 ZONE)
                              /* End of zone description. */.
ENDUCL                        /* End tgt zone update.  */.
```

Because there is a blank explicitly coded between the word ending in column 72 and the word starting in column 1, SMP/E will not run the words together.

UTILITY entry (global zone)

The UTILITY entry contains information that SMP/E uses when invoking a particular system utility program, such as the load module name of the program and the parameters that should be passed. This information is used only if the UTILITY entry is named in the OPTIONS entry that is in effect. For example, to have SMP/E use certain parameters when it calls a specific link-edit utility, you must do the following:

1. Define a UTILITY entry that names the program and specifies the parameters.

2. Define an OPTIONS entry that specifies that UTILITY entry as the one to use for the link-edit utility.

3. Put that OPTIONS entry into effect, either by specifying it on the SET command or by defining it as the default OPTIONS entry for the zone to be processed.

If the OPTIONS entry does not point to a UTILITY entry for a particular system utility, SMP/E uses default values for that utility. Table 7 on page 290 lists the default values for the various types of utility programs.

Table 7. Default values for UTILITY entries

Utility	NAME (see note 1)	RC	PRINT	PARM
Access method services	IDCAMS	0	SYSPRINT	
Assembler	ASMA90	4	SYSPRINT	XREF, NOOBJECT, DECK
Compress	IEBCOPY	0	SYSPRINT	
Copy	IEBCOPY	0	SYSPRINT	SPCLCMOD and CMWA=256K (for program elements and copied load modules)
Hierarchical file system copy	BPXCOPY	0	SYSPRINT (see note 3)	
Link-edit utility	IEWBLINK	8	SYSPRINT	LET, LIST, NCAL, XREF (see note 2)
Retry after *x*37 abends	IEBCOPY	0	SYSPRINT	
Update	IEBUPDTE	0	SYSPRINT	Determined by SMP/E during processing
Superzap	IMASPZAP	4	SYSPRINT	

Notes:

1. If you replace a default utility program, the replacement utility program must be compatible with the default utility it replaces, both in the way it processes any control statements and execution parameters generated by SMP/E and in the return codes that it returns to SMP/E.

2. When the load module being link-edited contains a CALLLIBS subentry list, SMP/E does not always use NCAL by default. In this case, SMP/E uses CALL for the link to the actual target library or NCAL for the link to the SMPLTS library. SMP/E always uses NCAL for ACCEPT processing.

3. If SYSTSPRT is specified as the PRINT value, it is ignored and the default of SYSPRINT is used instead.

If a UTILITY entry defines some, but not all, of these subentries, SMP/E uses the default values for the subentries not specified. For more information about OPTIONS entries, see

The names of the utilities SMP/E may call include all of SMP/E's default utility program names (BPXCOPY, ASMA90, IEBCOPY, IEBUPDTE, IEWBLINK, IEWL, and IMASPZAP) and any utility program name specified in an active UTILITY entry. Because SMP/E runs authorized, all the utility programs called by SMP/E must reside in an authorized library. Also, if a particular utility program is to be restricted, the z/OS Security Server must be used to control its execution.

Subentries

These are the subentries for the UTILITY entry as they appear in the LIST output:

name
> is the name of the UTILITY entry.
>
> The name can contain from 1 to 8 alphanumeric characters.

LIST

indicates whether member names should be listed when SMP/E invokes a copy utility to perform compress processing, retry processing, or element installation.

The UCL operand is LIST(YES|NO).

- YES indicates that member names should be listed during copy processing. This is the default.

- NO indicates that the list of member names should be suppressed during copy processing.

- The LIST value (including the default value) is ignored when the UTILITY entry is not for a copy utility.

NAME

is the name of the load module for the utility program that SMP/E is to call. Table 7 on page 290 lists the default names used by SMP/E if the OPTIONS entry in effect does not point to a UTILITY entry for a specific system utility.

The UCL operand is NAME(*prog*).

The name can contain from 1 to 8 alphanumeric characters. The first character must be alphabetic.

PARM

specifies the parameters to be passed to the utility program. SMP/E may add other parameters to this list. Table 7 on page 290 lists the default parameters used by SMP/E if the OPTIONS entry in effect does not point to a UTILITY entry for a specific system utility.

The UCL operand is PARM(*string*).

- If you specify PARM, you must specify all the parameters to be passed. You cannot specify a single value to add or change just one parameter.

- The parameter string is usually a hexadecimal character string with a length of up to 100 characters. This 100-character limit can be exceeded for a binder link-edit step using the OPTION option. For a more detailed explanation, see the GENERATE command chapter in *z/OS SMP/E Commands*. No validity checking is done on the string. If any blanks are specified between the parentheses, they will be deleted by SMP/E during processing.

- If the string contains parentheses, they must be in matched pairs. SMP/E assumes it has reached the end of the specified string when it encounters a closing parenthesis for the opening parenthesis. For example, the following is valid because there are matching parentheses within the string:

```
PARM(LIST,LET,SIZE=(1526K,80K),NCAL)
```

Parentheses within quotation marks are still considered parentheses. The next example is not valid because there are unmatched parentheses. SMP/E would continue scanning after the last parenthesis looking for another ")".

```
PARM(PRINT(A),SPECIALCHAR'(')
```

Likewise, the following example is not valid. SMP/E would stop at the ")" after SPECIALCHAR and give a syntax error on the closing quotation mark.

```
PARM(PRINT(A),SPECIALCHAR')')
```

- **Assembler utility:** If no PARM subentry is specified for an assembler utility, SMP/E passes the default PARM values shown for assembler utilities in Table 7 on page 290.

  If you change the PARM subentry, either include DECK in the parameters chosen or ensure that the assembler program produces an object deck.

- **HFS copy utility:** If the UTILITY entry for the HFS copy utility specifies a PARM value, those parameters are passed to the utility in addition to any parameters saved in the hierarchical file system element entry.

- **Link-edit utility:** SMP/E constructs the string of link-edit parameters to be passed to the link-edit utility in the following steps:

1. If a PARM subentry is specified for a link-edit utility, SMP/E starts the string with the parameters specified in the PARM subentry. Otherwise, SMP/E starts the string with the default parameters LET, LIST, NCAL, and XREF.

2. If SMP/E finds any link edit parameters in the LMOD entry for the load module being link-edited, SMP/E adds those parameters to the string.

3. If the load module being link-edited contains a CALLLIBS subentry list, then:

   – SMP/E adds CALL for the link to the actual target library

   – SMP/E adds NCAL for the link to the SMPLTS library

   Any parameter added by this step will override any CALL or NCAL parameters added in previous steps, whether they were explicitly specified or assumed by default.

4. Finally, if SMP/E determines that the binder is available on the system and SMP/E is processing CSECT deletes, SMP/E adds the STORENX link-edit parameter to the parameter list. (If you use the linkage-editor instead of the binder, STORENX will cause "INVALID" to appear in the "invocation parameters list" output of the linkage-editor. This has no effect on processing.)

The consolidated parameter list is passed to the link-edit utility.

LET, LIST, and sometimes either CALL or NCAL are ordinarily required for maintenance of IBM operating systems. If you change the PARM subentry, be sure to include LET, LIST, and any other required parameters in the parameters chosen. For more information, see the APPLY command chapter in *z/OS SMP/E Commands*.

Note: Before using the GENERATE command to create JCL that will update files in a UNIX file system, you may want to add UID(0) to the PARM value so that the JCL created by GENERATE will include UID(0) in the execution parameter string for the link-edit utility. Specify UID(0) if all the following are true:

1. Your UID is not 0 but you are authorized to the BPX.SUPERUSER facility class profile. The UID(0) option, in this case, causes the binder to set an effective UID of 0 for its execution.

2. The binder invoked by the generated JCL is at the proper level to understand the UID option.

3. UID 0 authority is needed for the binder's execution in the JCL created by the GENERATE command in order to update files in a UNIX file system.

If you do specify UID(0) for GENERATE, you must remove it after GENERATE has run, so that it is not used for other SMP/E commands (such as APPLY, which handles setting the effective UID itself prior to invoking the binder).

- **Update utility:** The PARM subentry for an update utility must not specify MOD or NEW.

PRINT
specifies the ddname that is to contain output from the utility (for example, SYSPRINT). lists the default ddnames used by SMP/E if the OPTIONS entry in effect does not point to a UTILITY entry for a specific system utility.

The UCL operand is PRINT(*ddname*).

The value specified can contain from 1 to 8 alphanumeric characters. The first character must be alphabetic.

Note:

1. The ddname specified for PRINT can affect whether you receive print output from this utility. For example, if you specify a DDDEF for a DUMMY data set or a DDDEF for a data set that is sent to a SYSOUT class that suppresses output, you do not receive print output from this utility.

2. If SYSTSPRT is specified as the PRINT value for the HFS copy utility, it is ignored and the default of SYSPRINT is used instead.

3. To exploit utility multi-tasking in SMP/E, ensure that the ddname that is to contain the link edit utility output is defined with either a DDDEF entry that identifies a SYSOUT class or an entry in the SMPPARM GIMDDALC member that identifies a SYSOUT class. SMP/E's default ddname for utility output is SYSPRINT, but can be changed using the PRINT subentry of the LKED UTILITY entry.

When multi-tasking, SMP/E will invoke multiple instances of the link edit utility at the same time, thus decreasing the total time required to complete an ACCEPT, APPLY, LINK LMODS or RESTORE command. Multi-tasking of link edit can occur when there are different target libraries and there are no dependencies on previous and subsequent link edits. If you do not define the print ddname using either a DDDEF entry or an SMPPARM GIMDDALC member, or if the print ddname definition identifies something other than SYSOUT class, or if you override the SYSPRINT DDDEF with a ddname in your JCL, then SMP/E will not multi-task link edit utility operations.

RC

specifies the maximum acceptable return code from this utility. If the return code is higher than this value, SMP/E normally assumes that the requested processing failed. However, the RETURN CODE value in the LMOD entry overrides (for that load module) the value in the UTILITY entry. Also, the success of a link-edit in the SMPLTS library is based on the following:

- Link-edits into the SMPLTS with a return code of 8 or less are considered successful regardless of the threshold return code specified in the UTILITY entry.
- Link-edits into the SMPLTS with a return code of greater than 8 are considered successful or failures based on the threshold (normal processing). If the SMPLTS link-edit return code is less than or equal to the threshold return code, it is considered successful. If the SMPLTS link-edit return code is greater than the threshold return code, it is considered a failure.

Table 7 on page 290 lists the default return codes used by SMP/E if the OPTIONS entry in effect does not point to a UTILITY entry for a specific system utility.

The UCL operand is RC(*rc*).

- The value can be from 0 to 16.
- The installation information for a product (such as the product program directory) may state the utility return codes you should expect during SMP/E processing. For example, it may state the expected link-edit return codes for its load modules during SMP/E processing. The expected return code may be 4 or 8, because post-SMP/E link-edit work is required (for example, the load modules may require interface routines or compiler library routines). Such return codes allow the SYSMODs to be installed, but they also require you to check the actual link-edit return code in the GIM23903I or GIM23903W messages in order to determine the actual success of utility processing.

  Before using the default SMP/E return codes (especially for link-edit processing), check the installation information for the products you plan to install and determine the appropriate maximum acceptable return codes for utility processing. You may find that you need to set up more than one UTILITY entry for a particular utility program in order to accommodate different maximum return codes for various products. (As a result, you may need additional OPTIONS entries to point to the appropriate UTILITY entries.)

LIST Examples

To list all the UTILITY entries in a global zone, you can

```
SET       BDY(GLOBAL)          /* Set to requested zone.     */.
LIST      UTILITY              /* List all UTILITY entries.  */.
```

To list specific UTILITY entries in a global zone, you can use these commands:

```
SET       BDY(GLOBAL)          /* Set to requested zone.     */.
LIST      UTILITY(IEUASM       /* List only these three      */
               IEWL            /* entries.                   */
               MYX37)          /*                            */.
```

The format of the LIST output for each UTILITY entry is the same for both of these commands. The only difference is the number of UTILITY entries listed.

Figure 61 on page 294 shows an example of LIST output for UTILITY entries.

UTILITY entry (global zone)

```
PAGE nnnn  - NOW SET TO zzzzzz ZONE nnnnnnn  DATE mm/dd/yy TIME hh:mm:ss SMP/E 36.nn SMPLIST OUTPUT

GLOBAL        UTILITY ENTRIES


  NAME

IEUASM     PRINT        = MYSYSPRT

IEWL       PARM         = SIZE=(1526K,80K)

MYX37      NAME         = USERRCVR
           PARM         = TYPE=FAST
           PRINT        = X37PRINT
           RC           = 4
           LIST         = NO
```

Figure 61. UTILITY entry: sample LIST output

UCLIN Examples

You can use the ADD, REP, and DEL UCL statements to change subentries in a UTILITY entry. When you use UCLIN to update a UTILITY entry, keep these points in mind:

- After the UCLIN changes are done, the UTILITY entry must contain at least one of the following subentries:
  - NAME
  - PARM
  - PRINT
  - RC

  Otherwise, there is not enough information in the entry for SMP/E to use the entry.
- If a DEL statement deletes all the existing subentries in the entry, SMP/E deletes the entire entry.

The following examples are provided to help you use the UTILITY entry.

Example 1: Changing the DD statement for SYSPRINT output

Assume that you want to direct all the SYSPRINT output for assemblies to the MYSYSPRT DD statement rather than to the SYSPRINT DD statement. To do this you need to build a UTILITY entry for the assembler and then specify that UTILITY entry in the OPTIONS entry that will be in effect. Assume that the UTILITY entry to be defined is named IEUASM, and that the OPTIONS entry to be used is MYOPT1. The following UCL will accomplish your objectives:

```
SET     BDY(GLOBAL)      /* Set to global zone.    */.
UCLIN                    /*                        */.
ADD     UTILITY(IEUASM)  /* Assembler utility.     */
        PRINT(MYSYSPRT)  /* Alternate SYSPRINT     */
                         /* Other values remain as
                            in SMP/E defaults.     */.
REP     OPTIONS(MYOPT1)  /* Connect to OPTIONS.    */
        ASM(IEUASM)      /* Use IEUASM for ASM.    */
                         /*                        */.
ENDUCL                   /*                        */.
```

Note: Remember, if you specify a DDDEF for a DUMMY data set or a DDDEF for a data set that is sent to a SYSOUT class that suppresses output, you do not receive SYSPRINT output from this utility.

Example 2: Defining link-edit utility parameters

Assume that you want to always pass the SIZE=(1526K,80K) parameter to the link-edit utility. To do this you need to build a UTILITY entry for the link-edit utility and then specify that UTILITY entry in the

OPTIONS entry that will be in effect. Assume that the UTILITY entry to be defined is named IEWL, and that the OPTIONS entry to be used is MYOPT1. The following UCL will accomplish your objectives:

```
SET     BDY(GLOBAL)             /* Set to global zone.     */.
UCLIN                           /*                         */.
ADD     UTILITY(IEWL)           /* Link-edit utility.      */
        PARM(LET,               /* PARM values.            */
            SIZE=(1526K,80K),   /*                         */
            NCAL)               /*                         */
                                /* Other values remain as
                                   in SMP/E defaults.      */.
REP     OPTIONS(MYOPT1)         /* Connect to OPTIONS.     */
        LKED(IEWL)              /* Use IEWL for LKED.      */
                                /*                         */.
ENDUCL                          /*                         */.
```

Example 3: Defining a user utility program

Assume that you want SMP/E to call a user routine, USERRCVR, rather than IEBCOPY, in order to recover from *x*37 abends. In addition to the program name change, the program must also receive parameter TYPE=FAST. You want to indicate that the output should go to X37PRINT rather than SYSPRINT, and that a return code of 4 or less is acceptable. You also want to suppress the listing of member names during retry processing done by your program. To do this you need to build a UTILITY entry for the program, and then specify the UTILITY entry in the OPTIONS entry that will be in effect. Assume that the UTILITY entry to be defined is named MYX37, and that the OPTIONS entry to be used is MYOPT1. The following UCL will accomplish your objectives:

```
SET     BDY(GLOBAL)          /* Set to global zone.      */.
UCLIN                        /*                          */.
ADD     UTILITY(MYX37)       /* Retry/recovery program.  */
        NAME(USERRCVR)       /* Program name.            */
        PARM(TYPE=FAST)      /* PARM value.              */
        PRINT(X37PRINT)      /* SYSPRINT ddname.         */
        RC(4)                /* Highest acceptable       */
                             /* return code.             */
        LIST(NO)             /* No list of member names. */
                             /*                          */.
REP     OPTIONS(MYOPT1)      /* Connect to OPTIONS.      */
        RETRY(MYX37)         /* Use MYX37 for RETRY.     */
                             /*                          */.
ENDUCL                       /*                          */.
```

ZONESET entry (global zone)

The ZONESET entry defines a group of zones to be used to limit the SYSMODs processed by an SMP/E command. For example, you can specify a ZONESET on the ZONESET operand of the REPORT command. This defines which zones SMP/E should check for installed SYSMODs that specify conditional requisites that might be needed for functions in other zones in the ZONESET. A ZONESET may also define a group of zones to be checked or ignored by the REJECT command.

Subentries

These are the subentries for the ZONESET entry as they appear in the LIST output:

name
> is the name of the ZONESET. The name can contain from 1 to 8 alphanumeric characters.

> To avoid confusion and undesired results, you may want to avoid giving a ZONESET the same name as any of the target or distribution zones defined to the global zone that will contain the ZONESET entry. This is because, on some SMP/E command operands, you can specify zones and ZONESETs. When ZONESETs and zones have the same name, you might not get the results you wanted.

> For example, suppose you have a ZONESET named Z001 and a zone named Z001. If you specify Z001 on an SMP/E command operand, SMP/E assumes that you want to use the zones defined in ZONESET Z001 (which might or might not include zone Z001), and not the individual zone Z001.

XZREQCHK
indicates whether this ZONESET should be used when establishing the default zone group for the APPLY, ACCEPT, and RESTORE commands.

The UCL operand is XZREQCHK(YES|NO).

- YES indicates that the ZONESET should be used when establishing the default zone group.
- NO indicates that the ZONESET should not be used when establishing the default zone group. NO is the default.
- The XZREQCHK value (including the default value) is ignored when the XZGROUP operand is specified on the APPLY, ACCEPT, or RESTORE command.

ZONE
lists the target or distribution zones that are to be part of this ZONESET.

The UCL operand is ZONE(*zone...*).

- Each value can contain from 1 to 7 alphanumeric characters.
- A ZONESET can contain both target and distribution zones.
- All the zones in a ZONESET must be defined in the same global zone as the ZONESET entry.
- The zones cannot be defined in global zones that are in different SMPCSI data sets. For an example of defining a ZONESET in order to report on zones controlled by different global zones, see the REPORT CROSSZONE command in *z/OS SMP/E Commands*.

LIST Examples

To list all the ZONESET entries in a global zone, you can use the following commands:

```
SET       BDY(GLOBAL)        /* Set to global zone.       */.
LIST      ZONESET            /* List all ZONESET entries. */.
```

To list specific ZONESET entries in a global zone, you can use these commands:

```
SET       BDY(GLOBAL)        /* Set to global zone.       */.
LIST      ZONESET(ZST2       /* List only these           */
                ZST4)        /* entries.                  */.
```

The format of the LIST output for each ZONESET entry is the same for both of these commands. The only difference is the number of ZONESET entries listed.

Figure 62 on page 296 shows an example of LIST output for ZONESET entries.

```
PAGE nnnn  - NOW SET TO zzzzzz ZONE nnnnnnn  DATE mm/dd/yy TIME hh:mm:ss SMP/E 36.nn SMPLIST OUTPUT

GLOBAL                   ZONESET ENTRIES


ZST2           ZONE    = ZONE21   ZONE22   ZONE23

ZST4           ZONE    = ZONE41   ZONE42   ZONE43
```

Figure 62. ZONESET entry: sample LIST output

UCLIN Examples

You can use the ADD, REP, and DEL UCL statements to change subentries in a ZONESET entry. When you use UCLIN to update a ZONESET entry, keep these points in mind:

- After the UCLIN changes are done, the ZONESET entry must contain at least a ZONE subentry. Otherwise, the entry contains so little information that SMP/E cannot use it.
- If a DEL statement deletes all the existing subentries in the entry, SMP/E deletes the entire entry.

Example: Defining a ZONESET entry

Assume that you have a system with four target zones: BP111, PROD111, BP999, and PROD999. BP111 and BP999 define two different base control programs. PROD111 and PROD999 are two versions of the same product that must be synchronized with each other and with their base control programs (PROD111 with BP111 and PROD999 with BP999). To keep service for these products at the same level, you can group BP111, PROD111, and PROD999 in one ZONESET (S111) and BP999, PROD999, and PROD111 in a second ZONESET (S999). The following UCL will define the ZONESET entries:

```
SET      BDY(GLOBAL)        /* Set to global zone.    */.
UCLIN                       /*                        */.
ADD      ZONESET(S111)      /* ZONESET S111.          */
         ZONE(BP111,        /* Include these target   */
             PROD111,       /* zones.                 */
             PROD999)       /*                        */
                            /*                        */.
ADD      ZONESET(S999)      /* ZONESET S999.          */
         ZONE(BP999,        /* Include these target   */
             PROD999,       /* zones.                 */
             PROD111)       /*                        */
                            /*                        */.
ENDUCL                      /*                        */.
```

Chapter 6. SMP/E CSI application programming interface

This chapter documents intended Programming Interfaces that allow the customer to write programs to obtain the services of IBM SMP/E for z/OS, V3R6.

This chapter describes the SMP/E CSI application program interface (GIMAPI), which provides access to data stored and used by SMP/E. This data is stored in a VSAM data set called the CSI (Consolidated Software Inventory). The purpose of GIMAPI is to provide you with read-only access to data maintained by SMP/E in the CSI.

Note: GIMAPI cannot be called from within an SMP/E user exit routine.

Overview of GIMAPI

GIMAPI is a program (load module) provided with IBM SMP/E for z/OS, V3R6 that can be called by a user-written application program to query the contents of the CSI.

GIMAPI uses data structures both to receive the query parameters from the caller and to return the query output to the caller. To request a query, you must set up the data structure containing the appropriate query parameters and call GIMAPI.

The general form of the GIMAPI call is:

```
GIMAPI(apicmd,parmptr,outptr,language,rc,cc,msgbuff)
```

where *apicmd* specifies the API command that GIMAPI is to process (either QUERY, FREE or VERSION). See "QUERY command" on page 299, "FREE command" on page 335 and "VERSION command" on page 336 for descriptions of these commands and a detailed description of how the remaining parameters *parmptr, outptr, language, rc, cc*, and *msgbuff*) are used on those commands.

QUERY command

Use the QUERY command to retrieve data from a zone or set of zones within a CSI. You specify the details of your query by setting parameters on the call to the QUERY command. These parameters determine what data is returned from the query. The data is returned to the calling program in a set of data structures.

QUERY command parameters

Here is an example of a call to the QUERY command:

```
GIMAPI('QUERY ',parmptr,outptr,'language',rc,cc,msgbuff)
```

QUERY
is a keyword that specifies that GIMAPI is to process a query request. Note that it is necessary to pad the string containing the QUERY command with blanks to bring it to a length of eight characters.

parmptr
A 4-byte pointer to an area of storage owned by the calling program that contains a data structure whose elements are the parameters for the QUERY command. A generic data structures definition can be found in "Data structures for QUERY command" on page 333.

outptr
A pointer variable that is to be set by GIMAPI with the address of an area of storage owned by GIMAPI that contains the output from the QUERY command. For the QUERY command, the output is a set

of linked lists that both contains and describes the requested data. See section "QUERY command output" on page 329 for more details on the QUERY command output.

language
A 3-byte character string specifying the language that QUERY command is to use when issuing messages. The valid values are:

ENU
US English

JPN
Japanese

Note: Specifying *language* as three blanks is also valid and is equivalent to specifying "ENU".

rc
A storage area owned by the calling program representing a 4-byte numeric variable. The value of the variable is set by GIMAPI to the return code of the QUERY command processing. The return code represents the severity of the condition raised by the GIMAPI processing. These codes may be returned:

0
Indicates the function terminated successfully. The *cc* value is also set to zero (0) when the return code is 0.

4
Indicates a warning condition was detected during execution of GIMAPI. The *cc* value given indicates the specific warning condition. When a warning occurs, the GIMAPI command continues processing. Appropriate warning messages are written to the message buffer.

8
Indicates an error condition was detected during execution of GIMAPI. The *cc* value given indicates the specific error condition. When an error occurs, the GIMAPI command terminates after storing the error condition in the message buffer. Other GIMAPI commands may be invoked from the calling program.

12
Indicates a severe error condition was detected during execution of GIMAPI. The *cc* value given indicates the specific error condition. When a severe error occurs, GIMAPI terminates after storing the severe condition in the message buffer. The calling program should not make any more calls to GIMAPI.

16
Indicates a terminating error condition was detected during execution of GIMAPI. The *cc* value given indicates the specific error condition. When a terminating error occurs, GIMAPI terminates after storing the condition in the message buffer. The calling program should not make any more calls to GIMAPI.

cc
A storage area owned by the calling program representing a 4-byte numeric variable. The value of the variable is set by GIMAPI to the condition code of the GIMAPI processing. The condition code, also known as a reason code, indicates the specific situation that occurred during GIMAPI processing. The condition code corresponds to the error message stored for the condition. For example, if error message GIM30400S is to be issued, the return code value is 8 and the condition code is 30400.

More than one error condition may occur during a call to GIMAPI. This is most likely during syntax checking, because GIMAPI checks every parameter for syntax errors. (GIMAPI will stop checking a parameter if it finds an error, so GIMAPI will not report any additional errors that may exist for that parameter.) The return code value set is the one with the highest severity of the raised conditions. If more than one "highest condition" is raised, the first condition of that severity encountered is the value returned as the condition code to the calling program.

GIMAPI sets the *cc* value to zero (0) for a successful run. The *rc* value is also zero (0) in this case.

msgbuff

A pointer variable that is set by GIMAPI to the head of a linked list of messages issued by the QUERY command processing. The elements of the list are ITEM_LIST structures. See section "Data structures for QUERY command" on page 333 for a description of that structure.

Building the QUERY parameters data structure

The QUERY command parameter list must include a pointer to a data structure (*parmptr*) that defines the query to be processed. This data structure, in turn, contains pointers to character strings that specify from where data is to be retrieved (that is, which CSI, zones, and entry types) and, optionally, a list of subentries and a *filter*, which allow you to specify conditions that will be used to determine which entries are retrieved.

The following parameters are part of the structure pointed to by the *parmptr* parameter of the QUERY command. All the parameters are required except the *subentry* and *filter* parameters.

csi

Input is accepted in mixed case. A character string that specifies the name of the global CSI to be searched by the QUERY command. The DDNAME used when allocating the data set is SMPCSI.

Input is accepted in mixed case. It is folded to uppercase when the command is processed.

csilen

A decimal number specifying the size of the character string containing the *csi* parameter. This value cannot be greater than 44.

A zero (0) value indicates that the calling program has already allocated a global CSI data set as the SMPCSI DDNAME. That allocation is used and the *csi* parameter is ignored.

zone

A character string that specifies the zones from which data is to be retrieved. You may enter one or more specific zone names separated by commas or blanks, or any of these values:

GLOBAL

Use the global zone

ALLTZONES

Use all target zones

ALLDZONES

Use all DLIB zones

*

Use all zones defined in the GLOBAL zone index

Except for the asterisk (*), these values can be used together and with specific zone names, provided that all values are separated by commas or blanks. The asterisk (*) must be used alone.

For example, these specifications are valid:

```
GLOBAL, ALLTZONES
ALLTZONES, ALLDZONES
ALLTZONES, MY1DLIB, MY2DLIB
ALLTZONES MY1DLIB, MY2DLIB
```

whereas this is not valid:

```
GLOBAL, *, ALLTZONES
```

A zone can be specified only once in the *zone* parameter. A specific DLIB zone name along with the ALLDZONES keyword or a specific target zone name along with the ALLTZONES keyword does not cause a message to be issued. However, processing for the duplicated zone names will be done only once.

Spaces can be used freely between values and commas.

Input is accepted in mixed case. It is folded to uppercase when the command is processed.

The specific zone name can be a ZONESET name. The command will first check whether the name is a ZONESET. If not, it will check the ZONEINDEX.

zonelen
A decimal number specifying the size of the character string containing the *zone* parameter.

entry
A character string that indicates the entry types from the specified zones to be searched. You can specify one or more entries, separated by commas or blanks, or you can specify an asterisk (*) to indicate that all entries are to be searched. If an asterisk is used, no other values may be specified.

An entry can be specified only once in the *entry* parameter.

Spaces can be used freely between entry names and commas.

Input is accepted in mixed case. It is folded to uppercase when the command is processed.

The entry types specified must be valid CSI entry types. Refer to "Valid entry types" on page 305 for a list of the valid types.

entrylen
A decimal number specifying the size of the character string containing the *entry* parameter.

subentry
Each CSI entry has a set of subentries associated with it. The *subentry* parameter is a pointer to a character string used to indicate the subentries for which data is retrieved. Several subentries can be specified separated by commas or blanks, or you can specify an asterisk (*) to indicate that all subentries are to be searched. If an asterisk is used, no other values may be specified.

A subentry can be specified only once in the *subentry* parameter.

Spaces can be used freely between subentry names and commas.

Input is accepted in mixed case. It is folded to uppercase when the command is processed.

Note: The *entry* parameter allows multiple entries to be specified. It may be that a subentry to be returned does not apply to all the specified entries. In this case, the subentry is simply ignored for a particular entry to which it does not apply.

Refer to "Valid subentry types" on page 306 for more details.

subentrylen
A decimal number specifying the size of the character string containing the *subentry* parameter.

filter
A character string that specifies the set of conditions with which to limit the set of entries being retrieved. A condition is in the form

```
subentry   operator   'value'
```

For example, FMID = 'HP10230' or INSTALLDATE >= '07203'.

Input is accepted in mixed case. The subentry types are folded to uppercase when the command is processed. The fixed values are not changed.

Filter is an optional parameter. When no filtering is desired, set the length field (filterlen) to zero (0).

Refer to "Filter parameter syntax" on page 302 for a detailed description of the syntax of the *filter* parameter.

filterlen
A decimal number specifying the size of the character string containing the *filter* parameter.

Filter parameter syntax

The *filter* parameter is used to limit the entries returned for a particular query. You specify a filter by listing a set of conditions. These conditions specify a comparison between two values; a subentry value in the CSI is compared with a fixed value specified in the condition.

Multiple filter conditions may be specified using the operators, & (logical AND) and | (logical OR). Parentheses may be used to group search conditions to ensure a desired evaluation sequence for the filter. The parentheses, when used, may be preceded by one or more blanks and also may be followed by one or more blanks.

When a subentry name represents a list of values, such as PRE, which is a list of all the prerequisite SYSMODs, each value in the list from the CSI is compared to the fixed value. If any match, the condition is evaluated as true.

These subentry types may be specified in the *filter* parameter only with the = or != operators and a null value (for example, ASMIN='' or LEPARM!=''):

- ASMIN
- HFSPARM
- HOLDDATA
- LECNTL
- LEPARM
- LMODALIAS
- LMODSYMLINK
- ORDERSERVER
- UTILPARM
- ZDESC

The fixed value must be enclosed within single quotation marks. Any set of characters may be used in the fixed value. If the fixed value is shorter than the subentry value, GIMAPI left-justifies the fixed value and pads it with blanks before doing the comparison.

Any apostrophes specified as part of a DESCRIPTION, VENDOR, or URL subentry (excluding delimiters) must be doubled. Double apostrophes count as two characters in the filter length.

Blanks may be freely used to separate the subentry, operator, and fixed value.

The values of some subentry types are lists of composite values. An example of this is the CIFREQ parameter whose value is *causer,requisite*. The fixed value for conditions containing these types of subentries must match the format of the output provided by the QUERY command. For example:

```
CIFREQ='PTF0001,PTF0002'
```

Refer to "Valid subentry types" on page 306 for the format of each composite subentry type.

A null value for a subentry is specified by using two single quotation marks with no blanks between them(''). This can be used to find all entries whose value is not blank, for example, PATH!=''.

The operators that can be used to compare a subentry value to a fixed value, or to join a set of conditions together, are:

Operator
    Function

=
    Equal to

!=
    Not equal to

<
    Less than

<=
    Less than or equal to

>
    Greater than

>=

Greater than or equal to

|

Logical OR, at least one of two conditions are true

&

Logical AND, both conditions must be true

Subentry names that may appear to the left of operators are the subentries defined in "Valid subentry types" on page 306. Potential values for each subentry can be found in Chapter 5, "SMP/E data set entries," on page 141. When a subentry is an indicator, such as PROTECT for the DDDEF entry, its value is either YES or NO.

A requested subentry may not have a value for a particular entry. This value is considered null. This may occur for one of the following reasons:

- The subentry has not been set for the entry.
- The subentry does not apply to that entry.

An example of the latter case is a request for the SYSLIB subentry with the SYSMOD entry type. SYSMOD entries do not contain SYSLIB subentries, causing a request for such a value to always return a null value.

When the not equal sign is used with some value within a filter, for example SYSLIB!='MACLIB', the case where no data is found for the subentry is considered a match, because "MACLIB" does not equal a null value. Therefore, if this filter is used when requesting the SYSMOD entry type, all the SYSMODs in the selected zones are returned. To exclude entries having no value, write the filter as (SYSLIB!='MACLIB' & SYSLIB!='').

The >, >=, < and <= operators can be used only with the subentry types that have date, time, or return code values, which are:

- HOLDDATE
- INSTALLDATE
- INSTALLTIME
- RC
- RECDATE
- RECTIME
- RESDATE
- RESTIME
- UCLDATE
- UCLTIME

The >, >=, < and <= operators have no meaning when applied to a subentry that has a null value. GIMAPI will always resolve such comparisons to false.

A date fixed value must use the format *yyddd*. It must be five numeric characters. For example, specify 01365 to represent December 31, 2001.

A time fixed value must use the format *hh:mm:ss*, where *hh* is a two-digit decimal number in the range 00 through 23, and *mm* and *ss* are both two-digit decimal numbers in the range 00 through 59. (That is, in 24-hour clock format.)

The *filter* parameter is made up of a combination of these parts:

- Subentry name
- Operator
- Fixed value
- Conjunction (& or |)
- Left parenthesis

- Right parenthesis

The following list indicates for each part, what other part or parts can follow it. Any other sequence of parts is considered a syntax error.

Filter part
> **Can be followed by**

Subentry
> Operator

Operator
> Fixed value

Fixed value
> Conjunction, Right parenthesis, End of filter

Conjunction
> Subentry Left parenthesis

Left parenthesis
> Subentry, Left parenthesis

Right parenthesis
> Conjunction, Right parenthesis, End of filter

Valid entry types

Table 8 on page 305 lists the values that are valid as *entry* parameters and the zones to which they are applicable.

Table 8. Valid entry values			
Entry type	**Global zone**	**DLIB zone**	**Target zone**
ASSEM		Valid	Valid
Data element[1]		Valid	Valid
DDDEF	Valid	Valid	Valid
DLIB		Valid	Valid
DLIBZONE		Valid	
ELEMENT[2]		Valid	Valid
FEATURE	Valid		
FMIDSET	Valid		
GLOBALZONE	Valid		
Hierarchical file system element[3]		Valid	Valid
HOLDDATA	Valid		
JAR		Valid	Valid
LMOD		Valid	Valid
MAC		Valid	Valid
MOD		Valid	Valid
OPTIONS	Valid		
ORDER	Valid		
PRODUCT	Valid		

Table 8. Valid entry values (continued)

Entry type	Global zone	DLIB zone	Target zone
PROGRAM		Valid	Valid
SRC		Valid	Valid
SYSMOD	Valid	Valid	Valid
TARGETZONE			Valid
UTILITY	Valid		
ZONESET	Valid		

Note:

1. The data element entry type has many possible values. Refer to "Data element entry (distribution and target zone)" on page 158 in Chapter 5, "SMP/E data set entries," on page 141 for the complete list. Any one of these values can be supplied as part of the *entry* parameter.

2. ELEMENT is a pseudo-entry value indicating that JAR, MOD, MAC, PROGRAM, SRC, and all data element and hierarchical file system element entries are returned.

3. The hierarchical file system element entry type has many possible values. Refer to "Hierarchical File System Element Entry (Distribution and Target Zone)" in Chapter 5 for the complete list. Any one of these values can be supplied as part of the *entry* parameter.

When you specify an entry that is not applicable to a particular zone, the result is simply that no data is returned. This is not an error condition.

Valid subentry types

For each entry type there is a set of valid subentry types that may be specified as part of the *subentry* and *filter* parameters. The following sections contain a table for each entry. Each table lists the valid subentries for that entry. The subentry types generally correspond to the types used for UCLIN processing. However, some additional types are used for GIMAPI processing (for example, the SMODTYPE subentry of the SYSMOD entry type). This specifies what type the SYSMOD is and is used instead of having a separate subentry for each SYSMOD type.

Note:

1. ENAME and zone name are automatically returned for each instance of an entry as part of the ENTRY structure so they do not have to be requested on the *subentry* parameter. ENAME can be used on the *filter* parameter.

   For a PRODUCT entry, ENAME has no meaning and is always blank. PRODID and VRM values are automatically returned instead of ENAME for PRODUCT entries. PRODID and VRM can be used on the *filter* parameter.

2. A maximum length of (*) indicates that a linked list of that subentry is returned to the caller. Details of the *return* structures passed back to the caller can be found in "Data structures for QUERY command" on page 333.

ASSEM

Table 9. Valid subentries for the ASSEM entry

Subentry name	Maximum length (decimal)	Description
ASMIN	*	A linked list of records containing the assembler statements that were saved for a module during JCLIN processing. Each record may be up to 80 bytes in length.

Table 9. Valid subentries for the ASSEM entry (continued)		
Subentry name	**Maximum length (decimal)**	**Description**
ENAME	8	Name of the entry.
LASTUPD	7	Cause of last change to assembler entry.
LASTUPDTYPE	8	Indicates how the entry was last changed.

Data element

Table 10. Valid subentries for the data element entries		
Subentry name	**Maximum length (decimal)**	**Description**
ALIAS	*	A linked list containing alias names associated with the data element. Each alias name may be up to 8 bytes in length.
DISTLIB	8	Identifies the ddname of the distribution library for the data element
ENAME	8	Name of the entry.
FMID	7	Specifies the functional owner of the data element.
LASTUPD	7	Cause of last change to this data element entry.
LASTUPDTYPE	3	Indicates how the entry was last changed.
RMID	7	Identifies the last SYSMOD that replaced this data element.
SYSLIB	8	Identifies the ddname of the target library for the data element.

DDDEF

Table 11. Valid subentries for the DDDEF entry		
Subentry name	**Maximum length (decimal)**	**Description**
CONCAT	*	A linked list containing the names of the DDDEF entries existing in the same zone that should be concatenated during SMP/E processing. Each entry name may be up to 8 bytes in length.
DATACLAS	8	Specifies the name of the data class to be used for allocating a new data set managed by SMS.
DATASET	44	Specifies the name of the data set to be allocated.
DIR	4	Specifies the number of directory blocks to allocate.
DISP	7	Specifies the final disposition of the data set. • CATALOG • DELETE • KEEP
DSNTYPE	7	Specifies the type of partitioned data set to be created. • LIBRARY • PDS

Table 11. Valid subentries for the DDDEF entry (continued)

Subentry name	Maximum length (decimal)	Description
DSPREFIX	26	Specifies the data set prefix to be used to construct the full data set name for SMPTLIB data sets. GLOBAL zone only.
ENAME	8	Name of the entry.
INITDISP	3	Specifies the initial disposition of the data set. • MOD • NEW • OLD • SHR
MGMTCLAS	8	Specifies the name of the management class to be used for allocating a new data set managed under SMS.
PATH	255	Identifies the name of the path to be allocated in a UNIX file system.
PROTECT	3	Indicates whether RACF® PROTECT option should be used when a new data set is allocated. (YES or NO).
SPACE	9	Specifies the primary and secondary space allocation in the format *nnnn,nnnn*.
STORCLAS	8	Specifies the name of a storage class used for allocating a new data set managed by SMS.
SYSOUT	1	Specifies the output class for SYSOUT data sets.
UNIT	8	Specifies the UNIT type the data set resides on if it is not cataloged.
UNITS	12	Specifies the space units to be used in allocating the data set. • BLK(size) • CYL • TRK
VOLUME	*	A linked list containing volume serial numbers of the device the data set resides on if it is not cataloged. Each volume serial number may be up to 6 bytes in length.
WAITFORDSN	3	Indicates whether SMP/E should wait for the data set to be allocated if the volume is not mounted or if the data set is already in use. (YES or NO).

DLIB

Table 12. Valid subentries for the DLIB entry

Subentry name	Maximum length (decimal)	Description
ENAME	8	Name of the entry.
LASTUPD	7	Cause of last change to DLIB entry.
LASTUPDTYPE	3	Indicates how the entry was last changed.

Table 12. Valid subentries for the DLIB entry (continued)

Subentry name	Maximum length (decimal)	Description
SYSLIB	*	A linked list containing the ddnames of the target libraries into which the distribution library should be copied. Each ddname may be up to 8 bytes in length.

DLIBZONE

Table 13. Valid subentries for the DLIBZONE entry

Subentry name	Maximum length (decimal)	Description
ACCJCLIN	3	Indicates whether JCLIN is to be saved in the distribution zone whenever a SYSMOD containing inline JCLIN is accepted. (YES or NO).
ENAME	8	Name of the entry.
OPTIONS	8	The name of the OPTIONS entry in the global zone that should be used when processing this distribution zone.
RELATED	7	The name of the target zone to which this distribution zone is related.
SREL	*	A linked list containing the names of the system releases supported within this distribution zone. Each name may be up to four bytes in length.
UPGLEVEL	9	Indicates the highest SMP/E release level that is allowed to make incompatible changes to the zone. If SMP/E attempts to make an incompatible change to a zone and the release level of SMP/E is higher than the UPGRADE level in that zone, then the incompatible change will not be made. This value is in the form *vr.ppp*, where *vr* represents the version and release of SMP/E and *ppp* represents the PTF level of SMP/E. The PTF level may be 2 to 3 characters in length.
ZDESC	500	A user-written description for this zone.

FEATURE

Table 14. Valid subentries for the FEATURE entry

Subentry name	Maximum length (decimal)	Description
DESCRIPTION	64	Describes the software feature in readable text.
ENAME	8	Name of the entry.
FMID	*	A linked list containing 7-byte names of function SYSMODS that make up the feature.
PRODUCT	17	The product identifier and version, release, and modification level (*prodid,vv.rr.mm*) of the feature's associated product.
RECDATE	5	Specifies the date on which this feature was received.
RECTIME	8	Specifies the time at which the ++FEATURE MCS was received.
REWORK	8	Specifies the level of this FEATURE entry, which was received again for minor changes.

Table 14. Valid subentries for the FEATURE entry (continued)

Subentry name	Maximum length (decimal)	Description
UCLDATE	5	Specifies the date on which this FEATURE entry was last modified through UCLIN.
UCLTIME	8	Specifies the time at which this FEATURE entry was last modified through UCLIN.

FMIDSET

Table 15. Valid subentries for the FMIDSET entry

Subentry name	Maximum length (decimal)	Description
ENAME	8	Name of the entry.
FMID	*	A linked list containing 7-byte names of the function SYSMODS (that is, FMIDs) that are to be part of this FMIDSET.

GLOBALZONE

Table 16. Valid subentries for the GLOBALZONE entry

Subentry name	Maximum length (decimal)	Description
ENAME	8	Name of the entry.
FMID	*	A linked list containing 7-byte names of the function SYSMODS for which SMP/E is to receive service.
OPTIONS	8	The name of the OPTIONS entry in the global zone that should be used when processing this global zone.
SREL	*	A linked list containing the names of the system releases supported within this global zone. Each name may be up to four bytes in length.
UPGLEVEL	9	Indicates the highest SMP/E release level that is allowed to make incompatible changes to the zone. If SMP/E attempts to make an incompatible change to a zone and the release level of SMP/E is higher than the UPGRADE level in that zone, then the incompatible change will not be made. This value is in the form *vr.ppp*, where *vr* represents the version and release of SMP/E and *ppp* represents the PTF level of SMP/E. The PTF level may be 2 to 3 characters in length.
ZDESC	500	A user-written description for this zone.
ZONEINDEX	*	A linked list of entries containing the names, data set names, and types of all the target zones and distribution zones associated with this global zone. The format of each entry is *name,dsn,type*. Each entry in the list may be up to 59 bytes in length.

Hierarchical file system elements

Table 17. Valid subentries for hierarchical file system element entries

Subentry name	Maximum length (decimal)	Description
DISTLIB	8	Specifies the DDNAME of the distribution library for the hierarchical file system element.
ENAME	8	Name of the entry.
FMID	7	Specifies the functional owner of the hierarchical file system element.
HFSPARM	300	Specifies a character string that is to be passed to the hierarchical file system copy utility as an execution-time parameter.
INSTMODE	6	Indicates the installation mode to be used when the HFS copy utility is invoked to install the element into a UNIX system. • BINARY • TEXT
LASTUPD	7	Cause of last change to this hierarchical file system element entry.
LASTUPDTYPE	3	Indicates how the entry was last changed.
LINK	*	A linked list containing the alternate names by which this element can be known in a UNIX file system. Each alternate name may be up to 1023 bytes in length.
RMID	7	Identifies the last SYSMOD that replaced this hierarchical file system element.
SHSCRIPT	17	Specifies a shell script, along with the point in processing in which SMP/E passes control to the shell script. The format is *script_name,*PRE,POST, where PRE and POST are optional.
SYMLINK	*	A linked list containing the names of symbolic links associated with a hierarchical file system element. Each symbolic link may be up to 1023 bytes in length.
SYMPATH	*	A linked list containing the path values for symbolic links associated with a hierarchical file system element. Each path value may be up to 1023 bytes in length.
SYSLIB	8	Specifies the ddname of the target library within a UNIX file system for the element.

HOLDDATA

Table 18. Valid subentries for the HOLDDATA entry

Subentry name	Maximum length (decimal)	Description
ENAME	7	Name of the entry. This is the name of the held SYSMOD.
HOLDCLASS	7	A 1- to 7-character string indicating an alternative reason to release an exception SYSMOD for processing.
HOLDDATA	*	A linked list containing records that contain the entire HOLDDATA entry. Each record may be up to 80 bytes in length.
HOLDDATE	5	Specifies the date that the ++HOLD statement was generated.

Table 18. Valid subentries for the HOLDDATA entry (continued)

Subentry name	Maximum length (decimal)	Description
HOLDFIXCAT	*	A linked list that contains Fix Category values. Each value can be up to 64 characters in length.
HOLDFMID	7	Specifies the FMID to which the held SYSMOD is applicable.
HOLDREASON	7	A 1- to 7-character string used to help you identify the reason why the SYSMOD has been put into exception SYSMOD status.
HOLDRESOLVER	7	Specifies the SYSMOD that resolves the exception condition, if known.
HOLDTYPE	9	Specifies the hold category into which the SYSMOD has been put. ERROR, FIXCAT, SYSTEM, or USER are the only valid values.

Refer to the "++HOLD MCS" section for specific details for each of the subentries listed in the HOLDDATA entry.

JAR

Table 19. Valid subentries for the JAR entry

Subentry name	Maximum length (decimal)	Description
DISTLIB	8	Specifies the ddname of the distribution library for the JAR element.
ENAME	8	Name of the entry.
FMID	7	Specifies the functional owner of the JAR element.
HFSPARM	300	Specifies a character string that is to be passed to the hierarchical file system copy utility as an execution-time parameter.
JARPARM	300	Specifies a character string that is to be passed to the jar command when updating the JAR file.
LASTUPD	7	Cause of last change to this JAR element entry.
LASTUPDTYP	3	Indicates how the entry was last changed.
LINK	*	A linked list containing alternate names by which this element can be known in a UNIX file system. Each alternate name may be up to 1023 bytes in length.
RMID	7	Identifies the last SYSMOD that replaced this JAR element.
SHSCRIPT	17	Specifies a shell script, along with the point in processing in which SMP/E passes control to the shell script. The format is script_name,PRE,POST where PRE and POST are optional.
SYMLINK	*	A linked list containing the names of symbolic links associated with this element. Each symbolic link may be up to 1023 bytes in length.
SYMPATH	*	A linked list containing the path values for symbolic links associated with this element. Each path value may be up to 1023 bytes in length.
SYSLIB	*	Specifies the ddname of the target library within a UNIX file system for the element.
UMID	*	A linked list containing 7-byte names of all SYSMODs that have updated this JAR file since it was last replaced.

LMOD

Subentry name	Maximum length (decimal)	Description
CALLLIBS	*	A linked list containing the names of the DDDEF entries, existing in the same zone, that compose the SYSLIB allocation to be used when the load module is link-edited. Each entry name may be up to 8 bytes in length.
COPIED	3	A special SMP/E indicator meaning that the load module was copied during system generation, and that there is a one-to-one correspondence between the distribution library module (the MOD entry) and the target system load module (the LMOD entry) (YES or NO).
ENAME	8	Name of the entry.
LASTUPD	7	Cause of last change to this LMOD entry.
LASTUPDTYPE	3	Indicates how the entry was last changed.
LEPARM	*	A character string containing the link-edit attributes that must be used when this load module is link-edited. This is not a linked list.
LECNTL	*	A linked list containing records of link-edit control cards necessary to relink-edit this load module. Each record may be up to 80 bytes in length.
LMODALIAS	*	A linked list containing the names of the aliases associated with this LMOD. Each alias name may be up to 1023 bytes in length.
LMODSYMLINK	*	A linked list containing the names of symbolic links associated with a load module. The values are the values specified on ALIAS (SYMLINK,*symlink_value*) link-edit control statements. Each symbolic link may be up to 1023 bytes in length.
MODDEL	*	A linked list containing the names of modules that were once part of this load module, but have been deleted. Each module name may be up to 8 bytes in length.
RC	2	Highest acceptable link edit return code for this load module.
SIDEDECKLIB	8	Specifies the ddname of the library where the load module's definition side deck resides.
SYSLIB	*	A linked list containing the ddnames of the target system libraries in which this load module resides. Each ddname may be up to 8 bytes in length.

Table 20. Valid subentries for the LMOD entry

Table 20. Valid subentries for the LMOD entry (continued)

Subentry name	Maximum length (decimal)	Description
UTIN	*	A linked list containing utility input (UTIN) subentries, each of which contains the filename of the input to be included during a link-edit operation, and the library ddname where the file resides. The format is *filename,ddname*. Each UTIN subentry may be up to 1032 bytes in length. (The filename may be up to 1023 characters in length. The ddname may be up to 8 characters in length. And then there is the delimiting comma.) **Note:** It should be noted that a comma is a valid character within the filename value. The comma is also the delimiter separating the filename from the ddname. Therefore, the utility input data should be scanned from the last byte to the first byte in order to accurately find the comma that separates the ddname from the *filename*.
XZMOD	*	A linked list containing the XZMOD subentries containing the names of the modules that were added to the load module by the LINK MODULE command. The name of the zone supplying each module is also indicated. The format is *module,zone*. Each XZMOD subentry may be up to 16 bytes in length.
XZMODP	3	Indicates whether the load module contains one or more modules from another zone and that XZMOD subentries exist in this LMOD entry. (YES or NO).

MAC

Table 21. Valid subentries for the MAC entry

Subentry name	Maximum length (decimal)	Description
DISTLIB	8	Specifies the ddname of the macro distribution library.
ENAME	8	Name of the entry.
FMID	7	Specifies the functional owner of this macro.
GENASM	*	A linked list containing the names of the assemblies that have to be done during APPLY each time this macro is modified. Each name may be up to 8 bytes in length.
LASTUPD	7	Cause of last change to this MAC entry.
LASTUPDTYPE	3	Indicates how the entry was last changed.
MALIAS	*	A linked list containing the alias names for this macro. Each name may be up to 8 bytes in length.
RMID	7	Identifies that last SYSMOD that replaced this macro.
SYSLIB	8	Specifies the ddname of the target library for the macro.
UMID	*	A linked list containing the 7-byte names of all those SYSMODs that updated this macro since it was last replaced.

MOD

Table 22. Valid subentries for the MOD entry

Subentry name	Maximum length (decimal)	Description
ASSEMBLE	3	Indicates whether the source for this module must always be assembled (YES or NO).
CSECT	*	A linked list containing the names of the CSECTs present in this module. Each name may be up to 8 bytes in length.
DALIAS	8	Specifies the alias name for the module, where the alias exists only in the distribution library.
DISTLIB	8	Specifies the ddname of the module distribution library.
ENAME	8	Name of the entry.
FMID	7	Identifies the functional owner of this module.
LASTUPD	7	Cause of last change to this MOD entry.
LASTUPDTYPE	3	Indicates how the entry was last changed.
LEPARM	*	A character string containing the link-edit attributes that must be used when this module is link-edited. This is not a linked list.
LMOD	*	A linked list containing the names of the load modules into which this module was copied or link-edited. Each name may be up to 8 bytes in length.
RMID	7	Identifies the last SYSMOD that replaced this module.
RMIDASM	3	Specifies that the last replacement (RMID) to the module was done by a SYSMOD that caused an assembly of the module as a result of a source or macro modification (YES or NO).
TALIAS	*	A linked list containing the alias names for the module, where the alias exists in the distribution library and, for copied modules, also in the target library. Each name may be up to 8 bytes in length.
UMID	*	A linked list containing 7-byte names of all those SYSMODs that have updated this module since it was last replaced.
XZLMOD	*	A linked list containing XZLMOD subentries that contain the name of each load module in another zone into which this module was added by the LINK MODULE command. The name of the zone containing each load module is also indicated. Each subentry has the format *load module,zone* and may be up to 16 bytes in length.
XZLMODP	3	Indicates whether this module has been linked into one or more load modules controlled by a different target zone, and that XZLMOD subentries exist in this MOD entry (YES or NO).

OPTIONS

Table 23. Valid subentries for the OPTIONS entry

Subentry name	Maximum length (decimal)	Description
AMS	8	Specifies the name of the UTILITY entry that SMP/E is to use to obtain information when calling the Access Method Services (AMS) utility.

Table 23. Valid subentries for the OPTIONS entry (continued)

Subentry name	Maximum length (decimal)	Description
ASM	8	Specifies the name of the UTILITY entry that SMP/E is to use to obtain information when calling the assembler utility.
CHANGEFILE	3	Specifies whether library change file records should be created during APPLY and RESTORE processing (YES or NO)
COMP	8	Specifies the name of the UTILITY entry that SMP/E is to use to obtain information when calling the utility program to compress data sets.
COMPACT	3	Specifies whether SMPPTS members should be compacted during RECEIVE and GZONEMERGE command processing (YES or NO)
COPY	8	Specifies the name of the UTILITY entry that SMP/E is to use to obtain information when calling the copy utility.
DSPREFIX	26	Specifies the data set prefix to be used to construct the full data set name when SMPTLIB data sets are being allocated for RELFILEs.
DSSPACE	14	Specifies the primary and secondary space allocation (in tracks) and the number of directory blocks to be allocated for each SMPTLIB data set. It is in the format *primary,secondary,directory, nnnn,nnnn,nnnn.*
ENAME	8	Name of the entry.
EXRTYDD	*	A linked list containing ddnames ineligible for retry processing after an x37 abend. Each ddname may be up to 8 bytes in length.
FIXCAT	*	A linked list that contains the Fix Categories whose HOLDDATA is considered during APPLY, ACCEPT, and REPORT MISSINGFIX command processing. Each value can be up to 64 bytes in length.
HFSCOPY	8	Specifies the name of the UTILITY entry that SMP/E is to use to obtain information when calling the hierarchical file system copy utility.
IOSUP	8	Specifies the name of the UTILITY entry that SMP/E is to use to obtain information when calling the IEHIOSUP utility program to process maintenance for an OS/VS1 system.
LKED	8	Specifies the name of the UTILITY entry that SMP/E is to use to obtain information when calling the link-edit utility.
MSGFILTER	3	Specifies whether messages issued to SMPOUT should be filtered.
MSGWIDTH	7	Specifies the width in which the message issued to SMPOUT should be formatted.
NOPURGE	3	Indicates whether after SMP/E accepts SYSMODS, it should not delete the associated global zone SYSMOD entries, SMPPTS MCS entries, or SMPTLIB data sets (YES or NO).
NOREJECT	3	Specifies that the global zone SYSMOD entry and the associated MCS entry should not be deleted after the SYSMOD is restored. (YES or NO).
ORDERRET	4	Indicates the retention period, in days, that ORDER entries are kept in the global zone before being deleted.
PAGELEN	4	Specifies the page length for SMPHRPT, SMPLIST, SMPOUT, and SMPRPT.
PEMAX	4	Specifies the maximum number of subentries that can be present in any CSI entry.

Table 23. Valid subentries for the OPTIONS entry (continued)

Subentry name	Maximum length (decimal)	Description
RECZGRP	*	A linked list containing 8-byte zone and zoneset names eligible for APPLYCHECK and ACCEPTCHECK processing during the SYSMOD selection phase of RECEIVE.
RECEXCGRP	*	A linked list containing 8-byte zone and zoneset names **not** eligible for APPLYCHECK and ACCEPTCHECK processing during the SYSMOD selection phase of RECEIVE.
RETRY	8	Specifies the name of the UTILITY entry that SMP/E is to use to obtain information when calling the utility program to compress a data set after an x37 abend.
RETRYDDN	*	A linked list containing ddnames eligible for retry processing after an x37 abend. Each ddname may be up to 8 bytes in length.
SAVEMTS	3	Indicates whether MTSMAC entries should be deleted from the SMPMTS after the SYSMODs that affect those macros have been successfully accepted (YES or NO).
SAVESTS	3	Indicates whether STSSRC entries should be deleted from the SMPSTS after the SYSMODs that affect the source have been successfully accepted (YES or NO).
SUPPHOLD	*	A linked list containing 7-byte HOLD Reason IDs for which the HOLDDATA card image is not displayed in Unresolved HOLD Reason Report and Bypassed HOLD Reason Report for Apply and Accept command processing, and the SYSMOD Comparison HOLDDATA Report for REPORT SYSMODS command processing.
UPDATE	8	Specifies the name of the UTILITY entry that SMP/E is to use to obtain information when calling the update utility.
ZAP	8	Specifies the name of the UTILITY entry that SMP/E is to use to obtain information when calling the superzap utility.

ORDER

Table 24. Valid subentries for the ORDER entry

Subentry name	Maximum length (decimal)	Description
APARS	*	A linked list containing 7-byte names of all the APARs for which resolving PTFs are to be in the order.
CONTENT	8	Indicates the content selected for the order. The value may be **ALL**, **APARS**, **CRITICAL**, **HOLDDATA**, **PTFS**, or a Recommended Service Update SOURCEID value (**RSU**$yymm$).
DOWNLDATE	5	Indicates the date ($yyddd$) on which the order was downloaded and stored in the SMPNTS directory.
DOWNLTIME	8	Indicates the time ($hh:mm:ss$) at which the order was downloaded and stored in the SMPNTS directory.
ENAME	8	Name of the entry.
ORDERDATE	5	Indicates the date ($yyddd$) on which the order request was made to the server.

Table 24. Valid subentries for the ORDER entry (continued)

Subentry name	Maximum length (decimal)	Description
ORDERID	10	Specifies the order identifier assigned by the server when the order was created. This value is used to correlate ORDER entries in the global zone with orders processed by the server.
ORDERSERVER	*	A linked list containing the *<ORDERSERVER>* tags used to identify the server used to fulfill the order request.
ORDERTIME	8	Indicates the time (*hh:mm:ss*) at which the order request was made to the IBM Automated Delivery Request server.
PKGID	50	Specifies the order package id value used to create a package subdirectory within the SMPNTS directory to contain the order's package files.
PTFS	*	A linked list containing 7-byte names of all the PTF SYSMODs to be in the order.
STATUS	10	Specifies the current status of the ORDER entry. Possible values are **PENDING**, **DOWNLOADED**, and **ERROR**.
USERID	8	Specifies the userid associated with the address space where SMP/E was executing when the ORDER entry was created.
ZONES	*	A linked list containing 8-byte names of all the target zones used to define the scope of the software inventory associated with this order.

PRODUCT

Table 25. Valid subentries for the PRODUCT entry

Subentry name	Maximum length (decimal)	Description
DESCRIPTION	64	Specifies the name of this product.
PRODID	8	Product identifier
PRODSUP	*	A linked list containing the *prodid,vv.rr.mm* values for the products superseded by this product.
RECDATE	5	Specifies the date on which this product was received.
RECTIME	8	Specifies the time at which the ++PRODUCT MCS was received.
REWORK	8	Specifies the level of this PRODUCT, which was received again for minor changes.
SREL	*	A linked list containing 4-byte names of the system or subsystem releases on which the PRODUCT can be installed.
UCLDATE	5	Specifies the date on which this PRODUCT entry was last modified through UCLIN.
UCLTIME	8	Specifies the time at which this PRODUCT entry was last modified through UCLIN.
URL	256	Specifies a uniform resource locator (URL) that can be accessed to obtain additional information about this product.
VENDOR	64	Specifies the name of the vendor supplying the product.

Table 25. Valid subentries for the PRODUCT entry (continued)

Subentry name	Maximum length (decimal)	Description
VRM	8	Specifies the version, release, and modification level (*vv.rr.mm*) of this product.

Note: For a PRODUCT entry, the ENAME value is always blank, because the PRODID and VRM subentries, taken together, make up the "entry name" normally provided by ENAME. For each PRODUCT entry satisfying the QUERY request, the PRODID and VRM subentries are returned instead of ENAME.

PROGRAM

Table 26. Valid subentries for the PROGRAM entry

Subentry name	Maximum length (decimal)	Description
ALIAS	*	A linked list containing alias names associated with the program element. Each alias name may be up to 8 bytes in length.
DISTLIB	8	Identifies the ddname of the distribution library for the program element
ENAME	8	Name of the entry.
FMID	7	Specifies the functional owner of the data element.
LASTUPD	7	Cause of last change to this program element entry.
LASTUPDTYPE	3	Indicates how the entry was last changed.
RMID	7	Identifies the last SYSMOD that replaced this program element.
SYSLIB	8	Identifies the ddname of the target library for the program element.

SRC

Table 27. Valid subentries for the SRC entry

Subentry name	Maximum length (decimal)	Description
DISTLIB	8	Specifies the ddname of the distribution library for the source.
ENAME	8	Name of the entry.
FMID	7	Specifies the functional owner of this source.
LASTUPD	7	Cause of last change to this SRC entry.
LASTUPDTYPE	3	Indicates how the entry was last changed.
RMID	7	Identifies that last SYSMOD that replaced this source.
SYSLIB	8	Specifies the ddname of the target library for this source.
UMID	*	A linked list containing 7-byte names of all those SYSMODs that updated this source since it was last replaced.

SYSMOD (GLOBAL zone)

Table 28. Valid subentries for the SYSMOD entry (GLOBAL zone)

Subentry name	Maximum length (decimal)	Description
ACCID	*	A linked list containing 7-byte names of the distribution zones into which the SYSMOD has been successfully accepted.
APPID	*	A linked list containing the 7-byte names of the target zones into which the SYSMOD has been successfully applied.
Data Element[1]	*	A linked list containing the names of data element replacements contained in the SYSMOD. Each name may be up to 8 bytes in length.
DELETE[2]	*	A linked list containing 7-byte names of SYSMODs deleted by this SYSMOD.
DELLMOD	3	Indicates whether the SYSMOD contained a ++DELETE statement (YES or NO).
DESCRIPTION	64	Descriptive name associated with a SYSMOD
DLMOD	*	A linked list containing the names of load modules to be deleted by ++DELETE statements contained in this SYSMOD. Each name may be up to 8 bytes in length.
ELEMENT	*	A "pseudo-subentry" name that refers to all element types. It is the same as individually requesting each of these elements: • All data elements • All hierarchical file system elements • JAR and JARUPD • MAC and MACUPD • MOD • PROGRAM, SZAP, and XZAP • SRC and SRCUPD
ELEMMOV	3	Indicates whether the SYSMOD contained a ++MOVE statement (YES or NO).
EMOVE	*	A linked list containing names of elements and load modules to be moved by the ++MOVE statements contained in this SYSMOD. Each name may be up to 8 bytes in length.
ENAME	8	Name of the entry.
ERROR	3	Indicates whether an error has occurred during the processing of this SYSMOD (YES or NO).
FEATURE	*	A linked list containing the 8-byte names of software features that contain this SYSMOD.
FESN	7	Identifies the field engineering (FE) service number.
FMID	7	Identifies the function SYSMOD to which this SYSMOD is applicable.
Hierarchical File System Element[3]	*	A linked list containing the names of hierarchical file system element replacements (++_hfs_element_ statements) in the SYSMOD. Each name may be up to 8 bytes in length.

Table 28. Valid subentries for the SYSMOD entry (GLOBAL zone) (continued)

Subentry name	Maximum length (decimal)	Description
HOLDDATA	*	A linked list of the hold information associated with the SYSMOD. Each data value in the list is in the format *holdtype,reason,sysmod* with the following definitions: *holdtype* is one of the following values: • ERROR • FIXCAT • SYSTEM(EXT) • SYSTEM(INT) • USER *reason* is the 1- to 7-character reason ID. *sysmod* is the SYSMOD ID specified on the ++HOLD.
JAR	*	A linked list containing the names of JAR file replacements (++JAR statements) in the SYSMOD. Each name may be up to 8 bytes in length.
JARUPD	*	A linked list containing the names of JAR file updates (++JARUPD statements) in the SYSMOD. Each name may be up to 8 bytes in length.
JCLIN	3	Indicates whether this SYSMOD contained inline JCLIN (YES or NO).
MAC	*	A linked list containing the names of macro replacements (++MAC statements) in the SYSMOD. Each name may be up to 8 bytes in length.
MACUPD	*	A linked list containing the names of macro updates (++MACUPD statements) in the SYSMOD. Each name may be up to 8 bytes in length.
MOD	*	A linked list containing the names of modules replacements (++MOD statements) in the SYSMOD. Each name may be up to 8 bytes in length.
NPRE[2]	*	A linked list containing the 7-byte names of negative prerequisite SYSMODS-that is, SYSMODs that must not be present in the system at the same time as this SYSMOD.
PRE[2]	*	A linked list containing the 7-byte names of prerequisite SYSMODS-that is, SYSMODs that must be present before this SYSMOD can be installed.
PROGRAM	*	A linked list containing the names of program element replacements (++PROGRAM statements) in the SYSMOD. Each name may be up to 8 bytes in length.
RECDATE	5	Specifies the date on which this SYSMOD was received.
RECTIME	8	Specifies the time at which this SYSMOD was received.
RENLMOD	3	Indicates whether the SYSMOD contained a ++RENAME statement (YES or NO).
REQ[2]	*	A linked list containing the 7-byte names of requisite SYSMODS-that is, SYSMODs that must be installed concurrent with this SYSMOD.

Table 28. Valid subentries for the SYSMOD entry (GLOBAL zone) (continued)

Subentry name	Maximum length (decimal)	Description
REWORK	8	Identifies the level of the SYSMOD, which was received again for minor changes.
RLMOD	*	A linked list containing the names of the load modules to be renamed by ++RENAME statements in this SYSMOD. Each name may be up to 8 bytes in length.
SMODTYPE	8	Indicates the type of SYSMOD (APAR, FUNCTION, PTF, UNKNOWN, or USERMOD).
SOURCEID	*	A linked list containing the names of strings assigned to this SYSMOD during RECEIVE processing. Each name can be up to 64 characters in length.
SRC	*	A linked list containing the names of source replacements (++SRC statements) in the SYSMOD. Each name may be up to 8 bytes in length.
SRCUPD	*	A linked list containing the names of source updates (++SRCUPD statements) in the SYSMOD. Each name may be up to 8 bytes in length.
SREL[2]	*	A linked list containing the names of the system or subsystem releases on which the SYSMOD can be installed. Each name may be up to four bytes in length.
SUPING[2]	*	A linked list containing the 7-byte names of SYSMODs superseded by this SYSMOD.
SZAP	*	A linked list containing the names of module superzaps (++ZAP statements) in the SYSMOD. Each name may be up to 8 bytes in length.
TLIBPREFIX	26	Specifies the high-level data set name qualifier of the SMPTLIB data sets used to receive the SYSMOD, which was packaged in RELFILEs
VERSION[2]	*	A linked list containing 7-byte names of function SYSMODs that are versioned by this SYSMOD.

Note:

1. For the *data element* subentry, the data is returned to the caller in separate linked lists for each unique data element type requested.

2. When PRE, NPRE, REQ, DELETE, SUPING, VERSION, FMID, or SREL subentries are requested, a subentry structure of type 'VER' is returned, along with the necessary associated VER pseudo-subentries. See "VER" on page 334 for more details.

3. For the *hierarchical file system element* subentry the data is returned to the caller in separate linked lists for each unique hierarchical file system element type requested.

4. For SYSMODs that are only superseded, a value of "UNKNOWN" is returned for the SMODTYPE subentry.

SYSMOD

Subentry name	Maximum length (decimal)	Description
ACCEPT	3	Indicates whether the SYSMOD has been successfully accepted (YES or NO). This subentry appears only for DLIB zones.
APPLY	3	Indicates whether the SYSMOD has been successfully applied (YES or NO). This subentry appears only for target zones.
ASSEM	*	A linked list containing the names of the assemblies done during the installation of this SYSMOD. Each name may be up to 8 bytes in length.
BYPASS	3	Indicates whether the BYPASS operand was specified when this SYSMOD was installed (YES or NO).
CIFREQ	*	A linked list of entries containing CIFREQ data that list the conditional requisites that must be installed when this function SYSMOD is installed. Each CIFREQ entry has the format *causer,requisite* and may be up to 15 bytes in length.
DELBY	7	Specifies the SYSMOD that deleted this SYSMOD.
DELETE[2]	*	A linked list containing 7-byte names of SYSMODs deleted by this SYSMOD.
DELLMOD	3	Indicates whether the SYSMOD contained a ++DELETE statement (YES or NO).
DESCRIPTION	64	Descriptive name associated with a SYSMOD
DLMOD	*	A linked list containing the names of load modules to be deleted by ++DELETE statements contained in this SYSMOD. Each name may be up to 8 bytes in length.
Data Element[1]	*	A linked list containing the 8-byte names of data element replacements contained in the SYSMOD.
ELEMENT	*	A "pseudo-subentry" name that refers to all element types. It is the same as individually requesting each of these elements: • All data elements • All hierarchical file system elements • JAR and JARUPD • MAC and MACUPD • MOD • PROGRAM, SZAP, and XZAP • SRC and SRCUPD
ELEMMOV	3	Indicates whether the SYSMOD contained a ++MOVE statement (YES or NO).
EMOVE	*	A linked list containing the names of elements and load modules to be moved by the ++MOVE statements contained in this SYSMOD. Each name may be up to 8 bytes in length.
ENAME	8	Name of the entry.

Table 29. Valid subentries for the SYSMOD entry (DLIB and target zones)

Table 29. Valid subentries for the SYSMOD entry (DLIB and target zones) (continued)

Subentry name	Maximum length (decimal)	Description
ERROR	3	Indicates whether an error has occurred during the processing of this SYSMOD (YES or NO).
FEATURE	*	A linked list containing the 8-byte names of software features that contain this SYSMOD.
FESN	7	Identifies the field engineering (FE) service number.
FMID	7	Identifies the function SYSMOD to which this SYSMOD is applicable.
Hierarchical File System Element[3]	*	A linked list containing the names of hierarchical file system element replacements (++*hfs_element* statements) in the SYSMOD. Each name may be up to 8 bytes in length.
IFREQ	*	A linked list containing 7-byte names of conditional requisites that were installed with this SYSMOD.
INSTALLDATE	5	Specifies the date on which this SYSMOD was installed.
INSTALLTIME	8	Specifies the time at which this SYSMOD was installed.
JAR	*	A linked list containing the names of JAR file replacements (++JAR statements) in the SYSMOD. Each name may be up to 8 bytes in length.
JARUPD	*	A linked list containing the names of JAR file updates (++JARUPD statements) in the SYSMOD. Each name may be up to 8 bytes in length.
JCLIN	3	Indicates whether this SYSMOD contained inline JCLIN (YES or NO).
LASTSUP	7	Specifies the most recent SYSMOD that superseded this SYSMOD. All previous superseding SYSMODs are saved in the SUPBY subentry list.
LASTUPD	7	Cause of last change to this SYSMOD entry.
LASTUPDTYPE	3	Indicates how the entry was last changed.
MAC	*	A linked list containing the names of macro replacements (++MAC statements) in the SYSMOD. Each name may be up to 8 bytes in length.
MACUPD	*	A linked list containing the names of macro updates (++MACUPD statements) in the SYSMOD. Each name may be up to 8 bytes in length.
MOD	*	A linked list containing the names of module replacements (++MOD statements) in the SYSMOD. Each name may be up to 8 bytes in length.
NPRE[2]	*	A linked list containing the 7-byte names of negative prerequisite SYSMODS-that is, SYSMODs that must not be present in the system at the same time as this SYSMOD.
PRE[2]	*	A linked list containing the 7-byte names of prerequisite SYSMODS-that is, SYSMODs that must be present before this SYSMOD can be installed.

Table 29. Valid subentries for the SYSMOD entry (DLIB and target zones) (continued)

Subentry name	Maximum length (decimal)	Description
PROGRAM	*	A linked list containing the names of program element replacements (++PROGRAM statements) in the SYSMOD. Each name may be up to 8 bytes in length.
RECDATE	5	Specifies the date on which this SYSMOD was received.
RECTIME	8	Specifies the time at which this SYSMOD was received.
REGEN	3	Indicates how the SYSMOD was installed in the target libraries (YES or NO).
RENLMOD	3	Indicates whether the SYSMOD contained a ++RENAME statement (YES or NO).
REQ[2]	*	A linked list containing the 7-byte names of requisite SYSMODS-that is, SYSMODs that must be installed concurrent with this SYSMOD.
RESDATE	5	Specifies the date on which this SYSMOD was restored.
RESTIME	8	Specifies the time at which this SYSMOD was restored.
RESTORE	3	Indicates whether a restore attempt has been made for this SYSMOD (YES or NO).
REWORK	8	Identifies the level of the SYSMOD, which was received again for minor changes.
RLMOD	*	A linked list containing the names of the load modules to be renamed by ++RENAME statements in this SYSMOD. Each name may be up to 8 bytes in length.
SMODTYPE	8	Indicates the type of SYSMOD (APAR, FUNCTION, PTF, UNKNOWN, or USERMOD).
SOURCEID	*	A linked list containing the names of strings assigned to this SYSMOD during RECEIVE processing. Each name can be up to 64 characters in length.
SRC	*	A linked list containing the names of source replacements (++SRC statements) in the SYSMOD. Each name may be up to 8 bytes in length.
SRCUPD	*	A linked list containing the names of source updates (++SRCUPD statements) in the SYSMOD. Each name may be up to 8 bytes in length.
SUPBY	*	A linked list containing the 7-byte names of SYSMODs that superseded this SYSMOD. The most recent SYSMOD to supersede this SYSMOD is not included in the SUPBY list. It is saved in the LASTSUP field. Therefore, to simply determine if a SYSMOD is superseded or not, the LASTSUP subentry should be interrogated instead of SUPBY.
SUPING[2]	*	A linked list containing the 7-byte names of SYSMODs superseded by this SYSMOD.
SZAP	*	A linked list containing the 8-byte names of module superzaps (++ZAP statements) in the SYSMOD. Each name may be up to 8 bytes in length.

Table 29. Valid subentries for the SYSMOD entry (DLIB and target zones) (continued)

Subentry name	Maximum length (decimal)	Description
UCLDATE	5	Specifies the date on which this SYSMOD was last modified through UCLIN.
UCLTIME	8	Specifies the time at which this SYSMOD was last modified through UCLIN.
VERSION[2]	*	A linked list containing 7-byte names of function SYSMODs that are versioned by this SYSMOD.
XZAP	*	A linked list containing the names of module superzaps the SYSMOD (++ZAP statements) that contain an EXPAND statement (indicating that the module should be expanded before it is updated). Each name may be up to 8 bytes in length.

Note:

1. For the *data element* subentry, the data is returned to the caller in separate linked lists for each unique data element type requested.

2. When PRE, NPRE, REQ, DELETE, SUPING, or VERSION subentries are requested, a subentry structure of type 'VER' is returned for each unique version statement, along with its associated pseudo-subentries. See "VER" on page 334 for more details.

3. For the *hierarchical file system element* subentry the data is returned to the caller in separate linked lists for each unique hierarchical file system element type requested.

4. For SYSMODs that are only superseded, a value of "UNKNOWN" is returned for the SMODTYPE subentry.

TARGETZONE

Table 30. Valid subentries for the TARGETZONE entry

Subentry name	Maximum length (decimal)	Description
ENAME	8	Name of the entry.
OPTIONS	8	The name of the OPTIONS entry in the global zone that should be used when processing this target zone.
RELATED	7	The name of the distribution zone to which this target zone is related.
SREL	*	A linked list containing the names of the system releases supported within this target zone. Each name may be up to four bytes in length.
TIEDTO	*	A linked list containing the names of other target zones that either supplied modules for load modules controlled by this target zone or control load modules that have been link-edited with modules supplied by this target zone. Each name may be up to 8 bytes in length.
UPGLEVEL	9	Indicates the highest SMP/E release level that is allowed to make incompatible changes to the zone. If SMP/E attempts to make an incompatible change to a zone and the release level of SMP/E is higher than the UPGRADE level in that zone, then the incompatible change will not be made. This value is in the form *vr.ppp*, where *vr* represents the version and release of SMP/E and *ppp* represents the PTF level of SMP/E. The PTF level may be 2 to 3 characters in length.

Table 30. Valid subentries for the TARGETZONE entry (continued)

Subentry name	Maximum length (decimal)	Description
XZLINK	9	Specifies whether APPLY and RESTORE processing in another zone should automatically update load modules in this zone when cross-zone modules previously added to those load modules by the LINK command are changed. This subentry value can contain one of these values: • DEFERRED • AUTOMATIC
ZDESC	500	A user-written description for this zone.

UTILITY

Table 31. Valid subentries for the UTILITY entry

Subentry name	Maximum length (decimal)	Description
ENAME	8	Name of the entry.
LIST	3	Indicates whether member names should be listed when SMP/E invokes a copy utility to perform compress processing, retry processing, or element installation (YES or NO).
NAME	8	The name of the load module for the utility program that SMP/E is to call.
UTILPARM	100	Specifies the parameters to be passed to the utility program.
PRINT	8	Specifies the ddname that is to contain output from the utility.
RC	2	Specifies the maximum acceptable return code from this utility. This value can be from 0 to 16.

ZONESET

Table 32. Valid subentries for the ZONESET entry

Subentry name	Maximum length (decimal)	Description
ENAME	8	Name of the entry.
XZREQCHK	3	Indicates whether this ZONESET should be used when establishing the default zone group for the APPLY, ACCEPT, and RESTORE commands. (YES or NO).
ZONENAME	*	A linked list containing the 7-byte names of the target or distribution zones that are to be part of this ZONESET.

QUERY command processing

GIMAPI checks to see if the SMPCSI DDNAME is already allocated. If it is, GIMAPI notes the data set name. GIMAPI then determines if the caller expects GIMAPI to allocate the global CSI data set. The caller indicates this by setting the *csilen* parameter to zero. The data set name previously noted is used so a shared enqueue can be done. The enqueue is to protect GIMAPI processing from updates by any other SMP/E processing that might be going on at the same time.

If the *csilen* parameter is not zero, GIMAPI verifies that the *csi* parameter provided is a valid CSI data set name.

If the SMPCSI DDNAME has already been allocated, the data set name of the allocation must be the same as the *CSI* parameter value. If they match, GIMAPI treats it as if the calling program already did the allocation. The data set is not allocated again (and is not deallocated at the end of the command).

If the CSI data set must be allocated, GIMAPI dynamically allocates the data set with DISP=SHR.

GIMAPI continues by validating the rest of the required QUERY parameters that it has received. Syntax checking is done for each QUERY parameter received. If GIMAPI encounters an error, it stops syntax checking for that parameter list and moves on to begin syntax checking for the next QUERY parameter. If any errors are found, the QUERY command terminates after syntax checking is done. The command continues if warnings are the highest severity found.

GIMAPI enforces these syntax rules:

- When multiple values are specified on the *zone, entry* or *subentry* parameters, they must be appropriately separated by commas or spaces.
- When an asterisk (*) is used on the *zone*, *entry* or *subentry* parameters, no other values can be specified.
- Entry types entered on the *entry* parameter must be a valid CSI entry type.
- A specific value can be specified only once on the *zone, entry* or *subentry* parameters.
- A condition using the >, >=, < or <= operator with a subentry type that is not a date or time field is not allowed.
- There are a set of subentry types that are not allowed to be used in conditions on the *filter* parameter. See section for the list.

  A null value for a subentry may be specified by using two single quotation marks with no other characters or blanks between them (''). A null value may be used with the "not equal" (!=) operator to find all entries that have some value — for example, PATH!=''.
- A date fixed value must use the format 'yyddd'. GIMAPI checks the characters entered to make sure they are all numeric.
- A time fixed value must use the format *'hh:mm:ss'*, where *hh* is a two-digit decimal number in the range 00 through 23, and *mm* and *ss* are both two-digit decimal numbers in the range 00 through 59. (That is, in 24-hour clock format.)
- A null value can be used with date and time fields with the = and != operators to check if the values have been set or not.
- A closing quotation mark "'" must be found by the end of the *filter* parameter.
- All parentheses must be closed before the *filter* parameter string ends.
- When any other syntax errors are found in the *filter* parameter, GIMAPI issues a general syntax error message specifying the character position with the error.

Once the syntax is checked and zones other than the GLOBAL zone are specified, specific zone names are checked to see if they are ZONESETs. If so, the ZONESET name is removed from the list of zones passed by the caller and replaced with the zone names from the ZONESET. An attempt is then made to read the ZONEINDEX for each of the zones to be processed to see if any other CSI data sets need to be dynamically allocated. If one of the zones specified is not a ZONESET and cannot be found in the ZONEINDEX, the QUERY command terminates.

If any of the target zones or DLIB zones are contained in separate CSI data sets than the global CSI, those data sets must be allocated. The command first determines if they are already allocated by the calling routine. If so, that allocation is used and the data sets are not deallocated later. If not, the data sets are allocated with DISP=SHR. In either case, a shared enqueue is done.

The command now attempts to process the query. The requested entries are read from the CSI and compared with the filter. If there is a match, the requested subentries and their values are attached to the output storage. This is a set of data structures linked together, with storage being allocated for each structure used.

Once all entries are processed, any CSI data sets that were dynamically allocated are dequeued and deallocated.

Note: Any successful allocations and enqueues are undone, even if errors are found. The data sets are always freed when the program ends.

If the query requested from the calling program did not result in any data being returned, the command *output* parameter is set to zero.

QUERY command output

The CSI data is made up of entries, which can be thought of as record types, and subentries, which can be considered the fields within those records. Each entry has a set of subentries. Those subentry values may be a single value, such as FMID subentry of the MOD entry. Other subentries are made up of a list of values, such as the LMOD subentry of the MOD entry, which lists all the LMODs that contain the MOD. Refer to Chapter 5, "SMP/E Data Set Entries" for a complete description of the CSI entry types and their associated subentries.

The data returned from the query follows this principle in that it is organized by entries and subentries.

What is returned from the QUERY command is a pointer to storage containing the output data. The storage is a set of linked lists that describe the output as well as contain the values.

The pointer returned points to the head of a linked list of records describing each of the entry types returned from the query. The entry type structures (ENTRY_LIST) contain a 12-character name of the entry type (padded with blanks) and a pointer to a link list of entry structures (ENTRY) representing the instances of that entry type. It also contains a pointer to the next entry type record.

The entry record points to a linked list of subentry type structures (SUBENTRY) and to the next instance of that entry type. It also contains a string with the name of the entry and a string with the name of the zone where the entry was found.

Note: If no subentries are requested on the *subentry* parameter, the entry name and zone name are still returned with the entry structure, with no subentry structures attached.

The subentry type structure contains a 12-character name of the subentry type and a pointer to the value structure (ITEM_LIST), as well as a pointer to the next subentry type.

The value structure contains the length of the value, a pointer to the character string value and a pointer to the next value structure, for subentries that have lists of values.

Refer to "Data structures for QUERY command" on page 333 for details on the structures defined. "Example of QUERY command" on page 330 contains pictures of what the output looks like in storage.

Some subentries are lists of values, while others are single values. There are some subentries that are composite values as well as being lists. For example, the ZONEINDEX subentry of the GLOBALZONE entry contains the zone name, the CSI data set containing the zone name, and the zone type. Subentries such as these are returned with their parts combined into one string separated by commas, as shown here:

```
TARG1,UID.TARG1.CSI,TARGET
```

The order of the parts for subentries whose values have multiple parts follow the description of that subentry in Chapter 5, "SMP/E Data Set Entries".

Other subentries are indicator variables. This means they do not have a value, but the fact that they do (or do not) exist in the CSI indicates something. The PROTECT subentry of the DDDEF entry is an example of this. These types of subentries will have values in the QUERY output. If the subentry exists, the value is set to YES, if it does not, the value is set to NO.

If an asterisk is used in either the *entry* or *subentry* parameters, QUERY creates a complete list, internally, of the specified type (entries, subentries, or both). If all entries and subentries are requested, then a list for each type is generated. Each list is processed as if the user had entered it from the command. Table 33 on page 330 shows the possible outputs.

Table 33. Entry and subentry combinations/output

ENTRY/SUBENTRY contents	Output
• ENTRY set to * • SUBENTRY set to *	All valid subentries are returned for all instances of all valid entry types that fulfill the filter restrictions.
• ENTRY set to specific entry types • SUBENTRY set to *	All valid subentries are returned all instances of each of the specified entry types that fulfill the filter restrictions.
• ENTRY set to * • SUBENTRY set to specific subentries	The specific subentries are returned for all instances of all entry types that fulfill the filter restrictions.

Example of QUERY command

Suppose you want to write an application that will perform the following query:

• Find all PTF and USERMOD type SYSMODs having an FMID of HMP1E00 that were received after December 1, 2008 and that were applied with BYPASS(anything). Return the SYSMOD name, the zone, the MODs that are replaced by the SYSMOD, and the installation date.

The storage containing the input parameters may be represented like this:
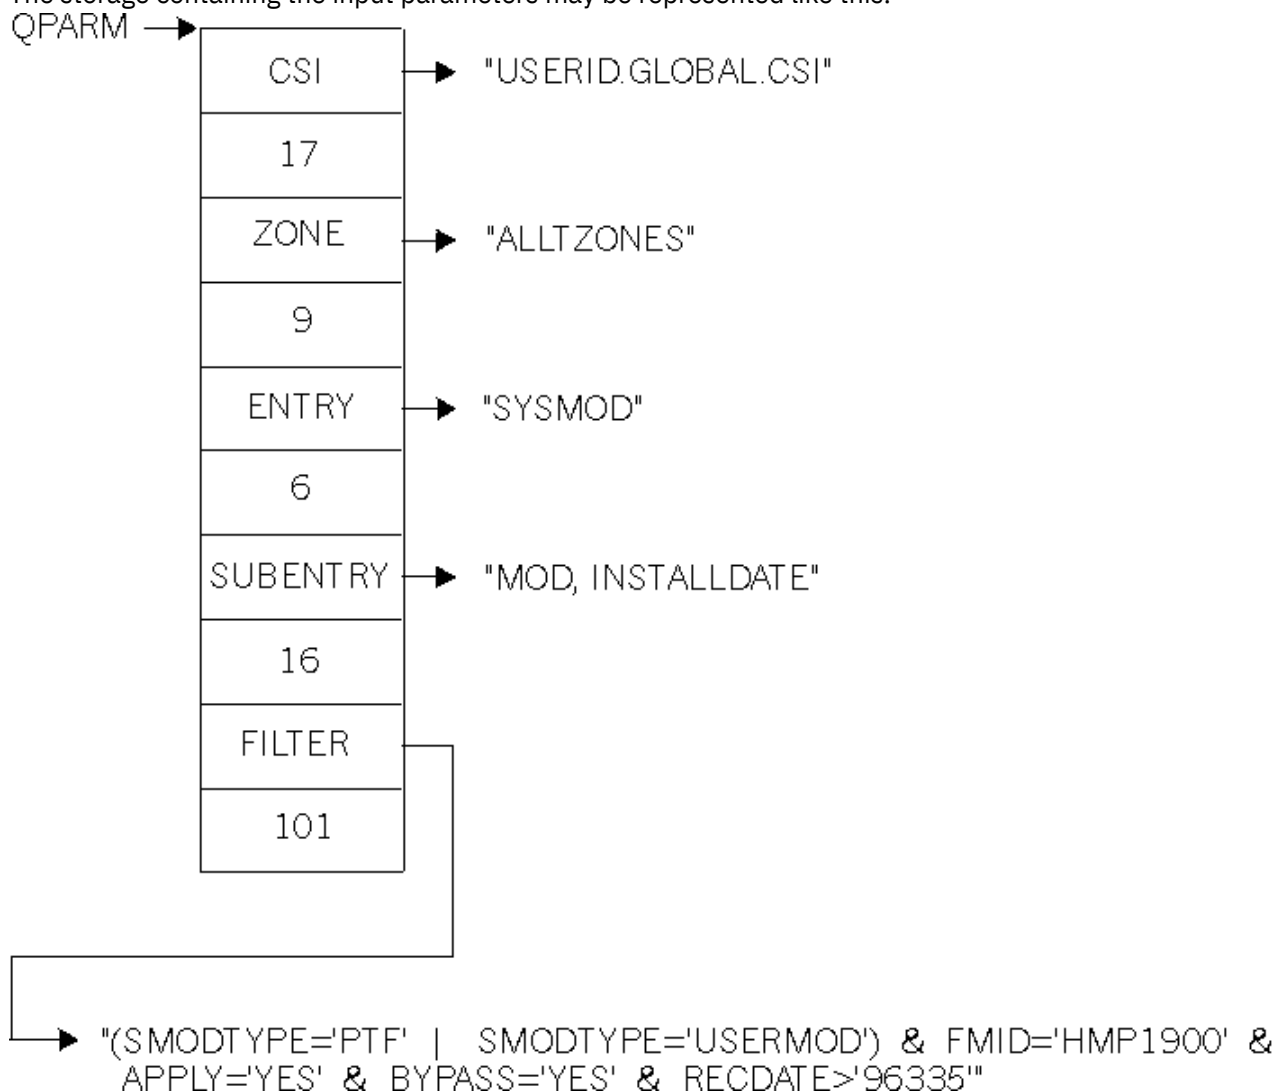
Note: The SYSMOD name and zone are automatically returned without your asking for them.

QPARM is a variable containing a pointer to a QUERY_PARMS data structure. The elements of the data structure point to strings in storage containing the values of the parameters.

Your program would then execute this call:

```
GIMAPI('QUERY',QPARM,QRESULT,'ENU',RC,CC,MSGBUFF)
```

The result of the query is two SYSMODS, SMOD019 and SMOD022. SMOD019 is in target zone TARG1 with two MODs, MOD01 and MOD02 and was installed on December 6, 2008. SMOD022 is in target zone MYTARG and has no MODs. It was installed on December 14, 2008.

provides a picture of what the output storage looks like. See section for a description of the data structures used in the picture.

Figure 63. Picture of storage for query output

Note: There are no values for the MOD subentry for SYSMOD SMOD022. Therefore, no SUBENTRY structure is added to the output for that subentry type.

Data structures for QUERY command

The data structures shown in this section are defined in a language-independent format. See "Programming in C" on page 338, "Programming in PL/I" on page 340, and "Programming in assembler" on page 342 to see the syntax of the data structures for those languages. Headers or macros with the language-specific structure definitions are provided in members named GIMHC370 (for C/370), GIMMPLI (for PL/I), and GIMMASM (for assembler) in libraries with these DDDEFs:

MACLIB
> Target library

AMACLIB
> Distribution library

Note: The pointer and numeric values found in the tables in this section are 4-byte fields.

QUERY_PARMS

Table 34 on page 333 lists the elements of the data structure used as input to the QUERY command. Only a brief description of each parameter is given here. See "QUERY command" on page 299 for a more complete description.

Table 34. QUERY command input parameters

Element	Data type	Description
csi	Pointer	Global CSI data set name to be searched by the query.
csilen	Numeric	Number of characters in the *csi* parameter.
zone	Pointer	Pointer to character string representing zones to search during the query.
zonelen	Numeric	Number of characters in the *zone* parameter
entrytype	Pointer	Pointer to character string representing the list of CSI entry types to be returned from the query.
entrylen	Numeric	Number of characters in the *entry* parameter
subentrytype	Pointer	Pointer to character string representing the list of subentries to be returned for each entry selected during the query.
subentrylen	Numeric	Number of characters in the *subentry* parameter
filter	Pointer	Pointer to character string representing the set of conditions used to limit the entries being returned by the query.
filterlen	Numeric	Number of characters in the *filter* parameter

ENTRY_LIST

Table 35. General SMP/E entry list structure

Element	Data type	Description
next	Pointer	Pointer to the next entry list structure. This value is zero (0) for the last item in the list.
entries	Pointer	Pointer to a linked list of entry structures. This value is zero (0) if there are no entries.
entrytype	Character(12)	Text representation of the type of entry pointed to by the entry list structure.

ENTRY

Element	Data type	Description
Table 36. General SMP/E entry structure		
Element	**Data type**	**Description**
next	Pointer	Pointer to the next entry structure. This value is zero (0) for the last item in the list.
subentries	Pointer	Pointer to a linked-list of subentry structures. This value is zero (0) if there are no subentries.
entryname	Character(8)	The name of an SMP/E entry.
zonename	Character(7)	The name of an SMP/E zone.

SUBENTRY

Element	Data type	Description
Table 37. General SMP/E subentry structure		
Element	**Data type**	**Description**
next	Pointer	Pointer to the next subentry structure. This value is zero (0) for the last item in the list.
subentrydata	Pointer	Pointer to a linked list of item list structures. This value is zero (0) if there is no subentry data.
type	Character(12)	Text representation of the type of subentry pointed to by the subentry structure.

VER

The VER structure is used when SYSMOD entry information is requested from the GLOBAL, target, or distribution zone. More specifically, this structure is used to return information that is typically associated with a ++VER statement. Refer to the "Valid subentry types" on page 306 for valid subentry and pseudo-subentry specification for the SYSMOD entry.

For each request for data associated with any of the information that is contained in a ++VER statement, a VER pseudo-subentry structure is returned.

The type value in the subentry structure is set to the 12-character string 'VER' and the subentrydata pointer points to a VER pseudo-subentry structure.

For each VER pseudo-subentry returned, a linked list of subentry structures are returned.

Each subentry structure returned will be associated with one of the types of data related to a ++VER statement. In other words, the 'type' field in the subentry structure will contain one of these text strings: SREL, FMID, PRE, NPRE, REQ, DELETE, SUP, or VERSION.

Each subentry structure returned will point to an item_list structure containing the actual VER-related data.
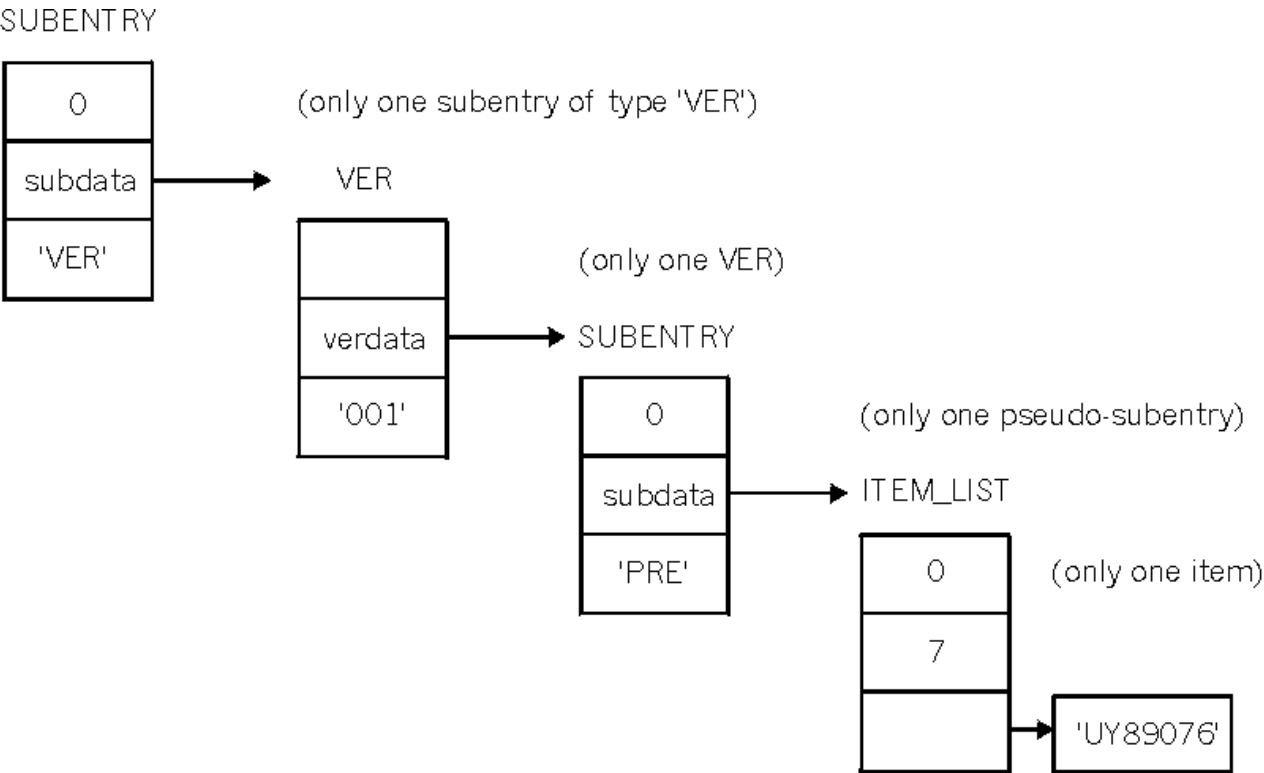
Figure 64. Illustration of VER data structure

Table 38. VER pseudo-subentry structure (GLOBAL, target, and DLIB zone)		
Element	**Data type**	**Description**
next	Pointer	Pointer to the next VER pseudo-subentry structure. This value is zero (0) for the last item in the list.
verdata	Pointer	Pointer to a linked list of subentry structures. This value is zero (0) if there is no data associated with the VER.
vernum	Character(3)	Contains the relative number of the ++VER statement when the SYSMOD was installed.

ITEM_LIST

This structure is used by many of the entry structures to hold an element in a list of subentry values attached to the entry. For example, it may contain a list of SYSMOD values, SRELs, or ALIAS names.

Table 39. Item list structure		
Element	**Data type**	**Description**
next	Pointer	Pointer to the next item in the list. This value is zero (0) when there are no more items.
datalen	Numeric	Number of characters in the item string.
data	Pointer	Pointer to the data associated with this item.

FREE command

The FREE command frees all the storage allocated by GIMAPI for query output, message buffer, and any other data that is saved between different calls to GIMAPI. The FREE command will free storage allocated by any previous queries that has not already been freed. It is suggested that the FREE command be run

after each query, but this is not required. It is good practice to use the FREE command prior to terminating the calling program.

FREE command parameters

Here is an example of a call to the FREE command:

```
GIMAPI('FREE ',0,0,'ENU',rc,cc,msgbuff)
```

FREE
is a keyword that specifies that GIMAPI is to process a free request. Note that it is necessary to pad the string containing the FREE command with blanks to bring it to a length of eight characters.

0,0
The *parmptr* and *outptr* parameters must both be set to zero (0) on the FREE command.

ENU
The *language* parameter must be specified as either ENU, JPN, or three blanks. The FREE command does not return any messages, so it does not matter which of these choices is specified.

rc
A storage area owned by the calling program representing a 4-byte numeric variable. This variable will contain the return code for the FREE command. GIMAPI will always set this variable to zero (0) after processing the FREE command.

cc
A storage area owned by the calling program representing a 4-byte numeric variable. This variable will contain the condition code for the FREE command. GIMAPI will always set this variable to zero (0) after processing the FREE command.

msgbuff
A pointer variable. No messages are issued by the FREE command, but this pointer variable must be passed to GIMAPI.

The *parmptr*, *outptr*, *language*, *rc*, *cc*, and *msgbuff* parameters must all be specified on the GIMAPI call, even though none of these parameters are used by the FREE command.

FREE command processing

The FREE command frees any storage obtained by GIMAPI on previous QUERY requests made by this caller.

VERSION command

The VERSION command is used to query the version of GIMAPI module being accessed by the calling program.

The version is returned as 11 bytes worth of data representing the current version of SMP/E in the format *vvrrmmppeee*, where:

vv
version

rr
release

mm
modification level

pp
PTF

eee
extendedPTF. When the PTF value (pp) exceeds 2 characters, the true PTF value can be found in this position.

The version of IBM SMP/E for z/OS, V3R6 is displayed as 36.00. This translates to an 11 byte value of "03060000000".

VERSION command parameters

Here is an example of a call to the VERSION command:

```
GIMAPI('VERSION',0,verout,'ENU',rc,cc,msgbuff)
```

No input parameters are defined for this command. A parameter passed to GIMAPI for this command would be ignored.

The output is a pointer to an API_VERSION structure containing the version of the current running GIMAPI module. Refer to

VERSION command processing

GIMAPI allocates the storage for an API_VERSION structure. If the storage cannot be allocated, existing message GIM30700E is issued and the module terminates.

If allocation was successful, GIMAPI sets the elements of the structure with the values of the current version, release, modification and PTF level then returns to the calling program.

For example, the level of IBM SMP/E for z/OS, V3R6 is 36.00. As the result of a call to GIMAPI with the VERSION command, the elements of the API_VERSION structure are set as follows:

```
apiver="03"
apirel="06"
apimod="00"
apiptf="00"
apieptf="000"
```

The version can be accessed by the user as the complete string, "03060000000" or as the previously listed individual parts.

The calling program needs to issue the FREE command at some point after calling the VERSION command.

VERSION command output

The data structure shown in this section are defined in a language independent format. See "Programming in C" on page 338, "Programming in PL/I" on page 340, and "Programming in assembler" on page 342 to see the syntax of the data structures for those languages. Headers or macros with the language-specific structure definitions are provided in members named GIMHC370 (for C/370), GIMMPLI (for PL/I), and GIMMASM (for assembler) in libraries with these DDDEFs:

MACLIB
Target Library

AMACLIB
Distribution Library

Table 40. API version		
Element	**Data type**	**Description**
apiver	Character(2)	A 2-digit character string representing the version of the API. Leading zeros are included.
apirel	Character(2)	A 2-digit character string representing the release of the API. Leading zeros are included.
apimod	Character(2)	A 2-digit character string representing the modification of the API. Leading zeros are included.

Table 40. API version (continued)		
Element	Data type	Description
apiptf	Character(2)	A 2-digit character string representing the PTF associated with the API. Leading zeros are included.
apieptf	Character(3)	A 3-digit character string representing the PTF associated with the API. Leading zeros are included. This is the true PTF level when apiptf exceeds 99.

Programming in C

For GIMAPI to be invoked from a C/370 program, GIMAPI must first be loaded into storage. The calling program must identify the parameter linkage as standard OS linkage and must declare the routine. shows the statements to load and unload the module, the #pragma statement that identifies the linkage, the function declaration, and the calling syntax.

```
typedef void APIPGM();
typedef void CFUNC();
#pragma linkage(APIPGM,OS)
⋮
APIPGM * gimapi;
gimapi = (APIPGM *) fetch("GIMAPI");
⋮
(*gimapi) (apicmd,&parmptr,&outptr,language,&rc,&cc,&msgbuff);
⋮
release((CFUNC*) gimapi)
```

Figure 65. C syntax of GIMAPI invocation

The FETCH must be done once, then GIMAPI can be invoked any number of times before it is released. The release function requires a pointer to a C function as its parameter. APIPGM becomes an OS program. The typedef of CFUNC is used to cast the *gimapi* parameter so the program compiles correctly.

apicmd
> The *apicmd* parameter is a string of length 8 that contains the name of the command to be passed to GIMAPI.

&parmptr
> A *qparm* variable is declared of type QUERY_PARMS. The *parmptr* pointer variable is set to the address of the query parameter structure. The address of the pointer variable is passed to GIMAPI.

&outptr
> The *outptr* variable is a pointer variable that will be set to the address of the beginning of storage containing the output of the command processing. The address of *outptr* is passed to GIMAPI.

language
> A character string of length 3 to indicate the language to use for messaging by GIMAPI. Valid values are ENU and JPN.

&rc
> Address of a variable defined as *long* to be set to the command's return code by GIMAPI.

&cc
> Address of a variable defined as *long* to be set to the command's condition code by GIMAPI.

&msgbuff
> The msgbuff variable is a pointer variable that is set to the head of a linked list of messages that could be created by GIMAPI processing. The elements to the link list are ITEM_LIST structures. The address of the pointer variable is passed to GIMAPI.

Data structures in C

Many of the structures contain character string data. The strings are not null terminated.

A header file will be provided in the MACLIB library. The member name for C/370 is GIMHC370.

QUERY_PARMS

```
typedef _Packed struct QUERY_PARMS
{
  char    *csi;
  long    csilen;
  char    *zone;
  long    zonelen;
  char    *entrytype;
  long    entrylen;
  char    *subentrytype;
  long    subentrylen;
  char    *filter;
  long    filterlen;
} QUERY_PARMS, * P_QUERY_PARMS;
```

API_VERSION

```
typedef _Packed struct API_Version
{
  char  apiver[2];
  char  apirel[2];
  char  apimod[2];
  char  apiptf[2];
  char  apieptf[3];
} API_VERSION, * P_API_VERSION;
```

ENTRY

```
typedef _Packed struct CSI_ENTRY
{
  _Packed struct CSI_ENTRY  *next;
  _Packed struct SUBENTRY   *subentries;
  char                      entryname[8];
  char                      zonename[7];
} CSI_ENTRY, * P_CSI_ENTRY;
```

ENTRY_LIST

```
typedef _Packed struct ENTRY_LIST
{
  _Packed struct ENTRY_LIST *next;
  _Packed struct CSI_ENTRY  *entries;
  char                      type[12];
} ENTRY_LIST, * P_ENTRY_LIST;
```

ITEM_LIST

```
typedef _Packed struct ITEM_LIST
{
  _Packed struct ITEM_LIST  *next;
  long                      datalen;
  char                      *data;
} ITEM_LIST, * P_ITEM_LIST;
```

SUBENTRY

```
typedef _Packed struct SUBENTRY
{
  _Packed struct SUBENTRY   *next;
```

```
   void                      *subentrydata;
   char                      type[12];
} SUBENTRY, * P_SUBENTRY;
```

Note: Given that C requires pointers to indicate the data type they will point to, a void pointer must be used here since more than one structure may be attached, the ITEM_LIST or the VER structure.

VER

```
typedef _Packed struct VER
{
   _Packed struct VER        *next;
   _Packed struct SUBENTRY   *verdata;
   char                      vernum[3];
} VER, * P_VER;
```

Programming in PL/I

Before the GIMAPI program can be invoked from a PL/I program, it must be declared. This identifies the variable as a callable routine and indicates to PL/I that standard OS linkage is used to pass parameters. The program must be loaded into storage before it can be invoked. The declaration, linkage, load and call syntax are shown in Figure 66 on page 340.

```
   DECLARE GIMAPI(CHAR(8), PTR, PTR, CHAR(3), FIXED, FIXED, PTR)
    ENTRY OPTIONS(ASSEMBLER INTER);
   ⋮
   FETCH GIMAPI;
   CALL GIMAPI(APICMD,PARMPTR,OUTPTR,LANGUAGE,RC,CC,MSGBUFF);
   ⋮
   RELEASE GIMAPI;
```

Figure 66. PL/I syntax of GIMAPI invocation

The FETCH must be done once, then GIMAPI can be invoked any number of times before it is released.

APICMD
A variable defined as CHAR(8) that contains the string representing the command GIMAPI is to process.

PARMPTR
A command parameter structure (QUERY_PARMS) that contains the parameters of the command being processed is declared in the calling program. PARMPTR is a pointer variable that contains the address of that structure. The variable is passed to GIMAPI.

OUTPTR
A pointer variable that will be set to the address of the beginning of storage containing the output of the command processing.

LANGUAGE
3-character national language identifier to use for messages. Valid values are ENU and JPN.

RC
A return code variable declared as FIXED BIN(31). The value of the variable is set to the return code of GIMAPI call by GIMAPI program.

CCPTR
A condition code variable declared as FIXED BIN(31). The value of the variable is set to the condition code of GIMAPI call by GIMAPI program.

MSGBUFF
The MSGBUFF variable is a pointer variable that is set to the head of a linked list of messages that could be created by GIMAPI processing. The elements to the link list are ITEM_LIST structures.

Data structures in PL/I

This section shows the data structures defined in "Data structures for QUERY command" on page 333 as they would appear in PL/I. The calling program must declare all the structures it is going to use before calling GIMAPI.

The data structures defined for the entries are BASED variables. This means no storage is allocated for the variables, but they are used to reference data allocated by the GIMAPI command. The calling program must declare pointer variables to reference the storage.

A macro file for PL/I is provided in member GIMMPLI of the MACLIB library.

QUERY_PARMS

```
DECLARE
 1 QUERY_PARMS UNALIGNED,
    2 CSI              POINTER,
    2 CSILEN           FIXED BIN(31),
    2 ZONE             POINTER,
    2 ZONELEN          FIXED BIN(31),
    2 ENTRYTYPE        POINTER,
    2 ENTRYLEN         FIXED BIN(31),
    2 SUBENTRYTYPE     POINTER,
    2 SUBENTRYLEN      FIXED BIN(31),
    2 FILTER           POINTER,
    2 FILTERLEN        FIXED BIN(31);
```

API_VERSION

```
DECLARE
  1 API_VERSION BASED UNALIGNED,
    2 APIVER           CHAR(2),
    2 APIREL           CHAR(2),
    2 APIMOD           CHAR(2),
    2 APIPTF           CHAR(2),
    2 APIEPTF          CHAR(3);
```

ENTRY

```
DECLARE
 1 CSI_ENTRY BASED UNALIGNED,
    2 NEXT             POINTER,
    2 SUBENTRIES       POINTER,
    2 ENTRYNAME        CHAR(8),
    2 ZONENAME         CHAR(7);
```

ENTRY_LIST

```
DECLARE
 1 ENTRY_LIST BASED UNALIGNED,
    2 NEXT             POINTER,
    2 ENTRIES          POINTER,
    2 TYPE             CHAR(12);
```

ITEM_LIST

```
DECLARE
 1 ITEM_LIST BASED UNALIGNED,
    2 NEXT             POINTER,
    2 DATALEN          FIXED BIN(31),
    2 DATA             POINTER;
```

SUBENTRY

```
DECLARE
 1 SUBENTRY BASED UNALIGNED,
   2 NEXT           POINTER,
   2 SUBENTRYDATA   POINTER,
   2 TYPE           CHAR(12);
```

VER

```
DECLARE
 1 VER BASED UNALIGNED,
   2 NEXT           POINTER,
   2 VERDATA        POINTER,
   2 VERNUM         CHAR(3);
```

Programming in assembler

Before the GIMAPI program can be invoked from an assembler program, the load module must be brought into virtual storage using the LOAD macro. It can then be invoked using the CALL macro passing the required parameters. An example of the declares, LOAD, CALL and FREE syntax are shown in Figure 67 on page 342.

```
              LA    2,APIPGM
         LOAD EPLOC=(2),LOADPT=PGMADR

         L    15,PGMADR
         CALL (15),(QUERYCMD,QUERY_PARMS@,CMDOUT,APILANG,RC,CC,MSG@)

         L    15,PGMADR
         CALL (15),(FREECMD,0,CMDOUT,APILANG,RC,CC,MSG@)

 APIPGM    DC    CL8'GIMAPI  '
 QUERYCMD  DC    CL8'QUERY   '
 FREECMD   DC    CL8'FREE    '
 APILANG   DC    CL3'ENU'
           DS    0F
 QUERY_PARMS@ DC AL4(QUERY_PARMS)
 PGMADR    DC    AL4(0)
 CMDOUT    DC    AL4(0)
 MSG@      DC    AL4(0)
 RC        DS    F'0'
 CC        DS    F'0'
```

Figure 67. Assembler syntax of GIMAPI invocation

QUERYCMD
   A variable defined as CHAR(8) that contains the string representing the QUERY command GIMAPI is to process.

FREECMD
   A variable defined as CHAR(8) that contains the string representing the FREE command GIMAPI is to process.

QUERY_PARMS@
   A command parameter structure (QUERY_PARMS) that contains the parameters of the command being processed is declared in the calling program. QUERY_PARMS@ is a pointer variable that contains the address of that structure. The variable is passed to GIMAPI.

CMDOUT
   A pointer variable that will be set to the address of the beginning of storage containing the output of the command processing.

APILANGE
   3-character national language identifier to use for messages. Valid values are ENU and JPN.

RC
> A return code variable declared as a fullword. The value of the variable is set to the return code of GIMAPI call by GIMAPI program.

CC
> A condition code variable declared as a fullword. The value of the variable is set to the condition code of GIMAPI call by GIMAPI program.

MSG@
> The MSG@ variable is an address variable that is set to the head of a linked list of messages that could be created by GIMAPI processing. The elements to the link list are ITEM_LIST structures.

Data structures in assembler

This section shows the data structures defined in "Data structures for QUERY command" on page 333 as they would appear in assembler. The calling program must declare all the structures it is going to use before calling GIMAPI.

The data structures defined for the entries are DSECTS. This means no storage is allocated for the variables, but they are used to reference data allocated by the GIMAPI command. The calling program must declare address variables to reference the storage.

A macro file for assembler code is provided in member GIMMASM of the MACLIB library.

QUERY_PARMS

```
QUERY_PARMS DS 0CL40       PARAMETERS FOR QUERY COMMAND
PCSI        DS AL4         PTR TO GLOBAL CSI DATASET
CSILEN      DS FL4         LENGTH OF DATA SET NAME
PZONE       DS AL4         PTR TO LIST OF ZONES
ZONELEN     DS FL4         LENGTH OF ZONE LIST
PENTRY      DS AL4         PTR TO LIST OF ENTRIES
ENTRYLEN    DS FL4         LENGTH OF ENTRY LIST
PSUBENTRY   DS AL4         PTR TO LIST OF SUBENTRIES
SUBENTRYLEN DS FL4         LENGTH OF SUBENTRY LIST
PFILTER     DS AL4         PTR TO QUERY FILTER
FILTERLEN   DS FL4         LENGTH OF FILTER
```

ENTRY_LIST

```
ENTRY_LIST  DSECT          LIST OF ENTRY TYPES
NEXT        DS AL4         PTR TO NEXT ITEM IN LINKED LIST
ENTRIES     DS AL4         PTR TO HEAD OF LINKED LIST OF
*                             CSIENTRY STRUCTURES LISTING THE
*                             INSTANCES OF THIS ENTRY TYPE
TYPE        DS CL12        ENTRY TYPE
```

ENTRY

```
CSIENTRY    DSECT          DESCRIPTION OF A SPECIFIC ENTRY
CSINEXT     DS AL4         PTR TO NEXT ITEM IN LINKED LIST
SUBENTRIES  DS AL4         PTR TO HEAD OF LINKED LIST OF
*                          SUBENTRY STRUCTURES LISTING THE
*                          SUBENTRY VALUES RETURNED FOR
*                             THIS ENTRY
ENTRYNAME   DS CL8         NAME OF ENTRY
ZONENAME    DS CL7         ZONE WHERE ENTRY WAS RETRIEVED
```

SUBENTRY

```
SUBENTRY    DSECT          CONTAINS SUBENTRY DATA
SUBNEXT     DS AL4         PTR TO NEXT ITEM IN LINKED LIST
SUBENTDATA  DS AL4         PTR TO HEAD OF LINKED LIST OF
*                             DATA VALUES FOR THIS SUBENTRY
SUBTYPE     DS CL12        SUBENTRY TYPE
```

VER

```
VER             DSECT           PLACEHOLDER FOR SUBENTRIES
*                                 ASSOCIATED WITH ++VER RECORDS
*                                 OF A SYSMOD
VERNEXT         DS AL4          PTR TO NEXT ITEM IN LINKED LIST
VERDATA         DS AL4          PTR TO HEAD OF LINKED LIST OF
*                                 SUBENTRIES ASSOC WITH A ++VER
VERNUM          DS CL3          ++VER STATEMENT NUMBER
```

ITEM_LIST

```
ITEM_LIST       DSECT           HOLDS PIECE OF SUBENTRY DATA
ITMNEXT         DS AL4          PTR TO NEXT ITEM IN LINKED LIST
DATALEN         DS FL4          NUMBER OF CHARS OF REAL DATA
DATA            DS AL3          PTR TO STORAGE CONTAINING DATA
```

API_VERSION

```
API_VERSION DSECT              API VERSION COMMAND OUTPUT
APIVER          DS CL2          CURRENT SMP/E LEVEL FOR VERSION
APIREL          DS CL2             "          "    RELEASE
APIMOD          DS CL2             "          "    MODIFICAT
APIPTF          DS CL2             "          "    2 CHAR PTF
APIEPTF         DS CL2             "          "    3 CHAR PTF
```

Additional programming considerations

GIMAPI is installed in SYS1.LINKLIB and SYS1.MIGLIB with the GIMSMP module. Because GIMAPI resides in LINKLIB, which is considered an authorized library, GIMAPI is available to both authorized and unauthorized callers.

The GIMAPI load module cannot be reused. A new copy must be brought into virtual storage for each use. Each task must load its own copy of GIMAPI. However, a single task can load GIMAPI once and then call GIMAPI several times before deleting GIMAPI.

GIMAPI runs in AMODE 31. It will obtain storage from either above or below the 16MB line, based on the AMODE of the caller upon entry to GIMAPI.

GIMAPI will not provide an ESTAE environment. If an ESTAE environment is required, the calling program must establish it before invoking GIMAPI.

Callers of GIMAPI should not link-edit a copy of the load module into the application program. GIMAPI should be called as an external routine.

The caller of GIMAPI must first load the GIMAPI load module into virtual storage using the appropriate syntax for the language in which the application is written. For example, an assembler application may use the LOAD or LINK macros. Once the GIMAPI module is loaded into virtual storage, it can be called by the application.

The caller of GIMAPI must adhere to standard linkage conventions. For more information regarding standard linkage conventions, refer to *z/OS MVS Programming: Assembler Services Guide* and the appropriate documentation for the language you are using for your application.

GIMAPI can only be invoked from a user-written program. GIMAPI cannot be invoked from JCL.

GIMAPI exposes certain control blocks to the user. These are the data structures defined in section "Data structures for QUERY command" on page 333. A macro data set is provided for you to include the definition of the control blocks in your program. The DDDEF of the macro library is MACLIB for the target library and AMACLIB for the distribution library. Refer to "Programming in C" on page 338, "Programming in PL/I" on page 340, and "Programming in assembler" on page 342 for the member names used for each language.

Sample programs that use GIMAPI

This section contains three sample programs that use GIMAPI. Each program does the same basic function, where differences are mostly due to the syntax of the languages involved.

Note: In addition to the programs documented here, SYS1.SAMPLIB provides other sample programs that use GIMAPI, such as GIMCRSAM and GIMPRSAM.

Each program follows these steps:

1. The query described in "Example of QUERY command" on page 330 is set up to be invoked. The global CSI data set used is 'SMP.VSAM.CSI'.
2. GIMAPI is called with the QUERY command to retrieve the data.
3. If the query is successful, it calls a generic result print function that prints the data retrieved from any query.
4. Once these are done, the storage is freed by calling GIMAPI with the FREE command.

The following is an example of the output of the sample programs:

```
Entry Type: SYSMOD
----------------------------------------
  ENAME       : SMOD19
  ZONE        : TARG1
  INSTALLDATE : 07340
  MOD         : MOD01
                MOD02
----------------------------------------
  ENAME       : SMOD22
  ZONE        : MYTARG
  INSTALLDATE : 07348
```

Sample C/370 program

This sample program can be found in the SAMPLIB library as member name GIMCSAMP.

```c
#include <stdio.h>
#include <stdlib.h>
#include <DD:SYSLIB(GIMHC370)> /* Contains API structure definitions */

#define FREE       "FREE     "
#define QUERY      "QUERY    "
#define APILANG    "ENU"
#define TXT_VER    "VER"
#define LEN_ETYPE       12
#define LEN_ENAME        8
#define LEN_ZNAME        7
#define LEN_VERNUM       3
#define LEN_TXTVER       3
#define LEN_MSG        256
typedef void APIPGM();
typedef void cfunc();
#pragma linkage(APIPGM,OS)

static void errprint(char *, long, long, ITEM_LIST *);
static void valprint(ITEM_LIST *);
static void resprint(ENTRY_LIST *);

void main(int argc, char *argv[])
{
  long  rc,cc;
  QUERY_PARMS qparms;
  P_QUERY_PARMS pqparms = &qparms;
  ENTRY_LIST  *qreslt;
  ITEM_LIST   *msgbuff;
  APIPGM      *gimapi;

  char csi[45];
  char zone[100];
  char ent[100];
  char subent[100];
```

```
  char filter[150];

 rc = 0;
 cc = 0;

 /************************************************/
 /* Load the GIMAPI load module for use later    */
 /************************************************/
 gimapi = (APIPGM *) fetch("GIMAPI");

 /*****************************************************/
 /* Create the QUERY. Put the parameter strings into  */
 /* variables and put the addresses of those variables*/
 /* in the query parameter structure along with the   */
 /* length of those strings.                          */
 /*****************************************************/
 strcpy(csi,"SMP.VSAM.CSI");
 strcpy(zone,"ALLTZONES");
 strcpy(ent,"SYSMOD");
 strcpy(subent,"MOD,INSTALLDATE");
 strcpy(filter,"(SMODTYPE='PTF' | SMODTYPE='USERMOD')");
 strcat(filter," & FMID='HMP1E00' & APPLY='YES'");
 strcat(filter," & BYPASS='YES' & RECDATE>'07335'");

 qparms.csi          = csi;
 qparms.csilen       = strlen(csi);
 qparms.zone         = zone;
 qparms.zonelen      = strlen(zone);
 qparms.entrytype    = ent;
 qparms.entrylen     = strlen(ent);
 qparms.subentrytype = subent;
 qparms.subentrylen  = strlen(subent);
 qparms.filter       = filter;
 qparms.filterlen    = strlen(filter);

 gimapi(QUERY,&pqparms,(void**) &qreslt,APILANG,&rc,&cc,&msgbuff);

 if (rc!=0)
 {
   errprint(QUERY, rc, cc, msgbuff);
   if (rc>4) goto EXIT;
 }

 /******************************************/
 /* Call routine to print results of query */
 /******************************************/
 resprint(qreslt);

 /******************************************/
 /* Free storage returned from the QUERY  */
 /******************************************/
 gimapi(FREE,0,0,APILANG,&rc,&cc,&msgbuff);

EXIT:

 release ((cfunc*)gimapi);
 }

 /*****************************/
 /* Print results of the query */
 /*****************************/
 static void resprint(ENTRY_LIST *head)
 {
 ENTRY_LIST  *curetype;
 CSI_ENTRY   *curentry;
 SUBENTRY    *cursubent;
 VER         *curver;
 SUBENTRY    *curversub;
 char        etype[13];
 char        vernumber[13];
 char        versubtype[13];
 char        stEname[LEN_ENAME+1];
 char        stZname[LEN_ZNAME+1];

  /******************************/
  /* Loop through each entry type */
  /******************************/
  for (curetype=head; curetype!=0 ; curetype=curetype->next)
  {
    /*******************************************/
    /* Print name of entry type being processed */
    /*******************************************/
```

```
      strncpy(etype,curetype->type,LEN_ETYPE);
      etype[LEN_ETYPE] = '\0';
      printf("Entry Type: %s\n",etype);

      /**********************************************************/
      /* Loop through each entry printing the ename and zone   */
      /* then the list of subentry values.                     */
      /**********************************************************/
      for (curentry=curetype->entries;
           curentry!=0;
           curentry=curentry->next)
      {
        printf("---------------------------------------\n");
        strncpy(stEname,curentry->entryname,LEN_ENAME);
        stEname[LEN_ENAME]='\0';
        strncpy(stZname,curentry->zonename,LEN_ZNAME);
        stZname[LEN_ZNAME]='\0';
        printf("  ENAME           : %s\n",stEname);
        printf("  ZONE            : %s\n",stZname);
        for (cursubent=curentry->subentries;
             cursubent!=0;
             cursubent=cursubent->next)
        {
          strncpy(etype,cursubent->type,LEN_ETYPE);
          etype[LEN_ETYPE] = '\0';
          if ((strncmp(etype,TXT_VER,LEN_TXTVER)) == 0)
          {
            for (curver=(P_VER) cursubent->subentrydata;
                 curver!=0;
                 curver=curver->next)
            {
              strncpy(vernumber,curver->vernum,LEN_VERNUM);
              vernumber[LEN_VERNUM]='\0';
              for (curversub=curver->verdata;
                   curversub!=0;
                   curversub=curversub->next)
              {

                /********************************/
                /* Now print ver subentry values */
                /********************************/
                strncpy(versubtype,curversub->type,LEN_ETYPE);
                versubtype[LEN_ETYPE]='\0';
                printf("  %.6s  VER(%s): ",versubtype,vernumber);
                valprint(curversub->subentrydata);
              }
            }
          } /* end ver subentry */
          else /* not a ver structure */
          {
            printf("  %s     : ",etype);
            valprint(cursubent->subentrydata);
          } /* end non-ver subentry */
        } /* end subentry for */
      } /* end entries for */
  } /* end entry type for */
}

static void valprint(ITEM_LIST *item1)
{
  char       databuff[500];
  ITEM_LIST  *curitem;

  for (curitem=item1;
       curitem!=0;
       curitem=curitem->next)
  {
    strncpy(databuff,curitem->data,curitem->datalen);
    databuff[curitem->datalen] = '\0';
    printf("%s\n",databuff);
    if (curitem->next!=0)
      printf("\n                       ");
  } /* end item for */
}

static void errprint(char *cmd, long rc, long cc, ITEM_LIST *msgs)
{
  char   msgout[LEN_MSG+1];
  ITEM_LIST *curmsg;
  unsigned short i;

  printf("Error processing command: %s. RC=%d CC=%d\n",
```

```
    cmd,rc,cc);
  printf("Messages follow:\n");
  /**********************************************/
  /* Loop through a linked list of error messages  */
  /* printing them out.                            */
  /**********************************************/
  for (curmsg=msgs; curmsg!=0; curmsg=curmsg->next)
  {
    strncpy(msgout,curmsg->data,curmsg->datalen);
    msgout[curmsg->datalen] = '\0';
    printf("%s\n",msgout);
  }
}
```

Sample PL/I program

This sample program can be found in the SAMPLIB library as member name GIMPSAMP.

```
 MAIN:
  PROC OPTIONS(MAIN) REORDER;

 %INCLUDE GIMMPLI;

 DCL QUERY CHAR(8)  INIT('QUERY');
 DCL FREE  CHAR(8)  INIT('FREE');
 DCL TXT_VER   CHAR(12) INIT('VER');
 DCL APILANG CHAR(3) INIT('ENU');

 DCL GIMAPI ENTRY(CHAR(8),PTR,PTR,CHAR(3),FIXED BIN(31),
               FIXED BIN(31),PTR)
             EXTERNAL OPTIONS(ASSEMBLER,INTER);

 DCL SYSNULL BUILTIN;
 DCL ADDR    BUILTIN;
 DCL SUBSTR  BUILTIN;

 DCL (RC,CC) FIXED BIN(31) INIT(0);

 DCL QPARMS POINTER;
 DCL MSGBUFF  POINTER;

 DCL QRESULT POINTER;
 DCL NULLPTR POINTER;

 DCL CSISTR     CHAR(44);
 DCL ZONESTR    CHAR(100);
 DCL ENTRYSTR   CHAR(100);
 DCL SUBENTSTR  CHAR(100);
 DCL FILTERSTR  CHAR(150);
 /******************************************************************/
 /*                                                              */
 /*  SET ADDRESS OF QUERY PARAMETERS TO THE QPARMS VAR           */
 /*                                                              */
 /******************************************************************/

 QPARMS = ADDR(QUERY_PARMS);

 /******************************************************************/
 /*                                                              */
 /*  LOAD GIMAPI LOAD MODULE                                     */
 /*                                                              */
 /******************************************************************/

 FETCH GIMAPI;

 /******************************************************************/
 /*                                                              */
 /*  PLUG QUERY PARMS INTO THE QUERY STRUCTURE                   */
 /*                                                              */
 /******************************************************************/
 CSISTR    = 'SMP.VSAM.CSI';
 ZONESTR   = 'ALLTZONES';
 ENTRYSTR  = 'SYSMOD';
 SUBENTSTR = 'MOD, INSTALLDATE';
 SUBSTR(FILTERSTR,1,41) = '(SMODTYPE=''PTF'' | SMODTYPE=''USERMOD'')';
 SUBSTR(FILTERSTR,42,19) = ' & FMID=''HMP1E00''';
 SUBSTR(FILTERSTR,61,16) = ' & APPLY=''YES''';
 SUBSTR(FILTERSTR,77,37) = ' & BYPASS=''YES'' & RECDATE>''07335''';
```

```
CSI         = ADDR(CSISTR);
CSILEN      = 12;
ZONE        = ADDR(ZONESTR);
ZONELEN     = 9;
ENTRYTYPE   = ADDR(ENTRYSTR);
ENTRYLEN    = 6;
SUBENTRYTYPE = ADDR(SUBENTSTR);
SUBENTRYLEN  = 16;
FILTER      = ADDR(FILTERSTR);
FILTERLEN   = 113;

CALL GIMAPI(QUERY,QPARMS,QRESULT,APILANG,RC,CC,MSGBUFF);

/***************************************/
/* PRINT ANY ERROR MESSAGES ENCOUNTERED */
/***************************************/
IF RC ¬=0
  THEN CALL ERRPRINT(QUERY);

/*******************************************************************/
/* CALL ROUTINE TO PRINT RESULTS OF QUERY IF QUERY WAS SUCCESSFUL */
/*******************************************************************/
IF RC<=4
  THEN CALL RESPRINT;

/***************************************/
/* FREE STORAGE RETURNED FROM THE QUERY */
/***************************************/
CALL GIMAPI(FREE,NULLPTR,NULLPTR,APILANG,RC,CC,MSGBUFF);

EXIT:
  RELEASE GIMAPI;

/*******************************************************************/
/* INTERNAL SUBROUTINES FOLLOW                                    */
/*******************************************************************/
RESPRINT:  PROCEDURE;

 DCL CURETYPE    POINTER;
 DCL CURENTRY    POINTER;
 DCL CURSUBENT   POINTER;
 DCL CURVER      POINTER;
 DCL CURVSUB     POINTER;
 DCL PRTITEM     POINTER;

 /********************************/
 /* LOOP THROUGH EACH ENTRY TYPE */
 /********************************/
 CURETYPE = QRESULT;                    /* POINT TO HEAD OF LIST */
 DO WHILE (CURETYPE¬=SYSNULL);
   /* PRINT NAME OF ENTRY BEING PROCESSED */
   PUT EDIT('Entry Type: ',CURETYPE->ENTRY_LIST.TYPE)
         (SKIP,A(12),A(15));
   /********************************************************/
   /* LOOP THROUGH EACH ENTRY PRINTING THE ENAME AND ZONE  */
   /* THEN THE LIST OF SUBENTRY VALUES.                    */
   /********************************************************/
   CURENTRY = CURETYPE->ENTRIES;
   DO WHILE (CURENTRY¬=SYSNULL);
     PUT SKIP LIST('-----------------------------------------');
     PUT EDIT('ENAME',':',CURENTRY->CSI_ENTRY.ENTRYNAME)
           (SKIP,X(2),A(5),X(10),A(1),X(1),A(8));
     PUT EDIT('ZONE',':',CURENTRY->CSI_ENTRY.ZONENAME)
           (SKIP,X(2),A(4),X(11),A(1),X(1),A(7));

     CURSUBENT=CURENTRY->SUBENTRIES;
     DO WHILE (CURSUBENT¬=SYSNULL);
       IF CURSUBENT->SUBENTRY.TYPE=TXT_VER THEN DO;
         CURVER=CURSUBENT->SUBENTRYDATA;
         DO WHILE (CURVER¬=SYSNULL);
           CURVSUB=CURVER->VERDATA;
           DO WHILE (CURVSUB¬=SYSNULL);
             PUT EDIT(CURVSUB->SUBENTRY.TYPE,'VER(',
                   CURVER->VERNUM,'):')
                 (SKIP,X(2),A(6),X(1),A(4),A(3),A(2),X(1));
             PRTITEM=CURVSUB->SUBENTRYDATA;
             CALL VALPRINT(PRTITEM);
             CURVSUB=CURVSUB->SUBENTRY.NEXT;
           END;
           CURVER=CURVER->VER.NEXT;
         END;
       END; /* End Process VER type subentries */
```

```
          ELSE DO;
            PUT EDIT(CURSUBENT->SUBENTRY.TYPE,':')
                    (SKIP,X(2),A(15),A(1),X(1));
            PRTITEM=CURSUBENT->SUBENTRYDATA;
            CALL VALPRINT(PRTITEM);
          END; /* End non-VER type subentries */


          CURSUBENT = CURSUBENT->SUBENTRY.NEXT;
        END; /* END SUBENT TYPE LOOP */

        CURENTRY = CURENTRY->CSI_ENTRY.NEXT; /* GET NEXT ENTRY */
      END; /* END ENTRY LOOP */
      PUT SKIP;
      CURETYPE = CURETYPE->ENTRY_LIST.NEXT;   /* GET NEXT ENTRY TYPE */
    END; /* END ENTRY TYPE LOOP */

  END  RESPRINT;

  VALPRINT:  PROCEDURE(ITEM1);
   DCL ITEM1       POINTER;
   DCL CURITEM     POINTER;
   DCL BUFFPTR     POINTER;
   DCL DATABUFF    CHAR(500) BASED(BUFFPTR);

   CURITEM = ITEM1;
   DO WHILE (CURITEM ¬=SYSNULL);
     BUFFPTR = CURITEM->DATA;
     PUT EDIT(SUBSTR(BUFFPTR->DATABUFF,1,CURITEM->DATALEN))
             (X(1),A);

     CURITEM = CURITEM->ITEM_LIST.NEXT; /* GET NEXT DATA VALUE */
     IF CURITEM¬=SYSNULL /* LINE UP NEXT VALUE IF THERE IS ONE */
       THEN PUT SKIP LIST('                 ');
   END; /* END DATA ITEM LOOP */

  END VALPRINT;

  ERRPRINT:  PROCEDURE(CMD);
   DCL CMD  CHAR(8);

   DCL CURMSG POINTER;
   DCL TEXTPTR     POINTER;
   DCL MSGTEXT     CHAR(256) BASED(TEXTPTR);

   PUT EDIT('Error processing command: ',CMD,'.  ','RC=',RC,'CC=',CC)
           (SKIP,A(26),A(8),A(3),A(3),F(5),X(2),A(3),F(5));
   IF MSGBUFF¬=SYSNULL THEN
   DO;
    PUT SKIP LIST('MESSAGES FOLLOW:');
    CURMSG = MSGBUFF;
    DO WHILE (CURMSG¬=SYSNULL);
      TEXTPTR = CURMSG->DATA;
      PUT SKIP LIST(SUBSTR(TEXTPTR->MSGTEXT,1,CURMSG->DATALEN));
      CURMSG = CURMSG->ITEM_LIST.NEXT;
    END;
   END;
    ELSE
     PUT SKIP LIST('NO MESSAGES RETURNED');
  END  ERRPRINT;
  END MAIN;
```

Sample assembler program

This sample program can be found in the SAMPLIB library as member name GIMASAMP.

```
  MAIN:
BALAPI   CSECT
         STM   14,12,12(13)
         LR    12,15
@PSTART  EQU   BALAPI
         USING @PSTART,12
         ST    13,SAVE+4
         LA    14,SAVE
         ST    14,8(13)
         LR    13,14
*
* END OF STANDARD LINKAGE
```

```
*
BALAPI   AMODE 31
*
* SET UP PARAMETERS FOR QUERY
*
         XC    QUERY_PARMS,QUERY_PARMS
         LA    3,MYCSI
         ST    3,PCSI
         LA    3,19
         ST    3,CSILEN
         LA    3,MYZONE
         ST    3,PZONE
         LA    3,3
         ST    3,ZONELEN
         LA    3,MYENTRY
         ST    3,PENTRY
         LA    3,6
         ST    3,ENTRYLEN
         LA    3,MYSUBNTY
         ST    3,PSUBENTRY
         LA    3,1
         ST    3,SUBENTRYLEN
         LA    3,MYFILTER
         ST    3,PFILTER
         LA    3,16
         ST    3,FILTERLEN
*
* NOW LOAD THE API
*
         LA    2,APIPGM
         LOAD  EPLOC=(2),LOADPT=PGMADR
*
* NOW DO THE QUERY
*
         L     15,PGMADR
         CALL  (15),(QUERYCMD,QUERY_PARMS@,CMDOUT,APILANG,RC,CC,MSG@)
*
*
* NOW SEE WHAT WAS RETURNED
*
         L     3,RC
         LTR   3,3
         BNZ   ERRPRINT
*
* ESTABLISH ADDRESSABILITY
*
         OPEN  (SYSPRINT,OUTPUT)
         L     3,CMDOUT
         USING ENTRY_LIST,3
         L     4,ENTRIES
         USING CSIENTRY,4
         L     5,SUBENTRIES
         USING SUBENTRY,5
         L     6,SUBENTDATA
         USING ITEM_LIST,6
         L     7,DATA
         USING RETDATA,7
*
* PRINT ENTRIES
*
PRTENT   LTR   4,4
         BZ    CLOSEOUT
         MVI   BUFFER,C' '
         MVC   BUFFER+1(119),BUFFER
         MVC   BUFFER(8),ENTRYNAME
         MVC   BUFFER+10(7),ZONENAME
         LA    2,ANSICHAR
         PUT   SYSPRINT,(2)
         B     PRTSUB
MOREENT  L     4,CSINEXT
         B     PRTENT
*
* PRINT SUBENTRIES
*
PRTSUB   LTR   5,5
         BZ    MOREENT
         MVI   BUFFER,C' '
         MVC   BUFFER+1(119),BUFFER
         MVC   BUFFER(11),SUBTYPE
         PUT   SYSPRINT,(2)
         B     PRTDATA
MORESUB  L     5,SUBNEXT
```

```
         B       PRTSUB
*
* PRINT DATA
*
PRTDATA  LTR     7,7
         BZ      MORESUB
         L       8,DATALEN
CHECKLEN C       8,OUTRECLN
         BNH     SETLEN
         L       8,OUTRECLN
SETLEN   LA      15,1
         SLR     8,15
         MVI     BUFFER,C' '
         MVC     BUFFER+1(119),BUFFER
         EX      8,@MOVDATA
         LA      2,ANSICHAR
         PUT     SYSPRINT,(2)
         L       8,DATALEN
         C       8,OUTRECLN
         BNH     MORESUB
         S       8,OUTRECLN
         ST      8,DATALEN
         A       7,OUTRECLN
         B       CHECKLEN
*
* PRINT ERROR MESSAGES
*
ERRPRINT OPEN    (SYSPRINT,OUTPUT)
         L       6,MSG@
         L       7,DATA
         L       8,DATALEN
CHKMSGLN C       8,OUTRECLN
         BNH     SETMSGLN
         L       8,OUTRECLN
SETMSGLN LA      15,1
         SLR     8,15
         MVI     BUFFER,C' '
         MVC     BUFFER+1(119),BUFFER
         EX      8,@MOVDATA
         LA      2,ANSICHAR
         PUT     SYSPRINT,(2)
         L       8,DATALEN
         C       8,OUTRECLN
         BNH     CLOSEOUT
         S       8,OUTRECLN
         ST      8,DATALEN
         A       4,OUTRECLN
         B       CHKMSGLN
*
CLOSEOUT CLOSE   SYSPRINT
*
* NOW FREE THE STORAGE OBTAINED DURING THE QUERY
*
CLEANUP  L       15,PGMADR
         CALL (15),(FREECMD,0,CMDOUT,APILANG,RC,CC,MSG@)
*
* NOW DELETE GIMAPI
*
         DELETE EPLOC=APIPGM
*
* EXIT
*
EXIT     LA      15,0
         L       13,4(,13)
         L       14,12(,13)
         LM      00,12,20(13)
         BR      14
*
*
@MOVDATA MVC     BUFFER(0),RETDATA
SYSPRINT DCB     DDNAME=SYSPRINT,DSORG=PS,MACRF=PM,LRECL=121,RECFM=FBA
ANSICHAR DC      CL1' '
BUFFER   DS      CL120
MYCSI    DC      CL19'TOPGUN.WAG.VSAM.CSI'
MYZONE   DC      CL3'TZ1'
MYENTRY  DC      CL6'SYSMOD'
MYFILTER DC      CL16'SMODTYPE=''PTF'''
MYSUBNTY DC      CL3'*'
APIPGM   DC      CL8'GIMAPI  '
QUERYCMD DC      CL8'QUERY   '
FREECMD  DC      CL8'FREE    '
APILANG  DC      CL3'ENU'
```

```
           DS      0F
QUERY_PARMS@ DC  AL4(QUERY_PARMS)
PGMADR    DC      AL4(0)
CMDOUT    DC      AL4(0)
MSG@      DC      AL4(0)
OUTRECLN  DC      F'120'
RC        DS      F'0'
CC        DS      F'0'
SAVE      DC      18F'0'
          GIMMASM
RETDATA   DSECT   CL0
          END
```

Chapter 7. Writing UNIX shell scripts

This chapter documents intended Programming Interfaces that allow the customer to write programs to obtain the services of IBM SMP/E for z/OS, V3R6.

To simplify the post-install work for z/OS UNIX Services application programs, some UNIX applications include shell scripts. These scripts perform additional processing when SMP/E installs elements into a UNIX file system. A product packager normally includes any necessary shell scripts with the product.

For example, if the hierarchical file system element is a TAR or PAX file, you can provide a shell script that performs the necessary steps to restore the file. As with other products, you use SMP/E to copy the element (a TAR or PAX file) to a directory in a UNIX file system. However, you rely on the element's shell script to actually explode the file into its component subdirectories and files.

SMP/E provides IBM and vendor product packagers with a generic interface for writing UNIX shell scripts. This chapter describes the interface and includes suggestions for designing a shell script for SMP/E processing.

You define the shell script to SMP/E through a hierarchical file system MCS statement, as described in "Hierarchical file system element MCS" on page 21. Once defined, the shell script receives control whenever SMP/E installs or deletes the element.

Designing a shell script for SMP/E processing

To process a file in a UNIX file system, a shell script must be able to cope with both of the actions that SMP/E can potentially perform on the file: copy and delete. That is, SMP/E can copy the file to a directory in a UNIX file system (as a new file or a replacement for an existing file) and, later, SMP/E can delete the file from a UNIX file system directory. A shell script must be able to detect either of these conditions (copy or delete) and respond accordingly.

Assume, for example, that as part of deleting or replacing a product on your system, you delete a function that was shipped in the product's tar file. SMP/E deletes only the original tar file from the directory in a UNIX file system. It is the responsibility of the shell script to clean up (delete) the tar file's exploded component subdirectories and files.

Fortunately, SMP/E provides shell scripts with the necessary input. This input comes in the form of environment variables that SMP/E sets, as follows:

Variable
> **Description**

SMP_Directory
> Directory in a UNIX file system that contains the file to be processed by your shell script. This directory should be considered the working directory for the shell script, and the shell script should not update any directories or files which do not reside within the working directory.

SMP_File
> Name of the file to be processed by your shell script. This is the hierarchical file system element processed by SMP/E.

SMP_Action
> Action that SMP/E is performing on the file. This value can be either of the following:
>
> **COPY**
>> SMP/E is copying the file into a UNIX file system directory.
>
> **DELETE**
>> SMP/E is deleting the file from a UNIX file system directory.

SMP_Phase
> Point in processing in which SMP/E calls the shell script. This value can be either of the following:

PRE
SMP/E is calling the shell script before performing the action indicated by the SMP_Action variable.

POST
SMP/E is calling the shell script after performing the action indicated by the SMP_Action variable.

SMP_Phase is always set to PRE when SMP_Action is DELETE to ensure that SMP/E calls the shell script before deleting its corresponding element.

LC_ALL
The default value for all of the UNIX locale environment variables. This value is always set to IBM-1047.

IBM-1047 defines the United States English character set. By specifying a specific locale, all shell scripts invoked by SMP/E receive the same environment, regardless of the locale being used on any particular driving system.

PATH
The Java runtime directory concatenated with the base PATH value (/bin:.). The Java runtime directory is obtained from either the SMPJHOME DD statement or the SMPJHOME DDDEF entry. The Java runtime directory is required only if the UNIX shell script invokes Java commands.

For example, if the SMPJHOME DD is allocated to "/usr/lpp/java/j1.4", the PATH environment variable value will be set to "/usr/lpp/java/j1.4/bin/:/bin:.".

If neither the SMPJHOME DD statement nor the SMPJHOME DDDEF entry is specified, SMP/E will not set the PATH environment variable.

This input allows you to design a shell script that can respond to SMP/E's actions, as described in the sections that follow.

Designing for copy actions

For a copy action, SMP/E sets the SMP_Action variable to COPY (as well as setting the other environmental variables) and passes control to the shell script.

The shell script should not assume a first-time environment for the file. For example, if APPLY REDO is in effect, previous copy processing for a file might have already completed successfully. Therefore the shell script should begin by performing some kind of clean-up processing before taking any copy-related actions.

Designing for delete actions

For a delete action, SMP/E sets the SMP_Action variable to DELETE (as well as setting the other environmental variables) and passes control to the shell script.

The shell script should not assume that the file has previously been installed. For example, if during an APPLY, a file is being installed for the first time, and is also being deleted by a subsequent SYSMOD, SMP/E might invoke the shell script to delete the file before it has actually been copied to a UNIX file system. Therefore, the shell script should determine whether the file exists before taking any delete-related actions.

Designing for diagnosis

Code your shell script to write status information to the STDOUT and STDERR files during the course of its processing. Doing so will aid users in verifying the successful completion of shell script processing, and, if necessary, diagnosing problems when failures occur.

SMP/E copies this information from the STDOUT and STDERR files to the print data set specified in the active HFSCOPY UTILITY entry where it can be viewed by the user. By default, SYSPRINT is the print data set.

It is recommended that shell scripts provide sufficient information about the completion of any functions they use, and especially about any detected error conditions.

Returning control to SMP/E

When your shell script completes processing, it must return control to SMP/E. Your shell script returns control to SMP/E through the exit shell command, specifying an appropriate exit status (for example, *exit 0* or *exit −1*). SMP/E uses your shell script's exit status value to determine whether the shell script completed successfully.

SMP/E recognizes the following exit statuses from shell scripts:

Exit status
> **Meaning**

00
> Shell script has completed successfully.

Any other value
> Shell script processing has failed.

Example shell script

In the following example, the shell script is designed to process a UNIX tar file. The shell script is written so that it:

1. Explodes the tar file in response to a COPY action by SMP/E
2. Deletes all files from the directory in response to a DELETE action by SMP/E.

```
# This script will either explode a tar file file or delete all files in
# the directory.  The script uses the following environment variables
# for input:
#
#   SMP_Directory - directory in which the tar file resides
#   SMP_File      - name of the tar file
#   SMP_Phase     - indicates whether the shell script is being called
#                   before or after SMP/E has processed the file
#   SMP_Action    - the action that SMP/E is performing: COPY or DELETE
#
echo "Starting script processing…"
status=0                    # initialize script status
#
# Verify that the required input was received by the shell script.
#
if test ! "$SMP_Directory"
then
  echo "** No SMP_Directory parameter specified."
  status=-1
elif test ! "$SMP_File"
then
  echo "** No SMP_File parameter specified."
  status=-1
elif test ! "$SMP_Phase"
then
  echo "** No SMP_Phase parameter specified."
  status=-1
elif test ! "$SMP_Action"
then
  echo "** No SMP_Action parameter specified."
  status=-1
fi
#
# If a parameter error was detected, then exit now.
#
if test $status -ne 0             # If status is not 0, an error was detected
then
  echo "   Ensure the input environment variables for this script have been provided."
  echo "   If SMP/E was not used to invoke the script, correct the caller to specify "
  echo "   the input environment variables.  If SMP/E invoked this script, contact "
  echo "   the IBM support center."
  echo "Parameter error.  Exiting script with status $status"
  exit $status
fi
#
# If the script is being invoked after the tar file has been copied
# to the directory (SMP/E phase is post-copy), and the desired action
# is COPY, explode the tar file.
```

```
#
if test $SMP_Phase = POST
then
  if test $SMP_Action = COPY
  then
    cd $SMP_Directory                # Set the working directory for pax
    echo "Exploding all components of $SMP_Directory$SMP_File using the pax command"
    pax -rvf $SMP_Directory$SMP_File # Explode the tar file
    status=$?                        # Get the status of the pax command
    if test $status -ne 0            # If pax failed, indicate so to the user
    then
      echo "** pax command failure: pax ended with status $status"
    else
      echo "pax command completed successfully."
    fi
    ls -l $SMP_Directory             # List the contents of the directory
  fi
fi
#
# If the action is DELETE, delete all component files.
#
if test $SMP_Action = DELETE
then
  echo "Deleting the following from $SMP_Directory using the rm command:"
  ls -l $SMP_Directory               # List the contents of the directory
  cd $SMP_Directory                  # Set the working directory for rm
  rm *                               # Delete the directory and all files
  status=$?                          # Get the status of the delete
  if test $status -ne 0              # If rm failed, indicate so to the user
  then
    echo "** rm command failure: rm ended with status $status"
  else
    echo "rm command completed successfully."
  fi
  if test $status = 0                # If rm was OK, then…
  then
    mkdir $SMP_Directory             # …recreate the directory, albeit empty
    status=$?                        # Get the status of the mkdir
    if test $status -ne 0            # If mkdir failed, indicate so to the user
   then
      echo "** mkdir command failure: mkdir ended with status $status"
    fi
  fi
fi
echo "Exiting script with status $status"
exit $status
```

Chapter 8. Library change file records

This chapter documents intended Programming Interfaces that allow the customer to write programs to obtain the services of IBM SMP/E for z/OS, V3R6.

This chapter documents the various record types that are produced and written to the SMPDATA1 and SMPDATA2 data sets as a result of APPLY or RESTORE processing. You can use these records to propagate the libraries and members modified by SMP/E APPLY and RESTORE processing to other systems that require the same changes.

Note: The CHANGEFILE subentry in the OPTIONS entry must be set to YES to instruct SMP/E to create these records.

SMP/E provides macros containing the mappings of these records:

- GIMMALC for assembler
- GIMMCLC for C
- GIMMPLC for PL/I.

Library change file record structure

The following sections provide samples of the various library change file records that are written to the SMPDATA1 and SMPDATA2 data sets as a result of APPLY or RESTORE processing.

Note:

1. SMP/E COMPRESS processing does not create any library change records.
2. The library change file records are of varying lengths.
3. All character data in all library change records are in uppercase, except for the following, which contain mixed-case character data:

   - Aliases in Alias Record Type 0 records
   - Link names in Alias Record Type 0 records
   - Symbolic link names in Alias Record Type 1 records
   - Path names in Element Record Type 1 records
   - Path names in Library Record Type 1 records

   See "A0 - Alias record type 0" on page 360 and "L1 - Library record type 1" on page 371 for details.
4. The scale lines shown in the examples do not appear in the actual SMPDATA1 and SMPDATA2 data sets.
5. In Figure 68 on page 362, the binary zeros after the alias name are not shown in the example.
6. In Figure 69 on page 363, the binary zeros after the symbolic link name are not shown in the example.

Library change file record types

The following are the valid record types produced by the library change interface:

Record type
  Description

A0
  Alias Record Type 0

A1
  Alias Record Type 1

C0
> Continuation Record Type 0

E0
> Element Record Type 0

E1
> Element Record Type 1

H0
> Header Record Type 0

L0
> Library Record Type 0

L1
> Library Record Type 1

L2
> Library Record Type 2

P0
> SYSMOD Status Record Type 0

S0
> SMP/E Environment Record Type 0

T0
> Trailer Record Type 0

A0 - Alias record type 0

An Alias Record Type 0 (A0) is created for each alias for an element or LMOD processed during APPLY or RESTORE processing. For hierarchical file system elements, the linkname is indicated in the record.

Multiple A0 records may be produced for the same element or LMOD, because more than one alias can exist for the same element or LMOD.

The purpose of the A0 record is to identify the aliases associated with the changed elements or LMODs associated with this execution of the APPLY or RESTORE command.

The format and contents of this record type are shown in .

Table 41. Alias record type 0			
Field name	**Position (decimal)**	**Length (decimal)**	**Description**
Record type	1	2	The characters 'A0'.
Name	3	8	The name of an SMP/E element or LMOD processed during APPLY or RESTORE processing. The data is left-justified and padded with blanks.

Table 41. Alias record type 0 (continued)			
Field name	**Position (decimal)**	**Length (decimal)**	**Description**
Type	11	12	The type of SMP/E element or LMOD processed during APPLY or RESTORE processing. The valid element types for an Alias Record Type 0 are: • Data elements • Hierarchical file system elements • JAR • MAC • LMOD • PROGRAM • SIDEDECK The data is left-justified and padded with blanks.
Action	23	8	The type of action SMP/E took against this alias of this element or LMOD during APPLY or RESTORE processing. Valid values are ADDREP and DELETE. This data is left-justified and padded with blanks.
DD name	31	8	The ddname associated with the target library associated with the named element or LMOD in this A0 record. For action type DELETE, this field indicates the target library ddname from which the alias has been deleted. For action type ADDREP, , this field indicates the target library ddname into which the alias has been added or replaced. An A0 record is created for every unique library that has an alias change. Elements with multiple SYSLIBs will have the appropriate A0 records created for each unique SYSLIB changed. Therefore, it is possible to have the same element have the same change in alias structure in multiple SYSLIBs. The data is left-justified and padded with blanks. This field contains the ddname of the associated target library regardless of whether the allocation was done by a DD statement or DDDEF.
Alias	39	1023	This is an alias name associated with this SMP/E element or LMOD. If the element is a hierarchical file system element, then the linkname is placed in this field. The data is mixed-case characters and is left-justified and padded with binary zeros.

For LMODs that have an associated side deck, there is one A0 record created for each unique alias associated with the LMOD (indicated by LMOD in the element type field of the A0 record) and one A0 record created for each unique alias associated with the LMOD's side deck (indicated by SIDEDECK in the element type field of the A0 record). Both the A0 record with element type LMOD and the A0 record with element type SIDEDECK contain the same name (the LMOD's name) in the Name field of the A0 record.

For example, suppose that SMP/E adds load module LMODA, which has two aliases, LMA and LMDA. Load module LMODA has a side deck. shows the A0 records created for LMODA:

```
              1         2         3         4         5         6         7         8
----+----0----+----0----+----0----+----0----+----0----+----0----+----0----+----0
A0LMODA    LMOD        ADDREP  LINKLIB LMA
A0LMODA    LMOD        ADDREP  LINKLIB LMDA
A0LMODA    SIDEDECK    ADDREP  SIDELIB LMA
A0LMODA    SIDEDECK    ADDREP  SIDELIB LMDA
```

Figure 68. Example of alias record type 0 records

A1 - Alias record type 1

An Alias Record Type 1 (A1) is created for each symbolic link associated with a hierarchical file system element and for each symbolic link associated with a load module within a UNIX file system that is processed during APPLY or RESTORE processing. For hierarchical file system elements, the symbolic link names are placed in the records. For load modules, values from the ALIAS (SYMLINK,*symlink*) control statements are placed in the records.

Multiple A1 records may be produced for the same hierarchical file system element or load module, because more than one symbolic link can exist for the same hierarchical file system element or load module.

The purpose of the A1 record is to identify the symbolic links associated with the changed hierarchical file system elements and the symbolic links associated with load modules processed with this execution of the APPLY or RESTORE command.

The format and contents of this record type are shown in .

Table 42. Alias record type 1			
Field name	**Position (decimal)**	**Length (decimal)**	**Description**
Record type	1	2	The characters 'A1'.
Name	3	8	The name of an SMP/E hierarchical file system element or load module processed during APPLY or RESTORE processing. The data is left-justified and padded with blanks.
Type	11	12	The type of SMP/E hierarchical file system element or LMOD processed during APPLY or RESTORE processing. The valid values are: • a hierarchical file system element type • JAR • LMOD The data is left-justified and padded with blanks.
Action	23	8	The type of action SMP/E took against this symbolic link during APPLY or RESTORE processing. Valid values are ADDREP and DELETE. The data is left-justified and padded with blanks.

Table 42. Alias record type 1 (continued)			
Field name	**Position (decimal)**	**Length (decimal)**	**Description**
DD name	31	8	The ddname associated with the target library for the named hierarchical file system element or load module in this A1 record. The data is left-justified and padded with blanks.
			For action type DELETE, this field indicates the target library ddname **from** which the symbolic link has been deleted.
			For action type ADDREP, this field indicates the target library ddname **into** which the symbolic link has been added or replaced.
			An A1 record is created for all libraries modified within a UNIX file system by having a hierarchical file system element replaced or deleted or by having a load module replaced or deleted.
			This field contains the ddname of the associated target library regardless of whether the allocation was done by a DD statement or DDDEF.
Symbolic link	39	1023	This is a symbolic link associated with this SMP/E hierarchical file system element or LMOD. The data is mixed-case characters and is left-justified and padded with binary zeros.

For example, suppose that load module LMODA, which has two symbolic links, LMA and LMDA, is added to the BPXUSER library. shows the A1 records created for LMODA:

```
          1         2         3         4         5         6         7         8
----+----0----+----0----+----0----+----0----+----0----+----0----+----0----+----0
A1LMODA    LMOD       ADDREP  BPXUSER LMA
A1LMODA    LMOD       ADDREP  BPXUSER LMDA
```

Figure 69. Example of alias record type 1

C0 - Continuation record type 0

One Continuation Record Type 0 (C0) is created for each delta produced by the results of APPLY or RESTORE processing when spill processing has occurred. The purpose of the C0 record is to uniquely identify the continuation of a set of library change records.

The format and contents of this record type are shown in .

Table 43. Continuation record type 0			
Field name	**Position (decimal)**	**Length (decimal)**	**Description**
Record type	1	2	The characters C0.
Target zone	3	7	The target zone name that was operated on by the APPLY or RESTORE command. It is left-justified and padded with blanks.

Table 43. Continuation record type 0 (continued)

Field name	Position (decimal)	Length (decimal)	Description
Time stamp	10	13	The date and time that APPLY or RESTORE processing completed for the entire delta. This field is in the form *yyyydddhhmmss*, where: **yyyy** year **ddd** day **hh** hours **mm** minutes **ss** seconds Time is represented in 24-hour clock format (military time). The time value matches the time value in the T0 record and is the time that the GIM20501I message is issued for the associated APPLY or RESTORE command.
ERROR count	23	6	The character representation of the count of SYSMODs applied or restored in error. The data is padded on the left with character zeroes (0).
INCMPLT count	29	6	The character representation of the count of SYSMODs left with a status of INCMPLT by the APPLY or RESTORE. The data is padded on the left with character zeroes (0).
APPLIED or RESTORED count	35	6	The character representation of the count of SYSMODs with a status of APPLIED or RESTORED. The data is padded on the left with character zeroes (0).
DELETED count	41	6	The character representation of the count of SYSMODs with a status of DELETED. The data is padded on the left with character zeroes (0).
SUPD count	47	6	The character representation of the count of SYSMODs with a status of SUPD. The data is padded on the left with character zeroes (0).

For example, suppose that a delta is produced for target zone TZONEA on January 16, 2008 at 9:57:22 AM. During this execution of SMP/E, two SYSMODs were applied in error, one SYSMOD was left incomplete, 123 SYSMODs were applied successfully, seven SYSMODs were deleted, and three SYSMODs were superseded. Figure 70 on page 364 shows the C0 records created that for delta:

```
             1         2         3         4         5         6         7         8
----+----0----+----0----+----0----+----0----+----0----+----0----+----0----+----0
C0TZONEA 2008016095722000002000001000123000007000003
```

Figure 70. Example of continuation record type 0

E0 - Element record type 0

An Element Record Type 0 (E0) is created for each element or LMOD that changed in a target library during APPLY or RESTORE processing. The term changed in this context refers to an element or LMOD that been deleted or replaced in a target library as a result of SMP/E processing. See "Valid action types" on page 377 for more information.

The purpose of the E0 record is to identify the changed elements or LMODs in the associated target libraries identified by the L0 and L1 records.

The format and contents of this record type are shown in Table 44 on page 365.

Table 44. Element record type 0			
Field name	**Position (decimal)**	**Length (decimal)**	**Description**
Record type	1	2	The characters 'E0'.
Name	3	8	The name of an SMP/E element or LMOD processed during APPLY or RESTORE processing. The data is left-justified and padded with blanks.
Type	11	12	The type of SMP/E element or LMOD processed during APPLY or RESTORE processing. The valid element types for an Element Record Type 0 are: • data elements • hierarchical file system elements • JAR • LMOD • MAC • PROGRAM • SIDEDECK • SRC The data is left-justified and padded with blanks.
Action	23	8	The type of action SMP/E took against this element or LMOD during APPLY or RESTORE processing. Valid values are ADDREP and DELETE. This data is left-justified and padded with blanks.

Table 44. Element record type 0 (continued)			
Field name	**Position (decimal)**	**Length (decimal)**	**Description**
DD name	31	8	The ddname associated with the target library associated with the named element or LMOD in this E0 record. The data is left-justified and padded with blanks.
			For cross-zone elements, this field indicates the SMP/E generated ddname associated with the SYSLIB for this element. The associated L0 record also contains the related true name for the SYSLIB in the cross-zone.
			For action type DELETE, this field indicates the target library ddname from which the element has been deleted.
			For action type ADDREP, this field indicates the target library ddname into which the element has been added or replaced.
			An E0 record is created for every unique library that has an element change. Elements with multiple SYSLIBs will have the appropriate E0 records created for each unique SYSLIB changed. Therefore, it is possible to have the same element have the same change take place in multiple SYSLIBs.
			This field contains the ddname of the associated target library regardless of whether the allocation was done by a DD statement or DDDEF.

For LMODs that have an associated side deck, there is one E0 record created for each LMOD (indicated by LMOD in the element type field of the E0 record) and one E0 record created for the LMOD's side deck (indicated by SIDEDECK in the element type field of the E0 record). Both the E0 record with element type LMOD and the E0 record with element type SIDEDECK contain the same name (the LMOD's name) in the *name* field of the E0 record.

For example, suppose that SMP/E adds load module LMODA, which has a SIDEDECK. Figure 71 on page 366 shows the E0 records created for LMODA:

```
            1         2         3         4         5         6         7         8
----+----0----+----0----+----0----+----0----+----0----+----0----+----0----+----0
E0LMODA    LMOD         ADDREP  LINKLIB
E0LMODA    SIDEDECK     ADDREP  SIDELIB
```

Figure 71. Example of element record type 0

E1 - Element record type 1

One Element Record Type 1 (E1) is created for each element having a UNIX shell script that changed in a target library during APPLY or RESTORE processing. The term changed in this context refers to an element that has been deleted or replaced in a target library as a result of SMP/E processing. See "Valid action types" on page 377 for more information.

The purpose of the E1 record is to identify the changed elements in the associated target libraries identified by the L0 and L1 records.

The format and contents of this record type are shown in Table 45 on page 367.

Table 45. Element record type 1			
Field name	**Position (decimal)**	**Length (decimal)**	**Description**
Record type	1	2	The characters 'E1'.
Name	3	8	The name of an SMP/E element processed during APPLY or RESTORE processing. The data is left-justified and padded with blanks.
Type	11	12	The type of hierarchical file system element processed by the APPLY or RESTORE command. The data is left-justified and padded with blanks.
Action	23	8	The type of action SMP/E took against this element during APPLY or RESTORE processing. Valid values are ADDREP and DELETE. This data is left-justified and padded with blanks.
DD name	31	8	The ddname associated with the target library associated with the named element in this E1 record. The data is left-justified and padded with blanks.
			For cross-zone elements, this field indicates the SMP/E generated ddname associated with the SYSLIB for this element. The associated L1 record also contains the related true name for the SYSLIB in the cross-zone.
			For action type DELETE, this field indicates the target library ddname from which the element has been deleted.
			For action type ADDREP, this field indicates the target library ddname into which the element has been added or replaced.
			An E1 record is created for every unique library that has an element change. Elements with multiple SYSLIBs will have the appropriate E1 records created for each unique SYSLIB changed. Therefore, it is possible to have the same element have the same change take place in multiple SYSLIBs.
			It is also possible to have multiple E1 records for the same element in the same SYSLIB if the element is both deleted and replaced. Normally, only the net effect would be represented in the E1 record. However, when a shell script is run, two records are produced because the shell script is run for each action.
			This field contains the ddname of the associated target library regardless of whether the allocation was done by a DD statement or DDDEF.

Table 45. Element record type 1 (continued)

Field name	Position (decimal)	Length (decimal)	Description

Phase	39	8	The phase of processing when the shell script was invoked to complete the installation of the identified element. Valid values are:
			
			PRE
			    Shell script was invoked before the action was performed.
			
			POST
			    Shell script was invoked after the action was performed.
			
			PRE,POST
			    Shell script was invoked both before and after the action was performed.
			
			The data is left-justified and padded with blanks.

Path name	47	263	The full path name of the shell script. The data is left-justified and padded with blanks. The value is of the form /*directory/file*, in which:
			
			/*directory*/
			    Directory in a UNIX file system in which the shell script resides. This value can be up to 255 characters.
			
			file
			    Name of the shell script file. This value can be 1 to 8 characters long.

For example, suppose that a ++HFS element named XMPHFS was added to target zone TZONEB using a shell script named MYSHELL in directory */usr/lpp/abc/* and that shell script MYSHELL was invoked after XMPHFS was installed. Figure 72 on page 368 shows the E1 records created for LMODA:

```
              1         2         3         4         5         6         7         8
    ----+----0----+----0----+----0----+----0----+----0----+----0----+----0----+----0
    E1XMPHFS  HFS          ADDREP  TZONEB  POST     /usr/lpp/abc/MYSHELL
```

Figure 72. Example of element record type 1

H0 - Header record type 0

One Header Record Type 0 (H0) is created for each delta produced by the results of APPLY or RESTORE processing. The purpose of the H0 record is to uniquely identify the start of a set of library change records.

The format and contents of this record type are shown in Table 46 on page 368.

Table 46. Header record type 0

Field name	Position (decimal)	Length (decimal)	Description
Record type	1	2	The characters 'H0'.
Zone name	3	7	The target zone name that was operated on by the APPLY or RESTORE command. It is left-justified and padded with blanks.

Table 46. Header record type 0 (continued)

Field name	Position (decimal)	Length (decimal)	Description

Time stamp	10	13	The date and time that APPLY or RESTORE processing completed for the entire delta. This field is in the form *yyyydddhhmmss*, where:
			
			yyyy
			    year
			
			ddd
			    day
			
			hh
			    hours
			
			mm
			    minutes
			
			ss
			    seconds
			
			Time is represented in 24-hour clock format (military time). The time value matches the time value in the T0 record and is the time that the GIM20501I message is issued for the associated APPLY or RESTORE command.

ERROR count	23	6	The character representation of the count of SYSMODs applied or restored in error. The data is padded on the left with character zeroes (0).
INCMPLT count	29	6	The character representation of the count of SYSMODs left with a status of INCMPLT by the APPLY or RESTORE. The data is padded on the left with character zeroes (0).
APPLIED or RESTORED count	35	6	The character representation of the count of SYSMODs with a status of APPLIED or RESTORED. The data is padded on the left with character zeroes (0).
DELETED count	41	6	The character representation of the count of SYSMODs with a status of DELETED. The data is padded on the left with character zeroes (0).
SUPD count	47	6	The character representation of the count of SYSMODs with a status of SUPD. The data is padded on the left with character zeroes (0).

For example, suppose that a delta is produced for target zone TZONEA on January 16 of 2008 at 9:57:22 AM. During this execution of SMP/E, two SYSMODs were applied in error, one SYSMOD was left incomplete, 123 SYSMODs were applied successfully, seven SYSMODs were deleted, and three SYSMODs were superseded. Figure 73 on page 369 shows the H0 records created for that delta:

```
          1         2         3         4         5         6         7         8
----+----0----+----0----+----0----+----0----+----0----+----0----+----0----+----0
H0TZONEA 2008016095722000002000001000123000007000003
```

Figure 73. Example of header record type 0

L0 - Library record type 0

One Library Record Type 0 (L0) is created for each target library changed during APPLY or RESTORE processing that is not associated with a pathname. A Library Record Type 1 (L1) is created for pathnames.

The purpose of the L0 record is to identify the target libraries (excluding pathnames) that were changed by this execution of APPLY or RESTORE.

SYSLIB concatenations used during the APPLY or RESTORE do not produce any L0 records.

The format and contents of this record type are shown in Table 47 on page 370.

Field name	Position (decimal)	Length (decimal)	Description
Record type	1	2	The characters 'L0'.
DD name	3	8	The SMP/E ddname associated with this L0 record. The data is left-justified and padded with blanks. This field contains the ddname of the allocated target library regardless of whether the allocation was done by a DD statement or DDDEF. For cross-zone libraries, this field indicates the SMP/E-generated ddname for this library.
SYSLIB ddname	11	8	This is the SMP/E SYSLIB ddname associated with cross-zone libraries. This is the true ddname of the cross-zone library in the associated cross-zone. The data is left-justified and padded with blanks. For libraries that are not cross-zone libraries, this field is set to blanks.
Cross-zone name	19	7	The cross-zone name associated with this cross-zone library. The data is left-justified and padded with blanks. For libraries that are not cross-zone libraries, this field is set to blanks.
Library type	26	8	The type of library that is associated with this L0 record. The valid library types are LIBRARY, PDS, and SEQ. The data is left-justified and padded with blanks.
Volume name	34	6	The volume name associated with the SMP/E target library named in this record.
Data set name	40	44	The data set name of an SMP/E target library that was changed during APPLY or RESTORE processing. The data is left-justified and padded with blanks.
Catalog data set name	84	44	The catalog data set name associated with the target library data set name identified in this record. The data is left-justified and padded with blanks. If the data set name identified in this record is not cataloged, this field is set to blanks.

Table 47. Library record type 0

For example, Figure 74 on page 371 shows the L0 record generated for a PDS library named SYS1.LINKLIB on volume MVSRES in catalog CATALOG.MVSICFM.VMVSRES. The library has a ddname of LINKLIB.

```
                                                                        1         1
           1         2         3         4         5         6         7         8         9         0         1
----+----0----+----0----+----0----+----0----+----0----+----0----+----0----+----0----+----0----+----0----+----0
L0LINKLIB             PDS     MVSRESSYS1.LINKLIB                              CATALOG.MVSICFM.VMVSRES
```

Figure 74. Example of library record type 0

L1 - Library record type 1

One Library Record Type 1 (L1) is created for each target library changed during APPLY or RESTORE processing that is associated with a pathname. An L0 record is created for libraries that are not associated with pathnames.

The purpose of the L1 record is to identify the target libraries containing pathnames that were changed by the associated APPLY or RESTORE command.

The format and contents of this record type are shown in Table 48 on page 371.

Table 48. Library record type 1

Field name	Position (decimal)	Length (decimal)	Description
Record type	1	2	The characters 'L1'.
DD name	3	8	The SMP/E ddname associated with this Library Type 1 record. The data is left-justified and padded with blanks. This field contains the ddname of the allocated target library regardless of whether the allocation was done by a DD statement or DDDEF. For cross-zone libraries, this field indicates the SMP/E-generated ddname for this library.
SYSLIB name	11	8	This is the SMP/E SYSLIB ddname associated with cross-zone libraries. This is the true ddname of the cross-zone library in the associated cross-zone. The data is left-justified and padded with blanks. For libraries that are not cross-zone libraries, this field is set to blanks.
Cross-zone name	19	7	The cross-zone name associated with this cross-zone library. The data is left-justified and padded with blanks. For libraries that are not cross-zone libraries, this field is set to blanks.
Library type	26	8	The type of library that is associated with this L1 record. The only valid type is HFS. The data is left-justified and padded with blanks.
Path name	34	255	The path name of an SMP/E target library that was changed during APPLY or RESTORE processing. The data is left-justified and padded with blanks.

For example, Figure 75 on page 372 shows the L1 record generated for a pathname of '*/hfs_path_name/*', which is an HFS library with the ddname of HFSLIB1.

```
                 1         2         3         4         5         6         7         8
        ----+----0----+----0----+----0----+----0----+----0----+----0----+----0----+----0
        L1HFSLIB1               HFS/hfs_path_name/
```

Figure 75. Example of library record type 1

L2 - Library record type 2

A Library Record Type 2 (L2) is created for each physical data set associated with a directory in a UNIX file system that was changed during APPLY or RESTORE processing. The purpose of the L2 record is to identify the physical data set in which a particular directory is located, how the data set is related to the directory, and the ddname that was used to allocate the directory.

The format and contents of this record type are shown in .

Table 49. Library record type 2			
Field name	**Position (decimal)**	**Length (decimal)**	**Description**
Record type	1	2	The characters 'L2'.
DD name	3	8	The SMP/E ddname associated with this Library Type 2 record. The data is left-justified and padded with blanks. This field contains the ddname that was used to allocate the directory in a UNIX file system that was updated. This value can be used to correlate the Library Type 2 records with the Library Type 1 records that describe the directory.
Relationship	11	8	This field describes the relationship between the updated directory and the physical data set identified in the record. Valid relationships are: **PATHHFS** The HFS data set specified in this record contains the directory that contains a file updated by SMP/E. **SYMHFS** The HFS data set specified in this record contains the directory that contains a symbolic link associated with a file updated by SMP/E. Library Type 2 records with a type of SYMHFS are created only when the SYMHFS data set names are different than the PATHHFS data set name. In addition, records only for unique SYMHFS data set names are produced. For example, if a file updated by SMP/E has two symbolic link values, and those symbolic link values reside in two directories that reside in the same physical data set, then only one Library Type 2 record with a type of SYMHFS is produced to describe this data set. In addition, if the data set that contains the symbolic links is the same physical data set that contains the file updated by SMP/E, then no Library Type 2 records with a type of SYMHFS are produced; a single record with a type of PATHHFS describes the physical data set that contains the file and its symbolic links. The data is left-justified and padded with blanks.

Table 49. Library record type 2 (continued)			
Field name	**Position (decimal)**	**Length (decimal)**	**Description**
Physical data set name	19	44	data set name of the physical data set that was updated when SMP/E updated a file in a UNIX file system. The data is left-justified and padded with blanks.

For example, suppose the following PTF is selected for APPLY processing and that the SBKSBIN library is allocated to directory /service/usr/lpp/booksrv/cgi-bin/IBM/ in a UNIX file system.

```
++PTF(UW12345).
++VER(Z038) FMID(HYY2900).
++HFS(BKSMAIN) SYSLIB(SBKSBIN) DISTLIB(ABKSBIN)
  PARM(PATHMODE(7,5,5))
  LINK('../bksmain')
  SYMLINK('../../../../../bin/bksmain')
  SYMPATH('../usr/lpp/booksrv/cgi-bin/bksmain').
```

The resulting APPLY processing replaces the BKSMAIN file in the /service/usr/lpp/booksrv/cgi-bin/IBM/ directory with the copy of the HFS element supplied in the PTF. Also, when the file's symbolic link value specified on the MCS is concatenated with the file's directory, the symbolic link value resolves to /service/bin/bksmain. This file is the symbolic link and resides in the /service/bin/ directory. In this example, the directory containing the symbolic link is in a different physical data set than the directory that contains the file BKSMAIN.

At the end of APPLY processing, SMP/E produces L1 and L2 records to summarize the updates performed. Figure 76 on page 373 shows the L1 record that describes the directory that contains the file, as well as the L2 records that describe the data sets that contain the file and its symbolic links.

```
         1         2         3         4         5         6         7         8
----+----0----+----0----+----0----+----0----+----0----+----0----+----0----+----0
L1SBKSBIN               HFS      /service/usr/lpp/booksrv/cgi-bin/IBM/
L2SBKSBIN PATHHFS OMVS.HFS.BOOKSRV
L2SBKSBIN SYMHFS  OMVS.HFS.ROOT.ZOS130
```

Figure 76. Example of library record type 2

P0 - SYSMOD status record type 0

Status records are created for SYSMODs for which some utility work was done during the command and for SYSMODs superseded by SYSMODs that were successfully applied or restored.

Note: The superseded SYSMODs are included to aid in analyzing when a set of deltas is ready for distribution. When a SYSMOD in ERROR status is noted as superseded in a subsequent delta, it's error status may be considered resolved.

The format and contents of this record type are shown in Table 50 on page 373.

Table 50. SYSMOD status record type 0			
Field name	**Position (decimal)**	**Length (decimal)**	**Description**
Record type	1	2	The characters 'P0'.
SYSMOD ID	3	7	The SYSMOD ID.

Field name	Position (decimal)	Length (decimal)	Description
Status	10	8	The status of the SYSMOD. The following text may appear in this field. Where the text is less than 8 characters, it is left justified and padded with blanks. • APPLIED • DELETED • ERROR • INCMPLT • RESTORED • SUPD The status in this library change record is consistent with the status of the SYSMOD as reported in the SYSMOD STATUS REPORT. **Note:** SYSMODs identified in the SYSMOD STATUS REPORT with the following status are not included in the library change records, because they failed during SYSMOD selection processing. • EXCLUDED • HELD • NOGO • NOGO(E) • NOGO(H)
FMID	18	7	The FMID of the SYSMOD. If the status is SUPD and the SYSMOD is not in process, this field contains blanks.
Type	25	8	The SYSMOD type. The following text is valid in this field: • FUNCTION • PTF • APAR • USERMOD If the status is SUPD and the SYSMOD is not in process, this field contains blanks.

Table 50. SYSMOD status record type 0 (continued)

The SYSMOD Status (P0) records follow the E0 record in the set of library change records for the command. Figure 77 on page 374 shows examples of P0 records written for APPLY processing.

```
          1         2         3         4         5         6         7         8
----+----0----+----0----+----0----+----0----+----0----+----0----+----0----+----0
P0UZ00002ERROR    H000001PTF
P0UZ00001APPLIED  H000001PTF
P0UZ00000SUPD
```

Figure 77. Example of SYSMOD status records

Additionally, fields are added to the header and trailer records to represent the count of SYSMODs represented in the P0 records with each status.

This information tells you whether any errors are contained in the current delta. However, just because a delta is error free, it does not mean that it can be distributed without resolving errors in previous deltas. You must perform the task of error resolution among deltas prior to distribution in order to ensure an error free package. Refer to "Usage recommendations" on page 378 for usage information.

S0 - SMP/E environment record type 0

One SMP/E Environment Record Type 0 (S0) is created for each delta produced by the results of APPLY or RESTORE processing.

The purpose of the S0 record is to identify the SMP/E environment from which this set of library change records was derived.

The format and contents of this record type are shown in Table 51 on page 375.

Table 51. SMP/E environment record type 0

Field name	Position (decimal)	Length (decimal)	Description
Record type	1	2	The characters 'S0'.
Command	3	7	The SMP/E command that was executed. Valid values are APPLY or RESTORE. The data is left-justified and padded with blanks.

SMP/E level	10	8	The level of SMP/E under which the APPLY or RESTORE command was executed. This subfield is in the form *vvrrmmpp*, where:
			
			vv
			    version of SMP/E
			
			rr
			    release of SMP/E
			
			mm
			    release of SMP/E
			
			pp
			    PTF level of SMP/E

Data set name	18	44	The data set name of the SMP/E target zone CSI that contains the target zone operated on by the SMP/E APPLY or RESTORE function. The data is left-justified and padded with blanks.
Catalog data set name	62	44	The catalog data set name of the SMP/E target zone CSI that contains the target zone operated on by the SMP/E APPLY or RESTORE function. The data is left-justified and padded with blanks.
Volume name	106	6	The volume name of the SMP/E target zone CSI that contains the target zone operated on by the SMP/E APPLY or RESTORE function.
SMP/E extendedPTF level	112	3	PTF level of SMP/E, to replace the PTF value in the SMP/E level appearing in column 16 when the PTF value exceeds 2 characters.

For example, Figure 78 on page 376 shows an S0 record for a target zone CSI data set named SAMPLE.ZOSYS.TARGET.CSI on volume SMPVOL in catalog CATALOG.MVSICF1.VMVSRES. The SMP/E level is 03.06.00, and the SMP/E command executed was APPLY.

```
              1         2         3         4         5         6         7         8         9         1         1         1
                                                                                              0         1         2
----+----0----+----0----+----0----+----0----+----0----+----0----+----0----+----0----+----0----+----0----+----0----+----0
S0APPLY  03060000000
SAMPLE.ZOSYS.TARGET.CSI                                      CATALOG.MVSICF1.VMVSRES                    SMPVOL
```

Figure 78. Example of SMP/E environment record type 0

T0 - Trailer record type 0

One Trailer Record Type 0 (T0) is created for each delta produced by the results of APPLY or RESTORE processing. The purpose of the T0 record is to uniquely identify the end of a set of library change records.

The format and contents of this record type are shown in Table 52 on page 376.

Table 52. Trailer record type 0

Field name	Position (decimal)	Length (decimal)	Description
Record type	1	2	The characters 'T0'.
Zone name	3	7	The target zone name that was operated on by the APPLY or RESTORE command. It is left-justified and padded with blanks.

Time stamp	10	13	The date and time that APPLY or RESTORE processing completed for the entire delta. This field is in the form *yyyydddhhmmss*, where:
			
			yyyy
			    year
			ddd
			    day
			hh
			    hours
			mm
			    minutes
			ss
			    seconds
			
			Time is represented in 24-hour clock format (military time). The time value matches the time value in the T0 record and is the time that the GIM20501I message is issued for the associated APPLY or RESTORE command.

ERROR count	23	6	The character representation of the count of SYSMODs applied or restored in error. The data is padded on the left with character zeroes (0).
INCMPLT count	29	6	The character representation of the count of SYSMODs left with a status of INCMPLT by the APPLY or RESTORE. The data is padded on the left with character zeroes (0).
APPLIED or RESTORED count	35	6	The character representation of the count of SYSMODs with a status of APPLIED or RESTORED. The data is padded on the left with character zeroes (0).

Table 52. Trailer record type 0 (continued)			
Field name	**Position (decimal)**	**Length (decimal)**	**Description**
DELETED count	41	6	The character representation of the count of SYSMODs with a status of DELETED. The data is padded on the left with character zeroes (0).
SUPD count	47	6	The character representation of the count of SYSMODs with a status of SUPD. The data is padded on the left with character zeroes (0).

For example, shows a T0 record for a target zone named MYTZN. The delta processing for this zone ended on January 16, 2008 at 10:12:34 AM. During this execution of SMP/E, one SYSMOD was applied in error, four SYSMODs were left incomplete, ten SYSMODs were applied successfully, one SYSMOD was deleted, and two SYSMODs were superseded.

```
          1         2         3         4         5         6         7         8
----+----0----+----0----+----0----+----0----+----0----+----0----+----0----+----0
T0MYTZN   200801610123400000100000400001000000010000002
```

Figure 79. Trailer record type 0

Valid action types

Action
> **Meaning**

ADDREP
> The ADDREP action indicates either that a new element or LMOD was added to a target library during APPLY or RESTORE processing, or that an existing element or LMOD was replaced in a target library during APPLY or RESTORE processing
>
> This includes elements added or replaced by way of:
>
> - Data element MCS
> - Hierarchical file system element MCS
> - Link-edits or copies of load modules:
>   - ++MOD MCS (add or replace)
> - ++MAC MCS
> - ++JAR MCS
> - ++JARUPD MCS
> - ++MACUPD MCS
> - ++MOVE MCS for MACs, SRC, and LMODs
> - ++PROGRAM MCS
> - ++RENAME MCS for load modules
> - ++SRC MCS
> - ++SRCUPD MCS
> - ++ZAP MCS
>
> This also includes aliases added or replaced by way of:
>
> - Data element MCS with ALIAS operand
> - Hierarchical file system element MCS with LINK operand

- Link-edits or copies of load modules:
  - ++MOD MCS with TALIAS operand
  - LMODs defined in JCLIN
- ++JAR or ++JARUPD MCS with LINK operand
- ++MAC MCS with MALIAS operand
- ++MACUPD MCS with MALIAS operand
- ++MOVE MCS for MACs and LMODs
- ++PROGRAM MCS with ALIAS operand
- ++RENAME MCS for load modules

DELETE

The DELETE action indicates that an element or LMOD was deleted from a target library during APPLY processing.

This includes elements and their aliases deleted by way of:

- DELETE operand on data element MCS
- DELETE operand on hierarchical file system element MCS
- DELETE operand on ++JAR MCS
- DELETE operand on ++MAC MCS
- DELETE operand on ++PROGRAM MCS
- DELETE operand on ++SRC MCS
- DELETE operand on ++VER MCS
- ++DELETE MCS for load modules

This also includes aliases deleted by way of:

- ALIAS operand on ++DELETE MCS for load modules

Usage recommendations

Users cannot simply look at the records created by the library change interface and move the identified parts. As is the case whenever SMP/E is executed, each processed SYSMOD has a specific completion status when the particular APPLY or RESTORE completes. Therefore, the results of each APPLY or RESTORE must be looked at collectively. Until a set of SYSMODs has been successfully processed, you cannot distribute the library change records associated with those SYSMODs.

Therefore, you must first get a successful APPLY CHECK or RESTORE CHECK before attempting to use the records produced by library change processing for a specified target zone. Since there are no records produced during CHECK processing, there are no records to be lost until delta processing is activated.

Once activated, delta processing produces records for an execution of APPLY or RESTORE. Since all SYSMODs may not successfully complete in one run, the records from each run against the same set of SYSMODs must be kept.

Once a successful APPLY or RESTORE is obtained, the net result of the processing must be analyzed by elements, LMODs, or both to determine the appropriate action to take to distribute the elements and LMODs.

Alternatively, once a successful APPLY or RESTORE is obtained, the net result of the libraries updated can be analyzed to determine the appropriate action to take to distribute the libraries affected by the APPLY/RESTORE activity.

Chapter 9. SMP/E exit routines

This chapter describes intended Programming Interfaces that allow the customer to write programs to obtain the services of IBM SMP/E for z/OS, V3R6.

This chapter describes how to write exit routines that:

- Process statements in SMPPTFIN at RECEIVE time
- Control retry processing when data sets run out of space during ACCEPT, APPLY, GZONEMERGE, LINK LMODS, LINK MODULE, RECEIVE, or RESTORE processing.

LINK LMODS,

A common parameter list is used to pass information between SMP/E and the exit routines. It is pointed to by register 1 and is mapped by macro GIMMPUXP in SYS1.MACLIB. Table 53 on page 379 shows the format and contents of the parameter list.

Table 53. GIMMPUXP: exit routine parameter list				
Field name	**Offset (DEC)**	**Offset (HEX)**	**Length**	**Description**
UXPUXNUM	+ 0	+ 0	2	Exit routine number: X'0001' – RECEIVE X'0002' – RETRY
	+ 2	+ 2	2	Not used
UXPUXNAM	+ 4	+ 4	8	Name of exit routine
UXPUXAD	+12	+ C	4	Address of exit routine
UXPFUNCT	+16	+10	8	SMP/E command
UXPPRMAD	+24	+18	4	Address of exit routine parameter list
UXPLOJAD	+28	+1C	4	Not used
UXPLOEAD	+32	+20	4	Not used
UXPCTBAD	+36	+24	4	Reserved for future use
UXPMODAD	+40	+28	4	Reserved for future use

The following sections describe the exit routines for RECEIVE and retry processing.

RECEIVE exit routine

The RECEIVE exit routine allows you to scan statements in the SMPPTFIN data set at RECEIVE time. This exit routine must be defined by a GIMEXITS control statement in the GIMEXITS member of SMPPARM, which tells SMP/E RECEIVE processing that an exit routine exists and should be called. Optionally, the GIMEXITS control statement may also specify the data set containing the exit routine. This exit routine must be a load module residing in an authorized library. For more information about specifying exit routines in GIMEXITS, see Chapter 3, "Defining control statements in SMPPARM members," on page 103. The RECEIVE exit routine is loaded at the start of RECEIVE command processing and is deleted at the end of RECEIVE processing.

When this exit routine is called, the parameter list contains the values shown in Table 54 on page 380 and Table 55 on page 380.

Table 54. RECEIVE exit routine: parameter list values

Field name	Description
UXPUXNUM	X'0001' (exit routine number)
UXPUXNAM	Name of exit routine
UXPUXAD	Address of exit routine
UXPFUNCT	RECEIVE
UXPPRMAD	Address of 81-byte buffer area (see Table 55 on page 380)

Table 55. RECEIVE exit routine: buffer passed by UXPPRMAD

Field name	Offset (DEC)	Offset (HEX)	Length	Description
UX001RC	+ 0	+ 0	1	X'00' – Buffer contains record to be processed X'04' – End-of-file on SMPPTFIN
UX001RCD	+ 1	+ 1	80	Record from SMPPTFIN

According to the input record, the RECEIVE exit may decide to continue RECEIVE processing, change the record, insert data after the record, or skip the record. Or, it may choose to stop processing for the SYSMOD, for the RECEIVE command, or for SMP/E.

When the exit routine returns control to SMP/E, it must set register 15 to one of the following values:

Value	Meaning
0	Continue normal RECEIVE processing.
8	Stop SYSMOD processing. SMP/E does not receive this SYSMOD, but continues to pass records from the SYSMOD to the exit routine.
12	Stop RECEIVE processing.
16	Stop SMP/E processing.
20	Insert a record after the current one in the buffer.
24	Skip the record in the buffer area.

If any other value is returned, SMP/E issues an error message and fails.

Changing records within SMPPTFIN

To change the current record within the SMPPTFIN data set, the exit routine must change the record currently in UX001RCD and set the return code in register 15 to 0 or 20. A return code of 0 indicates to SMP/E that the changed record should be processed and normal processing should continue. A return code of 20 indicates to SMP/E that the exit routine inserts records after this changed record. For more information, see "Inserting records within SMPPTFIN" on page 380.

Inserting records within SMPPTFIN

To insert one or more records within the SMPPTFIN data set after the current record in UX001RCD, the exit routine must set the return code in register 15 to 20.

SMP/E first processes the current record in UX001RCD; then it calls the exit routine again. On this call, the exit routine must place the record to be inserted into UX001RCD. If more records are to be inserted, the

exit routine must set the return code in register 15 to 20. If this is the last record to be inserted, the exit routine must set the return code to 0.

When the exit routine passes a return code of 0, SMP/E processes the record in UX001RCD; then it continues processing with the next record in the SMPPTFIN data set.

Inserting records at the end of SMPPTFIN

To insert one or more records at the end of the SMPPTFIN data set, the exit routine must set the return code in register 15 to 20.

SMP/E then calls the exit routine again. On this call, the exit routine must place the record to be inserted into UX001RCD. If more records are to be inserted, the exit routine must set the return code in register 15 to 20. If this is the last record to be inserted, the exit routine must set the return code to 0. SMP/E processes the inserted records as if they were at the end of the SMPPTFIN data set.

When the exit routine passes a return code of 0, SMP/E processes the record in UX001RCD; then it continues with normal end-of-file processing.

Note: When you specify the SELECT operand during RECEIVE processing, and all selected SYSMODs have been processed, SMP/E indicates end-of-file to the exit routine. If the exit routine passes a return code of 20 at this point, the add request is ignored, and SMP/E continues with end-of-file processing.

Skipping records in SMPPTFIN

To skip a record in SMPPTFIN, the exit routine must set the return code in register 15 to 24. SMP/E does not process the current record in UX001RCD; instead, it processes the next record in SMPPTFIN.

Retry exit routine

The retry exit routine enables you to control retry processing when a data set runs out of space during ACCEPT, APPLY, GZONEMERGE, LINK LMODS, LINK MODULE, RECEIVE, or RESTORE processing. (In retry processing, the data set is compressed and the utility that failed is called again.) This exit routine must be defined by a GIMEXITS control statement in the GIMEXITS member of SMPPARM, which tells SMP/E retry processing that an exit routine exists and should be called. Optionally, the GIMEXITS control statement may also specify the data set containing the exit routine. This exit routine must be a load module residing in an authorized library. For more information about specifying exit routines in GIMEXITS, see Chapter 3, "Defining control statements in SMPPARM members," on page 103. The retry exit routine is loaded at the start of ACCEPT, APPLY, GZONEMERGE, LINK LMODS, LINK MODULE, RECEIVE, or RESTORE command processing and is deleted at the end of command processing.

Note: The processing of this routine is not affected by the debatching SMP/E does after retry processing fails, because this routine is called as part of the initial retry processing, **before** debatching is attempted.

When a data set runs out of space, an *x*37 abend occurs. If SMP/E determines that a retry can be attempted, it cancels the abend dump and calls the retry exit routine. The parameter list contains these values:

Table 56. Retry exit routine: parameter list values	
Field name	**Description**
UXPUXNUM	X'0002' (exit routine number)
UXPUXNAM	Name of exit routine
UXPUXAD	Address of exit routine
UXPFUNCT	ACCEPT, APPLY, GZMRG, LINK, RECEIVE, or RESTORE
UXPPRMAD	Address of 25-byte parameter list (see Table 57 on page 382)

Table 57. Retry exit routine: parameter list passed by UXPPRMAD				
Field name	**Offset (DEC)**	**Offset (HEX)**	**Length**	**Description**
UX002DDN	+ 0	+ 0	8	ddname of data set that ran out of space
UX002PGM	+ 8	+ 8	8	Name of the utility program that failed
UX002ACH	+16	+10	3	ABEND code (hex) in the same format as field SDWACMPC in the SDWA control block
UX002RCH	+19	+13	1	ABEND reason code (hex)
UX002ACP	+20	+14	3	ABEND code (EBCDIC)
UX002RCP	+23	+17	2	ABEND reason code (EBCDIC)

According to the input, the retry exit can either cancel retry processing or perform some other method of recovery.

When the exit routine returns control to SMP/E, it must set register 15 to one of the following values:

Value	Meaning
0	Continue normal retry processing.
12	Stop command processing and perform no retry.
16	Stop SMP/E processing and perform no retry.
20	Perform modified retry processing. Call the failing utility, but do not compress the failing data set.

If any other value is returned, SMP/E converts it to 12, and command processing fails.

Chapter 10. JCL statements required to invoke SMP/E

Unless you are using the SMP/E dialogs, you must provide the following JCL statements to invoke SMP/E:

- A JOB statement
- An EXEC statement
- DD (data definition) statements

JOB statement

The JOB statement describes your installation-dependent parameters. The JOB statement (or the EXEC statement, or both) can also include the REGION parameter to set the size of the region in which SMP/E runs. For details, see *z/OS MVS JCL User's Guide* or *z/OS MVS JCL Reference*.

Note: To enable the SMP/E job step to get the maximum space above 16 megabytes. you can specify REGION=0M.

EXEC statement

The EXEC statement must specify PGM=GIMSMP or the name of your cataloged procedure for calling SMP/E. (For an example of a cataloged procedure, see *z/OS SMP/E User's Guide*.) The following can be specified in the EXEC statement PARM parameter:

COMPAT=WARNBYPASS or
COMPAT=NOWARNBYPASS
    The COMPAT parameter is used to control incompatible behaviors of SMP/E processing.

    COMPAT(WARNBYPASS)
        indicates that the APPLY and ACCEPT commands will issue warning messages to identify bypassed SYSTEM HOLD exceptions. This is the behavior for releases of SMP/E prior to V3R5.

    COMPAT(NOWARNBYPASS)
        indicates that the APPLY and ACCEPT commands will issue informational messages to identify bypassed SYSTEM HOLD exceptions. If neither COMPAT(WARNBYPASS) or COMPAT(NOWARNBYPASS) is specified, the default is COMPAT(NOWARNBYPASS).

CSI=*dsname*
    where *dsname* is the name of the CSI data set containing the global zone. (This data set is also known as the *master CSI*.) This parameter is used to enable SMP/E to allocate the master CSI data set dynamically.

    Note: If there is an SMPCSI DD statement, the CSI=*dsname* operand is not allowed. If both are specified, SMP/E does **not** run.

DATE=*date*
    where *date* can be one of the following:

    U or IPL
        To use the IPL date of the system.

    REPLY
        To request the date from the operator. As a result, SMP/E issues message GIM399I.

    yyddd
        To specify a specific date, where *yy* is the year and *ddd* is the day of the year (the Julian date).

    If DATE is not specified, the IPL date of the system is used.

LANGUAGE=*xxx*
    where *xxx* can be one of the following:

ENU
US English

JPN
Japanese

The LANGUAGE option defines which language to use for SMP/E messages.

LANGUAGE can also be specified as L. If LANGUAGE is not specified, the default is LANGUAGE=ENU.

You can specify LANGUAGE=JPN only if you have installed the Japanese language feature of SMP/E. If you have installed the Japanese language feature, you can specify LANGUAGE=ENU or LANGUAGE=JPN. (You do not need to install the English feature along with the Japanese feature.)

The output devices used must support the selected language and English single-byte characters. SMP/E does not check to verify that the output devices provide this support.

PROCESS=WAIT or
PROCESS=END
The PROCESS parameter is used to control how long a job should wait if a CSI or PTS data set is not immediately available because it is currently being used either by another job or by a dialog.

- WAIT causes the job to wait until the data set is available. A message is issued to the system operator every 30 minutes while the job is waiting.

- END causes the job to wait for 10 minutes. If the data set is still not available after the 10-minute wait, the command requiring the data set is stopped.

If PROCESS is not specified, the default is PROCESS=WAIT.

For more information about obtaining and sharing CSI data sets, see the "Sharing SMP/E Data Sets" appendix in *z/OS SMP/E Commands*.

Processing of the PTS data set is also affected by the WAITFORDSN value specified in its DDDEF entry. WAITFORDSN determines whether SMP/E should wait to allocate a data set that is not immediately available. If the DDDEF entry specifies WAITFORDSN=NO (or lets this value default to NO) and the data set is not available, allocation of the data set fails, regardless of the PROCESS value specified on the EXEC statement. If WAITFORDSN=NO, SMP/E does not wait to retry allocation of the data set.

For example, suppose a PTS with a disposition of OLD is already being used by a job, and a second job tries to access the same PTS data set by allocating it through a DDDEF entry. The DDDEF entry used by the second job for the PTS specifies WAITFORDSN=NO. As a result, allocation of the PTS fails for the second job.

DD statements

DD statements define the data sets that can be used in SMP/E processing. For information about the data sets required for each command, see the chapters on individual SMP/E commands in *z/OS SMP/E Commands*.

Note: You can use DDDEF entries, rather than DD statements, to allocate many of the necessary data sets. For more information, see "DDDEF entry (distribution, target, and global zone)" on page 163.

Chapter 11. Service routines

SMP/E provides these service routines:

- GIMCPTS
- GIMDTS
- GIMGTPKG
- GIMUNZIP
- GIMXSID
- GIMXTRX
- GIMZIP

GIMCPTS: SYSMOD compaction service routine

The GIMCPTS service routine is used to compact or expand inline element data within SYSMODs. The inline element data can be compacted prior to receiving a SYSMOD or SYSMODs already received can be compacted.

SYSMODs with their inline element data compacted require less space, so the space requirements of the SMPPTS data set should be reduced. Of course, SMP/E APPLY and ACCEPT processing will expand any compacted inline element data prior to installing the element in a target or distribution library respectively.

GIMCPTS can also be used to expand the inline element data within a SYSMOD. This could be useful if you need to send the SYSMOD to another system where the SMP/E installed on the other system cannot handle compacted inline element data. It could also be useful for viewing inline data. Note that the SMP/E Query Dialogs can be used to view MCS entries from the SMPPTS data set and that SMP/E expands the inline element data before displaying it.

Note: GIMCPTS is a separate load module residing in the MIGLIB library that runs independently from the rest of SMP/E processing.

Calling GIMCPTS

Following are the JCL statements needed to call GIMCPTS:

```
//JOBx     JOB …
//STEP1    EXEC PGM=GIMCPTS,PARM='options'
//SYSPRINT DD SYSOUT=*
//SYSUT1   DD DSN=aaaaaaaa,DISP=SHR
//SYSUT2   DD DSN=bbbbbbbb,DISP=OLD
```

Figure 80. JCL to call GIMCPTS

EXEC

    is the statement used to call GIMCPTS. The EXEC statement must specify PGM=GIMCPTS. The following options may be specified on the EXEC statement PARM operand:

    COMPACT

        The COMPACT option indicates to GIMCPTS the input data is to be compacted. COMPACT is mutually exclusive with EXPAND. If neither COMPACT nor EXPAND is specified, the default is COMPACT.

EXPAND

The EXPAND option indicates to GIMCPTS the input data is to be expanded. EXPAND is mutually exclusive with COMPACT. If neither EXPAND or COMPACT is specified, the default is COMPACT.

LANGUAGE=*xxx*

where *xxx* can be one of the following:

ENU

US English

JPN

Japanese

The LANGUAGE option defines which language to use for GIMCPTS messages. LANGUAGE can also be specified as L. If LANGUAGE is not specified, the default is LANGUAGE=ENU.

SYSPRINT

is used by GIMCPTS for messages.

The record format (RECFM) of SYSPRINT must be FBA, and the record length (LRECL) may be either 81 or 121. GIMCPTS will format messages to 80 characters or 120 characters in length, based on the record length of the SYSPRINT data set.

SYSUT1

points to the sequential or partitioned input data. Sequential input data may be a sequential data set or a single member of a partitioned data set. Partitioned input data is an entire partitioned data set (PDS) whose members each contain input data. The actual input data must be the MCS of one or more SYSMODs. The inline element data within the SYSMODs will be either compacted or expanded, as determined by the COMPACT and EXPAND options.

The record format (RECFM) of SYSUT1 must be fixed (F), fixed-block (FB), or fixed-block standard (FBS), and the record length (LRECL) must be 80.

SYSUT2

points to the sequential or partitioned output destination. If the input is sequential, then SYSUT2 must also be sequential. If the input is partitioned, then SYSUT2 must also be partitioned. The output will contain the input SYSMODs with either compacted or expanded inline element data.

The record format (RECFM) of SYSUT2 must be fixed (F), fixed-block (FB), or fixed-block standard (FBS), the record length (LRECL) must be 80, and the BLKSIZE value must be a multiple of 80.

IBM recommends that SYSUT2 be allocated the same size as the original SMPPTS data set. The amount of compaction that can be achieved will vary depending on the content and packaging method of the SYSMOD in the originating SMPPTS data set. For example, if the originating SMPPTS data set contained only functions packaged in RELFILE format, no compaction would occur.

After GIMCPTS is run, you can release the unused space using ISPF (or other methods). You could also specify RLSE on the SPACE parameter of the SYSUT2 data set (if it is being allocated when GIMCPTS is being run) and z/OS will free unused space when it is closed.

Example

Suppose your system has an existing SMPPTS data set that is quite large and requires the space of an entire physical volume. You can use the SMP/E GIMCPTS compaction service routine to compact data within the members of the SMPPTS. In this example, the existing SMPPTS data set is compacted in place. However, a new data set could be allocated and used to receive the compacted data.

```
//JOBx     JOB …
//COMPACT  EXEC PGM=GIMCPTS,PARM='COMPACT,LANGUAGE=ENU'
//*
//SYSPRINT DD SYSOUT=*
//SYSUT1   DD DSN=SMP.ZOSR1.SMPPTS,DISP=SHR
//SYSUT2   DD DSN=SMP.ZOSR1.SMPPTS,DISP=OLD
```

Figure 81. Sample GIMCPTS job stream

When this is executed, the GIMCPTS service routine will compact the inline element data within all SYSMOD members in the SMPPTS data set and rewrite them to the SMPPTS data set.

Note: You should only attempt to compact a partitioned data set in place if it is a PDSE. If the data set is a simple PDS however, it is likely the data set will get an out of space condition during the operation, unless a large amount of free space is available.

Return codes

To help you diagnose errors, GIMCPTS issues messages and return codes. The messages are documented in *z/OS SMP/E Messages, Codes, and Diagnosis*. Here is a description of the return codes:

Return code	Meaning
0	The input data was processed successfully.
8	One of the following:
	• Syntax errors were found in an input SYSMOD.
	• An unsupported compaction dictionary was used to compact data.
	• Compacted data is incomplete.
12	One of the following:
	• Data sets were missing.
	• Data sets could not be opened.
	• Data set organization of SYSUT1 and SYSUT2 is not sequential or partitioned.
	• Data set organizations of SYSUT1 and SYSUT2 did not match.
	• SYSUT2 data set had incorrect RECFM or LRECL.
	• SYSUT2 and SYSUT1 data sets are the same data set and they are sequential or members of a PDS.
	• Directory space for SYSUT2 was exceeded.
	• Compression and expansion services are not supported.
16	One of the following:
	• An I/O error occurred.
	• Syntax error was found on the EXEC statement parameters.
20	SYSPRINT data set is missing.
> 20	Internal error. Report the error to the IBM Support Center.

GIMDTS: Data transformation service routine

Elements may have any of a variety of record formats, depending on how they are meant to be used. However, when elements are packaged inline in a SYSMOD, they must contain fixed-block 80 records. To help you package inline elements, SMP/E provides the GIMDTS service routine. GIMDTS is a background utility program that transforms data into fixed-block 80 records. For example, it can be used to format

inline replacements for data elements. Although GIMDTS is packaged as part of SMP/E, it is a separate load module residing in SYS1.MIGLIB and runs independently from the rest of SMP/E processing.

The input for GIMDTS must meet these requirements:

- It must be a sequential data set or a member of a partitioned data set (PDS).
- It can contain either variable-length or fixed-length records.

The output from GIMDTS is in this format:

- It is a sequential data set or a member of a partitioned data set.
- It has these attributes:

  RECFM=FB
  LRECL=80
  BLKSIZE=a multiple of 80

After using GIMDTS to transform an element into the required format, you can package the transformed data inline in a SYSMOD, following the associated data element MCS. Later, when the element is installed, it is changed back to its original format.

The following sections describe:

- Statements used to call GIMDTS
- Processing done by GIMDTS
- Return codes issued by GIMDTS

Calling GIMDTS

Here are the JCL statements that are required to call GIMDTS:

```
//JOBx     JOB 'account','name',MSGLEVEL=(1,1)
//STEP1    EXEC PGM=GIMDTS
//SYSPRINT DD SYSOUT=A
//SYSUT1   DD DSN=aaaaaaa,DISP=SHR
//SYSUT2   DD DSN=bbbbbbb,DISP=SHR
```

Figure 82. Sample GIMDTS job stream

EXEC

is the statement used to call GIMDTS.

SYSPRINT
is used by GIMDTS for messages.

SYSUT1
points to the sequential data set or PDS member containing the data to be transformed.

A file in a file system can be copied to a sequential data set to be used as input to GIMDTS.

Note: A binary file must be copied with the binary option.

The record format (RECFM) must be F, FA, FM, FB, FBA, FBM, V, VA, VM, VB, VBA, VBM, VS, or VBS.

SYSUT2
points to the sequential data set or PDS member that will contain the transformed data. This data set must be on DASD.

The record format of this file should be fixed-block (RECFM=FB) with a BLKSIZE value that is a multiple of 80.

Processing

GIMDTS first checks for the required data sets, opens the SYSUT1 and SYSUT2 data sets, and makes sure these data sets have the correct attributes. If necessary, GIMDTS changes the RECFM value of the SYSUT2 data set to FB and the BLKSIZE value to 3200. Next, GIMDTS reads records from the SYSUT1 data set and transforms the data.

After processing the data, GIMDTS writes the output records to the SYSUT2 data set. GIMDTS also issues messages to SYSPRINT indicating the input data set, the output data set, and the return code. If any errors occur during processing, GIMDTS issues messages describing the errors and stops processing. If an abend occurs, normal system abend processing is done.

If SYSUT1 is already fixed-block with LRECL=80 and the data has no "++" characters in columns 1 and 2, then the data needs no transformation and is simply copied from SYSUT1 to SYSUT2.

Return codes

GIMDTS may issue the following return codes:

Return code	Meaning
00	The input data was processed successfully.
04	The SYSUT2 data set was reblocked.
08	The input data was not processed successfully.
16	An I/O error occurred.

GIMGTPKG service routine

The GIMGTPKG service routine is used to get a GIMZIP package from a remote FTP or HTTP(S) server in a TCP/IP network and store the package on a local z/OS host. GIMGTPKG performs the functions of the SMP/E RECEIVE FROMNETWORK TRANSFERONLY command, but does so independently of SMP/E. It either uses FTP or HTTP(S) to transport the files of a GIMZIP package from a remote server to a local host, thus it will provide the following capabilities (which are also available with the RECEIVE FROMNETWORK command):

- Industry standard FTP and HTTP(S) protocols.
- Secure transmission using the capabilities of the z/OS FTP client and Java 2 Technology Edition Version 6 (or higher).
- Ensured integrity of the transported files.

GIMGTPKG requires input to define the FTP or HTTP(S) server and package to be obtained and to control local host processing. GIMGTPKG depends on either the z/OS FTP client or Java 6 for its transfer operations, and it also depends on either ICSF or Java 6 to perform SHA-1 hash calculations. For Java 6, SMP/E specifically requires the IBM 31-bit SDK for z/OS, Java Technology Edition, V6 (5655-R31), or the IBM 64-bit SDK for z/OS, Java Technology Edition, V6 (5655-R32), or a logical successor.

GIMGTPKG is an independent load module residing in the MIGLIB library. It has an addressing mode of 31 (AMODE=31), a residence mode of 24 (RMODE=24), and requires no special authorization (AC=0). GIMGTPKG runs independently from the rest of SMP/E processing.

Calling GIMGTPKG

Following are the JCL statements required to call GIMGTPKG.

GIMGTPKG

```
//job      JOB ...
//step     EXEC PGM=GIMGTPKG,PARM='options'
//SMPOUT   DD SYSOUT=*
//SYSPRINT DD SYSOUT=*
//SMPNTS   DD PATH='package_directory',PATHDISP=KEEP
//SMPCPATH DD PATH='smpclasses_directory', PATHDISP=KEEP
//SMPJHOME DD PATH='javaruntime_directory', PATHDISP=KEEP
//SMPSRVR  DD *
  server tag and attributes
/*
//SMPCLNT  DD *
  client tag and attributes
/*
```

Figure 83. JCL to call GIMGTPKG

EXEC

is the statement used to call GIMGTPKG. The EXEC statement must specify PGM=GIMGTPKG. The following option may be specified on the EXEC statement PARM operand:

LANGUAGE=*xxx*

where *xxx* can be either ENU (US English) or JPN (Japanese). The LANGUAGE option defines which language to use for GIMGTPKG messages. LANGUAGE can also be specified as L. If LANGUAGE is not specified, the default is LANGUAGE=ENU.

SMPOUT

used for GIMGTPKG messages. SMPOUT may be sequential, a member of a partitioned data set, or a file in the UNIX file system. See "SMPOUT" on page 126 for information about allocating SMPOUT. GIMGTPKG will format messages to 80 characters in width, regardless of the LRECL for SMPOUT.

SYSPRINT

is used for FTP and HTTP(S) output. SYSPRINT may be sequential, a member of a partitioned data set, or a file in the UNIX file system. See "SYSPRINT" on page 137 for information about allocating SYSPRINT. GIMGTPKG will format output to 80-characters in width, regardless of the LRECL for SYSPRINT.

SMPNTS

specifies a directory in the UNIX file system where the package to be obtained is to be stored.

SMPCPATH

specifies a directory in the UNIX file system where the SMP/E Java classes reside. The SMPCPATH DD statement is optional and only required to calculate SHA-1 hash values if ICSF is not available.

SMPJHOME

specifies a directory in the UNIX file system where the Java runtime resides. The SMPJHOME DD statement is optional and only required to calculate SHA-1 hash values if ICSF is not available.

SMPSRVR

provides the necessary information about either the FTP or HTTP(S) server from which a package is to be obtained and about the package itself. This package and server information is specified with the <SERVER> tag. For information about using the <SERVER> tag, see *z/OS SMP/E Commands*.

SMPSRVR may be sequential, a member of a partitioned data set, or a file in the UNIX file system. See "SMPSRVR" on page 130 for information about allocating SMPSRVR.

SMPCLNT

provides information about the local FTP or HTTP(S) client host, such as choice of download method (FTP or HTTP(S)), and information about local firewall navigation. The local client host information is specified with the <CLIENT> tag. For information about using the <CLIENT> tag, see *z/OS SMP/E Commands*.

SMPCLNT is an optional DD statement. SMPCLNT may be sequential, a member of a partitioned data set, or a file in the UNIX file system. See "SMPCLNT" on page 115 for information about allocating SMPCLNT.

Example of using GIMGTPKG

```
//job       JOB ...
//step      EXEC PGM=GIMGTPKG
//SMPOUT    DD SYSOUT=*
//SYSPRINT  DD SYSOUT=*
//SMPNTS    DD PATH='/u/usr01/pkgs/',PATHDISP=KEEP
//SMPCPATH  DD PATH='/usr/lpp/smp/classes/',PATHDISP=KEEP
//SMPJHOME  DD PATH='/usr/lpp/java/J6.0/',PATHDISP=KEEP
//SMPSRVR   DD *
<SERVER host="host.sample.com"
        user="usr01"
        pw="n0peekng">
  <PACKAGE file="CBPROC/012345/RIMTAPE/GIMPAF.XML"
           hash="12345678901234567890123456789012345678​90"
           id="012345">
  </PACKAGE>
</SERVER>
/*
//SMPCLNT   DD *
<CLIENT retry="3">
</CLIENT>
/*
```

Figure 84. GIMGTPKG example

In this example, a GIMZIP package residing on the host.sample.com FTP server is transferred onto the local z/OS host and stored in the `/u/usr01/pkgs/` directory. The <SERVER> tag and attributes in the SMPSRVR data set describe the FTP server and the package to be obtained. The <CLIENT> tag and attributes in the optional SMPCLNT data set indicate FTP operations should be retried a maximum of three times if necessary in order to obtain full and accurate package files.

GIMGTPKG processing

GIMGTPKG checks the required and optional DD statements that were specified, opens the data sets, and ensures they have the correct attributes.

The remainder of GIMGTPKG processing is much like that for the RECEIVE FROMNETWORK command with TRANSFERONLY (see the RECEIVE Command chapter in *z/OS SMP/E Commands*). Briefly, this processing includes the following activities:

- The SMPSRVR and SMPCLNT data sets are read and analyzed. For comparison to the RECEIVE FROMNETWORK command, SMPSRVR is analogous to the SERVER data set and SMPCLNT is analogous to the CLIENT data set.

- The package directory is a specific subdirectory of the SMPNTS directory and is the destination on the local host for the GIMZIP package files to be stored. The name of this subdirectory is the package-id value from the SMPSRVR data set. If the package directory does not already exist, then it is created.

- Each of the files in the GIMZIP package on the FTP or HTTP(S) server are retrieved and stored in the package directory created in the previous step. Each file is verified by computing an SHA-1 hash value for the file.
- After all files have been transferred successfully, then, if any archive files in the package have been segmented, the segments are combined to create complete archive files.

Return codes

GIMGTPKG may complete with any of the following return codes.

Return code
Meaning

00

GIMGTPKG processing completed successfully.

12

One of the following:

- Required data sets were missing or could not be opened.
- Required modules could not be loaded.
- Syntax error was found in the SMPSRVR or SMPCLNT tags.
- An FTP or HTTP(S) error was detected.
- Neither ICSF or Java 2 Version 1 Release 4 is available.

16

One of the following:

- An I/O error occurred.
- Syntax error was found on the EXEC statement parameters.

20

The SMPOUT data set is missing.

GIMUNZIP file extraction service routine

The GIMUNZIP service routine is used to extract data sets, files, and directories from archive files in GIMZIP packages created by the GIMZIP service routine. These packages typically contain software and associated materials in the form of SYSMODs, RELFILE data sets, HOLDDATA, and other materials such as documentation, samples, and text files. These GIMZIP packages may be transported through a network, processed by the GIMUNZIP service routine, and then processed by the SMP/E RECEIVE command.

More specifically, the GIMUNZIP service routine extracts data sets, files, and directories from the archive files that compose the GIMZIP package. An archive file consists of a portable image of a sequential, partitioned, or VSAM data set, or a file or directory in a UNIX file system, and the information needed to create that data set, file, or directory from the portable image. The data set, file, or directory into which the archive file is to be extracted can already exist or GIMUNZIP can create a new one of the appropriate type.

Note:

1. GIMUNZIP is a separate load module residing in the MIGLIB library and runs independently from the rest of SMP/E processing.
2. GIMUNZIP optionally requires either the Integrated Cryptographic Services Facility (ICSF) One-Way Hash Generate callable service or Java 8, or a logical successor, to be available for its use in order to compute an SHA-1 hash value.

Calling GIMUNZIP

The following figure shows the job control statements for GIMUNZIP:

```
//JOBx     JOB …
//STEP1      EXEC PGM=GIMUNZIP,PARM='options'
//SMPDIR     DD PATH='package_directory',PATHDISP=KEEP
//SMPWKDIR   DD PATH='work_directory',PATHDISP=KEEP
//SMPCPATH DD PATH='smpclasses_directory',PATHDISP=KEEP
//SMPJHOME DD PATH='javaruntime_directory',PATHDISP=KEEP
//SMPOUT     DD SYSOUT=*
//SYSPRINT   DD SYSOUT=*
//SYSUT3     DD UNIT=SYSALLDA,SPACE=(CYL,(50,10))
//SYSUT4     DD UNIT=SYSALLDA,SPACE=(CYL,(25,5))
//SMPCLNT    DD *
  client tag and attributes
/*
//SYSIN      DD *
  package control statements
/*
```

Figure 85. JCL to call GIMUNZIP

EXEC

is the statement used to call GIMUNZIP. The EXEC statement must specify PGM=GIMUNZIP. The following options may be specified on the EXEC statement PARM operand:

LANGUAGE=*xxx*

where *xxx* can be one of the following:

ENU
US English

JPN
Japanese

The LANGUAGE option defines the language to use for GIMUNZIP messages. LANGUAGE can also be specified as L. If LANGUAGE is not specified, the default is LANGUAGE=ENU.

HASH=*xxx*

where *xxx* can be one of the following:

YES
indicates GIMUNZIP is to perform SHA-1 hash checking for the archive files specified in the SYSIN data set.

Note: GIMUNZIP requires either the Integrated Cryptographic Services Facility (ICSF) One-Way Hash Generate callable service or Java 2 Version 1 Release 4 to be available in order to compute an SHA-1 hash value.

NO
indicates GIMUNZIP is not to perform SHA-1 hash checking for the archive files specified in the SYSIN data set.

HASH can also be specified as H. If HASH is not specified, the default is HASH=NO.

VERIFYSIG=*xxx*

where *xxx* can be one of the following:

YES
indicates if the package has been signed, then GIMUNZIP will verify the signature for the GIMPAF2.XML file found in the directory specified by the SMPDIR ddname.

NO

indicates GIMUNZIP will not verify the signature for the GIMPAF2.XML file found in the directory specified by the SMPDIR ddname. This is the default if VERIFYSIG is not specified.

SYSIN

specifies a control data set that contains the package control tags to control GIMUNZIP processing (see "GIMUNZIP package control tags" on page 395). The control data set may be a sequential data set or a single member of a partitioned data set, and its attributes must be LRECL=80 and RECFM=F or FB.

SMPDIR

specifies a directory in a UNIX file system that contains a GIMZIP package. This directory is referred to as the package directory.

SMPWKDIR

specifies a directory in a UNIX file system that is used by GIMUNZIP for temporary work files. This is an optional DD statement. If the SMPWKDIR DD statement is not provided, GIMUNZIP will use the package directory specified on the SMPDIR DD statement for temporary work files.

SMPCPATH

specifies a directory in the UNIX file system where the SMP/E Java classes reside. The SMPCPATH DD statement is optional and only required to calculate SHA-1 hash values if ICSF is not available, or to verify a digital signature.

Alternatively, the Java class path may be specified by the classpath attribute of the <CLIENT> tag in the SMPCLNT DD statement.

SMPJHOME

specifies a directory in the UNIX file system where the Java runtime resides. The SMPJHOME DD statement is optional and only required to calculate SHA-1 hash values if ICSF is not available, or to verify a digital signature.

Alternatively, the Java runtime location may be specified by the javahome attribute of the <CLIENT> tag in the SMPCLNT DD statement.

SMPCLNT

provides information about the local client host environment, including the Java home directory, Java class path, Java debug options, and the signaturekeyring value. The local client host information is specified with the <CLIENT> tag. For information about using the <CLIENT> tag, see *z/OS SMP/E Commands*.

SMPCLNT is an optional DD statement. SMPCLNT may be sequential, a member of a partitioned data set, or a file in the z/OS UNIX file system. See "SMPCLNT" on page 115 for information about allocating SMPCLNT.

SMPOUT

specifies an output data set that will contain SMP/E messages issued during GIMUNZIP processing. The data set attributes are LRECL=81 or LRECL=121 and RECFM=FBA.

SYSUT3 and SYSUT4

specifies work areas for the IEBCOPY utility to use for spill data sets on DASD.

SYSPRINT

specifies a data set to contain output and messages from the utilities called by GIMUNZIP.

GIMUNZIP package control tags

GIMUNZIP service routine processing is controlled by package control tags. The package control tags identify the archive files to be extracted by GIMUNZIP and are specified in the SYSIN data set. The package control tags follow XML syntax rules and their format is as follows:

GIMUNZIP package control tags
GIMUNZIP

```
►►─ <GIMUNZIP> ──┬───────────┬──┬── ARCHDEF ◄─┬── </GIMUNZIP> ─►◄
                 └── TEMPDS ──┘  └────────────┘
```

TEMPDS

```
►►─ <TEMPDS ──┬─────────────────────────────┬──┬──────────────────────┬──►
              └─ volume= ─ " ─ data set volume ─ " ─┘  └─ unit= ─ " ─ unit type ─ " ─┘

   ►── > ── </TEMPDS> ─►◄
```

ARCHDEF

```
►►─ <ARCHDEF ──┬── name=" archive name " ──┬──►
               └── archid= ─ " ─ archive id ─ " ─┘

   ►──┬────────────────────────────────────────┬──┬─────────────────────┬──►
      └─ volume=" ─┬── data set volume ──┬─ " ─┘  ├─ newname=" new name " ─┤
                   └──────── * ──────────┘        └─ prefix=" name prefix " ─┘

   ►──┬─────────────────────────┬──┬──────────────────────────┬──►
      └─ unit= ─ " ─ unit type ─ " ─┘  └─ dataclas= ─ " ─ data class ─ " ─┘

   ►──┬──────────────────────────────────────┬──►
      └─ mgmtclas= ─ " ─ management class ─ " ─┘

   ►──┬────────────────────────────────┬──┬─────────────────────┬──►
      └─ storclas= ─ " ─ storage class ─ " ─┘  └─ catalog= ─┬── yes ──┬─┘
                                                            └── no ───┘

   ►──┬─────────────────────┬──┬───────────────────────┬── > ── </ARCHDEF> ─►◄
      └─ replace= ─┬── no ──┬─┘  └─ preserveid= ─┬── no ──┬─┘
                   └── yes ─┘                    └── yes ─┘
```

<GIMUNZIP> Tag syntax

The <GIMUNZIP> and corresponding </GIMUNZIP> tags identify the beginning and end of a set of archive retrieval requests.

<TEMPDS> Tag syntax

The <TEMPDS> and corresponding </TEMPDS> tags identify the beginning and end specification for temporary work data sets. Temporary work data sets are used by GIMUNZIP to contain the IEBCOPY sequential unload form of archived partitioned data sets, and the IDCAMS REPRO sequential unload form of VSAM data sets. The <TEMPDS> tag is optional, but if specified may contain the following attributes:

volume="*volser*"

specifies the volume serial number for a volume to be used when allocating temporary data sets. The `volume` attribute is optional. If specified, the `volume` value must be from 1 to 6 alphanumeric characters.

unit="*unit type*"

specifies the device unit type to use when allocating temporary data sets. The `unit` attribute is optional. If specified, the `unit` value must be from 1 to 8 non-blank characters.

<ARCHDEF> Tag syntax

The <ARCHDEF> and corresponding </ARCHDEF> tags identify the beginning and end of a specific archive retrieval request. The following attributes may be found on the <ARCHDEF> tag:

name="*archive name*"

specifies the path name for an archive file to be extracted by GIMUNZIP. The path name may include a subdirectory name (provided by the `subdir` attribute when the archive was created by GIMZIP). The path name value is a relative value and it is relative to the package directory specified on the SMPDIR DD statement. The name attribute is mutually exclusive with the `archid` attribute.

Note: If the archive file has been segmented in the package, you should specify the name of the archive file, not the names of the archive segment files. GIMUNZIP will combine the associated archive segment files to extract the complete archive.

archid="*archive id*"

specifies a unique archive id associated with an archive file to be extracted by GIMUNZIP. GIMUNZIP will search the package attribute file looking for the archive file with this unique `archid`. The `archid` attribute is mutually exclusive with the archive name attribute.

volume="*data set volume | *"*

specifies the volume serial number of the volume on which GIMUNZIP is to request the allocation of the data set to be extracted from the archive file. The volume identifier must be from 1 to 6 alphanumeric characters or an asterisk (*).

Note:

1. If you are extracting a sequential, partitioned, or VSAM data set into an existing cataloged data set, you should not specify the `volume` attribute.

2. If you are extracting a sequential or partitioned data set:

   - Into a new data set to be allocated on a specific volume, or into an existing, uncataloged data set on a specific volume, you must specify `volume=`*data set volume*. If you specify `volume=*` it is ignored.

   - Into an existing, uncataloged data set on a specific volume, you must specify `volume=`*data set volume*. If you specify `volume=*` it is ignored.

   The data set is first allocated with a disposition of (OLD,KEEP) using the specified volume. If the data set is not on the volume, and if the catalog attribute is not specified or indicates the data set should be cataloged, then the data set is allocated again with a disposition of (OLD,KEEP) without specifying a volume, to see if there is an existing cataloged data set that should be used for the extraction. If an existing cataloged data set is allocated, the archive is extracted into the cataloged data set, even though the volume on which it is allocated is different from the volume specified.

3. If you are extracting a VSAM data set:

   - Into a new data set, you can specify `volume=`*data set volume*,

     Or,

     You can specify `volume=*`. GIMUNZIP uses the value on the VOLUMES parameter on the IDCAMS DEFINE command that GIMUNZIP will generate to create the new VSAM data set. Using the VOLUMES(*) parameter on the generated IDCAMS DEFINE command may allow SMS to

assign a volume for the new VSAM data set if the automatic class selection (ACS) routines allow it.

- Into an existing destination data set, a data set volume you specify on the `volume` attribute is ignored.

    Since VSAM data sets must be cataloged, the catalog is checked first to see if the VSAM destination data set already exists.

4. If you are extracting a UNIX file or directory, you do not need to specify the `volume` attribute. It is ignored. However, if you specify a `volume` attribute, it must be syntactically correct.

newname="*name*"

specifies the data set name or absolute pathname to use for the data set or file to be extracted from the archive file. This name replaces the original name of the data set or file that is recorded in the archive's file attribute file.

The newname specified must identify the same type of data as the file being unzipped. That is, if the file being unzipped is a data set, newname must represent a data set. If the file being unzipped is a file or directory in the UNIX file system, newname must represent a file or directory in the UNIX file system.

The first character of an absolute pathname of a file or directory in the UNIX file system must be a slash (/). If the newname value does not begin with a slash, GIMUNZIP processing will assume that it is an MVS data set name and that it will conform to the MVS data set naming conventions. If the newname value does begin with a slash, it is assumed to be a file or directory in the UNIX file system. When a file or directory in the UNIX file system is specified, the newname value can be from 1 to 1023 bytes long with 255 characters between delimiters (/). The value can be any character from X'40' through X'FE', except less than (<) and greater than (>) signs, ampersands (&), and double quotation marks ("). All data beyond column 72 is ignored, including blanks. The pathname of a file or directory in the UNIX file system is case sensitive and will not be converted to uppercase alphabetic during processing of the GIMUNZIP package control tags.

The newname specified may identify an existing data set, directory, or file. In this case, the archived file is unzipped into the existing data structure specified on the newname attribute. If the newname attribute is not specified for a file or directory in the UNIX file system, the original name for the file or directory is used. If neither the `prefix` nor the newname attributes are specified when a data set is being extracted, the original name for the data set is used.

prefix="*data set prefix*"

specifies a data set prefix to use for the data set to be extracted from the archive file. This prefix replaces the high-level qualifier of the original name for the data set as recorded in the archive's file attribute file. The specified data set prefix can be up to 26 characters long and must conform to standard data set naming conventions.

The resolved name using the `prefix` attribute may be for an existing data set. In this case, the archived data is extracted into the existing data set specified using the `prefix` attribute.

A prefix value has no meaning when the file being unzipped is a file or directory in the UNIX file system. In this case, the prefix value is parsed but ignored.

If neither the `prefix` nor the newname attributes are specified, the original name for the data set is used.

unit="*unit type*"

specifies the device unit type value to use when allocating a new data set to contain data extracted from the archive file. The `unit` attribute is optional. If specified, the `unit` value must be from 1 to 8 non-blank characters.

The `unit` attribute applies only when extracting a new partitioned or sequential data set from an archive. The `unit` attribute is parsed for syntax but otherwise ignored if the target data set already exists, is not a partitioned or sequential data set, or when extracting a UNIX file or directory.

dataclas="*data class*"

specifies the data class to use when allocating a new data set to contain data extracted from the archive file. The `dataclas` attribute is optional. If specified, the `dataclas` value must be from 1 to 8 alphanumeric or national characters, (A-Z, 0-9, #, @, $), and the first cannot be numeric.

The `dataclas` attribute applies only when extracting a new data set from an archive. The `dataclas` attribute is parsed for syntax but otherwise ignored if the target data set already exists, or if extracting a UNIX file or directory.

mgmtclas="*management class*"

specifies the management class to use when allocating a new data set to contain data extracted from the archive file. The `mgmtclas` attribute is optional. If specified, the `mgmtclas` value must be from 1 to 8 alphanumeric or national characters, (A-Z, 0-9, #, @, $), and the first cannot be numeric.

The `mgmtclas` attribute applies only when extracting a new data set from an archive. The `mgmtclas` attribute is parsed for syntax but otherwise ignored if the target data set already exists, or if extracting a UNIX file or directory.

storclas="*storage class*"

specifies the storage class to use when allocating a new data set to contain data extracted from the archive file. The `storclas` attribute is optional. If specified, the `storclas` value must be from 1 to 8 alphanumeric or national characters, (A-Z, 0-9, #, @, $), and the first cannot be numeric.

The `storclas` attribute applies only when extracting a new data set from an archive. The `storclas` attribute is parsed for syntax but otherwise ignored if the target data set already exists, or if extracting a UNIX file or directory.

catalog=="YES | NO"

indicates when allocating a new data set, whether that data set should be cataloged or not. A value of YES indicates the new data set should be cataloged, and a value of NO indicates the new data set should not be cataloged. The `catalog` attribute is optional. If the `catalog` attribute is not specified, the new target data set will be cataloged.

The `catalog` attribute applies only when extracting a new partitioned or sequential data set from an archive. The `catalog` attribute is parsed for syntax but otherwise ignored if the target data set already exists, is not a partitioned or sequential data set, or when extracting a UNIX file or directory.

replace="YES | NO"

indicates whether data in an existing data set, file or directory should be replaced by data from the archive file. A value of YES indicates the data in the existing data set, file or directory should be replaced. A value of NO is equivalent to not specifying the replace attribute and indicates the data should not be replaced. The `replace` attribute is not meaningful when the data set, file or directory does not already exist.

- For an archive of a partitioned data set or directory in the UNIX file system, `replace="YES"` indicates that existing members in a partitioned data set and files in a directory will be replaced by members or files with the same name from the archive file. If `replace="NO"` is specified, or if the replace attribute is not specified at all, existing members and files will not be replaced.

- For an archive of a sequential data set or a file in the UNIX file system, `replace="YES"` indicates the existing data set or file and its attributes will be replaced with the data from the archive file. If `replace="NO"` is specified, or if the replace attribute is not specified at all, the existing sequential data set or file will not be replaced with the data from the archive file.

Notes:

1. GIMUNZIP will not check if a sequential data set or file is empty. If the data set or file exists, it will not be replaced unless `replace="YES"` is specified.

2. For an archive of a sequential data set or a file in the UNIX file system, if the original data set or file was empty and `replace="YES"` is specified, then the contents of the existing data set or file will be lost.

- For a VSAM data set (cluster), `replace="YES"` indicates that an existing VSAM cluster should be populated with the data from the archive file. If `replace="NO"` is specified, or if the replace attribute is not specified at all, there will be no attempt to populate the existing VSAM cluster with the data from the archive.

  Note:

  1. The existing VSAM cluster must be of the same type as the archived cluster (ESDS, KSDS, LDS, or RRDS).

  2. The existing VSAM cluster must have characteristics that are compatible with the archived cluster (such as, record size, key size, and key offset). GIMUNZIP does not verify the compatibility of these characteristics.

  3. To replace the contents of an existing cluster, the cluster is altered to a reusable state by using an IDCAMS ALTER command, if necessary, before the data from the VSAM archive is copied into the cluster by using an IDCAMS REPRO command. The REPRO command will use both the REPLACE and REUSE operands. Following the REPRO operation, the cluster is altered back to a nonreusable state if that was its state to begin with.

  4. If the original VSAM data set was empty and `replace="YES"` is specified, then the contents of the existing VSAM data set will be lost.

preserveid="YES | NO"

indicates whether the UID and GID of the extracted file or directory should:

- Remain as defined when the archive was created, or
- Inherit the UID and GID of the userid performing the extract operation (GIMUNZIP).

A value of YES indicates that the UID and GID from when the archive was created will be preserved. A value of NO is equivalent to not specifying the `preserveid` attribute and indicates that the UID and GID should be inherited from the userid performing the extract operation (GIMUNZIP).

Note:

1. Preserving the UID and GID defined when the archive was created may cause GIMUNZIP to fail if the installer does not have the proper permissions.

2. The `preserveid` attribute applies only when extracting a UNIX file or directory from an archive. When extracting a data set, `preserveid` is parsed for syntax but otherwise ignored.

Syntax notes

1. GIMUNZIP ignores columns 73 through 80. If data is specified beyond column 72, GIMUNZIP ignores it, which may lead to the diagnosis of an error in a following tag.

2. Package control tags may contain comments. Comments start with `<!--` (hex 4C5A6060) and end with `-->` (hex 60606E). The first `-->` encountered after the initial `<!--` will end the comment. A comment may appear between a start-tag and its matching end-tag, but never within a tag.

Examples of using GIMUNZIP

In this first example, suppose a GIMZIP package contains a set of archive files and you want to extract the data sets and files from those archive files. In addition, you want to verify the integrity of the archive files and rename the destination data sets and files in the process. The following job stream can be used to perform such an operation.

```
//JOBx     JOB …
//STEP1    EXEC PGM=GIMUNZIP,PARM="HASH=YES"
//SMPDIR   DD PATH='/u/smpe/ORDER123/',PATHDISP=KEEP
//SMPCPATH DD PATH='/usr/lpp/smp/classes/',PATHDISP=KEEP
//SMPJHOME DD PATH='/usr/lpp/java/J6.0/',PATHDISP=KEEP
//SMPOUT   DD SYSOUT=*
//SYSPRINT DD SYSOUT=*
//SYSUT3   DD UNIT=SYSALLDA,SPACE=(CYL,(50,10))
//SYSUT4   DD UNIT=SYSALLDA,SPACE=(CYL,(25,5))
//SYSIN    DD *
<GIMUNZIP>

  <ARCHDEF archid="SMPMCS"
           replace="YES"
           preserveid="YES"
           newname="/userid/IBM/DO000029/SMPMCS">
  </ARCHDEF>

  <ARCHDEF archid="FMID001.F1"
           newname="userid.IBM.FMID001.F1">
  </ARCHDEF>

  <ARCHDEF archid="FMID002.F2"
           newname="userid.IBM.FMID002.F2">
  </ARCHDEF>

  <ARCHDEF archid="GLOBAL"
           prefix="userid.IBM">
  </ARCHDEF>

  <ARCHDEF name="S0006.2003030335435443234.pax.Z"
           replace="YES"
           newname="/userid/IBM/DO000029/RootHFS/">
  </ARCHDEF>

</GIMUNZIP>
/*
```

Figure 86. GIMUNZIP example 1

Note: Blank lines and spaces have been added to the package control tags for clarity, but are not required.

This second example of the GIMUNZIP operation uses the new storclas and catalog attributes:

```
//jobname JOB ...
//UNZIP EXEC PGM=GIMUNZIP,PARM="HASH=NO"
//SMPDIR DD PATH='/u/zosmf/swiExport/',PATHDISP=KEEP
//SMPCPATH DD PATH='/usr/lpp/smp/classes/',PATHDISP=KEEP
//SMPJHOME DD PATH='/usr/lpp/java/J6.0/',PATHDISP=KEEP
//SMPOUT DD SYSOUT=*
//SYSPRINT DD SYSOUT=*
//SYSUT3 DD UNIT=SYSALLDA,SPACE=(CYL,(50,10))
//SYSUT4 DD UNIT=SYSALLDA,SPACE=(CYL,(25,5))
//SYSIN DD *
<GIMUNZIP>
  <TEMPDS unit="SYSALLDA"> </TEMPDS>
  <ARCHDEF archid="LV1022.SYS1.MIGLIB"
           newname="SYS1.MIGLIB"
           volume="STG833"
           catalog="NO" >
  </ARCHDEF>
  <ARCHDEF archid="ML5443.SYS1.AOS12"
           newname="AOS.AOS12"
           storclas="STANDARD"
           catalog="YES" >
  </ARCHDEF>
</GIMUNZIP>
/*
```

Figure 87. GIMUNZIP example 2

In this example, the <TEMPDS> tag indicates the device unit type SYSALLDA will be used when allocating temporary work data sets. If the target MIGLIB data set does not already exist on volume STG833, then it will be allocated on that volume and will not be cataloged. In addition, if the target AOS12 data set does not already exist, it will be allocated using storage class STANDARD and will be cataloged.

GIMUNZIP processing

The GIMUNZIP service routine extracts data sets and files from the archive files that compose the GIMZIP package. An archive file consists of a portable image of a data set or file and the information needed to reload the data from the portable image (see "Archive files" on page 434 for details on archive file contents and construction).

GIMUNZIP uses the UNIX System Services pax command to expand the following component files from an archive file temporarily into a UNIX file system:

1. the portable image of the original data set or file, and

2. the file attribute file that contains the information necessary to reload the data from the archived data set or file.

GIMUNZIP then reads the file attribute file and uses the information recorded there, along with information specified on the <ARCHDEF> tag, such as `volume`, `unit`, `data class`, `management class`, `storage class`, `catalog`, `newname`, or `prefix`, to allocate a new or existing data set or file. The portable image of the original data set or file is stored in the data set or file just allocated for that purpose.

When GIMUNZIP encounters a package that contains archive segments, GIMUNZIP reassembles the archive segments into the original archive form, optionally verifies the integrity of the archive, and recreates the original data.

If the HASH=YES option was specified, then GIMUNZIP performs SHA-1 hash checking for the archive files. Specifically, GIMUNZIP reads the package attribute file for the package. The package attribute file, GIMPAF.XML, is found in the package directory, and contains the known hash values for the archives. These hash values were recorded by GIMZIP when the package and archives were created. GIMUNZIP then uses either ICSF services or Java 2 programs to compute the current hash values for the archive files, and compares this value to the known hash values for the archive files.

If the computed hash value matches the known hash value, then the integrity of the archive file is ensured. However, if the computed hash value does not match the known hash value, GIMUNZIP processing stops. This condition indicates that the archive file was corrupted after it was produced by GIMZIP when the package was created.

If the VERIFYSIG parameter was specified, and a signaturekeyring was provided in the SMPCLNT data set, then GIMUNZIP verifies the package signature using the following process:

1. The certificate chain is extracted from the GIMPAF2.XML file. This includes the signing certificate, intermediate CA certificates (if any), and the root CA certificate.

2. The signer's certificate is validated with its certificate path against a well-known and trusted Certificate Authority root certificate in the truststore identified by the signaturekeyring attribute specified in the SMPCLNT data set.

3. If the signer's certificate is trusted, GIMUNZIP then uses the public key contained in the certificate and the signature value from the <?PKGSIG> statement to verify the signature for the content of the GIMPAF2.XML file between and including the <PKGDEF> and </PKGDEF> tags.

Determining the required size of SMPWKDIR

SMPWKDIR is a directory in the UNIX file system used by GIMUNZIP for storing temporary work files. If this working directory is not identified by an SMPWKDIR DD statement, GIMUNZIP will use the package directory specified on the SMPDIR DD statement for storing temporary work files. In either case, the required size for the working directory varies depending on the type of information contained within an archive, and whether or not the archive is segmented. If an archive file is segmented, the file name of each segment will have a file type that identifies the segment position and the number of segments that make up the archive. For example, if an archive file is split into three segments, the segment files will

have file types of 1of3, 2of3 and 3of3 in the package directory. Additionally, in the Package Attribute File, each segment will be identified by an <ARCHSEG> tag and matching </ARCHSEG> end tag. The required working directory size can be calculated as the maximum of the following:

- the archive size for the largest segmented UNIX file system data set archive. That is, the sum of the size of all the segments. All the segments are copied from SMPDIR into the working directory to create a complete archive file. Using the z/OS UNIX pax command, the archive is expanded into its end point location.
- twice the original data set size for the largest segmented PDSE, VSAM or sequential data set archive. All the segments are copied from SMPDIR into the working directory to create a complete archive file. Using the z/OS UNIX pax command, the archive is expanded into a file in the working directory. The end point data set is then loaded from the contents of the expanded file.
- the original data set size for the largest unsegmented PDSE, VSAM or sequential data set archive. Using the z/OS UNIX pax command, the archive in SMPDIR is expanded into a file in the working directory. The end point data set is then loaded from the contents of the expanded file.

Sample RECEIVE job for GIMUNZIP output

After GIMUNZIP has been run, the SMP/E RECEIVE command can be used to process the data sets and files extracted by GIMUNZIP. Here is a sample RECEIVE job that would process the SMPMCS and RELFILE data sets and files extracted by GIMUNZIP in the sample in Figure 2:

```
//JOBx      JOB …
//STEP1     EXEC PGM=GIMSMP
//SMPCSI    DD DSN=SMPE.GLOBAL.CSI,DISP=SHR
//SMPPTFIN DD DSN=USERID.FMID001.SMPMCS,DISP=SHR
//SMPCNTL   DD *
  SET BDY(GLOBAL).
  RECEIVE SYSMODS
          RFPREFIX(USERID).
/*
```

Figure 88. Sample RECEIVE job for GIMUNZIP

Return codes

GIMUNZIP may end with the following return codes:

Return code
    Meaning

00
    The input data was processed successfully.

04
    A call to a system service may not have completed successfully.

12
    One of the following:

- Required file attributes could not be obtained.
- Input files were not sequential, partitioned, or VSAM data sets, or files and directories in the UNIX file system
- Call to a required system service failed.
- Syntax errors in the SYSIN data set.
- Data sets could not be opened.
- SMPDIR was not allocated to a UNIX directory.
- Data sets, files, or directories were missing.
- Required modules could not be loaded.
- Existing data set, file, or directory is not compatible.

- Existing data set, file, or directory could not be replaced.

16

One of the following:

- An I/O error occurred.

- A syntax error was found on the EXEC statement parameters.

20

SMPOUT data set is missing.

> 20

Internal error. Report the error to the IBM Support Center.

GIMXSID software inventory data service routine

GIMXSID is an SMP/E service routine to produce software inventory data. It is used as part of the IBM ShopzSeries offering, and can be used more generally to produce software inventory data. GIMXSID creates a single data source required by ShopzSeries to place customized software product and service orders. The data source created by GIMXSID, the software inventory data, is a composite of three kinds of information as follows:

Feature list

a list of FEATUREs found in the SMPCSI data sets. The Feature List is used by ShopzSeries to perform product requisite checking and also to prime the order checklist when ordering a ServerPac.

FMID list

a list of the FMIDs found in the SMPCSI data sets. The FMID List is used by ShopzSeries to scope service orders to the PTFs applicable solely to the user's desired configuration of target and global zones.

PTF list

a list of the PTFs found in the specified target zones and global zones. The PTF list is used by ShopzSeries to produce service packages that do not contain PTFs that are already present in the user's configuration.

GIMXSID is an independent load module residing in the MIGLIB library. It has an addressing mode of 31 (AMODE=31), a residence mode of 24 (RMODE=24), and is not authorized (AC=0). GIMXSID runs independently from the rest of SMP/E processing.

Calling GIMXSID

The JCL statements needed to call GIMXSID are:

```
//job       JOB ...
//step      EXEC PGM=GIMXSID,PARM='options'
//SMPOUT    DD SYSOUT=*
//SMPXTOUT  DD DSN=output.dataset,DISP=(NEW,CATLG),
//             SPACE=(TRK,(1,1)),UNIT=SYSALLDA
//SYSIN     DD *
input control statements
/*
```

Figure 89. JCL to call GIMXSID

EXEC

is the statement used to call GIMXSID. The EXEC statement must specify PGM=GIMXSID. The following options may be specified on the EXEC statement PARM operand:

INVENTORY=IBM | ALL

GIMXSID

Specifies which form of the software inventory will be produced. A value of IBM indicates the software inventory will be produced in the form expected by IBM's ShopzSeries. This form of the software inventory identifies all installed FMIDs and only PTFs whose IDs match the naming convention used by IBM. This is the default value.

A value of ALL indicates the software inventory will be produced in a generic form that identifies all installed FEATUREs, FMIDs and PTFs regardless of their naming convention.

LANGUAGE=ENU | JPN

The LANGUAGE option defines which language to use for GIMXSID messages. ENU is for US English, and JPN is for Japanese. LANGUAGE can also be specified as L. If LANGUAGE is not specified, the default is LANGUAGE=ENU.

WAIT=*minutes*MIN

where *minutes* is a decimal number between 0 and 9999. The WAIT option is used to control how long GIMXSID should wait if an SMPCSI data set is unavailable because it is in use by another SMP/E task. WAIT=*minutes*MIN causes GIMXSID to wait for a data set for the specified number of minutes. If a data set is still not available after the specified number of minutes, then GIMXSID stops. If WAIT is not specified, the default is WAIT=60MIN.

WAIT may be specified as WAIT=*minutes*MIN or WAIT=*minutes*. The MIN keyword is optional.

SMPOUT

used for GIMXSID messages. The RECFM must be FBA and the LRECL must be 81 or 121. GIMXSID formats messages to 80 characters in length, regardless of the LRECL for SMPOUT.

SMPXTOUT

the output specification for the software inventory data collected by GIMXSID. The output can be a sequential data set, a member of a partitioned data set, or a file in a UNIX file system. The required attributes for the output data set or file are dependent upon the inventory format being produced.

If INVENTORY=IBM is specified on the EXEC statement (or defaulted), then a data set specified on SMPXTOUT must have a RECFM of FB and the LRECL must be 12560. If a file in a UNIX file system is specified, then FILEDATA=BINARY and PATHOPTS(OWRONLY) must be specified on the DD statement.

If INVENTORY=ALL is specified on the EXEC statement, then a data set specified on SMPXTOUT must have a RECFM of VB and the LRECL must be 255 or greater. If a file in a UNIX file system is specified, then FILEDATA=TEXT and PATHOPTS(OWRONLY) must be specified on the DD statement.

SYSIN

specifies an input data set that contains the GIMXSID control statements. The input data set may be sequential, or a member of a partitioned data set. The RECFM must be F or FB and the LRECL must be 80.

GIMXSID control statements

Following are the control statements used to direct the processing for GIMXSID. These control statements are specified in the SYSIN data set.

CSI=*dsname*
indicates the data set name for an SMPCSI data set containing a global zone. This is a required control statement. There can be multiple CSI statements to identify multiple SMPCSI data sets and global zones to be processed.

dsname can be up to 44 characters long and must conform to standard data set naming conventions

TARGET=*zone, zone, ...*

 indicates the target zones controlled by the global zone identified on the preceding CSI statement that are to be processed. There can be multiple instances of the TARGET statement to allow many target zones to be specified. The *zone* values can be from 1 to 7 characters and must be separated by a comma and zero or more blank characters, or, one or more blank characters.

 This is an optional control statement. If a TARGET statement is not specified following a CSI control statement, all target zones controlled by the global zone identified on the preceding CSI statement are processed.

Syntax notes

1. Control statements must each start on a new line.

2. Control statements can begin in any column of a line.

3. Control statements must end on the same line on which they begin (control statements cannot span records).

4. SMP/E ignores columns 73 through 80 of lines in the SYSIN data set.

Example of using GIMXSID

```
//job      JOB ...
//BITMAP   EXEC PGM=GIMXSID
//SMPOUT   DD SYSOUT=*
//SMPXTOUT DD DSN=userid.GIMXSID.OUTPUT,DISP=(NEW,CATLG),
//           SPACE=(TRK,(1,1)),UNIT=SYSALLDA
//SYSIN    DD *
CSI=SMPE.ZOS.GLOBAL.CSI
TARGET=ZOS14, JES314
CSI=SMPE.DB2.GLOBAL.CSI
/*
```

Figure 90. GIMXSID example

In this example, the ZOS14 and JES314 target zones are processed from the global zone in the SMPE.ZOS.GLOBAL.CSI data set. In addition, all target zones from the global zone in the SMPE.DB2.GLOBAL.CSI data set are processed.

Processing

There are several phases of GIMXSID processing.

Initialization

GIMXSID first verifies that the SMPOUT data set, which is used for messages, is available. If SMPOUT is not allocated or cannot be opened, there is nowhere to write a message. Therefore, GIMXSID terminates with a return code of 20.

GIMXSID then checks for the required data sets, opens these data sets, and ensures they have the correct attributes.

Read and verify control statements

After the required data sets have been verified, GIMXSID reads the input control statements from the SYSIN data set. The CSI control statement is a required control statement and must precede any TARGET control statements. For each data set specified on a CSI control statement, GIMXSID ensures that the data set exists.

Verifying the target zones

For each SMPCSI data set specified on a CSI control statement, a list of target zones may be specified on one or more TARGET control statements. If target zones are specified, GIMXSID ensures that the specified zones exist and are in fact target zones.

Accessing resources

When reading SMPCSI data sets, it is possible one or more data sets are unavailable for use by GIMXSID because the data sets are being used by another SMP/E task. If so, GIMXSID waits for the data sets to become available based on the value of the WAIT option (see "Calling GIMXSID" on page 403).

If a global zone data set or a target zone data set is not available, then GIMXSID will wait until the data set is available, or until the number of minutes specified on the WAIT option have elapsed (or 60 minutes by default). During the WAIT time, GIMXSID attempts to read the data set every 30 seconds, just in case it becomes available. If the specified number of minutes elapses and one or more of the data sets are still not available, processing stops.

If WAIT=0MIN was specified, GIMXSID does not wait for SMPCSI data sets to become available. The initial attempt is made to access the data sets, and if they are not available, GIMXSID does not wait or make additional access attempts.

The number of minutes specified on the WAIT option (or defaulted) is the total time GIMXSID waits for all data sets. For example, if WAIT=60MIN was specified and the first SMPCSI data set causes a wait of 40 minutes, GIMXSID waits only a maximum of 20 additional minutes for all subsequent SMPCSI data sets.

Building the GIMXSID information

To identify itself as the creator of the software inventory data, GIMXSID creates GIMXSID information and writes it to the output data set. This GIMXSID information is written first in the output data set, but will be written only if there is software inventory data to produce (that is, if there are other records to be written). This GIMXSID information provides the name and service level of the program that produced the software inventory data, as well as the date and time it was produced. Table 58 on page 406 shows the format of the GIMXSID record written when INVENTORY=IBM is specified (or defaulted) on the EXEC statement.

Table 58. GIMXSID record format		
Field	**Length**	**Description**
Header	60	The value is GIMXSID padded to the right with blanks, to indicate the record type, and to identify the name of the program used to create the software inventory data.
Version	2	Two-character representation of the version level of GIMXSID.
Release	2	Two-character representation of the release level of GIMXSID.
Modification	2	Two-character representation of the modification level of GIMXSID.
PTF	3	Three-character representation of the PTF level of GIMXSID..
Date	9	The date the software inventory data was created, padded to the right with a blank. The date value is of the form *mm/dd/yy*.
Time	8	The time the software inventory data was created. The time value is of the form *hh:mm:ss*.
Reserved	12474	Reserved field, padded with blanks.

When INVENTORY=ALL is specified on the EXEC statement, the GIMXSID information is represented in the output data set in XML format as shown in the following example:

```
<gimxsidVersion>vv.rr.mm.ppp</gimxsidVersion>
<date>yyyy-mm-ddThh:mm:ss</date>
```

Building the feature list

The Feature List segment of the software inventory data contains the combined Feature information from all specified SMPCSI data sets. For each SMPCSI data set specified on a CSI control statement:

- If target zones were specified on one or more TARGET control statements for an SMPCSI data set, then the Feature List for that SMPCSI contains all FEATURE subentries found in FUNCTION SYSMOD entries from all specified target zones.
- If no target zones were specified for an SMPCSI data set, then the Feature List for that SMPCSI contains the name of all FEATURE entries found in the global zone, in addition to all FEATURE subentries found in FUNCTION SYSMOD entries from **all** target zones in the SMPCSI data set.

No regard is paid to the ERROR status of the FUNCTION SYSMOD entries; if a FUNCTION SYSMOD entry happens to be in ERROR, then it is assumed an overt attempt was made to install the FUNCTION and therefore its presence is desired. However, if a FUNCTION SYSMOD has been superseded or deleted, then it is ignored.

The individual Feature Lists from each specified SMPCSI are then merged to create the total Feature List. That is, the total Feature List set is the union of the individual sets.

When INVENTORY=IBM is specified (or defaulted) on the EXEC statement, the total Feature List is represented in the output data set using the definition for FEATURE records shown in . Multiple FEATURE records may be produced in order to contain the total FEATURE information.

Table 59. FEATURE record format		
Field	**Length**	**Description**
Header	8	Value is FEATURE to indicate the record type, padded to the right with blanks.
Feature Name	52	Feature name, padded to the right with blanks.
Product-id	9	The product identifier of the feature's associated product, padded to the right with blanks. For IBM products, this field is an IBM program product number (5647-A01, for example). This field may be all blank if the FEATURE entry is not found in a global zone in an SMPCSI data set.
VRM	9	The version, release, and modification level of the feature's associated product, padded to the right with blanks. This field is of the form *vv.rr.mm* where *vv* is the version, *rr* is the release, and *mm* is the modification level. This field may be all blank if the FEATURE entry is not found in a global zone in an SMPCSI data set.
reserved	5	This field is reserved for future use.
Description	65	A text description for the feature, padded to the right with blanks. This field may be all blank if the FEATURE entry is not found in a global zone in an SMPCSI data set.
FMID List	12412	List of function SYSMODs that make up this feature. This field is a list of 7-byte FMID names separated by one blank space, padded to the right with blanks. This field may be all blank if the FEATURE entry is not found in a global zone in an SMPCSI data set, or the FMID subentry is not found for the FEATURE entry.

A single FEATURE may be represented in the output data set with more than one FEATURE record. Although an unlikely case, it is necessary only to accommodate a FEATURE that contains more FMIDs than can fit in the FMID List field of a single FEATURE record.

When INVENTORY=ALL is specified on the EXEC statement, the total Feature List is represented in the output data set using the <featureList> XML tag as shown in the following code example. Each FEATURE is defined by a separate <feature> tag.

```
<featureList>
  <feature name="feature-name" prodid="product-id"
  vrm="vv.rr.mm">
  <description>feature-description</description>
  <fmidList>
    <id>fmid</id>
    <id>fmid</id>
    ...
  </fmidList>
  </feature>
</featureList>
```

The definition of the values in the <feature> tag are the same as for the fields in the Feature record described in Table 59 on page 407.

Building the FMID list

The FMID List segment of the software inventory data contains the combined FMID information from all specified SMPCSI data sets. For each SMPCSI data set specified on a CSI control statement:

- If target zones were specified on one or more TARGET control statements for an SMPCSI data set, then the FMID List for that SMPCSI contains the names of all FUNCTION SYSMODs found in all specified target zones.
- If no target zones were specified for an SMPCSI data set, then the FMID List for that SMPCSI contains the names of all FUNCTION SYSMODs found in the global zone, in addition to the names of all FUNCTION SYSMODs found in **all** target zones for the SMPCSI data set.

No regard is paid to the ERROR status of the FUNCTION SYSMOD entries; if a FUNCTION SYSMOD entry happens to be in ERROR, then it is assumed an overt attempt was made to install the FUNCTION and therefore its presence should be noted. However, if a FUNCTION SYSMOD in a target zone has been superseded or deleted, then it is ignored.

The individual FMID Lists from each specified SMPCSI are then merged to create the total FMID List. That is, the total FMID List set is the union of the individual sets.

When INVENTORY=IBM is specified (or defaulted) on the EXEC statement, the total FMID List is represented in the output data set using the definition for PRODLIST records shown in Table 60 on page 408. Multiple PRODLIST records may be produced in order to contain the complete list of FMIDs.

Table 60. PRODLIST record format		
Field	**Length**	**Description**
Header	60	Value is "PRODLIST" and padded to the right with blanks.
FMID List	12500	List of 7-byte FMID names separated by one blank space, padded to the right with blanks.

When INVENTORY=ALL is specified on the EXEC statement, the total FMID List is represented in the output data set using the <fmidList> XML tag as shown in the following code example. Each FMID is listed in a separate <id> tag.

```
<fmidList>
  <id>fmid</id>
  <id>fmid</id>
  ...
</fmidList>
```

Building the PTF list

The PTF list segment of the software inventory data contains a composite of PTF information from all specified SMPCSI data sets. The point of the PTF list is to identify all PTFs that are already on-site and therefore do not need to be included in a service order. That is, there is no need to deliver PTFs that are already applied in all target zones in which they could be applied, or are already in the global zones and could be applied to any target zone as needed.

For each SMPCSI data set specified on a CSI control statement, the individual zones to be processed when creating the PTF information are determined as follows:

- If target zones were specified on one or more TARGET control statements for an SMPCSI data set, then those target zones are the current set of zones to be processed.
- If no target zones were specified for an SMPCSI data set, then all target zones in the global zone, and the global zone itself, are the current set of zones to be processed.

For this discussion, a PTF is considered applicable to a target zone if the FMID of the PTF is applied in that target zone. A PTF is considered applicable to a global zone if the FMID of the PTF is received in the global zone. A PTF is considered applicable to an SMPCSI data set if the PTF is applicable to any of the zones in the current set, as determined previously.

The PTF information for an individual SMPCSI data set contains all PTFs that are received in the global zone, or are applied in all target zones of the current set in which they are applicable. Such PTFs are determined to be "present" in an SMPCSI data set. That is, such PTFs can be considered already present and therefore no need exists to deliver such PTFs on behalf of this SMPCSI, because they are either already in the global zone and can be applied to any target zone, or they are already applied in all target zones in which they could be applied.

A PTF is considered applied in a target zone whether or not it is superseded. A PTF is not considered applied in a target zone, or received in a global zone, if it's SYSMOD entry is in ERROR status.

When INVENTORY=IBM is specified (or defaulted) on the EXEC statement, a PTF's presence is noted in the PTF list only if its name follows the IBM naming convention for PTFs. Specifically, if the last five characters of the PTF name are numeric, then it is noted in the PTF list. When INVENTORY=ALL is specified on the EXEC statement, the PTF naming convention does not matter, and all PTFs will be noted in the PTF list.

Further, it is not always possible to know the SYSMOD TYPE for a SYSMOD that has been superseded in a target zone. Therefore, when INVENTORY=IBM is specified (or defaulted) on the EXEC statement, if the name of such a SYSMOD follows the IBM naming convention for PTFs, then it is assumed to be a PTF and its presence is noted. Specifically, if the first character of the name of such a SYSMOD is a character in the set "LMNOPQRSTU" and the last five characters of the name are numeric, then it is assumed to be a PTF and its presence will be noted. When INVENTORY =ALL is specified on the EXEC statement, the naming convention of SYSMODs does not matter, and all such SYSMODs will be assumed to be a PTF and their presence will be noted.

The PTF information from each individual SMPCSI data set is combined to create the total PTF list. A PTF is included in the PTF list only if it is present in all SMPCSI data sets in which it is applicable. Stated another way, if a PTF is applicable to *n* SMPCSI data sets, but it is present in only *n-1* or fewer SMPCSI data sets, then that PTF cannot be included in the PTF list.

When INVENTORY=IBM is specified (or defaulted) on the EXEC statement, the total PTF list is represented in the output data set using the definition for BITPTF records. Multiple BITPTF records can be produced in order to contain the total PTF bitmap.

Table 61. BITPTF record format		
Field	**Length**	**Description**
Header	6	Value is BITPTF to indicate the record type. BITPTF indicates that the record contains a PTF identification bit string.

Table 61. BITPTF record format (continued)		
Field	**Length**	**Description**
PTF prefix	54	A two-character PTF prefix value, padded to the right with blanks. The PTF prefix value is the first two alphabetic characters in a PTF name. For example, UR or UW.
Bit string	12500	Bit-string where each bit represents a number from 00000 to 99999. The numbers from 00000 to 99999 correspond to the last five numeric characters in a PTF name. The offset of a particular bit corresponds to a particular number from 00000 to 99999. This number, combined with the PTF prefix, represent a specific PTF name. For example, if the PTF prefix is "UW" and the bit at offset 12345 is on, then the PTF with the name "UW12345" is represented.

When INVENTORY=ALL is specified on the EXEC statement, the total PTF list is represented in the output data set using the <ptfList> XML tag as shown in the following code example. Each PTF is listed in on a separate <id> tag.

```
<ptfList>
   <id>ptf</id>
   <id>ptf</id>
   ...
</ptfList>
```

Sample software inventory when INVENTORY=ALL

When INVENTORY=ALL is specified on the EXEC statement, the software inventory is represented in XML format. The following shows a sample XML software inventory.

```
<?xml version="1.0" encoding="UTF-8"?>
<?xml-stylesheet type="text/xsl" href="gimxsid.swinventory.xsl"?>
<gimxsidSwInventory>
<gimxsidVersion>03.06.00.56</gimxsidVersion>
<date>2014-09-03T11:44:22</date>
<featureList>
<feature name="IECBA510" prodid="5655-W32" vrm="05.01.00">
<description>Ent COBOL z/OS V5</description>
<fmidList>
<id>HADB510</id>
<id>JADB51H</id>
<id>JADB511</id>
</fmidList>
</feature>
<feature name="IDBTBC10" prodid="5655-W70" vrm="12.01.00">
<description>Debug Tool V12</description>
<fmidList>
<id>HADRC10</id>
<id>HVWR160</id>
<id>H09F210</id>
</fmidList>
</feature>
</featureList>
<fmidList>
<id>HADB510</id>
<id>JADB51H</id>
<id>JADB511</id>
<id>HADRC10</id>
<id>HVWR160</id>
<id>H09F210</id>
</fmidList>
<ptfList>
<id>UI14448</id>
<id>UI16133</id>
<id>UI16134</id>
<id>UI18034</id>
<id>UI18055</id>
<id>UI18382</id>
```

```
<id>UI19849</id>
<id>UI19850</id>
<id>UK96988</id>
<id>UK97247</id>
<id>UK98481</id>
<id>UK98482</id>
<id>UK98499</id>
<id>AI05656</id>
<id>AI05657</id>
<id>AI05658</id>
<id>AI06128</id>
<id>AI06899</id>
<id>AI08238</id>
<id>AM99261</id>
</ptfList>
</gimxsidSwInventory>
```

Zone and data set sharing considerations

GIMXSID uses the system enqueue facility to control access to SMPCSI data sets. The following identifies the phases of GIMXSID processing and the zones and data sets GIMXSID may require for shared use during those phases.

- For the initialization phase:

Global zones (as specified on CSI control statements)	Read with shared enqueue.

- For the processing phase:

Global zones (as specified on CSI control statements)	Read with shared enqueue.
Target zones (as specified on TARGET control statements)	Read with shared enqueue.

- Termination

  All resources are freed.

Return codes

GIMXSID may complete with any of the following return codes.

Return code
  Meaning

00
  GIMXSID processing completed successfully.

12
  One of the following:

  - Required data sets were missing or could not be opened.
  - Required modules could not be loaded.
  - Syntax error was found in the SYSIN control statements.
  - SMPCSI data sets were not available.
  - Specified zones were not defined or are not target zones.

16
  One of the following:

  - An I/O error occurred.
  - Syntax error was found on the EXEC statement parameters.

20
  The SMPOUT data set is missing.

GIMXTRX service routine

GIMXTRX is intended for use as part of an offering called ShopzSeries. It provides two basic functions:

List target zone names (LSTTZN)
Generate a list of target zone names associated with a given GLOBAL zone SMPCSI data set name.

Create a BITMAP of an SMPCSI (BMPTZN)
Generate a BITMAP representation of FUNCTION and PTF SYSMODs found in a given list of target zone names associated with a given GLOBAL zone SMPCSI data set name.

The input to GIMXTRX is the name of a data set that contains the input parameters. See "Input parameter data set contents" on page 412.

The results of each function is written to a data set created by GIMXTRX. See "Processing" on page 414.

Note: GIMXTRX is a separate load module residing in the MIGLIB library and runs independently from the rest of SMP/E, and is executable on a z/OS, OS/390, or MVS system; it will not run on VM.

Calling GIMXTRX

GIMXTRX has normal standard linkage and expects an input parameter that is a character string that is a data set name. Following are the JCL statements needed to call GIMXTRX:

```
//JOBx     JOB ...
//XTRX EXEC PGM=GIMXTRX,PARM='parmdsn'
//SYSTSPRT DD SYSOUT=(*)
//SMPXTOUT DD DSN=outputdsn,DISP=(NEW,CATLG),
//           UNIT=unit,SPACE=(TRK,(2,1))
```

Figure 91. JCL to call GIMXTRX

EXEC

is the statement used to call GIMXTRX. The EXEC statement must specify PGM=GIMXTRX. The following option must be specified on the EXEC statement PARM operand:

parmdsn
name of the data set that contains the input parameters required for the GIMXTRX function. See "Input parameter data set contents" on page 412 for a description of the data contained in the input parameter data set.

SYSTSPRT

is used by GIMXTRX for messages. The record format (RECFM) of SYSTSPRT must be FBA, and the record length (LRECL) must be 121. GIMXTRX will format messages to 120 characters in length.

Note: GIMXTRX messages are only issued in the U.S. English language.

SMPXTOUT

is the output data set. This DD statement is optional. If SMPXTOUT is specified, GIMXTRX writes the output to the identified data set. If SMPXTOUT is not specified, GIMXTRX allocates a new data set for the output. SMPXTOUT must be a sequential data set, the record format (RECFM) must be FB, and the record length (LRECL) must be 12560. If incorrect attributes are specified for SMPXTOUT, GIMXTRX will change the attributes to the required values.

Input parameter data set contents

The input parameter data set is a fixed-block 80-byte record sequential data set. Each input parameter consists of a keyword and an associated value. The rules for the input parameter data set are as follows:

- There must be one keyword with its associated parameter data per 80-byte record.
- All keywords must begin in column one of the records.
- No blanks are allowed after the keyword and before the equal sign
- No blanks are allowed after the equal sign and before the first parameter value
- Any data read from the input parameter data set that is not a recognized keyword parameter is ignored and processing continues.

The keywords and parameters passed to the GIMXTRX program are as follows:

GIMXTRX input parameters
SUFFNC keyword

▶▶─ SUFFNC= ──┬── LSTTZN ──┬─ ▶◀
              └── BMPTZN ──┘

SUFGBL keyword

▶▶─ SUFGBL= ── *globaldsn* ─ ▶◀

SUFUNIT keyword

▶▶─ SUFUNIT= ── *unit* ─ ▶◀

SUFVOL keyword

▶▶─ SUFVOL= ── *volume* ─ ▶◀

SUFTGT keyword

```
              ◄──────────┐
▶▶─ SUFTGT= ──┬─ tzone name ─┴─ ▶◀
```

SUFFNC
This is the name of the function to be processed by the GIMXTRX program. The two allowable values are

LSTTZN
Generate a list of target zone names.

BMPTZN
Create a BITMAP of a set of target zones.

This is a required parameter.

There can only be one function executed per invocation of GIMXTRX.

SUFGBL
This is the data set name containing the GLOBAL zone to be processed. This parameter is 44 characters long. This is a required parameter.

SUFUNIT
This is the UNIT value used to allocate a new output data set. The value is up to 8 characters long. This parameter is optional. If no UNIT value is specified, GIMXTRX uses a default UNIT value of SYSALLDA when allocating a new output data set.

SUFVOL
This is the work volume used to allocate the output data set. The value is six characters long, and must be a valid MVS volume name. This is an optional parameter. If no volume is specified, GIMXTRX will attempt to allocate the output data set without a volume serial.

SUFTGT
This is the list of target zone names that will be queried. Each target zone name is 7 characters long separated by one blank space. If a target zone name is less than 7 characters long it must be left-justified and padded with blanks. Each SUFTGT record can contain up to nine target zones. This is a required parameter when SUFFNC=BMPTZN. It is not required when SUFFNC=LSTTZN.

There can be multiple instances of the SUFTGT record to allow for situations where a single 80-byte record cannot contain the required input data.

Examples of using GIMXTRX

In the following examples, the input data set is created in the job stream. In the first example, the output of GIMXTRX is written to a new data set created by GIMXTRX.

```
//jobname  JOB ...
//*
//* Create the input data set.
//*
//step1     EXEC  PGM=IEBGENER
//SYSPRINT DD SYSOUT=*
//SYSIN     DD DUMMY
//SYSUT2    DD DSN=userid.dataset.INPUT,DISP=(NEW,CATLG),
//            DCB=(RECFM=FB,LRECL=80),
//            UNIT=unit,SPACE=(TRK,(1,0))
//SYSUT1    DD DATA,DLM=$$
SUFFNC=BMPTZN
SUFGBL=userid.GLOBAL.CSI
SUFTGT=ZOSTGT
$$
//step2     EXEC PGM=GIMXTRX,PARM='userid.dataset.INPUT'
//*
//* The output Bit-Map is written to a new data set
//* allocated with the name "userid.dataset.OUT"
//*
//SYSTSPRT DD SYSOUT=*
//SYSPRINT DD SYSOUT=*
```

In the second example, the output of GIMXTRX is written to the data set specified on the SMPXTOUT DD statement.

```
/jobname  JOB ...
//step1     EXEC  PGM=IEBGENER
//*
//* Create the input data set.
//*
//SYSPRINT DD SYSOUT=*
//SYSIN     DD DUMMY
//SYSUT2    DD DSN=userid.dataset.INPUT,DISP=(NEW,CATLG),
//            DCB=(RECFM=FB,LRECL=80),
//            UNIT=unit,SPACE=(TRK,(1,0))
//SYSUT1    DD DATA,DLM=$$
SUFFNC=BMPTZN
SUFGBL=userid.GLOBAL.CSI
SUFTGT=ZOSTGT
$$
//step2     EXEC PGM=GIMXTRX,PARM='userid.dataset.INPUT'
//*
//* The output Bit-Map is written to the data set specified
//* on the SMPXTOUT DD statement.
//*
//SMPXTOUT DD DSN=userid.dataset.OUT,DISP=(NEW,CATLG),
//            UNIT=unit,SPACE=(TRK,(2,1))
//SYSTSPRT DD SYSOUT=*
//SYSPRINT DD SYSOUT=*
```

Processing

GIMXTRX obtains the input parameter data set name from the PARM value passed to it on the EXEC statement.

If the SMPXTOUT DD statement is specified, GIMXTRX uses the data set specified for the output. If the SMPXTOUT DD statement is not specified, GIMXTRX allocates a new data set for the output. GIMXTRX constructs the name of the output data set using the input parameter data set name as a base. The last qualifier of the input parameter data set name is changed to 'OUT' and this becomes the name of the output data set. For example, if the input parameter data set name is TEST.GIMXTRX.DATA.IN, the output data set name will be TEST.GIMXTRX.DATA.OUT.

The output data set is allocated and cataloged as follows:

```
DSNAME=dsn.OUT
DSORG=PS
RECFM=FB
```

```
LRECL=12560
BLKSIZE=0
SPACE=(TRK,(10,1),RLSE)
UNIT=unit (supplied as SUFUNIT= on input) or default of SYSALLDA
VOLUME=volume name (supplied as SUFVOL= on input).
```

If the SMPXTOUT DD statement is not specified, GIMXTRX attempts to create the output data set, which therefore must not exist at the time that GIMXTRX is initiated. If there is no volume supplied as input, GIMXTRX will attempt the allocation without a volume serial. If there is no unit supplied as input, GIMXTRX will attempt the allocation with UNIT=SYSALLDA.

GIMXTRX then allocates and opens all of the required data sets and extracts the input parameters from the input parameter data set.

Once the input parameters have been read and the output data set created, GIMXTRX continues processing based on the function requested (LSTTZN or BMPTZN).

List target zone names (LSTTZN)

This function of GIMXTRX produces a list of target zone names associated with a given GLOBAL zone SMPCSI data set name. The results are written to the output data set.

Inputs

The inputs to the List Target Zone Names function are:

SUFFNC
Name of the function to be executed (LSTTZN)

SUFGBL
GLOBAL zone SMPCSI data set name

SUFUNIT
UNIT to be used for output data set (optional)

SUFVOL
VOLUME to be used for output data set (optional)

A sample of the parameters contained in the input parameter data set for the List Target Zone Names function is:

```
          1         2         3         4         5         6         7         8
----+----0----+----0----+----0----+----0----+----0----+----0----+----0----+----0
SUFFNC=LSTTZN
SUFGBL=TEST.GLOBAL.VSAM.CSI
SUFVOL=VOL111
```

Figure 92. Sample input parameter data set for function LSTTZN

Using the SMP/E GIMAPI, GIMXTRX obtains the names of all of the target zones associated with the specified global zone. GIMXTRX then writes the list of target zone names to the output data set in 12560-byte records defined as follows:

Field
  Description

header
60-byte header character string starting with the literal "TARGETZN" in uppercase. The remaining 52 bytes are set to blanks.

Target zone name list
7-character target zone names each separated by one blank space, up to 12500 bytes in length. For target zone names that are less than 7 characters long, the field is left-justified and padded to the right with blanks. When the list is less than 12500 bytes long, the remaining bytes contain blanks.

Example of a target zone name record

The following example indicates an SMPCSI that has target zones MVST1, JESTGT1, CICTGA, DB2TG2.

```
            1         2         3         4         5         6         7         8         9
----+----0----+----0----+----0----+----0----+----0----+----0----+----0----+----0----+----0

TARGETZN                                              MVST1    JESTGT1 CICTGA   DB2TG2
```

Figure 93. TARGETZN record

Create a BITMAP of an SMPCSI (BMPTZN)

This function of GIMXTRX produces a BITMAP representation of the PTF service present in a user's SMPCSI. This information is intended for shipment to IBM where it can be analyzed to create a custom-built service package to be delivered to the user over the internet.

The intent of the Create a BITMAP of an SMPCSI function is to identify and record in a set of BITMAP records:

- All FMIDs and PTFs applied in the target zones identified in the input data set (SUFTGT parameter).
- All PTFs received in the GLOBAL zone identified in the input data set (SUFGBL parameter) that are associated with the FMIDs applied in the specified target zones.

A PTF is considered applicable to a target zone if its FMID is applied in that target zone. If a PTF is applied in all selected target zones to which it is applicable, it is recorded in a BITMAP record. If a PTF is not applied in all of the selected target zones to which it is applicable, then it is not recorded in a BITMAP record. For example, if an FMID is applied in 4 of the selected target zones, but a PTF associated with that FMID is applied in only 3 of those zones, the PTF is not recorded in the BITMAP record.

All PTFs that are received in the GLOBAL zone are checked against the FMIDs applied in the selected target zones to determine if the PTFs are associated with the applied FMIDs. All PTFs received in the GLOBAL zone that are applicable to the FMIDs applied in the selected target zones are recorded in a BITMAP record.

Inputs

The inputs to the Create a BITMAP of an SMPCSI function are:

SUFFNC
Name of the function to be executed (BMPTZN)

SUFGBL
GLOBAL zone SMPCSI data set name

SUFUNIT
UNIT to be used for output data set (optional)

SUFVOL
VOLUME to be used for output data set (optional)

SUFTGT
Target zone names selected by the user

A sample of the input parameters contained in the input parameter data set for the create a BITMAP of an SMPCSI function is:

```
         1         2         3         4         5         6         7         8
----+----0----+----0----+----0----+----0----+----0----+----0----+----0----+----0

    SUFFNC=BMPTZN
    SUFGBL=TEST.GLOBAL.VSAM.CSI
    SUFVOL=VOL111
    SUFTGT=TARGET1 TGT2    TARGET3 TGT4    TARGET5 TARGET6 TGT7
```

Figure 94. Sample input parameter data set for function BMPTZN

The output records created by GIMXTRX for the BMPTZN function represent the functions applied in the target zones selected by the user (PRODLIST record) and the PTFs that are either applied or received (BITMAP records). GIMXTRX then writes the PRODLIST and BITMAP records to the output data set in 12560-byte records defined as follows:

PRODLIST record

The PRODLIST record contains the list of FMIDs applied in the target zones identified by the user in the input data set. Following is the definition of the PRODLIST record that is placed in the output data set by GIMXTRX.

Field
> **Description**

header
> 60-byte header starting with the literal "PRODLIST" in uppercase. The remaining 52 bytes are set to blank.

FMID list
> list of 7-character FMIDs each separated by one blank space. This field is 12500 bytes long. When the FMID list is less than 12500 bytes long, the remaining bytes will be made blanks.

The following example of a PRODLIST record indicates an SMPCSI that has FMIDs HBB6602, HMP1B00, JMP1B01 present.

```
         1         2         3         4         5         6         7         8
----+----0----+----0----+----0----+----0----+----0----+----0----+----0----+----0----+

PRODLIST                                              HBB6602 HMP1B00 JMP1B01
```

Figure 95. PRODLIST record

The PRODLIST record will be placed before any BITMAP bitstring records in the output data set.

BITMAP records

The BITMAP records represent the PTFs applied in the target zones or received in the GLOBAL zone that are associated with the FMIDs listed in the PRODLIST record.

The BITMAP bitstring logic is based on the following PTF naming convention:

- PTF names are architected by IBM to consist of a 2 character prefix followed by 5 decimal digits providing 100000 possible PTF numbers per prefix (00000-99999).
- A BITMAP bitstring record simply contains 100000 bits and a 2-character PTF prefix in the record header.
- For example, if the prefix is "UR", then the first bit in the bitstring represents PTF UR00000, the second bit represents PTF UR00001, the third represents PTF UR00002 and so on. The last bit in the bitstring represents UR99999.
- The bits representing the UR*nnnnn* PTF numbers received or applied on the customer system are turned on while the remaining bits are set to zero.
- One bitstring is required for each PTF prefix in the SMPCSI

Following is the definition of the BITMAP records that are placed in the output data set by GIMXTRX.

Field
> **Description**

header
> 60-byte header starting with the literal "BITPTF" in uppercase, followed by the 2-character PTF prefix for each unique PTF prefix present in the zones being queried. This field is in uppercase. The remaining 52 bytes in the header are set to blank.

bitstring
> 12500 bytes with 100000 bits representing PTF numbers present in the SMPCSI associated with the two-character PTF prefix in this record. A bit is set on to indicate a PTF number exists. All remaining bit positions are set to zero.

The following example of a BITMAP bitstring record indicates an SMPCSI that has PTFs UR00001 and UR00005 present.

```
          1         2         3         4         5         6         7         8
----+----0----+----0----+----0----+----0----+----0----+----0----+----0----+----0

 BITPTFUR                                              bitstring...........
```
Figure 96. BITMAP record

The bitstring section of the BITMAP record is expanded as follows (where each mark of the scale line represents one bit position of the BITMAP record):

```
6       6       6       6       6       6       6       6       6       7
1-------2-------3-------4-------5-------6-------7-------8-------9-------0

01000100000000000000000000000000000000000000000000.........................
```
Figure 97. Bitstring section of BITMAP record

Return codes

To help you diagnose errors, GIMXTRX issues messages and return codes. The messages are documented in *z/OS SMP/E Messages, Codes, and Diagnosis*.

All nonzero return codes indicate unsuccessful completion of GIMXTRX. Here is a description of the return codes:

Return code
> **Meaning**

0
> The requested GIMXTRX function was processes successfully.

4
> A warning condition was encountered in a program called by GIMXTRX.

8
> An error condition was encountered in a program called by GIMXTRX.

12
> One of the following
> - A required parameter is missing or incorrect.
> - An allocation error occurred.
> - An I/O error occurred.
> - An error occurred while opening or closing a data set.

- GIMXTRX could not load either GIMAPI, GIMSGENU, or GIMKXENU.
- SMP/E could not obtain enough storage for processing.

If an error occurs with SYSTSPRT, there will be no message issued. Only a return code will be returned through standard linkage.

GIMXTRX does not provide an ESTAE routine for recovery when an ABEND occurs. Normal system ABEND processing occurs.

GIMZIP packaging service routine

The GIMZIP service routine creates portable packages of software and associated materials. Typically the packages will contain SYSMODs, RELFILE data sets, HOLDDATA, and associated materials such as documentation, samples, and text files. These GIMZIP packages may be transported through a network, processed by the GIMUNZIP service routine, and then processed by the SMP/E RECEIVE command.

More specifically, a GIMZIP package consists of one or more package definition files, a set of archive files, and text files. The package definition files describe the total package and identifies the archive files and text files contained in the package. An archive file consists of :

1. a portable image of any of the following:

   - a sequential data set
   - a partitioned data set
   - a VSAM data set
   - a file in the UNIX file system
   - a directory in the UNIX file system

2. and the information necessary to reload the data from the portable image.

A single GIMZIP package typically consists of several archive files.

Note:

1. GIMZIP is a separate load module residing in the MIGLIB library and runs independently from the rest of SMP/E processing.

2. GIMZIP requires either the Integrated Cryptographic Services Facility (ICSF) One-Way Hash Generate callable service or Java 8, or a logical successor, to be available for its use in order to compute an SHA-1 hash value or to digitally sign the GIMZIP package.

Calling GIMZIP

The following figure shows the job control statements for GIMZIP:

```
//JOBx      JOB …
//STEP1     EXEC PGM=GIMZIP,PARM='option'
//SMPDIR    DD PATH='package_directory',PATHDISP=KEEP
//SMPWKDIR  DD PATH='work_directory',PATHDISP=KEEP
//SMPCPATH  DD PATH='smpclasses_directory',PATHDISP=KEEP
//SMPJHOME  DD PATH='javaruntime_directory',PATHDISP=KEEP
//SMPOUT    DD SYSOUT=*
//SYSPRINT  DD SYSOUT=*
//SYSUT2    DD UNIT=SYSALLDA,SPACE=(CYL,(200,20))
//SYSUT3    DD UNIT=SYSALLDA,SPACE=(CYL,(50,10))
//SYSUT4    DD UNIT=SYSALLDA,SPACE=(CYL,(25,5))
//SMPCLNT   DD *
  client tag and attributes
/*
//SYSIN     DD *
  package control tags
/*
```

Figure 98. JCL to call GIMZIP

EXEC

is the statement used to start GIMZIP. The EXEC statement must specify PGM=GIMZIP. The following options may be specified on the EXEC statement PARM operand:

LANGUAGE=*xxx*

where *xxx* can be one of the following:

ENU

US English

JPN

Japanese

The LANGUAGE option defines which language to use for GIMZIP messages. LANGUAGE can also be specified as L. If LANGUAGE is not specified, the default is LANGUAGE=ENU.

SEGMENT=*nnnn*M

where *nnnn* is a decimal number between 1 and 9999. The SEGMENT option indicates the archive files produced by GIMZIP are to be divided into archive segments no larger than nnnn megabytes. If SEGMENT is not specified, the default is to perform no archive segmenting.

SYSIN

specifies a control data set that contains the package control tags to direct GIMZIP processing (see). The control data set may be a sequential data set or a single member of a partitioned data set and its attributes must be LRECL=80 and RECFM=F or FB.

SMPDIR

specifies a directory in a UNIX file system that is to contain the package created by GIMZIP. All files and archives created by GIMZIP are stored in this directory. This directory is referred to as the package directory.

The package directory must be empty at the start of GIMZIP processing. If any files, including a previous package, already exist in the package directory, you must move them to another directory, delete them, or specify a different directory on the SMPDIR DD statement for the package directory.

SMPWKDIR

specifies a directory in a UNIX file system that is used by GIMZIP for temporary work files. This is an optional DD statement. If the SMPWKDIR DD statement is not provided, GIMZIP will use the package directory specified on the SMPDIR DD statement for temporary work files.

SMPCPATH

specifies a directory in the UNIX file system where the SMP/E Java classes reside. The SMPCPATH DD statement is optional and only required to calculate SHA-1 hash values if ICSF is not available, or to digitally sign the GIMZIP package.

Alternatively, the Java class path may be specified by the classpath attribute of the <CLIENT> tag in the SMPCLNT DD statement.

SMPJHOME

specifies a directory in the UNIX file system where the Java runtime resides. The SMPJHOME DD statement is optional and only required to calculate SHA-1 hash values if ICSF is not available, or to digitally sign the GIMZIP package.

Alternatively, the Java runtime location may be specified by the javahome attribute of the <CLIENT> tag in the SMPCLNT DD statement.

SMPOUT

specifies an output data set that will contain SMP/E messages issued during GIMZIP processing. The data set attributes are LRECL=81 or LRECL=121 and RECFM=FBA.

SMPCLNT

provides information about the local client host environment. This includes the Java home directory, Java class path, and Java debug options. The local client host information is specified with the <CLIENT> tag. For information about using the <CLIENT> tag, see *z/OS SMP/E Commands*.

SMPCLN

SMPCLNT is an optional DD statement. SMPCLNT may be sequential, a member of a partitioned data set, or a file in the z/OS UNIX file system. See for information about allocating SMPCLNT.

SYSUT2

specifies a work area for the IEBCOPY utility to use for the sequential unload data sets placed on DASD as the result of an IEBCOPY unload operation. This data set must be large enough to accommodate the largest partitioned data set being processed by the current execution of GIMZIP.

SYSUT3 and SYSUT4

specifies work areas for the IEBCOPY utility to use for spill data sets on DASD.

SYSPRINT

specifies a data set to contain output and messages from the utilities called by GIMZIP.

GIMZIP package control tags

GIMZIP service routine processing is controlled by package control tags. The package control tags define the GIMZIP package to be created and are specified in the SYSIN data set. The package control tags follow XML syntax rules and their format is as follows:

GIMZIP package control tags

GIMZIP

```
►►─ <GIMZIP ─┬──────────────────────────────────────┬─►
             └─ description=" package description " ─┘

►─┬──────────────────────────────────────────────────────────────┬─►
  └─ signingcertificate=" certificate label " ─ keyring=" ─┬─ keyring_name ─┬─ " ─┘
                                                           └─ userid/keyring name ─┘

►─┬─ > ─┬◄─────────────┬─┬─ </GIMZIP> ─►◄
        └─── FILEDEF ───┘
```

FILEDEF
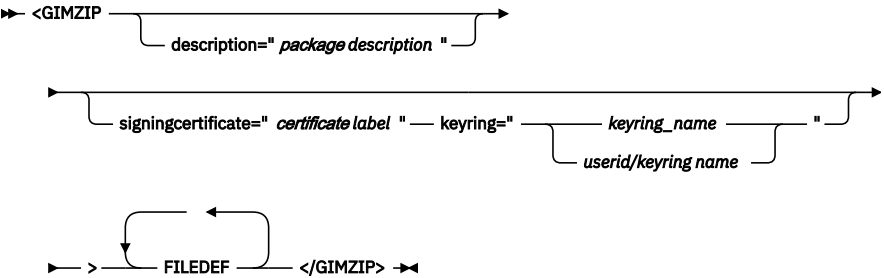
```
►►─ <FILEDEF ─ name=" file name " ─┬──────────────────────┬─►
                                   └─ archid=" archive id " ─┘

►─┬──────────────────────────────┬─┬─────────────────────────────┬─►
  └─ subdir=" subdirectory name " ─┘ └─ description=" file description " ─┘

►─┬───────────────────────────────────┬─┬─────────────────────────┬─ > ─►
  └─ type=" ─┬─ README ──┬─ " ─┘        └─ volume=" file volume " ─┘
             ├─ SMPHOLD ──┤
             ├─ SMPPTFIN ─┤
             └─ SMPRELF ──┘

►─ </FILEDEF> ─►◄

►►─►◄
```
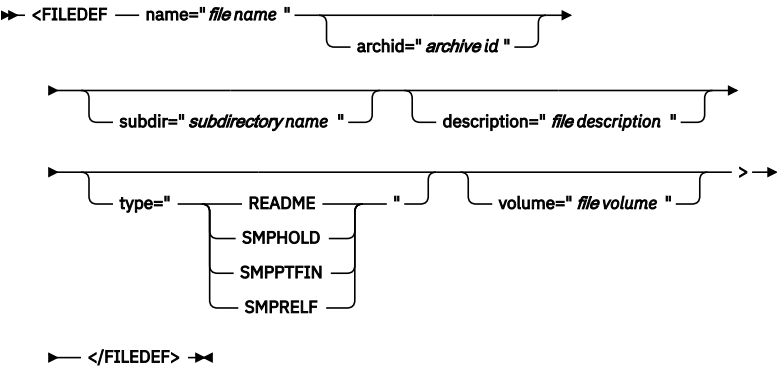
<GIMZIP> Tag syntax

The <GIMZIP> and corresponding </GIMZIP> tags identify the beginning and end of a GIMZIP package build request. The following attributes may be found on the <GIMZIP> tag:

> description="*package description*"
>
> > specifies a text description for the package. The description can be up to 500 characters long. This attirbute is optional.

signingcertificate

> identifies the label for the certificate that will be used to sign the GIMZIP package. This certificate is bound to the private key that is used to sign the GIMZIP package, and contains the matching public key used to verify the package signature. This attribute is optional, but if specified, the GIMZIP package will be digitally signed.
>
> Certificate labels are from 1 to 32 characters in length and may include only characters X'40' through X'FE' excluding the reserved XML characters, less than (<), greater than (>), double quotation mark ("), and ampersand (&).

keyring

> specifies the security manager key ring that contains the specified signing certificate. The keyring must also contain all certificates in the certification path from the signing certificate to the issuing root CA certificate.
>
> Key ring names are from 1 to 237 characters in length and may include only characters X'40' through X'FE' excluding the forward slash (/) and the reserved XML characters, less than (<), greater than (>), double quotation mark ("), and ampersand (&).
>
> If the key ring is associated with a user ID other than that used to execute program GIMZIP, then the key ring name must be prefixed with the associated user ID. User IDs are from 1 to 8 alphanumeric characters in length, can consist entirely of numbers and need not begin with any particular character. The user ID is separated from the key ring value by a forward slash (userid/keyring).

`<FILEDEF>` Tag syntax

The `<FILEDEF>` and corresponding `</FILEDEF>` tags identify the beginning and end of a file definition. Each file described by a `<FILEDEF>` tag is input to GIMZIP and is to be a component of the GIMZIP package. The following attributes may be found on the `<FILEDEF>` tag:

name="*file name*"

> specifies the name of a data set or file that is to be a component of the GIMZIP package. The name attribute can specify either a sequential, partitioned, or VSAM data set, or the absolute pathname for a file or directory in the UNIX file system.
>
> The first character of the absolute pathname of a file or directory in the UNIX file system must be a slash (/). If the name does not begin with a slash, GIMZIP processing will assume that it is an MVS data set name and that it will conform to the MVS data set naming conventions. When a file or directory in the UNIX file system is specified, the name can be from 1 to 1023 bytes long with 255 characters between delimiters (/). The value can contain any character from X'40' through X'FE', except '<', '>', '&', and the double quotation mark ("). All data beyond column 72 is ignored, including blanks. The pathname of a file or directory in the UNIX file system is case sensitive and will not be converted to uppercase alphabetic during GIMZIP processing. Data set names are not case sensitive and will be converted to uppercase alphabetic during GIMZIP processing.
>
> Note the following facts:
>
> 1. When a VSAM data set is being specified, the true cluster name must be used. Do not reference a VSAM data set by a path name. Although an alternate index may be defined to the cluster, the alternate index does not become part of the archive. If an alternate index is desired at the destination site after the archive is unzipped, then the alternate index must be defined and built at the destination site.
>
> 2. GIMZIP processing uses the z/OS UNIX System Services pax command for creating archive files. Although GIMZIP will attempt to process UNIX file and directory names of up to 1023 bytes in

length, this processing is subject to the limitations of the pax command. See *z/OS UNIX System Services Command Reference* for information about the pax command.

archid="*archive id*"

specifies a unique archive id value to be associated with the archive file created for the data set, file, or directory identified by the associated name attribute. The value of this attribute will most likely be something like a ddname or the low level qualifier of an MVS library name or some generic specification for a file or directory in the UNIX file system.

The value can be from 1 to 243 characters in length. The value can contain any character from X'40' through X'FE', except '<', '>', '&', '/' and the double quotation mark ("). The archid value is case sensitive and will not be converted to uppercase alphabetic by GIMZIP processing.

The same archid value may not be used in more than one file definition group in the same package.

description="*file description*"

specifies a text description for the file. The description can be up to 500 characters long.

subdir="*subdirectory name*"

specifies the subdirectory in which to store the file specified by the corresponding name attribute. The subdirectory name can be from 1 to 500 bytes long with 255 characters between delimiters (/).

This attribute can only be specified if the file type attribute is either not specified or is README.

The subdirectory name is a relative pathname for the desired subdirectory, so the first character of the subdirectory name cannot be a slash (/). However, the subdirectory name may be specified with or without an ending slash.

"SMP", all capitalized, cannot be specified as the first three characters of the subdir attribute value. Other case variations are allowed.

type="*file type*"

specifies a file type indicator that describes the contents of the file, identifies how the file is to be processed by GIMZIP, and also how the file will be processed by the SMP/E RECEIVE command. Any of the following values may be specified:

README
indicates that the associated file is a sequential text data set or a file in the UNIX file system that is intended to be viewable text after it is placed in the package. These files are not placed in archives by GIMZIP and are not compressed. They are only copied into the package by GIMZIP. The file name specified on the associated name attribute must identify a file in the UNIX file system or a sequential data set that does not have a record format of VS (variable spanned).

SMPHOLD
indicates that the associated data set or file contains SMP/E HOLDDATA statements. These files are placed in archives by GIMZIP and compressed. The name specified on the associated name attribute must identify a sequential data set with a logical record length of eighty (80) bytes and a record format of FB (fixed block), or it must identify a file in the UNIX file system.

SMPPTFIN
indicates that the associated data set or file contains SMP/E modification control statements (MCS). These files are placed in archives by GIMZIP and compressed. The name specified on the associated name attribute must identify a sequential data set with a logical record length of eighty (80) bytes and a record format of FB (fixed block), or it must identify a file in the UNIX file system.

SMPRELF
indicates the data set specified on the name attribute is an SMP/E RELFILE data set associated with a SYSMOD contained in the SMPPTFIN file. Data sets with the SMPRELF file type must be

either a partitioned data set or an IEBCOPY unloaded sequential image of a partitioned data set. These data sets are placed in archives by GIMZIP and compressed.

The ʈype attribute must not be specified when a VSAM data set or a UNIX file system directory is specified on the name attribute. If the ʈype attribute is not specified, then the data set specified on the name attribute can be either a sequential, partitioned, or VSAM data set, or a file or directory in the UNIX file system.

volume="*file volume***"**

specifies the volume serial number of the volume containing the data set specified on the name attribute. The volume identifier must be from 1 to 6 alphanumeric characters and should be specified only if the name attribute identifies an uncataloged data set.

If the associated name attribute identifies a file or directory in the UNIX file system (the name begins with a slash), the volume attribute will be checked for syntax, but it will be ignored otherwise.

Note: The volume is used to allocate data sets. However, VSAM data sets must be catalogued. Information needed for the archival of VSAM data sets is gathered from the catalog and related areas, such as the VVDS - VSAM volume data set.

Syntax notes

1. GIMZIP ignores columns 73 through 80. If data is specified beyond column 72, GIMZIP ignores it, which may lead to the diagnosis of an error in a following tag.
2. Package control tags may contain comments. Comments start with <!-- (hex 4C5A6060) and end with --> (hex 60606E). The first --> encountered after the initial <!-- will end the comment. A comment may appear between a start-tag and its matching end-tag, but never within a tag.

Example of using GIMZIP

Suppose a GIMZIP package is to be created containing the following data sets, files, and directories:

- /SAMPLE/ORDER123/readme.html - A README file.
- /SAMPLE/ORDER123/SMPMCS - an SMP/E MCS file
- SAMPLE.IBM.FMID001.F1 - An SMP/E PDS relative file
- SAMPLE.IBM.FMID001.F2 - An SMP/E PDSE relative file
- SAMPLE.ORDER123.DOCLIB - A PDS containing documents
- SAMPLE.ORDER123.RIMLIB - A PDS containing related installation materials
- SAMPLE.ORDER123.MVS.GLOBAL.CSI - A VSAM cluster
- /SAMPLE/ORDER123/RootHFS/ - the root directory

The following job stream can be used to create such a GIMZIP package:

```
//JOBx     JOB …
//STEP1    EXEC PGM=GIMZIP,PARM='SEGMENT=12M'
//SMPDIR   DD PATH='/u/smpe/GIMZIP/ORDER123/',PATHDISP=KEEP
//SMPCPATH DD PATH='/usr/lpp/smp/classes/',PATHDISP=KEEP
//SMPJHOME DD PATH='/usr/lpp/java/J6.0/',PATHDISP=KEEP
//SMPOUT   DD SYSOUT=*
//SYSPRINT DD SYSOUT=*
//SYSUT2   DD UNIT=SYSALLDA,SPACE=(CYL,(200,20))
//SYSUT3   DD UNIT=SYSALLDA,SPACE=(CYL,(50,10))
//SYSUT4   DD UNIT=SYSALLDA,SPACE=(CYL,(25,5))
//SYSIN    DD *
<GIMZIP description="This is a sample software package.">

   <!-- This sample package contains the files shown later. -->
   <FILEDEF name="/SAMPLE/ORDER123/readme.html"
            archid="readme.html"
            description="This is the README file for the package."
            subdir="README"
            type="README">
   </FILEDEF>

<FILEDEF name="/SAMPLE/ORDER123/SMPMCS"
            description="This is the SMPMCS file for ORDER123."
            archid="SMPMCS"
            type="SMPPTFIN">
   </FILEDEF>

   <FILEDEF name="SAMPLE.IBM.FMID001.F1"
            archid="FMID001.F1"
            description="This is SMP/E RELFILE 1 for FMID001."
            type="SMPRELF">
   </FILEDEF>

   <FILEDEF name="SAMPLE.IBM.FMID001.F2"
            archid="FMID001.F2"
            description="This is SMP/E RELFILE 2 for FMID001."
            type="SMPRELF">
   </FILEDEF>

   <FILEDEF name="SAMPLE.ORDER123.DOCLIB"
            subdir="DOCLIB"
            description="This is the Document Library for the package.">
   </FILEDEF>

   <FILEDEF name="SAMPLE.ORDER123.RIMLIB"
            subdir="RIMLIB"
            description="This is the Related Installation Materials
                         Library for the package.">
   </FILEDEF>

   <FILEDEF name="SAMPLE.ORDER123.MVS.GLOBAL.CSI"
            archid="GLOBAL"
            description="This is a sample VSAM cluster.">
   </FILEDEF>

   <FILEDEF name="/SAMPLE/ORDER123/RootHFS/"
            description="This is the entire root directory.">
   </FILEDEF>

   </GIMZIP>
/*
```

Figure 99. GIMZIP example

Note: Blank lines and spaces have been added to the package control tags for clarity, but are not required.

GIMZIP processing

The GIMZIP service routine processes the following data structures as input to create GIMZIP packages:

- sequential data sets
- partitioned (PDS and PDSE) data sets

- VSAM data sets
- UNIX files
- UNIX directory files

The packages produced are stored in the package directory in a UNIX file system. The package directory is specified by the SMPDIR DD statement.

GIMZIP collects various attributes for the input data sets and files specified in the package control tags. These attributes, along with portable images of the input data sets, are temporarily stored in a UNIX file system. These images are then archived and compressed to create archive files. Each GIMZIP package will contain one or more archive files.

For each package there are at least two additional files:

- The package attribute file, GIMPAF.XML, which describes the contents of the package itself. The package attribute file identifies the archives for a package and contains a hash value for each archive within the package. The hash value for each archive is used for data integrity purposes and can be checked when the original data set is recreated from the archive using the GIMUNZIP service routine.
- An extensible stylesheet language (XSL) file, which describes how to format the information found in the package attribute file. The XSL file is used for rendering (displaying on a browser) the package attribute file.
- The extended package attribute file, GIMPAF2.XML,is produced only if the GIMZIP package is signed. Like the GIMPAF.XML file, the GIMPAF2.XML file describes the contents of the package, but identifies more secure hash values for each archive as well as digital signature information for the package.

Preparing to sign GIMZIP packages

GIMZIP package signing is implemented using public/private key technology; a private key is used to generate a digital signature for package files, and the corresponding public key is used to verify the signatures. The key pair is bound to an X.509 certificate and GIMZIP uses this certificate for the signing operation. This signing certificate and public/private key pair must meet the following requirements:

1. The public/private key pair must be generated using the RSA algorithm and can be from 1024 to 4096 bits long.
2. The signing certificate must have the Digital Signature key usage certificate extension.
3. The signing certificate must not be expired.
4. The signing certificate should be issued by a well known and trusted certificate authority (CA), whose root certificate is easily obtained by the consumers that will verify the digital signature of the signed GIMZIP package.
5. The signing certificate, the issuing root CA certificate, and intermediate CA certificates (if any), must be stored in the z/OS security manager database and connected to a key ring.

Any appropriate tool or method may be used to generate a signing certificate, but to illustrate these requirements, the following example z/OS Security Server RACF command will generate an acceptable signing certificate for use by the GIMZIP service routine:

```
RACDCERT GENCERT ID(cert-owner) +
  SUBJECTSDN(CN('My Package Signing Cert') +
             O('My Company') +
             C('US')) +
  RSA +
  SIZE(2048) +
  KEYUSAGE(HANDSHAKE) +
  NOTAFTER( DATE(2033-04-01) ) +
  SIGNWITH(CERTAUTH LABEL('My Root CA')) +
  WITHLABEL('My Package Signing Cert')
```

where *cert-owner* is the user ID associated with the signing certificate.

If the signing certificate is generated elsewhere, it must be added to the security manager database. Once the signing certificate, root CA certificate, and intermediate CA certificates (if any) are defined in

the security manager database, they must be connected to a key ring. A key ring is a named collection of certificates associated with a specific user. A key ring may be created using a RACF command like this:

```
RACDCERT ID(keyring-owner) ADDRING(keyring-name)
```

where *keyring-owner* is the user ID that will own the key ring and *keyring-name* is a name that you chose for the key ring.

Connect the certificates to the key ring using a RACF command like this:

```
RACDCERT ID(keyring-owner) +
CONNECT(CERTAUTH LABEL('My Root CA') +
  RING(keyring-name) USAGE(CERTAUTH) )
RACDCERT ID(keyring-owner) +
  CONNECT(ID(cert-owner) LABEL('My Package Signing Cert') +
  RING(keyring-name) )
```

Once connected to a key ring, to access those certificates and the signing certificate's private key, the user ID under which GIMZIP runs must be properly authorized to the key ring. One of the following conditions must be true:

- If the user ID under which GIMZIP runs is the user ID associated with the signing certificate, then that user ID must have READ authority to either the IRR.DIGTCERT.LISTRING resource in the FACILITY class, or the <keyring-owner>.<keyring-name>.LST resource in the RDATALIB class. READ access enables retrieving one's own private key.

- If the user ID under which GIMZIP runs is not the user ID associated with the signing certificate, then that user ID must have UPDATE authority to the <keyring-owner>.<keyring-name>.LST resource in the RDATALIB class. UPDATE access enables retrieving others' private key.

The following example RACF commands can be used to authorize the user ID under which GIMZIP runs:

```
PERMIT IRR.DIGTCERT.LISTRING CLASS(FACILITY) ID(gimzip-userid) ACCESS(READ)
```

Or

```
RDEFINE RDATALIB keyring-owner.keyring-name.LST UACC(NONE)
PERMIT keyring-owner.keyring-name.LST +
  CLASS(RDATALIB) ID(gimzip-userid) ACCESS(UPDATE)
SETROPTS RACLIST(RDATALIB) CLASSACT(RDATALIB)
SETROPTS RACLIST(RDATALIB) REFRESH
```

where *gimzip-userid* is the user ID under which GIMZIP runs.

Note: If using z/OS Security Server RACF, to enable the RDATALIB authorization checking, the RDATALIB class must be RACLISTed.

The key ring and certificate label for the signing certificate are specified on the keyring and signingcertificate attributes on the <GIMZIP> tag in the input for the GIMZIP service routine, like this:

```
<GIMZIP
 signingcertificate="My Package Signing Cert"
 keyring="keyring-owner/keyring-name" >
<FILEDEF …/>
</GIMZIP>
```

If the user ID under which GIMZIP runs is not the user ID that owns the keyring, then specify both the keyring owner user ID and the keyring name on the keyring attribute.

Package attribute file

Each package contains a package attribute file, GIMPAF.XML. If the package is digitally signed it will also contain an extended package attribute file, GIMPAF2.XML. Both files are stored in the package directory identified by the SMPDIR DD statement. These files contain package definition control tags that describe the contents of the package and how the package was created. The package definition control tags follow XML syntax rules.

For the format of the package definition control tags for GIMPAF.XML, refer to "GIMPAF.XML" on page 428. For the format of the package definition control tags for GIMPAF2.XML, refer to "GIMPAF2.XML" on page 428.

GIMPAF.XML

The format of the package definition control tags for GIMPAF.XML is as follows:

Package definition control tags for GIMPAF.XML (Part 1 of 2)
PKGDEF
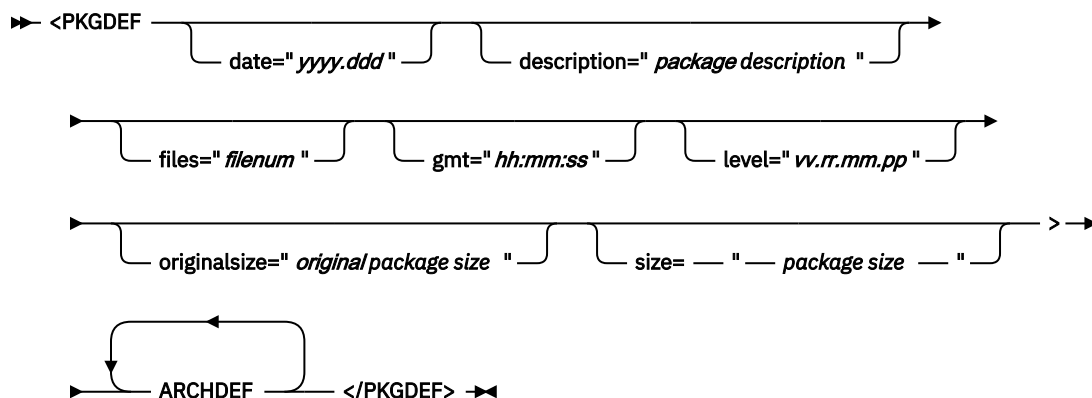
ARCHDEF
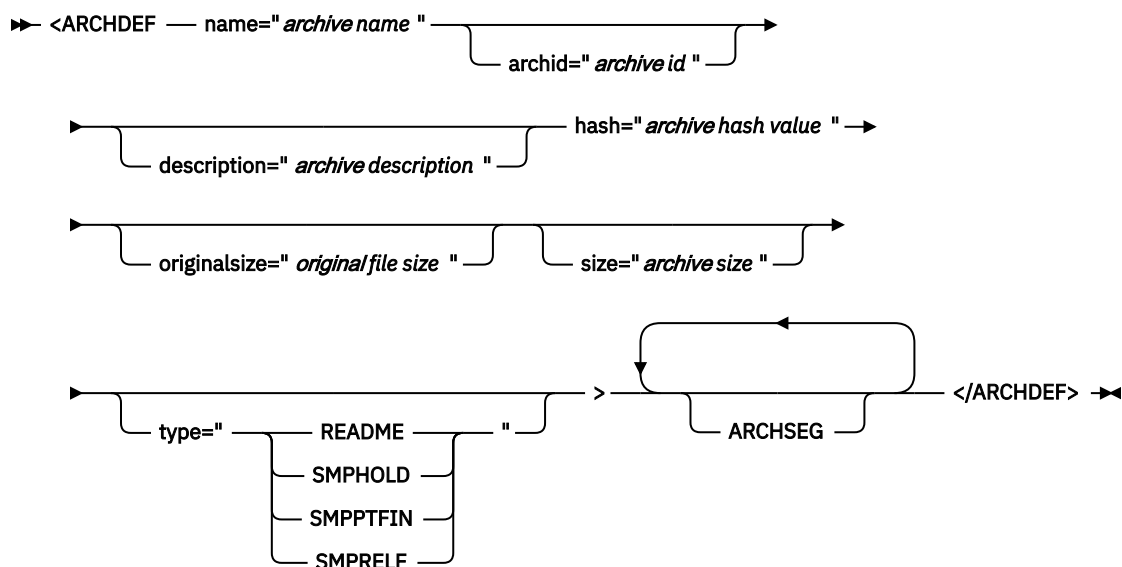
ARCHSEG

Package definition control tags for GIMPAF.XML (Part 2 of 2)
PKGHASH

GIMPAF2.XML

The format of the package definition control tags for GIMPAF2.XML is as follows:

Package definition control tags for GIMPAF2.XML (Part 1 of 2)
PKGDEF

►►— <PKGDEF ──►
　　　　　└─ date=" *yyyy.ddd* " ─┘ └─ description=" *package description* " ─┘

►──►
　└─ files=" *filenum* " ─┘ └─ gmt=" *hh:mm:ss* " ─┘ └─ level=" *vv.rr.mm.pp* " ─┘

►── > ──►
　└─ originalsize=" *original package size* " ─┘ └─ size= " *package size* " ─┘

►─┬─────────────┬─ SignatureInfo ─ </PKGDEF> ►◄
　└─ ARCHDEF ◄─┘

ARCHDEF

►►— <ARCHDEF — name=" *archive name* " ──────────────────────►
　　　　　　　└─ archid=" *archive id* " ─┘

►──►
　└─ description=" *archive description* " ─┘ └─ originalsize=" *original file size* " ─┘

►── > ─ hash ─►
　└─ size=" *archive size* " ─┘ └─ type=" ─┬─ README ──┬─ " ─┘
　　　　　　　　　　　　　　　　　　　　　　├─ SMPHOLD ──┤
　　　　　　　　　　　　　　　　　　　　　　├─ SMPPTFIN ─┤
　　　　　　　　　　　　　　　　　　　　　　└─ SMPRELF ──┘

►─┬─────────────┬─ </ARCHDEF> ►◄
　└─ ARCHSEG ◄─┘
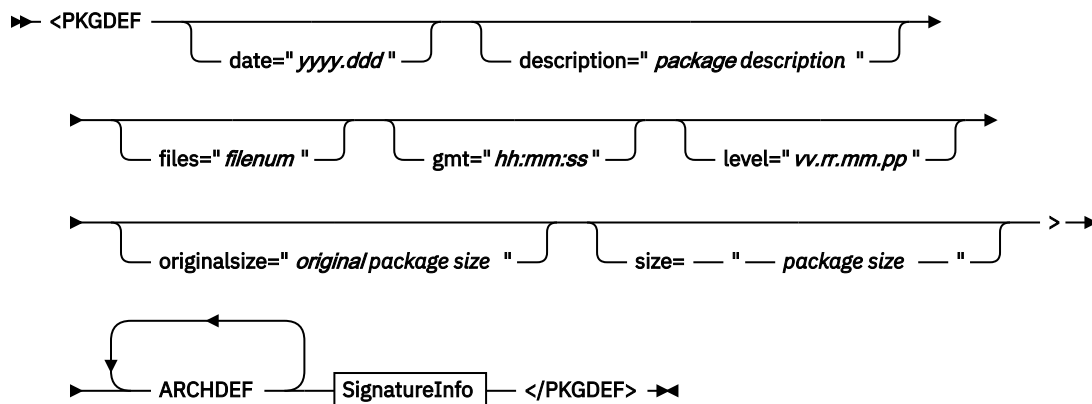
hash

►►— <hash — algorithm=" *hash algorithm* " — > — *archive file hash* — </hash> ►◄

ARCHSEG

►►— <ARCHSEG — name=" *segment name* " ──────────────── > ─ hash ─ </ARCHSEG> ►◄
　　　　　　　　└─ size=" *segment size* " ─┘

SignatureInfo

►►— <SignatureInfo> ─ SignatureAlgorithm ─ SignerSubjectName ─ CertPath ─ </SignatureInfo> ►◄

SignatureAlgorithm

►►— <SignatureAlgorithm> — *signature algorithm* — </SignatureAlgorithm> ►◄

SignerSubjectName

►►— <SignerSubjectName> — *certificate subject name* — </SignerSubjectName> ►◄

CertPath

►►— <CertPath> ─ X509Data ─ </CertPath> ►◄

X509Data

►► <X509Data> ── | X509SubjectName | | X509Certificate | ── </X509Data> ►◄

X509SubjectName

►► <X509SubjectName> ── *certificate subject name* ── </X509SubjectName> ►◄

X509Certificate

►► <X509Certificate> ── *certificate data* ── </X509Certificate> ►◄

Package definition control tags for GIMPAF2.XML (Part 2 of 2)
PKGSIG

►► <?PKGSIG ── signature=" *package signature* " ── ?> ►◄

<PKGDEF> Tag syntax

The <PKGDEF> and corresponding </PKGDEF> tags identify the beginning and end of a GIMZIP package definition. The following attributes may be found on the <PKGDEF> tag:

date="*yyyy.ddd*"
indicates the year and Julian day on which the package was created by the GIMZIP service routine.

description="*package description*"
specifies a text description for the package. The value GIMZIP places here, is the same as the description value, if specified, on the input <GIMZIP> tag.

files="*filenum*"
indicates the total number of files that compose the package. This number includes all archives, and any and all control files, such as the Package Attribute File.

gmt="*hh:mm:ss*"
indicates the time at which the package was created by the GIMZIP service routine. The time value is expressed as Greenwich Mean Time (GMT).

level="*vv.rr.mm.pp*"
indicates the service level of the GIMZIP service routine used to create this package. The individual values are as follows:

vv
version

rr
release

mm
modification level

pp
PTF level

Note: When the PTF level (pp) of the GIMZIP service routine exceeds 2 characters, the first character will be dropped. For example, when the PTF level of the GIMZIP service routine reaches 100, the PTF level will start back at '00'.

originalsize="*original package size*"
indicates the sum of all the values of *originalsize* attributes on ARCHDEF tags in this package definition. It is used to determine the amount of auxiliary storage required to retrieve all the archives in this package.

size="*package size*"
indicates the total size of the package, in bytes. This size is the sum of the sizes of all files that compose the package.

<ARCHDEF> Tag syntax

The <ARCHDEF> and corresponding </ARCHDEF> tags identify the beginning and end of a specific archive file definition. The following attributes may be found on the <ARCHDEF> tag:

name="*archive name*"
    specifies the relative path name for an archive file in the GIMZIP package. The path name value is relative to the package directory specified on the SMPDIR DD statement. See "Archive files" on page 434 for information about the proper naming convention for archive files.

archid="*archive id*"
    specifies a unique archive id associated with the archive specified in the archive name attribute. There can be only one `archid` attribute in an archive definition group. The value is taken from the `archid` value found on the file definition tag in the GIMZIP package control tags (if one was specified). This is an optional attribute.

description="*archive description*"
    specifies a text description for the archive. The value GIMZIP produces is copied from the description value, if specified, on the corresponding input <FILEDEF> tag.

hash="*archive hash*"
    indicates an SHA-1 hash value for the archive file.

originalsize="*original file size*"
    indicates the total size, in bytes, of the original data set, directory, or file stored in the archive file.

size="*archive size*"
    indicates the total size of the archive file, in bytes.

type="*file type*"
    specifies an archive type indicator that describes the contents of the archive and how the file will be processed by the SMP/E RECEIVE command. The value is copied from the file type value, if specified, on the corresponding <FILEDEF> statement, and may be any of the following values:

    README
        indicates the archive contains text data.

    SMPHOLD
        indicates the archive contains SMP/E HOLDDATA statements.

    SMPPTFIN
        indicates the archive contains SMP/E modification control statements (MCS).

    SMPRELF
        indicates the archive contains an SMP/E RELFILE data set associated with a SYSMOD contained in the SMPPTFIN file.

<hash>

The <hash> tag identifies a hash value for the subject file. The following attribute may be found on the <hash> tag:

algorithm="*hash algorithm*"
    Identifies the hashing algorithm used to calculate the value specified on the <hash> tag.

<ARCHSEG> Tag syntax

The <ARCHSEG> tag and matching </ARCHSEG> end tag identify a segment of an archive file. The following attributes may be found on the <ARCHSEG> tag:

name="*segment name*"
    specifies the path name of a segment file for an archive file in the GIMZIP package. The path name value is a relative value and it is relative to the package directory specified on the SMPDIR DD statement for GIMZIP. This attribute and corresponding attribute value is always produced when archives are to be segmented.

size="*segment size*"
indicates the total size of the segment file, in bytes. While GIMZIP always produces a byte count when archives are to be segmented, this attribute and attribute value is optional and is not used for GIMUNZIP or RECEIVE processing.

hash="*segment hash*"
indicates an SHA-1 hash value for the segment file. This attribute and corresponding attribute value is always produced when archives are to be segmented.

\<hash\>

The \<hash\> tag identifies a hash value for the subject file segment. The following attribute may be found on the\<hash\>tag:

algorithm="*hash algorithm*"
Identifies the hashing algorithm used to calculate the value specified on the \<hash\> tag.

\<?PKGHASH\> Tag syntax

hash="*package hash*"
indicates an SHA-1 hash value for the information in the package attribute file between the starting \<PKGDEF\> tag and the ending \</PKGDEF\> tag.

\<SignatureInfo\> Tag syntax

The \<SignatureInfo\> and corresponding \</SignatureInfo\> tags describes information required for verifying the GIMZIP package signatures. The following tags may be found in the \<SignatureInfo\> tag:

\<SignatureAlgorithm\>
Identifies the algorithm used to calculate the value specified on the signature attribute of the \<? PKGSIG\> statement.

\<SignerSubjectName\>
Identifies the subject name from the signer's certificate.

\<CertPath\>
The list of certificates in the certification path from the trusted root CA certificate to the end entity certificate, including any intermediate CA certificates. The end entity certificate, or signer's certificate, contains the public key used to verify the GIMZIP package signature.

\<X509Data\>
Information about a single certificate in the certificate path.

\<X509SubjectName\>
The subject name for the certificate.

\<X509Certificate\>
A base64 encoded x.509 certificate (PEM file format).

\<?PKGSIG\> Tag syntax

The \<?PKGSIG\> tag contains the signature for the content of the GIMPAF2.XML file between and including the \<PKGDEF\> and \</PKGDEF\> tags. The following attribute may be found on the \<?PKGSIG\> tag:

signature
Contains the signature value.

Syntax notes

1. Data in columns 73 through 80 is ignored. If data is specified beyond column 72, it is ignored and an error in a following tag may be indicated.

2. Package definition control tags may contain comments. Comments start with <!-- (hex 4C5A6060) and end with --> (hex 60606E). The first --> encountered after the initial <!-- will end the comment. A comment may appear between a start-tag and its matching end-tag, but never within a tag.

Package attribute file example

The package attribute file for the package produced by GIMZIP in would look like the following:

```
<?xml version="1.0" ?>
<?xml-stylesheet type="text/xsl" href="GIMPAF.XSL"?>
<PKGDEF files="8"
          originalsize="18755980"
          size="14568080"
          date="2003.090"
          gmt="22:10:20"
          level="03.03.00.00">

  <ARCHDEF name="SMPPTFIN/S0001.SMPMCS.pax.Z"
          archid="SMPMCS"
          type="SMPPTFIN"
          description="This is the SMPMCS file for FMID001"
          originalsize="18432000"
          size="14336000"
          hash="22CA8A741CA2EE1BC57FAA2D0EBA22895727B8E9">
    <ARCHSEG name="SMPPTFIN/2003030341435443234.1of2"
            size="12288000"
            hash="23EB39F41CA2EE1BC57FAA2D0EBA22895727B8E9">
    </ARCHSEG>
    <ARCHSEG name="SMPPTFIN/2003030341435443234.2of2"
            size="2048000"
            hash="58FA9A741CA2EE1BC57FAA2D0EBA22895727B8E9">
    </ARCHSEG>
  </ARCHDEF>

  <ARCHDEF name="README/S0002.packlist.html"
          type="README"
          description="This is the packing list for your order"
          originalsize="32620"
          size="18256"
          hash="CA2EE1BC57F2EE1BC57FAA2D0EBA2289C57FAA2D">
  </ARCHDEF>

  <ARCHDEF name="SMPRELF/CPACTST.FMID001.F1.pax.Z"
          archid="FMID001.F1"
          type="SMPRELF"
          description="This is an SMPRELF file for FMID001.F1"
          originalsize="55840"
          size="32256"
          hash="D0EBAA741CA2EE1BC57FAA2D0EBA2289EECE1BC5">
  </ARCHDEF>

  <ARCHDEF name="SMPRELF/CPACTST.FMID002.F2.pax.Z"
          archid="FMID002.F2"
          type="SMPRELF"
          description="This is an SMPRELF file for FMID002.F2"
          originalsize="55840"
          size="32256"
          hash="B7DBAA741CA2EE1BC57FAA2D0EBA2289EECE1BC5">
  </ARCHDEF>

  <ARCHDEF name="S0005.CPACTST.D0000029.MVS.GLOBAL.CSI.pax.Z"
          archid="GLOBAL"
          description="This is an example of a VSAM cluster
          for the product."
          originalsize="23218"
          size="17556"
          hash="D0DB27641CA2EE1BC5BFAA2D0EBA2EE9EECE1BC5">

  </ARCHDEF>
```

Figure 100. Package attribute file (GIMPAF.XML) example (Part 1)

```
  <ARCHDEF name="S0006.2003030335435443234.pax.Z"
          description="This is the entire root directory."
          originalsize="155840"
          size="132256"
          hash="D0EB49385CA2EE1BC57FAA2D0EBA2289EECE1BC5">
  </ARCHDEF>

  <ARCHDEF name="GIMPAF.XSL"
          description="This is an Extensible Stylesheet Language
          (XSL) document used to render the package attribute file
          (GIMPAF.XML) on an internet  browser.  This file has not
          been archived and should not be included on the input
          control statements for GIMUNZIP processing."
          size="4800"
          hash="6DD284BE148DB4FA39CD5C3798F2171482D29553">
  </ARCHDEF>
</PKGDEF>

<?PKGHASH hash="1C4DC48997B66271202887E3A1628EA2ADDC8DDE"?>
```

Figure 101. Package attribute file (GIMPAF.XML) example (Part 2)

Note: Spaces have been added to the example for clarity, but are not necessarily produced by GIMZIP, nor are they required. Some of the values, such as size, original size, and hash, are for illustration purposes only, and may not be accurate.

Displaying the package attribute file

The package attribute file (GIMPAF.XML) can be rendered on an internet browser that supports XML and XSL (for example, Microsoft Internet Explorer 5). The package directory containing the package attribute file must be accessible to the browser through your z/OS web server. The accessibility of the package directory is determined by your specific web server configuration.

For example, suppose */u/userid/public/* is defined by your z/OS web server to be your publicly accessible directory. Further suppose, to view the contents of this directory, you must specify an URL of the form `http://host name/~userid`. Then, in order to view the package attribute file, your package directory must be contained within the */u/userid/public/* directory, or a symbolic link that points to the package directory must be in the */u/userid/public/* directory in a UNIX file system. Furthermore, your z/OS web server must be configured to recognize `.XML` and `.XSL` as text MIME types and, if appropriate, convert EBCDIC text to ASCII text.

The package attribute file can then be viewed on an internet browser by entering an URL similar to the following:

```
http://host name/~userid/GIMPAF.XML
```

Archive files

The actual software and associated materials for a package are stored in archive files. GIMZIP creates an archive file for each input data set or file specified by a <FILEDEF> tag. An archive file consists of a portable image of the original data set or file and attributes of the original data set or file needed to reload the data from the portable image.

The attributes of the original data set or file are recorded in a file attribute file in the UNIX file system called GIMFAF.XML. For data sets, the portable image of the original data set is stored temporarily as a file in the UNIX file system (this step is not required for UNIX files and directories). The name of this temporary file is the type attribute value for the archive (SMPPTFIN, SMPHOLD, SMPRELF) or the name MVSFILE if the archive has no type attribute specified. An archive file is then created by using the UNIX System Services pax command to combine the GIMFAF.XML file and the original data into an archive. Although a UNIX file is not first stored as a temporary file with a name matching its type attribute value, its name within the archive will be its type attribute value, or the name MVSFILE, if the archive has no type attribute specified. The GIMUNZIP service routine and the RECEIVE command both expect an archive file to contain the GIMFAF.XML component file, and the following, depending on the original content:

For data sets:
a file named SMPPTFIN, SMPHOLD, SMPRELF, or MVSFILE

For UNIX files:
a file named SMPPTFIN, SMPHOLD, or MVSFILE

For UNIX directories:
the directory and its contents are stored in the archive using the original names.

Data sets and files with a type of README are an exception to the previously described archive processing. README data sets and files have no file attribute file and are not archived using pax, but rather are stored unchanged in a file in the package directory. Likewise, README data sets and files are not subject to archive segmentation.

The absolute names for archive and README files in the package directory use the following format:

```
/package_directory/subdir/Snnnn.original__name.pax.Z
```

The absolute names for archive segment files in the package directory use the following format:

```
/package_directory/subdir/date_timeofday.nofm
```

package_directory
> indicates the package directory specified on the SMPDIR DD statement.

subdir
> indicates the subdirectory into which the archive file is stored. GIMZIP creates this subdirectory based on the filetype or the subdir attributes specified on the <FILEDEF> tag for the data set. If neither attribute is specified, then no subdirectory is used. If the file type is README, the subdirectory is created only if the subdir attribute has been specified.

S*nnnn*
> specifies the sequence indicator for the archive file. The sequence indicator is necessary so the correct order can be determined for the SMPHOLD and SMPPTFIN files when eventually processed by the SMP/E RECEIVE command. GIMZIP processes the input data sets or files in the sequence in which they are specified in the package control tags and the indicator is assigned accordingly. All archive files, except those specified with a file type of SMPRELF, are assigned a sequence indicator. If the file type is README and an archid tag was specified, then no sequence number will be assigned.

original_name
> specifies the original name of the data set, file, or directory, as indicated on the <FILEDEF> tag. This could also be the archid for UNIX directories and README data sets or files, or the date and time for UNIX directories, if the archid is not specified.

pax.Z
> is the file extension for all archive files, except those with a file type of README. The pax.Z extension indicates a file that has been processed by the UNIX System Services pax command with the compress option.

date_timeofday
> the date and time of day value, unique for each archive, but the same for each segment file of a particular archive.

n
> segment number

m
> the total number of segments for the archive.

File attribute file

The File Attribute File (FAF) is included in the archive file along with the data set, file or directory when it is archived. The FAF contains control tags that:

- Describe the attributes of the original data set, file or directory
- Provide information needed by GIMUNZIP to process the archive

The GIMUNZIP service routine uses the information in an FAF to reload the data from an archive file. The FAF for a file or directory in the UNIX file system is created in the temporary working directory, just as it is for data sets. The FAF control tags follow XML syntax rules, and their format is as follows:

File attribute file control tags (Part 1 of 2)
FILELIST

▶▶── <FILELIST> ──┬── FILEATTR for MVS data set ──┬── </FILELIST> ──◄►
　　　　　　　　　├── FILEATTR for UNIX file system ──┤
　　　　　　　　　└── FILEATTR for VSAM data set ──┘

FILEATTR for MVS data set
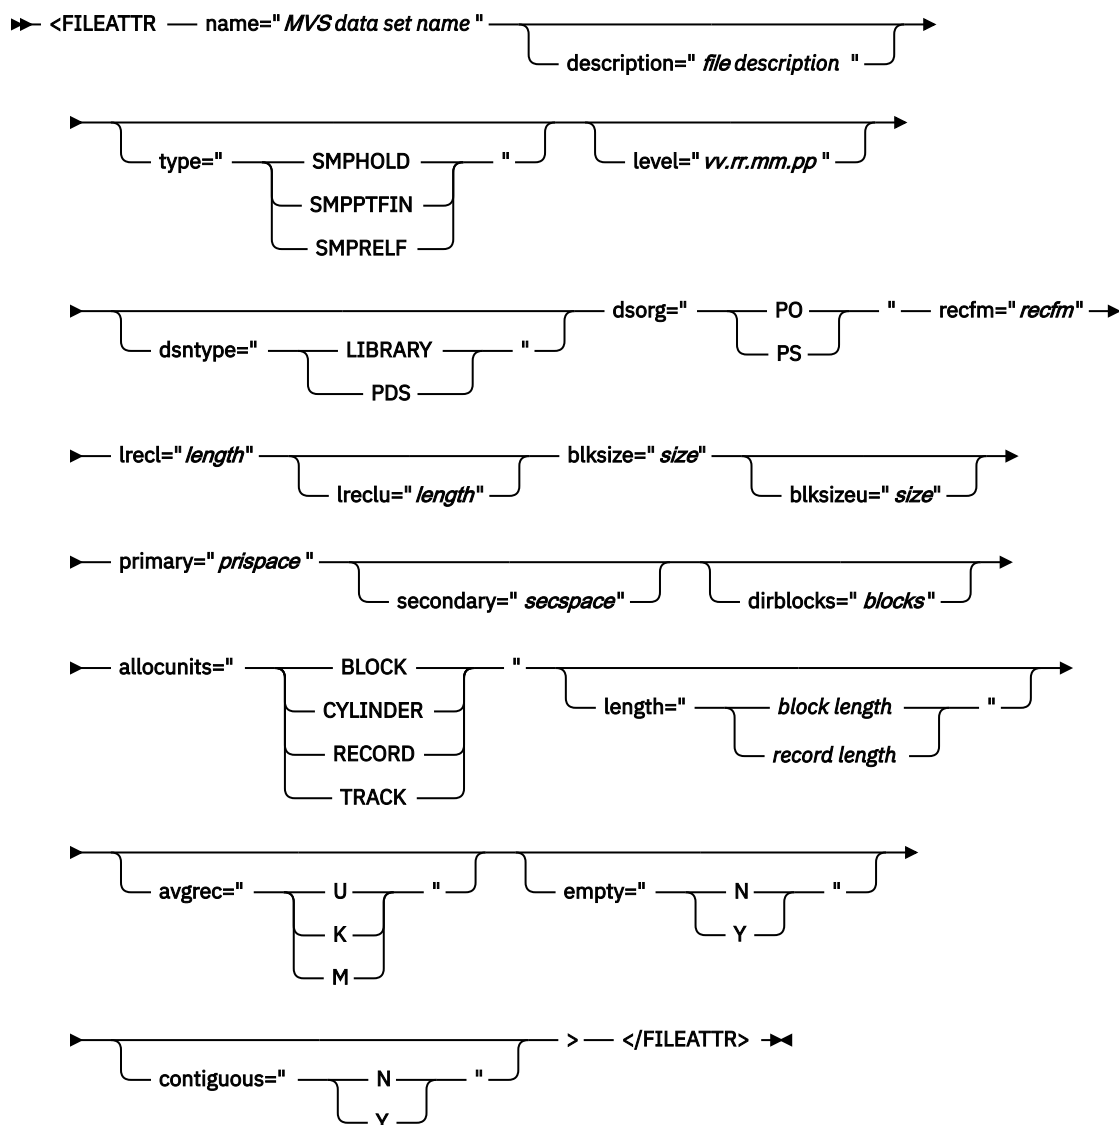
File attribute file control tags (Part 2 of 2)
FILEATTR for UNIX file system
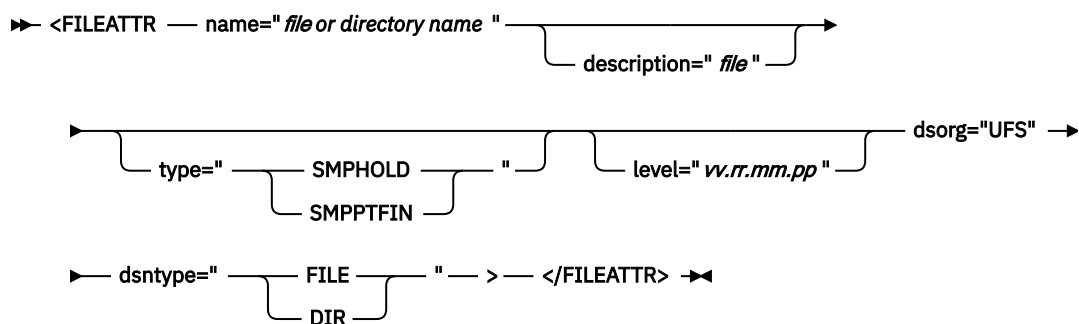
FILEATTR for VSAM data set

```
►►─ <FILEATTR ─── name=" VSAM cluster name "─┬──────────────────────────────┬─►
                                              └─ description=" file description "─┘

►─┬─────────────────────┬─ dsorg="VSAM" ── dsntype="─┬─ ESDS ─┬─"─►
  └─ level=" vv.rr.mm.pp "─┘                            ├─ KSDS ─┤
                                                        ├─ LDS ──┤
                                                        └─ RRDS ─┘

►── lreclu=" LRECL of REPROed data" ── blksizeu=" BLKSIZE of REPROed data" ── recfm=" recfm" ─►

►── primary=" prispace "─┬─────────────────────────┬─►
                          └─ secondary=" secspace"─┘

►── allocunits="─┬─ BLOCK ───┬─"─┬──────────────────────────────┬─►
                 ├─ CYLINDER ─┤   └─ length="─┬─ block length ──┬─"─┘
                 ├─ RECORD ───┤               └─ record length ─┘
                 └─ TRACK ────┘

►─┬──────────────────────┬─┬────────────────────┬─►
  └─ avgrec="─┬─ U ─┬─"─┘   └─ empty="─┬─ N ─┬─"─┘
              ├─ K ─┤                   └─ Y ─┘
              └─ M ─┘

►─┬───────────────────────────────┬─ > ─►
  └─ unloadformat="─┬─ REPRO ─┬─"─┘
                    └─ DUMP ──┘

►─┬──────────────────────────────────────────────────────────────────────────────────────────────┬─►
  └─ <DEFINECLUSTER ── component="─┬─ DATA ──┬─"> ─┬─ <OPERAND> IDCAMS DEFINE </OPERAND> ─┬─ </DEFINECLUSTER> ─┘
                                   └─ INDEX ─┘      └◄──────────────────────────────────┘

►── </FILEATTR> ─►◄
```
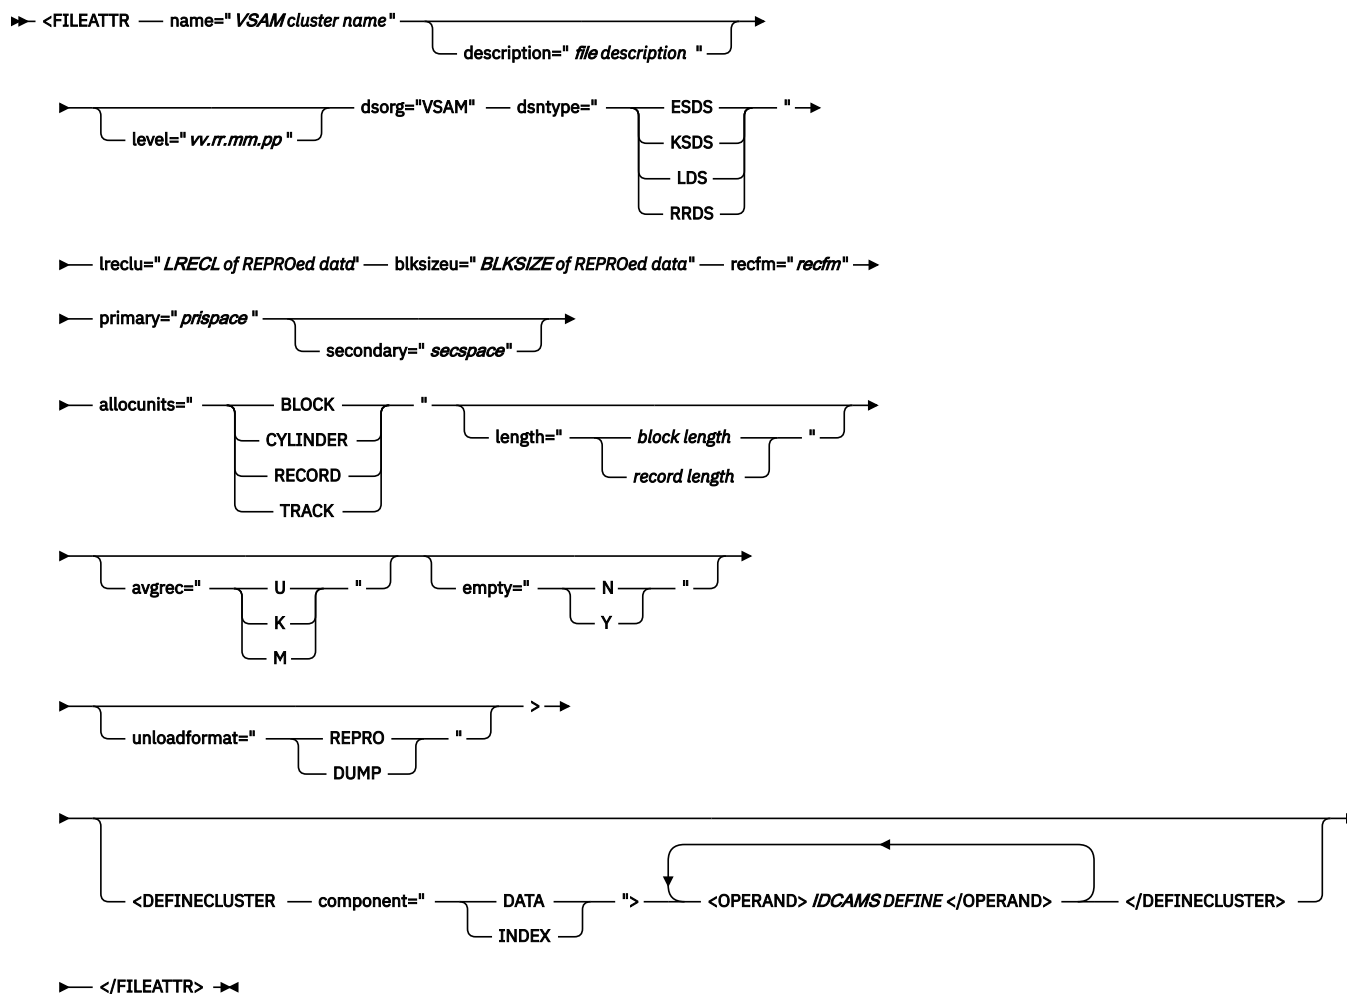
<FILELIST> Tag syntax

The <FILELIST> and corresponding </FILELIST> tags identify the beginning and end of a group of file definitions.

<FILEATTR> Tag syntax

The <FILEATTR> and corresponding </FILEATTR> tags identify the beginning and end of a file attribute definition for an archive file. The following attributes may be found on the <FILEATTR> tag:

name="*file name | pathname*"
    specifies one of the following:

- The name of the original MVS data set or VSAM cluster in this archive
- The pathname for the original UNIX file or directory in this archive

Note:

1. The first character of the absolute pathname of a file or directory in the UNIX file system must be a slash (/).
2. If the value of the name attribute does not begin with a slash, it is assumed to be an MVS data set name (sequential, partitioned, and VSAM data sets) and must conform to the MVS data set naming conventions.
3. When a file or directory in the UNIX file system is specified, the name can be from 1 to 1023 bytes long with 255 characters between delimiters (/). The value can be any character from X'40' through X'FE' except '<', '>', '&', and double quotation mark ("). All data beyond column 72 is

ignored, including blanks. The pathname of a file or directory in the UNIX file system is case sensitive and will not to be converted to uppercase alphabetic during GIMUNZIP processing. (Data set names are not case sensitive. They will be converted to uppercase alphabetic during GIMUNZIP processing.)

description="*file description*"
specifies a text description for the file. The value GIMZIP produces is copied from the description value, if specified, on the corresponding input <FILEDEF> tag. The description can be up to 500 characters long.

type="*file type*"
specifies a file type indicator that describes the contents of the file, identifies how the file is to be processed by GIMZIP, and also how the file will be processed by the SMP/E RECEIVE command. Any of the following values may be specified:

SMPHOLD
indicates that the data set or file contains SMP/E HOLDDATA statements.

SMPPTFIN
indicates that the data set or file contains SMP/E modification control statements (MCS).

SMPRELF
indicates that the data set contains an SMP/E RELFILE data set associated with a SYSMOD contained in the SMPPTFIN file.

If the type attribute is not specified, then the data set specified on the name attribute can be either sequential or partitioned.

level="*vv.rr.mm.pp*"
indicates the service level of the GIMZIP service routine used to create this file attribute file. The individual values are as follows:

vv
version

rr
release

mm
modification level

pp
PTF level

Note: When the PTF level (pp) of the GIMZIP service routine exceeds 2 characters, the first character will be dropped. For example, when the PTF level of the GIMZIP service routine reaches 100, the PTF level will start back at '00'.

dsntype="*type*"
specifies the type of partitioned data set, VSAM cluster, or UNIX file or directory for the original data.

DIR
indicates a directory in the UNIX file system.

FILE
indicates a file in the UNIX file system.

ESDS
indicates a VSAM cluster that is an entry sequence data set.

KSDS
indicates a VSAM cluster that is a keyed sequence data set.

LDS
indicates a VSAM cluster that is a linear data set.

RRDS
indicates a VSAM cluster that is a relative record data set.

LIBRARY
>   indicates a PDSE partitioned data set.

PDS
>   indicates a PDS partitioned data set.

dsorg="*organization*"
>   specifies the organization of the original data set, file, or directory.

PO
>   indicates a partitioned data set

PS
>   indicates a sequential data set

UFS
>   indicates a file or directory in the UNIX file system

VSAM
>   indicates a VSAM cluster

recfm="*format*"
>   specifies:
>
>   - When `dsorg="PS"` or `dsorg="PO"`, the record format of the sequential or partitioned MVS data set in an archive
>   - When `dsorg="VSAM"`, the record format of the data set used to contain the REPROed VSAM data
>
>   The value can be from 1 to 3 alphabetic characters.

lrecl="*length*"
>   specifies the logical record length of the original data set. The value can be from 1 to 5 decimal digits.

lreclu="*length*"
>   specifies:
>
>   - When `dsorg="PO"`, the logical record length of the IEBCOPY unloaded format of the original data set
>   - When `dsorg="VSAM"`, the logical record length of the data set used to contain the REPROed VSAM data from the cluster
>
>   In either case, the value can be from 1 to 5 decimal digits.

blksize="*size*"
>   specifies the block size of the original data set. The value can be from 1 to 5 decimal digits in the range of 00000 through 32760.

blksizeu="*size*"
>   specifies:
>
>   - When `dsorg="PO"`, the block size of the IEBCOPY unloaded format of the original data set
>   - When `dsorg="VSAM"`, the block size of the data set used to contain the REPROed VSAM data from the cluster
>
>   In either case, the value can be from 1 to 5 decimal digits in the range of 00000 through 32760.

primary="*prispace*"
>   specifies:
>
>   - When `dsorg="PS"` or `dsorg="PO"`, the primary space needed to reload the sequential or partitioned MVS data set in an archive
>   - When `dsorg="VSAM"`, the primary space needed to create a temporary data set to contain the REPROed VSAM data
>
>   The value must be non-zero and can be from 1 to 8 decimal digits.

secondary="*secspace*"
>   specifies:

- When dsorg="PS" or dsorg="PO", the secondary space needed to reload the sequential or partitioned MVS data set in an archive
- When dsorg="VSAM", the secondary space needed to create a temporary data set to contain the REPROed VSAM data

The value must be non-zero and can be from 1 to 8 decimal digits.

dirblocks="*blocks*"
specifies the number of directory blocks in the original data set. The value must be non-zero and can be from 1 to 5 decimal digits.

allocunits="*units*"
specifies:

- When dsorg="PS" or dsorg="PO", the units to be used when allocating the data set used to reload the sequential or partitioned MVS data set in an archive
- When dsorg="VSAM", the units to be used to create a temporary data set to contain the REPROed VSAM data

The value can be one of the following:

BLOCK
indicates space is to be allocated in blocks. The size of the blocks is specified on the length attribute.

CYLINDER
indicates space is to be allocated in cylinders.

RECORD
indicates space is to be allocated in records. The average length of these records is specified on the length attribute.

TRACK
indicates space is to be allocated in tracks.

length="*average record length | block length*"
specifies either the block length, or the average record length, used to allocate the primary and secondary space for reloading the data from the archive. If allocunits="BLOCK" is specified, then this value represents the block length. If allocunits="RECORD" is specified, then this value represents the average record length. The value of the length attribute must be from 1 to 5 decimal digits and must be non-zero.

avgrec="*multiplier*"
indicates the multiplier for the primary and secondary space allocation values only when allocunits="RECORD".

U
indicates that the primary and secondary space quantities represent the number of records in units (multiplier of 1).

K
indicates that the primary and secondary space quantities represent the number of records in thousands (multiplier of 1024 or 1K).

M
indicates that the primary and secondary space quantities represent the number of records in millions (multiplier of 1048576 or 1M).

- When dsorg="PS" or dsorg="PO", this value represents the multiplier to be used for allocating the sequential or partitioned data set in an archive.
- When dsorg="VSAM", this value represents the multiplier to be used for allocating the temporary data set to contain the REPROed VSAM data.

contiguous
indicates whether the originating data set's space was allocated contiguously or not. When the value is Y, then the originating data set's space was allocated contiguously. When the value is N or not

specified at all in the File Attribute File, then the originating data set's space was not allocated contiguously.

unloadformat
specifies the method used by GIMZIP to unload VSAM data sets

> **REPRO**
> The data set was unloaded using IDCAMS REPRO. This is the default value if **unloadformat** is not specified in the File Attribute File.

> **DUMP**
> The data set was unloaded using DFSMSdss ADRDSSU DUMP.

empty="Y | N"
indicates whether the original data set was empty. When the value is 'Y', the original data set was empty. When the value is 'N', or if the empty attribute is not specified, the original data set was not empty. If an original data set is empty, it is not included in the resulting archive file. Only its File Attribute File is included in the archive file.

<DEFINECLUSTER> Tag syntax

The <DEFINECLUSTER> tag and corresponding </DEFINECLUSTER> tag identify the beginning and end of information that is needed to define a new VSAM cluster during GIMUNZIP processing. The <DEFINECLUSTER> tag may be specified for the data component of a VSAM cluster. The <DEFINECLUSTER> tag may also be specified for the index component of a VSAM cluster if the cluster is a KSDS. The following attribute is found on the <DEFINECLUSTER> tag:

component=
specifies the name of the component type of a VSAM cluster, and may be one of the following values:

> **"DATA"**
> indicates that the following <OPERAND> tags are providing information to define the data component of a VSAM cluster.

> **"INDEX"**
> indicates that the following <OPERAND> tags are providing information to define the index component of a VSAM KSDS cluster.

Following the <DEFINECLUSTER> tag is the <OPERAND> tag.

<OPERAND> Tag syntax

The <OPERAND> tag and corresponding </OPERAND> tag identify the beginning and end of a valid IDCAMS DEFINE command operand for GIMUNZIP processing when a new VSAM cluster must be created in order to extract a VSAM archive. The <OPERAND> and </OPERAND>tags can only be specified within a <DEFINECLUSTER> tag and corresponding </DEFINECLUSTER> tag.

A maximum of 50 characters is allowed between the <OPERAND> tag and its corresponding </OPERAND> end tag. Multiple <OPERAND> tags may be specified to provide IDCAMS DEFINE command information.

Building packages manually

Note: A portable package of software can be built either using GIMZIP, or manually (without using GIMZIP). If a packager is building a GIMZIP-like package manually and always intends that the users of the package will predefine a VSAM cluster for usage by GIMUNZIP, the <DEFINECLUSTER> tag and its embedded <OPERAND> tags may be omitted from the FAF.

Syntax notes

1. Data in columns 73 through 80 is ignored. If data is specified beyond column 72, it is ignored and an error in a following tag may be indicated.
2. File definition control tags may contain comments. Comments start with <!-- (hex 4C5A6060) and end with --> (hex 60606E). The first --> encountered after the initial <!-- will end the comment. A comment may appear between a start-tag and its matching end-tag, but never within a tag.

3. The space allocation for an archived data set is stored in the File Attribute File in terms of AVGREC regardless of how the data set was originally allocated. GIMZIP converts the current allocation information for a data set (as obtained from its DSCBs) into record type allocation information using AVGREC (usually set to U, although K could be used for extremely large data sets) and an average record length equal to the data set block size. The data set allocated using this information when an archive is retrieved by GIMUNZIP will approximately equal in size (number of bytes) the original data set, but will not use the same allocation units as the original data set that was archived by GIMZIP.

File attribute file examples

The following is an example of a file attribute file for a sequential data set containing modification control statements.

```
<?xml version="1.0" ?>
<FILELIST>

  <FILEATTR name="SAMPLE.FMID001.SMPMCS"
            description="This is the SMPMCS file for FMID001."
            level="02.07.00.35"
            type="SMPPTFIN"
            dsorg="PS"
            recfm="FB"
            lrecl="80"
            blksize="6160"
            allocunits="RECORD"
            length="6160"
            avgrec="U"
            primary="24000">

  </FILEATTR>

</FILELIST>
```

Figure 102. File attribute file (GIMFAF.XML) example for a sequential data set

The file attribute file included with the archive for file /CPACTST/DO000029/SMPMCS from the example in Figure 100 on page 433 would look something like this:

```
<FILELIST>
<FILEATTR
  name="/CPACTST/DO000029/SMPMCS"
  description="This is the SMPMCS file for DO000029"
  level="03.03.00.00"
  type="SMPPTFIN"
  dsntype="FILE"
  dsorg="UFS">
</FILEATTR>
</FILELIST>
```

Figure 103. File attribute file example for a UNIX file

The file attribute file included with the archive for VSAM cluster CPACTST.DO000029.MVS.GLOBAL.CSI from the example in Figure 100 on page 433 would look something like this:

```
<FILELIST>
<FILEATTR
  name="CPACTST.DO000029.MVS.GLOBAL.CSI"
  description="This is an example VSAM cluster for the product."
  level="03.03.00.00"
  dsorg="VSAM"
  dsntype="KSDS"
  lreclu="32760"
  blksizeu="27998"
  recfm="VBS"
  primary="2004"
  secondary="1"
  allocunits="BLOCK"
  length="27998">
  <DEFINECLUSTER component="DATA">
    <OPERAND>KEYS(24 0)</OPERAND>
    <OPERAND>RECORDSIZE(24 143)</OPERAND>
    <OPERAND>FREESPACE(10 5)</OPERAND>
    <OPERAND>CONTROLINTERVALSIZE(8192)</OPERAND>
    <OPERAND>CYLINDERS(5 2)</OPERAND>
    <OPERAND>SHAREOPTIONS(2 3)</OPERAND>
  </DEFINECLUSTER>
  <DEFINECLUSTER component="INDEX">
    <OPERAND>CONTROLINTERVALSIZE(4096)</OPERAND>
    <OPERAND>TRACKS(1 1)</OPERAND>
    <OPERAND>SHAREOPTIONS(2 3)</OPERAND>
  </DEFINECLUSTER>
</FILEATTR>
</FILELIST>
```

Figure 104. File attribute file example for a VSAM cluster

Note: Blank lines and spaces have been added to the file attribute file for clarity, but are not necessarily produced by the GIMZIP service routine, nor are they required.

Return codes

GIMZIP may end with the following return codes:

Return code
Meaning

00

The input data was processed successfully.

04

A call to a system service may not have completed successfully.

12

One of the following:

- Required file attributes could not be obtained.
- Input data sets were not sequential, partitioned, or VSAM data sets, nor files and directories in the UNIX file system.
- Call to a required system service failed.
- Syntax errors in the SYSIN data set.
- Data sets could not be opened.
- SMPDIR was not allocated to a UNIX directory.
- SMPDIR directory was not empty.
- Data sets, directories, or files were missing.
- Required modules could not be loaded.

16

One of the following:

- An I/O error occurred.
- A syntax error was found on the EXEC statement parameters.

20

SMPOUT data set is missing.

> 20

Internal error. Report the error to the IBM Support Center.

Chapter 12. GIMIAP: Copy utility invocation program

The SMP/E GENERATE command creates a job stream that builds a set of target libraries from a set of distribution libraries. For data elements or hierarchical file system elements, the GENERATE command builds a job (DEIINST for data elements or HFSINST for hierarchical file system elements) to invoke the appropriate copy utility for those elements. Each job step in the job installs multiple elements from multiple distribution libraries into a single target library. This is done by invoking the SMP/E program GIMIAP, which calls the appropriate copy utility to do the actual installation.

Although GIMIAP is to be used only by SMP/E, you may need to understand the control statements passed to the program. For example, you may need to diagnose errors detected by GIMIAP (such as syntax errors or missing information). To help you with this task, this chapter describes:

- The control statements used to invoke GIMIAP
- The return codes issued by GIMIAP
- The JCL statements used in the job that invokes GIMIAP

Control statements used to invoke GIMIAP

The GIMIAP program runs as a background job and is driven by control statements that identify the following:

- The copy utility to be invoked
- The distribution and target libraries to be used for each invocation of GIMIAP
- The elements to be installed
- The parameters to be passed to the copy utility

These control statements are created by SMP/E during GENERATE processing and are for use during GIMIAP processing. Each invocation of GIMIAP can install many elements through multiple invocations of the copy utility. Each invocation of the copy utility by GIMIAP installs a single element. These are the control statements used for input to a GIMIAP step:

- INVOKE
- COPY
- SELECT

The INVOKE control statement

The INVOKE control statement identifies the copy utility that GIMIAP should invoke.

The INVOKE control statement is always produced. There is one INVOKE control statement for each job step (that is, one INVOKE control statement for each invocation of GIMIAP). The INVOKE control statement must be the first control statement and must be a single-card image (one 80-byte record).

INVOKE control statement

The following operands are required on the INVOKE statement:

COPY
Specifies the name of the copy utility to be used when copying data elements that do not need reformatting and are being installed into a partitioned data set. The value of COPY is either the value of the NAME operand in the UTILITY entry in effect at GENERATE time or the default of IEBCOPY.

HFSCOPY
Specifies the name of the copy utility to be used when copying hierarchical file system elements. The value of HFSCOPY is either the value of the NAME operand in the UTILITY entry in effect at GENERATE time or the default of BPXCOPY.

LIST
Applies only when copying data elements. It specifies whether member names should be listed during copy processing. This value is derived from the copy utility entry in effect when the GENERATE command was issued.

YES
Indicates that member names should be listed during copy processing. This is the default.

NO
Indicates that the list of member names should be suppressed during copy processing.

PRINT
Specifies the ddname that is to be used for print output generated by the specified copy utility. The value of PRINT is either the value of the PRINT operand in the UTILITY entry in effect at GENERATE time or the default of SYSPRINT.

Note: If SYSTSPRT is specified as the PRINT value for the copy utility, it is ignored and the default of SYSPRINT is used instead.

RC
Specifies the highest acceptable return code from the specified copy utility. The value of RC is either the value of the RC operand in the UTILITY entry in effect at GENERATE time or the default of 0.

REPLACE
Applies only when copying data elements. It specifies whether existing members are to be replaced when copying an element. If REPLACE is specified on the GENERATE command, then it is specified on the INVOKE control statement for data elements.

COPY control statement

A COPY control statement identifies the input and output libraries to be used for the SELECT control statements that follow it. Each COPY statement identifies a distribution library that has elements to be copied to a single target library. Multiple COPY statements can identify the same target library.

A COPY control statement must follow an INVOKE or SELECT control statement. Also, each COPY control statement must be followed by at least one SELECT control statement.

COPY control statement

►►─ COPY ── FROMLIB(*ddname*) ── TOLIB(*ddname*) ── • ─►◄

The following operands are required on the COPY statement:

FROMLIB
is the ddname used by the copy utility as its input file for the source of the element (that is, the distribution library). The name can be 1 to 8 characters and must be composed of uppercase alphabetic, numeric, or national ($, #, @) characters.

TOLIB
is the ddname used by the copy utility as its output file identifying where the element is to be installed (that is, the target library). The name can be 1 to 8 characters and must be composed of uppercase alphabetic, numeric, or national ($, #, @) characters.

The SELECT control statement

The SELECT control statement identifies the element to be installed and the operands to be passed to the copy utility or SMP/E to enable its installation. The SELECT control statement must follow a COPY control statement and may span multiple 80-byte records.

SELECT control statement

type
> is the type of the element to be processed. Any valid data element type or hierarchical file system element type, including those with national language identifiers, may be specified. An element type name is formed by stripping the ++ from the beginning of the name of the MCS for that element. For example, the element type for an element defined by a ++CLIST MCS is CLIST, while the element type for a ++HFSRMS MCS is HFSRMS. See "Data element MCS" on page 8 for a list of data element MCS and "Hierarchical file system element MCS" on page 21 for a list of hierarchical file system element MCS.

name
> is the name of element to be processed.

ALIAS
> specifies the list of alias names for the data element. This operand applies only to data elements.

EPARM
> is the parameter list to be passed for this invocation of the HFSCOPY utility. This operand applies only to hierarchical file system elements. The parameter list consists of an LL value and the actual execution parameters.

> • LL represents a halfword hexadecimal length of the character string (the HFSCOPY utility execution parameters) immediately following it as part of the EPARM value. The length of the character string described does not include the opening parenthesis preceding the LL value nor the closing parenthesis following the last option specified.

> No blanks are allowed between the opening parenthesis and the LL value. The opening parenthesis and the LL value must be in the same record.

> Because the LL value is nondisplayable (ready-to-use) hexadecimal, it may appear as blanks or odd characters. This is valid data and should not be removed or modified.

> • The *option* values are the execution parameters to be used by the HFSCOPY utility for this invocation of the utility. SMP/E always supplies execution parameters to the HFSCOPY utility, and the parameters are separated by commas with no intervening blanks.

> If the HFSCOPY UTILITY entry that is in effect supplies execution parameters, these values precede the SMP/E-generated information. For example, suppose the UTILITY entry has the following values:

> **NAME**
> > MYHFSCPY

> **PRINT**
> > MYPRINT

> **PARM**
> > A-PARM-FOR-MYHFSCPY

> The character string for the execution parameters would be generated as:

```
LLA-PARM-FOR-MYHFSCPY,ELEMENT(hfselm1),TYPE(TEXT),LINK('linknm01'))
```

If a PARM value of *user_info* is specified in an element's MCS, the character string for the execution parameters would be generated as:

```
LLA-PARM-FOR-MYHFSCPY,user_info,ELEMENT(hfselm1),TYPE(TEXT),LINK('linknm01'))
```

These are the parameters that are generated by SMP/E, using information in the hierarchical file system element entry:

ELEMENT(element_name)
> the name of the element to be installed in a UNIX file system. The *element_name* is an unquoted character string 1 to 8 bytes long. It is composed of uppercase alphabetic, numeric, or national ($, #, @) characters.

LINK('linkname1','linkname2','linkname3'...)
> the alternate names by which this element can be known in the target library within a UNIX file system. A *linkname* can be up to 1023 characters long and can contain special characters other than just uppercase alphabetic, numeric, or national ($, #, and @) characters.

> SMP/E always puts apostrophes around each *linkname* and separates multiple values with commas with no intervening blanks.

SYMLINK('linkname1','linkname2','linkname3'...)
> a list of symbolic links, which are file names that can be used as alternate names for referring to this element in a UNIX file system. Each linkname listed here is associated with a pathname listed in the SYMPATH operand. A symbolic link can be up to 1023 characters long and can contain characters in the range X'40' through X'FE'.

> SMP/E always puts apostrophes around each symbolic link and separates multiple values with commas with no intervening blanks.

SYMPATH('pathname1','pathname2','pathname3'...)
> a list of pathnames that are associated with symbolic links identified by the SYMLINK operand. A symbolic pathname can be up to 1023 characters long and can contain characters in the range X'40' through X'FE'.

> SMP/E always puts apostrophes around each symbolic pathname and separates multiple values with commas with no intervening blanks.

> SYMPATH appears if SYMLINK appears, otherwise it is omitted.

> For information about how the pathnames and linknames are associated, see the description of the SYMPATH operand and "Example 3: Packaging a SYSMOD with a symbolic link" on page 30 in "Hierarchical file system element MCS" on page 21.

TYPE(TEXT|BINARY)
> the installation format for the element in a UNIX file system. SMP/E generates this parameter from information stored in the element entry in the target zone. If no setting exists for TEXT mode or BINARY mode, SMP/E does not pass this parameter to the HFS copy utility.

SHSCRIPT
> indicates that a UNIX shell script is to be invoked to complete the installation of the selected element. This operand applies only to hierarchical file system elements whose ELEMENT entry contains a SHSCRIPT subentry.

/directory/file
> is the full path specification for the shell script to be invoked. The */directory/file* is derived from the SHELLSCR entry that matches the name in the hierarchical file system element's SHSCRIPT subentry. The SYSLIB subentry of the shell script identifies the ddname of the directory and the shell script name itself is the file.

PRE
> indicates the shell script is to be invoked before the selected element is copied to the directory in a UNIX file system. This value is obtained from the SHSCRIPT subentry of the selected element.

POST
indicates the shell script is to be invoked after the selected element is copied to the directory in a UNIX file system. This value is obtained from the SHSCRIPT subentry of the selected element.

Return codes

To help you diagnose errors, GIMIAP issues messages and return codes. The messages are documented in *z/OS SMP/E Messages, Codes, and Diagnosis*. Here is a description of the return codes:

Return code	Meaning
00	GIMIAP processing was successful. The message issued by GIMIAP states that the invocation of GIMIAP was successful.
04	GIMIAP processing was successful. The message issued by GIMIAP states that the invocation of GIMIAP was successful, but one or more of the following has occurred: • An ALIAS was not installed during element installation. • The COPY or HFSCOPY utility issued a non-zero return code that is less than or equal to the acceptable return code as defined in the RC operand of the INVOKE statement.
08	SMP/E processing errors occurred or a return code from the copy utility was higher than the acceptable return code specified on the INVOKE control statement. At least one element was not installed correctly, although an attempt was made to install all elements to the target library. Review the copy utility print output and SMPOUT output to determine the error. Correct the problem and rerun the job step.
12	GIMIAP encountered an invalid control statement. As a result, the elements in that job step were not installed. Once a control statement error is encountered, no more elements are processed for that job step. The statement must be corrected and the job step must be rerun. However, subsequent job steps will be processed. The message issued by GIMIAP indicates how many 80-byte records had been processed when the syntax error was encountered.
16	A severe error was encountered. As a result, the elements in that job step were not installed. After correcting the problem, the job step must be rerun. However, subsequent job steps will be processed. The message issued by GIMIAP indicates the cause of the problem, such as an I/O error on SYSIN, an open failure on SYSIN or another ddname, a syntax error on the EXEC parm, a syntax error on the control cards, or the absence of the copy utility.
20	A terminating error was encountered. The SMPOUT data set is not allocated or cannot be opened. GIMIAP is terminated.

JCL statements used in the DEIINST or HFSINST job

JCL statements are created for the DEIINST or HFSINST job during the GENERATE process. Along with the GIMIAP control statements for each job step, the JCL is composed of:

• A JOB statement. The JOB statement describes current installation-dependent parameters. The jobname is "DEIINST" when GIMIAP is to install data elements or "HFSINST" for hierarchical file system elements.

• One or more EXEC statements. The EXEC `statement` specifies PGM=GIMIAP,PARM='*option*'.

The step name is the name of the target library.

The following is an example of the EXEC statement to call GIMIAP:

```
//jobname  JOB …
//stepname EXEC PGM=GIMIAP,PARM='option'
```

Figure 105. JCL to call GIMIAP

EXEC
> is the statement used to call GIMIAP. The EXEC statement must specify PGM=GIMIAP. The following option may be specified on the EXEC statement PARM operand:

> **LANGUAGE=*xxx***
>> where *xxx* can be one of the following:

>> **ENU**
>>> US English

>> **JPN**
>>> Japanese

>> The LANGUAGE option defines which language to use for GIMIAP messages. LANGUAGE can also be specified as L. The LANGUAGE value will be the same as was used on GIMSMP while the GENERATE command was executed. If LANGUAGE is not specified, the default is LANGUAGE=ENU.

- **DD statements.** The DD statements define the data sets to be used by GIMIAP processing. The following ddnames are required:

dlib
> The file identified by a FROMLIB operand on a COPY control statement.

syslib
> The file identified by a TOLIB operand on a COPY control statement.

print
> This ddname is used by the copy utility for messages and processing information. If no DDDEF is found for the print file during GENERATE processing, the default of SYSOUT=* is generated. The ddname is the same as is specified for the PRINT operand of the INVOKE statement being used for the current invocation of GIMIAP.

SYSIN
> Defines the input control stream for GIMIAP.

SMPOUT
> The file to be used by the GIMIAP for message writing.

SYSPRINT
> Contains output when GIMIAP is processing data elements or hierarchical file system elements with shell scripts.

The following DD statements are required when processing data elements:

SYSUT1
> The file identified for containing the input statements to the copy program.

SYSUT3
> A work file to be used by the copy program.

SYSUT4
> A work file to be used by the copy program.

The following DD statement is required when processing file system elements that require the invocation of a UNIX shell script, and the shell script uses Java commands:

SMPJHOME
> Directory in the UNIX file system that contains the Java runtime.

During GIMIAP processing (invoked from the JCL produced by the GENERATE command), SMP/E can invoke a UNIX shell script to perform installation activities on behalf of a ++HFS file. If the UNIX shell script issues a Java command, the SMPJHOME DD statement is required.

An SMPJHOME DD statement is created in the HFSINST job by the GENERATE command, if an SMPJHOME DDDEF entry is defined at the time the GENERATE command is run.

Figure 106 on page 451 illustrates the JCL statements that could be generated to invoke GIMIAP processing for data elements and Figure 107 on page 452 illustrates the JCL statements that could be generated to invoke GIMIAP processing for hierarchical file system elements.

```
        ----+----1----+----2----+----3----+----4----+----5----+----6----+----7----+----8
 1
       //DEIINST  JOB 'accounting info',MSGLEVEL=(1,1)
 2
       //tgtlib01  EXEC PGM=GIMIAP,PARM='LANGUAGE=ENU'                    EXECDEI
       //distlib1 DD distribution library DD info.                       distlib1
       //distlib2 DD distribution library DD info.                       distlib2
       //tgtlib01 DD target library 1 DD info                            tgtlib01
       //* ------ work data sets ----------------
       //SYSUT1   DD UNIT=SYSDA,SPACE=(TRK,(x,y)),DISP=(NEW,DELETE)       SYSUT1
       //SYSUT3   DD UNIT=SYSDA,SPACE=(TRK,(x,y)),DISP=(NEW,DELETE)       SYSUT3
       //SYSUT4   DD UNIT=SYSDA,SPACE=(TRK,(x,y)),DISP=(NEW,DELETE)       SYSUT4
       //SYSPRINT DD SYSOUT=*                                            DEFAULT
       //SMPOUT   DD SYSOUT=*                                            SMPOUT
       //SYSIN    DD *                                                   DEFAULT
       INVOKE COPY(IEBCOPY) RC(0) PRINT(SYSPRINT) REPLACE.
       COPY FROMLIB(distlib1) TOLIB(tgtlib01).
       SELECT CLIST(elem1).
       SELECT CLIST(elem2) ALIAS(elemA,elemB).
       COPY FROMLIB(distlib2) TOLIB(tgtlib01).
       SELECT CLIST(elem3) ALIAS(elemC).
       SELECT CLIST(elem4).
       /*
 2
       //tgtlib02  EXEC PGM=GIMIAP,PARM='LANGUAGE=ENU'                    EXECDEI
       //distlib3 DD distribution library DD info.                       distlib3
       //distlib4 DD distribution library DD info.                       distlib4
       //tgtlib02 DD target library 2 DD info.                           tgtlib02
       //* ------ work data sets ----------------
       //SYSUT1   DD UNIT=SYSDA,SPACE=(TRK,(x,y)),DISP=(NEW,DELETE)       SYSUT1
       //SYSUT3   DD UNIT=SYSDA,SPACE=(TRK,(x,y)),DISP=(NEW,DELETE)       SYSUT3
       //SYSUT4   DD UNIT=SYSDA,SPACE=(TRK,(x,y)),DISP=(NEW,DELETE)       SYSUT4
       //SYSPRINT DD SYSOUT=*                                            DEFAULT
       //SMPOUT   DD SYSOUT=*                                            SMPOUT
       //SYSIN    DD *                                                   DEFAULT
       INVOKE COPY(IEBCOPY) RC(0) PRINT(SYSPRINT) REPLACE.
       COPY FROMLIB(distlib3) TOLIB(tgtlib02).
       SELECT HELPENU(elem3).
       SELECT HELPENU(elem4) ALIAS(elemD,elemE).
       COPY FROMLIB(distlib4) TOLIB(tgtlib02).
       SELECT HELPENU(elem5) ALIAS(elemF).
       SELECT PNLENU(elem6).
       /*
```

Additional information:

1

The **job name** "DEIINST" is generated during the job creation process. A single job installs the data elements.

2

The **job step** name is the target library. Each job step installs multiple data elements from multiple distribution libraries into a single target library.

Figure 106. Sample DEIINST job for GIMIAP

```
        ----+----1----+----2----+----3----+----4----+----5----+----6----+----7----+----8
 1
        //HFSINST   JOB 'accounting info',MSGLEVEL=(1,1)
 2
        //tghfs01  EXEC PGM=GIMIAP                                        EXECIAP
        //distlib1 DD distribution library DD info.                      distlib1
        //distlib2 DD distribution library DD info.                      distlib2
        //tghfs01  DD PATH='/Hierarchical/File/System/library1/DD/info/'  tghfs01
        //SYSPRINT DD SYSOUT=*                                           DEFAULT
        //SYSIN    DD *                                                  DEFAULT
        INVOKE HFSCOPY(BPXCOPY) RC(0) PRINT(SYSPRINT)
        COPY FROMLIB(distlib1) TOLIB(tgthfs01)
        SELECT HFS(hfselm1) EPARM(LLELEMENT(hfselm1),LINK('linkname01','linkname
        thatislongerthanthefirstoneandshouldbemoreuserfriendly'),TYPE(BINARY))
        SELECT HFS(hfselm2) EPARM(LLELEMENT(hfselm2),LINK('linknm01'),TYPE(TEXT)
        )
        COPY FROMLIB(distlib2) TOLIB(tgthfs01)
        SELECT HFS(hfselm3) EPARM(LLELEMENT(hfselm3),TYPE(TEXT))
        SELECT HFS(hfselm4) EPARM(LLELEMENT(hfselm4),LINK('linknm01','linknm02',
        'linknm03','linknm04','linknm05'))
        /*
 2
        //tghfs02  EXEC PGM=GIMIAP                                        EXECIAP
        //distlib3 DD distribution library DD info.                      distlib3
        //distlib4 DD distribution library DD info.                      distlib4
        //tghfs02  DD PATH='/Hierarchical/File/System/library2/DD/info/'  tghfs02
        //SYSPRINT DD SYSOUT=*                                           DEFAULT
        //SYSIN    DD *                                                  DEFAULT
        INVOKE HFSCOPY(BPXCOPY) RC(0) PRINT(SYSPRINT)
        COPY FROMLIB(distlib3) TOLIB(tgthfs02)
        SELECT HFS(hfselm3) EPARM(LLELEMENT(hfselm3),LINK('linknm01','linknm02',
        'linknm03','linknm04'),TYPE(TEXT))
        SELECT HFS(hfselm4) EPARM(LLELEMENT(hfselm4),LINK('linknm01'),TYPE(TEXT)
        )
        COPY FROMLIB(distlib4) TOLIB(tgthfs02)
        SELECT HFS(hfselm5) EPARM(LLELEMENT(hfselm5),TYPE(TEXT))
        SELECT HFS(hfselm6) EPARM(LLELEMENT(hfselm6),LINK('linknm01','linknm02',
        'linknm03','linknm04','linknm05'),TYPE(TEXT))
        /*
```

Additional information:

1

The **job name** "HFSINST" is generated during the job creation process. A single job installs all hierarchical file system elements.

2

The **job step name** on the EXEC statement is the target library within a UNIX file system. Each job step installs multiple hierarchical file system elements from multiple distribution libraries into a single target library within a UNIX file system.

Figure 107. Sample HFSINST job for invoking GIMIAP

Appendix A. SMP/E naming conventions

This appendix describes the naming conventions used by SMP/E. Table 62 on page 453 summarizes the naming conventions you need to follow when using SMP/E. After this table are details on the naming conventions IBM follows for:

- HOLD reason IDs and classes
- Source IDs
- SYSMODs

Table 62. Summary of SMP/E naming conventions

Entry or value	Number of characters	Other requirements
DDDEF entry	1–8 alphabetic (A–Z), national (@, #, or $), or numeric (0–9) characters	- The first character must be alphabetic or national - Must match ddname of data set
DLIB zone or DLIBZONE entry	1–7 alphanumeric (A–Z, 0–9) or national ($, #, @) characters	The first character must be alphabetic.

Element entry	1–8 alphanumeric (A–Z, 0–9) characters	Naming convention for first character:
		A–I
		    Used by IBM
		J–Z
		    Available for users

FMID *see* SYSMOD ID		
FMIDSET entry	1–8 alphanumeric (A–Z, 0–9) characters	
Global zone or GLOBALZONE entry		Must be GLOBAL

Hold class	1–7 alphanumeric (A–Z, 0–9) characters	Naming convention for first character:
		A
		    Reserved for IBM use
		V
		    Reserved for user-assigned values
		For a description of the hold classes used by IBM, see "Class values" on page 456.

Hold reason ID: error	1–7 alphanumeric (A–Z, 0–9) characters	Number of the APAR used to report an error in a PTF For details on the conventions IBM follows, see "Error reason IDs" on page 455.

Table 62. Summary of SMP/E naming conventions (continued)

Entry or value	Number of characters	Other requirements

Hold reason ID: system	1–7 alphanumeric (A–Z, 0–9) characters	Naming convention for first character:
		
		A–U
		    Reserved for IBM use
		
		V–Z
		    Reserved for user-assigned values
		
		For a description of the system reason IDs used by IBM, see "System reason IDs" on page 455.

Hold reason ID: user	1–7 alphanumeric (A–Z, 0–9) characters	To avoid conflicts with IBM reason IDs, follow the conventions for system reason IDs. For details on the conventions IBM follows, see "User reason IDs" on page 456.
OPTIONS entry	1–8 alphanumeric (A–Z, 0–9) characters	
Source ID	1–64 characters in length, consisting of any nonblank character (X'41' through X'FE') except single quotation mark ('), asterisk (*), percent (%), comma (,), left parenthesis ((), and right parenthesis ()).	For a description of the source IDs used by IBM, see "Naming conventions for source IDs" on page 457.

SYSMOD ID	7 alphanumeric (A–Z, 0–9) characters	Must start with a letter. Naming convention for first character:
		
		A–K
		    Used by IBM for functions
		
		L–T
		    Available for users
		
		U
		    Used by IBM for PTFs
		
		V–Z
		    Used by IBM for APARs
		
		For details on the conventions IBM follows, see "Naming conventions for SYSMODs" on page 458.

Target zone or TARGETZONE entry	1–7 alphanumeric (A–Z, 0–9) or national ($, #, @) characters	The first character must be alphabetic.
UTILITY entry	1–8 alphanumeric (A–Z, 0–9) characters	Must match the associated name specified in the appropriate OPTIONS entry Should match the name of the associated utility program
ZONESET entry	1–8 alphanumeric (A–Z, 0–9) characters	Avoid using the same name as any of the target or DLIB zones defined in the global zone containing the ZONESET entry

Naming conventions for HOLD reason IDs and HOLD classes

The ++HOLD statement prevents SMP/E from installing a SYSMOD until some special action is taken. The type of action is indicated by the reason ID or class specified on the ++HOLD statement. A reason ID or class value can contain from 1 to 7 alphanumeric characters. To prevent conflicts between IBM- and user-specified values, there are naming conventions for the three types of HOLD reason IDs (error, system, and user), as well as for HOLD classes.

Error reason IDs

The reason ID for an error HOLD is the number of the APAR used to report an error in a PTF. Therefore, error reason IDs follow the naming conventions for APARs. These are described under "PTF, APAR, and USERMOD SYSMOD IDs" on page 458.

System reason IDs

The reason ID for a system HOLD is generally a brief indication of the kind of processing the SYSMOD requires. These are the values currently used by IBM:

ID
> **Explanation**

ACTION
> The SYSMOD needs special handling before or during APPLY processing, ACCEPT processing, or both.

AO
> The SYSMOD may require action to change automated operations procedures and associated data sets and user exits in products or in customer applications. The PTF cover letter describes any changes (such as to operator message text, operator command syntax, or expected actions for operator messages and commands) that can affect automation routines.

DB2BIND
> A DB2 application REBIND is required for the SYSMOD to become effective.

DDDEF
> Data set changes or additions as required.

DELETE
> The SYSMOD contains a ++DELETE MCS, which deletes a load module from the system.

DEP
> The SYSMOD has a software dependency.

DOC
> The SYSMOD has a documentation change that should be read before the SYSMOD is installed.

DOWNLD
> Code that is shipped with maintenance that needs to be downloaded.

DYNACT
> The changes supplied by the SYSMOD may be activated dynamically without requiring an IPL. The HOLD statement describes the instructions required for dynamic activation. If those instructions are not followed, then an IPL is required for the SYSMOD to take effect.

EC
> The SYSMOD needs a related engineering change.

ENH
> The SYSMOD contains an enhancement, new option or function. The HOLD statement provides information to the user regarding the implementation and use of the enhancement.

EXIT
> The SYSMOD contains changes that may affect a user exit. For example, the interface for an exit may be changed, an exit may need to be reassembled, or a sample exit may be changed.

EXRF

The SYSMOD must be installed in both the active and the alternative Extended Recovery Facility (XRF) systems at the same time to maintain system compatibility. (If you are not running XRF, you should bypass this reason ID.)

FULLGEN

The SYSMOD needs a complete system or subsystem generation to take effect.

IOGEN

The SYSMOD needs a system or subsystem I/O generation to take effect.

IPL

The SYSMOD requires an IPL to become effective. For example, the SYSMOD may contain changes to LPA or NUCLEUS, the changes may require a CLPA, or a failure to perform an IPL might lead to catastrophic results, such as could be caused by activation of a partial fix.

Note: If you plan to perform an IPL with CLPA after the SYSMOD has been applied, then no further investigation of the HOLD is required; simply bypass the IPL reason ID. However, if you are not planning to perform an IPL with CLPA, then the details of the HOLD statement must be investigated to determine what kind of actions are required to activate the SYSMOD.

MSGSKEL

This SYSMOD contains message changes that must be compiled for translated versions of the message changes to become operational on extended TSO consoles.

If you want to use translated versions of the messages, you must run the message compiler once for the library containing the English message outlines, and once for each additional language you want to be available on your system. For details, see *z/OS MVS Planning: Operations*.

If you want to use **only** the English version of the messages, you do not need to run the message compiler. You should bypass this reason ID.

MULTSYS

Identifies fixes that need to be applied to multiple systems, in one of three cases: preconditioning, coexistence, or exploitation.

RESTART

To become effective, the SYSMOD requires a special subsystem restart operation. The HOLD statement contains information regarding the required restart actions.

User reason IDs

The reason ID for a user HOLD is whatever you think is appropriate to describe why the SYSMOD should be held. Because IBM does not use these reason IDs, there are no restrictions on the values you can use. However, to prevent possible confusion with other IBM reason IDs, follow the naming conventions for system reason IDs.

Class values

A class value indicates an alternative way to release a held SYSMOD. These are the values currently used by IBM:

Class
    Explanation

ERREL

The SYSMOD is held for an error reason ID but should be installed anyway. IBM has determined that the problem the SYSMOD resolves is significantly more critical than the error reflected by the holding APAR.

HIPER

The SYSMOD is held with a hold class of HIPER (High Impact)

PE
The SYSMOD is held with a hold class of "PTF in Error".

UCLREL
UCLIN needed for the SYSMOD has been handled by IBM and no longer requires your attention.

YR2000
Identifies PTFs that provide Year 2000 function, or fix a Year 2000-related problem.

Naming conventions for source IDs

With the SOURCEID operand, you can associate a name (a source ID) with each SYSMOD processed by a single RECEIVE command, and then later process those SYSMODs as a group by using that source ID. The following values are currently used by IBM.

RSU*yymm*
This is the indicator used to identify a Recommended Service Upgrade. For more information about RSU, see *z/OS SMP/E User's Guide*.

PUT*yymm*
Each PTF is assigned a source ID in the format PUT*yymm*.

- *yy* is the year the PTF was made available as preventive service.
- *mm* is the month the PTF was made available as preventive service.

The PUT*yymm* source ID represents a monthly accumulation of PTFs that have been COR-closed (made available for corrective service) within a given month and assigned a RETAIN VOLID indicating this. The RETAIN VOLID reflects the monthly PUT*yymm* level no later than the 15th of the following month. For example, PUT0706 represents PTFs that were COR-closed in June 2008 and assigned a RETAIN VOLID of 0706 between July 1 and July 15.

SMCCOR
This indicates a PTF that is approved for distribution, but has not been assigned a monthly PUT*yymm* source ID at the time your order was created. SMCCOR PTFs are not necessary for installation; however, they are made available for corrective service. If you plan to install any of the PTFs with this source ID, refer to the CORPE PSP bucket or latest available PUT bucket for the most current HOLD information. These PTFs are provided in case you experience the problem they fix and need to install them as corrective service.

SMCREC
This indicates a PTF that is approved for distribution and is recommended for installation, but has not been assigned a monthly PUT*yymm* source ID at the time your order was created. These PTFs should be installed on your system. They were assigned a source ID of SMCREC for one of the following reasons:

- The PTF is needed to install a product.
- The PTF is needed to support new hardware.
- The PTF resolves an error for another PTF (PE-PTF).

To further enhance your ability to selectively install PUT*yymm* and SMCREC PTFs, the additional SOURCEIDs listed in this section are assigned to the applicable PTFs. (There will be PUT*yymm* and SMCREC PTFs that do not fit into the following categories and, therefore, are not assigned multiple source IDs.)

HIPER
Identifies PTFs resolving a high-impact APAR

SPE
Identifies PTFs that are small programming enhancements

PRP
Identifies PTFs that resolve PTFs in error

OS390R*n*

    This indicates a PTF that is at the integration tested level for OS/390 Release *n*.

YR2000

    This identifies PTFs that provide Year 2000 function, or fix a Year 2000-related problem.

Naming conventions for SYSMODs

The specific naming conventions for a SYSMOD ID depend on the type of SYSMOD: function, PTF, APAR, or USERMOD. The following sections define IBM's naming conventions for SYSMOD IDs. This information is provided to help you develop a naming scheme for your own SYSMODs and avoid conflicts with IBM-written SYSMODs. However, these conventions are **not** requirements for user-written SYSMODs, and SMP/E does not check whether a SYSMOD ID follows these conventions.

Function SYSMOD IDs

The IBM convention for the SYSMOD ID of a function SYSMOD is *tcccrrr*, where:

t

    is the type of function. These are the values used by IBM:

    A

        Licensed vendor or business partner base function

    B

        Licensed vendor or business partner dependent function

    C, D

        Reserved for future use

    E

        Unlicensed base function

    F

        Unlicensed dependent function

    G

        Reserved for future use

    H

        Licensed base function

    I

        Reserved for future use

    J

        Licensed dependent function

    K

        Reserved for future use

    Names of user-written functions can start with any letters other than these.

ccc

    For functions provided by IBM, this is the product code. It must be three alphanumeric characters (no vowels).

rrr

    For functions provided by IBM, this is the release number. It identifies a specific function within a product. These three characters identify a specific release of a product function and must be alphanumeric.

Note: The former convention for the SYSMOD ID of a function was *tvvvrrr*.

PTF, APAR, and USERMOD SYSMOD IDs

The IBM convention for the SYSMOD ID of a service SYSMOD (APAR, APAR fix, PTF) or USERMOD is *tannnnn*, where:

t

    identifies the type of SYSMOD. It is a single alphanumeric character. These are the values used by IBM:

    A–K

        Used by IBM for various levels of an APAR fix

    L–T

        Available for users

    U

        Used by IBM for PTFs

    V–Z

        Used by IBM for various levels of an APAR fix

a

    is any alphabetic character (A-Z). Any valid character can be used for user SYSMODs.

nnnnn

    is an additional identifier for the SYSMOD. For PTFs and APAR fixes supplied by IBM, it is a number from 00001 to 99999. Any valid characters can be used for user SYSMODs.

Appendix B. Accessibility

Accessible publications for this product are offered through IBM Documentation (www.ibm.com/docs/en/zos).

If you experience difficulty with the accessibility of any z/OS information, send a detailed message to the Contact the z/OS team web page (www.ibm.com/systems/campaignmail/z/zos/contact_z) or use the following mailing address.

IBM Corporation
Attention: MHVRCFS Reader Comments
Department H6MA, Building 707
2455 South Road
Poughkeepsie, NY 12601-5400
United States

Notices

This information was developed for products and services that are offered in the USA or elsewhere.

IBM may not offer the products, services, or features discussed in this document in other countries. Consult your local IBM representative for information on the products and services currently available in your area. Any reference to an IBM product, program, or service is not intended to state or imply that only that IBM product, program, or service may be used. Any functionally equivalent product, program, or service that does not infringe any IBM intellectual property right may be used instead. However, it is the user's responsibility to evaluate and verify the operation of any non-IBM product, program, or service.

IBM may have patents or pending patent applications covering subject matter described in this document. The furnishing of this document does not grant you any license to these patents. You can send license inquiries, in writing, to:

IBM Director of Licensing
IBM Corporation
North Castle Drive, MD-NC119
Armonk, NY 10504-1785
United States of America

For license inquiries regarding double-byte character set (DBCS) information, contact the IBM Intellectual Property Department in your country or send inquiries, in writing, to:

Intellectual Property Licensing
Legal and Intellectual Property Law
IBM Japan Ltd.
19-21, Nihonbashi-Hakozakicho, Chuo-ku
Tokyo 103-8510, Japan

The following paragraph does not apply to the United Kingdom or any other country where such provisions are inconsistent with local law: INTERNATIONAL BUSINESS MACHINES CORPORATION PROVIDES THIS PUBLICATION "AS IS" WITHOUT WARRANTY OF ANY KIND, EITHER EXPRESS OR IMPLIED, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF NON-INFRINGEMENT, MERCHANTABILITY OR FITNESS FOR A PARTICULAR PURPOSE. Some states do not allow disclaimer of express or implied warranties in certain transactions, therefore, this statement may not apply to you.

This information could include technical inaccuracies or typographical errors. Changes are periodically made to the information herein; these changes will be incorporated in new editions of the publication. IBM may make improvements and/or changes in the product(s) and/or the program(s) described in this publication at any time without notice.

This information could include missing, incorrect, or broken hyperlinks. Hyperlinks are maintained in only the HTML plug-in output for IBM Documentation. Use of hyperlinks in other output formats of this information is at your own risk.

Any references in this information to non-IBM websites are provided for convenience only and do not in any manner serve as an endorsement of those websites. The materials at those websites are not part of the materials for this IBM product and use of those websites is at your own risk.

IBM may use or distribute any of the information you supply in any way it believes appropriate without incurring any obligation to you.

Licensees of this program who wish to have information about it for the purpose of enabling: (i) the exchange of information between independently created programs and other programs (including this one) and (ii) the mutual use of the information which has been exchanged, should contact:

IBM Corporation
Site Counsel
2455 South Road

463

Such information may be available, subject to appropriate terms and conditions, including in some cases, payment of a fee.

The licensed program described in this document and all licensed material available for it are provided by IBM under terms of the IBM Customer Agreement, IBM International Program License Agreement or any equivalent agreement between us.

Any performance data contained herein was determined in a controlled environment. Therefore, the results obtained in other operating environments may vary significantly. Some measurements may have been made on development-level systems and there is no guarantee that these measurements will be the same on generally available systems. Furthermore, some measurements may have been estimated through extrapolation. Actual results may vary. Users of this document should verify the applicable data for their specific environment.

Information concerning non-IBM products was obtained from the suppliers of those products, their published announcements or other publicly available sources. IBM has not tested those products and cannot confirm the accuracy of performance, compatibility or any other claims related to non-IBM products. Questions on the capabilities of non-IBM products should be addressed to the suppliers of those products.

All statements regarding IBM's future direction or intent are subject to change or withdrawal without notice, and represent goals and objectives only.

This information contains examples of data and reports used in daily business operations. To illustrate them as completely as possible, the examples include the names of individuals, companies, brands, and products. All of these names are fictitious and any similarity to the names and addresses used by an actual business enterprise is entirely coincidental.

COPYRIGHT LICENSE:

This information contains sample application programs in source language, which illustrate programming techniques on various operating platforms. You may copy, modify, and distribute these sample programs in any form without payment to IBM, for the purposes of developing, using, marketing or distributing application programs conforming to the application programming interface for the operating platform for which the sample programs are written. These examples have not been thoroughly tested under all conditions. IBM, therefore, cannot guarantee or imply reliability, serviceability, or function of these programs. The sample programs are provided "AS IS", without warranty of any kind. IBM shall not be liable for any damages arising out of your use of the sample programs.

Terms and conditions for product documentation

Permissions for the use of these publications are granted subject to the following terms and conditions.

Applicability

These terms and conditions are in addition to any terms of use for the IBM website.

Personal use

You may reproduce these publications for your personal, noncommercial use provided that all proprietary notices are preserved. You may not distribute, display or make derivative work of these publications, or any portion thereof, without the express consent of IBM.

Commercial use

You may reproduce, distribute and display these publications solely within your enterprise provided that all proprietary notices are preserved. You may not make derivative works of these publications, or

reproduce, distribute or display these publications or any portion thereof outside your enterprise, without the express consent of IBM.

Rights

Except as expressly granted in this permission, no other permissions, licenses or rights are granted, either express or implied, to the publications or any information, data, software or other intellectual property contained therein.

IBM reserves the right to withdraw the permissions granted herein whenever, in its discretion, the use of the publications is detrimental to its interest or, as determined by IBM, the above instructions are not being properly followed.

You may not download, export or re-export this information except in full compliance with all applicable laws and regulations, including all United States export laws and regulations.

IBM MAKES NO GUARANTEE ABOUT THE CONTENT OF THESE PUBLICATIONS. THE PUBLICATIONS ARE PROVIDED "AS-IS" AND WITHOUT WARRANTY OF ANY KIND, EITHER EXPRESSED OR IMPLIED, INCLUDING BUT NOT LIMITED TO IMPLIED WARRANTIES OF MERCHANTABILITY, NON-INFRINGEMENT, AND FITNESS FOR A PARTICULAR PURPOSE.

IBM Online Privacy Statement

IBM Software products, including software as a service solutions, ("Software Offerings") may use cookies or other technologies to collect product usage information, to help improve the end user experience, to tailor interactions with the end user, or for other purposes. In many cases no personally identifiable information is collected by the Software Offerings. Some of our Software Offerings can help enable you to collect personally identifiable information. If this Software Offering uses cookies to collect personally identifiable information, specific information about this offering's use of cookies is set forth below.

Depending upon the configurations deployed, this Software Offering may use session cookies that collect each user's name, email address, phone number, or other personally identifiable information for purposes of enhanced user usability and single sign-on configuration. These cookies can be disabled, but disabling them will also eliminate the functionality they enable.

If the configurations deployed for this Software Offering provide you as customer the ability to collect personally identifiable information from end users via cookies and other technologies, you should seek your own legal advice about any laws applicable to such data collection, including any requirements for notice and consent.

For more information about the use of various technologies, including cookies, for these purposes, see IBM's Privacy Policy at ibm.com®/privacy and IBM's Online Privacy Statement at ibm.com/privacy/details in the section entitled "Cookies, Web Beacons and Other Technologies," and the "IBM Software Products and Software-as-a-Service Privacy Statement" at ibm.com/software/info/product-privacy.

Policy for unsupported hardware

Various z/OS elements, such as DFSMSdfp, JES2, JES3, and MVS, contain code that supports specific hardware servers or devices. In some cases, this device-related element support remains in the product even after the hardware devices pass their announced End of Service date. z/OS may continue to service element code; however, it will not provide service related to unsupported hardware devices. Software problems related to these devices will not be accepted for service, and current service activity will cease if a problem is determined to be associated with out-of-support devices. In such cases, fixes will not be issued.

Minimum supported hardware

The minimum supported hardware for z/OS releases identified in z/OS announcements can subsequently change when service for particular servers or devices is withdrawn. Likewise, the levels of other software products supported on a particular release of z/OS are subject to the service support lifecycle of those

products. Therefore, z/OS and its product publications (for example, panels, samples, messages, and product documentation) can include references to hardware and software that is no longer supported.

- For information about software support lifecycle, see: IBM Lifecycle Support for z/OS (www.ibm.com/software/support/systemsz/lifecycle)
- For information about currently-supported IBM hardware, contact your IBM representative.

Programming interface information

This publication is intended to help you understand the output from SMP/E processing.

This publication also documents intended Programming Interfaces that allow the customer to write programs to obtain the services of IBM SMP/E for z/OS, V3R6. This Programming Interface information is identified where it occurs by an introductory statement to a chapter or section.

This publication also documents information that is **not** intended to be used as a Programming Interface of IBM SMP/E for z/OS, V3R6. This information is identified where it occurs by an introductory statement to a chapter or section.

Trademarks

IBM, the IBM logo, and ibm.com are trademarks or registered trademarks of International Business Machines Corp., registered in many jurisdictions worldwide. Other product and service names might be trademarks of IBM or other companies. A current list of IBM trademarks is available on the Web at Copyright and Trademark information (www.ibm.com/legal/copytrade.shtml).

Special attributions

Microsoft, Windows, Windows NT, and the Windows logo are trademarks of Microsoft Corporation in the United States, other countries, or both.

UNIX is a registered trademark of The Open Group in the United States and other countries.

Java and all Java-based trademarks and logos are trademarks or registered trademarks of Oracle and/or its affiliates.

Index

BYPASS
SYSMOD subentry
distribution zone 265
target zone 265

C

C
data structures in 339
sample GIMAPI program 345
CALL
effect of CALLLIBS subentry on 290
LMOD subentry 207
CALLLIBS
LMOD subentry 205
specifying on JCLIN 49
CASE
LMOD subentry 208
catalog
ARCHDEF attribute 398
CATALOG
DDDEF entry 163
CATEGORY
++HOLD MCS operand 31
CBPDO
Recommended service upgrade 457
ccc value of FMID 458
CHANGEFILE
OPTIONS subentry 242
CHANGEFILE(YES) subentry
required for library change file records 359
CIFREQ
SYSMOD subentry
distribution zone 265
target zone 265
CLASS
++HOLD MCS operand 32
naming conventions 456
values
ERREL 32, 456
HIPER 32, 456
PE 32, 457
SECINT 32
UCLIN 32, 457
YR2000 32, 457
CLIENT, defining 113
CLIST data element
installing in variable block library 8
COMMENT
++HOLD MCS operand 32
COMP
OPTIONS subentry 242
UTILITY entry 289
COMPACT
OPTIONS subentry 242
COMPACT option
of GIMCPTS 385
compacting SMPPTS data set
with GIMCPTS 385
COMPAT
LMOD subentry 208
MOD entry 229
COMPAT option
of GIMSMP 383

compress utility
default values 242, 289, 290
OPTIONS entry 242
UTILITY entry for 242, 289
CONCAT
DDDEF entry 164
concatenating data sets
DDDEF entry 164
not allowed for SMPPTS spill data sets 128
SYSLIB data sets
++MAC 56
++MACUPD 59
order 136
SMPMTS 125
use 136
conditional requisites
defining 37, 45
contact
z/OS 461
CONTENT
ALL value
ORDER subentry 250
ORDER subentry 250
COPY
++JCLIN MCS operand 49
INVOKE control statement operand 445
LMOD subentry 206
OPTIONS subentry 242
UTILITY entry for 289
COPY control statement
for GIMIAP 446
copy utility
default values 242, 289, 290
OPTIONS entry 242
specifying on JCLIN 49
UTILITY entry for 242, 289
corequisite SYSMODs
++IF MCS 38
++JAR MCS 44
defining 94
CSECT
++MOD MCS operand 61
MOD entry 227
CSI
application programming interface for 299
ASSEM entry
distribution zone 152
target zone 152
data element entry
distribution zone 158
target zone 158
DDDEF entry
distribution zone 163
global zone 163
target zone 163
defining 116
DLIB entry
distribution zone 174
target zone 174
DLIBZONE entry 177
FEATURE entry 181
FMIDSET entry 182
GIMAPI for 299
GLOBALZONE entry 184

Y

YR2000 class value 32, 457
YR2000 SOURCEID
    naming conventions 458

Z

z/OS SecureWay Security Server protection for data sets 166
z/OS SMP/E Reference
    content, changed xxv, xxvi
    content, new xxv
    summary of changes xxv
ZAP
    MCS statement 97
    OPTIONS subentry 248
    UTILITY entry for 289
ZONE
    ZONESET entry 296
ZONEDESCRIPTION
    DLIBZONE subentry 179
    GLOBALZONE subentry 185
    TARGETZONE subentry 288
ZONEINDEX
    GLOBALZONE subentry 186
zones
    defining 139
    naming conventions 177, 286
ZONES
    ORDER subentry 251
ZONESET entry
    listing 295
    subentries
        XZREQCHK 296
        ZONE 296
    summary 295
    UCLIN for 295

IBM®

Product Number: 5650-ZOS